

```
diff --git a/.dockerignore b/.dockerignore
deleted file mode 100644
index 6c85e6f6..00000000
--- a/.dockerignore
+++ /dev/null
@@ -1,9 +0,0 @@
-.ci
-.github
-.gitignore
-.golangci.yml
-.idea
-.vscode
-
-LICENSE
-*.md
diff --git a/.github/workflows/pull_ci.yml b/.github/workflows/pull_ci.yml
deleted file mode 100644
index 1ddb8471..00000000
--- a/.github/workflows/pull_ci.yml
+++ /dev/null
@@ -1,52 +0,0 @@
-name: PR CI
-on: [pull_request]
-
-jobs:
-  lint:
-    name: Lint
-    runs-on: ubuntu-18.04
-    steps:
-      - uses: actions/checkout@v2
-      - name: golangci-lint
-        uses: golangci/golangci-lint-action@v2
-        with:
-          version: latest
-          working-directory: .
-          args: --timeout 3m
-  test:
-    name: Golang Unit Tests v${{ matrix.go }} (${{ matrix.os }})
-    runs-on: ${{ matrix.os }}
-    strategy:
-      matrix:
-        go: ['1.16']
-        os: [ubuntu-20.04]
-    steps:
-      - uses: actions/checkout@v2
-      - uses: actions/setup-go@v1
-        with:
-          go-version: ${{ matrix.go }}
-      - run: go mod download
-        shell: bash
-      - run: ./scripts/build.sh evm
-        shell: bash
-      - run: ./scripts/build_test.sh
-        shell: bash
-  e2e:
-    name: Golang E2E Tests v${{ matrix.go }} (${{ matrix.os }})
-    runs-on: ${{ matrix.os }}
-    strategy:
-      matrix:
-        go: ['1.16']
-        os: [ubuntu-20.04]
-    steps:
-      - uses: actions/checkout@v2
-      - uses: actions/setup-go@v1
-        with:
-          go-version: ${{ matrix.go }}
-      - run: .github/workflows/run_e2e_tests.sh --parallelism 1 --client-id $KURTOSIS_CLIENT_ID --client-secret $KURTOSIS_CLIENT_SECRET
-        shell: bash
-      - env:
-          DOCKER_USERNAME: ${ secrets.DOCKER_USERNAME }
-          DOCKER_PASS: ${ secrets.DOCKER_PASS }
-          KURTOSIS_CLIENT_ID: ${ secrets.KURTOSIS_CLIENT_ID }
-          KURTOSIS_CLIENT_SECRET: ${ secrets.KURTOSIS_CLIENT_SECRET }
diff --git a/.github/workflows/push_ci.yml b/.github/workflows/push_ci.yml
deleted file mode 100644
index 1ac73f18..00000000
--- a/.github/workflows/push_ci.yml
+++ /dev/null
@@ -1,52 +0,0 @@
-name: Branch Push CI
-on: [push]
-
-jobs:
-  lint:
-    name: Lint
-    runs-on: ubuntu-18.04
-    steps:
-      - uses: actions/checkout@v2
-      - name: golangci-lint
-        uses: golangci/golangci-lint-action@v2
-        with:
-          version: latest
-          working-directory: .
-          args: --timeout 3m
-  test:
-    name: Golang Unit Tests v${{ matrix.go }} (${{ matrix.os }})
-    runs-on: ${{ matrix.os }}
-    strategy:
-      matrix:
-        go: ['1.16']
-        os: [macos-11.0, ubuntu-18.04, ubuntu-20.04, windows-latest]
-    steps:
-      - uses: actions/checkout@v2
-      - uses: actions/setup-go@v1
-        with:
-          go-version: ${{ matrix.go }}
-      - run: go mod download
-        shell: bash
-      - run: ./scripts/build.sh evm
-        shell: bash
-      - run: ./scripts/build_test.sh
-        shell: bash
-  e2e:
-    name: Golang E2E Tests v${{ matrix.go }} (${{ matrix.os }})
-    runs-on: ${{ matrix.os }}
-    strategy:
-      matrix:
-        go: ['1.16']
-        os: [ubuntu-20.04]
-    steps:
-      - uses: actions/checkout@v2
-      - uses: actions/setup-go@v1
-        with:
-          go-version: ${{ matrix.go }}
-      - run: .github/workflows/run_e2e_tests.sh --parallelism 1 --client-id $KURTOSIS_CLIENT_ID --client-secret $KURTOSIS_CLIENT_SECRET
-        shell: bash
-      - env:
-          DOCKER_USERNAME: ${ secrets.DOCKER_USERNAME }
-          DOCKER_PASS: ${ secrets.DOCKER_PASS }
-          KURTOSIS_CLIENT_ID: ${ secrets.KURTOSIS_CLIENT_ID }
-          KURTOSIS_CLIENT_SECRET: ${ secrets.KURTOSIS_CLIENT_SECRET }
diff --git a/.github/workflows/run_e2e_tests.sh b/.github/workflows/run_e2e_tests.sh
deleted file mode 100755
```

```
index 04f297fc..00000000
--- a/.github/workflows/run_e2e_tests.sh
+++ /dev/null
@@ -1,86 +0,0 @@
-set -o errexit
-set -o nounset
-set -o pipefail
-
-# If Docker Credentials are not available fail
-if [[ -z ${DOCKER_USERNAME} ]]; then
-    echo "Skipping Tests because Docker Credentials were not present."
-    exit 1
-fi
-
-# Testing specific variables
-avalanche_testing_repo="avaplatform/avalanche-testing"
-avalanchego_repo="avaplatform/avalanchego"
-# Define default avalanche testing version to use
-avalanche_testing_image="${avalanche_testing_repo}:master"
-
-# Avalanche root directory
-CORETH_PATH=$( cd "$( dirname "${BASH_SOURCE[0]}" )" && pwd )
-
-# Load the versions
-source "$CORETH_PATH"/scripts/versions.sh
-
-# Load the constants
-source "$CORETH_PATH"/scripts/constants.sh
-
-# Login to docker
-echo "$DOCKER_PASS" | docker login --username "$DOCKER_USERNAME" --password-stdin
-
-# Checks available docker tags exist
-function docker_tag_exists() {
-    TOKEN=$(curl -s -H "Content-Type: application/json" -X POST -d '{"username": "'${DOCKER_USERNAME}'", "password": "'${DOCKER_PASS}'"}' https://hub.docker.com/v2/users/login/ | jq -r .token)
-    curl --silent -H "Authorization: JWT ${TOKEN}" -f --head -L https://hub.docker.com/v2/repositories/$1/tags/$2/ > /dev/null
-}
-
-# Defines the avalanche-testing tag to use
-# Either uses the same tag as the current branch or uses the default
-if docker_tag_exists $avalanche_testing_repo $current_branch; then
-    echo "$avalanche_testing_repo:$current_branch exists; using this image to run e2e tests"
-    avalanche_testing_image="$avalanche_testing_repo:$current_branch"
-else
-    echo "$avalanche_testing_repo $current_branch does NOT exist; using the default image to run e2e tests"
-fi
-
-echo "Using $avalanche_testing_image for e2e tests"
-
-# Defines the avalanchego tag to use
-# Either uses the same tag as the current branch or uses the default
-# Disable matchup in favor of explicit tag
-# TODO re-enable matchup when our workflow better supports it.
-# if docker_tag_exists $avalanchego_repo $current_branch; then
-#     echo "$avalanchego_repo:$current_branch exists; using this avalanchego image to run e2e tests"
-#     AVALANCHE_VERSION=$current_branch
-# else
-#     echo "$avalanchego_repo $current_branch does NOT exist; using the default image to run e2e tests"
-# fi
-
-# pulling the avalanche-testing image
-docker pull $avalanche_testing_image
-
-# Setting the build ID
-git_commit_id=$( git rev-list -1 HEAD )
-
-# Build current avalanchego
-source "$CORETH_PATH"/scripts/build_image.sh
-
-# Target built version to use in avalanche-testing
-avalanche_image="avaplatform/avalanchego:$build_image_id"
-
-echo "Running Avalanche Image: ${avalanche_image}"
-echo "Running Avalanche Testing Image: ${avalanche_testing_image}"
-echo "Git Commit ID : ${git_commit_id}"
-
-# >>>>>> avalanche-testing custom parameters <<<<<<<<<<<<
-custom_params_json="{
-    \"isKurtosisCoreDevMode\": false,
-    \"avalanchegoImage\": \"${avalanche_image}\",
-    \"testBatch\": \"avalanchego\"
-}"
-# >>>>>> avalanche-testing custom parameters <<<<<<<<<<<<
-
-bash "$CORETH_PATH/.kurtosis/kurtosis.sh" \
-    --tests "C-Chain Bombard WorkFlow" \
-    --custom-params "${custom_params_json}" \
-    "${avalanche_testing_image}" \
-    $@
diff --git a/.gitignore b/.gitignore
index 87ff040a..84c048a7 100644
--- a/.gitignore
+++ b/.gitignore
@@ -1,46 +1 @@
-./main
-
-*.log
-*.~
-.DS_Store
-
-awscpu
-
-# Binaries for programs and plugins
-*.exe
-*.exe~
-*.dll
-*.so
-*.dylib
-*.profile
-
-# Test binary, build with `go test -c`
-*.test
-
-# Output of the go coverage tool, specifically when used with LiteIDE
-*.out
-
-# ignore GoLand metafiles directory
-.idea/
-
-*.logs/
-
-.vscode*
-
-*.pb*
-
-*.cpu[0-9]*
-*.mem[0-9]*
-*.lock[0-9]*
-*.profile
-*.swp
-*.aux
-*.fdb*
-*.fls
```

```

*.pdf
-coverage
-
-bin/
-build/
+/build/
diff --git a/.golangci.yml b/.golangci.yml
deleted file mode 100644
index d2a5e38c..00000000
--- a/.golangci.yml
+++ /dev/null
@@ -1,48 +0,0 @@
-# This file configures github.com/golangci/golangci-lint.
-
-run:
-  timeout: 3m
-  tests: true
-  # default is true. Enables skipping of directories:
-  #   vendor$, third_party$, testdata$, examples$, Godeps$, builtin$
-  skip-dirs-use-default: true
-  skip-files:
-    - core/genesis_alloc.go
-
-linters:
-  disable-all: true
-  enable:
-    - deadcode
-    - goconst
-    - goimports
-    - gosimple
-    - govet
-    - ineffassign
-    - misspell
-    - unconvert
-    - varcheck
-
-linters-settings:
-  gofmt:
-    simplify: true
-  goconst:
-    min-len: 3 # minimum length of string constant
-    min-occurrences: 6 # minimum number of occurrences
-
-issues:
-  exclude-rules:
-    - path: crypto/blake2b/
-      linters:
-        - deadcode
-    - path: crypto/bn256/cloudflare
-      linters:
-        - deadcode
-    - path: p2p/discv5/
-      linters:
-        - deadcode
-    - path: core/vm/instructions_test.go
-      linters:
-        - goconst
-    - path: cmd/faucet/
-      linters:
-        - deadcode
diff --git a/kurtosis/kurtosis.sh b/kurtosis/kurtosis.sh
deleted file mode 100755
index a3d1cd85..00000000
--- a/kurtosis/kurtosis.sh
+++ /dev/null
@@ -1,226 +0,0 @@
-# !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!! WARNING !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
-#
-#       Do not modify this file! It will get overwritten when you upgrade Kurtosis!
-#
-# !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!! WARNING !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
-#
-set -euo pipefail
-
-# =====
-#                               Constants
-# =====
-# The directory where Kurtosis will store files it uses in between executions, e.g. access tokens
-# Can make this configurable if needed
-KURTOSIS_DIRPATH="${HOME}/.kurtosis"
-
-KURTOSIS_CORE_TAG="1.8"
-KURTOSIS_DOCKERHUB_ORG="kurtosistech"
-INITIALIZER_IMAGE="${KURTOSIS_DOCKERHUB_ORG}/kurtosis-core_initializer:${KURTOSIS_CORE_TAG}"
-API_IMAGE="${KURTOSIS_DOCKERHUB_ORG}/kurtosis-core_api:${KURTOSIS_CORE_TAG}"
-
-POSITIONAL_ARG_DEFINITION_FRAGMENTS=2
-
-# =====
-#                               Arg Parsing
-# =====
-function print_help_and_exit() {
-  echo ""
-  echo "${basename "$0"}" [--custom-params custom_params_json] [--client-id client_id] [--client-secret client_secret] [--help] [--kurtosis-log-level kurtosis_log_level] [--list] [--parallelism parallelism]
-  echo ""
-  echo "  --custom-params custom_params_json           JSON string containing arbitrary data that will be passed as-is to your testsuite, so it can modify its behaviour based on input (default: {})"
-  echo "  --client-id client_id                         An OAuth client ID which is needed for running Kurtosis in CI, and should be left empty when running Kurtosis on a local machine"
-  echo "  --client-secret client_secret                 An OAuth client secret which is needed for running Kurtosis in CI, and should be left empty when running Kurtosis on a local machine"
-  echo "  --help                                         Display this message"
-  echo "  --kurtosis-log-level kurtosis_log_level       The log level that all output generated by the Kurtosis framework itself should log at (panic|fatal|error|warning|info|debug|trace) (default: info)"
-  echo "  --list                                         Rather than running the tests, lists the tests available to run"
-  echo "  --parallelism parallelism                     The number of texts to execute in parallel (default: 4)"
-  echo "  --tests test_names                           List of test names to run, separated by ',' (default or empty: run all tests)"
-  echo "  --test-suite-log-level test_suite_log_level   A string that will be passed as-is to the test suite container to indicate what log level the test suite container should output at; this string must be one of the following: panic, fatal, error, warning, info, debug, trace"
-  echo "  test_suite_image                             The Docker image containing the testsuite to execute"
-
-  echo ""
-  exit 1 # Exit with an error code, so that if it gets accidentally called in parent scripts/CI it fails loudly
-}
-
-# =====
-#                               Arg Parsing
-# =====
-client_id=""
-client_secret=""
-custom_params_json="{}"
-do_list="false"
-kurtosis_log_level="info"
-parallelism="4"
-show_help="false"
-test_names=""
-test_suite_image=""
-test_suite_log_level="info"
```

```
-POSITIONAL=()
-while [ ${#} -gt 0 ]; do
-    key="${1}"
-    case "${key}" in
-
-        --custom-params)
-
-            custom_params_json="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --client-id)
-
-            client_id="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --client-secret)
-
-            client_secret="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --help)
-            show_help="true"
-            shift    # Shift to clear out the flag
-
-            ;;
-
-        --kurtosis-log-level)
-
-            kurtosis_log_level="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --list)
-            do_list="true"
-            shift    # Shift to clear out the flag
-
-            ;;
-
-        --parallelism)
-
-            parallelism="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --tests)
-
-            test_names="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        --test-suite-log-level)
-
-            test_suite_log_level="${2}"
-            shift    # Shift to clear out the flag
-            shift    # Shift again to clear out the value
-            ;;
-
-        *)
-            echo "ERROR: Unrecognized flag '${key}'" >&2
-            exit 1
-            ;;
-
-        *)
-            POSITIONAL+=("${1}")
-            shift
-            ;;
-    esac
-done
-
-# Restore positional parameters and assign them to variables
-set -- "${POSITIONAL[@]}"
-test_suite_image="${1:-}"
-
-
-
-
-# =====
-#                               Arg Validation
-# =====
-
-# [ "${#}" -ne 1 ]; then
-#     echo "ERROR: Expected 1 positional variables but got ${#}" >&2
-#     print_help_and_exit
-# fi
-
-# [ -z "$test_suite_image" ]; then
-#     echo "ERROR: Variable 'test_suite_image' cannot be empty" >&2
-#     exit 1
-# fi
-
-
-
-# =====
-#                               Main Logic
-# =====
-# =====# Because Kurtosis X.Y.Z tags are normalized to X.Y so that minor patch updates are transparently
-# used, we need to pull the latest API & initializer images
-# "Pulling latest versions of API & initializer image..."
-# echo "pulling latest versions of API & initializer image..."
-# if ! docker pull "${INITIALIZER_IMAGE}"; then
-#     echo "WARN: An error occurred pulling the latest version of the initializer image (${INITIALIZER_IMAGE}); you may be running an out-of-date version" >&2
-# else
-#     echo "Successfully pulled latest version of initializer image"
-# fi
-# if ! docker pull "${API_IMAGE}"; then
-#     echo "WARN: An error occurred pulling the latest version of the API image (${API_IMAGE}); you may be running an out-of-date version" >&2
-# else
-#     echo "Successfully pulled latest version of API image"
-# fi
-
-# Kurtosis needs a Docker volume to store its execution data in
-# To learn more about volumes, see: https://docs.docker.com/storage/volumes/
-sanitized_image=$(echo "${test_suite_image}" | sed 's/[^a-zA-Z0-9_-]/_/g')
-suite_execution_volume=$(date +%Y-%m-%dT%H.%M.%S)_${sanitized_image}
-if ! docker volume create "${suite_execution_volume}" > /dev/null; then
-    echo "ERROR: Failed to create a Docker volume to store the execution files in" >&2
-    exit 1
-fi
-
-# if ! mkdir -p "${KURTOSIS_DIRPATH}"; then
-#     echo "ERROR: Failed to create the Kurtosis directory at '${KURTOSIS_DIRPATH}'" >&2
-#     exit 1
-# fi
```

```
-docker run \
- '# The Kurtosis initializer runs inside a Docker container, but needs to access to the Docker engine; this is how to do it` \
- '# For more info, see the bottom of: http://jpetazzo.github.io/2015/09/03/do-not-use-docker-in-docker-for-ci/` \
- --mount "type=bind,source=/var/run/docker.sock,target=/var/run/docker.sock" \
- \
- '# Because the Kurtosis initializer runs inside Docker but needs to persist & read files on the host filesystem between execution,` \
- '# the container expects the Kurtosis directory to be bind-mounted at the special "/kurtosis" path` \
- --mount "type=bind,source=${KURTOSIS_DIRPATH},target=/kurtosis" \
- \
- '# The Kurtosis initializer image requires the volume for storing suite execution data to be mounted at the special "/suite-execution" path` \
- --mount "type=volume,source=${suite_execution_volume},target=/suite-execution" \
- \
- \
- '# Keep these sorted alphabetically` \
- --env CLIENT_ID="${client_id}" \
- --env CLIENT_SECRET="${client_secret}" \
- --env CUSTOM_PARAMS_JSON="${custom_params_json}" \
- --env DO_LIST="${do_list}" \
- --env KURTOSIS_API_IMAGE="${API_IMAGE}" \
- --env KURTOSIS_LOG_LEVEL="${kurtosis_log_level}" \
- --env PARALLELISM="${parallelism}" \
- --env SUITE_EXECUTION_VOLUME="${suite_execution_volume}" \
- --env TEST_NAMES="${test_names}" \
- --env TEST_SUITE_IMAGE="${test_suite_image}" \
- --env TEST_SUITE_LOG_LEVEL="${test_suite_log_level}" \
- \
- "${INITIALIZER_IMAGE}"
diff --git a/CHANGES.pdf b/CHANGES.pdf
new file mode 100644
index 00000000..7864125d
Binary files /dev/null and b/CHANGES.pdf differ
diff --git a/Dockerfile b/Dockerfile
deleted file mode 100644
index ec7b1720..00000000
--- a/Dockerfile
+++ /dev/null
@@ -1,30 +0,0 @@
-# syntax=docker/dockerfile:experimental
-
-# ===== Setting up base Stage =====
-# Set required AVALANCHE_VERSION parameter in build image script
-ARG AVALANCHE_VERSION
-
-# ===== Compilation Stage =====
-# golang:1.17.1-buster AS builder
-RUN apt-get update && apt-get install -y --no-install-recommends bash=5.0-4 git=1:2.20.1-2+deb10u3 make=4.2.1-1.2 gcc=4:8.3.0-1 musl-dev=1.1.21-2 ca-certificates=20200601-deb10u2 linux-headers-amd64
-
-WORKDIR /build
-# Copy and download avalanche dependencies using go mod
-COPY go.mod .
-COPY go.sum .
-RUN go mod download
-
-# Copy the code into the container
-COPY . .
-
-# Pass in CORETH_COMMIT as an arg to allow the build script to set this externally
-ARG CORETH_COMMIT
-ARG CURRENT_BRANCH
-
-RUN export CORETH_COMMIT=${CORETH_COMMIT} && export CURRENT_BRANCH=${CURRENT_BRANCH} && ./scripts/build.sh /build/evm
-
-# ===== Cleanup Stage =====
-FROM avaplatform/avalanchego:${AVALANCHE_VERSION} AS builtImage
-
-# Copy the evm binary into the correct location in the container
-COPY --from=builder /build/evm /avalanchego/build/plugins/evm
diff --git a/accounts/abi/bind/auth.go b/accounts/abi/bind/auth.go
index 859083c0..b50ac70d 100644
--- a/accounts/abi/bind/auth.go
+++ b/accounts/abi/bind/auth.go
@@ -34,13 +34,23 @@ import (
    "io/ioutil"
    "math/big"

    "github.com/ava-labs/coreth/accounts"
    "github.com/ava-labs/coreth/accounts/external"
    "github.com/ava-labs/coreth/accounts/keystore"
    "github.com/ava-labs/coreth/core/types"
+<<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+
+// =====
+>>>>>>> upstream-v0.8.5-rc.2
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/accounts/external"
+    "github.com/flare-foundation/coreth/accounts/keystore"
+    "github.com/flare-foundation/coreth/core/types"
+<<<<<<< HEAD
+// =====
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+>>>>>>> upstream-v0.8.5-rc.2
+    )
+
+// ErrNoChainID is returned whenever the user failed to specify a chain id.
diff --git a/accounts/abi/bind/backend.go b/accounts/abi/bind/backend.go
index 29a4b3cb..179a043a 100644
--- a/accounts/abi/bind/backend.go
+++ b/accounts/abi/bind/backend.go
@@ -31,9 +31,16 @@ import (
    "errors"
    "math/big"

    "github.com/ava-labs/coreth/core/types"
    "github.com/ava-labs/coreth/interfaces"
+<<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+// =====
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    "github.com/ethereum/go-ethereum/common"
+>>>>>>> upstream-v0.8.5-rc.2
+    )
+
+var (
diff --git a/accounts/abi/bind/backends/simulated.go b/accounts/abi/bind/backends/simulated.go
index 4860cb34..e4d214e1 100644
--- a/accounts/abi/bind/backends/simulated.go
+++ b/accounts/abi/bind/backends/simulated.go
@@ -34,27 +34,42 @@ import (
    "sync"
    "time"

    "github.com/ava-labs/coreth/eth"
    "github.com/ava-labs/coreth/accounts/abi"
    "github.com/ava-labs/coreth/accounts/abi/bind"
    "github.com/ava-labs/coreth/consensus/dummy"
```

```
- "github.com/ava-labs/coreth/core/bloombits"
- "github.com/ava-labs/coreth/core/rawdb"
- "github.com/ava-labs/coreth/core/state"
- "github.com/ava-labs/coreth/core/types"
- "github.com/ava-labs/coreth/core/vm"
- "github.com/ava-labs/coreth/eth/filters"
- "github.com/ava-labs/coreth/ethdb"
- "github.com/ava-labs/coreth/interfaces"
- "github.com/ava-labs/coreth/params"
- "github.com/ava-labs/coreth/rpc"
+<<<<<< HEAD
+ "github.com/ethereum/go-ethereum/common"
+ "github.com/ethereum/go-ethereum/common/hexutil"
+ "github.com/ethereum/go-ethereum/common/math"
+ "github.com/ethereum/go-ethereum/event"
+ "github.com/ethereum/go-ethereum/log"
+=====
+ "github.com/flare-foundation/coreth/eth"
+>>>>>> upstream-v0.8.5-rc.2
+
+ "github.com/flare-foundation/coreth/accounts/abi"
+ "github.com/flare-foundation/coreth/accounts/abi/bind"
+ "github.com/flare-foundation/coreth/consensus/dummy"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/flare-foundation/coreth/core/bloombits"
+ "github.com/flare-foundation/coreth/core/rawdb"
+ "github.com/flare-foundation/coreth/core/state"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+<<<<<< HEAD
+ "github.com/flare-foundation/coreth/eth"
+=====
+>>>>>> upstream-v0.8.5-rc.2
+ "github.com/flare-foundation/coreth/eth/filters"
+ "github.com/flare-foundation/coreth/ethdb"
+ "github.com/flare-foundation/coreth/interfaces"
+ "github.com/flare-foundation/coreth/params"
+ "github.com/flare-foundation/coreth/rpc"
+<<<<<< HEAD
+=====
+ "github.com/ethereum/go-ethereum/common"
+ "github.com/ethereum/go-ethereum/common/hexutil"
+ "github.com/ethereum/go-ethereum/common/math"
+ "github.com/ethereum/go-ethereum/event"
+ "github.com/ethereum/go-ethereum/log"
+>>>>>> upstream-v0.8.5-rc.2
)

// Verify that SimulatedBackend implements required interfaces
@@ -107,7 +122,11 @@ type SimulatedBackend struct {
func NewSimulatedBackendWithDatabase(database ethdb.Database, alloc core.GenesisAlloc, gasLimit uint64) *SimulatedBackend {
    cpcf := params.TestChainConfig
    cpcf.ChainID = big.NewInt(1337)
+<<<<<< HEAD
+ genesis := core.Genesis{Config: cpcf, GasLimit: gasLimit, Alloc: alloc, Coinbase: common.HexToAddress("0x0100000000000000000000000000000000")}
+=====
+ genesis := core.Genesis{Config: cpcf, GasLimit: gasLimit, Alloc: alloc}
+>>>>>> upstream-v0.8.5-rc.2
+ genesis.MustCommit(database)
+ cacheConfig := &core.CacheConfig{
+     blockchain, _ := core.NewBlockchain(database, cacheConfig, genesis.Config, dummy.NewFaker(), vm.Config{}, common.Hash{})
@@ -499,6 +518,12 @@ func (b *SimulatedBackend) AcceptedNonceAt(ctx context.Context, account common.A
// SuggestGasPrice implements ContractTransactor.SuggestGasPrice. Since the simulated
// chain doesn't have miners, we just return a gas price of 1 for any call.
func (b *SimulatedBackend) SuggestGasPrice(ctx context.Context) (*big.Int, error) {
+<<<<<< HEAD
+=====
+ b.mu.Lock()
+ defer b.mu.Unlock()
+
+>>>>>> upstream-v0.8.5-rc.2
+ if b.acceptedBlock.Header().BaseFee != nil {
+     return b.acceptedBlock.Header().BaseFee, nil
+ }
diff --git a/accounts/abi/bind/backends/simulated_test.go b/accounts/abi/bind/backends/simulated_test.go
index d979e791..8f3ff124 100644
--- a/accounts/abi/bind/backends/simulated_test.go
+++ b/accounts/abi/bind/backends/simulated_test.go
@@ -37,14 +37,14 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/accounts/abi"
-    "github.com/ava-labs/coreth/accounts/abi/bind"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/accounts/abi"
+    "github.com/flare-foundation/coreth/accounts/abi/bind"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    "github.com/flare-foundation/coreth/params"
)

func TestSimulatedBackend(t *testing.T) {
@@ -512,7 +512,7 @@ func TestEstimateGas(t *testing.T) {
    GasPrice: big.NewInt(0),
    Value:     nil,
    Data:      common.Hex2Bytes("b9b046f9"),
-    }, 0, errors.New("invalid opcode: opcode 0xfe not defined"), nil},
+    }, 0, errors.New("invalid opcode: INVALID"), nil},
)

{"Valid", interfaces.CallMsg{
    From: addr,
diff --git a/accounts/abi/bind/base.go b/accounts/abi/bind/base.go
index 58ac1ea0..3428bb04 100644
--- a/accounts/abi/bind/base.go
+++ b/accounts/abi/bind/base.go
@@ -34,12 +34,12 @@ import (
    "strings"
    "sync"

-    "github.com/ava-labs/coreth/accounts/abi"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
+<<<<<< HEAD
+ "github.com/ethereum/go-ethereum/common"
+ "github.com/ethereum/go-ethereum/crypto"
+ "github.com/ethereum/go-ethereum/event"
+
+ "github.com/flare-foundation/coreth/accounts/abi"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/interfaces"
+=====
+ "github.com/flare-foundation/coreth/accounts/abi"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+ "github.com/flare-foundation/coreth/interfaces"
+ "github.com/ethereum/go-ethereum/common"
```

```

+       "github.com/ethereum/go-ethereum/crypto"
+       "github.com/ethereum/go-ethereum/event"
+)
+
+var (
+    ErrNilAssetAmount      = errors.New("cannot specify nil asset amount for native asset call")
+    errNativeAssetDeployContract = errors.New("cannot specify native asset params while deploying a contract")
+>>>>>> upstream-v0.8.5-rc.2
+)
+
+// SignerFn is a signer function callback when a contract requires a method to
@@ -54,6 +70,15 @@ type CallOpts struct {
+    Context      context.Context // Network context to support cancellation and timeouts (nil = no timeout)
+}
+
+<<<<<< HEAD
+=====
+// NativeAssetCallOpts contains params for native asset call
+type NativeAssetCallOpts struct {
+    + AssetID      common.Hash // Asset ID
+    + AssetAmount *big.Int   // Asset amount
+}
+
+>>>>>> upstream-v0.8.5-rc.2
+// TransactOpts is the collection of authorization data required to create a
+// valid Ethereum transaction.
+type TransactOpts struct {
@@ -70,6 +95,17 @@ type TransactOpts struct {
+    Context context.Context // Network context to support cancellation and timeouts (nil = no timeout)
+
+    NoSend bool // Do all transact steps but do not send the transaction
+
+<<<<<<< HEAD
+=====
+
+    + // If set, the transaction is transformed to perform the requested call through the native asset
+    + // precompile. This will update the to address of the transaction to that of the native asset precompile
+    + // and pack the requested [to] address, [assetID], [assetAmount], and [input] data for the transaction
+    + // into the call data of the transaction. When executed within the EVM, the precompile will parse the input
+    + // data and attempt to atomically transfer [assetAmount] of [assetID] to the [to] address and invoke the
+    + // contract at [to] if present, passing in the original [input] data.
+    + NativeAssetCall *NativeAssetCallOpts
+>>>>>> upstream-v0.8.5-rc.2
+}
+
+// FilterOpts is the collection of options to fine tune filtering for events
@@ -240,6 +276,42 @@ func (c *BoundContract) Transfer(opts *TransactOpts) (*types.Transaction, error)
+    // or not, reject invalid transaction at the first place
+    return c.transact(opts, &c.address, nil)
+}
+
+<<<<<<< HEAD
+=====
+
+// wrapNativeAssetCall preprocesses the arguments to transform the requested call to go through the
+// native asset call precompile if it is specified on [opts].
+func wrapNativeAssetCall(opts *TransactOpts, contract *common.Address, input []byte) (*common.Address, []byte, error) {
+    + if opts.NativeAssetCall != nil {
+    +     // Prevent the user from sending a non-zero value through native asset call precompile as this will
+    +     // transfer the funds to the precompile address and essentially burn the funds.
+    +     if opts.Value != nil && opts.Value.Cmp(common.Big0) != 0 {
+    +         return nil, nil, fmt.Errorf("value must be 0 when performing native asset call, found %d", opts.Value)
+    +     }
+    +     if opts.NativeAssetCall.AssetAmount == nil {
+    +         return nil, nil, ErrNilAssetAmount
+    +     }
+    +     if opts.NativeAssetCall.AssetAmount.Cmp(common.Big0) < 0 {
+    +         return nil, nil, fmt.Errorf("asset value cannot be < 0 when performing native asset call, found %d", opts.NativeAssetCall.AssetAmount)
+    +     }
+    +     // Prevent potential panic if [contract] is nil in the case that transact is called through DeployContract.
+    +     if contract == nil {
+    +         return nil, nil, errNativeAssetDeployContract
+    +     }
+    +     // wrap input with native asset call params
+    +     input = vm.PackNativeAssetCallInput(
+    +         *contract,
+    +         opts.NativeAssetCall.AssetID,
+    +         opts.NativeAssetCall.AssetAmount,
+    +         input,
+    +     )
+    +     // target addr is now precompile
+    +     contract = &vm.NativeAssetCallAddr
+    + }
+    return contract, input, nil
+}
+
+>>>>>>> upstream-v0.8.5-rc.2
+func (c *BoundContract) createDynamicTx(opts *TransactOpts, contract *common.Address, input []byte, head *types.Header) (*types.Transaction, error) {
+    // Normalize value
+    value := opts.Value
@@ -375,6 +447,14 @@ func (c *BoundContract) transact(opts *TransactOpts, contract *common.Address, i
+    rawTx *types.Transaction
+    err    error
+
+    )
+
+<<<<<<<< HEAD
+=====
+
+    + // Preprocess native asset call arguments if present
+    + contract, input, err = wrapNativeAssetCall(opts, contract, input)
+    + if err != nil {
+    +     return nil, err
+    + }
+>>>>>>> upstream-v0.8.5-rc.2
+    + if opts.GasPrice != nil {
+    +     rawTx, err = c.createLegacyTx(opts, contract, input)
+    + } else {
diff --git a/accounts/abi/bind/base_test.go b/accounts/abi/bind/base_test.go
index d7988516..6ca2eee8 100644
--- a/accounts/abi/bind/base_test.go
+++ b/accounts/abi/bind/base_test.go
@@ -28,20 +28,38 @@ package bind_test

import (
    "context"
+<<<<<<< HEAD
+=====
+
+    + "fmt"
+>>>>>>> upstream-v0.8.5-rc.2
+    + "math/big"
+    + "reflect"
+    + "strings"
+    + "testing"
+
+    - "github.com/ava-labs/coreth/accounts/abi"
+    - "github.com/ava-labs/coreth/accounts/abi/bind"
+    - "github.com/ava-labs/coreth/core/types"
+    - "github.com/ava-labs/coreth/interfaces"
+<<<<<<< HEAD
+    + "github.com/stretchr/testify/assert"
+
+=====
+
+    + "github.com/flare-foundation/coreth/accounts/abi"
+    + "github.com/flare-foundation/coreth/accounts/abi/bind"
+    + "github.com/flare-foundation/coreth/core/types"
+    + "github.com/flare-foundation/coreth/core/vm"
+    + "github.com/flare-foundation/coreth/interfaces"
+>>>>>>> upstream-v0.8.5-rc.2

```

```

    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/hexutil"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/rlp"
+<<<<<< HEAD
+
+    "github.com/flare-foundation/coreth/accounts/abi"
+    "github.com/flare-foundation/coreth/accounts/abi/bind"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+=====
+    "github.com/stretchr/testify/assert"
+>>>>>> upstream-v0.8.5-rc.2
+
+
+func mockSign(addr common.Address, tx *types.Transaction) (*types.Transaction, error) { return tx, nil }
@@ -277,6 +295,83 @@ func TestUnpackIndexedBytesTyLogIntoMap(t *testing.T) {
    unpackAndCheck(t, bc, expectedReceivedMap, mockLog)
}

+<<<<<< HEAD
+=====
+func TestTransactNativeAssetCallNilAssetAmount(t *testing.T) {
+    assert := assert.New(t)
+    mt := &mockTransactor{}
+    bc := bind.NewBoundContract(common.Address{}, abi.ABI{}, nil, mt, nil)
+    opts := &bind.TransactOpts{
+        Signer: mockSign,
+    }
+    // fails if asset amount is nil
+    opts.NativeAssetCall = &bind.NativeAssetCallOpts{
+        AssetID:    common.Hash{},
+        AssetAmount: nil,
+    }
+    _, err := bc.Transact(opts, "")
+    assert.ErrorIs(err, bind.ErrNilAssetAmount)
+}
+
+func TestTransactNativeAssetCallNonZeroValue(t *testing.T) {
+    assert := assert.New(t)
+    mt := &mockTransactor{}
+    bc := bind.NewBoundContract(common.Address{}, abi.ABI{}, nil, mt, nil)
+    opts := &bind.TransactOpts{
+        Signer: mockSign,
+    }
+    opts.NativeAssetCall = &bind.NativeAssetCallOpts{
+        AssetID:    common.Hash{},
+        AssetAmount: big.NewInt(11),
+    }
+    // fails if value > 0
+    opts.Value = big.NewInt(11)
+    _, err := bc.Transact(opts, "")
+    assert.Equal(err.Error(), fmt.Sprintf("value must be 0 when performing native asset call, found %v", opts.Value))
+    // fails if value < 0
+    opts.Value = big.NewInt(-11)
+    _, err = bc.Transact(opts, "")
+    assert.Equal(err.Error(), fmt.Sprintf("value must be 0 when performing native asset call, found %v", opts.Value))
+}
+
+func TestTransactNativeAssetCall(t *testing.T) {
+    assert := assert.New(t)
+    json := ` [{"type": "function", "name": "method", "inputs": [{"type": "uint256"}, {"type": "string"}] } ]`
+    parsed, err := abi.JSON(strings.NewReader(json))
+    assert.Nil(err)
+    mt := &mockTransactor{}
+    contractAddr := common.Address{11}
+    bc := bind.NewBoundContract(contractAddr, parsed, nil, mt, nil)
+    opts := &bind.TransactOpts{
+        Signer: mockSign,
+    }
+    // normal call tx
+    methodName := "method"
+    arg1 := big.NewInt(22)
+    arg2 := "33"
+    normalCallTx, err := bc.Transact(opts, methodName, arg1, arg2)
+    assert.Nil(err)
+    // native asset call tx
+    assetID := common.Hash{44}
+    assetAmount := big.NewInt(55)
+    opts.NativeAssetCall = &bind.NativeAssetCallOpts{
+        AssetID:    assetID,
+        AssetAmount: assetAmount,
+    }
+    nativeCallTx, err := bc.Transact(opts, methodName, arg1, arg2)
+    assert.Nil(err)
+    // verify transformations
+    assert.Equal(vm.NativeAssetCallAddr, *nativeCallTx.To())
+    unpackedAddr, unpackedAssetID, unpackedAssetAmount, unpackedData, err := vm.UnpackNativeAssetCallInput(nativeCallTx.Data())
+    assert.Nil(err)
+    assert.NotEmpty(unpackedData)
+    assert.Equal(unpackedData, normalCallTx.Data())
+    assert.Equal(unpackedAddr, contractAddr)
+    assert.Equal(unpackedAssetID, assetID)
+    assert.Equal(unpackedAssetAmount, assetAmount)
+}
+
+>>>>>> upstream-v0.8.5-rc.2
+func TestTransactGasFee(t *testing.T) {
+    assert := assert.New(t)
+
diff --git a/accounts/abi/bind/bind.go b/accounts/abi/bind/bind.go
index 22e1a8cb..0158d7e3 100644
--- a/accounts/abi/bind/bind.go
+++ b/accounts/abi/bind/bind.go
@@ -40,8 +40,14 @@ import (
    "text/template"
    "unicode"

-    "github.com/ava-labs/coreth/accounts/abi"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/log"
+
+    "github.com/flare-foundation/coreth/accounts/abi"
+=====
+    "github.com/flare-foundation/coreth/accounts/abi"
+    "github.com/ethereum/go-ethereum/log"
+>>>>>> upstream-v0.8.5-rc.2
+
+
+// Lang is a target programming language selector to generate bindings for.
@@ -98,6 +104,16 @@ func Bind(types []string, abis []string, bytecodes []string, fsigs []map[string]
    transactIdentifiers = make(map[string]bool)
    eventIdentifiers    = make(map[string]bool)
}

+<<<<<< HEAD
+=====
+
+    for _, input := range evmABI.Constructor.Inputs {
+        if hasStruct(input.Type) {
+            bindStructType[lang](input.Type, structs)
+        }
+    }
+
+>>>>>> upstream-v0.8.5-rc.2

```



```

        for _, original := range evmABI.Methods {
            // Normalize the method for capital cases and non-anonymous inputs/outputs
            normalized := original
        }
    }
    diff --git a/accounts/abi/bind/bind_test.go b/accounts/abi/bind/bind_test.go
    index 3cd9f3f1..95103425 100644
    --- a/accounts/abi/bind/bind_test.go
    +++ b/accounts/abi/bind/bind_test.go
    @@ -298,9 +298,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -353,9 +353,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -399,9 +399,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -456,10 +456,10 @@ var bindTests = []struct {
        ,
        "math/big"
        "reflect"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/ethereum/go-ethereum/common"
    -     "github.com/ava-labs/coreth/core"
    -     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -505,9 +505,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -571,9 +571,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    @@ -616,10 +616,10 @@ var bindTests = []struct {
        []string{"6060604052609f8060106000396000f3606060405260e060020a6000350463f97a60058114601a575b005b600060605260c0604052600d60809081527f4920646f6e277420657869737400000000000000000000000000000000"},
        []string{[{"constant":true,"inputs":[],"name":"String","outputs":[{"name":"","type":"string"}],"type":"function"}]},
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/ethereum/go-ethereum/common"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/core"
    ,
    }
    ,
    }
    // Create a simulator and wrap a non-deployed contract
    @@ -655,10 +655,10 @@ var bindTests = []struct {
        []string{"6080604052348015600f57600080fd5b5060888061001e6000396000f3fe6080604052348015600f57600080fd5b5060043610600285760003560e01c8063d5f6622514602d575b600080fd5b6033604c565b6040805192835"},
        []string{[{"inputs":[],"name":"Struct","outputs":[{"internalType":"uint256","name":"a","type":"uint256"}, {"internalType":"uint256","name":"b","type":"uint256"}], "stateMutability":"pure"},
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/ethereum/go-ethereum/common"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/core"
    ,
    }
    ,
    }
    // Create a simulator and wrap a non-deployed contract
    @@ -703,9 +703,9 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    -     "github.com/ava-labs/coreth/core"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/core"
    +     "github.com/ethereum/go-ethereum/crypto"
    ,
    }
    ,
    }
    @@ -752,10 +752,10 @@ var bindTests = []struct {
        ,
        "math/big"
    -
    -     "github.com/ava-labs/coreth/accounts/abi/bind"
    -     "github.com/ava-labs/coreth/accounts/abi/bind/backends"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    +     "github.com/flare-foundation/coreth/accounts/abi/bind"
    ,
    }
    ,
    }

```

```

+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/ethereum/go-ethereum/common"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -779,7 +779,11 @@ var bindTests = []struct {
        t.Errorf("Invalid address returned, want: %x, got: %x", (common.Address{}), res)
    }
}

+<<<<<< HEAD
+             for _, addr := range []common.Address{common.Address{}, common.Address{2}, common.Address{4}} {
+=====
+             for _, addr := range []common.Address{common.Address{}, common.Address{1}, common.Address{2}} {
+>>>>>> upstream-v0.8.5-rc.2
+                 if res, err := callfrom.CallFrom(&bind.CallOpts{From: addr}); err != nil {
+                     t.Fatalf("Failed to call constant function: %v", err)
+                 } else if res != addr {
@@ -828,9 +832,9 @@ var bindTests = []struct {
    "fmt"
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -921,10 +925,10 @@ var bindTests = []struct {
    "math/big"
    "time"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/ethereum/go-ethereum/common"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -1112,9 +1116,9 @@ var bindTests = []struct {
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -1247,9 +1251,9 @@ var bindTests = []struct {
    "math/big"
    "reflect"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -1389,9 +1393,9 @@ var bindTests = []struct {
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
,
@@ -1455,11 +1459,11 @@ var bindTests = []struct {
    "math/big"
    "time"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/core"
+             "github.com/ethereum/go-ethereum/crypto"
+             "github.com/flare-foundation/coreth/params"
,
    // Initialize test accounts
@@ -1566,10 +1570,10 @@ var bindTests = []struct {
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/ethereum/go-ethereum/crypto"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/core"
,
    // Initialize test accounts
@@ -1629,10 +1633,10 @@ var bindTests = []struct {
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"
-             "github.com/ava-labs/coreth/accounts/abi/bind/backends"
+             "github.com/flare-foundation/coreth/accounts/abi/bind"
+             "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+             "github.com/ethereum/go-ethereum/crypto"
-             "github.com/ava-labs/coreth/core"
+             "github.com/flare-foundation/coreth/core"
,
    key, _ := crypto.GenerateKey()
@@ -1691,9 +1695,9 @@ var bindTests = []struct {
    "math/big"

-             "github.com/ava-labs/coreth/accounts/abi/bind"

```

```

- "github.com/ava-labs/coreth/accounts/abi/bind/backends"
- "github.com/ava-labs/coreth/core"
+ "github.com/flare-foundation/coreth/accounts/abi/bind"
+ "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/ethereum/go-ethereum/crypto"

@@ -1752,9 +1756,9 @@ var bindTests = []struct {
    "bytes"
    "math/big"

- "github.com/ava-labs/coreth/accounts/abi/bind"
- "github.com/ava-labs/coreth/accounts/abi/bind/backends"
- "github.com/ava-labs/coreth/core"
+ "github.com/flare-foundation/coreth/accounts/abi/bind"
+ "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/ethereum/go-ethereum/crypto"

@@ -1840,9 +1844,9 @@ var bindTests = []struct {
    "math/big"

- "github.com/ava-labs/coreth/accounts/abi/bind"
- "github.com/ava-labs/coreth/accounts/abi/bind/backends"
- "github.com/ava-labs/coreth/core"
+ "github.com/flare-foundation/coreth/accounts/abi/bind"
+ "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/ethereum/go-ethereum/crypto"

@@ -1909,9 +1913,9 @@ var bindTests = []struct {
    "math/big"

- "github.com/ava-labs/coreth/accounts/abi/bind"
- "github.com/ava-labs/coreth/accounts/abi/bind/backends"
- "github.com/ava-labs/coreth/core"
+ "github.com/flare-foundation/coreth/accounts/abi/bind"
+ "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/ethereum/go-ethereum/crypto"

@@ -1942,6 +1946,52 @@ var bindTests = []struct {
    nil,
    nil,
},
+<<<<<< HEAD
+=====
+{
+    name: `ConstructorWithStructParam`,
+    contract:
+    pragma solidity >=0.8.0 <0.9.0;
+    contract ConstructorWithStructParam {
+        struct StructType {
+            uint256 field;
+        }
+        constructor(StructType memory st) {}
+    }
+    bytecode: []string{"0x608060405234801561001057600080fd5b506040516101c43803806101c48339818101604052810190610032919061014a565b50610177565b6000604051905090565b600080fd5b600080fd5b6000601f196"}
+    abi: []string[{"inputs":[{"components":[{"internalType":"uint256","name":"field","type":"uint256"}],"internalType":"struct ConstructorWithStructParam.StructType","name":"st","type"}]}]
+    imports:
+        "math/big"
+        "github.com/flare-foundation/coreth/accounts/abi/bind"
+        "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+        "github.com/flare-foundation/coreth/core"
+        "github.com/ethereum/go-ethereum/crypto"
+    tester: `
+    var (
+        key, _ = crypto.GenerateKey()
+        user, _ = bind.NewKeyedTransactorWithChainID(key, big.NewInt(1337))
+        sim = backends.NewSimulatedBackend(core.GenesisAlloc{user.From: {Balance: big.NewInt(1000000000000000000)}}), 10000000)
+    )
+    defer sim.Close()
+    _, tx, _, err := DeployConstructorWithStructParam(user, sim, ConstructorWithStructParamStructType{Field: big.NewInt(42)})
+    if err != nil {
+        t.Fatalf("DeployConstructorWithStructParam() got err %v; want nil err", err)
+    }
+    sim.Commit(true)
+    if _, err = bind.WaitDeployed(nil, sim, tx); err != nil {
+        t.Logf("Deployment tx: %v", tx)
+        t.Errorf("bind.WaitDeployed(nil, %T, <deployment tx>) got err %v; want nil err", sim, err)
+    }
+    },
+>>>>>> upstream-v0.8.5-rc.2
+}

// The binding tests have been modified to run in two separate test
@@ -1975,6 +2025,7 @@ func golangBindings(t *testing.T, overload bool) {
}
// Generate the test suite for all the contracts
for i, tt := range bindTests {
+<<<<<< HEAD
// Skip the "Overload" test if [!overload]
if !overload && tt.name == "Overload" {
continue
@@ -1999,6 +2050,33 @@ func golangBindings(t *testing.T, overload bool) {
}
// Generate the test file with the injected test code
code := fmt.Sprintf(`
+=====
+    t.Run(tt.name, func(t *testing.T) {
+        // Skip the "Overload" test if [!overload]
+        if !overload && tt.name == "Overload" {
+            return
+        }
+        // Skip all tests except for "Overload" if [overload]
+        if overload && tt.name != "Overload" {
+            return
+        }
+        var types []string
+        if tt.types != nil {
+            types = tt.types
+        } else {
+            types = []string{tt.name}
+        }
+        // Generate the binding and create a Go source file in the workspace
+        bind, err := Bind(types, tt.abi, tt.bytecode, tt.fsigs, "bindtest", LangGo, tt.libs, tt.alias)
+        if err != nil {
+            t.Fatalf("test %d: failed to generate binding: %v", i, err)
+        }
+        if err = ioutil.WriteFile(filepath.Join(pkg, strings.ToLower(tt.name)+".go"), []byte(bind), 0600); err != nil {

```

```

+         t.Fatalf("test %d: failed to write binding: %v", i, err)
+     }
+     // Generate the test file with the injected test code
+     code := fmt.Sprintf(`
+>>>>>> upstream-v0.8.5-rc.2
+package bindtest
+
+import (
@@ -2010,9 +2088,16 @@ func golangBindings(t *testing.T, overload bool) {
+    %s
+    }
+    `, tt.imports, tt.name, tt.testers)
+<<<<<< HEAD
+if err := ioutil.WriteFile(filepath.Join(pkg, strings.ToLower(tt.name)+"_test.go"), []byte(code), 0600); err != nil {
+    t.Fatalf("test %d: failed to write tests: %v", i, err)
+}
+
+=====
+if err := ioutil.WriteFile(filepath.Join(pkg, strings.ToLower(tt.name)+"_test.go"), []byte(code), 0600); err != nil {
+    t.Fatalf("test %d: failed to write tests: %v", i, err)
+}
+
+    })
+>>>>>> upstream-v0.8.5-rc.2
+    }
+    // Convert the package to go modules and use the current source for go-ethereum
+    moder := exec.Command(gocmd, "mod", "init", "bindtest")
@@ -2021,7 +2106,7 @@ func golangBindings(t *testing.T, overload bool) {
+    t.Fatalf("failed to convert binding test to modules: %v\n%s", err, out)
+}
+pwd, _ := os.Getwd()
+replacer := exec.Command(gocmd, "mod", "edit", "-x", "-require", "github.com/ava-labs/coreth@v0.0.0", "-replace", "github.com/ava-labs/coreth="+filepath.Join(pwd, "..", "..", "..")) // Repo root
+replacer := exec.Command(gocmd, "mod", "edit", "-x", "-require", "github.com/flare-foundation/coreth@v0.0.0", "-replace", "github.com/flare-foundation/coreth="+filepath.Join(pwd, "..", "..", "..")
+replacer.Dir = pkg
+if out, err := replacer.CombinedOutput(); err != nil {
+    t.Fatalf("failed to replace binding test dependency to current source tree: %v\n%s", err, out)
diff --git a/accounts/abi/bind/template.go b/accounts/abi/bind/template.go
index f66f175c..c88f741f 100644
--- a/accounts/abi/bind/template.go
+++ b/accounts/abi/bind/template.go
@@ -26,7 +26,13 @@
package bind

-import "github.com/ava-labs/coreth/accounts/abi"
+<<<<<< HEAD
+import (
+    "github.com/flare-foundation/coreth/accounts/abi"
+)
+=====
+import "github.com/flare-foundation/coreth/accounts/abi"
+>>>>>> upstream-v0.8.5-rc.2
+
+// tmplData is the data structure required to fill the binding template.
+type tmplData struct {
@@ -102,12 +108,21 @@ import (
+    "strings"
+    "errors"
+
+    "github.com/ava-labs/coreth/accounts/abi"
+    "github.com/ava-labs/coreth/accounts/abi/bind"
+    "github.com/ava-labs/coreth/core/types"
+    "github.com/ava-labs/coreth/interfaces"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+
+=====
+>>>>>> upstream-v0.8.5-rc.2
+    "github.com/flare-foundation/coreth/accounts/abi"
+    "github.com/flare-foundation/coreth/accounts/abi/bind"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+>>>>>> upstream-v0.8.5-rc.2
+    )
+
+// Reference imports to suppress errors if they are not otherwise used.
diff --git a/accounts/abi/bind/util.go b/accounts/abi/bind/util.go
index 3fb6b5c7..315d0626 100644
--- a/accounts/abi/bind/util.go
+++ b/accounts/abi/bind/util.go
@@ -31,9 +31,16 @@ import (
+    "errors"
+    "time"
+
+    "github.com/ava-labs/coreth/core/types"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+
+    "github.com/flare-foundation/coreth/core/types"
+=====
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+>>>>>> upstream-v0.8.5-rc.2
+    )
+
+// WaitMined waits for tx to be mined on the blockchain.
diff --git a/accounts/abi/bind/util_test.go b/accounts/abi/bind/util_test.go
index aab4ae89..37778769 100644
--- a/accounts/abi/bind/util_test.go
+++ b/accounts/abi/bind/util_test.go
@@ -33,12 +33,21 @@ import (
+    "testing"
+    "time"
+
+    "github.com/ava-labs/coreth/accounts/abi/bind"
+    "github.com/ava-labs/coreth/accounts/abi/bind/backends"
+    "github.com/ava-labs/coreth/core"
+    "github.com/ava-labs/coreth/core/types"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+
+=====
+>>>>>> upstream-v0.8.5-rc.2
+    "github.com/flare-foundation/coreth/accounts/abi/bind"
+    "github.com/flare-foundation/coreth/accounts/abi/bind/backends"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/types"
+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+>>>>>> upstream-v0.8.5-rc.2
+    )
+
+var testKey, _ = crypto.HexToECDSA("b71c71a67e1177ad4e901695e1b4b9ee17ae16c6668d313eac2f96dbca3f291")
diff --git a/accounts/abi/unpack.go b/accounts/abi/unpack.go
index 9bf2c0da..2eda7f68 100644
--- a/accounts/abi/unpack.go

```

```

+++ b/accounts/abi/unpack.go
@@ -300,7 +300,7 @@ func tuplePointsTo(index int, output []byte) (start int, err error) {
    offset := big.NewInt(0).SetBytes(output[index : index+32])
    outputLen := big.NewInt(int64(len(output)))

-    if offset.Cmp(big.NewInt(int64(len(output)))) > 0 {
+    if offset.Cmp(outputLen) > 0 {
        return 0, fmt.Errorf("abi: cannot marshal in to go slice: offset %v would go over slice boundary (len=%v)", offset, outputLen)
    }
    if offset.BitLen() > 63 {
diff --git a/accounts/accounts.go b/accounts/accounts.go
index dd7df0be..58d901c2 100644
--- a/accounts/accounts.go
+++ b/accounts/accounts.go
@@ -31,8 +31,8 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/event"
    "golang.org/x/crypto/sha3"
@@ -186,7 +186,11 @@ type Backend interface {
    // TextHash is a helper function that calculates a hash for the given message that can be
    // safely used to calculate a signature from.
    //
+<<<<<<< HEAD
    // The hash is calculated as
+<=====
+// The hash is calculated as
+>>>>>>> upstream-v0.8.5-rc.2
    // keccak256("\x19Ethereum Signed Message:\n"${message length}${message}).
    //
    // This gives context to the signed message and prevents signing of transactions.
@@ -198,7 +202,11 @@ func TextHash(data []byte) []byte {
    // TextAndHash is a helper function that calculates a hash for the given message that can be
    // safely used to calculate a signature from.
    //
+<<<<<<< HEAD
    // The hash is calculated as
+<=====
+// The hash is calculated as
+>>>>>>> upstream-v0.8.5-rc.2
    // keccak256("\x19Ethereum Signed Message:\n"${message length}${message}).
    //
    // This gives context to the signed message and prevents signing of transactions.
diff --git a/accounts/external/backend.go b/accounts/external/backend.go
index 558700fc..1c047667 100644
--- a/accounts/external/backend.go
+++ b/accounts/external/backend.go
@@ -31,15 +31,15 @@ import (
    "math/big"
    "sync"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ava-labs/coreth/rpc"
-    "github.com/ava-labs/coreth/signer/core/apitypes"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/event"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    "github.com/flare-foundation/coreth/rpc"
+    "github.com/flare-foundation/coreth/signer/core/apitypes"
)

type ExternalBackend struct {
diff --git a/accounts/keystore/account_cache.go b/accounts/keystore/account_cache.go
index 4c35aa74..d35d11ff 100644
--- a/accounts/keystore/account_cache.go
+++ b/accounts/keystore/account_cache.go
@@ -37,10 +37,10 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/accounts"
+    mapset "github.com/deckarep/golang-set"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
)

// Minimum amount of time between cache reloads. This limit applies if the platform does
diff --git a/accounts/keystore/account_cache_test.go b/accounts/keystore/account_cache_test.go
index ad7ab51f..01671ddf 100644
--- a/accounts/keystore/account_cache_test.go
+++ b/accounts/keystore/account_cache_test.go
@@ -36,10 +36,10 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/accounts"
+    "github.com/cespare/cp"
    "github.com/davecgh/go-spew/spew"
    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/accounts"
)

var (
diff --git a/accounts/keystore/key.go b/accounts/keystore/key.go
index 71402d36..bb7a5436 100644
--- a/accounts/keystore/key.go
+++ b/accounts/keystore/key.go
@@ -39,9 +39,9 @@ import (
    "strings"
    "time"

-    "github.com/ava-labs/coreth/accounts"
+    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/accounts"
    "github.com/google/uuid"
)

diff --git a/accounts/keystore/keystore.go b/accounts/keystore/keystore.go
index ff82ef88..e63e891f 100644
--- a/accounts/keystore/keystore.go
+++ b/accounts/keystore/keystore.go
@@ -42,10 +42,10 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/event"

```

```

+      "github.com/flare-foundation/coreth/accounts"
+      "github.com/flare-foundation/coreth/core/types"
+    )
+
+    var (
diff --git a/accounts/keystore/keystore_test.go b/accounts/keystore/keystore_test.go
index 651ab709..c04370df 100644
--- a/accounts/keystore/keystore_test.go
+++ b/accounts/keystore/keystore_test.go
@@ -38,10 +38,10 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/event"
+    "github.com/flare-foundation/coreth/accounts"
+  )

  var testSigData = make([]byte, 32)
diff --git a/accounts/keystore/passphrase.go b/accounts/keystore/passphrase.go
index 5da41f63..abe81df9 100644
--- a/accounts/keystore/passphrase.go
+++ b/accounts/keystore/passphrase.go
@@ -48,10 +48,10 @@ import (
    "os"
    "path/filepath"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/math"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/google/uuid"
+    "golang.org/x/crypto/pbkdf2"
+    "golang.org/x/crypto/scrypt"
diff --git a/accounts/keystore/presale.go b/accounts/keystore/presale.go
index 1dfbd9c2..49cb183b 100644
--- a/accounts/keystore/presale.go
+++ b/accounts/keystore/presale.go
@@ -35,8 +35,8 @@ import (
    "errors"
    "fmt"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/google/uuid"
+    "golang.org/x/crypto/pbkdf2"
+  )

  // keystoreWallet implements the accounts.Wallet interface for the original
diff --git a/accounts/scwallet/hub.go b/accounts/scwallet/hub.go
index 7a630fac..c037cebd 100644
--- a/accounts/scwallet/hub.go
+++ b/accounts/scwallet/hub.go
@@ -51,10 +51,10 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/event"
-    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    pcsc "github.com/gballet/go-libpcsclite"
+  )

diff --git a/accounts/scwallet/wallet.go b/accounts/scwallet/wallet.go
index fcecc10b..bdd7b3e2 100644
--- a/accounts/scwallet/wallet.go
+++ b/accounts/scwallet/wallet.go
@@ -43,12 +43,12 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    pcsc "github.com/gballet/go-libpcsclite"
+    "github.com/status-im/keycard-go/derivationpath"
+  )

diff --git a/chain/chain_test.go b/chain/chain_test.go
index 680e91b7..d418cd20 100644
--- a/chain/chain_test.go
+++ b/chain/chain_test.go
@@ -9,20 +9,21 @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/accounts/keystore"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/eth"
-    "github.com/ava-labs/coreth/eth/ethconfig"
-    "github.com/ava-labs/coreth/node"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
-    "github.com/ethereum/go-ethereum/log"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/accounts/keystore"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/eth"
+  )

```

```

+      "github.com/flare-foundation/coreth/eth/ethconfig"
+      "github.com/flare-foundation/coreth/node"
+      "github.com/flare-foundation/coreth/params"
+      "github.com/flare-foundation/flare/utls/timer/mockable"
+    )
+
+    type testChain struct {
+@@ -63,6 +64,7 @@ func newTestChain(name string, config *eth.Config,
+        },
+        },
+        common.Hash{},
+        &mockable.Clock{},
+    )
+    if err != nil {
+        t.Fatal(err)
+    }
diff --git a/chain/coreth.go b/chain/coreth.go
index 34543a98..d8e64045 100644
--- a/chain/coreth.go
+++ b/chain/coreth.go
@@ -7,15 +7,16 @@ import (
    "fmt"
    "time"

-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/eth"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/node"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/eth"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/node"
+    "github.com/flare-foundation/coreth/rpc"
+    "github.com/flare-foundation/flare/utls/timer/mockable"
+
+    var (
+@@ -34,12 +35,12 @@ type ETHChain struct {
+    }

+    // NewETHChain creates an Ethereum blockchain with the given configs.
+    -func NewETHChain(config *eth.Config, nodecfg *node.Config, chainDB ethdb.Database, settings eth.Settings, consensusCallbacks *dummy.ConsensusCallbacks, lastAcceptedHash common.Hash) (*ETHChain, error) {
+    +func NewETHChain(config *eth.Config, nodecfg *node.Config, chainDB ethdb.Database, settings eth.Settings, consensusCallbacks *dummy.ConsensusCallbacks, lastAcceptedHash common.Hash, clock *mockable.Clock) (*ETHChain, error) {
+        node, err := node.New(nodecfg)
+        if err != nil {
+            return nil, err
+        }
+        backend, err := eth.New(node, config, consensusCallbacks, chainDB, settings, lastAcceptedHash)
+    + backend, err := eth.New(node, config, consensusCallbacks, chainDB, settings, lastAcceptedHash, clock)
+        if err != nil {
+            return nil, fmt.Errorf("failed to create backend: %w", err)
+        }
+@@ -165,16 +166,33 @@ func (self *ETHChain) NewRPCHandler(maximumDuration time.Duration) *rpc.Server {
+    return rpc.NewServer(maximumDuration)
+    }

+    -func (self *ETHChain) AttachEthService(handler *rpc.Server, namespaces []string) {
+    -    nsmap := make(map[string]bool)
+    -    for _, ns := range namespaces {
+    -        nsmap[ns] = true
+    +// AttachEthService registers the backend RPC services provided by Ethereum
+    +// to the provided handler under their assigned namespaces.
+    +func (self *ETHChain) AttachEthService(handler *rpc.Server, namespaces []string) error {
+        enabledServicesSet := make(map[string]struct{})
+        for _, ns := range namespaces {
+            enabledServicesSet[ns] = struct{}{}
+        }

+        apiSet := make(map[string]rpc.API)
+        for _, api := range self.backend.APIs() {
+            - if nsmap[api.Namespace] {
+            -         handler.RegisterName(api.Namespace, api.Service)
+            + if existingAPI, exists := apiSet[api.Name]; exists {
+                return fmt.Errorf("duplicated API name: %s, namespaces %s and %s", api.Name, api.Namespace, existingAPI.Namespace)
+            }
+            apiSet[api.Name] = api
+        }

+        for name := range enabledServicesSet {
+            api, exists := apiSet[name]
+            if !exists {
+                return fmt.Errorf("API service %s not found", name)
+            }
+            if err := handler.RegisterName(api.Namespace, api.Service); err != nil {
+                return err
+            }
+        }

+        return nil
+    }

+    func (self *ETHChain) GetTxSubmitCh() <-chan core.NewTxEvent {
diff --git a/chain/counter_test.go b/chain/counter_test.go
index 7209ac51..af77d569 100644
--- a/chain/counter_test.go
+++ b/chain/counter_test.go
@@ -15,8 +15,8 @@ import (
    "testing"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/ethereum/go-ethereum/log"
+
+    )
+@@ -30,7 +30,7 @@ func TestCounter(t *testing.T) {
+    // NOTE: use precompiled `counter.sol` for portability, do not remove the
+    // following code (for debug purpose)
+    - //counterSrc, err := filepath.Abs(gopath + "/src/github.com/ava-labs/coreth/examples/counter/counter.sol")
+    + //counterSrc, err := filepath.Abs(gopath + "/src/github.com/flare-foundation/coreth/examples/counter/counter.sol")
+    // if err != nil {
+    //     t.Fatal(err)
+    // }
diff --git a/chain/multicoins_test.go b/chain/multicoins_test.go
index a2ce4cb5..042c9c6e 100644
--- a/chain/multicoins_test.go
+++ b/chain/multicoins_test.go
@@ -28,19 +28,20 @@ import (
    "strings"
    "testing"

-    "github.com/ava-labs/coreth/accounts/keystore"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/accounts/keystore"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+
+    )

```

```

- "github.com/ava-labs/coreth/core/types"
- "github.com/ava-labs/coreth/core/vm"
- "github.com/ava-labs/coreth/eth"
- "github.com/ava-labs/coreth/eth/ethconfig"
- "github.com/ava-labs/coreth/node"
    "github.com/ethereum/go-ethereum/accounts/abi"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/log"
+ "github.com/flare-foundation/coreth/accounts/keystore"
+ "github.com/flare-foundation/coreth/consensus/dummy"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/flare-foundation/coreth/core/rawdb"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+ "github.com/flare-foundation/coreth/eth"
+ "github.com/flare-foundation/coreth/eth/ethconfig"
+ "github.com/flare-foundation/coreth/node"
+ "github.com/flare-foundation/flare/utils/timer/mockable"
)

// TestMulticoIn tests multicoIn low-level state management and regular
@@ -72,7 +73,7 @@ func TestMulticoIn(t *testing.T) {
    //if GOPATH == "" {
    //    GOPATH = build.Default.GOPATH
    //}
-    //counterSrc, err := filepath.Abs(GOPATH + "/src/github.com/ava-labs/coreth/examples/multicoIn/mc_test.sol")
+    //counterSrc, err := filepath.Abs(GOPATH + "/src/github.com/flare-foundation/coreth/examples/multicoIn/mc_test.sol")
    //if err != nil {
    //    t.Fatal(err)
    //}
@@ -106,6 +107,7 @@ func TestMulticoIn(t *testing.T) {
    eth.DefaultSettings,
    new(dummy.ConsensusCallbacks),
    common.Hash{},
+    &mockable.Clock{},
    )
    if err != nil {
    t.Fatal(err)
diff --git a/chain/payment_test.go b/chain/payment_test.go
index 5d9da4d8..7ae81fcf 100644
--- a/chain/payment_test.go
+++ b/chain/payment_test.go
@@ -7,8 +7,8 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/types"
)

// TestPayment tests basic payment (balance, not multi-coin)
diff --git a/chain/subscribe_accepted_heads_test.go b/chain/subscribe_accepted_heads_test.go
index 0a94bfde..db69c8fa 100644
--- a/chain/subscribe_accepted_heads_test.go
+++ b/chain/subscribe_accepted_heads_test.go
@@ -4,10 +4,10 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/types"
)

func TestAcceptedHeadSubscriptions(t *testing.T) {
diff --git a/chain/subscribe_block_logs_test.go b/chain/subscribe_block_logs_test.go
index 26661dad..186e6b66 100644
--- a/chain/subscribe_block_logs_test.go
+++ b/chain/subscribe_block_logs_test.go
@@ -6,10 +6,10 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/eth/filters"
+    "github.com/flare-foundation/coreth/eth/filters"
)

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

func TestBlockLogsAllowUnfinalized(t *testing.T) {
diff --git a/chain/subscribe_transactions_test.go b/chain/subscribe_transactions_test.go
index aac6db4a..9ed43937 100644
--- a/chain/subscribe_transactions_test.go
+++ b/chain/subscribe_transactions_test.go
@@ -4,10 +4,10 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/eth/filters"
+    "github.com/flare-foundation/coreth/eth/filters"
)

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

func TestSubscribeTransactions(t *testing.T) {
diff --git a/chain/test_chain.go b/chain/test_chain.go
index e4420de1..800316f6 100644
--- a/chain/test_chain.go
+++ b/chain/test_chain.go
@@ -8,17 +8,18 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/accounts/keystore"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/eth"
-    "github.com/ava-labs/coreth/eth/ethconfig"
-    "github.com/ava-labs/coreth/node"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/flare-foundation/coreth/accounts/keystore"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/eth"
+    "github.com/flare-foundation/coreth/eth/ethconfig"
+    "github.com/flare-foundation/coreth/node"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/flare/utils/timer/mockable"
)

```



```

var (
@@ -88,6 +89,7 @@ func NewDefaultChain(t *testing.T) (*ETHChain, chan core.NewTxPoolHeadEvent, <-c
    eth.DefaultSettings,
    new(dummy.ConsensusCallbacks),
    common.Hash{},
+    &mockable.Clock{},
)
if err != nil {
t.Fatal(err)
diff --git a/changes.sh b/changes.sh
new file mode 100755
index 00000000..b2b8d301
--- /dev/null
+++ b/changes.sh
@@ -0,0 +1,4 @@
+#!/bin/bash
+# Requires wkhtmltopdf and aha tools
+# Install using: sudo apt install wkhtmltopdf aha
+git diff --color upstream-v0.7.4-rc.1 | aha > CHANGES.html && wkhtmltopdf CHANGES.html CHANGES.pdf && rm CHANGES.html
\ No newline at end of file
diff --git a/cmd/abigen/main.go b/cmd/abigen/main.go
index ebb5b590..a7f2837f 100644
--- a/cmd/abigen/main.go
+++ b/cmd/abigen/main.go
@@ -35,13 +35,13 @@ import (
    "regexp"
    "strings"

-    "github.com/ava-labs/coreth/accounts/abi"
-    "github.com/ava-labs/coreth/accounts/abi/bind"
-    "github.com/ava-labs/coreth/internal/flags"
+    "github.com/ethereum/go-ethereum/cmd/utlils"
+    "github.com/ethereum/go-ethereum/common/compiler"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts/abi"
+    "github.com/flare-foundation/coreth/accounts/abi/bind"
+    "github.com/flare-foundation/coreth/internal/flags"
    "gopkg.in/urfave/cli.v1"
)

diff --git a/consensus/consensus.go b/consensus/consensus.go
index 675be10f..cf9e4869 100644
--- a/consensus/consensus.go
+++ b/consensus/consensus.go
@@ -30,11 +30,11 @@ package consensus
import (
    "math/big"

-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
)

// ChainHeaderReader defines a small collection of methods needed to access the local
@@ -73,8 +72,7 @@ type Engine interface {
    Author(header *types.Header) (common.Address, error)

    // VerifyHeader checks whether a header conforms to the consensus rules of a
-    // given engine. Verifying the seal may be done optionally here, or explicitly
-    // via the VerifySeal method.
+    // given engine.
    //
    // NOTE: VerifyHeader does not validate the correctness of fields that rely
    // on the contents of the block (as opposed to the current and/or parent
@@ -85,10 +83,6 @@ type Engine interface {
    // rules of a given engine.
    VerifyUncles(chain ChainReader, block *types.Block) error

-    // VerifySeal checks whether the crypto seal on a header is valid according to
-    // the consensus rules of the given engine.
-    VerifySeal(chain ChainHeaderReader, header *types.Header) error
-
    // Prepare initializes the consensus fields of a block header according to the
    // rules of a particular engine. The changes are executed inline.
    Prepare(chain ChainHeaderReader, header *types.Header) error
@@ -112,9 +106,6 @@ type Engine interface {
    // that a new block should have.
    CalcDifficulty(chain ChainHeaderReader, time uint64, parent *types.Header) *big.Int

-    // APIs returns the RPC APIs this consensus engine provides.
-    APIs(chain ChainHeaderReader) []rpc.API
-
    // Close terminates any background threads maintained by the consensus engine.
    Close() error
}

diff --git a/consensus/dummy/consensus.go b/consensus/dummy/consensus.go
index 08cc229e..f965a019 100644
--- a/consensus/dummy/consensus.go
+++ b/consensus/dummy/consensus.go
@@ -10,13 +10,13 @@ import (
    "math/big"
    "time"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/rpc"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/rpc"
+    "github.com/flare-foundation/coreth/trie"
)

var (
@@ -31,20 +31,26 @@ var (
    errExtDataGasUsedTooLarge = errors.New("extDataGasUsed is not uint64")
)

+type Mode uint
+
+const (
+    ModeSkipHeader    Mode = 1 // Skip over header verification
+    ModeSkipBlockFee  Mode = 2 // Skip block fee verification
+)
+
type (
    OnFinalizeAndAssembleCallbackType = func(header *types.Header, state *state.StateDB, txs []*types.Transaction) (extraData []byte, blockFeeContribution *big.Int, extDataGasUsed *big.Int, err error)
    OnAPISCallbackType                = func(consensus.ChainHeaderReader) []rpc.API
    OnExtraStateChangeType            = func(block *types.Block, statedb *state.StateDB) (blockFeeContribution *big.Int, extDataGasUsed *big.Int, err error)

    ConsensusCallbacks struct {
-        OnAPIS
+        OnAPISCallbackType
    }

```

```

        OnFinalizeAndAssemble OnFinalizeAndAssembleCallbackType
        OnExtraStateChange    OnExtraStateChangeType
    }

    DummyEngine struct {
-       cb          *ConsensusCallbacks
-       ethFaker bool
+       cb          *ConsensusCallbacks
+       consensusMode Mode
    }
}

@@ -56,15 +62,15 @@ func NewDummyEngine(cb *ConsensusCallbacks) *DummyEngine {

func NewETHFaker() *DummyEngine {
    return &DummyEngine{
-       cb:          new(ConsensusCallbacks),
-       ethFaker: true,
+       cb:          new(ConsensusCallbacks),
+       consensusMode: ModeSkipBlockFee,
    }
}

func NewComplexETHFaker(cb *ConsensusCallbacks) *DummyEngine {
    return &DummyEngine{
-       cb:          cb,
-       ethFaker: true,
+       cb:          cb,
+       consensusMode: ModeSkipBlockFee,
    }
}

@@ -72,19 +78,29 @@ func NewFaker() *DummyEngine {
    return NewDummyEngine(new(ConsensusCallbacks))
}

+func NewFullFaker() *DummyEngine {
+    return &DummyEngine{
+        cb:          new(ConsensusCallbacks),
+        consensusMode: ModeSkipHeader,
+    }
+}

func (self *DummyEngine) verifyHeaderGasFields(config *params.ChainConfig, header *types.Header, parent *types.Header) error {
    timestamp := new(big.Int).SetUint64(header.Time)

    // Verify that the gas limit is <= 2^63-1
    cap := uint64(0x7fffffffffffffff)
-    if header.GasLimit > cap {
-        return fmt.Errorf("invalid gasLimit: have %v, max %v", header.GasLimit, cap)
+    if header.GasLimit > params.MaxGasLimit {
+        return fmt.Errorf("invalid gasLimit: have %v, max %v", header.GasLimit, params.MaxGasLimit)
    }
    // Verify that the gasUsed is <= gasLimit
    if header.GasUsed > header.GasLimit {
        return fmt.Errorf("invalid gasUsed: have %d, gasLimit %d", header.GasUsed, header.GasLimit)
    }
-    if config.IsApricotPhase1(timestamp) {
+    if config.IsApricotPhase5(timestamp) {
+        if header.GasLimit != params.ApricotPhase5GasLimit {
+            return fmt.Errorf("expected gas limit to be %d, but found %d", params.ApricotPhase5GasLimit, header.GasLimit)
+        }
+    } else if config.IsApricotPhase1(timestamp) {
        if header.GasLimit != params.ApricotPhase1GasLimit {
            return fmt.Errorf("expected gas limit to be %d, but found %d", params.ApricotPhase1GasLimit, header.GasLimit)
        }
    }

@@ -138,12 +154,16 @@ func (self *DummyEngine) verifyHeaderGasFields(config *params.ChainConfig, heade
    return nil
}

-    // Enforce Apricot Phase 4 constraints
+    // Enforce BlockGasCost constraints
+    blockGasCostStep := ApricotPhase4BlockGasCostStep
+    if config.IsApricotPhase5(timestamp) {
+        blockGasCostStep = ApricotPhase5BlockGasCostStep
+    }
    expectedBlockGasCost := calcBlockGasCost(
        ApricotPhase4TargetBlockRate,
        ApricotPhase4MinBlockGasCost,
        ApricotPhase4MaxBlockGasCost,
-        ApricotPhase4BlockGasCostStep,
+        blockGasCostStep,
        parent.BlockGasCost,
        parent.Time, header.Time,
    )

@@ -203,8 +223,7 @@ func (self *DummyEngine) verifyHeader(chain consensus.ChainHeaderReader, header
    if diff := new(big.Int).Sub(header.Number, parent.Number); diff.Cmp(big.NewInt(1)) != 0 {
        return consensus.ErrInvalidNumber
    }
-    // Verify the engine specific seal securing the block
-    return self.VerifySeal(chain, header)
+    return nil
}

func (self *DummyEngine) Author(header *types.Header) (common.Address, error) {
@@ -212,6 +231,10 @@ func (self *DummyEngine) Author(header *types.Header) (common.Address, error) {
}

func (self *DummyEngine) VerifyHeader(chain consensus.ChainHeaderReader, header *types.Header) error {
+    // If we're running a full engine faking, accept any input as valid
+    if self.consensusMode == ModeSkipHeader {
+        return nil
+    }
    // Short circuit if the header is known, or it's parent not
    number := header.Number.Uint64()
    if chain.GetHeader(header.Hash(), number) != nil {
@@ -232,10 +255,6 @@ func (self *DummyEngine) VerifyUncles(chain consensus.ChainReader, block *types.
    return nil
}

-func (self *DummyEngine) VerifySeal(chain consensus.ChainHeaderReader, header *types.Header) error {
-    return nil
-}
-

func (self *DummyEngine) Prepare(chain consensus.ChainHeaderReader, header *types.Header) error {
    header.Difficulty = big.NewInt(1)
    return nil
}

@@ -248,7 +267,7 @@ func (self *DummyEngine) verifyBlockFee(
    receipts []*types.Receipt,
    extraStateChangeContribution *big.Int,
) error {
-    if self.ethFaker {
+    if self.consensusMode == ModeSkipBlockFee {
        return nil
    }
    if baseFee == nil || baseFee.Sign() <= 0 {
@@ -324,11 +343,15 @@ func (self *DummyEngine) Finalize(chain consensus.ChainHeaderReader, block *type
    if blockExtDataGasUsed := block.ExtDataGasUsed(); blockExtDataGasUsed == nil || !blockExtDataGasUsed.IsUint64() || blockExtDataGasUsed.Cmp(extDataGasUsed) != 0 {
        return fmt.Errorf("invalid extDataGasUsed: have %d, want %d", blockExtDataGasUsed, extDataGasUsed)
    }
    blockGasCostStep := ApricotPhase4BlockGasCostStep
+    if chain.Config().IsApricotPhase5(new(big.Int).SetUint64(block.Time())) {
+        blockGasCostStep = ApricotPhase5BlockGasCostStep
+    }
}

```

```

        blockGasCost := calcBlockGasCost(
            ApricotPhase4TargetBlockRate,
            ApricotPhase4MinBlockGasCost,
            ApricotPhase4MaxBlockGasCost,
            ApricotPhase4BlockGasCostStep,
            blockGasCostStep,
            parent.BlockGasCost,
            parent.Time, block.Time(),
        )
    @@ -362,19 +385,20 @@ func (self *DummyEngine) FinalizeAndAssemble(chain consensus.ChainHeaderReader,
        return nil, err
    }
}
- if self.ethFaker {
-     extDataGasUsed = new(big.Int).Set(common.Big0)
- }
if chain.Config().IsApricotPhase4(new(big.Int).SetUint64(header.Time)) {
    header.ExtDataGasUsed = extDataGasUsed
    if header.ExtDataGasUsed == nil {
        header.ExtDataGasUsed = new(big.Int).Set(common.Big0)
    }
    blockGasCostStep := ApricotPhase4BlockGasCostStep
    if chain.Config().IsApricotPhase5(new(big.Int).SetUint64(header.Time)) {
        blockGasCostStep = ApricotPhase5BlockGasCostStep
    }
    header.BlockGasCost = calcBlockGasCost(
        ApricotPhase4TargetBlockRate,
        ApricotPhase4MinBlockGasCost,
        ApricotPhase4MaxBlockGasCost,
        ApricotPhase4BlockGasCostStep,
        blockGasCostStep,
        parent.BlockGasCost,
        parent.Time, header.Time,
    )
}
@@ -402,14 +426,6 @@ func (self *DummyEngine) CalcDifficulty(chain consensus.ChainHeaderReader, time
    return big.NewInt(1)
}

-func (self *DummyEngine) APIs(chain consensus.ChainHeaderReader) (res []rpc.API) {
-    res = nil
-    if self.cb.OnAPIs != nil {
-        res = self.cb.OnAPIs(chain)
-    }
-    return
-}

func (self *DummyEngine) Close() error {
    return nil
}

diff --git a/consensus/dummy/consensus_test.go b/consensus/dummy/consensus_test.go
index 64a84398..5aa7a66a 100644
--- a/consensus/dummy/consensus_test.go
+++ b/consensus/dummy/consensus_test.go
@@ -8,8 +8,8 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

func TestVerifyBlockFee(t *testing.T) {
diff --git a/consensus/dummy/dynamic_fees.go b/consensus/dummy/dynamic_fees.go
index e9442487..95365092 100644
--- a/consensus/dummy/dynamic_fees.go
+++ b/consensus/dummy/dynamic_fees.go
@@ -8,24 +8,28 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/avalanchego/utils/wrappers"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/flare/params/wrappers"
)

var (
-    ApricotPhase3MinBaseFee      = big.NewInt(params.ApricotPhase3MinBaseFee)
-    ApricotPhase3MaxBaseFee      = big.NewInt(params.ApricotPhase3MaxBaseFee)
-    ApricotPhase4MinBaseFee      = big.NewInt(params.ApricotPhase4MinBaseFee)
-    ApricotPhase4MaxBaseFee      = big.NewInt(params.ApricotPhase4MaxBaseFee)
-    TargetGas                    uint64 = 10_000_000
+    ApricotPhase3MinBaseFee = big.NewInt(params.ApricotPhase3MinBaseFee)
+    ApricotPhase3MaxBaseFee = big.NewInt(params.ApricotPhase3MaxBaseFee)
+    ApricotPhase4MinBaseFee = big.NewInt(params.ApricotPhase4MinBaseFee)
+    ApricotPhase4MaxBaseFee = big.NewInt(params.ApricotPhase4MaxBaseFee)
+    TargetGas                uint64 = 10_000_000
+    ApricotPhase3MinBaseFee = big.NewInt(params.ApricotPhase3MinBaseFee)
+    ApricotPhase3MaxBaseFee = big.NewInt(params.ApricotPhase3MaxBaseFee)
+    ApricotPhase4MinBaseFee = big.NewInt(params.ApricotPhase4MinBaseFee)
+    ApricotPhase4MaxBaseFee = big.NewInt(params.ApricotPhase4MaxBaseFee)
+
+    ApricotPhase4BaseFeeChangeDenominator = new(big.Int).SetUint64(params.ApricotPhase4BaseFeeChangeDenominator)
+    ApricotPhase5BaseFeeChangeDenominator = new(big.Int).SetUint64(params.ApricotPhase5BaseFeeChangeDenominator)
+
+    ApricotPhase3BlockGasFee    uint64 = 1_000_000
+    ApricotPhase4MinBlockGasCost = new(big.Int).Set(common.Big0)
+    ApricotPhase4MaxBlockGasCost = big.NewInt(1_000_000)
+    ApricotPhase4BlockGasCostStep = big.NewInt(50_000)
+    ApricotPhase4TargetBlockRate uint64 = 2 // in seconds
+    ApricotPhase4TargetBlockRate uint64 = 1 // in seconds
+    ApricotPhase5BlockGasCostStep = big.NewInt(200_000)
+    rollupWindow                  uint64 = 10
)

@@ -40,6 +44,7 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    var (
        isApricotPhase3 = config.IsApricotPhase3(bigTimestamp)
        isApricotPhase4 = config.IsApricotPhase4(bigTimestamp)
        isApricotPhase5 = config.IsApricotPhase5(bigTimestamp)
    )
    if !isApricotPhase3 || parent.Number.Cmp(common.Big0) == 0 {
        initialSlice := make([]byte, params.ApricotPhase3ExtraDataSize)
    @@ -62,12 +67,18 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
        return nil, nil, err
    }
}

// If APS, use a less responsive [BaseFeeChangeDenominator] and a higher gas
// block limit
var (
-    parentGasTarget      = TargetGas
-    parentGasTargetBig    = new(big.Int).SetUint64(parentGasTarget)
-    baseFeeChangeDenominator = new(big.Int).SetUint64(params.BaseFeeChangeDenominator)
-    baseFee               = new(big.Int).Set(parent.BaseFee)
+    parentGasTarget      = TargetGas
+    parentGasTargetBig    = new(big.Int).SetUint64(parentGasTarget)
+    baseFeeChangeDenominator = ApricotPhase4BaseFeeChangeDenominator
+    parentGasTarget        = params.ApricotPhase3TargetGas
+
+    if isApricotPhase5 {
+        baseFeeChangeDenominator = ApricotPhase5BaseFeeChangeDenominator
+        parentGasTarget = params.ApricotPhase5TargetGas
+    }
+    parentGasTargetBig := new(big.Int).SetUint64(parentGasTarget)

    // Add in the gas used by the parent block in the correct place

```

```

// If the parent consumed gas within the rollup window, add the consumed
@@ -75,7 +86,14 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    if roll < rollupWindow {
        var blockGasCost, parentExtraStateGasUsed uint64
        switch {
-           // If ApricotPhase4 is enabled, use the updated block fee calculation.
+           case isApricotPhase5:
+               // [blockGasCost] has been removed in AP5, so it is left as 0.
+
+               // At the start of a new network, the parent
+               // may not have a populated [ExtDataGasUsed].
+               if parent.ExtDataGasUsed != nil {
+                   parentExtraStateGasUsed = parent.ExtDataGasUsed.Uint64()
+               }
+           case isApricotPhase4:
+               // The [blockGasCost] is paid by the effective tips in the block using
+               // the block's value of [baseFee].
@@ -88,12 +106,11 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
        ).Uint64()

-           // On the boundary of AP3 and AP4, the parent may not have a populated
-           // [ExtDataGasUsed].
+           // On the boundary of AP3 and AP4 or at the start of a new network, the parent
+           // may not have a populated [ExtDataGasUsed].
+           if parent.ExtDataGasUsed != nil {
+               parentExtraStateGasUsed = parent.ExtDataGasUsed.Uint64()
+           }
-           // Otherwise, we must be in ApricotPhase3 and use the constant [ApricotPhase3BlockGasFee].
        default:
            blockGasCost = ApricotPhase3BlockGasFee
        }
@@ -103,10 +120,15 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    if overflow {
        addedGas = math.MaxUint64
    }
-   addedGas, overflow = math.SafeAdd(addedGas, blockGasCost)
-   if overflow {
-       addedGas = math.MaxUint64
-   }
+   // Only add the [blockGasCost] to the gas used if it isn't AP5
+   if !isApricotPhase5 {
+       addedGas, overflow = math.SafeAdd(addedGas, blockGasCost)
+       if overflow {
+           addedGas = math.MaxUint64
+       }
+   }
+   }

    slot := rollupWindow - 1 - roll
    start := slot * wrappers.LongLen
    updateLongWindow(newRollupWindow, start, addedGas)
@@ -129,7 +151,6 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    common.Big1,
)

-   // Gas price is increasing, so ensure it does not increase past the maximum
    baseFee.Add(baseFee, baseFeeDelta)
} else {
    // Otherwise if the parent block used less gas than its target, the baseFee should decrease.
@@ -152,7 +173,10 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    baseFee.Sub(baseFee, baseFeeDelta)
}

+   // Ensure that the base fee does not increase/decrease outside of the bounds
+   switch {
+   case isApricotPhase5:
+       baseFee = selectBigWithinBounds(ApricotPhase4MinBaseFee, baseFee, nil)
+   case isApricotPhase4:
+       baseFee = selectBigWithinBounds(ApricotPhase4MinBaseFee, baseFee, ApricotPhase4MaxBaseFee)
+   default:
+       baseFee = selectBigWithinBounds(ApricotPhase4MinBaseFee, baseFee, ApricotPhase4MaxBaseFee)
+   }
@@ -162,14 +186,26 @@ func CalcBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uin
    return newRollupWindow, baseFee, nil
}

+// EstimateNextBaseFee attempts to estimate the next base fee based on a block with [parent] being built at
+// [timestamp].
+// If [timestamp] is less than the timestamp of [parent], then it uses the same timestamp as parent.
+// Warning: This function should only be used in estimation and should not be used when calculating the canonical
+// base fee for a subsequent block.
+func EstimateNextBaseFee(config *params.ChainConfig, parent *types.Header, timestamp uint64) ([]byte, *big.Int, error) {
+   if timestamp < parent.Time {
+       timestamp = parent.Time
+   }
+   return CalcBaseFee(config, parent, timestamp)
+}

+// selectBigWithinBounds returns [value] if it is within the bounds:
+// lowerBound <= value <= upperBound or the bound at either end if [value]
+// is outside of the defined boundaries.
+func selectBigWithinBounds(lowerBound, value, upperBound *big.Int) *big.Int {
    switch {
-   case value.Cmp(lowerBound) < 0:
+   case lowerBound != nil && value.Cmp(lowerBound) < 0:
        return new(big.Int).Set(lowerBound)
-   case value.Cmp(upperBound) > 0:
+   case upperBound != nil && value.Cmp(upperBound) > 0:
        return new(big.Int).Set(upperBound)
    default:
        return value
    }
}

diff --git a/consensus/dummy/dynamic_fees_test.go b/consensus/dummy/dynamic_fees_test.go
index 9bbfcc12..a41d9def 100644
--- a/consensus/dummy/dynamic_fees_test.go
+++ b/consensus/dummy/dynamic_fees_test.go
@@ -8,10 +8,10 @@ import (
    "math/big"
    "testing"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/params"
+   "github.com/ethereum/go-ethereum/common/math"
+   "github.com/ethereum/go-ethereum/log"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/stretchr/testify/assert"
)

diff --git a/consensus/misc/dao.go b/consensus/misc/dao.go
index a0ab4029..be7c8df8 100644
--- a/consensus/misc/dao.go
+++ b/consensus/misc/dao.go
@@ -31,9 +31,9 @@ import (
    "errors"
    "math/big"

-   "github.com/ava-labs/coreth/core/state"
-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/params"
+   "github.com/flare-foundation/coreth/core/state"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/params"
)

var (
diff --git a/core/bench_test.go b/core/bench_test.go

```

```
new file mode 100644
index 00000000..bec5d025
--- /dev/null
+++ b/core/bench_test.go
@@ -0,0 +1,300 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2015 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "crypto/ecdsa"
+    "io/ioutil"
+    "math/big"
+    "os"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
+)
+
+func BenchmarkInsertChain_empty_memdb(b *testing.B) {
+    benchInsertChain(b, false, nil)
+}
+
+func BenchmarkInsertChain_empty_diskdb(b *testing.B) {
+    benchInsertChain(b, true, nil)
+}
+
+func BenchmarkInsertChain_valueTx_memdb(b *testing.B) {
+    benchInsertChain(b, false, genValueTx(0))
+}
+
+func BenchmarkInsertChain_valueTx_diskdb(b *testing.B) {
+    benchInsertChain(b, true, genValueTx(0))
+}
+
+func BenchmarkInsertChain_valueTx_100kB_memdb(b *testing.B) {
+    benchInsertChain(b, false, genValueTx(100*1024))
+}
+
+func BenchmarkInsertChain_valueTx_100kB_diskdb(b *testing.B) {
+    benchInsertChain(b, true, genValueTx(100*1024))
+}
+
+func BenchmarkInsertChain_ring200_memdb(b *testing.B) {
+    benchInsertChain(b, false, genTxRing(200))
+}
+
+func BenchmarkInsertChain_ring200_diskdb(b *testing.B) {
+    benchInsertChain(b, true, genTxRing(200))
+}
+
+func BenchmarkInsertChain_ring1000_memdb(b *testing.B) {
+    benchInsertChain(b, false, genTxRing(1000))
+}
+
+func BenchmarkInsertChain_ring1000_diskdb(b *testing.B) {
+    benchInsertChain(b, true, genTxRing(1000))
+}
+
+var (
+    // This is the content of the genesis block used by the benchmarks.
+    benchRootKey, _ = crypto.HexToECDSA("b71c71a67e1177ad4e901695e1b4b9ee17ae16c6668d313eac2f96dbcd3af291")
+    benchRootAddr = crypto.PubkeyToAddress(benchRootKey.PublicKey)
+    benchRootFunds = math.BigPow(2, 100)
+)
+
+// genValueTx returns a block generator that includes a single
+// value-transfer transaction with n bytes of extra data in each
+// block.
+func genValueTx(nbytes int) func(int, *BlockGen) {
+    return func(i int, gen *BlockGen) {
+        toaddr := common.Address{
+            data := make([]byte, nbytes)
+            gas, _ := IntrinsicGas(data, nil, false, false, false)
+            tx, _ := types.SignTx(types.NewTransaction(gen.TxNonce(benchRootAddr), toaddr, big.NewInt(1), gas, big.NewInt(2250000000000)), data), types.HomesteadSigner{}, benchRootKey)
+            gen.AddTx(tx)
+        }
+    }
+}
+
+var (
+    ringKeys = make([]*ecdsa.PrivateKey, 1000)
+    ringAddr = make([]common.Address, len(ringKeys))
+)
+
+func init() {
+    ringKeys[0] = benchRootKey
+    ringAddr[0] = benchRootAddr
+    for i := 1; i < len(ringKeys); i++ {
+        ringKeys[i], _ = crypto.GenerateKey()
+        ringAddr[i] = crypto.PubkeyToAddress(ringKeys[i].PublicKey)
+    }
+}
+
+// genTxRing returns a block generator that sends ether in a ring
+// among n accounts. This creates n entries in the state database
+// and fills the blocks with many small transactions.
+func genTxRing(naccounts int) func(int, *BlockGen) {
+    from := 0
+    fee := big.NewInt(0).SetUint64(params.TxGas * 2250000000000)
+    amount := big.NewInt(0).Set(benchRootFunds)
+    return func(i int, gen *BlockGen) {
+        block := gen.PrevBlock(i - 1)
+        gas := block.GasLimit()
+        for {
+            gas -= params.TxGas
+            if gas < params.TxGas {
+                break
+            }
+            to := (from + 1) % naccounts
+            tx := types.NewTransaction(

```

```

+         gen.TxNonce(ringAddrs[from]),
+         ringAddrs[to],
+         amount.Sub(amount, fee),
+         params.TxGas,
+         big.NewInt(225000000000),
+         nil,
+     )
+     tx, _ = types.SignTx(tx, types.HomesteadSigner{}, ringKeys[from])
+     gen.AddTx(tx)
+     from = to
+ }
+ }
+}
+
+func benchInsertChain(b *testing.B, disk bool, gen func(int, *BlockGen)) {
+    // Create the database in memory or in a temporary directory.
+    var db ethdb.Database
+    if !disk {
+        db = rawdb.NewMemoryDatabase()
+    } else {
+        dir, err := ioutil.TempDir("", "eth-core-bench")
+        if err != nil {
+            b.Fatalf("cannot create temporary directory: %v", err)
+        }
+        defer os.RemoveAll(dir)
+        db, err = rawdb.NewLevelDBDatabase(dir, 128, 128, "", false)
+        if err != nil {
+            b.Fatalf("cannot create temporary database: %v", err)
+        }
+        defer db.Close()
+    }
+
+    // Generate a chain of b.N blocks using the supplied block
+    // generator function.
+    gspect := Genesis{
+        Config: params.TestChainConfig,
+        Alloc:  GenesisAlloc{benchRootAddr: {Balance: benchRootFunds}},
+    }
+    genesis := gspect.MustCommit(db)
+    chain, _ := GenerateChain(gspect.Config, genesis, dummy.NewFaker(), db, b.N, 10, gen)
+
+    // Time the insertion of the new chain.
+    // State and blocks are stored in the same DB.
+    chainman := NewBlockChain(db, DefaultCacheConfig, gspect.Config, dummy.NewFaker(), vm.Config{}, common.Hash{})
+    defer chainman.Stop()
+    b.ReportAllocs()
+    b.ResetTimer()
+    if i, err := chainman.InsertChain(chain); err != nil {
+        b.Fatalf("insert error (block %d): %v\n", i, err)
+    }
+}
+
+func BenchmarkChainRead_header_10k(b *testing.B) {
+    benchReadChain(b, false, 10000)
+}
+
+func BenchmarkChainRead_full_10k(b *testing.B) {
+    benchReadChain(b, true, 10000)
+}
+
+func BenchmarkChainRead_header_100k(b *testing.B) {
+    benchReadChain(b, false, 100000)
+}
+
+func BenchmarkChainRead_full_100k(b *testing.B) {
+    benchReadChain(b, true, 100000)
+}
+
+func BenchmarkChainRead_header_500k(b *testing.B) {
+    benchReadChain(b, false, 500000)
+}
+
+func BenchmarkChainRead_full_500k(b *testing.B) {
+    benchReadChain(b, true, 500000)
+}
+
+func BenchmarkChainWrite_header_10k(b *testing.B) {
+    benchWriteChain(b, false, 10000)
+}
+
+func BenchmarkChainWrite_full_10k(b *testing.B) {
+    benchWriteChain(b, true, 10000)
+}
+
+func BenchmarkChainWrite_header_100k(b *testing.B) {
+    benchWriteChain(b, false, 100000)
+}
+
+func BenchmarkChainWrite_full_100k(b *testing.B) {
+    benchWriteChain(b, true, 100000)
+}
+
+func BenchmarkChainWrite_header_500k(b *testing.B) {
+    benchWriteChain(b, false, 500000)
+}
+
+func BenchmarkChainWrite_full_500k(b *testing.B) {
+    benchWriteChain(b, true, 500000)
+}
+
+// makeChainForBench writes a given number of headers or empty blocks/receipts
+// into a database.
+func makeChainForBench(db ethdb.Database, full bool, count uint64) {
+    var hash common.Hash
+    for n := uint64(0); n < count; n++ {
+        header := &types.Header{
+            Coinbase:    common.Address{},
+            Number:      big.NewInt(int64(n)),
+            ParentHash:   hash,
+            Difficulty:   big.NewInt(1),
+            UncleHash:    types.EmptyUncleHash,
+            TxHash:       types.EmptyRootHash,
+            ReceiptHash:  types.EmptyRootHash,
+        }
+        hash = header.Hash()
+
+        rawdb.WriteHeader(db, header)
+        rawdb.WriteCanonicalHash(db, hash, n)
+
+        if full || n == 0 {
+            block := types.NewBlockWithHeader(header)
+            rawdb.WriteBody(db, hash, n, block.Body())
+            rawdb.WriteReceipts(db, hash, n, nil)
+        }
+    }
+}
+
+func benchWriteChain(b *testing.B, full bool, count uint64) {
+    for i := 0; i < b.N; i++ {
+        dir, err := ioutil.TempDir("", "eth-chain-bench")
+        if err != nil {
+            b.Fatalf("cannot create temporary directory: %v", err)
+        }
+        db, err := rawdb.NewLevelDBDatabase(dir, 128, 1024, "", false)
+        if err != nil {
+            b.Fatalf("error opening database at %v: %v", dir, err)
+        }
+        makeChainForBench(db, full, count)
+        db.Close()
+        os.RemoveAll(dir)
+    }
+}
+
+func benchReadChain(b *testing.B, full bool, count uint64) {
+    dir, err := ioutil.TempDir("", "eth-chain-bench")

```

```

+         if err != nil {
+             b.Fatalf("cannot create temporary directory: %v", err)
+         }
+         defer os.RemoveAll(dir)
+
+         db, err := rawdb.NewLevelDBDatabase(dir, 128, 1024, "", false)
+         if err != nil {
+             b.Fatalf("error opening database at %v: %v", dir, err)
+         }
+         makeChainForBench(db, full, count)
+         db.Close()
+
+         b.ReportAllocs()
+         b.ResetTimer()
+
+         for i := 0; i < b.N; i++ {
+             db, err := rawdb.NewLevelDBDatabase(dir, 128, 1024, "", false)
+             if err != nil {
+                 b.Fatalf("error opening database at %v: %v", dir, err)
+             }
+             chain, err := NewBlockChain(db, DefaultCacheConfig, params.TestChainConfig, dummy.NewFaker(), vm.Config{}, common.Hash{})
+             if err != nil {
+                 b.Fatalf("error creating chain: %v", err)
+             }
+
+             for n := uint64(0); n < count; n++ {
+                 header := chain.GetHeaderByNumber(n)
+                 if full {
+                     hash := header.Hash()
+                     rawdb.ReadBody(db, hash, n)
+                     rawdb.ReadReceipts(db, hash, n, chain.Config())
+                 }
+             }
+             chain.Stop()
+             db.Close()
+         }
+     }
+}
diff --git a/core/block_validator.go b/core/block_validator.go
index 287a42fb..59545436 100644
--- a/core/block_validator.go
+++ b/core/block_validator.go
@@ -29,11 +29,11 @@
import (
    "fmt"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/trie"
)

// BlockValidator is responsible for validating block headers, uncles and
diff --git a/core/block_validator_test.go b/core/block_validator_test.go
new file mode 100644
index 00000000..1ea33274
--- /dev/null
+++ b/core/block_validator_test.go
@@ -0,0 +1,63 @@
// (c) 2019-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
//
// *****
// Copyright 2015 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "testing"
+)
+
+func TestCalcGasLimit(t *testing.T) {
+    for i, tc := range []struct {
+        pGasLimit uint64
+        max        uint64
+        min        uint64
+    }{
+        {200000000, 20019530, 19980470},
+        {400000000, 40039061, 39960939},
+    } {
+        // Increase
+        if have, want := CalcGasLimit(0, tc.pGasLimit, 2*tc.pGasLimit, 2*tc.pGasLimit), tc.max; have != want {
+            t.Errorf("test %d: have %d want <%d", i, have, want)
+        }
+        // Decrease
+        if have, want := CalcGasLimit(0, tc.pGasLimit, 0, 0), tc.min; have != want {
+            t.Errorf("test %d: have %d want >%d", i, have, want)
+        }
+        // Small decrease
+        if have, want := CalcGasLimit(0, tc.pGasLimit, tc.pGasLimit-1, tc.pGasLimit-1), tc.pGasLimit-1; have != want {
+            t.Errorf("test %d: have %d want %d", i, have, want)
+        }
+        // Small increase
+        if have, want := CalcGasLimit(0, tc.pGasLimit, tc.pGasLimit+1, tc.pGasLimit+1), tc.pGasLimit+1; have != want {
+            t.Errorf("test %d: have %d want %d", i, have, want)
+        }
+        // No change
+        if have, want := CalcGasLimit(0, tc.pGasLimit, tc.pGasLimit, tc.pGasLimit), tc.pGasLimit; have != want {
+            t.Errorf("test %d: have %d want %d", i, have, want)
+        }
+    }
+}
diff --git a/core/blockchain.go b/core/blockchain.go
index 9e46bcd1..a6179a34 100644
--- a/core/blockchain.go
+++ b/core/blockchain.go
@@ -37,18 +37,18 @@
import (
    "sync/atomic"
    "time"

```

```

-     "github.com/ava-labs/coreth/consensus"
-     "github.com/ava-labs/coreth/core/rawdb"
-     "github.com/ava-labs/coreth/core/state"
-     "github.com/ava-labs/coreth/core/state/snapshot"
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/core/vm"
-     "github.com/ava-labs/coreth/ethdb"
-     "github.com/ava-labs/coreth/params"
-     "github.com/ava-labs/coreth/trie"
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/event"
-     "github.com/ethereum/go-ethereum/log"
+     "github.com/flare-foundation/coreth/consensus"
+     "github.com/flare-foundation/coreth/core/rawdb"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/state/snapshot"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/vm"
+     "github.com/flare-foundation/coreth/ethdb"
+     "github.com/flare-foundation/coreth/params"
+     "github.com/flare-foundation/coreth/trie"
+     lru "github.com/hashicorp/golang-lru"
)

@@ -231,11 +231,12 @@ func NewBlockChain(
    var nilBlock *types.Block
    bc.currentBlock.Store(nilBlock)

+    // Create the state manager
+    bc.stateManager = NewTrieWriter(bc.stateCache.TrieDB(), cacheConfig)
+
    if err := bc.loadLastState(lastAcceptedHash); err != nil {
        return nil, err
    }

-    // Create the state manager
-    bc.stateManager = NewTrieWriter(bc.stateCache.TrieDB(), cacheConfig)

    // Make sure the state associated with the block is available
    head := bc.CurrentBlock()
@@ -264,17 +265,6 @@ func (bc *BlockChain) SenderCacher() *TxSenderCacher {
    return bc.senderCacher
}

-// empty returns an indicator whether the blockchain is empty.
-// func (bc *BlockChain) empty() bool {
-//     genesis := bc.genesisBlock.Hash()
-//     for i, hash := range []common.Hash{rawdb.ReadHeadBlockHash(bc.db), rawdb.ReadHeadHeaderHash(bc.db), rawdb.ReadHeadFastBlockHash(bc.db)} {
-//         if hash != genesis {
-//             return false
-//         }
-//     }
-//     return true
-// }

// loadLastState loads the last known chain state from the database. This method
// assumes that the chain manager mutex is held.
func (bc *BlockChain) loadLastState(lastAcceptedHash common.Hash) error {
@@ -305,11 +295,8 @@ func (bc *BlockChain) loadLastState(lastAcceptedHash common.Hash) error {
    }
    bc.hc.SetCurrentHeader(currentHeader)

-    headerTd := bc.GetTd(currentHeader.Hash(), currentHeader.Number.Uint64())
-    blockTd := bc.GetTd(currentBlock.Hash(), currentBlock.Number.Uint64())

-    log.Info("Loaded most recent local header", "number", currentHeader.Number, "hash", currentHeader.Hash(), "td", headerTd, "age", common.PrettyAge(time.Unix(int64(currentHeader.Time), 0)))
-    log.Info("Loaded most recent local full block", "number", currentBlock.Number(), "hash", currentBlock.Hash(), "td", blockTd, "age", common.PrettyAge(time.Unix(int64(currentBlock.Time), 0)))
+    log.Info("Loaded most recent local header", "number", currentHeader.Number, "hash", currentHeader.Hash(), "age", common.PrettyAge(time.Unix(int64(currentHeader.Time), 0)))
+    log.Info("Loaded most recent local full block", "number", currentBlock.Number(), "hash", currentBlock.Hash(), "age", common.PrettyAge(time.Unix(int64(currentBlock.Time), 0)))

    // Otherwise, set the last accepted block and perform a re-org.
    bc.lastAccepted = bc.GetBlockByHash(lastAcceptedHash)
@@ -339,10 +330,10 @@ func (bc *BlockChain) loadLastState(lastAcceptedHash common.Hash) error {
    return fmt.Errorf("failed to set preference to last accepted block while loading last state: %w", err)
}

-    // reprocessState as necessary to ensure that the last accepted state is
+    // reprocessState is necessary to ensure that the last accepted state is
// available. The state may not be available if it was not committed due
// to an unclean shutdown.
-    return bc.reprocessState(bc.lastAccepted, 2*commitInterval, true)
+    return bc.reprocessState(bc.lastAccepted, 2*commitInterval)
}

// removeIndices removes all transaction lookup entries for the transactions contained in the canonical chain
@@ -369,7 +356,6 @@ func (bc *BlockChain) removeIndices(from, to uint64) (int, error) {
    func (bc *BlockChain) loadGenesisState() error {
        // Prepare the genesis block and reinitialise the chain
        batch := bc.db.NewBatch()
-        rawdb.WriteTd(batch, bc.genesisBlock.Hash(), bc.genesisBlock.Number.Uint64(), bc.genesisBlock.Difficulty())
        rawdb.WriteBlock(batch, bc.genesisBlock)
        if err := batch.Write(); err != nil {
            log.Crit("Failed to write genesis block", "err", err)
        }
    }
@@ -418,32 +404,24 @@ func (bc *BlockChain) ExportN(w io.Writer, first uint64, last uint64) error {

    // writeHeadBlock injects a new head block into the current block chain. This method
    // assumes that the block is indeed a true head. It will also reset the head
-    // header and the head fast sync block to this very same block if they are older
-    // or if they are on a different side chain.
+    // header to this very same block if they are older or if they are on a different side chain.
    //
    // Note, this function assumes that the `mu` mutex is held!
    func (bc *BlockChain) writeHeadBlock(block *types.Block) {
        // If the block is on a side chain or an unknown one, force other heads onto it too
        updateHeads := rawdb.ReadCanonicalHash(bc.db, block.Number.Uint64()) != block.Hash()

-        // Add the block to the canonical chain number scheme and mark as the head
        batch := bc.db.NewBatch()
        rawdb.WriteCanonicalHash(batch, block.Hash(), block.Number.Uint64())

+        rawdb.WriteHeadBlockHash(batch, block.Hash())
+        rawdb.WriteHeadHeaderHash(batch, block.Hash())

-        // If the block is better than our head or is on a different chain, force update heads
-        if updateHeads {
-            rawdb.WriteHeadHeaderHash(batch, block.Hash())
-            rawdb.WriteHeadFastBlockHash(batch, block.Hash())
-        }
        // Flush the whole batch into the disk, exit the node if failed
        if err := batch.Write(); err != nil {
            log.Crit("Failed to update chain indexes and markers", "err", err)
        }
        // Update all in-memory chain markers in the last step
        if updateHeads {
            bc.hc.SetCurrentHeader(block.Header())
        }
        bc.hc.SetCurrentHeader(block.Header())
        bc.currentBlock.Store(block)
    }

@@ -561,15 +549,12 @@ func (bc *BlockChain) Stop() {
    log.Info("Blockchain stopped")
}

```



```

-// WriteStatus status of write
-type WriteStatus byte
-
-const (
-    NonStatTy WriteStatus = iota
-    CanonStatTy
-    SideStatTy
-)
-
// SetPreference attempts to update the head block to be the provided block and
// emits a ChainHeadEvent if successful. This function will handle all reorg
// side effects, if necessary.
@@ -607,7 +576,7 @@ func (bc *BlockChain) setPreference(block *types.Block) error {
    return fmt.Errorf("unable to invoke writeKnownBlock: %w", err)
}

-
-    // Send an ChainHeadEvent if we end up altering
+    // Send a ChainHeadEvent if we end up altering
    // the head block. Many internal aysnc processes rely on
    // receiving these events (i.e. the TxPool).
    bc.chainHeadFeed.Send(ChainHeadEvent{Block: block})
@@ -752,30 +721,38 @@ func (bc *BlockChain) newTip(block *types.Block) bool {
    return block.ParentHash() == bc.CurrentBlock().Hash()
}

+
+// writeBlockAndSetHead persists the block and associated state to the database
+// and optimistically updates the canonical chain if [block] extends the current
+// canonical chain.
+// writeBlockAndSetHead expects to be the last verification step during InsertBlock
+// since it creates a reference that will only be cleaned up by Accept/Reject.
+func (bc *BlockChain) writeBlockAndSetHead(block *types.Block, receipts []*types.Receipt, logs []*types.Log, state *state.StateDB) error {
+    if err := bc.writeBlockWithState(block, receipts, logs, state); err != nil {
+        return err
+    }
+
+    // If [block] represents a new tip of the canonical chain, we optimistically add it before
+    // setPreference is called. Otherwise, we consider it a side chain block.
+    if bc.newTip(block) {
+        bc.writeCanonicalBlockWithLogs(block, logs)
+    } else {
+        bc.chainSideFeed.Send(ChainSideEvent{Block: block})
+    }
+
+    return nil
+}
+
// writeBlockWithState writes the block and all associated state to the database,
// but it expects the chain mutex to be held.
-// writeBlockWithState expects to be the last verification step during InsertBlock
-// since it creates a reference that will only be cleaned up by Accept/Reject.
-func (bc *BlockChain) writeBlockWithState(block *types.Block, receipts []*types.Receipt, logs []*types.Log, state *state.StateDB) (WriteStatus, error) {
+func (bc *BlockChain) writeBlockWithState(block *types.Block, receipts []*types.Receipt, logs []*types.Log, state *state.StateDB) error {
    bc.wg.Add(1)
    defer bc.wg.Done()

-    // Calculate the total difficulty of the block
-    ptd := bc.GetTd(block.ParentHash(), block.NumberU64()-1)
-    if ptd == nil {
-        return NonStatTy, consensus.ErrUnknownAncestor
-    }
-    // Make sure no inconsistent state is leaked during insertion
-    // currentBlock := bc.CurrentBlock()
-    // localTd := bc.GetTd(currentBlock.Hash(), currentBlock.NumberU64())
-    externTd := new(big.Int).Add(block.Difficulty(), ptd)
-
-    // Irrelevant of the canonical status, write the block itself to the database.
-    //
-    // Note all the components of block(td, hash->number map, header, body, receipts)
-    // should be written atomically. BlockBatch is used for containing all components.
-    blockBatch := bc.db.NewBatch()
-    rawdb.WriteTd(blockBatch, block.Hash(), block.NumberU64(), externTd)
-    rawdb.WriteBlock(blockBatch, block)
-    rawdb.WriteReceipts(blockBatch, block.Hash(), block.NumberU64(), receipts)
-    rawdb.WritePreimages(blockBatch, state.Preimages())
@@ -792,7 +769,7 @@ func (bc *BlockChain) writeBlockWithState(block *types.Block, receipts []*types.
    _, err = state.CommitWithSnap(bc.chainConfig.IsEIP158(block.Number()), bc.snaps, block.Hash(), block.ParentHash())
}
if err != nil {
-    return NonStatTy, err
+    return err
}

// Note: if InsertTrie must be the last step in verification that can return an error.
@@ -807,18 +784,18 @@ func (bc *BlockChain) writeBlockWithState(block *types.Block, receipts []*types.
    log.Debug("failed to discard snapshot after being unable to insert block trie", "block", block.Hash(), "root", block.Root())
}
return NonStatTy, err
}

-
-    // If [block] represents a new tip of the canonical chain, we optimistically add it before
-    // setPreference is called. Otherwise, we consider it a side chain block.
-    if bc.newTip(block) {
-        bc.writeCanonicalBlockWithLogs(block, logs)
-        return CanonStatTy, nil
+    return err
+
+    bc.chainSideFeed.Send(ChainSideEvent{Block: block})
+    return SideStatTy, nil
+    return nil
}

// InsertChain attempts to insert the given batch of blocks in to the canonical
@@ -992,42 +961,42 @@ func (bc *BlockChain) insertBlock(block *types.Block, writes bool) error {
}

// Write the block to the chain and get the status.
-// writeBlockWithState creates a reference that will be cleaned up in Accept/Reject
-// so we need to ensure an error cannot occur later in verification, since that would
-// cause the referenced root to never be dereferenced.
-    status, err := bc.writeBlockWithState(block, receipts, logs, statedb)
-    if err != nil {
+    // writeBlockWithState (called within writeBlockAndSethead) creates a reference that
+    // will be cleaned up in Accept/Reject so we need to ensure an error cannot occur
+    // later in verification, since that would cause the referenced root to never be dereferenced.
+    if err := bc.writeBlockAndSetHead(block, receipts, logs, statedb); err != nil {
        return err
    }
    log.Debug("Inserted new block", "number", block.Number(), "hash", block.Hash(),
        "parentHash", block.ParentHash(),
        "uncles", len(block.Uncles()), "txs", len(block.Transactions()), "gas", block.GasUsed(),
        "elapsed", common.PrettyDuration(time.Since(start)),
        "root", block.Root(), "baseFeePerGas", block.BaseFee(), "blockGasCost", block.BlockGasCost(),
    )

-    switch status {
-    case CanonStatTy:
-        log.Debug("Inserted new block", "number", block.Number(), "hash", block.Hash(),
-            "parentHash", block.ParentHash(),
-            "uncles", len(block.Uncles()), "txs", len(block.Transactions()), "gas", block.GasUsed(),
-            "elapsed", common.PrettyDuration(time.Since(start)),
-            "root", block.Root(), "baseFeePerGas", block.BaseFee(), "blockGasCost", block.BlockGasCost(),
-        )
-    }

```

```

- // Only count canonical blocks for GC processing time
- case SideStatTy:
-     log.Debug("Inserted forked block", "number", block.Number(), "hash", block.Hash(),
-         "parentHash", block.ParentHash(),
-         "diff", block.Difficulty(), "elapsed", common.PrettyDuration(time.Since(start)),
-         "txs", len(block.Transactions()), "gas", block.GasUsed(), "uncles", len(block.Uncles()),
-         "root", block.Root(), "baseFeePerGas", block.BaseFee(), "blockGasCost", block.BlockGasCost(),
-     )
-     default:
-         // This in theory is impossible, but lets be nice to our future selves and leave
-         // a log, instead of trying to track down blocks imports that don't emit logs.
-         log.Warn("Inserted block with unknown status", "number", block.Number(), "hash", block.Hash(),
-             "parentHash", block.ParentHash(),
-             "diff", block.Difficulty(), "elapsed", common.PrettyDuration(time.Since(start)),
-             "txs", len(block.Transactions()), "gas", block.GasUsed(), "uncles", len(block.Uncles()),
-             "root", block.Root(), "baseFeePerGas", block.BaseFee(), "blockGasCost", block.BlockGasCost(),
-         )
+     return nil
+}
+
+// collectLogs collects the logs that were generated or removed during
+// the processing of the block that corresponds with the given hash.
+// These logs are later announced as deleted or reborn.
+func (bc *BlockChain) collectLogs(hash common.Hash, removed bool) []*types.Log {
+     number := bc.Hc.GetBlockNumber(hash)
+     if number == nil {
+         return nil
+     }
+     return bc.gatherBlockLogs(hash, *number, removed)
+}
-
-     return err
+// mergeLogs returns a merged log slice with specified sort order.
+func mergeLogs(logs []*types.Log, reverse bool) []*types.Log {
+     var ret []*types.Log
+     if reverse {
+         for i := len(logs) - 1; i >= 0; i-- {
+             ret = append(ret, logs[i]...)
+         }
+     } else {
+         for i := 0; i < len(logs); i++ {
+             ret = append(ret, logs[i]...)
+         }
+     }
+     return ret
+}
-
- // reorg takes two blocks, an old chain and a new chain and will reconstruct the
@@ -1044,45 +1017,17 @@ func (bc *BlockChain) reorg(oldBlock, newBlock *types.Block) error {
-     deletedLogs []*types.Log
-     rebirthLogs []*types.Log
-
-     // collectLogs collects the logs that were generated or removed during
-     // the processing of the block that corresponds with the given hash.
-     // These logs are later announced as deleted or reborn
-     collectLogs = func(hash common.Hash, removed bool) {
-         number := bc.Hc.GetBlockNumber(hash)
-         if number == nil {
-             return
-         }
-         logs := bc.gatherBlockLogs(hash, *number, removed)
-         if len(logs) > 0 {
-             if removed {
-                 deletedLogs = append(deletedLogs, logs)
-             } else {
-                 rebirthLogs = append(rebirthLogs, logs)
-             }
-         }
-     }
-     // mergeLogs returns a merged log slice with specified sort order.
-     mergeLogs = func(logs []*types.Log, reverse bool) []*types.Log {
-         var ret []*types.Log
-         if reverse {
-             for i := len(logs) - 1; i >= 0; i-- {
-                 ret = append(ret, logs[i]...)
-             }
-         } else {
-             for i := 0; i < len(logs); i++ {
-                 ret = append(ret, logs[i]...)
-             }
-         }
-         return ret
-     }
-
-     // Reduce the longer chain to the same number as the shorter one
-     if oldBlock.NumberU64() > newBlock.NumberU64() {
-         // Old chain is longer, gather all transactions and logs as deleted ones
-         for ; oldBlock != nil && oldBlock.NumberU64() != newBlock.NumberU64(); oldBlock = bc.GetBlock(oldBlock.ParentHash(), oldBlock.NumberU64()-1) {
-             oldChain = append(oldChain, oldBlock)
-             collectLogs(oldBlock.Hash(), true)
-             // Collect deleted logs for notification
-             logs := bc.collectLogs(oldBlock.Hash(), true)
-             if len(logs) > 0 {
-                 deletedLogs = append(deletedLogs, logs)
-             }
-         }
-     } else {
-         // New chain is longer, stash all blocks away for subsequent insertion
@@ -1106,7 +1051,11 @@ func (bc *BlockChain) reorg(oldBlock, newBlock *types.Block) error {
-         // Remove an old block as well as stash away a new block
-         oldChain = append(oldChain, oldBlock)
-         collectLogs(oldBlock.Hash(), true)
-         // Collect deleted logs for notification
-         logs := bc.collectLogs(oldBlock.Hash(), true)
-         if len(logs) > 0 {
-             deletedLogs = append(deletedLogs, logs)
-         }
-     }
-     newChain = append(newChain, newBlock)
-
@@ -1131,15 +1080,15 @@ func (bc *BlockChain) reorg(oldBlock, newBlock *types.Block) error {
-     // Ensure the user sees large reorgs
-     if len(oldChain) > 0 && len(newChain) > 0 {
-         logFn := log.Info
-         msg := "Chain reorg detected"
-         msg := "Resetting chain preference"
-         if len(oldChain) > 63 {
-             msg = "Large chain reorg detected"
-             msg = "Large chain preference change detected"
-             logFn = log.Warn
-         }
-         logFn(msg, "number", commonBlock.Number(), "hash", commonBlock.Hash(),
-             "drop", len(oldChain), "dropfrom", oldChain[0].Hash(), "add", len(newChain), "addfrom", newChain[0].Hash())
-     } else {
-         log.Warn("Unlikely reorg (rewind to ancestor) occurred", "oldnum", oldHead.Number(), "oldhash", oldHead.Hash(), "newnum", newHead.Number(), "newhash", newHead.Hash())
-         log.Warn("Unlikely preference change (rewind to ancestor) occurred", "oldnum", oldHead.Number(), "oldhash", oldHead.Hash(), "newnum", newHead.Number(), "newhash", newHead.Hash())
-     }
-     // Insert the new chain(except the head block(reverse order)),
-     // taking care of the proper incremental order.
@@ -1148,7 +1097,10 @@ func (bc *BlockChain) reorg(oldBlock, newBlock *types.Block) error {
-     bc.writeHeadBlock(newChain[i])

```

```

// Collect reborn logs due to chain reorg
collectLogs(newChain[i].Hash(), false)
logs := bc.collectLogs(newChain[i].Hash(), false)
if len(logs) > 0 {
    rebirthLogs = append(rebirthLogs, logs)
}
}
// Delete any canonical number assignments above the new head
indexesBatch := bc.db.NewBatch()
@@ -1255,7 +1207,7 @@ func (bc *BlockChain) RemoveRejectedBlocks(start, end uint64) error {
// it reaches a block with a state committed to the database. reprocessState does not use
// snapshots since the disk layer for snapshots will most likely be above the last committed
// state that reprocessing will start from.
-func (bc *BlockChain) reprocessState(current *types.Block, reexec uint64, report bool) error {
+func (bc *BlockChain) reprocessState(current *types.Block, reexec uint64) error {
    var (
        origin = current.NumberU64()
    )
    @@ -1300,7 +1252,7 @@ func (bc *BlockChain) reprocessState(current *types.Block, reexec uint64, report
    log.Info("Re-executing blocks to generate state for last accepted block", "from", current.NumberU64()+1, "to", origin)
    for current.NumberU64() < origin {
        // Print progress logs if long enough time elapsed
        if time.Since(logged) > 8*time.Second && report {
+            if time.Since(logged) > 8*time.Second {
                log.Info("Regenerating historical state", "block", current.NumberU64()+1, "target", origin, "remaining", origin-current.NumberU64(), "elapsed", time.Since(start))
                logged = time.Now()
            }
    @@ -1335,12 +1287,82 @@ func (bc *BlockChain) reprocessState(current *types.Block, reexec uint64, report
    }
    previousRoot = root
}
- if report {
-     nodes, imgs := triedb.Size()
-     log.Info("Historical state regenerated", "block", current.NumberU64(), "elapsed", time.Since(start), "nodes", nodes, "preimages", imgs)
- }
+ nodes, imgs := triedb.Size()
+ log.Info("Historical state regenerated", "block", current.NumberU64(), "elapsed", time.Since(start), "nodes", nodes, "preimages", imgs)
+ if previousRoot != (common.Hash{}) {
+     return triedb.Commit(previousRoot, report, nil)
+     return triedb.Commit(previousRoot, true, nil)
+ }
+ return nil
+})
+
+// CleanBlockRootsAboveLastAccepted gathers the blocks that may have previously been in processing above the
+// last accepted block and wipes their block roots from disk to mark their tries as inaccessible.
+// This is used prior to pruning to ensure that all of the tries that may still be in processing are marked
+// as inaccessible and mirrors the handling of middle roots in the geth offline pruning implementation.
+// This is not strictly necessary, but maintains a soft assumption.
+func (bc *BlockChain) CleanBlockRootsAboveLastAccepted() error {
+    targetRoot := bc.LastAcceptedBlock().Root()
+
+    // Clean up any block roots above the last accepted block before we start pruning.
+    // Note: this takes the place of middleRoots in the geth implementation since we do not
+    // track processing block roots via snapshot journals in the same way.
+    processingRoots := bc.gatherBlockRootsAboveLastAccepted()
+    // If there is a block above the last accepted block with an identical state root, we
+    // explicitly remove it from the set to ensure we do not corrupt the last accepted trie.
+    delete(processingRoots, targetRoot)
+    for processingRoot := range processingRoots {
+        // Delete the processing root from disk to mark the trie as inaccessible (no need to handle this in a batch).
+        if err := bc.db.Delete(processingRoot[:]); err != nil {
+            return fmt.Errorf("failed to remove processing root (%s) preparing for offline pruning: %w", processingRoot, err)
+        }
+    }
+
+    return nil
+}
+
+// gatherBlockRootsAboveLastAccepted iterates forward from the last accepted block and returns a list of all block roots
+// for any blocks that were inserted above the last accepted block.
+// Given that we never insert a block into the chain unless all of its ancestors have been inserted, this should gather
+// all of the block roots for blocks inserted above the last accepted block that may have been in processing at some point
+// in the past and are therefore potentially still acceptable.
+// Note: there is an edge case where the node dies while the consensus engine is rejecting a branch of blocks since the
+// consensus engine will reject the lowest ancestor first. In this case, these blocks will not be considered acceptable in
+// the future.
+// Ex.
+//   A
+//  /  \
+// B    C
+// |
+// D
+// |
+// E
+// |
+// F
+//
+// The consensus engine accepts block C and proceeds to reject the other branch in order (B, D, E, F).
+// If the consensus engine dies after rejecting block D, block D will be deleted, such that the forward iteration
+// may not find any blocks at this height and will not reach the previously processing blocks E and F.
+func (bc *BlockChain) gatherBlockRootsAboveLastAccepted() map[common.Hash]struct{} {
+    blockRoots := make(map[common.Hash]struct{})
+    for height := bc.lastAccepted.NumberU64() + 1; ; height++ {
+        blockHashes := rawdb.ReadAllHashes(bc.db, height)
+        // If there are no block hashes at [height], then there should be no further acceptable blocks
+        // past this point.
+        if len(blockHashes) == 0 {
+            break
+        }
+
+        // Fetch the blocks and append their roots.
+        for _, blockHash := range blockHashes {
+            block := bc.GetBlockByHash(blockHash)
+            if block == nil {
+                continue
+            }
+
+            blockRoots[block.Root()] = struct{}{}
+        }
+    }
+
+    return blockRoots
+})
diff --git a/core/blockchain_reader.go b/core/blockchain_reader.go
index 33adaea4..5b2646b8 100644
--- a/core/blockchain_reader.go
+++ b/core/blockchain_reader.go
@@ -27,17 +27,15 @@
package core

import (
-     "math/big"
-
-     "github.com/ava-labs/coreth/consensus"
-     "github.com/ava-labs/coreth/core/rawdb"
-     "github.com/ava-labs/coreth/core/state"
-     "github.com/ava-labs/coreth/core/state/snapshot"
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/core/vm"
-     "github.com/ava-labs/coreth/params"
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/event"

```

```

+      "github.com/flare-foundation/coreth/consensus"
+      "github.com/flare-foundation/coreth/core/rawdb"
+      "github.com/flare-foundation/coreth/core/state"
+      "github.com/flare-foundation/coreth/core/state/snapshot"
+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/core/vm"
+      "github.com/flare-foundation/coreth/params"
+    )

// CurrentHeader retrieves the current head header of the canonical chain. The
@@ -208,12 +206,6 @@ func (bc *BlockChain) GetTransactionLookup(hash common.Hash) *rawdb.LegacyTxLook
    return lookup
}

-// GetTd retrieves a block's total difficulty in the canonical chain from the
-// database by hash and number, caching it if found.
-// func (bc *BlockChain) GetTd(hash common.Hash, number uint64) *big.Int {
-//     return bc.hc.GetTd(hash, number)
-// }
-
// HasState checks if state trie is fully present in the database or not.
func (bc *BlockChain) HasState(hash common.Hash) bool {
    _, err := bc.stateCache.OpenTrie(hash)
}

diff --git a/core/blockchain_repair_test.go b/core/blockchain_repair_test.go
new file mode 100644
index 00000000..a58822a7
--- /dev/null
+++ b/core/blockchain_repair_test.go
@@ -0,0 +1,601 @@
// (c) 2019-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
// *****
// Copyright 2020 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
// Tests that abnormal program termination (i.e. crash) and restart doesn't leave
// the database in some strange state with gaps in the chain, nor with block data
// dangling in the future.
+
+package core
+
+import (
+    "io/ioutil"
+    "math/big"
+    "os"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
+)
+
+// rewindTest is a test case for chain rollback upon user request.
+type rewindTest struct {
+    canonicalBlocks int    // Number of blocks to generate for the canonical chain (heavier)
+    sidechainBlocks int    // Number of blocks to generate for the side chain (lighter)
+    commitBlock      uint64 // Block number for which to commit the state to disk
+
+    expCanonicalBlocks int    // Number of canonical blocks expected to remain in the database (excl. genesis)
+    expSidechainBlocks int    // Number of sidechain blocks expected to remain in the database (excl. genesis)
+    expHeadBlock      uint64 // Block number of the expected head full block
+}
+
+// Tests a recovery for a short canonical chain where a recent block was already
+// committed to disk and then the process crashed. In this case we expect the full
+// chain to be rolled back to the committed block, but the chain data itself left
+// in the database for replaying.
+func TestShortRepair(t *testing.T) { testShortRepair(t, false) }
+func TestShortRepairWithSnapshots(t *testing.T) { testShortRepair(t, true) }
+
+func testShortRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8
+    //
+    // Expected head block   : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 8,
+        sidechainBlocks: 0,
+        commitBlock:     4,
+        expCanonicalBlocks: 8,
+        expSidechainBlocks: 0,
+        expHeadBlock:    0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a short canonical chain and a shorter side chain, where a
+// recent block was already committed to disk and then the process crashed. In this
+// test scenario the side chain is below the committed block. In this case we expect
+// the canonical chain to be rolled back to the committed block, but the chain data
+// itself left in the database for replaying.
+func TestShortOldForkedRepair(t *testing.T) { testShortOldForkedRepair(t, false) }
+func TestShortOldForkedRepairWithSnapshots(t *testing.T) { testShortOldForkedRepair(t, true) }
+
+func testShortOldForkedRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //   L->S1->S2->S3
+    //

```

```

+ // Commit: G, C4
+ //
+ // CRASH
+ //
+ // -----
+ //
+ // Expected in leveldb:
+ // G->C1->C2->C3->C4->C5->C6->C7->C8
+ // L->S1->S2->S3
+ //
+ // Expected head block      : C4 (C0 with no snapshots)
+ rt := &rewindTest{
+     canonicalBlocks: 8,
+     sidechainBlocks: 3,
+     commitBlock: 4,
+     expCanonicalBlocks: 8,
+     expSidechainBlocks: 3,
+     expHeadBlock: 0,
+ }
+ if snapshots {
+     rt.expHeadBlock = 4
+ }
+ testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a short canonical chain and a shorter side chain, where a
+// recent block was already committed to disk and then the process crashed. In this
+// test scenario the side chain reaches above the committed block. In this case we
+// expect the canonical chain to be rolled back to the committed block, but the
+// chain data itself left in the database for replaying.
+func TestShortNewlyForkedRepair(t *testing.T) { testShortNewlyForkedRepair(t, false) }
+func TestShortNewlyForkedRepairWithSnapshots(t *testing.T) { testShortNewlyForkedRepair(t, true) }
+
+func testShortNewlyForkedRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    // L->S1->S2->S3->S4->S5->S6
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8
+    // L->S1->S2->S3->S4->S5->S6
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 8,
+        sidechainBlocks: 6,
+        commitBlock: 4,
+        expCanonicalBlocks: 8,
+        expSidechainBlocks: 6,
+        expHeadBlock: 0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a short canonical chain and a longer side chain, where a
+// recent block was already committed to disk and then the process crashed. In this
+// case we expect the canonical chain to be rolled back to the committed block, but
+// the chain data itself left in the database for replaying.
+func TestShortReorgedRepair(t *testing.T) { testShortReorgedRepair(t, false) }
+func TestShortReorgedRepairWithSnapshots(t *testing.T) { testShortReorgedRepair(t, true) }
+
+func testShortReorgedRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 8,
+        sidechainBlocks: 10,
+        commitBlock: 4,
+        expCanonicalBlocks: 8,
+        expSidechainBlocks: 10,
+        expHeadBlock: 0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain where a recent block was already
+// committed to disk and then the process crashed. In this case we expect the chain
+// to be rolled back to the committed block, but the chain data itself left in the
+// database for replaying.
+func TestLongShallowRepair(t *testing.T) { testLongShallowRepair(t, false) }
+func TestLongShallowRepairWithSnapshots(t *testing.T) { testLongShallowRepair(t, true) }
+
+func testLongShallowRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18 (HEAD)
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 18,
+        sidechainBlocks: 0,
+        commitBlock: 4,
+        expCanonicalBlocks: 18,
+        expSidechainBlocks: 0,
+        expHeadBlock: 0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}

```

```

+
+// Tests a recovery for a long canonical chain where a recent block was already committed
+// to disk and then the process crashed. In this case we expect the chain to be rolled
+// back to the committed block, but the chain data itself left in the database for replaying.
+func TestLongDeepRepair(t *testing.T) { testLongDeepRepair(t, false) }
+func TestLongDeepRepairWithSnapshots(t *testing.T) { testLongDeepRepair(t, true) }
+
+func testLongDeepRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24 (HEAD)
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb: none
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 24,
+        sidechainBlocks:  0,
+        commitBlock:      4,
+        expCanonicalBlocks: 24,
+        expSidechainBlocks: 0,
+        expHeadBlock:      0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain with a shorter side chain, where a recent
+// block was already committed to disk and then the process crashed. In this test scenario
+// the side chain is below the committed block. In this case we expect the chain to be
+// rolled back to the committed block, but the chain data itself left in the database
+// for replaying.
+func TestLongOldForkedShallowRepair(t *testing.T) {
+    testLongOldForkedShallowRepair(t, false)
+}
+func TestLongOldForkedShallowRepairWithSnapshots(t *testing.T) {
+    testLongOldForkedShallowRepair(t, true)
+}
+
+func testLongOldForkedShallowRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18 (HEAD)
+    //   L->S1->S2->S3
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18
+    //   L->S1->S2->S3
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 18,
+        sidechainBlocks:  3,
+        commitBlock:      4,
+        expCanonicalBlocks: 18,
+        expSidechainBlocks: 3,
+        expHeadBlock:      0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain a shorter side chain, where a recent block
+// was already committed to disk and then the process crashed. In this test scenario the side
+// chain is below the committed block. In this case we expect the canonical chain to be
+// rolled back to the committed block, but the chain data itself left in the database for replaying.
+func TestLongOldForkedDeepRepair(t *testing.T) { testLongOldForkedDeepRepair(t, false) }
+func TestLongOldForkedDeepRepairWithSnapshots(t *testing.T) { testLongOldForkedDeepRepair(t, true) }
+
+func testLongOldForkedDeepRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24 (HEAD)
+    //   L->S1->S2->S3
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24
+    //   L->S1->S2->S3
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 24,
+        sidechainBlocks:  3,
+        commitBlock:      4,
+        expCanonicalBlocks: 24,
+        expSidechainBlocks: 3,
+        expHeadBlock:      0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain with a shorter side chain, where a recent
+// block was already committed to disk and then the process crashed. In this test scenario
+// the side chain is above the committed block. In this case we expect the chain to be
+// rolled back to the committed block, but the chain data itself left in the database for replaying.
+func TestLongNewerForkedShallowRepair(t *testing.T) {
+    testLongNewerForkedShallowRepair(t, false)
+}
+func TestLongNewerForkedShallowRepairWithSnapshots(t *testing.T) {
+    testLongNewerForkedShallowRepair(t, true)
+}
+
+func testLongNewerForkedShallowRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18 (HEAD)
+    //   L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //

```

```

+ // -----
+ //
+ // Expected in leveledb:
+ // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18
+ // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12
+ //
+ // Expected head block      : C4 (C0 with no snapshots)
+ rt := &rewindTest{
+     canonicalBlocks: 18,
+     sidechainBlocks: 12,
+     commitBlock:     4,
+     expCanonicalBlocks: 18,
+     expSidechainBlocks: 12,
+     expHeadBlock:    0,
+ }
+ if snapshots {
+     rt.expHeadBlock = 4
+ }
+ testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain with a shorter side chain, where a recent block
+// was already committed to disk and then the process crashed. In this test scenario the side
+// chain is above the committed block. In this case we expect the canonical chain to be rolled
+// back to the committed block, but the chain data itself left in the database for replaying.
+func TestLongNewerForkedDeepRepair(t *testing.T) { testLongNewerForkedDeepRepair(t, false) }
+func TestLongNewerForkedDeepRepairWithSnapshots(t *testing.T) { testLongNewerForkedDeepRepair(t, true) }
+
+func testLongNewerForkedDeepRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24 (HEAD)
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 24,
+        sidechainBlocks: 12,
+        commitBlock:     4,
+        expCanonicalBlocks: 24,
+        expSidechainBlocks: 12,
+        expHeadBlock:    0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain with a longer side chain, where a recent block
+// was already committed to disk and then the process crashed. In this case we expect the chain to be
+// rolled back to the committed block, but the chain data itself left in the database for replaying.
+func TestLongReorgedShallowRepair(t *testing.T) { testLongReorgedShallowRepair(t, false) }
+func TestLongReorgedShallowRepairWithSnapshots(t *testing.T) { testLongReorgedShallowRepair(t, true) }
+
+func testLongReorgedShallowRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18 (HEAD)
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12->S13->S14->S15->S16->S17->S18->S19->S20->S21->S22->S23->S24->S25->S26
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12->S13->S14->S15->S16->S17->S18->S19->S20->S21->S22->S23->S24->S25->S26
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 18,
+        sidechainBlocks: 26,
+        commitBlock:     4,
+        expCanonicalBlocks: 18,
+        expSidechainBlocks: 26,
+        expHeadBlock:    0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+// Tests a recovery for a long canonical chain with a longer side chain, where a recent block
+// was already committed to disk and then the process crashed. In this case we expect the canonical
+// chains to be rolled back to the committed block, but the chain data itself left in the database
+// for replaying.
+func TestLongReorgedDeepRepair(t *testing.T) { testLongReorgedDeepRepair(t, false) }
+func TestLongReorgedDeepRepairWithSnapshots(t *testing.T) { testLongReorgedDeepRepair(t, true) }
+
+func testLongReorgedDeepRepair(t *testing.T, snapshots bool) {
+    // Chain:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24 (HEAD)
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12->S13->S14->S15->S16->S17->S18->S19->S20->S21->S22->S23->S24->S25->S26
+    //
+    // Commit: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    // G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10->C11->C12->C13->C14->C15->C16->C17->C18->C19->C20->C21->C22->C23->C24
+    // L->S1->S2->S3->S4->S5->S6->S7->S8->S9->S10->S11->S12->S13->S14->S15->S16->S17->S18->S19->S20->S21->S22->S23->S24->S25->S26
+    //
+    // Expected head block      : C4 (C0 with no snapshots)
+    rt := &rewindTest{
+        canonicalBlocks: 24,
+        sidechainBlocks: 26,
+        commitBlock:     4,
+        expCanonicalBlocks: 24,
+        expSidechainBlocks: 26,
+        expHeadBlock:    0,
+    }
+    if snapshots {
+        rt.expHeadBlock = 4
+    }
+    testRepair(t, rt, snapshots)
+}
+
+func testRepair(t *testing.T, tt *rewindTest, snapshots bool) {
+    // It's hard to follow the test case, visualize the input
+    //log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))

```

```

+ // fmt.Println(tt.dump(true))
+
+ // Create a temporary persistent database
+ datadir, err := ioutil.TempDir("", "")
+ if err != nil {
+     t.Fatalf("Failed to create temporary datadir: %v", err)
+ }
+ os.RemoveAll(datadir)
+
+ db, err := rawdb.NewLevelDBDatabase(datadir, 0, 0, "", false)
+ if err != nil {
+     t.Fatalf("Failed to create persistent database: %v", err)
+ }
+ defer db.Close() // Might double close, should be fine
+
+ // Initialize a fresh chain
+ var (
+     genesis = (&Genesis{Config: params.TestChainConfig, BaseFee: big.NewInt(params.ApricotPhase3InitialBaseFee)}).MustCommit(db)
+     engine   = dummy.NewFullFaker()
+     config   = &CacheConfig{
+         TrieCleanLimit: 256,
+         TrieDirtyLimit: 256,
+         SnapshotLimit:  0, // Disable snapshot by default
+     }
+ )
+ if snapshots {
+     config.SnapshotLimit = 256
+ }
+ chain, err := NewBlockChain(db, config, params.TestChainConfig, engine, vm.Config{}, common.Hash{})
+ if err != nil {
+     t.Fatalf("Failed to create chain: %v", err)
+ }
+ lastAcceptedHash := chain.GetBlockByNumber(0).Hash()
+
+ // If sidechain blocks are needed, make a light chain and import it
+ var sideblocks types.Blocks
+ if tt.sidechainBlocks > 0 {
+     sideblocks, _, _ = GenerateChain(params.TestChainConfig, genesis, engine, rawdb.NewMemoryDatabase(), tt.sidechainBlocks, 10, func(i int, b *BlockGen) {
+         b.SetCoinbase(common.Address{0x01})
+     })
+     if _, err := chain.InsertChain(sideblocks); err != nil {
+         t.Fatalf("Failed to import side chain: %v", err)
+     }
+ }
+
+ canonblocks, _, _ := GenerateChain(params.TestChainConfig, genesis, engine, rawdb.NewMemoryDatabase(), tt.canonicalBlocks, 10, func(i int, b *BlockGen) {
+     b.SetCoinbase(common.Address{0x02})
+     b.SetDifficulty(big.NewInt(1000000))
+ })
+ if _, err := chain.InsertChain(canonblocks[:tt.commitBlock]); err != nil {
+     t.Fatalf("Failed to import canonical chain start: %v", err)
+ }
+
+ if tt.commitBlock > 0 {
+     if snapshots {
+         for i := uint64(0); i < tt.commitBlock; i++ {
+             if err := chain.Accept(canonblocks[i]); err != nil {
+                 t.Fatalf("Failed to accept block %v: %v", i, err)
+             }
+             lastAcceptedHash = canonblocks[i].Hash()
+         }
+     }
+     if _, err := chain.InsertChain(canonblocks[tt.commitBlock:]); err != nil {
+         t.Fatalf("Failed to import canonical chain tail: %v", err)
+     }
+ }
+
+ // Pull the plug on the database, simulating a hard crash
+ db.Close()
+
+ // Start a new blockchain back up and see where the repair leads us
+ db, err = rawdb.NewLevelDBDatabase(datadir, 0, 0, "", false)
+ if err != nil {
+     t.Fatalf("Failed to reopen persistent database: %v", err)
+ }
+ defer db.Close()
+
+ chain, err = NewBlockChain(db, DefaultCacheConfig, params.TestChainConfig, engine, vm.Config{}, lastAcceptedHash)
+ if err != nil {
+     t.Fatalf("Failed to recreate chain: %v", err)
+ }
+ defer chain.Stop()
+
+ // Iterate over all the remaining blocks and ensure there are no gaps
+ verifyNoGaps(t, chain, true, canonblocks)
+ verifyNoGaps(t, chain, false, sideblocks)
+ verifyCutoff(t, chain, true, canonblocks, tt.expCanonicalBlocks)
+ verifyCutoff(t, chain, false, sideblocks, tt.expSidechainBlocks)
+
+ if head := chain.CurrentHeader(); head.Number.Uint64() != tt.expHeadBlock {
+     t.Errorf("Head header mismatch: have %d, want %d", head.Number, tt.expHeadBlock)
+ }
+ if head := chain.CurrentBlock(); head.NumberU64() != tt.expHeadBlock {
+     t.Errorf("Head block mismatch: have %d, want %d", head.NumberU64(), tt.expHeadBlock)
+ }
+}
diff --git a/core/blockchain_sethead_test.go b/core/blockchain_sethead_test.go
new file mode 100644
index 00000000..03cb332d
--- /dev/null
+++ b/core/blockchain_sethead_test.go
@@ -0,0 +1,142 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2020 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// Tests that setting the chain head backwards doesn't leave the database in some
+// strange state with gaps in the chain, nor with block data dangling in the future.
+
+package core
+
+import (
+     "testing"

```



```

+         "github.com/flare-foundation/coreth/core/types"
+)
+
+// verifyNoGaps checks that there are no gaps after the initial set of blocks in
+// the database and errors if found.
+func verifyNoGaps(t *testing.T, chain *BlockChain, canonical bool, inserted types.Blocks) {
+    t.Helper()
+
+    var end uint64
+    for i := uint64(0); i <= uint64(len(inserted)); i++ {
+        header := chain.GetHeaderByNumber(i)
+        if header == nil && end == 0 {
+            end = i
+        }
+        if header != nil && end > 0 {
+            if canonical {
+                t.Errorf("Canonical header gap between #%d-#%d", end, i-1)
+            } else {
+                t.Errorf("Sidechain header gap between #%d-#%d", end, i-1)
+            }
+            end = 0 // Reset for further gap detection
+        }
+    }
+    end = 0
+    for i := uint64(0); i <= uint64(len(inserted)); i++ {
+        block := chain.GetBlockByNumber(i)
+        if block == nil && end == 0 {
+            end = i
+        }
+        if block != nil && end > 0 {
+            if canonical {
+                t.Errorf("Canonical block gap between #%d-#%d", end, i-1)
+            } else {
+                t.Errorf("Sidechain block gap between #%d-#%d", end, i-1)
+            }
+            end = 0 // Reset for further gap detection
+        }
+    }
+    end = 0
+    for i := uint64(1); i <= uint64(len(inserted)); i++ {
+        receipts := chain.GetReceiptsByHash(inserted[i-1].Hash())
+        if receipts == nil && end == 0 {
+            end = i
+        }
+        if receipts != nil && end > 0 {
+            if canonical {
+                t.Errorf("Canonical receipt gap between #%d-#%d", end, i-1)
+            } else {
+                t.Errorf("Sidechain receipt gap between #%d-#%d", end, i-1)
+            }
+            end = 0 // Reset for further gap detection
+        }
+    }
+}
+
+// verifyCutoff checks that there are no chain data available in the chain after
+// the specified limit, but that it is available before.
+func verifyCutoff(t *testing.T, chain *BlockChain, canonical bool, inserted types.Blocks, head int) {
+    t.Helper()
+
+    for i := 1; i <= len(inserted); i++ {
+        if i <= head {
+            if header := chain.GetHeader(inserted[i-1].Hash(), uint64(i)); header == nil {
+                if canonical {
+                    t.Errorf("Canonical header   #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain header   #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+            if block := chain.GetBlock(inserted[i-1].Hash(), uint64(i)); block == nil {
+                if canonical {
+                    t.Errorf("Canonical block     #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain block     #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+            if receipts := chain.GetReceiptsByHash(inserted[i-1].Hash()); receipts == nil {
+                if canonical {
+                    t.Errorf("Canonical receipts  #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain receipts  #%2d [%x...] missing before cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+        } else {
+            if header := chain.GetHeader(inserted[i-1].Hash(), uint64(i)); header != nil {
+                if canonical {
+                    t.Errorf("Canonical header   #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain header   #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+            if block := chain.GetBlock(inserted[i-1].Hash(), uint64(i)); block != nil {
+                if canonical {
+                    t.Errorf("Canonical block     #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain block     #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+            if receipts := chain.GetReceiptsByHash(inserted[i-1].Hash()); receipts != nil {
+                if canonical {
+                    t.Errorf("Canonical receipts  #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                } else {
+                    t.Errorf("Sidechain receipts  #%2d [%x...] present after cap %d", inserted[i-1].Number(), inserted[i-1].Hash().Bytes()[3], head)
+                }
+            }
+        }
+    }
+}
+
+diff --git a/core/blockchain_snapshot_test.go b/core/blockchain_snapshot_test.go
+new file mode 100644
+index 00000000..5a92c123
+--- /dev/null
+++ b/core/blockchain_snapshot_test.go
+@@ -0,0 +1,613 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2020 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of

```

```

+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// Tests that abnormal program termination (i.e.crash) and restart can recovery
+// the snapshot properly if the snapshot is enabled.
+
+package core
+
+import (
+    "bytes"
+    "fmt"
+    "io/ioutil"
+    "math/big"
+    "os"
+    "strings"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
+)
+
+// snapshotTestBasic wraps the common testing fields in the snapshot tests.
+type snapshotTestBasic struct {
+    chainBlocks int    // Number of blocks to generate for the canonical chain
+    snapshotBlock uint64 // Block number of the relevant snapshot disk layer
+
+    expCanonicalBlocks int // Number of canonical blocks expected to remain in the database (excl. genesis)
+    expHeadBlock      uint64 // Block number of the expected head full block
+    expSnapshotBottom  uint64 // The block height corresponding to the snapshot disk layer
+
+    // share fields, set in runtime
+    datadir string
+    db      ethdb.Database
+    gendb   ethdb.Database
+    engine  consensus.Engine
+
+    lastAcceptedHash common.Hash
+}
+
+func (basic *snapshotTestBasic) prepare(t *testing.T) (*BlockChain, []*types.Block) {
+    // Create a temporary persistent database
+    datadir, err := ioutil.TempDir("", "")
+    if err != nil {
+        t.Fatalf("Failed to create temporary datadir: %v", err)
+    }
+    os.RemoveAll(datadir)
+
+    db, err := rawdb.NewLevelDBDatabase(datadir, 0, 0, "", false)
+    if err != nil {
+        t.Fatalf("Failed to create persistent database: %v", err)
+    }
+    // Initialize a fresh chain
+    var (
+        genesis = (&Genesis{Config: params.TestChainConfig, BaseFee: big.NewInt(params.ApricotPhase3InitialBaseFee)}).MustCommit(db)
+        engine  = dummy.NewFullFaker()
+        gendb   = rawdb.NewMemoryDatabase()
+
+        // Snapshot is enabled, the first snapshot is created from the Genesis.
+        // The snapshot memory allowance is 256MB, it means no snapshot flush
+        // will happen during the block insertion.
+        cacheConfig = DefaultCacheConfig
+    )
+    chain, err := NewBlockChain(db, cacheConfig, params.TestChainConfig, engine, vm.Config{}, common.Hash{})
+    if err != nil {
+        t.Fatalf("Failed to create chain: %v", err)
+    }
+    blocks, _, _ := GenerateChain(params.TestChainConfig, genesis, engine, gendb, basic.chainBlocks, 10, func(i int, b *BlockGen) {})
+
+    // genesis as last accepted
+    basic.lastAcceptedHash = chain.GetBlockByNumber(0).Hash()
+
+    // Insert the blocks with configured settings.
+    var breakpoints []uint64
+    breakpoints = append(breakpoints, basic.snapshotBlock)
+    var startPoint uint64
+    for _, point := range breakpoints {
+        if _, err := chain.InsertChain(blocks[startPoint:point]); err != nil {
+            t.Fatalf("Failed to import canonical chain start: %v", err)
+        }
+        startPoint = point
+
+        if basic.snapshotBlock > 0 && basic.snapshotBlock == point {
+            // Flushing from 0 to snapshotBlock into the disk
+            for i := uint64(0); i < point; i++ {
+                if err := chain.Accept(blocks[i]); err != nil {
+                    t.Fatalf("Failed to accept block %v: %v", i, err)
+                }
+                basic.lastAcceptedHash = blocks[i].Hash()
+            }
+
+            diskRoot, blockRoot := chain.snaps.DiskRoot(), blocks[point-1].Root()
+            if !bytes.Equal(diskRoot.Bytes(), blockRoot.Bytes()) {
+                t.Fatalf("Failed to flush disk layer change, want %x, got %x", blockRoot, diskRoot)
+            }
+        }
+    }
+    if _, err := chain.InsertChain(blocks[startPoint:]); err != nil {
+        t.Fatalf("Failed to import canonical chain tail: %v", err)
+    }
+
+    // Set runtime fields
+    basic.datadir = datadir
+    basic.db = db
+    basic.gendb = gendb
+    basic.engine = engine
+    return chain, blocks
+}
+
+func (basic *snapshotTestBasic) verify(t *testing.T, chain *BlockChain, blocks []*types.Block) {
+    // Iterate over all the remaining blocks and ensure there are no gaps
+    verifyNoGaps(t, chain, true, blocks)
+    verifyCutoff(t, chain, true, blocks, basic.expCanonicalBlocks)
+
+    if head := chain.CurrentHeader(); head.Number.Uint64() != basic.expHeadBlock {
+        t.Errorf("Head header mismatch: have %d, want %d", head.Number, basic.expHeadBlock)
+    }
+    if head := chain.CurrentBlock(); head.NumberU64() != basic.expHeadBlock {
+        t.Errorf("Head block mismatch: have %d, want %d", head.NumberU64(), basic.expHeadBlock)
+    }
+
+    // Check the disk layer, ensure they are matched
+    block := chain.GetBlockByNumber(basic.expSnapshotBottom)
+    if block == nil {
+        t.Errorf("The corresponding block[%d] of snapshot disk layer is missing", basic.expSnapshotBottom)
+    } else if !bytes.Equal(chain.snaps.DiskRoot().Bytes(), block.Root().Bytes()) {

```

```

+         t.Errorf("The snapshot disk layer root is incorrect, want %x, get %x", block.Root(), chain.snaps.DiskRoot())
+     } else if len(chain.snaps.Snapshots(block.Hash()), -1, false)) != 1 {
+         t.Errorf("The corresponding block[%d] of snapshot disk layer is missing", basic.expSnapshotBottom)
+     }
+
+     // Check the snapshot, ensure it's integrated
+     if err := chain.snaps.Verify(block.Root()); err != nil {
+         t.Errorf("The disk layer is not integrated %v", err)
+     }
+ }
+
+func (basic *snapshotTestBasic) dump() string {
+     buffer := new(strings.Builder)
+
+     fmt.Fprint(buffer, "Chain:\n G")
+     for i := 0; i < basic.chainBlocks; i++ {
+         fmt.Fprintf(buffer, "->C%d", i+1)
+     }
+     fmt.Fprint(buffer, " (HEAD)\n\n")
+
+     fmt.Fprintf(buffer, "Snapshot: G")
+     if basic.snapshotBlock > 0 {
+         fmt.Fprintf(buffer, ", C%d", basic.snapshotBlock)
+     }
+     fmt.Fprint(buffer, "\n")
+
+     //if crash {
+     //     fmt.Fprintf(buffer, "\nCRASH\n\n")
+     //} else {
+     //     fmt.Fprintf(buffer, "\nSetHead(%d)\n\n", basic.setHead)
+     //}
+     fmt.Fprint(buffer, "-----\n\n")
+
+     fmt.Fprint(buffer, "Expected in leveledb:\n G")
+     for i := 0; i < basic.expCanonicalBlocks; i++ {
+         fmt.Fprintf(buffer, "->C%d", i+1)
+     }
+     fmt.Fprint(buffer, "\n\n")
+     fmt.Fprintf(buffer, "Expected head header      : C%d\n", basic.expHeadBlock)
+     if basic.expHeadBlock == 0 {
+         fmt.Fprintf(buffer, "Expected head block       : G\n")
+     } else {
+         fmt.Fprintf(buffer, "Expected head block       : C%d\n", basic.expHeadBlock)
+     }
+     if basic.expSnapshotBottom == 0 {
+         fmt.Fprintf(buffer, "Expected snapshot disk    : G\n")
+     } else {
+         fmt.Fprintf(buffer, "Expected snapshot disk    : C%d\n", basic.expSnapshotBottom)
+     }
+     return buffer.String()
+ }
+
+func (basic *snapshotTestBasic) teardown() {
+     basic.db.Close()
+     basic.gendb.Close()
+     os.RemoveAll(basic.datadir)
+ }
+
+// snapshotTest is a test case type for normal snapshot recovery.
+// It can be used for testing that restart Geth normally.
+type snapshotTest struct {
+     snapshotTestBasic
+ }
+
+func (snaptest *snapshotTest) test(t *testing.T) {
+     // It's hard to follow the test case, visualize the input
+     // log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))
+     // fmt.Println(tt.dump())
+     chain, blocks := snaptest.prepare(t)
+
+     // Restart the chain normally
+     chain.Stop()
+     newchain, err := NewBlockChain(snaptest.db, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+     if err != nil {
+         t.Fatalf("Failed to recreate chain: %v", err)
+     }
+     defer newchain.Stop()
+
+     snaptest.verify(t, newchain, blocks)
+ }
+
+// crashSnapshotTest is a test case type for innormal snapshot recovery.
+// It can be used for testing that restart Geth after the crash.
+type crashSnapshotTest struct {
+     snapshotTestBasic
+ }
+
+func (snaptest *crashSnapshotTest) test(t *testing.T) {
+     // It's hard to follow the test case, visualize the input
+     // log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))
+     // fmt.Println(tt.dump())
+     chain, blocks := snaptest.prepare(t)
+
+     // Pull the plug on the database, simulating a hard crash
+     db := chain.db
+     db.Close()
+
+     // Start a new blockchain back up and see where the repair leads us
+     newdb, err := rawdb.NewLevelDBDatabase(snaptest.datadir, 0, 0, "", false)
+     if err != nil {
+         t.Fatalf("Failed to reopen persistent database: %v", err)
+     }
+     defer newdb.Close()
+
+     // The interesting thing is: instead of starting the blockchain after
+     // the crash, we do restart twice here: one after the crash and one
+     // after the normal stop. It's used to ensure the broken snapshot
+     // can be detected all the time.
+     newchain, err := NewBlockChain(newdb, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+     if err != nil {
+         t.Fatalf("Failed to recreate chain: %v", err)
+     }
+     newchain.Stop()
+
+     newchain, err = NewBlockChain(newdb, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+     if err != nil {
+         t.Fatalf("Failed to recreate chain: %v", err)
+     }
+     defer newchain.Stop()
+
+     snaptest.verify(t, newchain, blocks)
+ }
+
+// gappedSnapshotTest is a test type used to test this scenario:
+// - have a complete snapshot
+// - restart without enabling the snapshot
+// - insert a few blocks
+// - restart with enabling the snapshot again
+type gappedSnapshotTest struct {
+     snapshotTestBasic
+     gapped int // Number of blocks to insert without enabling snapshot
+ }
+
+func (snaptest *gappedSnapshotTest) test(t *testing.T) {

```

```

+ // It's hard to follow the test case, visualize the input
+ // log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))
+ // fmt.Println(tt.dump())
+ chain, blocks := snaptest.prepare(t)
+
+ // Insert blocks without enabling snapshot if gapping is required.
+ chain.Stop()
+ gappedBlocks, _, _ := GenerateChain(params.TestChainConfig, blocks[len(blocks)-1], snaptest.engine, snaptest.gendb, snaptest.gapped, 10, func(i int, b *BlockGen) {})
+
+ // Insert a few more blocks without enabling snapshot
+ var cacheConfig = &CacheConfig{
+     TrieCleanLimit: 256,
+     TrieDirtyLimit: 256,
+     SnapshotLimit: 0,
+ }
+ newchain, err := NewBlockChain(snaptest.db, cacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+ if err != nil {
+     t.Fatalf("Failed to recreate chain: %v", err)
+ }
+ newchain.InsertChain(gappedBlocks)
+ newchain.Stop()
+
+ // Restart the chain with enabling the snapshot
+ newchain, err = NewBlockChain(snaptest.db, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+ if err != nil {
+     t.Fatalf("Failed to recreate chain: %v", err)
+ }
+ defer newchain.Stop()
+
+ snaptest.verify(t, newchain, blocks)
+}
+
+// restartCrashSnapshotTest is the test type used to test this scenario:
+// - have a complete snapshot
+// - restart chain
+// - insert more blocks with enabling the snapshot
+// - commit the snapshot
+// - crash
+// - restart again
+type restartCrashSnapshotTest struct {
+    snapshotTestBasic
+    newBlocks int
+}
+
+func (snaptest *restartCrashSnapshotTest) test(t *testing.T) {
+    // It's hard to follow the test case, visualize the input
+    // log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))
+    // fmt.Println(tt.dump())
+    chain, blocks := snaptest.prepare(t)
+
+    // Firstly, stop the chain properly, with all snapshot journal
+    // and state committed.
+    chain.Stop()
+
+    newchain, err := NewBlockChain(snaptest.db, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+    if err != nil {
+        t.Fatalf("Failed to recreate chain: %v", err)
+    }
+    newBlocks, _, _ := GenerateChain(params.TestChainConfig, blocks[len(blocks)-1], snaptest.engine, snaptest.gendb, snaptest.newBlocks, 10, func(i int, b *BlockGen) {})
+    newchain.InsertChain(newBlocks)
+
+    // Commit the entire snapshot into the disk if requested. Note only
+    // (a) snapshot root and (b) snapshot generator will be committed,
+    // the diff journal is not.
+    for i := uint64(0); i < uint64(len(newBlocks)); i++ {
+        if err := newchain.Accept(newBlocks[i]); err != nil {
+            t.Fatalf("Failed to accept block %v: %v", i, err)
+        }
+        snaptest.lastAcceptedHash = newBlocks[i].Hash()
+    }
+
+    // Simulate the blockchain crash
+    // Don't call chain.Stop here, so that no snapshot
+    // journal and latest state will be committed
+
+    // Restart the chain after the crash
+    newchain, err = NewBlockChain(snaptest.db, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+    if err != nil {
+        t.Fatalf("Failed to recreate chain: %v", err)
+    }
+    defer newchain.Stop()
+
+    snaptest.verify(t, newchain, blocks)
+}
+
+// wipeCrashSnapshotTest is the test type used to test this scenario:
+// - have a complete snapshot
+// - restart, insert more blocks without enabling the snapshot
+// - restart again with enabling the snapshot
+// - crash
+type wipeCrashSnapshotTest struct {
+    snapshotTestBasic
+    newBlocks int
+}
+
+func (snaptest *wipeCrashSnapshotTest) test(t *testing.T) {
+    // It's hard to follow the test case, visualize the input
+    // log.Root().SetHandler(log.LvlFilterHandler(log.LvlTrace, log.StreamHandler(os.Stderr, log.TerminalFormat(true))))
+    // fmt.Println(tt.dump())
+    chain, blocks := snaptest.prepare(t)
+
+    // Firstly, stop the chain properly, with all snapshot journal
+    // and state committed.
+    chain.Stop()
+
+    config := &CacheConfig{
+        TrieCleanLimit: 256,
+        TrieDirtyLimit: 256,
+        SnapshotLimit: 0,
+    }
+    newchain, err := NewBlockChain(snaptest.db, config, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+    if err != nil {
+        t.Fatalf("Failed to recreate chain: %v", err)
+    }
+    newBlocks, _, _ := GenerateChain(params.TestChainConfig, blocks[len(blocks)-1], snaptest.engine, snaptest.gendb, snaptest.newBlocks, 10, func(i int, b *BlockGen) {})
+    newchain.InsertChain(newBlocks)
+    newchain.Stop()
+
+    // Restart the chain, the wiper should starts working
+    config = &CacheConfig{
+        TrieCleanLimit: 256,
+        TrieDirtyLimit: 256,
+        SnapshotLimit: 256,
+    }
+    newchain, err = NewBlockChain(snaptest.db, config, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+    if err != nil {
+        t.Fatalf("Failed to recreate chain: %v", err)
+    }
+    // Simulate the blockchain crash.
+
+    newchain, err = NewBlockChain(snaptest.db, DefaultCacheConfig, params.TestChainConfig, snaptest.engine, vm.Config{}, snaptest.lastAcceptedHash)
+    if err != nil {
+        t.Fatalf("Failed to recreate chain: %v", err)
+    }
+}

```

```

+         snaptest.verify(t, newchain, blocks)
+}
+
+// Tests a Geth restart with valid snapshot. Before the shutdown, all snapshot
+// journal will be persisted correctly. In this case no snapshot recovery is
+// required.
+func TestRestartWithNewSnapshot(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8
+    //
+    // Expected head header   : C8
+    // Expected head block    : C4
+    // Expected snapshot disk  : C4
+    test := &snapshotTest{
+        snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    4,
+            expCanonicalBlocks: 8,
+            expHeadBlock:     4,
+            expSnapshotBottom: 4, // Initial disk layer built from genesis
+        },
+    }
+    test.test(t)
+    test.teardown()
+}
+
+// Tests a Geth was crashed and restarts with a broken snapshot. In this case the
+// chain head should be rewound to the point with available state. And also the
+// new head should must be lower than disk layer. But there is no committed point
+// so the chain should be rewound to genesis and the disk layer should be left
+// for recovery.
+func TestNoCommitCrashWithNewSnapshot(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8
+    //
+    // Expected head block    : C4
+    // Expected snapshot disk  : C4
+    test := &crashSnapshotTest{
+        snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    4,
+            expCanonicalBlocks: 8,
+            expHeadBlock:     4,
+            expSnapshotBottom: 4, // Last committed disk layer, wait recovery
+        },
+    }
+    test.test(t)
+    test.teardown()
+}
+
+// Tests a Geth was crashed and restarts with a broken snapshot. In this case the
+// chain head should be rewound to the point with available state. And also the
+// new head should must be lower than disk layer. But there is only a low committed
+// point so the chain should be rewound to committed point and the disk layer
+// should be left for recovery.
+func TestLowCommitCrashWithNewSnapshot(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8
+    //
+    // Expected head block    : C4
+    // Expected snapshot disk  : C4
+    test := &crashSnapshotTest{
+        snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    4,
+            expCanonicalBlocks: 8,
+            expHeadBlock:     4,
+            expSnapshotBottom: 4, // Last committed disk layer, wait recovery
+        },
+    }
+    test.test(t)
+    test.teardown()
+}
+
+// Tests a Geth was crashed and restarts with a broken snapshot. In this case
+// the chain head should be rewound to the point with available state. And also
+// the new head should must be lower than disk layer. But there is only a high
+// committed point so the chain should be rewound to genesis and the disk layer
+// should be left for recovery.
+func TestHighCommitCrashWithNewSnapshot(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G, C4
+    //
+    // CRASH
+    //
+    // -----
+    //
+    // Expected in leveldb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8
+    //
+    // Expected head block    : C4
+    // Expected snapshot disk  : C4
+    test := &crashSnapshotTest{
+        snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    4,
+            expCanonicalBlocks: 8,
+            expHeadBlock:     4,
+            expSnapshotBottom: 4, // Last committed disk layer, wait recovery
+        },
+    }
+    test.test(t)
+    test.teardown()
+}
+
+// Tests a Geth was running with snapshot enabled. Then restarts without
+// enabling snapshot and after that re-enable the snapshot again. In this

```

```

+// case the snapshot should be rebuilt with latest chain head.
+func TestGappedNewSnapshot(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10
+    //
+    // Expected head block   : G
+    // Expected snapshot disk : G
+    test := &gappedSnapshotTest{
+        snapshotTestBasic: snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    0,
+            expCanonicalBlocks: 10,
+            expHeadBlock:     0,
+            expSnapshotBottom: 0, // Rebuilt snapshot from the latest HEAD
+        },
+        gapped: 2,
+    }
+    test.test(t)
+    test.teardown()
+}
+
+// Tests the Geth was running with a complete snapshot and then imports a few
+// more new blocks on top without enabling the snapshot. After the restart,
+// crash happens. Check everything is ok after the restart.
+func TestRecoverSnapshotFromWipingCrash(t *testing.T) {
+    // Chain:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8 (HEAD)
+    //
+    // Snapshot: G
+    //
+    // -----
+    //
+    // Expected in leveledb:
+    //   G->C1->C2->C3->C4->C5->C6->C7->C8->C9->C10
+    //
+    // Expected head block   : C4
+    // Expected snapshot disk : C4
+    test := &wipeCrashSnapshotTest{
+        snapshotTestBasic: snapshotTestBasic{
+            chainBlocks:      8,
+            snapshotBlock:    4,
+            expCanonicalBlocks: 10,
+            expHeadBlock:     4,
+            expSnapshotBottom: 4,
+        },
+        newBlocks: 2,
+    }
+    test.test(t)
+    test.teardown()
+}
+
+diff --git a/core/blockchain_test.go b/core/blockchain_test.go
+index 1685e04b..59909458 100644
+--- a/core/blockchain_test.go
++++ b/core/blockchain_test.go
+@@ -4,17 +4,19 @@
+ package core
+
+ import (
+     "fmt"
+     "math/big"
+     "testing"
+
+     "github.com/ava-labs/coreth/consensus/dummy"
+     "github.com/ava-labs/coreth/core/rawdb"
+     "github.com/ava-labs/coreth/core/state"
+     "github.com/ava-labs/coreth/core/types"
+     "github.com/ava-labs/coreth/core/vm"
+     "github.com/ava-labs/coreth/ethdb"
+     "github.com/ava-labs/coreth/params"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/coreth/consensus/dummy"
+     "github.com/flare-foundation/coreth/core/rawdb"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/state/pruner"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/vm"
+     "github.com/flare-foundation/coreth/ethdb"
+     "github.com/flare-foundation/coreth/params"
+ )
+
+ func TestArchiveBlockChain(t *testing.T) {
+@@ -323,3 +325,84 @@ func TestCorruptSnapshots(t *testing.T) {
+     })
+ }
+
+ +func TestBlockChainOfflinePruningUngracefulShutdown(t *testing.T) {
+     create := func(db ethdb.Database, chainConfig *params.ChainConfig, lastAcceptedHash common.Hash) (*BlockChain, error) {
+         // Import the chain. This runs all block validation rules.
+         blockchain, err := NewBlockChain(
+             db,
+             &CacheConfig{
+                 TrieCleanLimit: 256,
+                 TrieDirtyLimit: 256,
+                 Pruning:        true, // Enable pruning
+                 SnapshotLimit: 256,
+             },
+             chainConfig,
+             dummy.NewDummyEngine(&dummy.ConsensusCallbacks{
+                 OnExtraStateChange: func(block *types.Block, sdb *state.StateDB) (*big.Int, *big.Int, error) {
+                     sdb.SetBalanceMultiCoin(common.HexToAddress("0xdeadbeef"), common.HexToHash("0xdeadbeef"), big.NewInt(block.Number().Int64()))
+                     return nil, nil, nil
+                 },
+                 OnFinalizeAndAssemble: func(header *types.Header, sdb *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+                     sdb.SetBalanceMultiCoin(common.HexToAddress("0xdeadbeef"), common.HexToHash("0xdeadbeef"), big.NewInt(header.Number.Int64()))
+                     return nil, nil, nil, nil
+                 },
+             }),
+             vm.Config{
+                 lastAcceptedHash,
+             },
+         )
+         if err != nil {
+             return nil, err
+         }
+
+         // Overwrite state manager, so that Shutdown is not called.
+         // This tests to ensure that the state manager handles an ungraceful shutdown correctly.
+         blockchain.stateManager = &wrappedStateManager{TrieWriter: blockchain.stateManager}
+
+         if lastAcceptedHash == (common.Hash{}) {
+             return blockchain, nil
+         }
+
+         tempDir := t.TempDir()
+         if err := blockchain.CleanBlockRootsAboveLastAccepted(); err != nil {
+             return nil, err
+         }
+     }
+ }

```

```

+         pruner, err := pruner.NewPruner(db, tempDir, 256)
+         if err != nil {
+             return nil, fmt.Errorf("offline pruning failed (%s, %d): %w", tempDir, 256, err)
+         }
+
+         targetRoot := blockchain.LastAcceptedBlock().Root()
+         if err := pruner.Prune(targetRoot); err != nil {
+             return nil, fmt.Errorf("failed to prune blockchain with target root: %s due to: %w", targetRoot, err)
+         }
+         // Re-initialize the blockchain after pruning
+         return NewBlockChain(
+             db,
+             &CacheConfig{
+                 TrieCleanLimit: 256,
+                 TrieDirtyLimit: 256,
+                 Pruning:         true, // Enable pruning
+                 SnapshotLimit: 256,
+             },
+             chainConfig,
+             dummy.NewDummyEngine(&dummy.ConsensusCallbacks{
+                 OnExtraStateChange: func(block *types.Block, sdb *state.StateDB) (*big.Int, *big.Int, error) {
+                     sdb.SetBalanceMultiCoin(common.HexToAddress("0xdeadbeef"), common.HexToHash("0xdeadbeef"), big.NewInt(block.Number().Int64()))
+                     return nil, nil, nil
+                 },
+                 OnFinalizeAndAssemble: func(header *types.Header, sdb *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+                     sdb.SetBalanceMultiCoin(common.HexToAddress("0xdeadbeef"), common.HexToHash("0xdeadbeef"), big.NewInt(header.Number.Int64()))
+                     return nil, nil, nil, nil
+                 },
+             }),
+             vm.Config{},
+             lastAcceptedHash,
+         )
+     }
+     for _, tt := range tests {
+         t.Run(tt.Name, func(t *testing.T) {
+             tt.testFunc(t, create)
+         })
+     }
+ }
+ }
+ }
+ }
diff --git a/core/bloom_indexer.go b/core/bloom_indexer.go
index 4d8f58fd..9404a818 100644
--- a/core/bloom_indexer.go
+++ b/core/bloom_indexer.go
@@ -20,12 +20,12 @@ import (
     "context"
     "time"

-    "github.com/ava-labs/coreth/core/bloombits"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/bitutil"
+    "github.com/flare-foundation/coreth/core/bloombits"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
 )

const (
diff --git a/core/bloombits/generator.go b/core/bloombits/generator.go
index c0422caa..32e011a4 100644
--- a/core/bloombits/generator.go
+++ b/core/bloombits/generator.go
@@ -29,7 +29,7 @@ package bloombits
import (
    "errors"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/types"
)

var (
diff --git a/core/bloombits/generator_test.go b/core/bloombits/generator_test.go
index 067c1db6..345ec4e6 100644
--- a/core/bloombits/generator_test.go
+++ b/core/bloombits/generator_test.go
@@ -31,7 +31,7 @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/types"
)

// Tests that batched bloom bits are correctly rotated from the input bloom
diff --git a/core/chain_indexer.go b/core/chain_indexer.go
index 975f82b3..7163929c 100644
--- a/core/chain_indexer.go
+++ b/core/chain_indexer.go
@@ -34,12 +34,12 @@ import (
    "sync/atomic"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
)

// ChainIndexerBackend defines the methods needed to process chain segments in
diff --git a/core/chain_indexer_test.go b/core/chain_indexer_test.go
index 3edf175d..481f4e15 100644
--- a/core/chain_indexer_test.go
+++ b/core/chain_indexer_test.go
@@ -35,9 +35,9 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
)

// Runs multiple tests with randomized parameters.
diff --git a/core/chain_makers.go b/core/chain_makers.go
index 9830ff71..e65d58ab 100644
--- a/core/chain_makers.go
+++ b/core/chain_makers.go
@@ -30,15 +30,15 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/consensus/misc"
+    "github.com/ava-labs/coreth/consensus"
+    "github.com/ava-labs/coreth/consensus/state"
)

```

```

- "github.com/ava-labs/coreth/core/types"
- "github.com/ava-labs/coreth/core/vm"
- "github.com/ava-labs/coreth/ethdb"
- "github.com/ava-labs/coreth/params"
- "github.com/ethereum/go-ethereum/common"
+ "github.com/flare-foundation/coreth/consensus"
+ "github.com/flare-foundation/coreth/consensus/dummy"
+ "github.com/flare-foundation/coreth/consensus/misc"
+ "github.com/flare-foundation/coreth/core/state"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+ "github.com/flare-foundation/coreth/ethdb"
+ "github.com/flare-foundation/coreth/params"
)

// BlockGen creates blocks for testing.
@@ -111,7 +111,7 @@ func (b *BlockGen) AddTx(tx *types.Transaction) {
// the block in chain will be returned.
func (b *BlockGen) AddTxWithChain(bc *BlockChain, tx *types.Transaction) {
    if b.gasPool == nil {
-       b.SetCoinbase(common.Address{})
+       b.SetCoinbase(common.HexToAddress("0x0100000000000000000000000000000000000000"))
    }
    b.statedb.Prepare(tx.Hash(), len(b.txs))
    receipt, err := ApplyTransaction(b.config, bc, &b.header.Coinbase, b.gasPool, b.statedb, b.header, tx, &b.header.GasUsed, vm.Config{})
@@ -217,11 +217,15 @@ func GenerateChain(config *params.ChainConfig, parent *types.Block, engine conse
    b.header = makeHeader(chainreader, config, parent, gap, statedb, b.engine)

    // Mutate the state and block according to any hard-fork specs
    if daoBlock := config.DAOForkBlock; daoBlock != nil {
-       limit := new(big.Int).Add(daoBlock, params.DAOForkExtraRange)
-       if b.header.Number.Cmp(daoBlock) >= 0 && b.header.Number.Cmp(limit) < 0 {
-           if config.DAOForkSupport {
-               b.header.Extra = common.CopyBytes(params.DAOForkBlockExtra)
+       timestamp := new(big.Int).SetUint64(b.header.Time)
+       if !config.IsApricotPhase3(timestamp) {
+           // avoid dynamic fee extra data override
+           if daoBlock := config.DAOForkBlock; daoBlock != nil {
+               limit := new(big.Int).Add(daoBlock, params.DAOForkExtraRange)
+               if b.header.Number.Cmp(daoBlock) >= 0 && b.header.Number.Cmp(limit) < 0 {
+                   if config.DAOForkSupport {
+                       b.header.Extra = common.CopyBytes(params.DAOForkBlockExtra)
+                   }
+               }
+           }
+       }
    }
}

@@ -277,7 +281,9 @@ func makeHeader(chain consensus.ChainReader, config *params.ChainConfig, parent
    timestamp := new(big.Int).SetUint64(time)
    var gasLimit uint64
-    if config.IsApricotPhase1(timestamp) {
+    if config.IsApricotPhase5(timestamp) {
+        gasLimit = params.ApricotPhase5GasLimit
+    } else if config.IsApricotPhase1(timestamp) {
        gasLimit = params.ApricotPhase1GasLimit
    } else {
        gasLimit = CalcGasLimit(parent.GasUsed(), parent.GasLimit(), parent.GasLimit(), parent.GasLimit())
}

diff --git a/core/chain_makers_test.go b/core/chain_makers_test.go
index cd003932..84a4b7e6 100644
--- a/core/chain_makers_test.go
+++ b/core/chain_makers_test.go
@@ -30,13 +30,13 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
)

func ExampleGenerateChain() {
diff --git a/core/dao_test.go b/core/dao_test.go
new file mode 100644
index 00000000..16cb7de7
--- /dev/null
+++ b/core/dao_test.go
@@ -0,0 +1,197 @@
@@ -0,0 +1,197 @@
+// (c) 2021-2022, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2016 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "math/big"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
+)
+
+// Tests that DAO-fork enabled clients can properly filter out fork-commencing
+// blocks based on their extradata fields.
+func TestDAOForkRangeExtradata(t *testing.T) {
+    forkBlock := big.NewInt(32)
+
+    // Generate a common prefix for both pro-forkers and non-forkers
+    db := rawdb.NewMemoryDatabase()
+    gspec := &Genesis{
+        BaseFee: big.NewInt(params.ApricotPhase3InitialBaseFee),

```



```

+         Config: params.TestApricotPhase2Config,
+     }
+     genesis := gspec.MustCommit(db)
+     prefix, _, _ := GenerateChain(params.TestApricotPhase2Config, genesis, dummy.NewFaker(), db, int(forkBlock.Int64()-1), 10, func(i int, gen *BlockGen) {})
+
+     // Create the concurrent, conflicting two nodes
+     proDb := rawdb.NewMemoryDatabase()
+     gspec.MustCommit(proDb)
+
+     proConf := *params.TestApricotPhase2Config
+     proConf.DAOForkBlock = forkBlock
+     proConf.DAOForkSupport = true
+
+     proBc, _ := NewBlockChain(proDb, DefaultCacheConfig, &proConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+     defer proBc.Stop()
+
+     conDb := rawdb.NewMemoryDatabase()
+     gspec.MustCommit(conDb)
+
+     conConf := *params.TestApricotPhase2Config
+     conConf.DAOForkBlock = forkBlock
+     conConf.DAOForkSupport = false
+
+     conBc, _ := NewBlockChain(conDb, DefaultCacheConfig, &conConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+     defer conBc.Stop()
+
+     if _, err := proBc.InsertChain(prefix); err != nil {
+         t.Fatalf("pro-fork: failed to import chain prefix: %v", err)
+     }
+     if _, err := conBc.InsertChain(prefix); err != nil {
+         t.Fatalf("con-fork: failed to import chain prefix: %v", err)
+     }
+
+     // Try to expand both pro-fork and non-fork chains iteratively with other camp's blocks
+     for i := int64(0); i < params.DAOForkExtraRange.Int64(); i++ {
+         // Create a pro-fork block, and try to feed into the no-fork chain
+         db = rawdb.NewMemoryDatabase()
+         gspec.MustCommit(db)
+         bc, _ := NewBlockChain(db, DefaultCacheConfig, &conConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+         defer bc.Stop()
+
+         blocks := conBc.GetBlocksFromHash(conBc.CurrentBlock().Hash(), int(conBc.CurrentBlock().NumberU64()))
+         for j := 0; j < len(blocks)/2; j++ {
+             blocks[j], blocks[len(blocks)-1-j] = blocks[len(blocks)-1-j], blocks[j]
+         }
+         if _, err := bc.InsertChain(blocks); err != nil {
+             t.Fatalf("failed to import contra-fork chain for expansion: %v", err)
+         }
+         if err := bc.stateCache.TrieDB().Commit(bc.CurrentHeader().Root, true, nil); err != nil {
+             t.Fatalf("failed to commit contra-fork head for expansion: %v", err)
+         }
+         blocks, _, _ = GenerateChain(&proConf, conBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+         if _, err := conBc.InsertChain(blocks); err != nil {
+             t.Fatalf("contra-fork chain accepted pro-fork block: %v", blocks[0])
+         }
+         // Create a proper no-fork block for the contra-forker
+         blocks, _, _ = GenerateChain(&conConf, conBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+         if _, err := conBc.InsertChain(blocks); err != nil {
+             t.Fatalf("contra-fork chain didn't accepted no-fork block: %v", err)
+         }
+         // Create a no-fork block, and try to feed into the pro-fork chain
+         db = rawdb.NewMemoryDatabase()
+         gspec.MustCommit(db)
+         bc, _ = NewBlockChain(db, DefaultCacheConfig, &proConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+         defer bc.Stop()
+
+         blocks = proBc.GetBlocksFromHash(proBc.CurrentBlock().Hash(), int(proBc.CurrentBlock().NumberU64()))
+         for j := 0; j < len(blocks)/2; j++ {
+             blocks[j], blocks[len(blocks)-1-j] = blocks[len(blocks)-1-j], blocks[j]
+         }
+         if _, err := bc.InsertChain(blocks); err != nil {
+             t.Fatalf("failed to import pro-fork chain for expansion: %v", err)
+         }
+         if err := bc.stateCache.TrieDB().Commit(bc.CurrentHeader().Root, true, nil); err != nil {
+             t.Fatalf("failed to commit pro-fork head for expansion: %v", err)
+         }
+         blocks, _, _ = GenerateChain(&conConf, proBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+         if _, err := proBc.InsertChain(blocks); err != nil {
+             t.Fatalf("pro-fork chain accepted contra-fork block: %v", blocks[0])
+         }
+         // Create a proper pro-fork block for the pro-forker
+         blocks, _, _ = GenerateChain(&proConf, proBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+         if _, err := proBc.InsertChain(blocks); err != nil {
+             t.Fatalf("pro-fork chain didn't accepted pro-fork block: %v", err)
+         }
+     }
+
+     // Verify that contra-forkers accept pro-fork extra-datas after forking finishes
+     db = rawdb.NewMemoryDatabase()
+     gspec.MustCommit(db)
+     bc, _ := NewBlockChain(db, DefaultCacheConfig, &conConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+     defer bc.Stop()
+
+     blocks := conBc.GetBlocksFromHash(conBc.CurrentBlock().Hash(), int(conBc.CurrentBlock().NumberU64()))
+     for j := 0; j < len(blocks)/2; j++ {
+         blocks[j], blocks[len(blocks)-1-j] = blocks[len(blocks)-1-j], blocks[j]
+     }
+     if _, err := bc.InsertChain(blocks); err != nil {
+         t.Fatalf("failed to import contra-fork chain for expansion: %v", err)
+     }
+     if err := bc.stateCache.TrieDB().Commit(bc.CurrentHeader().Root, true, nil); err != nil {
+         t.Fatalf("failed to commit contra-fork head for expansion: %v", err)
+     }
+     blocks, _, _ = GenerateChain(&proConf, conBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+     if _, err := conBc.InsertChain(blocks); err != nil {
+         t.Fatalf("contra-fork chain didn't accept pro-fork block post-fork: %v", err)
+     }
+
+     // Verify that pro-forkers accept contra-fork extra-datas after forking finishes
+     db = rawdb.NewMemoryDatabase()
+     gspec.MustCommit(db)
+     bc, _ = NewBlockChain(db, DefaultCacheConfig, &proConf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+     defer bc.Stop()
+
+     blocks = proBc.GetBlocksFromHash(proBc.CurrentBlock().Hash(), int(proBc.CurrentBlock().NumberU64()))
+     for j := 0; j < len(blocks)/2; j++ {
+         blocks[j], blocks[len(blocks)-1-j] = blocks[len(blocks)-1-j], blocks[j]
+     }
+     if _, err := bc.InsertChain(blocks); err != nil {
+         t.Fatalf("failed to import pro-fork chain for expansion: %v", err)
+     }
+     if err := bc.stateCache.TrieDB().Commit(bc.CurrentHeader().Root, true, nil); err != nil {
+         t.Fatalf("failed to commit pro-fork head for expansion: %v", err)
+     }
+     blocks, _, _ = GenerateChain(&conConf, proBc.CurrentBlock(), dummy.NewFaker(), db, 1, 10, func(i int, gen *BlockGen) {})
+     if _, err := proBc.InsertChain(blocks); err != nil {
+         t.Fatalf("pro-fork chain didn't accept contra-fork block post-fork: %v", err)
+     }
+ }
+
+func TestDAOForkSupportPostApricotPhase3(t *testing.T) {
+     forkBlock := big.NewInt(0)
+
+     conf := *params.TestChainConfig
+     conf.DAOForkSupport = true
+     conf.DAOForkBlock = forkBlock

```

```

+ db := rawdb.NewMemoryDatabase()
+ gspec := &Genesis{
+     BaseFee: big.NewInt(params.ApricotPhase3InitialBaseFee),
+     Config: &conf,
+ }
+ genesis := gspec.MustCommit(db)
+ bc, _ := NewBlockChain(db, DefaultCacheConfig, &conf, dummy.NewFaker(), vm.Config{}, common.Hash{})
+ defer bc.Stop()
+
+ blocks, _, _ := GenerateChain(&conf, genesis, dummy.NewFaker(), db, 32, 10, func(i int, gen *BlockGen) {})
+
+ if _, err := bc.InsertChain(blocks); err != nil {
+     t.Fatalf("failed to import blocks: %v", err)
+ }
+}
diff --git a/core/error.go b/core/error.go
index 02654489..a938eef8 100644
--- a/core/error.go
+++ b/core/error.go
@@ -29,7 +29,7 @@ package core
import (
    "errors"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/types"
)

var (
@@ -56,6 +56,10 @@ var (
@@ -56,6 +56,10 @@ var (
    // next one expected based on the local chain.
    ErrNonceTooHigh = errors.New("nonce too high")

+    // ErrNonceMax is returned if the nonce of a transaction sender account has
+    // maximum allowed value and would become invalid if incremented.
+    ErrNonceMax = errors.New("nonce has max value")
+
+    // ErrGasLimitReached is returned by the gas pool if the amount of gas required
+    // by a transaction is higher than what's left in the block.
+    ErrGasLimitReached = errors.New("gas limit reached")
diff --git a/core/events.go b/core/events.go
index 4898dbc0..f71c2d40 100644
--- a/core/events.go
+++ b/core/events.go
@@ -27,8 +27,8 @@ package core

import (
-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

// NewTxEvent is posted when a batch of transactions enter the transaction pool.
diff --git a/core/evm.go b/core/evm.go
index d45f2413..19a85573 100644
--- a/core/evm.go
+++ b/core/evm.go
@@ -29,10 +29,10 @@ package core
import (
    "math/big"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    //"github.com/ethereum/go-ethereum/log"
)

@@ -125,15 +125,8 @@ func CanTransfer(db vm.StateDB, addr common.Address, amount *big.Int) bool {
    return db.GetBalance(addr).Cmp(amount) >= 0
}

-func CanTransferMC(db vm.StateDB, addr common.Address, to common.Address, coinID *common.Hash, amount *big.Int) bool {
-    if coinID == nil {
-        return true
-    }
-    if db.GetBalanceMultiCoin(addr, *coinID).Cmp(amount) >= 0 {
-        return true
-    }
-    // insufficient balance
-    return false
+func CanTransferMC(db vm.StateDB, addr common.Address, to common.Address, coinID common.Hash, amount *big.Int) bool {
+    return db.GetBalanceMultiCoin(addr, coinID).Cmp(amount) >= 0
+}

// Transfer subtracts amount from sender and adds amount to recipient using the given Db
@@ -143,10 +143,7 @@ func Transfer(db vm.StateDB, sender, recipient common.Address, amount *big.Int)
}

// Transfer subtracts amount from sender and adds amount to recipient using the given Db
-func TransferMultiCoin(db vm.StateDB, sender, recipient common.Address, coinID *common.Hash, amount *big.Int) {
-    if coinID == nil {
-        return
-    }
-    db.SubBalanceMultiCoin(sender, *coinID, amount)
-    db.AddBalanceMultiCoin(recipient, *coinID, amount)
+func TransferMultiCoin(db vm.StateDB, sender, recipient common.Address, coinID common.Hash, amount *big.Int) {
+    db.SubBalanceMultiCoin(sender, coinID, amount)
+    db.AddBalanceMultiCoin(recipient, coinID, amount)
+}
diff --git a/core/gen_genesis.go b/core/gen_genesis.go
index a4ec8f54..bc942f96 100644
--- a/core/gen_genesis.go
+++ b/core/gen_genesis.go
@@ -7,7 +7,7 @@ package core
import (
    "errors"
    "math/big"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/common/math"
diff --git a/core/genesis.go b/core/genesis.go
index 7e10249d..ce6b02dd 100644
--- a/core/genesis.go
+++ b/core/genesis.go
@@ -34,16 +34,16 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/ethereum/go-ethereum/log"

```

```

+ "github.com/flare-foundation/coreth/core/rawdb"
+ "github.com/flare-foundation/coreth/core/state"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/ethdb"
+ "github.com/flare-foundation/coreth/params"
+ "github.com/flare-foundation/coreth/trie"
+ )
+
+ //go:generate gencodec -type Genesis -field-override genesisSpecMarshaling -out gen_genesis.go
@@ -209,12 +209,14 @@ func SetupGenesisBlock(db ethdb.Database, genesis *Genesis) (*params.ChainConfig
+
+ // Check config compatibility and write the config. Compatibility errors
+ // are returned to the caller unless we're already at block zero.
- height := rawdb.ReadHeaderNumber(db, rawdb.ReadHeadHeaderHash(db))
- if height == nil {
-     return newcfg, fmt.Errorf("missing block number for head header hash")
+ headBlock := rawdb.ReadHeadBlock(db)
+ if headBlock == nil {
+     return newcfg, fmt.Errorf("missing head block")
+ }
- compatErr := storedcfg.CheckCompatible(newcfg, *height)
- if compatErr != nil && *height != 0 && compatErr.RewindTo != 0 {
+ height := headBlock.NumberU64()
+ timestamp := headBlock.Time()
+ compatErr := storedcfg.CheckCompatible(newcfg, height, timestamp)
+ if compatErr != nil && height != 0 && compatErr.RewindTo != 0 {
+     return newcfg, compatErr
+ }
+ rawdb.WriteChainConfig(db, stored, newcfg)
@@ -292,12 +294,10 @@ func (g *Genesis) Commit(db ethdb.Database) (*types.Block, error) {
+ if err := config.CheckConfigForkOrder(); err != nil {
+     return nil, err
+ }
- rawdb.WriteTd(db, block.Hash(), block.NumberU64(), g.Difficulty)
- rawdb.WriteBlock(db, block)
- rawdb.WriteReceipts(db, block.Hash(), block.NumberU64(), nil)
- rawdb.WriteCanonicalHash(db, block.Hash(), block.NumberU64())
- rawdb.WriteHeadBlockHash(db, block.Hash())
- rawdb.WriteHeadFastBlockHash(db, block.Hash())
- rawdb.WriteHeadHeaderHash(db, block.Hash())
- rawdb.WriteChainConfig(db, block.Hash(), config)
- return block, nil
@@ -316,6 +316,7 @@ func (g *Genesis) MustCommit(db ethdb.Database) *types.Block {
+ // GenesisBlockForTesting creates and writes a block in which addr has the given wei balance.
+ func GenesisBlockForTesting(db ethdb.Database, addr common.Address, balance *big.Int) *types.Block {
+     g := Genesis{
+         Config:    params.TestChainConfig,
+         Alloc:     GenesisAlloc{addr: {Balance: balance}},
+         BaseFee:   big.NewInt(params.ApricotPhase3InitialBaseFee),
+     }
+ }
diff --git a/core/genesis_test.go b/core/genesis_test.go
new file mode 100644
index 00000000..533bed18
--- /dev/null
+++ b/core/genesis_test.go
@@ -0,0 +1,167 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "embed"
+    "math/big"
+    "reflect"
+    "testing"
+
+    "github.com/davecgh/go-spew/spew"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
+)
+
+func setupGenesisBlock(db ethdb.Database, genesis *Genesis) (*params.ChainConfig, common.Hash, error) {
+    conf, err := SetupGenesisBlock(db, genesis)
+    stored := rawdb.ReadCanonicalHash(db, 0)
+    return conf, stored, err
+}
+
+func TestGenesisBlockForTesting(t *testing.T) {
+    genesisBlockForTestingHash := common.HexToHash("0xb378f22ccd9ad52c6c42f5d46ef2aad6d6866cfc778ea97a0b6dfde13387330")
+    block := GenesisBlockForTesting(rawdb.NewMemoryDatabase(), common.Address{1}, big.NewInt(1))
+    if block.Hash() != genesisBlockForTestingHash {
+        t.Errorf("wrong testing genesis hash, got %v, want %v", block.Hash(), genesisBlockForTestingHash)
+    }
+}
+
+func TestSetupGenesis(t *testing.T) {
+    apricotPhase1Config := *params.TestApricotPhase1Config
+    apricotPhase1Config.ApricotPhase1BlockTimestamp = big.NewInt(100)
+    var (
+        customghash = common.HexToHash("0x1099a11e9e454bd3ef31d688cf21936671966407bc330f051d754b5ce401e7ed")
+        customgm    = Genesis{
+            Config: &apricotPhase1Config,
+            Alloc:  GenesisAlloc{
+                {1}: {Balance: big.NewInt(1), Storage: map[common.Hash]common.Hash{{1}: {1}}},
+            },
+        }
+        oldcustomgm = customgm
+    )
+
+    rollbackApricotPhase1Config := apricotPhase1Config
+    rollbackApricotPhase1Config.ApricotPhase1BlockTimestamp = big.NewInt(90)
+    oldcustomgm.Config = &rollbackApricotPhase1Config
+    tests := []struct {
+        name      string
+        fn        func(ethdb.Database) (*params.ChainConfig, common.Hash, error)
+        wantConfig *params.ChainConfig
+    }

```

```

+ wantHash    common.Hash
+ wantErr     error
+
+ }{
+
+     {
+         name: "genesis without ChainConfig",
+         fn: func(db ethdb.Database) (*params.ChainConfig, common.Hash, error) {
+             return setupGenesisBlock(db, new(Genesis))
+         },
+         wantErr:    errGenesisNoConfig,
+         wantConfig: nil,
+     },
+     {
+         name: "no block in DB, genesis == nil",
+         fn: func(db ethdb.Database) (*params.ChainConfig, common.Hash, error) {
+             return setupGenesisBlock(db, nil)
+         },
+         wantErr:    ErrNoGenesis,
+         wantConfig: nil,
+     },
+     {
+         name: "custom block in DB, genesis == nil",
+         fn: func(db ethdb.Database) (*params.ChainConfig, common.Hash, error) {
+             customg.MustCommit(db)
+             return setupGenesisBlock(db, nil)
+         },
+         wantErr:    ErrNoGenesis,
+         wantHash:    customghash,
+         wantConfig: nil,
+     },
+     {
+         name: "compatible config in DB",
+         fn: func(db ethdb.Database) (*params.ChainConfig, common.Hash, error) {
+             oldcustomg.MustCommit(db)
+             return setupGenesisBlock(db, &customg)
+         },
+         wantHash:    customghash,
+         wantConfig: customg.Config,
+     },
+     {
+         name: "incompatible config for avalanche fork in DB",
+         fn: func(db ethdb.Database) (*params.ChainConfig, common.Hash, error) {
+             // Commit the 'old' genesis block with ApricotPhase1 transition at 90.
+             // Advance to block #4, past the ApricotPhase1 transition block of customg.
+             genesis := oldcustomg.MustCommit(db)
+
+             bc, _ := NewBlockChain(db, DefaultCacheConfig, oldcustomg.Config, dummy.NewFullFaker(), vm.Config{}, common.Hash{})
+             defer bc.Stop()
+
+             blocks, _, _ := GenerateChain(oldcustomg.Config, genesis, dummy.NewFullFaker(), db, 4, 25, nil)
+             bc.InsertChain(blocks)
+             bc.CurrentBlock()
+             // This should return a compatibility error.
+             return setupGenesisBlock(db, &customg)
+         },
+         wantHash:    customghash,
+         wantConfig: customg.Config,
+         wantErr:    &params.ConfigCompatError{
+             What:         "ApricotPhase1 fork block timestamp",
+             StoredConfig: big.NewInt(90),
+             NewConfig:      big.NewInt(100),
+             RewindTo:        89,
+         },
+     },
+ }
+
+ for _, test := range tests {
+     t.Run(test.name, func(t *testing.T) {
+         db := rawdb.NewMemoryDatabase()
+         config, hash, err := test.fn(db)
+         // Check the return values.
+         if !reflect.DeepEqual(err, test.wantErr) {
+             spew := spew.ConfigState{DisablePointerAddresses: true, DisableCapacities: true}
+             t.Errorf("returned error %#v, want %#v", spew.NewFormatter(err), spew.NewFormatter(test.wantErr))
+         }
+         if !reflect.DeepEqual(config, test.wantConfig) {
+             t.Errorf("returned %v\nwant      %v", config, test.wantConfig)
+         }
+         if hash != test.wantHash {
+             t.Errorf("returned hash %s, want %s", hash.Hex(), test.wantHash.Hex())
+         } else if err == nil {
+             // Check database content.
+             stored := rawdb.ReadBlock(db, test.wantHash, 0)
+             if stored.Hash() != test.wantHash {
+                 t.Errorf("block in DB has hash %s, want %s", stored.Hash(), test.wantHash)
+             }
+         }
+     })
+ }
+
+}
+
+diff --git a/core/headerchain.go b/core/headerchain.go
+index 092350a0..48b81e19 100644
+--- a/core/headerchain.go
++++ b/core/headerchain.go
+@@ -33,12 +33,12 @@ import (
+     mrand "math/rand"
+     "sync/atomic"
+
+     "github.com/ava-labs/coreth/consensus"
+     "github.com/ava-labs/coreth/core/rawdb"
+     "github.com/ava-labs/coreth/core/types"
+     "github.com/ava-labs/coreth/ethdb"
+     "github.com/ava-labs/coreth/params"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/coreth/consensus"
+     "github.com/flare-foundation/coreth/core/rawdb"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/ethdb"
+     "github.com/flare-foundation/coreth/params"
+     lru "github.com/hashicorp/golang-lru"
+
+ )
+
+@@ -55,9 +55,9 @@ const (
+     // HeaderChain is responsible for maintaining the header chain including the
+     // header query and updating.
+     //
+     -// The components maintained by headerchain includes: (1) total difficult
+     -// (2) header (3) block hash -> number mapping (4) canonical number -> hash mapping
+     -// and (5) head header flag.
+     +// The components maintained by headerchain includes:
+     +// (1) header (2) block hash -> number mapping (3) canonical number -> hash mapping
+     +// and (4) head header flag.
+     //
+     // It is not thread safe either, the encapsulating chain structures should do
+     // the necessary mutex locking/unlocking.
+@@ -131,22 +131,6 @@ func (hc *HeaderChain) GetBlockNumber(hash common.Hash) *uint64 {
+     return number
+
+ }
+
+ -// GetTd retrieves a block's total difficulty in the canonical chain from the
+ -// database by hash and number, caching it if found.
+ -func (hc *HeaderChain) GetTd(hash common.Hash, number uint64) *big.Int {
+ -    // Short circuit if the td's already in the cache, retrieve otherwise
+ -    if cached, ok := hc.tdCache.Get(hash); ok {
+ -        return cached.(*big.Int)
+     }

```

```

-     }
-     td := rawdb.ReadTd(hc.chainDb, hash, number)
-     if td == nil {
-         return nil
-     }
-     // Cache the found body for next time and return
-     hc.tdCache.Add(hash, td)
-     return td
-}
-
// GetHeader retrieves a block header from the database by hash and number,
// caching it if found.
func (hc *HeaderChain) GetHeader(hash common.Hash, number uint64) *types.Header {
diff --git a/core/headerchain_test.go b/core/headerchain_test.go
new file mode 100644
index 00000000..552dcc38
--- /dev/null
+++ b/core/headerchain_test.go
@@ -0,0 +1,123 @@
// (c) 2019-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
// *****
// Copyright 2020 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "errors"
+    "fmt"
+    "math/big"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
+)
+
+func verifyUnbrokenCanonchain(bc *BlockChain) error {
+    h := bc.hc.CurrentHeader()
+    for {
+        canonHash := rawdb.ReadCanonicalHash(bc.hc.chainDb, h.Number.Uint64())
+        if exp := h.Hash(); canonHash != exp {
+            return fmt.Errorf("Canon hash chain broken, block %d got %x, expected %x",
+                h.Number, canonHash[:8], exp[:8])
+        }
+        if h.Number.Uint64() == 0 {
+            break
+        }
+        h = bc.hc.GetHeader(h.ParentHash, h.Number.Uint64()-1)
+    }
+    return nil
+}
+
+func testInsert(t *testing.T, bc *BlockChain, chain []*types.Block, wantErr error) {
+    t.Helper()
+
+    _, err := bc.InsertChain(chain)
+    // Always verify that the header chain is unbroken
+    if err := verifyUnbrokenCanonchain(bc); err != nil {
+        t.Fatal(err)
+    }
+    if !errors.Is(err, wantErr) {
+        t.Fatalf("unexpected error from InsertHeaderChain: %v", err)
+    }
+}
+
+// This test checks status reporting of InsertHeaderChain.
+func TestHeaderInsertion(t *testing.T) {
+    var (
+        db      = rawdb.NewMemoryDatabase()
+        genesis = (&Genesis{
+            BaseFee: big.NewInt(params.ApricotPhase3InitialBaseFee),
+            Config:  params.TestChainConfig,
+        }).MustCommit(db)
+    )
+    chain, err := NewBlockChain(db, DefaultCacheConfig, params.TestChainConfig, dummy.NewFaker(), vm.Config{}, common.Hash{})
+    if err != nil {
+        t.Fatal(err)
+    }
+    // chain A: G->A1->A2...A128
+    chainA, _, _ := GenerateChain(params.TestChainConfig, types.NewBlockWithHeader(genesis.Header()), dummy.NewFaker(), db, 128, 10, func(i int, b *BlockGen) {
+        b.SetCoinbase(common.Address{0: byte(10), 19: byte(i)})
+    })
+    // chain B: G->A1->B2...B128
+    chainB, _, _ := GenerateChain(params.TestChainConfig, types.NewBlockWithHeader(chainA[0].Header()), dummy.NewFaker(), db, 128, 10, func(i int, b *BlockGen) {
+        b.SetCoinbase(common.Address{0: byte(10), 19: byte(i)})
+    })
+    log.Root().SetHandler(log.StdoutHandler)
+
+    // Inserting 64 headers on an empty chain
+    testInsert(t, chain, chainA[:64], nil)
+
+    // Inserting 64 identical headers
+    testInsert(t, chain, chainA[:64], nil)
+
+    // Inserting the same some old, some new headers
+    testInsert(t, chain, chainA[32:96], nil)
+
+    // Inserting side blocks, but not overtaking the canon chain
+    testInsert(t, chain, chainB[0:32], nil)
+
+    // Inserting more side blocks, but we don't have the parent
+    testInsert(t, chain, chainB[34:36], consensus.ErrUnknownAncestor)
+
+    // Inserting more sideblocks, overtaking the canon chain
+    testInsert(t, chain, chainB[32:97], nil)
+
+    // Inserting more A-headers, taking back the canonicity

```

```

+         testInsert(t, chain, chainA[90:100], nil)
+
+         // And B becomes canon again
+         testInsert(t, chain, chainB[97:107], nil)
+
+         // And B becomes even longer
+         testInsert(t, chain, chainB[107:128], nil)
+    }
+}
diff --git a/core/keeper.go b/core/keeper.go
new file mode 100644
index 00000000..a011f45c
--- /dev/null
+++ b/core/keeper.go
@@ -0,0 +1,149 @@
+// (c) 2021, Flare Networks Limited. All rights reserved.
+// Please see the file LICENSE for licensing terms.
+
+package core
+
+import (
+    "fmt"
+    "math/big"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+
+    "github.com/flare-foundation/coreth/core/vm"
+)
+
+// Define errors
+type ErrInvalidKeeperData struct{}
+
+func (e *ErrInvalidKeeperData) Error() string { return "invalid return data from keeper trigger" }
+
+type ErrKeeperDataEmpty struct{}
+
+func (e *ErrKeeperDataEmpty) Error() string { return "return data from keeper trigger empty" }
+
+type ErrMaxMintExceeded struct {
+    mintMax      *big.Int
+    mintRequest  *big.Int
+}
+
+func (e *ErrMaxMintExceeded) Error() string {
+    return fmt.Sprintf("mint request of %s exceeded max of %s", e.mintRequest.Text(10), e.mintMax.Text(10))
+}
+
+type ErrMintNegative struct{}
+
+func (e *ErrMintNegative) Error() string { return "mint request cannot be negative" }
+
+// Define interface for dependencies
+type EVMCaller interface {
+    Call(caller vm.ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error)
+    GetBlockNumber() *big.Int
+    GetGasLimit() uint64
+    AddBalance(addr common.Address, amount *big.Int)
+}
+
+// Define maximums that can change by block height
+func GetKeeperGasMultiplier(blockNumber *big.Int) uint64 {
+    switch {
+    default:
+        return 100
+    }
+}
+
+func GetSystemTriggerContractAddr(blockNumber *big.Int) string {
+    switch {
+    default:
+        return "0x1000000000000000000000000000000000000000000000000000000000000002"
+    }
+}
+
+func GetSystemTriggerSelector(blockNumber *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0x7f, 0xec, 0x8d, 0x38}
+    }
+}
+
+func GetPrioritisedFTS0Contract(blockTime *big.Int) string {
+    switch {
+    default:
+        return "0x1000000000000000000000000000000000000000000000000000000000000003"
+    }
+}
+
+func GetMaximumMintRequest(blockNumber *big.Int) *big.Int {
+    switch {
+    default:
+        maxRequest, _ := new(big.Int).SetString("50000000000000000000000000000000", 10)
+        return maxRequest
+    }
+}
+
+func triggerKeeper(evm EVMCaller) (*big.Int, error) {
+    bigZero := big.NewInt(0)
+    // Get the contract to call
+    systemTriggerContract := common.HexToAddress(GetSystemTriggerContractAddr(evm.GetBlockNumber()))
+    // Call the method
+    triggerRet, _, triggerErr := evm.Call(
+        vm.AccountRef(systemTriggerContract),
+        systemTriggerContract,
+        GetSystemTriggerSelector(evm.GetBlockNumber()),
+        GetKeeperGasMultiplier(evm.GetBlockNumber())*evm.GetGasLimit(),
+        bigZero)
+    // If no error and a value came back...
+    if triggerErr == nil && triggerRet != nil {
+        // Did we get one big int?
+        if len(triggerRet) == 32 {
+            // Convert to big int
+            // Mint request cannot be less than 0 as SetBytes treats value as unsigned
+            mintRequest := new(big.Int).SetBytes(triggerRet)
+            // return the mint request
+            return mintRequest, nil
+        } else {
+            // Returned length was not 32 bytes
+            return bigZero, &ErrInvalidKeeperData{}
+        }
+    } else {
+        if triggerErr != nil {
+            return bigZero, triggerErr
+        } else {
+            return bigZero, &ErrKeeperDataEmpty{}
+        }
+    }
+}
+
+func mint(evm EVMCaller, mintRequest *big.Int) error {
+    // If the mint request is greater than zero and less than max
+    max := GetMaximumMintRequest(evm.GetBlockNumber())
+    if mintRequest.Cmp(big.NewInt(0)) > 0 &&
+        mintRequest.Cmp(max) <= 0 {

```

[illegible]

```

+     mockEVMCallerData := &MockEVMCallerData{
+         blockNumber:      *big.NewInt(0),
+         gasLimit:         0,
+         mintRequestReturn: mintRequestReturn,
+     }
+     defaultEVMMock := &DefaultEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+
+     mintRequest, mintRequestError := triggerKeeper(defaultEVMMock)
+
+     if mintRequestError != nil {
+         t.Errorf("received unexpected error %s", mintRequestError)
+     }
+
+     if mintRequest.Sign() < 1 {
+         t.Errorf("unexpected negative")
+     }
+ }
+
+ // Define a bad mint request return size mock
+ type BadMintReturnSizeEVMMock struct {
+     mockEVMCallerData MockEVMCallerData
+ }
+
+ func (e *BadMintReturnSizeEVMMock) Call(caller vm.ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {
+     e.mockEVMCallerData.callCalls++
+     // Should be size 32 bytes
+     buffer := []byte{0}
+     return e.mockEVMCallerData.mintRequestReturn.FillBytes(buffer), 0, nil
+ }
+
+ func (e *BadMintReturnSizeEVMMock) GetBlockNumber() *big.Int {
+     return defaultGetBlockNumber(&e.mockEVMCallerData)
+ }
+
+ func (e *BadMintReturnSizeEVMMock) GetGasLimit() uint64 {
+     return defaultGetGasLimit(&e.mockEVMCallerData)
+ }
+
+ func (e *BadMintReturnSizeEVMMock) AddBalance(addr common.Address, amount *big.Int) {
+     defaultAddBalance(&e.mockEVMCallerData, addr, amount)
+ }
+
+ func TestKeeperTriggerValidatesMintRequestReturnValueSize(t *testing.T) {
+     var mintRequestReturn big.Int
+     // TODO: Compact with exponent?
+     buffer := []byte{255}
+     mintRequestReturn.SetBytes(buffer)
+
+     mockEVMCallerData := &MockEVMCallerData{
+         blockNumber:      *big.NewInt(0),
+         gasLimit:         0,
+         mintRequestReturn: mintRequestReturn,
+     }
+     badMintReturnSizeEVMMock := &BadMintReturnSizeEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+     // Call to return less than 32 bytes
+     _, err := triggerKeeper(badMintReturnSizeEVMMock)
+
+     if err != nil {
+         if err, ok := err.(*ErrInvalidKeeperData); !ok {
+             want := &ErrInvalidKeeperData{}
+             t.Errorf("got '%s' want '%s'", err.Error(), want.Error())
+         }
+     } else {
+         t.Errorf("no error returned as expected")
+     }
+ }
+
+ // Define a mock to simulate keeper trigger returning an error from Call
+ type BadTriggerCallEVMMock struct {
+     mockEVMCallerData MockEVMCallerData
+ }
+
+ func (e *BadTriggerCallEVMMock) Call(caller vm.ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {
+     e.mockEVMCallerData.callCalls++
+
+     buffer := []byte{0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0}
+     return e.mockEVMCallerData.mintRequestReturn.FillBytes(buffer), 0, errors.New("Call error happened")
+ }
+
+ func (e *BadTriggerCallEVMMock) GetBlockNumber() *big.Int {
+     return defaultGetBlockNumber(&e.mockEVMCallerData)
+ }
+
+ func (e *BadTriggerCallEVMMock) GetGasLimit() uint64 {
+     return defaultGetGasLimit(&e.mockEVMCallerData)
+ }
+
+ func (e *BadTriggerCallEVMMock) AddBalance(addr common.Address, amount *big.Int) {
+     defaultAddBalance(&e.mockEVMCallerData, addr, amount)
+ }
+
+ func TestKeeperTriggerReturnsCallError(t *testing.T) {
+     mockEVMCallerData := &MockEVMCallerData{}
+     badTriggerCallEVMMock := &BadTriggerCallEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+     // Call to return less than 32 bytes
+     _, err := triggerKeeper(badTriggerCallEVMMock)
+
+     if err == nil {
+         t.Errorf("no error received")
+     } else {
+         if err.Error() != "Call error happened" {
+             t.Errorf("did not get expected error")
+         }
+     }
+ }
+
+ type LoggerMock struct {
+     mockLoggerData MockLoggerData
+ }
+
+ func (l *LoggerMock) New(ctx ...interface{}) log.Logger {
+     return nil
+ }
+
+ func (l *LoggerMock) GetHandler() log.Handler {
+     return nil
+ }
+
+ func (l *LoggerMock) SetHandler(h log.Handler) {
+ }
+
+ func (l *LoggerMock) Trace(msg string, ctx ...interface{}) {}
+ func (l *LoggerMock) Debug(msg string, ctx ...interface{}) {}
+ func (l *LoggerMock) Info(msg string, ctx ...interface{}) {}
+ func (l *LoggerMock) Error(msg string, ctx ...interface{}) {}
+ func (l *LoggerMock) Crit(msg string, ctx ...interface{}) {}
+
+ func (l *LoggerMock) Warn(msg string, ctx ...interface{}) {}

```



```
l.mockLoggerData.warnCalls++  
+}  
+func TestKeeperTriggerAndMintLogsError(t *testing.T) {  
+    // Assemble  
+    // Set up mock EVM call to return an error  
+    mockEVMCallerData := &MockEVMCallerData{  
+        badTriggerCallEVMMock := &BadTriggerCallEVMMock{  
+            mockEVMCallerData: *mockEVMCallerData,  
+        }  
+    }  
+    // Set up a mock logger  
+    mockLoggerData := &MockLoggerData{}  
+    loggerMock := &LoggerMock{  
+        mockLoggerData: *mockLoggerData,  
+    }  
+  
+    // Act  
+    triggerKeeperAndMint(badTriggerCallEVMMock, loggerMock)  
+  
+    // Assert  
+    if loggerMock.mockLoggerData.warnCalls != 1 {  
+        t.Errorf("Logger.Warn not called as expected")  
+    }  
+}  
+// Define a mock to simulate keeper trigger returning nil for mint request  
+type ReturnNilMintRequestEVMMock struct {  
+    mockEVMCallerData MockEVMCallerData  
+}  
+func (e *ReturnNilMintRequestEVMMock) Call(caller vm.ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {  
+    e.mockEVMCallerData.callCalls++  
  
+    return nil, 0, nil  
+}  
+func (e *ReturnNilMintRequestEVMMock) GetBlockNumber() *big.Int {  
+    return defaultGetBlockNumber(&e.mockEVMCallerData)  
+}  
+func (e *ReturnNilMintRequestEVMMock) GetGasLimit() uint64 {  
+    return defaultGetGasLimit(&e.mockEVMCallerData)  
+}  
+func (e *ReturnNilMintRequestEVMMock) AddBalance(addr common.Address, amount *big.Int) {  
+    defaultAddBalance(&e.mockEVMCallerData, addr, amount)  
+}  
+func TestKeeperTriggerHandlesNilMintRequest(t *testing.T) {  
+    mockEVMCallerData := &MockEVMCallerData{  
+        returnNilMintRequestEVMMock := &ReturnNilMintRequestEVMMock{  
+            mockEVMCallerData: *mockEVMCallerData,  
+        }  
+    }  
+    // Call to return less than 32 bytes  
+    _, err := triggerKeeper(returnNilMintRequestEVMMock)  
  
+    if err != nil {  
+        if err, ok := err.(*ErrKeeperDataEmpty); !ok {  
+            want := &ErrKeeperDataEmpty{}  
+            t.Errorf("got '%s' want '%s'", err.Error(), want.Error())  
+        }  
+    } else {  
+        t.Errorf("no error returned as expected")  
+    }  
+}  
+func TestKeeperTriggerShouldNotMintMoreThanMax(t *testing.T) {  
+    mintRequest, _ := new(big.Int).SetString("500000000000000000000000000001", 10)  
+    mockEVMCallerData := &MockEVMCallerData{  
+        blockNumber:      *big.NewInt(0),  
+        gasLimit:          0,  
+        mintRequestReturn: *big.NewInt(0),  
+    }  
+    defaultEVMMock := &DefaultEVMMock{  
+        mockEVMCallerData: *mockEVMCallerData,  
+    }  
  
+    err := mint(defaultEVMMock, mintRequest)  
  
+    if err != nil {  
+        if err, ok := err.(*ErrMsgMintExceeded); !ok {  
+            want := &ErrMsgMintExceeded{  
+                mintRequest: mintRequest,  
+                mintMax:     GetMaximumMintRequest(big.NewInt(0)),  
+            }  
+            t.Errorf("got '%s' want '%s'", err.Error(), want.Error())  
+        }  
+    } else {  
+        t.Errorf("no error returned as expected")  
+    }  
+}  
+func TestKeeperTriggerShouldNotMintNegative(t *testing.T) {  
+    mintRequest := big.NewInt(-1)  
+    mockEVMCallerData := &MockEVMCallerData{  
+        blockNumber:      *big.NewInt(0),  
+        gasLimit:          0,  
+        mintRequestReturn: *big.NewInt(0),  
+    }  
+    defaultEVMMock := &DefaultEVMMock{  
+        mockEVMCallerData: *mockEVMCallerData,  
+    }  
  
+    err := mint(defaultEVMMock, mintRequest)  
  
+    if err != nil {  
+        if err, ok := err.(*ErrMintNegative); !ok {  
+            want := &ErrMintNegative{}  
+            t.Errorf("got '%s' want '%s'", err.Error(), want.Error())  
+        }  
+    } else {  
+        t.Errorf("no error returned as expected")  
+    }  
+}  
+func TestKeeperTriggerShouldMint(t *testing.T) {  
+    // Assemble  
+    mintRequest, _ := new(big.Int).SetString("500000000000000000000000000000", 10)  
+    mockEVMCallerData := &MockEVMCallerData{  
+        blockNumber:      *big.NewInt(0),  
+        gasLimit:          0,  
+        mintRequestReturn: *big.NewInt(0),  
+    }  
+    defaultEVMMock := &DefaultEVMMock{  
+        mockEVMCallerData: *mockEVMCallerData,  
+    }  
  
+    // Act  
+    err := mint(defaultEVMMock, mintRequest)  
  
+    // Assert  
+    if err == nil {  
+        if defaultEVMMock.mockEVMCallerData.addBalanceCalls != 1 {  
+            t.Errorf("AddBalance not called as expected")
```

```

+         }
+         if defaultEVMMock.mockEVMCallerData.lastAddBalanceAddr.String() != GetSystemTriggerContractAddr(big.NewInt(0)) {
+             t.Errorf("wanted addr %s; got addr %s", GetSystemTriggerContractAddr(big.NewInt(0)), defaultEVMMock.mockEVMCallerData.lastAddBalanceAddr)
+         }
+         if defaultEVMMock.mockEVMCallerData.lastAddBalanceAmount.Cmp(mintRequest) != 0 {
+             t.Errorf("wanted amount %s; got amount %s", mintRequest.Text(10), defaultEVMMock.mockEVMCallerData.lastAddBalanceAmount.Text(10))
+         }
+     } else {
+         t.Errorf("unexpected error returned; was = %s", err.Error())
+     }
+ }
+
+func TestKeeperTriggerShouldNotErrorMintingZero(t *testing.T) {
+     // Assemble
+     mintRequest := big.NewInt(0)
+     mockEVMCallerData := &MockEVMCallerData{
+         blockNumber:    *big.NewInt(0),
+         gasLimit:        0,
+         mintRequestReturn: *big.NewInt(0),
+     }
+     defaultEVMMock := &DefaultEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+
+     // Act
+     err := mint(defaultEVMMock, mintRequest)
+
+     // Assert
+     if err == nil {
+         if defaultEVMMock.mockEVMCallerData.addBalanceCalls != 0 {
+             t.Errorf("AddBalance called unexpectedly")
+         }
+     } else {
+         t.Errorf("unexpected error returned; was %s", err.Error())
+     }
+ }
+
+func TestKeeperTriggerFiredAndMinted(t *testing.T) {
+     mintRequestReturn, _ := new(big.Int).SetString("5000000000000000000000", 10)
+     mockEVMCallerData := &MockEVMCallerData{
+         blockNumber:    *big.NewInt(0),
+         gasLimit:        0,
+         mintRequestReturn: *mintRequestReturn,
+     }
+     defaultEVMMock := &DefaultEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+
+     log := log.New()
+     triggerKeeperAndMint(defaultEVMMock, log)
+
+     // EVM Call function calling the keeper should have been cqlled
+     if defaultEVMMock.mockEVMCallerData.callCalls != 1 {
+         t.Errorf("EVM Call count not as expected. got %d want 1", defaultEVMMock.mockEVMCallerData.callCalls)
+     }
+     // AddBalance should have been called on the state database, minting the request asked for
+     if defaultEVMMock.mockEVMCallerData.addBalanceCalls != 1 {
+         t.Errorf("Add balance call count not as expected. got %d want 1", defaultEVMMock.mockEVMCallerData.addBalanceCalls)
+     }
+ }
+
+func TestKeeperTriggerShouldNotMintMoreThanLimit(t *testing.T) {
+     mintRequestReturn, _ := new(big.Int).SetString("500000000000000000000001", 10)
+     mockEVMCallerData := &MockEVMCallerData{
+         blockNumber:    *big.NewInt(0),
+         gasLimit:        0,
+         mintRequestReturn: *mintRequestReturn,
+     }
+     defaultEVMMock := &DefaultEVMMock{
+         mockEVMCallerData: *mockEVMCallerData,
+     }
+
+     log := log.New()
+     triggerKeeperAndMint(defaultEVMMock, log)
+
+     // EVM Call function calling the keeper should have been called
+     if defaultEVMMock.mockEVMCallerData.callCalls != 1 {
+         t.Errorf("EVM Call count not as expected. got %d want 1", defaultEVMMock.mockEVMCallerData.callCalls)
+     }
+     // AddBalance should not have been called on the state database, as the mint request was over the limit
+     if defaultEVMMock.mockEVMCallerData.addBalanceCalls != 0 {
+         t.Errorf("Add balance call count not as expected. got %d want 1", defaultEVMMock.mockEVMCallerData.addBalanceCalls)
+     }
+ }
+
+diff --git a/core/mkalloc.go b/core/mkalloc.go
+index 76978a54..95fcafb9 100644
+--- a/core/mkalloc.go
++++ b/core/mkalloc.go
+@@ -45,8 +45,8 @@ import (
+     "sort"
+     "strconv"
+
+     "github.com/ava-labs/coreth/core"
+     "github.com/ethereum/go-ethereum/rlp"
+     "github.com/flare-foundation/coreth/core"
+ )
+
+ type allocItem struct{ Addr, Balance *big.Int }
+diff --git a/core/rawdb/accessors_chain.go b/core/rawdb/accessors_chain.go
+index 8b134e23..4bc082fc 100644
+--- a/core/rawdb/accessors_chain.go
++++ b/core/rawdb/accessors_chain.go
+@@ -30,14 +30,13 @@ import (
+     "bytes"
+     "encoding/binary"
+     "errors"
+     "math/big"
+
+     "github.com/ava-labs/coreth/core/types"
+     "github.com/ava-labs/coreth/ethdb"
+     "github.com/ava-labs/coreth/params"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/ethereum/go-ethereum/log"
+     "github.com/ethereum/go-ethereum/rlp"
+
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/ethdb"
+     "github.com/flare-foundation/coreth/params"
+ )
+
+ // ReadCanonicalHash retrieves the hash assigned to a canonical block number.
+@@ -202,77 +201,6 @@ func WriteHeadBlockHash(db ethdb.KeyValueWriter, hash common.Hash) {
+ }
+ }
+
+ // ReadHeadFastBlockHash retrieves the hash of the current fast-sync head block.
+func ReadHeadFastBlockHash(db ethdb.KeyValueReader) common.Hash {
+     data, _ := db.Get(headFastBlockKey)
+     if len(data) == 0 {
+         return common.Hash{}
+     }
+     return common.BytesToHash(data)
+ }
+
+ // WriteHeadFastBlockHash stores the hash of the current fast-sync head block.

```

```

-func WriteHeadFastBlockHash(db ethdb.KeyValueWriter, hash common.Hash) {
-    if err := db.Put(headFastBlockKey, hash.Bytes()); err != nil {
-        log.Crit("Failed to store last fast block's hash", "err", err)
-    }
-}
-
-
-// ReadFastTrieProgress retrieves the number of tries nodes fast synced to allow
-// reporting correct numbers across restarts.
-func ReadFastTrieProgress(db ethdb.KeyValueReader) uint64 {
-    data, _ := db.Get(fastTrieProgressKey)
-    if len(data) == 0 {
-        return 0
-    }
-    return new(big.Int).SetBytes(data).Uint64()
-}
-
-
-// WriteFastTrieProgress stores the fast sync trie process counter to support
-// retrieving it across restarts.
-func WriteFastTrieProgress(db ethdb.KeyValueWriter, count uint64) {
-    if err := db.Put(fastTrieProgressKey, new(big.Int).SetUint64(count).Bytes()); err != nil {
-        log.Crit("Failed to store fast sync trie progress", "err", err)
-    }
-}
-
-
-// ReadTxIndexTail retrieves the number of oldest indexed block
-// whose transaction indices has been indexed. If the corresponding entry
-// is non-existent in database it means the indexing has been finished.
-func ReadTxIndexTail(db ethdb.KeyValueReader) *uint64 {
-    data, _ := db.Get(txIndexTailKey)
-    if len(data) != 8 {
-        return nil
-    }
-    number := binary.BigEndian.Uint64(data)
-    return &number
-}
-
-
-// WriteTxIndexTail stores the number of oldest indexed block
-// into database.
-func WriteTxIndexTail(db ethdb.KeyValueWriter, number uint64) {
-    if err := db.Put(txIndexTailKey, encodeBlockNumber(number)); err != nil {
-        log.Crit("Failed to store the transaction index tail", "err", err)
-    }
-}
-
-
-// ReadFastTxLookupLimit retrieves the tx lookup limit used in fast sync.
-func ReadFastTxLookupLimit(db ethdb.KeyValueReader) *uint64 {
-    data, _ := db.Get(fastTxLookupLimitKey)
-    if len(data) != 8 {
-        return nil
-    }
-    number := binary.BigEndian.Uint64(data)
-    return &number
-}
-
-
-// WriteFastTxLookupLimit stores the txlookup limit used in fast sync into database.
-func WriteFastTxLookupLimit(db ethdb.KeyValueWriter, number uint64) {
-    if err := db.Put(fastTxLookupLimitKey, encodeBlockNumber(number)); err != nil {
-        log.Crit("Failed to store transaction lookup limit for fast sync", "err", err)
-    }
-}
-
-
-// ReadHeaderRLP retrieves a block header in its raw RLP database encoding.
-func ReadHeaderRLP(db ethdb.Reader, hash common.Hash, number uint64) rlp.RawValue {
-    // Then try to look up the data in leveldb.
@@ -408,48 +336,6 @@ func DeleteBody(db ethdb.KeyValueWriter, hash common.Hash, number uint64) {
-    }
-}
-
-
-// ReadTdRLP retrieves a block's total difficulty corresponding to the hash in RLP encoding.
-func ReadTdRLP(db ethdb.Reader, hash common.Hash, number uint64) rlp.RawValue {
-    // Then try to look up the data in leveldb.
-    data, _ := db.Get(headerTDKey(number, hash))
-    if len(data) > 0 {
-        return data
-    }
-    return nil // Can't find the data anywhere.
-}
-
-
-// ReadTd retrieves a block's total difficulty corresponding to the hash.
-func ReadTd(db ethdb.Reader, hash common.Hash, number uint64) *big.Int {
-    data := ReadTdRLP(db, hash, number)
-    if len(data) == 0 {
-        return nil
-    }
-    td := new(big.Int)
-    if err := rlp.Decode(bytes.NewReader(data), td); err != nil {
-        log.Error("Invalid block total difficulty RLP", "hash", hash, "err", err)
-        return nil
-    }
-    return td
-}
-
-
-// WriteTd stores the total difficulty of a block into the database.
-func WriteTd(db ethdb.KeyValueWriter, hash common.Hash, number uint64, td *big.Int) {
-    data, err := rlp.EncodeToBytes(td)
-    if err != nil {
-        log.Crit("Failed to RLP encode block total difficulty", "err", err)
-    }
-    if err := db.Put(headerTDKey(number, hash), data); err != nil {
-        log.Crit("Failed to store block total difficulty", "err", err)
-    }
-}
-
-
-// DeleteTd removes all block total difficulty data associated with a hash.
-func DeleteTd(db ethdb.KeyValueWriter, hash common.Hash, number uint64) {
-    if err := db.Delete(headerTDKey(number, hash)); err != nil {
-        log.Crit("Failed to delete block total difficulty", "err", err)
-    }
-}
-
-
-// HasReceipts verifies the existence of all the transaction receipts belonging
-// to a block.
-func HasReceipts(db ethdb.Reader, hash common.Hash, number uint64) bool {
@@ -653,7 +539,6 @@ func DeleteBlock(db ethdb.KeyValueWriter, hash common.Hash, number uint64) {
-    DeleteReceipts(db, hash, number)
-    DeleteHeader(db, hash, number)
-    DeleteBody(db, hash, number)
-    DeleteTd(db, hash, number)
-}
-
-
-// DeleteBlockWithoutNumber removes all block data associated with a hash, except
@@ -662,7 +547,6 @@ func DeleteBlockWithoutNumber(db ethdb.KeyValueWriter, hash common.Hash, number
-    DeleteReceipts(db, hash, number)
-    deleteHeaderWithoutNumber(db, hash, number)
-    DeleteBody(db, hash, number)
-    DeleteTd(db, hash, number)
-}
-
-
-// FindCommonAncestor returns the last common ancestor of two block headers
diff --git a/core/rawdb/accessors_chain_test.go b/core/rawdb/accessors_chain_test.go
index 9d35e7ba..4773d72f 100644
--- a/core/rawdb/accessors_chain_test.go
+++ b/core/rawdb/accessors_chain_test.go
@@ -25,10 +25,10 @@ import (

```

```

"reflect"
"testing"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/params"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/params"
+   "golang.org/x/crypto/sha3"
)

@@ -187,29 +187,6 @@ func TestPartialBlockStorage(t *testing.T) {
}

}

-// Tests block total difficulty storage and retrieval operations.
-func TestTdStorage(t *testing.T) {
-   db := NewMemoryDatabase()
-
-   // Create a test TD to move around the database and make sure it's really new
-   hash, td := common.Hash{}, big.NewInt(314)
-   if entry := ReadTd(db, hash, 0); entry != nil {
-       t.Fatalf("Non existent TD returned: %v", entry)
-   }
-   // Write and verify the TD in the database
-   WriteTd(db, hash, 0, td)
-   if entry := ReadTd(db, hash, 0); entry == nil {
-       t.Fatalf("Stored TD not found")
-   } else if entry.Cmp(td) != 0 {
-       t.Fatalf("Retrieved TD mismatch: have %v, want %v", entry, td)
-   }
-   // Delete the TD and verify the execution
-   DeleteTd(db, hash, 0)
-   if entry := ReadTd(db, hash, 0); entry != nil {
-       t.Fatalf("Deleted TD returned: %v", entry)
-   }
-}

// Tests that canonical numbers can be mapped to hashes and retrieved.
func TestCanonicalMappingStorage(t *testing.T) {
   db := NewMemoryDatabase()
@@ -239,7 +216,6 @@ func TestHeadStorage(t *testing.T) {

   blockHead := types.NewBlockWithHeader(&types.Header{Extra: []byte("test block header")})
   blockFull := types.NewBlockWithHeader(&types.Header{Extra: []byte("test block full")})
-   blockFast := types.NewBlockWithHeader(&types.Header{Extra: []byte("test block fast")})

   // Check that no head entries are in a pristine database
   if entry := ReadHeadHeaderHash(db); entry != (common.Hash{}) {
@@ -248,13 +224,9 @@ func TestHeadStorage(t *testing.T) {
@@ -248,13 +224,9 @@ func TestHeadStorage(t *testing.T) {
   if entry := ReadHeadBlockHash(db); entry != (common.Hash{}) {
       t.Fatalf("Non head block entry returned: %v", entry)
   }
-   if entry := ReadHeadFastBlockHash(db); entry != (common.Hash{}) {
-       t.Fatalf("Non fast head block entry returned: %v", entry)
-   }
-   // Assign separate entries for the head header and block
-   WriteHeadHeaderHash(db, blockHead.Hash())
-   WriteHeadBlockHash(db, blockFull.Hash())
-   WriteHeadFastBlockHash(db, blockFast.Hash())

   // Check that both heads are present, and different (i.e. two heads maintained)
   if entry := ReadHeadHeaderHash(db); entry != blockHead.Hash() {
@@ -263,9 +235,6 @@ func TestHeadStorage(t *testing.T) {
@@ -263,9 +235,6 @@ func TestHeadStorage(t *testing.T) {
   if entry := ReadHeadBlockHash(db); entry != blockFull.Hash() {
       t.Fatalf("Head block hash mismatch: have %v, want %v", entry, blockFull.Hash())
   }
-   if entry := ReadHeadFastBlockHash(db); entry != blockFast.Hash() {
-       t.Fatalf("Fast head block hash mismatch: have %v, want %v", entry, blockFast.Hash())
-   }
-}

// Tests that receipts associated with a single block can be stored and retrieved.
@@ -393,7 +362,6 @@ func TestCanonicalHashIteration(t *testing.T) {
// Fill database with testing data.
for i := uint64(1); i <= 8; i++ {
   WriteCanonicalHash(db, common.Hash{i}, i)
-   WriteTd(db, common.Hash{i}, i, big.NewInt(10)) // Write some interferential data
}
for i, c := range cases {
   numbers, _ := ReadAllCanonicalHashes(db, c.from, c.to, c.limit)
diff --git a/core/rawdb/accessors_indexes.go b/core/rawdb/accessors_indexes.go
index 0ac33214..a75dc386 100644
--- a/core/rawdb/accessors_indexes.go
+++ b/core/rawdb/accessors_indexes.go
@@ -30,12 +30,12 @@ import (
   "bytes"
   "math/big"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/ethdb"
-   "github.com/ava-labs/coreth/params"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/log"
-   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/ethdb"
+   "github.com/flare-foundation/coreth/params"
+   "golang.org/x/crypto/sha3"
)

// ReadTxLookupEntry retrieves the positional metadata associated with a transaction
diff --git a/core/rawdb/accessors_indexes_test.go b/core/rawdb/accessors_indexes_test.go
index e64818b1..0d2aaf7d 100644
--- a/core/rawdb/accessors_indexes_test.go
+++ b/core/rawdb/accessors_indexes_test.go
@@ -22,10 +22,10 @@ import (
   "math/big"
   "testing"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/ethdb"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/ethdb"
+   "golang.org/x/crypto/sha3"
)

diff --git a/core/rawdb/accessors_metadata.go b/core/rawdb/accessors_metadata.go
index 6b1a0f13..6c845dfc 100644
--- a/core/rawdb/accessors_metadata.go
+++ b/core/rawdb/accessors_metadata.go
@@ -28,12 +28,12 @@ package rawdb

import (
   "encoding/json"
+   "time"

-   "github.com/ava-labs/coreth/ethdb"
-   "github.com/ava-labs/coreth/params"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/log"
-   "github.com/ethereum/go-ethereum/rlp"

```

```

+         "github.com/flare-foundation/coreth/ethdb"
+         "github.com/flare-foundation/coreth/params"
+     )
+
+ // ReadDatabaseVersion retrieves the version number of the database.
@@ -89,3 +90,117 @@ func WriteChainConfig(db ethdb.KeyValueWriter, hash common.Hash, cfg *params.ChainConfig) error {
+     log.Crit("Failed to store chain config", "err", err)
+ }
+ }
+
+ // crashList is a list of unclean-shutdown-markers, for rlp-encoding to the
+ // database
+ type crashList struct {
+     Discarded uint64 // how many ucs have we deleted
+     Recent    []uint64 // unix timestamps of 10 latest unclean shutdowns
+ }
+
+ const crashesToKeep = 10
+
+ // PushUncleanShutdownMarker appends a new unclean shutdown marker and returns
+ // the previous data
+ // - a list of timestamps
+ // - a count of how many old unclean-shutdowns have been discarded
+ func PushUncleanShutdownMarker(db ethdb.KeyValueStore) ([]uint64, uint64, error) {
+     var uncleanShutdowns crashList
+     // Read old data
+     if data, err := db.Get(uncleanShutdownKey); err != nil {
+         log.Warn("Error reading unclean shutdown markers", "error", err)
+     } else if err := rlp.DecodeBytes(data, &uncleanShutdowns); err != nil {
+         return nil, 0, err
+     }
+     var discarded = uncleanShutdowns.Discarded
+     var previous = make([]uint64, len(uncleanShutdowns.Recent))
+     copy(previous, uncleanShutdowns.Recent)
+     // Add a new (but cap it)
+     uncleanShutdowns.Recent = append(uncleanShutdowns.Recent, uint64(time.Now().Unix()))
+     if count := len(uncleanShutdowns.Recent); count > crashesToKeep+1 {
+         numDel := count - (crashesToKeep + 1)
+         uncleanShutdowns.Recent = uncleanShutdowns.Recent[numDel:]
+         uncleanShutdowns.Discarded += uint64(numDel)
+     }
+     // And save it again
+     data, _ := rlp.EncodeToBytes(uncleanShutdowns)
+     if err := db.Put(uncleanShutdownKey, data); err != nil {
+         log.Warn("Failed to write unclean-shutdown marker", "err", err)
+         return nil, 0, err
+     }
+     return previous, discarded, nil
+ }
+
+ // PopUncleanShutdownMarker removes the last unclean shutdown marker
+ func PopUncleanShutdownMarker(db ethdb.KeyValueStore) {
+     var uncleanShutdowns crashList
+     // Read old data
+     if data, err := db.Get(uncleanShutdownKey); err != nil {
+         log.Warn("Error reading unclean shutdown markers", "error", err)
+     } else if err := rlp.DecodeBytes(data, &uncleanShutdowns); err != nil {
+         log.Error("Error decoding unclean shutdown markers", "error", err) // Should mos def _not_ happen
+     }
+     if l := len(uncleanShutdowns.Recent); l > 0 {
+         uncleanShutdowns.Recent = uncleanShutdowns.Recent[:l-1]
+     }
+     data, _ := rlp.EncodeToBytes(uncleanShutdowns)
+     if err := db.Put(uncleanShutdownKey, data); err != nil {
+         log.Warn("Failed to clear unclean-shutdown marker", "err", err)
+     }
+ }
+
+ // UpdateUncleanShutdownMarker updates the last marker's timestamp to now.
+ func UpdateUncleanShutdownMarker(db ethdb.KeyValueStore) {
+     var uncleanShutdowns crashList
+     // Read old data
+     if data, err := db.Get(uncleanShutdownKey); err != nil {
+         log.Warn("Error reading unclean shutdown markers", "error", err)
+     } else if err := rlp.DecodeBytes(data, &uncleanShutdowns); err != nil {
+         log.Warn("Error decoding unclean shutdown markers", "error", err)
+     }
+     // This shouldn't happen because we push a marker on Backend instantiation
+     count := len(uncleanShutdowns.Recent)
+     if count == 0 {
+         log.Warn("No unclean shutdown marker to update")
+         return
+     }
+     uncleanShutdowns.Recent[count-1] = uint64(time.Now().Unix())
+     data, _ := rlp.EncodeToBytes(uncleanShutdowns)
+     if err := db.Put(uncleanShutdownKey, data); err != nil {
+         log.Warn("Failed to write unclean-shutdown marker", "err", err)
+     }
+ }
+
+ // WriteOfflinePruning writes a marker of the last attempt to run offline pruning
+ // The marker is written when offline pruning completes and is deleted when the node
+ // is started successfully with offline pruning disabled. This ensures users must
+ // disable offline pruning and start their node successfully between runs of offline
+ // pruning.
+ func WriteOfflinePruning(db ethdb.KeyValueStore) error {
+     data, err := rlp.EncodeToBytes(uint64(time.Now().Unix()))
+     if err != nil {
+         return err
+     }
+     return db.Put(offlinePruningKey, data)
+ }
+
+ // ReadOfflinePruning reads to check if there is a marker of the last attempt
+ // to run offline pruning.
+ func ReadOfflinePruning(db ethdb.KeyValueStore) (uint64, error) {
+     data, err := db.Get(offlinePruningKey)
+     if err != nil {
+         return 0, err
+     }
+     var offlinePruningRun uint64
+     if err := rlp.DecodeBytes(data, &offlinePruningRun); err != nil {
+         return 0, err
+     }
+     return offlinePruningRun, nil
+ }
+
+ // DeleteOfflinePruning deletes any marker of the last attempt to run offline pruning.
+ func DeleteOfflinePruning(db ethdb.KeyValueStore) error {
+     return db.Delete(offlinePruningKey)
+ }
+
diff --git a/core/rawdb/accessors_snapshot.go b/core/rawdb/accessors_snapshot.go
index 2c7c6c74..0e97a99b 100644
--- a/core/rawdb/accessors_snapshot.go
+++ b/core/rawdb/accessors_snapshot.go
@@ -27,11 +27,9 @@ package rawdb

import (
-     "encoding/binary"
-
-     "github.com/ava-labs/coreth/ethdb"

```

```

        "github.com/ethereum/go-ethereum/common"
        "github.com/ethereum/go-ethereum/log"
+       "github.com/flare-foundation/coreth/ethdb"
    )

    // ReadSnapshotRoot retrieves the root of the block whose state is contained in
@@ -158,56 +156,3 @@ func DeleteSnapshotGenerator(db ethdb.KeyValueWriter) {
    log.Crit("Failed to remove snapshot generator", "err", err)
}

- }
-
- // ReadSnapshotRecoveryNumber retrieves the block number of the last persisted
- // snapshot layer.
- func ReadSnapshotRecoveryNumber(db ethdb.KeyValueReader) *uint64 {
-     data, _ := db.Get(snapshotRecoveryKey)
-     if len(data) == 0 {
-         return nil
-     }
-     if len(data) != 8 {
-         return nil
-     }
-     number := binary.BigEndian.Uint64(data)
-     return &number
- }
-
- // WriteSnapshotRecoveryNumber stores the block number of the last persisted
- // snapshot layer.
- func WriteSnapshotRecoveryNumber(db ethdb.KeyValueWriter, number uint64) {
-     var buf [8]byte
-     binary.BigEndian.PutUint64(buf[:], number)
-     if err := db.Put(snapshotRecoveryKey, buf[:]); err != nil {
-         log.Crit("Failed to store snapshot recovery number", "err", err)
-     }
- }
-
- // DeleteSnapshotRecoveryNumber deletes the block number of the last persisted
- // snapshot layer.
- func DeleteSnapshotRecoveryNumber(db ethdb.KeyValueWriter) {
-     if err := db.Delete(snapshotRecoveryKey); err != nil {
-         log.Crit("Failed to remove snapshot recovery number", "err", err)
-     }
- }
-
- // ReadSnapshotSyncStatus retrieves the serialized sync status saved at shutdown.
- func ReadSnapshotSyncStatus(db ethdb.KeyValueReader) []byte {
-     data, _ := db.Get(snapshotSyncStatusKey)
-     return data
- }
-
- // WriteSnapshotSyncStatus stores the serialized sync status to save at shutdown.
- func WriteSnapshotSyncStatus(db ethdb.KeyValueWriter, status []byte) {
-     if err := db.Put(snapshotSyncStatusKey, status); err != nil {
-         log.Crit("Failed to store snapshot sync status", "err", err)
-     }
- }
-
- // DeleteSnapshotSyncStatus deletes the serialized sync status saved at the last
- // shutdown
- func DeleteSnapshotSyncStatus(db ethdb.KeyValueWriter) {
-     if err := db.Delete(snapshotSyncStatusKey); err != nil {
-         log.Crit("Failed to remove snapshot sync status", "err", err)
-     }
- }
-
diff --git a/core/rawdb/accessors_state.go b/core/rawdb/accessors_state.go
index 38f663cd..f17ac27e 100644
--- a/core/rawdb/accessors_state.go
+++ b/core/rawdb/accessors_state.go
@@ -27,9 +27,9 @@
package rawdb

import (
-     "github.com/ava-labs/coreth/ethdb"
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/log"
+     "github.com/flare-foundation/coreth/ethdb"
)

// ReadPreimage retrieves a single preimage of the provided hash.
diff --git a/core/rawdb/database.go b/core/rawdb/database.go
index 55a42d15..abf286c8 100644
--- a/core/rawdb/database.go
+++ b/core/rawdb/database.go
@@ -32,10 +32,11 @@
@@ -32,10 +32,11 @@ import (
    "os"
    "time"

-     "github.com/ava-labs/coreth/ethdb"
-     "github.com/ava-labs/coreth/ethdb/memorydb"
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/log"
+     "github.com/flare-foundation/coreth/ethdb"
+     "github.com/flare-foundation/coreth/ethdb/leveldb"
+     "github.com/flare-foundation/coreth/ethdb/memorydb"
+     "github.com/olekukonko/tablewriter"
)

@@ -63,6 +64,16 @@
@@ -63,6 +64,16 @@ func NewMemoryDatabaseWithCap(size int) ethdb.Database {
    return NewDatabase(memorydb.NewWithCap(size))
}

+ // NewLevelDBDatabase creates a persistent key-value database without a freezer
+ // moving immutable chain segments into cold storage.
+ func NewLevelDBDatabase(file string, cache int, handles int, namespace string, readonly bool) (ethdb.Database, error) {
+     db, err := leveldb.New(file, cache, handles, namespace, readonly)
+     if err != nil {
+         return nil, err
+     }
+     return NewDatabase(db), nil
+ }
+
+ type counter uint64

func (c counter) String() string {
@@ -108,7 +119,6 @@
@@ -108,7 +119,6 @@ func InspectDatabase(db ethdb.Database, keyPrefix, keyStart []byte) error {
    headers      stat
    bodies       stat
    receipts     stat
-     tds          stat
    numHashPairings stat
    hashNumPairings stat
    tries        stat
@@ -145,8 +155,6 @@
@@ -145,8 +155,6 @@ func InspectDatabase(db ethdb.Database, keyPrefix, keyStart []byte) error {
    bodies.Add(size)
    case bytes.HasPrefix(key, blockReceiptsPrefix) && len(key) == (len(blockReceiptsPrefix)+8+common.HashLength):
        receipts.Add(size)
-     case bytes.HasPrefix(key, headerPrefix) && bytes.HasSuffix(key, headerTDSuffix):
-         tds.Add(size)
    case bytes.HasPrefix(key, headerPrefix) && bytes.HasSuffix(key, headerHashSuffix):
        numHashPairings.Add(size)
    case bytes.HasPrefix(key, headerNumberPrefix) && len(key) == (len(headerNumberPrefix)+common.HashLength):
@@ -182,10 +190,8 @@
@@ -182,10 +190,8 @@ func InspectDatabase(db ethdb.Database, keyPrefix, keyStart []byte) error {
default:
    var accounted bool
    for _, meta := range [][]byte{

```

```

-         databaseVersionKey, headHeaderKey, headBlockKey, headFastBlockKey, lastPivotKey,
-         fastTrieProgressKey, snapshotDisabledKey, snapshotRootKey, snapshotJournalKey,
-         snapshotGeneratorKey, snapshotRecoveryKey, txIndexTailKey, fastTxLookupLimitKey,
-         uncleanShutdownKey, badBlockKey,
+         databaseVersionKey, headHeaderKey, headBlockKey,
+         snapshotRootKey, snapshotGeneratorKey, uncleanShutdownKey,
+
    } {
        if bytes.Equal(key, meta) {
            metadata.Add(size)
@@ -208,7 +214,6 @@ func InspectDatabase(db ethdb.Database, keyPrefix, keyStart []byte) error {
    { "Key-Value store", "Headers", headers.Size(), headers.Count() },
    { "Key-Value store", "Bodies", bodies.Size(), bodies.Count() },
    { "Key-Value store", "Receipt lists", receipts.Size(), receipts.Count() },
-    { "Key-Value store", "Difficulties", tds.Size(), tds.Count() },
    { "Key-Value store", "Block number->hash", numHashPairings.Size(), numHashPairings.Count() },
    { "Key-Value store", "Block hash->number", hashNumPairings.Size(), hashNumPairings.Count() },
    { "Key-Value store", "Transaction index", txLookups.Size(), txLookups.Count() },
diff --git a/core/rawdb/schema.go b/core/rawdb/schema.go
index 67bdd634..6f6701fc 100644
--- a/core/rawdb/schema.go
+++ b/core/rawdb/schema.go
@@ -46,51 +46,23 @@ var (
    // headBlockKey tracks the latest known full block's hash.
    headBlockKey = []byte("LastBlock")

-    // headFastBlockKey tracks the latest known incomplete block's hash during fast sync.
-    headFastBlockKey = []byte("LastFast")
-
-    // lastPivotKey tracks the last pivot block used by fast sync (to reenale on sethead).
-    lastPivotKey = []byte("LastPivot")
-
-    // fastTrieProgressKey tracks the number of trie entries imported during fast sync.
-    fastTrieProgressKey = []byte("TrieSync")
-
-    // snapshotDisabledKey flags that the snapshot should not be maintained due to initial sync.
-    snapshotDisabledKey = []byte("SnapshotDisabled")
-
-    // snapshotRootKey tracks the hash of the last snapshot.
-    snapshotRootKey = []byte("SnapshotRoot")
-
-    // snapshotBlockHashKey tracks the block hash of the last snapshot.
-    snapshotBlockHashKey = []byte("SnapshotBlockHash")
-
-    // snapshotJournalKey tracks the in-memory diff layers across restarts.
-    snapshotJournalKey = []byte("SnapshotJournal")
-
-    // snapshotGeneratorKey tracks the snapshot generation marker across restarts.
-    snapshotGeneratorKey = []byte("SnapshotGenerator")
-
-    // snapshotRecoveryKey tracks the snapshot recovery marker across restarts.
-    snapshotRecoveryKey = []byte("SnapshotRecovery")
-
-    // snapshotSyncStatusKey tracks the snapshot sync status across restarts.
-    snapshotSyncStatusKey = []byte("SnapshotSyncStatus")
-
-    // txIndexTailKey tracks the oldest block whose transactions have been indexed.
-    txIndexTailKey = []byte("TransactionIndexTail")
-
-    // fastTxLookupLimitKey tracks the transaction lookup limit during fast sync.
-    fastTxLookupLimitKey = []byte("FastTransactionLookupLimit")
-
-    // badBlockKey tracks the list of bad blocks seen by local
-    badBlockKey = []byte("InvalidBlock")
-
-    // uncleanShutdownKey tracks the list of local crashes
    uncleanShutdownKey = []byte("unclean-shutdown") // config prefix for the db

+    // offlinePruningKey tracks runs of offline pruning
+    offlinePruningKey = []byte("OfflinePruning")
+
    // Data item prefixes (use single byte to avoid mixing data types, avoid `i`, used for indexes).
    headerPrefix      = []byte("h") // headerPrefix + num (uint64 big endian) + hash -> header
-    headerTDSuffix    = []byte("t") // headerPrefix + num (uint64 big endian) + hash + headerTDSuffix -> td
    headerHashSuffix   = []byte("n") // headerPrefix + num (uint64 big endian) + headerHashSuffix -> hash
    headerNumberPrefix = []byte("H") // headerNumberPrefix + hash -> num (uint64 big endian)

@@ -138,11 +110,6 @@ func headerKey(number uint64, hash common.Hash) []byte {
    return append(append(headerPrefix, encodeBlockNumber(number)...), hash.Bytes()...)
}

-// headerTDKey = headerPrefix + num (uint64 big endian) + hash + headerTDSuffix
-func headerTDKey(number uint64, hash common.Hash) []byte {
-    return append(headerKey(number, hash), headerTDSuffix...)
-}
-
// headerHashKey = headerPrefix + num (uint64 big endian) + headerHashSuffix
func headerHashKey(number uint64) []byte {
    return append(append(headerPrefix, encodeBlockNumber(number)...), headerHashSuffix...)
diff --git a/core/rawdb/table.go b/core/rawdb/table.go
index 179ebf39..3e4f4bdd 100644
--- a/core/rawdb/table.go
+++ b/core/rawdb/table.go
@@ -27,7 +27,7 @@ package rawdb

import (
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

// table is a wrapper around a database that prefixes each key access with a pre-
diff --git a/core/rawdb/table_test.go b/core/rawdb/table_test.go
index c7cac982..2f3bd97f 100644
--- a/core/rawdb/table_test.go
+++ b/core/rawdb/table_test.go
@@ -30,7 +30,7 @@ func TestTableDatabase(t *testing.T) {
    "bytes"
    "testing"

-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

func TestTableDatabase(t *testing.T) {
    { testTableDatabase(t, "prefix") }
diff --git a/core/rpl_test.go b/core/rpl_test.go
new file mode 100644
index 00000000..6815d1ab
--- /dev/null
+++ b/core/rpl_test.go
@@ -0,0 +1,219 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2019 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify

```

```

+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "fmt"
+    "math/big"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "golang.org/x/crypto/sha3"
+)
+
+func getBlock(transactions int, uncles int, dataSize int) *types.Block {
+    var (
+        aa = common.HexToAddress("0x0000000000000000000000000000000000000000000000000000000000000000")
+        // Generate a canonical chain to act as the main dataset
+        engine = dummy.NewFaker()
+        db      = rawdb.NewMemoryDatabase()
+        // A sender who makes transactions, has some funds
+        key, _  = crypto.HexToECDSA("b71c71a67e1177ad4e901695e1b4b9ee17ae16c6668d313eac2f96dbcda3f291")
+        address = crypto.PubkeyToAddress(key.PublicKey)
+        funds   = big.NewInt(50000 * 225000000000 * 200)
+        gsSpec  = &Genesis{
+            Config: params.TestChainConfig,
+            Alloc:  GenesisAlloc{address: {Balance: funds}},
+        }
+        genesis = gsSpec.MustCommit(db)
+    )
+
+    // We need to generate as many blocks +1 as uncles
+    blocks, _, _ := GenerateChain(params.TestChainConfig, genesis, engine, db, uncles+1, 10,
+        func(n int, b *BlockGen) {
+            if n == uncles {
+                // Add transactions and stuff on the last block
+                for i := 0; i < transactions; i++ {
+                    tx, _ := types.SignTx(types.NewTransaction(uint64(i), aa,
+                        big.NewInt(0), 50000, b.header.BaseFee, make([]byte, dataSize)), types.LatestSigner(params.TestChainConfig), key)
+                    b.AddTx(tx)
+                }
+                for i := 0; i < uncles; i++ {
+                    b.AddUncle(&types.Header{ParentHash: b.PrevBlock(n - 1 - i).Hash(), Number: big.NewInt(int64(n - i))})
+                }
+            }
+        })
+    block := blocks[len(blocks)-1]
+    return block
+}
+
+// TestRlpIterator tests that individual transactions can be picked out
+// from blocks without full unmarshalling/marshalling
+func TestRlpIterator(t *testing.T) {
+    for _, tt := range []struct {
+        txs      int
+        uncles   int
+        dataSize int
+    }{
+        {0, 0, 0},
+        {0, 2, 0},
+        {10, 0, 0},
+        {10, 2, 0},
+        {10, 2, 50},
+    } {
+        testRlpIterator(t, tt.txs, tt.uncles, tt.dataSize)
+    }
+}
+
+func testRlpIterator(t *testing.T, txs, uncles, dataSize int) {
+    desc := fmt.Sprintf("%d txs [%d dataSize] and %d uncles", txs, dataSize, uncles)
+    bodyRlp, _ := rlp.EncodeToBytes(getBlock(txs, uncles, dataSize).Body())
+    it, err := rlp.NewListIterator(bodyRlp)
+    if err != nil {
+        t.Fatal(err)
+    }
+    // Check that txs exist
+    if !it.Next() {
+        t.Fatal("expected four elems, got zero")
+    }
+    txdata := it.Value()
+    // Check that uncles exist
+    if !it.Next() {
+        t.Fatal("expected four elems, got one")
+    }
+    // Check that version exist
+    if !it.Next() {
+        t.Fatal("expected four elems, got two")
+    }
+    // Check that extdata exist
+    if !it.Next() {
+        t.Fatal("expected four elems, got three")
+    }
+    // No more after that
+    if it.Next() {
+        t.Fatal("expected only four elems, got more")
+    }
+    txIt, err := rlp.NewListIterator(txdata)
+    if err != nil {
+        t.Fatal(err)
+    }
+    var gotHashes []common.Hash
+    var expHashes []common.Hash
+    for txIt.Next() {
+        gotHashes = append(gotHashes, crypto.Keccak256Hash(txIt.Value()))
+    }
+
+    var expBody types.Body
+    err = rlp.DecodeBytes(bodyRlp, &expBody)
+    if err != nil {
+        t.Fatal(err)
+    }
+    for _, tx := range expBody.Transactions {
+        expHashes = append(expHashes, tx.Hash())
+    }
+    if gotLen, expLen := len(gotHashes), len(expHashes); gotLen != expLen {
+        t.Fatalf("testcase %v: length wrong, got %d exp %d", desc, gotLen, expLen)
+    }
+    // also sanity check against input
+    if gotLen := len(gotHashes); gotLen != txs {

```



```

+         t.Fatalf("testcase %v: length wrong, got %d exp %d", desc, gotLen, txs)
+     }
+     for i, got := range gotHashes {
+         if exp := expHashes[i]; got != exp {
+             t.Errorf("testcase %v: hash wrong, got %x, exp %x", desc, got, exp)
+         }
+     }
+ }
+}
+
+// BenchmarkHashing compares the speeds of hashing a rlp raw data directly
+// without the unmarshalling/marshalling step
+func BenchmarkHashing(b *testing.B) {
+    // Make a pretty fat block
+    var (
+        bodyRlp []byte
+        blockRlp []byte
+    )
+    {
+        block := getBlock(200, 2, 50)
+        bodyRlp, _ = rlp.EncodeToBytes(block.Body())
+        blockRlp, _ = rlp.EncodeToBytes(block)
+    }
+    var got common.Hash
+    var hasher = sha3.NewLegacyKeccak256()
+    b.Run("iteratorhashing", func(b *testing.B) {
+        b.ResetTimer()
+        for i := 0; i < b.N; i++ {
+            var hash common.Hash
+            it, err := rlp.NewListIterator(bodyRlp)
+            if err != nil {
+                b.Fatal(err)
+            }
+            it.Next()
+            txs := it.Value()
+            txIt, err := rlp.NewListIterator(txs)
+            if err != nil {
+                b.Fatal(err)
+            }
+            for txIt.Next() {
+                hasher.Reset()
+                hasher.Write(txIt.Value())
+                hasher.Sum(hash[:0])
+                got = hash
+            }
+        }
+    })
+    var exp common.Hash
+    b.Run("fullbodyhashing", func(b *testing.B) {
+        b.ResetTimer()
+        for i := 0; i < b.N; i++ {
+            var body types.Body
+            rlp.DecodeBytes(bodyRlp, &body)
+            for _, tx := range body.Transactions {
+                exp = tx.Hash()
+            }
+        }
+    })
+    b.Run("fullblockhashing", func(b *testing.B) {
+        b.ResetTimer()
+        for i := 0; i < b.N; i++ {
+            var block types.Block
+            rlp.DecodeBytes(blockRlp, &block)
+            for _, tx := range block.Transactions() {
+                tx.Hash()
+            }
+        }
+    })
+    if got != exp {
+        b.Fatalf("hash wrong, got %x exp %x", got, exp)
+    }
+}
+}
+
diff --git a/core/state/database.go b/core/state/database.go
index bef0af1e..eb98ad66 100644
--- a/core/state/database.go
+++ b/core/state/database.go
@@ -31,11 +31,11 @@ import (
     "fmt"

     "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
     lru "github.com/hashicorp/golang-lru"
 )

diff --git a/core/state/dump.go b/core/state/dump.go
index 8f30d482..e0cd9f82 100644
--- a/core/state/dump.go
+++ b/core/state/dump.go
@@ -31,12 +31,12 @@ import (
     "fmt"
     "time"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/trie"
 )

// DumpConfig is a set of options to control what portions of the state will be
diff --git a/core/state/iterator.go b/core/state/iterator.go
index 2ad4ed93..01ab5d55 100644
--- a/core/state/iterator.go
+++ b/core/state/iterator.go
@@ -30,10 +30,10 @@ import (
     "bytes"
     "fmt"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/trie"
 )

// NodeIterator is an iterator to traverse the entire state trie post-order,
diff --git a/core/state/pruner/bloom.go b/core/state/pruner/bloom.go
new file mode 100644
index 00000000..13fbeb0
--- /dev/null
+++ b/core/state/pruner/bloom.go
@@ -0,0 +1,142 @@
+@ -0,0 +1,142 @@
+
+// (c) 2019-2020, Ava Labs, Inc.

```

```

+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2020 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package pruner
+
+import (
+    "encoding/binary"
+    "errors"
+    "os"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    bloomfilter "github.com/holiman/bloomfilter/v2"
+)
+
+// stateBloomHasher is a wrapper around a byte blob to satisfy the interface API
+// requirements of the bloom library used. It's used to convert a trie hash or
+// contract code hash into a 64 bit mini hash.
+type stateBloomHasher []byte
+
+func (f stateBloomHasher) Write(p []byte) (n int, err error) { panic("not implemented") }
+func (f stateBloomHasher) Sum(b []byte) []byte { panic("not implemented") }
+func (f stateBloomHasher) Reset() { panic("not implemented") }
+func (f stateBloomHasher) BlockSize() int { panic("not implemented") }
+func (f stateBloomHasher) Size() int { return 8 }
+func (f stateBloomHasher) Sum64() uint64 { return binary.BigEndian.Uint64(f) }
+
+// stateBloom is a bloom filter used during the state conversion(snapshot->state).
+// The keys of all generated entries will be recorded here so that in the pruning
+// stage the entries belong to the specific version can be avoided for deletion.
+//
+// The false-positive is allowed here. The "false-positive" entries means they
+// actually don't belong to the specific version but they are not deleted in the
+// pruning. The downside of the false-positive allowance is we may leave some "dangling"
+// nodes in the disk. But in practice the it's very unlikely the dangling node is
+// state root. So in theory this pruned state shouldn't be visited anymore. Another
+// potential issue is for fast sync. If we do another fast sync upon the pruned
+// database, it's problematic which will stop the expansion during the syncing.
+// TODO address it @rjl493456442 @holiman @karalabe.
+//
+// After the entire state is generated, the bloom filter should be persisted into
+// the disk. It indicates the whole generation procedure is finished.
+type stateBloom struct {
+    bloom *bloomfilter.Filter
+}
+
+// newStateBloomWithSize creates a brand new state bloom for state generation.
+// The bloom filter will be created by the passing bloom filter size. According
+// to the https://hur.st/bloomfilter/?n=600000000&p=5m=2048MB&k=4, the parameters
+// are picked so that the false-positive rate for mainnet is low enough.
+func newStateBloomWithSize(size uint64) (*stateBloom, error) {
+    bloom, err := bloomfilter.New(size*1024*1024*8, 4)
+    if err != nil {
+        return nil, err
+    }
+    log.Info("Initialized state bloom", "size", common.StorageSize(float64(bloom.M())/8))
+    return &stateBloom{bloom: bloom}, nil
+}
+
+// NewStateBloomFromDisk loads the state bloom from the given file.
+// In this case the assumption is held the bloom filter is complete.
+func NewStateBloomFromDisk(filename string) (*stateBloom, error) {
+    bloom, _, err := bloomfilter.ReadFile(filename)
+    if err != nil {
+        return nil, err
+    }
+    return &stateBloom{bloom: bloom}, nil
+}
+
+// Commit flushes the bloom filter content into the disk and marks the bloom
+// as complete.
+func (bloom *stateBloom) Commit(filename, tempname string) error {
+    // Write the bloom out into a temporary file
+    _, err := bloom.bloom.WriteFile(tempname)
+    if err != nil {
+        return err
+    }
+    // Ensure the file is synced to disk
+    f, err := os.OpenFile(tempname, os.O_RDWR, 0666)
+    if err != nil {
+        return err
+    }
+    if err := f.Sync(); err != nil {
+        f.Close()
+        return err
+    }
+    f.Close()
+
+    // Move the temporary file into its final location
+    return os.Rename(tempname, filename)
+}
+
+// Put implements the KeyValueWriter interface. But here only the key is needed.
+func (bloom *stateBloom) Put(key []byte, value []byte) error {
+    // If the key length is not 32bytes, ensure it's contract code
+    // entry with new scheme.
+    if len(key) != common.HashLength {
+        isCode, codeKey := rawdb.IsCodeKey(key)
+        if !isCode {
+            return errors.New("invalid entry")
+        }
+        bloom.bloom.Add(stateBloomHasher(codeKey))
+        return nil
+    }
+    bloom.bloom.Add(stateBloomHasher(key))
+    return nil
+}
+
+// Delete removes the key from the key-value data store.
+func (bloom *stateBloom) Delete(key []byte) error { panic("not supported") }
+
+

```

```

+// Contain is the wrapper of the underlying contains function which
+// reports whether the key is contained.
+// - If it says yes, the key may be contained
+// - If it says no, the key is definitely not contained.
+func (bloom *stateBloom) Contain(key []byte) (bool, error) {
+    return bloom.bloom.Contains(stateBloomHasher(key)), nil
+}
diff --git a/core/state/pruner/pruner.go b/core/state/pruner/pruner.go
new file mode 100644
index 00000000..50537b4b
--- /dev/null
+++ b/core/state/pruner/pruner.go
@@ -0,0 +1,417 @@
+// (c) 2019-2020, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2020 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package pruner
+
+import (
+    "bytes"
+    "encoding/binary"
+    "errors"
+    "fmt"
+    "math"
+    "os"
+    "path/filepath"
+    "strings"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state/snapshot"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
+)
+
+const (
+    // stateBloomFilePrefix is the filename prefix of state bloom filter.
+    stateBloomFilePrefix = "statebloom"
+
+    // stateBloomFilePrefix is the filename suffix of state bloom filter.
+    stateBloomFileSuffix = ".bf.gz"
+
+    // stateBloomFileTempSuffix is the filename suffix of state bloom filter
+    // while it is being written out to detect write aborts.
+    stateBloomFileTempSuffix = ".tmp"
+
+    // rangeCompactionThreshold is the minimal deleted entry number for
+    // triggering range compaction. It's a quite arbitrary number but just
+    // to avoid triggering range compaction because of small deletion.
+    rangeCompactionThreshold = 100000
+)
+
+var (
+    // emptyRoot is the known root hash of an empty trie.
+    emptyRoot = common.HexToHash("56e81f171bcc55a6ff8345e692c0f86e5b48e01b9996cadc001622fb5e363b421")
+
+    // emptyCode is the known hash of the empty EVM bytecode.
+    emptyCode = crypto.Keccak256(nil)
+)
+
+// Pruner is an offline tool to prune the stale state with the
+// help of the snapshot. The workflow of pruner is very simple:
+//
+// - iterate the snapshot, reconstruct the relevant state
+// - iterate the database, delete all other state entries which
+//   don't belong to the target state and the genesis state
+//
+// It can take several hours(around 2 hours for mainnet) to finish
+// the whole pruning work. It's recommended to run this offline tool
+// periodically in order to release the disk usage and improve the
+// disk read performance to some extent.
+type Pruner struct {
+    db          ethdb.Database
+    stateBloom  *stateBloom
+    datadir     string
+    headHeader  *types.Header
+    snaptree    *snapshot.Tree
+}
+
+// NewPruner creates the pruner instance.
+func NewPruner(db ethdb.Database, datadir string, bloomSize uint64) (*Pruner, error) {
+    headBlock := rawdb.ReadHeadBlock(db)
+    if headBlock == nil {
+        return nil, errors.New("Failed to load head block")
+    }
+    // Note: we refuse to start a pruning session unless the snapshot disk layer exists, which should prevent
+    // us from ever needing to enter RecoverPruning in an invalid pruning session (a session where we do not have
+    // the protected trie in the triedb and in the snapshot disk layer).
+    snaptree, err := snapshot.New(db, trie.NewDatabase(db), 256, headBlock.Hash(), headBlock.Root(), false, false, false)
+    if err != nil {
+        return nil, fmt.Errorf("failed to create snapshot for pruning, must restart without offline pruning disabled to recover: %w", err) // The relevant snapshot(s) might not exist
+    }
+    // Sanitize the bloom filter size if it's too small.
+    if bloomSize < 256 {
+        log.Warn("Sanitizing bloomfilter size", "provided(MB)", bloomSize, "updated(MB)", 256)
+        bloomSize = 256
+    }
+    stateBloom, err := newStateBloomWithSize(bloomSize)
+    if err != nil {
+        return nil, err
+    }
+    return &Pruner{
+        db:          db,
+        stateBloom:  stateBloom,
+        datadir:     datadir,
+        headHeader:  headBlock.Header(),
+    }, nil
+}

```

```

+         snaptree: snaptree,
+     }, nil
+ }
+
+func prune(maindb ethdb.Database, stateBloom *stateBloom, bloomPath string, start time.Time) error {
+    // Delete all stale trie nodes in the disk. With the help of state bloom
+    // the trie nodes(and codes) belong to the active state will be filtered
+    // out. A very small part of state tries will also be filtered because of
+    // the false-positive rate of bloom filter. But the assumption is held here
+    // that the false-positive is low enough(~0.05%). The probability of the
+    // dangling node is the state root is super low. So the dangling nodes in
+    // theory will never ever be visited again.
+    var (
+        count int
+        size   common.StorageSize
+        pstart = time.Now()
+        logged = time.Now()
+        batch  = maindb.NewBatch()
+        iter   = maindb.NewIterator(nil, nil)
+    )
+    // We wrap iter.Release() in an anonymous function so that the [iter]
+    // value captured is the value of [iter] at the end of the function as opposed
+    // to incorrectly capturing the first iterator immediately.
+    defer func() {
+        iter.Release()
+    }()
+
+    for iter.Next() {
+        key := iter.Key()
+
+        // All state entries don't belong to specific state and genesis are deleted here
+        // - trie node
+        // - legacy contract code
+        // - new-scheme contract code
+        isCode, codeKey := rawdb.IsCodeKey(key)
+        if len(key) == common.HashLength || isCode {
+            checkKey := key
+            if isCode {
+                checkKey = codeKey
+            }
+            if ok, err := stateBloom.Contain(checkKey); err != nil {
+                return err
+            } else if ok {
+                continue
+            }
+            count += 1
+            size += common.StorageSize(len(key) + len(iter.Value()))
+            if err := batch.Delete(key); err != nil {
+                return err
+            }
+
+            var eta time.Duration // Realistically will never remain uninited
+            if done := binary.BigEndian.Uint64(key[:8]); done > 0 {
+                var (
+                    left  = math.MaxUint64 - binary.BigEndian.Uint64(key[:8])
+                    speed = done/uint64(time.Since(pstart)/time.Millisecond+1) + 1 // +1s to avoid division by zero
+                )
+                eta = time.Duration(left/speed) * time.Millisecond
+            }
+            if time.Since(logged) > 8*time.Second {
+                log.Info("Pruning state data", "nodes", count, "size", size,
+                    "elapsed", common.PrettyDuration(time.Since(pstart)), "eta", common.PrettyDuration(eta))
+                logged = time.Now()
+            }
+            // Recreate the iterator after every batch commit in order
+            // to allow the underlying compactor to delete the entries.
+            if batch.ValueSize() >= ethdb.IdealBatchSize {
+                if err := batch.Write(); err != nil {
+                    return err
+                }
+                batch.Reset()
+
+                iter.Release()
+                iter = maindb.NewIterator(nil, key)
+            }
+        }
+    }
+
+    if err := iter.Error(); err != nil {
+        return fmt.Errorf("failed to iterate db during pruning: %w", err)
+    }
+
+    if batch.ValueSize() > 0 {
+        if err := batch.Write(); err != nil {
+            return err
+        }
+        batch.Reset()
+    }
+    iter.Release()
+    log.Info("Pruned state data", "nodes", count, "size", size, "elapsed", common.PrettyDuration(time.Since(pstart)))
+
+    // Write marker to DB to indicate offline pruning finished successfully. We write before calling os.RemoveAll
+    // to guarantee that if the node dies midway through pruning, then this will run during RecoverPruning.
+    if err := rawdb.WriteOfflinePruning(maindb); err != nil {
+        return fmt.Errorf("failed to write offline pruning success marker: %w", err)
+    }
+
+    // Delete the state bloom, it marks the entire pruning procedure is
+    // finished. If any crashes or manual exit happens before this,
+    // 'RecoverPruning' will pick it up in the next restarts to redo all
+    // the things.
+    if err := os.RemoveAll(bloomPath); err != nil {
+        return fmt.Errorf("failed to remove bloom filter from disk: %w", err)
+    }
+
+    // Start compactions, will remove the deleted data from the disk immediately.
+    // Note for small pruning, the compaction is skipped.
+    if count >= rangeCompactionThreshold {
+        cstart := time.Now()
+        for b := 0x00; b <= 0xf0; b += 0x10 {
+            var (
+                start = []byte(byte(b))
+                end   = []byte(byte(b + 0x10))
+            )
+            if b == 0xf0 {
+                end = nil
+            }
+            log.Info("Compacting database", "range", fmt.Sprintf("%#x-%#x", start, end), "elapsed", common.PrettyDuration(time.Since(cstart)))
+            if err := maindb.Compact(start, end); err != nil {
+                log.Error("Database compaction failed", "error", err)
+                return err
+            }
+        }
+        log.Info("Database compaction finished", "elapsed", common.PrettyDuration(time.Since(cstart)))
+    }
+    log.Info("State pruning successful", "pruned", size, "elapsed", common.PrettyDuration(time.Since(start)))
+    return nil
+}
+
+// Prune deletes all historical state nodes except the nodes belong to the
+// specified state version. If user doesn't specify the state version, use
+// the bottom-most snapshot diff layer as the target.
+func (p *Pruner) Prune(root common.Hash) error {
+    // If the state bloom filter is already committed previously,
+    // reuse it for pruning instead of generating a new one. It's
+    // mandatory because a part of state may already be deleted,

```

```

+ // the recovery procedure is necessary.
+ _, stateBloomRoot, err := findBloomFilter(p.datadir)
+ if err != nil {
+     return err
+ }
+ if stateBloomRoot != (common.Hash{}) {
+     return RecoverPruning(p.datadir, p.db)
+ }
+
+ // If the target state root is not specified, return a fatal error.
+ if root == (common.Hash{}) {
+     return fmt.Errorf("cannot prune with an empty root: %s", root)
+ }
+ // Ensure the root is really present. The weak assumption
+ // is the presence of root can indicate the presence of the
+ // entire trie.
+ if blob := rawdb.ReadTrieNode(p.db, root); len(blob) == 0 {
+     return fmt.Errorf("associated state[%x] is not present", root)
+ } else {
+     log.Info("Selecting last accepted block root as the pruning target", "root", root)
+ }
+
+ // Traverse the target state, re-construct the whole state trie and
+ // commit to the given bloom filter.
+ start := time.Now()
+ if err := snapshot.GenerateTrie(p.snaptree, root, p.db, p.stateBloom); err != nil {
+     return err
+ }
+ // Traverse the genesis, put all genesis state entries into the
+ // bloom filter too.
+ if err := extractGenesis(p.db, p.stateBloom); err != nil {
+     return err
+ }
+ filterName := bloomFilterName(p.datadir, root)
+
+ log.Info("Writing state bloom to disk", "name", filterName)
+ if err := p.stateBloom.Commit(filterName, filterName+stateBloomFileTempSuffix); err != nil {
+     return err
+ }
+ log.Info("State bloom filter committed", "name", filterName)
+ return prune(p.db, p.stateBloom, filterName, start)
+}
+
+// RecoverPruning will resume the pruning procedure during the system restart.
+// This function is used in this case: user tries to prune state data, but the
+// system was interrupted midway because of crash or manual-kill. In this case
+// if the bloom filter for filtering active state is already constructed, the
+// pruning can be resumed. What's more if the bloom filter is constructed, the
+// pruning **has to be resumed**. Otherwise a lot of dangling nodes may be left
+// in the disk.
+func RecoverPruning(datadir string, db ethdb.Database) error {
+    stateBloomPath, stateBloomRoot, err := findBloomFilter(datadir)
+    if err != nil {
+        return err
+    }
+    if stateBloomPath == "" {
+        return nil // nothing to recover
+    }
+    headBlock := rawdb.ReadHeadBlock(db)
+    if headBlock == nil {
+        return errors.New("Failed to load head block")
+    }
+    stateBloom, err := NewStateBloomFromDisk(stateBloomPath)
+    if err != nil {
+        return err
+    }
+    log.Info("Loaded state bloom filter", "path", stateBloomPath)
+
+    // All the state roots of the middle layers should be forcibly pruned,
+    // otherwise the dangling state will be left.
+    if stateBloomRoot != headBlock.Root() {
+        return fmt.Errorf("cannot recover pruning to state bloom root: %s, with head block root: %s", stateBloomRoot, headBlock.Root())
+    }
+
+    return prune(db, stateBloom, stateBloomPath, time.Now())
+}
+
+// extractGenesis loads the genesis state and commits all the state entries
+// into the given bloomfilter.
+func extractGenesis(db ethdb.Database, stateBloom *stateBloom) error {
+    genesisHash := rawdb.ReadCanonicalHash(db, 0)
+    if genesisHash == (common.Hash{}) {
+        return errors.New("missing genesis hash")
+    }
+    genesis := rawdb.ReadBlock(db, genesisHash, 0)
+    if genesis == nil {
+        return errors.New("missing genesis block")
+    }
+    t, err := trie.NewSecure(genesis.Root(), trie.NewDatabase(db))
+    if err != nil {
+        return err
+    }
+    accIter := t.NodeIterator(nil)
+    for accIter.Next(true) {
+        hash := accIter.Hash()
+
+        // Embedded nodes don't have hash.
+        if hash != (common.Hash{}) {
+            stateBloom.Put(hash.Bytes(), nil)
+        }
+        // If it's a leaf node, yes we are touching an account,
+        // dig into the storage trie further.
+        if accIter.Leaf() {
+            var acc types.StateAccount
+            if err := rlp.DecodeBytes(accIter.LeafBlob(), &acc); err != nil {
+                return err
+            }
+            if acc.Root != emptyRoot {
+                storageTrie, err := trie.NewSecure(acc.Root, trie.NewDatabase(db))
+                if err != nil {
+                    return err
+                }
+                storageIter := storageTrie.NodeIterator(nil)
+                for storageIter.Next(true) {
+                    hash := storageIter.Hash()
+                    if hash != (common.Hash{}) {
+                        stateBloom.Put(hash.Bytes(), nil)
+                    }
+                }
+                if storageIter.Error() != nil {
+                    return storageIter.Error()
+                }
+            }
+            if !bytes.Equal(acc.CodeHash, emptyCode) {
+                stateBloom.Put(acc.CodeHash, nil)
+            }
+        }
+    }
+    return accIter.Error()
+}
+
+func bloomFilterName(datadir string, hash common.Hash) string {
+    return filepath.Join(datadir, fmt.Sprintf("%s.%s.%s", stateBloomFilePrefix, hash.Hex(), stateBloomFileSuffix))
+}

```

```

+func isBloomFilter(filename string) (bool, common.Hash) {
+    filename = filepath.Base(filename)
+    if strings.HasPrefix(filename, stateBloomFilePrefix) && strings.HasSuffix(filename, stateBloomFileSuffix) {
+        return true, common.HexToHash(filename[len(stateBloomFilePrefix)+1 : len(filename)-len(stateBloomFileSuffix)-1])
+    }
+    return false, common.Hash{}
+}
+func findBloomFilter(datadir string) (string, common.Hash, error) {
+    var (
+        stateBloomPath string
+        stateBloomRoot common.Hash
+    )
+    if err := filepath.Walk(datadir, func(path string, info os.FileInfo, err error) error {
+        if info != nil && !info.IsDir() {
+            ok, root := isBloomFilter(path)
+            if ok {
+                stateBloomPath = path
+                stateBloomRoot = root
+            }
+        }
+    }); err != nil {
+        return "", common.Hash{}, err
+    }
+    return stateBloomPath, stateBloomRoot, nil
+}
diff --git a/core/state/snapshot/conversion.go b/core/state/snapshot/conversion.go
index bfb157a1..01b3d23d 100644
--- a/core/state/snapshot/conversion.go
+++ b/core/state/snapshot/conversion.go
@@ -36,12 +36,12 @@ import (
     "sync"
     "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
 )

// trieKV represents a trie key-value pair
diff --git a/core/state/snapshot/difflayer_test.go b/core/state/snapshot/difflayer_test.go
index 245acc83..562929f2 100644
--- a/core/state/snapshot/difflayer_test.go
+++ b/core/state/snapshot/difflayer_test.go
@@ -32,9 +32,9 @@ import (
     "testing"

     "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
 )

func copyDestructs(destructs map[common.Hash]struct{}) map[common.Hash]struct{} {
diff --git a/core/state/snapshot/disklayer.go b/core/state/snapshot/disklayer.go
index 07add6be..932722d5 100644
--- a/core/state/snapshot/disklayer.go
+++ b/core/state/snapshot/disklayer.go
@@ -32,11 +32,11 @@ import (
     "time"

     "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
 )

// diskLayer is a low level persistent snapshot built on top of a key-value store.
diff --git a/core/state/snapshot/disklayer_test.go b/core/state/snapshot/disklayer_test.go
index adf3778e..310a730f 100644
--- a/core/state/snapshot/disklayer_test.go
+++ b/core/state/snapshot/disklayer_test.go
@@ -30,11 +30,11 @@ import (
     "bytes"
     "testing"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
 )

// reverse reverses the contents of a byte slice. It's used to update random accs
diff --git a/core/state/snapshot/generate.go b/core/state/snapshot/generate.go
index ae655aec..b1031b69 100644
--- a/core/state/snapshot/generate.go
+++ b/core/state/snapshot/generate.go
@@ -34,14 +34,14 @@ import (
     "time"

     "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
 )

var (
diff --git a/core/state/snapshot/generate_test.go b/core/state/snapshot/generate_test.go
index 8c480e12..ee77b9c8 100644
--- a/core/state/snapshot/generate_test.go
+++ b/core/state/snapshot/generate_test.go
@@ -33,13 +33,13 @@ import (
     "testing"
     "time"

-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
+    "github.com/flare-foundation/coreth/trie"
 )

```

```

"github.com/ethereum/go-ethereum/common"
"github.com/ethereum/go-ethereum/core/rawdb"
"github.com/ethereum/go-ethereum/log"
"github.com/ethereum/go-ethereum/rlp"
+ "github.com/flare-foundation/coreth/ethdb"
+ "github.com/flare-foundation/coreth/ethdb/memorydb"
+ "github.com/flare-foundation/coreth/trie"
"golang.org/x/crypto/sha3"
)

diff --git a/core/state/snapshot/iterator.go b/core/state/snapshot/iterator.go
index b07c2938..3bb36a54 100644
--- a/core/state/snapshot/iterator.go
+++ b/core/state/snapshot/iterator.go
@@ -31,9 +31,9 @@ import (
    "fmt"
    "sort"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

// Iterator is an iterator to step over all the accounts or the specific
diff --git a/core/state/snapshot/iterator_test.go b/core/state/snapshot/iterator_test.go
index 9d1b7440..fb56870f 100644
--- a/core/state/snapshot/iterator_test.go
+++ b/core/state/snapshot/iterator_test.go
@@ -33,8 +33,8 @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/core/rawdb"
    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
)

// TestAccountIteratorBasics tests some simple single-layer(diff and disk) iteration
diff --git a/core/state/snapshot/journal.go b/core/state/snapshot/journal.go
index 7feba072..43d2a660 100644
--- a/core/state/snapshot/journal.go
+++ b/core/state/snapshot/journal.go
@@ -33,12 +33,12 @@ import (
    "time"

    "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/log"
    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
)

// journalGenerator is a disk layer entry containing the generator progress marker.
diff --git a/core/state/snapshot/snapshot.go b/core/state/snapshot/snapshot.go
index 6c65f938..dad4e98f 100644
--- a/core/state/snapshot/snapshot.go
+++ b/core/state/snapshot/snapshot.go
@@ -36,9 +36,9 @@ import (
    "time"

    "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/trie"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/log"
    "github.com/ethereum/go-ethereum/metrics"
@@ -737,10 +737,13 @@ func (t *Tree) Rebuild(blockHash, root common.Hash) {
    t.lock.Lock()
    defer t.lock.Unlock()

+<<<<<<<< HEAD
+// Firstly delete any recovery flag in the database. Because now we are
+// building a brand new snapshot. Also reenale the snapshot feature.
+rawdb.DeleteSnapshotRecoveryNumber(t.diskdb)

+=====
+>>>>>>> upstream-v0.8.5-rc.2
+// Track whether there's a wipe currently running and keep it alive if so
+var wiper chan struct{}

diff --git a/core/state/snapshot/snapshot_test.go b/core/state/snapshot/snapshot_test.go
index 6ccee14d..4f483997 100644
--- a/core/state/snapshot/snapshot_test.go
+++ b/core/state/snapshot/snapshot_test.go
@@ -33,9 +33,9 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
)

// randomHash generates a random blob of data and returns it as a hash.
diff --git a/core/state/snapshot/wipe.go b/core/state/snapshot/wipe.go
index 298a77a1..01f3cec7 100644
--- a/core/state/snapshot/wipe.go
+++ b/core/state/snapshot/wipe.go
@@ -30,10 +30,10 @@ import (
    "bytes"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

// wipeSnapshot starts a goroutine to iterate over the entire key-value database
diff --git a/core/state/snapshot/wipe_test.go b/core/state/snapshot/wipe_test.go
index dbf009ea..5d4329a8 100644
--- a/core/state/snapshot/wipe_test.go
+++ b/core/state/snapshot/wipe_test.go
@@ -30,9 +30,9 @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
    "github.com/ethereum/go-ethereum/common"

```

```

+ "github.com/flare-foundation/coreth/core/rawdb"
+ "github.com/flare-foundation/coreth/ethdb/memorydb"
+ )
+
+ // Tests that given a database with random data content, all parts of a snapshot
@@ -40,57 +40,35 @@ import (
func TestWipe(t *testing.T) {
    // Create a database with some random snapshot data
    db := memorydb.New()

    for i := 0; i < 128; i++ {
        account := randomHash()
        rawdb.WriteAccountSnapshot(db, account, randomHash().Bytes())
        for j := 0; j < 1024; j++ {
            rawdb.WriteStorageSnapshot(db, account, randomHash(), randomHash().Bytes())
        }
        rawdb.WriteAccountSnapshot(db, randomHash(), randomHash().Bytes())
    }
    rawdb.WriteSnapshotBlockHash(db, randomHash())
    rawdb.WriteSnapshotRoot(db, randomHash())

    // Add some random non-snapshot data too to make wiping harder
    for i := 0; i < 65536; i++ {
        // Generate a key that's the wrong length for a state snapshot item
        var keysize int
        for keysize == 0 || keysize == 32 || keysize == 64 {
            keysize = 8 + rand.Intn(64) // +8 to ensure we will "never" randomize duplicates
        }
        // Randomize the suffix, dedup and inject it under the snapshot namespace
        keysuffix := make([]byte, keysize)
    + for i := 0; i < 500; i++ {
        // Generate keys with wrong length for a state snapshot item
        keysuffix := make([]byte, 31)
        rand.Read(keysuffix)

        if rand.Int31n(2) == 0 {
            db.Put(append(rawdb.SnapshotAccountPrefix, keysuffix...), randomHash().Bytes())
        } else {
            db.Put(append(rawdb.SnapshotStoragePrefix, keysuffix...), randomHash().Bytes())
        }
    }
    // Sanity check that all the keys are present
    var items int

    it := db.NewIterator(rawdb.SnapshotAccountPrefix, nil)
    defer it.Release()

    for it.Next() {
        key := it.Key()
        if len(key) == len(rawdb.SnapshotAccountPrefix)+common.HashLength {
            items++
        }
        db.Put(append(rawdb.SnapshotAccountPrefix, keysuffix...), randomHash().Bytes())
        keysuffix = make([]byte, 33)
        rand.Read(keysuffix)
        db.Put(append(rawdb.SnapshotAccountPrefix, keysuffix...), randomHash().Bytes())
    }
    it = db.NewIterator(rawdb.SnapshotStoragePrefix, nil)
    defer it.Release()

    for it.Next() {
        key := it.Key()
        if len(key) == len(rawdb.SnapshotStoragePrefix)+2*common.HashLength {
            items++
        }
    + count := func() (items int) {
    +     it := db.NewIterator(rawdb.SnapshotAccountPrefix, nil)
    +     defer it.Release()
    +     for it.Next() {
    +         if len(it.Key()) == len(rawdb.SnapshotAccountPrefix)+common.HashLength {
    +             items++
    +         }
    +     }
    +     return items
    + }
    + if items != 128+128*1024 {
    +     t.Fatalf("snapshot size mismatch: have %d, want %d", items, 128+128*1024)
    + }
    // Sanity check that all the keys are present
    if items := count(); items != 128 {
        t.Fatalf("snapshot size mismatch: have %d, want %d", items, 128)
    }
    if hash := rawdb.ReadSnapshotBlockHash(db); hash == (common.Hash{}) {
        t.Errorf("snapshot block hash marker mismatch: have %#x, want <not-nil>", hash)
    }
@@ -102,40 +80,24 @@ func TestWipe(t *testing.T) {
    <-wipeSnapshot(db, true)

    // Iterate over the database and ensure no snapshot information remains
    it = db.NewIterator(rawdb.SnapshotAccountPrefix, nil)
    defer it.Release()

    for it.Next() {
        key := it.Key()
        if len(key) == len(rawdb.SnapshotAccountPrefix)+common.HashLength {
            t.Errorf("snapshot entry remained after wipe: %x", key)
        }
    }
    + if items := count(); items != 0 {
    +     t.Fatalf("snapshot size mismatch: have %d, want %d", items, 0)
    + }
    + it = db.NewIterator(rawdb.SnapshotStoragePrefix, nil)
    // Iterate over the database and ensure miscellaneous items are present
    items := 0
    it := db.NewIterator(nil, nil)
    defer it.Release()

    for it.Next() {
        key := it.Key()
        if len(key) == len(rawdb.SnapshotStoragePrefix)+2*common.HashLength {
            t.Errorf("snapshot entry remained after wipe: %x", key)
        }
    + items++
    + }
    + if items != 1000 {
    +     t.Fatalf("misc item count mismatch: have %d, want %d", items, 1000)
    + }
    +
    if hash := rawdb.ReadSnapshotBlockHash(db); hash != (common.Hash{}) {
        t.Errorf("snapshot block hash marker remained after wipe: %#x", hash)
    }
    if hash := rawdb.ReadSnapshotRoot(db); hash != (common.Hash{}) {
        t.Errorf("snapshot block root marker remained after wipe: %#x", hash)
    }
    // Iterate over the database and ensure miscellaneous items are present
    items = 0

    it = db.NewIterator(nil, nil)
    defer it.Release()

    for it.Next() {
        items++
    }
    if items != 65536 {
        t.Fatalf("misc item count mismatch: have %d, want %d", items, 65536)
    }
}
diff --git a/core/state/state_object.go b/core/state/state_object.go

```



```
index 0e5be407..8d09e5f3 100644
--- a/core/state/state_object.go
+++ b/core/state/state_object.go
@@ -34,11 +34,11 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/metrics"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/types"
+ )

var emptyCodeHash = crypto.Keccak256(nil)
diff --git a/core/state/state_test.go b/core/state/state_test.go
index e7a05ef2..89c1de80 100644
--- a/core/state/state_test.go
+++ b/core/state/state_test.go
@@ -27,9 +27,9 @@
package state

import (
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+ )

type stateTest struct {
diff --git a/core/state/statedb.go b/core/state/statedb.go
index 44a9ee47..2b7f0fe8 100644
--- a/core/state/statedb.go
+++ b/core/state/statedb.go
@@ -34,15 +34,15 @@ import (
    "sort"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state/snapshot"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/trie"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/log"
-    "github.com/ethereum/go-ethereum/metrics"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state/snapshot"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/trie"
+ )

type revision struct {
diff --git a/core/state/statedb_test.go b/core/state/statedb_test.go
index 9c295f9e..a4a98878 100644
--- a/core/state/statedb_test.go
+++ b/core/state/statedb_test.go
@@ -39,11 +39,11 @@ import (
    "testing"
    "testing/quick"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state/snapshot"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state/snapshot"
+    "github.com/flare-foundation/coreth/core/types"
+ )

// Tests that updating a state trie does not leak any database writes prior to
diff --git a/core/state/trie_prefetcher_test.go b/core/state/trie_prefetcher_test.go
index 563a773d..183baca2 100644
--- a/core/state/trie_prefetcher_test.go
+++ b/core/state/trie_prefetcher_test.go
@@ -31,8 +31,8 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/rawdb"
+ )

func filledStateDB() *StateDB {
diff --git a/core/state_connector.go b/core/state_connector.go
new file mode 100644
index 00000000..18c9b9af
--- /dev/null
+++ b/core/state_connector.go
@@ -0,0 +1,232 @@
+// (c) 2021, Flare Networks Limited. All rights reserved.
+// Please see the file LICENSE for licensing terms.
+
+package core
+
+import (
+    "encoding/hex"
+    "math/big"
+    "os"
+    "strings"
+
+    "github.com/ethereum/go-ethereum/common"
+
+    "github.com/flare-foundation/coreth/core/vm"
+)
+
+var (
+    flareChainID      = new(big.Int).SetUint64(14) // https://github.com/ethereum-lists/chains/blob/master/_data/chains/eip155-14.json
+    songbirdChainID   = new(big.Int).SetUint64(19) // https://github.com/ethereum-lists/chains/blob/master/_data/chains/eip155-19.json
+
+    flareStateConnectorActivationTime = new(big.Int).SetUint64(1000000000000)
+    songbirdStateConnectorActivationTime = new(big.Int).SetUint64(1000000000000)
+)
+
+type AttestationVotes struct {
+    reachedMajority bool
+    majorityDecision string
+    majorityAttestors []common.Address
+    divergentAttestors []common.Address
+    abstainedAttestors []common.Address
+}
+
+func GetTestingChain(chainID *big.Int) bool {
+    return chainID.Cmp(flareChainID) != 0 && chainID.Cmp(songbirdChainID) != 0
+}
+
+func GetStateConnectorActivated(chainID *big.Int, blockTime *big.Int) bool {
+    if GetTestingChain(chainID) {
+        return true
+    }
+}
```

```

+ } else if chainID.Cmp(flareChainID) == 0 {
+     return blockTime.Cmp(flareStateConnectorActivationTime) >= 0
+ } else if chainID.Cmp(songbirdChainID) == 0 {
+     return blockTime.Cmp(songbirdStateConnectorActivationTime) >= 0
+ }
+ return false
+}
+func GetStateConnectorContract(chainID *big.Int, blockTime *big.Int) common.Address {
+    switch {
+    case GetStateConnectorActivated(chainID, blockTime) && chainID.Cmp(songbirdChainID) == 0:
+        return common.HexToAddress("0x6b5DEa84F71052c1302b5fe652e17FD442D126a9")
+    default:
+        return common.HexToAddress("0x1000000000000000000000000000000000000000000000000000000000000000")
+    }
+}
+func GetStateConnectorCoinbaseSignalAddr(chainID *big.Int, blockTime *big.Int) common.Address {
+    switch {
+    default:
+        return common.HexToAddress("0x0000000000000000000000000000000000000000000000000000000000000000")
+    }
+}
+func SubmitAttestationSelector(chainID *big.Int, blockTime *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0xcf, 0xd1, 0xfd, 0xad}
+    }
+}
+func GetAttestationSelector(chainID *big.Int, blockTime *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0x29, 0xbe, 0x4d, 0xb2}
+    }
+}
+func FinaliseRoundSelector(chainID *big.Int, blockTime *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0xea, 0xeb, 0xf6, 0xd3}
+    }
+}
+func GetVoterWhitelisterSelector(chainID *big.Int, blockTime *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0x71, 0xe1, 0xfa, 0xd9}
+    }
+}
+func GetFtsowWhitelistedPriceProvidersSelector(chainID *big.Int, blockTime *big.Int) []byte {
+    switch {
+    default:
+        return []byte{0x09, 0xfc, 0xb4, 0x00}
+    }
+}
+// The default attestors are the FTSO price providers
+func (st *StateTransition) GetDefaultAttestors(chainID *big.Int, timestamp *big.Int) ([]common.Address, error) {
+    if os.Getenv("TESTING_ATTESTATION_PROVIDERS") != "" && GetTestingChain(chainID) {
+        return GetEnvAttestationProviders("TESTING"), nil
+    } else {
+        // Get VoterWhitelister contract
+        voterWhitelisterContractBytes, _, err := st.evm.Call(
+            vm.AccountRef(st.msg.From()),
+            common.HexToAddress(GetPrioritisedFTSOContract(timestamp)),
+            GetVoterWhitelisterSelector(chainID, timestamp),
+            GetKeeperGasMultiplier(st.evm.Context.BlockNumber)*st.evm.Context.GasLimit,
+            big.NewInt(0))
+        if err != nil {
+            return []common.Address{}, err
+        }
+        // Get FTSO price providers
+        voterWhitelisterContract := common.BytesToAddress(voterWhitelisterContractBytes)
+        priceProvidersBytes, _, err := st.evm.Call(
+            vm.AccountRef(st.msg.From()),
+            voterWhitelisterContract,
+            GetFtsowWhitelistedPriceProvidersSelector(chainID, timestamp),
+            GetKeeperGasMultiplier(st.evm.Context.BlockNumber)*st.evm.Context.GasLimit,
+            big.NewInt(0))
+        if err != nil {
+            return []common.Address{}, err
+        }
+        NUM_ATTESTORS := len(priceProvidersBytes) / 32
+        var attestors []common.Address
+        for i := 0; i < NUM_ATTESTORS; i++ {
+            attestors = append(attestors, common.BytesToAddress(priceProvidersBytes[i*32:(i+1)*32]))
+        }
+        return attestors, nil
+    }
+}
+func GetEnvAttestationProviders(attestorType string) []common.Address {
+    envAttestationProvidersString := os.Getenv(attestorType + "_ATTESTATION_PROVIDERS")
+    if envAttestationProvidersString == "" {
+        return []common.Address{}
+    }
+    envAttestationProviders := strings.Split(envAttestationProvidersString, ",")
+    NUM_ATTESTORS := len(envAttestationProviders)
+    var attestors []common.Address
+    for i := 0; i < NUM_ATTESTORS; i++ {
+        attestors = append(attestors, common.HexToAddress(envAttestationProviders[i]))
+    }
+    return attestors
+}
+func (st *StateTransition) GetAttestation(attestor common.Address, instructions []byte) (string, error) {
+    merkleRootHash, _, err := st.evm.Call(vm.AccountRef(attestor), st.to(), instructions, 20000, big.NewInt(0))
+    return hex.EncodeToString(merkleRootHash), err
+}
+func (st *StateTransition) CountAttestations(attestors []common.Address, instructions []byte) (AttestationVotes, error) {
+    var attestationVotes AttestationVotes
+    hashFrequencies := make(map[string][]common.Address)
+    for i, a := range attestors {
+        h, err := st.GetAttestation(a, instructions)
+        if err != nil {
+            attestationVotes.abstainedAttestors = append(attestationVotes.abstainedAttestors, a)
+        }
+        hashFrequencies[h] = append(hashFrequencies[h], attestors[i])
+    }
+    // Find the plurality
+    var pluralityNum int
+    var pluralityKey string
+    for key, val := range hashFrequencies {
+        if len(val) > pluralityNum {
+            pluralityNum = len(val)
+            pluralityKey = key
+        }
+    }
+    if pluralityNum > len(attestors)/2 {
+        attestationVotes.reachedMajority = true
+    }
+}

```

```

+         attestationVotes.majorityDecision = pluralityKey
+         attestationVotes.majorityAttestors = hashFrequencies[pluralityKey]
+     }
+     for key, val := range hashFrequencies {
+         if key != pluralityKey {
+             attestationVotes.divergentAttestors = append(attestationVotes.divergentAttestors, val...)
+         }
+     }
+     return attestationVotes, nil
+}
+
+func (st *StateTransition) FinalisePreviousRound(chainID *big.Int, timestamp *big.Int, currentRoundNumber []byte) error {
+    getAttestationSelector := GetAttestationSelector(chainID, timestamp)
+    instructions := append(getAttestationSelector[:], currentRoundNumber[:])...
+    defaultAttestors, err := st.GetDefaultAttestors(chainID, timestamp)
+    if err != nil {
+        return err
+    }
+    defaultAttestationVotes, err := st.CountAttestations(defaultAttestors, instructions)
+    if err != nil {
+        return err
+    }
+    localAttestors := GetEnvAttestationProviders("LOCAL")
+    var finalityReached bool
+    if len(localAttestors) > 0 {
+        localAttestationVotes, err := st.CountAttestations(localAttestors, instructions)
+        if defaultAttestationVotes.reachedMajority && localAttestationVotes.reachedMajority && defaultAttestationVotes.majorityDecision == localAttestationVotes.majorityDecision {
+            finalityReached = true
+        } else if err != nil || (defaultAttestationVotes.reachedMajority && defaultAttestationVotes.majorityDecision != localAttestationVotes.majorityDecision) {
+            // Make a back-up of the current state database, because this node is about to branch from the default set
+        }
+    } else if defaultAttestationVotes.reachedMajority {
+        finalityReached = true
+    }
+    if finalityReached {
+        // Finalise defaultAttestationVotes.majorityDecision
+        finaliseRoundSelector := FinaliseRoundSelector(chainID, timestamp)
+        finalisedData := append(finaliseRoundSelector[:], currentRoundNumber[:])...
+        merkleRootHashBytes, err := hex.DecodeString(defaultAttestationVotes.majorityDecision)
+        if err != nil {
+            return err
+        }
+        finalisedData = append(finalisedData[:], merkleRootHashBytes[:])...
+        coinbaseSignal := GetStateConnectorCoinbaseSignalAddr(chainID, timestamp)
+        originalCoinbase := st.evm.Context.Coinbase
+        defer func() {
+            st.evm.Context.Coinbase = originalCoinbase
+        }()
+        st.evm.Context.Coinbase = coinbaseSignal
+
+        _, _, err = st.evm.Call(vm.AccountRef(coinbaseSignal), st.to(), finalisedData, st.evm.Context.GasLimit, new(big.Int).SetUint64(0))
+        if err != nil {
+            return err
+        }
+    }
+    // Issue rewards to defaultAttestationVotes.majorityAttestors here:
+    return nil
+}
+
diff --git a/core/state_manager.go b/core/state_manager.go
index ad312d1d..c4d85008 100644
--- a/core/state_manager.go
+++ b/core/state_manager.go
@@ -30,14 +30,14 @@ import (
     "fmt"
     "math/rand"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
 )

const (
    commitInterval = 4096
-    tipBufferSize  = 16
+    tipBufferSize  = 128
)

type TrieWriter interface {
diff --git a/core/state_manager_test.go b/core/state_manager_test.go
index 17e2b1d4..9a1e02ad 100644
--- a/core/state_manager_test.go
+++ b/core/state_manager_test.go
@@ -7,7 +7,7 @@ import (
     "math/big"
     "testing"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/stretchr/testify/assert"
)

diff --git a/core/state_prefetcher.go b/core/state_prefetcher.go
index ale0cde5..33a9f9d9 100644
--- a/core/state_prefetcher.go
+++ b/core/state_prefetcher.go
@@ -30,11 +30,11 @@ import (
     "math/big"
     "sync/atomic"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/vm"
-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
)

// statePrefetcher is a basic Prefetcher, which blindly executes a block on top
diff --git a/core/state_processor.go b/core/state_processor.go
index 636eb20f..c5865b47 100644
--- a/core/state_processor.go
+++ b/core/state_processor.go
@@ -30,14 +30,14 @@ import (
     "fmt"
     "math/big"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/consensus/misc"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/misc"
+    "github.com/flare-foundation/coreth/core/state"

```

```

+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/core/vm"
+      "github.com/flare-foundation/coreth/params"
+    )
+
+    // StateProcessor is a basic Processor, which takes care of transitioning
diff --git a/core/state_processor_test.go b/core/state_processor_test.go
new file mode 100644
index 00000000..73238af6
--- /dev/null
+++ b/core/state_processor_test.go
@@ -0,0 +1,351 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2020 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package core
+
+import (
+    "math/big"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/trie"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
+    "golang.org/x/crypto/sha3"
+)
+
+// TestStateProcessorErrors tests the output from the 'core' errors
+// as defined in core/error.go. These errors are generated when the
+// blockchain imports bad blocks, meaning blocks which have valid headers but
+// contain invalid transactions
+func TestStateProcessorErrors(t *testing.T) {
+    var (
+        config = &params.ChainConfig{
+            ChainID:          big.NewInt(1),
+            HomesteadBlock:    big.NewInt(0),
+            EIP150Block:       big.NewInt(0),
+            EIP150Hash:        common.Hash{},
+            EIP155Block:       big.NewInt(0),
+            EIP158Block:       big.NewInt(0),
+            ByzantiumBlock:    big.NewInt(0),
+            ConstantinopleBlock: big.NewInt(0),
+            PetersburgBlock:    big.NewInt(0),
+            IstanbulBlock:      big.NewInt(0),
+            MuirGlacierBlock:   big.NewInt(0),
+            ApricotPhase1BlockTimestamp: big.NewInt(0),
+            ApricotPhase2BlockTimestamp: big.NewInt(0),
+            ApricotPhase3BlockTimestamp: big.NewInt(0),
+            ApricotPhase4BlockTimestamp: big.NewInt(0),
+            ApricotPhase5BlockTimestamp: big.NewInt(0),
+        }
+        signer    = types.LatestSigner(config)
+        testKey, _ = crypto.HexToECDSA("b71c71a67e1177ad4e901695e1b4b9ee17ae16c6668d313eac2f96dbca3df291")
+    )
+    var makeTx = func(nonce uint64, to common.Address, amount *big.Int, gasLimit uint64, gasPrice *big.Int, data []byte) *types.Transaction {
+        tx, _ := types.SignTx(types.NewTransaction(nonce, to, amount, gasLimit, gasPrice, data), signer, testKey)
+        return tx
+    }
+    var mkDynamicTx = func(nonce uint64, to common.Address, gasLimit uint64, gasTipCap, gasFeeCap *big.Int) *types.Transaction {
+        tx, _ := types.SignTx(types.NewTx(&types.DynamicFeeTx{
+            Nonce:      nonce,
+            GasTipCap:  gasTipCap,
+            GasFeeCap:  gasFeeCap,
+            Gas:        gasLimit,
+            To:         &to,
+            Value:      big.NewInt(0),
+        }), signer, testKey)
+        return tx
+    }
+    { // Tests against a 'recent' chain definition
+        var (
+            db      = rawdb.NewMemoryDatabase()
+            gspec   = &Genesis{
+                Config: config,
+                Alloc:  GenesisAlloc{
+                    common.HexToAddress("0x71562b71999873085b286dF957af199Ec94617F7"): GenesisAccount{
+                        Balance: big.NewInt(2000000000000000000), // 2 ether
+                        Nonce:   0,
+                    },
+                },
+                GasLimit: params.ApricotPhase1GasLimit,
+            }
+            genesis      = gspec.MustCommit(db)
+            blockchain, _ = NewBlockChain(db, DefaultCacheConfig, gspec.Config, dummy.NewFaker(), vm.Config{}, common.Hash{})
+        )
+        defer blockchain.Stop()
+        bigNumber := new(big.Int).SetBytes(common.FromHex("0xffffffffffffffffffffffffffffffffffffffffffffffffffffffff"))
+        tooBigNumber := new(big.Int).Set(bigNumber)
+        tooBigNumber.Add(tooBigNumber, common.Big1)
+        for i, tt := range []struct {
+            txs []*types.Transaction
+            want string
+        }{
+            { // ErrNonceTooLow
+                txs: []*types.Transaction{
+                    makeTx(0, common.Address{}, big.NewInt(0), params.TxGas, big.NewInt(225000000000), nil),
+                    makeTx(0, common.Address{}, big.NewInt(0), params.TxGas, big.NewInt(225000000000), nil),
+                },
+                want: "could not apply tx 1 [0x734d821c990099c6ae42d78072aadd3931c35328cf03ef4cf5b2a4ac9c398522]: nonce too low: address 0x71562b71999873085b286dF957af199Ec94617F7, tx: 0 : ",
+            },
+            { // ErrNonceTooHigh
+                txs: []*types.Transaction{
+                    makeTx(100, common.Address{}, big.NewInt(0), params.TxGas, big.NewInt(225000000000), nil),
+                },
+                want: "could not apply tx 0 [0x0df36254cfbef8ed6961b38fc68aacc777177166144c8a56bc8919e23a559bf4]: nonce too high: address 0x71562b71999873085b286dF957af199Ec94617F7, tx: 1 ",
+            },
+        }
+    }
+}

```

```

    },
    { // ErrGasLimitReached
      txs: []*types.Transaction{
        makeTx(0, common.Address{}, big.NewInt(0), 8000001, big.NewInt(225000000000), nil),
      },
      want: "could not apply tx 0 [0xfbe38b817aaa760c2766b56c019fcd8a506560a28fd41c69ae96bdaa4569e317]: gas limit reached",
    },
    { // ErrInsufficientFundsForTransfer
      txs: []*types.Transaction{
        makeTx(0, common.Address{}, big.NewInt(2000000000000000000), params.TxGas, big.NewInt(225000000000), nil),
      },
      want: "could not apply tx 0 [0xae1601ef55b676ebb824ee7e16a0d14af725b7f9cf5ec79e21f14833c26b5b35]: insufficient funds for gas * price + value: address 0x71562b71999873085b2",
    },
    { // ErrInsufficientFunds
      txs: []*types.Transaction{
        makeTx(0, common.Address{}, big.NewInt(0), params.TxGas, big.NewInt(900000000000000000), nil),
      },
      want: "could not apply tx 0 [0x4a69690c4b0cd85e64d0d9ea06302455b01e10a83db964d60281739752003440]: insufficient funds for gas * price + value: address 0x71562b71999873085b2",
    },
    { // ErrGasUintOverflow
      // One missing 'core' error is ErrGasUintOverflow: "gas uint64 overflow",
      // In order to trigger that one, we'd have to allocate a huge chunk of data, such that the
      // multiplication len(data) * gas_per_byte overflows uint64. Not testable at the moment
      { // ErrIntrinsicGas
        txs: []*types.Transaction{
          makeTx(0, common.Address{}, big.NewInt(0), params.TxGas-1000, big.NewInt(225000000000), nil),
        },
        want: "could not apply tx 0 [0x2fc3e3b5cc26917d413e26983fe189475f47d4f0757e32aaa5561fcb9c9dc432]: intrinsic gas too low: have 20000, want 21000",
      },
    },
    { // ErrGasLimitReached
      txs: []*types.Transaction{
        makeTx(0, common.Address{}, big.NewInt(0), params.TxGas*381, big.NewInt(225000000000), nil),
      },
      want: "could not apply tx 0 [0x9ee548e001369418ae53aaa1b5d823f081cc7fa9c9a7ee71a978ae17a2aece0]: gas limit reached",
    },
    { // ErrFeeCapTooLow
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, big.NewInt(0), big.NewInt(0)),
      },
      want: "could not apply tx 0 [0xc4ab868fef0c82ae0387b742ae87907f2d0fc528fc6ea0a021459fb0fc4a4a8]: max fee per gas less than block base fee: address 0x71562b71999873085b286",
    },
    { // ErrTipVeryHigh
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, tooBigNumber, big.NewInt(1)),
      },
      want: "could not apply tx 0 [0x15b8391b9981f266b32f3ab7da564bb3d6c21628364ea9b32a21139f89f712]: max priority fee per gas higher than 2^256-1: address 0x71562b71999873085b2",
    },
    { // ErrFeeCapVeryHigh
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, big.NewInt(1), tooBigNumber),
      },
      want: "could not apply tx 0 [0x48bc299b83fdb345c57478f239e89814bb3063eb4e4b49f3b6057a69255c16bd]: max fee per gas higher than 2^256-1: address 0x71562b71999873085b286dF957",
    },
    { // ErrTipAboveFeeCap
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, big.NewInt(2), big.NewInt(1)),
      },
      want: "could not apply tx 0 [0xf987a31ff0c71895780a7612f965a0c8b056deb54e020bb44fa478092f14c9b4]: max priority fee per gas higher than max fee per gas: address 0x71562b719",
    },
    { // ErrInsufficientFunds
      // Available balance: 1000000000000000000
      // Effective cost: 18375000021000
      // FeeCap * gas: 1050000000000000000
      // This test is designed to have the effective cost be covered by the balance, but
      // the extended requirement on FeeCap*gas < balance to fail
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, big.NewInt(1), big.NewInt(1000000000000000)),
      },
      want: "could not apply tx 0 [0x3388378ed60640e75d2edf728d5528a305f599997abc4f23ec46b351b6197499]: insufficient funds for gas * price + value: address 0x71562b71999873085b2",
    },
    { // Another ErrInsufficientFunds, this one to ensure that feecap/tip of max u256 is allowed
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas, bigNumber, bigNumber),
      },
      want: "could not apply tx 0 [0xd82a0c2519acfeac9a948258c47e784acd20651d9d80f9a1c67b4137651c3a24]: insufficient funds for gas * price + value: address 0x71562b71999873085b2",
    },
  } {
    block := GenerateBadBlock(genesis, dummy.NewFaker(), tt.txs, gspect.Config)
    _, err := blockchain.InsertChain(types.Blocks(block))
    if err == nil {
      t.Fatal("block imported without errors")
    }
    if have, want := err.Error(), tt.want; have != want {
      t.Errorf("test %d:\nhave \"%v\"\\nwant \"%v\"\\n", i, have, want)
    }
  }
}

// ErrTxTypeNotSupported, For this, we need an older chain
{
  var (
    db = rawdb.NewMemoryDatabase()
    gspect = &Genesis{
      Config: &params.ChainConfig{
        ChainID: big.NewInt(1),
        HomesteadBlock: big.NewInt(0),
        EIP150Block: big.NewInt(0),
        EIP150Hash: common.Hash{},
        EIP155Block: big.NewInt(0),
        EIP158Block: big.NewInt(0),
        ByzantiumBlock: big.NewInt(0),
        ConstantinopleBlock: big.NewInt(0),
        PetersburgBlock: big.NewInt(0),
        IstanbulBlock: big.NewInt(0),
        MuirGlacierBlock: big.NewInt(0),
        ApricotPhase1BlockTimestamp: big.NewInt(0),
        ApricotPhase2BlockTimestamp: big.NewInt(0),
      },
      Alloc: GenesisAlloc{
        common.HexToAddress("0x71562b71999873085b286dF957af199Ec94617F7"): GenesisAccount{
          Balance: big.NewInt(1000000000000000000), // 1 ether
          Nonce: 0,
        },
      },
      GasLimit: params.ApricotPhase1GasLimit,
    }
    genesis = gspect.MustCommit(db)
    blockchain, _ = NewBlockChain(db, DefaultCacheConfig, gspect.Config, dummy.NewFaker(), vm.Config{}, common.Hash{})
  )
  defer blockchain.Stop()
  for i, tt := range []struct {
    txs []*types.Transaction
    want string
  } {
    { // ErrTxTypeNotSupported
      txs: []*types.Transaction{
        mkDynamicTx(0, common.Address{}, params.TxGas-1000, big.NewInt(0), big.NewInt(0)),
      },
      want: "could not apply tx 0 [0x88626ac0d53cb65308f2416103c62bb1f18b085573d4f96a3640bbbfff13c14f]: transaction type not supported",
    },
  } {
    block := GenerateBadBlock(genesis, dummy.NewFaker(), tt.txs, gspect.Config)
    _, err := blockchain.InsertChain(types.Blocks(block))
    if err == nil {

```

```

+         t.Fatal("block imported without errors")
+     }
+     if have, want := err.Error(), tt.want; have != want {
+         t.Errorf("test %d:\nhave \"%v\"\\nwant \"%v\"\\n", i, have, want)
+     }
+ }
+
+ // ErrSenderNoEOA, for this we need the sender to have contract code
+ {
+     var (
+         db      = rawdb.NewMemoryDatabase()
+         gspect = &Genesis{
+             Config: config,
+             Alloc:  GenesisAlloc{
+                 common.HexToAddress("0x71562b71999873DB5b286dF957af199Ec94617F7"): GenesisAccount{
+                     Balance: big.NewInt(1000000000000000000), // 1 ether
+                     Nonce:    0,
+                     Code:      common.FromHex("0xB0B0FACE"),
+                 },
+             },
+             GasLimit: params.ApricotPhase1GasLimit,
+         }
+         genesis = gspect.MustCommit(db)
+         blockchain, _ = NewBlockChain(db, DefaultCacheConfig, gspect.Config, dummy.NewFaker(), vm.Config{}, common.Hash{})
+     )
+     defer blockchain.Stop()
+     for i, tt := range []struct {
+         txs []*types.Transaction
+         want string
+     }{
+         { // ErrSenderNoEOA
+             txs: []*types.Transaction{
+                 mkDynamicTx(0, common.Address{}, params.TxGas-1000, big.NewInt(0), big.NewInt(0)),
+             },
+             want: "could not apply tx 0 [0x88626ac0d53cb65308f2416103c62bb1f18b805573d4f96a3640bbbf13c14f]: sender not an eoa: address 0x71562b71999873DB5b286dF957af199Ec94617F7, co",
+         },
+         {
+             block := GenerateBadBlock(genesis, dummy.NewFaker(), tt.txs, gspect.Config)
+             _, err := blockchain.InsertChain(types.Blocks{block})
+             if err == nil {
+                 t.Fatal("block imported without errors")
+             }
+             if have, want := err.Error(), tt.want; have != want {
+                 t.Errorf("test %d:\nhave \"%v\"\\nwant \"%v\"\\n", i, have, want)
+             }
+         }
+     }
+ }
+}
+
+// GenerateBadBlock constructs a "block" which contains the transactions. The transactions are not expected to be
+// valid, and no proper post-state can be made. But from the perspective of the blockchain, the block is sufficiently
+// valid to be considered for import:
+// - valid pow (fake), ancestry, difficulty, gaslimit etc
+func GenerateBadBlock(parent *types.Block, engine consensus.Engine, txs types.Transactions, config *params.ChainConfig) *types.Block {
+    header := &types.Header{
+        ParentHash: parent.Hash(),
+        Coinbase:   parent.Coinbase(),
+        Difficulty: engine.CalcDifficulty(&fakeChainReader{config}, parent.Time()+10, &types.Header{
+            Number:   parent.Number(),
+            Time:     parent.Time(),
+            Difficulty: parent.Difficulty(),
+            UncleHash: parent.UncleHash(),
+        }),
+        GasLimit:   parent.GasLimit(),
+        Number:     new(big.Int).Add(parent.Number(), common.Big1),
+        Time:       parent.Time() + 10,
+        UncleHash:  types.EmptyUncleHash,
+    }
+    if config.IsApricotPhase3(new(big.Int).SetUint64(header.Time)) {
+        header.Extra, header.BaseFee, _ = dummy.CalcBaseFee(config, parent.Header(), header.Time)
+    }
+    if config.IsApricotPhase4(new(big.Int).SetUint64(header.Time)) {
+        header.BlockGasCost = big.NewInt(0)
+        header.ExtDataGasUsed = big.NewInt(0)
+    }
+    var receipts []*types.Receipt
+    // The post-state result doesn't need to be correct (this is a bad block), but we do need something there
+    // Preferably something unique. So let's use a combo of blocknum + txhash
+    hasher := sha3.NewLegacyKeccak256()
+    hasher.Write(header.Number.Bytes())
+    var cumulativeGas uint64
+    for _, tx := range txs {
+        txh := tx.Hash()
+        hasher.Write(txh[:])
+        receipt := types.NewReceipt(nil, false, cumulativeGas+tx.Gas())
+        receipt.TxHash = tx.Hash()
+        receipt.GasUsed = tx.Gas()
+        receipts = append(receipts, receipt)
+        cumulativeGas += tx.Gas()
+    }
+    header.Root = common.BytesToHash(hasher.Sum(nil))
+    // Assemble and return the final block for sealing
+    return types.NewBlock(header, txs, nil, receipts, trie.NewStackTrie(nil), nil, true)
+}
+
+diff --git a/core/state_transition.go b/core/state_transition.go
index 381327cf..7eb1bca 100644
--- a/core/state_transition.go
+++ b/core/state_transition.go
@@ -33,10 +33,10 @@ import (
     "github.com/ethereum/go-ethereum/crypto"
 
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
 )
 
 var emptyCodeHash = crypto.Keccak256Hash(nil)
@@ -115,6 +115,23 @@ func (result *ExecutionResult) Return() []byte {
     return common.CopyBytes(result.ReturnData)
 }
 
 // Implement the EVMCaller interface on the state transition structure; simply delegate the calls
+func (st *StateTransition) Call(callder vm.ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {
+    return st.evm.Call(callder, addr, input, gas, value)
+}
+
+func (st *StateTransition) GetBlockNumber() *big.Int {
+    return st.evm.Context.BlockNumber
+}
+
+func (st *StateTransition) GetGasLimit() uint64 {
+    return st.evm.Context.GasLimit
+}
+
+func (st *StateTransition) AddBalance(addr common.Address, amount *big.Int) {
+    st.state.AddBalance(addr, amount)
+}
+

```

```

// Revert returns the concrete revert reason if the execution is aborted by `REVERT`
// opcode. Note the reason can be nil if no data supplied with revert opcode.
func (result *ExecutionResult) Revert() []byte {
@@ -232,12 +249,18 @@ func (st *StateTransition) preCheck() error {
    } else if stNonce > msgNonce {
        return fmt.Errorf("%w: address %v, tx: %d state: %d", ErrNonceTooLow,
            st.msg.From().Hex(), msgNonce, stNonce)
    } else if stNonce+1 < stNonce {
        return fmt.Errorf("%w: address %v, nonce: %d", ErrNonceMax,
            st.msg.From().Hex(), stNonce)
    }
    // Make sure the sender is an EOA
    if codeHash := st.state.GetCodeHash(st.msg.From()); codeHash != emptyCodeHash && codeHash != (common.Hash{}) {
        return fmt.Errorf("%w: address %v, codehash: %s", ErrSenderNoEOA,
            st.msg.From().Hex(), codeHash)
    }
    if st.msg.From() == st.evm.Context.Coinbase {
        return fmt.Errorf("%w: address %v", vm.ErrNoSenderBlackhole, st.msg.From())
    }
}
// Make sure that transaction gasFeeCap is greater than the baseFee (post london)
if st.evm.ChainConfig().IsApricotPhase3(st.evm.Context.Time) {
@@ -321,19 +344,69 @@ func (st *StateTransition) TransitionDb() (*ExecutionResult, error) {
    if rules := st.evm.ChainConfig().AvalancheRules(st.evm.Context.BlockNumber, st.evm.Context.Time); rules.IsApricotPhase2 {
        st.state.PrepareAccessList(msg.From(), msg.To(), vm.ActivePrecompiles(rules), msg.AccessList())
    }
}
var (
    ret []byte
    vmerr error // vm errors do not effect consensus and are therefore not assigned to err
    ret []byte
    vmerr error // vm errors do not affect consensus and are therefore not assigned to err
    chainID *big.Int
    timestamp *big.Int
    burnAddress common.Address
)
chainID = st.evm.ChainConfig().ChainID
timestamp = st.evm.Context.Time
burnAddress = st.evm.Context.Coinbase
if burnAddress != common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000000") {
    return nil, fmt.Errorf("Invalid value for block.coinbase")
}
if contractCreation {
    ret, _, st.gas, vmerr = st.evm.Create(sender, st.data, st.gas, st.value)
} else {
    // Increment the nonce for the next transaction
    st.state.SetNonce(msg.From(), st.state.GetNonce(sender.Address())+1)
    ret, st.gas, vmerr = st.evm.Call(sender, st.to(), st.data, st.gas, st.value)
    if vmerr == nil && *msg.To() == GetStateConnectorContract(chainID, timestamp) && len(st.data) >= 36 && len(ret) == 32 {
        if GetStateConnectorActivated(chainID, timestamp) &&
            bytes.Equal(st.data[0:4], SubmitAttestationSelector(chainID, timestamp)) &&
            binary.BigEndian.Uint64(ret[24:32]) > 0 {
            err = st.FinalisePreviousRound(chainID, timestamp, st.data[4:36])
            if err != nil {
                log.Warn("Error finalising state connector round", "error", err)
            }
        }
    }
}
}
st.refundGas(apricotPhase1)
- st.state.AddBalance(st.evm.Context.Coinbase, new(big.Int).Mul(new(big.Int).SetUint64(st.gasUsed()), st.gasPrice))
if vmerr == nil && msg.To() != nil && *msg.To() == common.HexToAddress(GetPrioritisedFTS0Contract(timestamp)) {
    nominalGasUsed := uint64(21000)
    nominalGasPrice := uint64(225_000_000_000)
    nominalFee := new(big.Int).Mul(new(big.Int).SetUint64(nominalGasUsed), new(big.Int).SetUint64(nominalGasPrice))
    actualGasUsed := st.gasUsed()
    actualGasPrice := st.gasPrice
    actualFee := new(big.Int).Mul(new(big.Int).SetUint64(actualGasUsed), actualGasPrice)
    if actualFee.Cmp(nominalFee) > 0 {
        feeRefund := new(big.Int).Sub(actualFee, nominalFee)
        st.state.AddBalance(st.msg.From(), feeRefund)
        st.state.AddBalance(burnAddress, nominalFee)
    } else {
        st.state.AddBalance(burnAddress, actualFee)
    }
} else {
    st.state.AddBalance(burnAddress, new(big.Int).Mul(new(big.Int).SetUint64(st.gasUsed()), st.gasPrice))
}
// Call the keeper contract trigger method if there is no vm error
if vmerr == nil {
    // Temporarily disable EVM debugging
    oldDebug := st.evm.Config.Debug
    st.evm.Config.Debug = false
    // Call the keeper contract trigger
    log := log.Root()
    triggerKeeperAndMint(st, log)
    st.evm.Config.Debug = oldDebug
}
return &ExecutionResult{
    UsedGas: st.gasUsed(),
diff --git a/core/test_blockchain.go b/core/test_blockchain.go
index 7b7a1109..1de25016 100644
--- a/core/test_blockchain.go
+++ b/core/test_blockchain.go
@@ -9,14 +9,14 @@ import (
    "strings"
    "testing"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
)
type ChainTest struct {
@@ -1334,7 +1334,7 @@ func TestGenerateChainInvalidBlockFee(t *testing.T, create func(db ethdb.Databas
    if err == nil {
        t.Fatal("should not have been able to build a block because of insufficient block fee")
    }
-    if !strings.Contains(err.Error(), "insufficient gas (0) to cover the block cost (100000)") {
+    if !strings.Contains(err.Error(), "insufficient gas (0) to cover the block cost (400000)") {
        t.Fatalf("should have gotten insufficient block fee error but got %v instead", err)
    }
}
@@ -1404,7 +1404,7 @@ func TestInsertChainInvalidBlockFee(t *testing.T, create func(db ethdb.Database,
    if err == nil {
        t.Fatal("should not have been able to build a block because of insufficient block fee")
    }
-    if !strings.Contains(err.Error(), "insufficient gas (0) to cover the block cost (100000)") {

```

```

+         if !strings.Contains(err.Error(), "insufficient gas (0) to cover the block cost (400000)") {
+             t.Fatalf("should have gotten insufficient block fee error but got %v instead", err)
+         }
+     }
@@ -1440,7 +1440,7 @@ func TestInsertChainValidBlockFee(t *testing.T, create func(db ethdb.Database, c
signer := types.LatestSigner(params.TestChainConfig)
// Generate chain of blocks using [genDB] instead of [chainDB] to avoid writing
// to the Blockchain's database while generating blocks.
-     tip := big.NewInt(2000 * params.GWei)
+     tip := big.NewInt(50000 * params.GWei)
transfer := big.NewInt(10000)
chain, _, err := GenerateChain(gspec.Config, genesis, blockchain.engine, genDB, 3, 0, func(i int, gen *BlockGen) {
    feeCap := new(big.Int).Add(gen.BaseFee(), tip)
diff --git a/core/tx_cacher.go b/core/tx_cacher.go
index 19b92806..bd71fbe4 100644
--- a/core/tx_cacher.go
+++ b/core/tx_cacher.go
@@ -27,7 +27,7 @@
package core

import (
-     "github.com/ava-labs/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/types"
)

// txSenderCacherRequest is a request for recovering transaction senders with a
diff --git a/core/tx_journal.go b/core/tx_journal.go
index b2bfa538..ee9db7ad 100644
--- a/core/tx_journal.go
+++ b/core/tx_journal.go
@@ -31,10 +31,10 @@
@@ -31,10 +31,10 @@ import (
    "io"
    "os"

-     "github.com/ava-labs/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/types"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/log"
    "github.com/ethereum/go-ethereum/rlp"
+     "github.com/flare-foundation/coreth/core/types"
)

// errNoActiveJournal is returned if a transaction is attempted to be inserted
diff --git a/core/tx_list.go b/core/tx_list.go
index 4a6999e3..da04b5d3 100644
--- a/core/tx_list.go
+++ b/core/tx_list.go
@@ -35,8 +35,8 @@
@@ -35,8 +35,8 @@ import (
    "sync/atomic"
    "time"

-     "github.com/ava-labs/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/types"
    "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/coreth/core/types"
)

// nonceHeap is a heap.Interface implementation over 64bit unsigned integers for
diff --git a/core/tx_list_test.go b/core/tx_list_test.go
index ecfa9154..74209920 100644
--- a/core/tx_list_test.go
+++ b/core/tx_list_test.go
@@ -31,8 +31,8 @@
@@ -31,8 +31,8 @@ import (
    "math/rand"
    "testing"

-     "github.com/ava-labs/coreth/core/types"
+     "github.com/flare-foundation/coreth/core/types"
    "github.com/ethereum/go-ethereum/crypto"
+     "github.com/flare-foundation/coreth/core/types"
)

// Tests that transactions can be added to strict lists and list contents and
diff --git a/core/tx_noncer.go b/core/tx_noncer.go
index 0dc31c4..c1c1e1f6 100644
--- a/core/tx_noncer.go
+++ b/core/tx_noncer.go
@@ -29,8 +29,8 @@
@@ -29,8 +29,8 @@ package core
import (
    "sync"

-     "github.com/ava-labs/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/state"
    "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/coreth/core/state"
)

// txNoncer is a tiny virtual state database to manage the executable nonces of
diff --git a/core/tx_pool.go b/core/tx_pool.go
index abc7f562..7bd7e1e3 100644
--- a/core/tx_pool.go
+++ b/core/tx_pool.go
@@ -36,15 +36,15 @@
@@ -36,15 +36,15 @@ import (
    "sync/atomic"
    "time"

-     "github.com/ava-labs/coreth/consensus/dummy"
-     "github.com/ava-labs/coreth/core/state"
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/params"
+     "github.com/flare-foundation/coreth/consensus/dummy"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/params"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/prque"
    "github.com/ethereum/go-ethereum/event"
    "github.com/ethereum/go-ethereum/log"
    "github.com/ethereum/go-ethereum/metrics"
+     "github.com/flare-foundation/coreth/consensus/dummy"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/params"
)

const (
@@ -1253,7 +1253,7 @@ func (pool *TxPool) runReorg(done chan struct{}, reset *txpoolResetRequest, dirt
if reset != nil {
    pool.demoteUnexecutables()
    if reset.newHead != nil && pool.chainconfig.IsApricotPhase3(new(big.Int).SetUint64(reset.newHead.Time)) {
-         _, baseFeeEstimate, err := dummy.CalcBaseFee(pool.chainconfig, reset.newHead, uint64(time.Now()).Unix())
+         _, baseFeeEstimate, err := dummy.CalcBaseFee(pool.chainconfig, reset.newHead, uint64(time.Now()).Unix())
+         _, baseFeeEstimate, err := dummy.EstimateNextBaseFee(pool.chainconfig, reset.newHead, uint64(time.Now()).Unix())
        if err == nil {
            pool.priced.SetBaseFee(baseFeeEstimate)
        }
    }
}
@@ -1696,7 +1696,7 @@ func (pool *TxPool) updateBaseFee() {
pool.mu.Lock()
defer pool.mu.Unlock()

-     _, baseFeeEstimate, err := dummy.CalcBaseFee(pool.chainconfig, pool.currentHead, uint64(time.Now()).Unix())
+     _, baseFeeEstimate, err := dummy.EstimateNextBaseFee(pool.chainconfig, pool.currentHead, uint64(time.Now()).Unix())
    if err == nil {
        pool.priced.SetBaseFee(baseFeeEstimate)
    } else {
diff --git a/core/tx_pool_test.go b/core/tx_pool_test.go
index 32725ea0..6ddcb230 100644
--- a/core/tx_pool_test.go
+++ b/core/tx_pool_test.go
@@ -1,3 +1,13 @@
+// (c) 2019-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original

```



```

+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2015 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
@@ -29,14 +39,14 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/trie"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/event"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/trie"
)

var (

diff --git a/core/types.go b/core/types.go
index aa8d9873..9ba59a65 100644
--- a/core/types.go
+++ b/core/types.go
@@ -27,9 +27,9 @@
package core

import (
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
)

// Validator is an interface which defines the standard for block validation. It
diff --git a/core/types/block.go b/core/types/block.go
index 58d2fc49..9b97bbf5 100644
--- a/core/types/block.go
+++ b/core/types/block.go
@@ -169,10 +169,6 @@
@@ -169,10 +169,6 @@ type Block struct {
    // caches
    hash atomic.Value
    size atomic.Value

-    // Td is used by package core to store the total difficulty
-    // of the chain up to and including the block.
-    td *big.Int
}

// "external" block encoding. used for eth protocol, etc.
@@ -195,7 +191,7 @@
@@ -195,7 +191,7 @@ func NewBlock(
    header *Header, txs []*Transaction, uncles []*Header, receipts []*Receipt,
    hasher TrieHasher, extdata []byte, recalc bool,
) *Block {
-    b := &Block{header: CopyHeader(header), td: new(big.Int)}
+    b := &Block{header: CopyHeader(header)}

    // TODO: panic if len(txs) != len(receipts)
    if len(txs) == 0 {
diff --git a/core/types/block_test.go b/core/types/block_test.go
index 892e5c4c..f00396ad 100644
--- a/core/types/block_test.go
+++ b/core/types/block_test.go
@@ -33,11 +33,11 @@
@@ -33,11 +33,11 @@ import (
    "reflect"
    "testing"

-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/math"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/params"
+    "golang.org/x/crypto/sha3"
)

diff --git a/core/types/dynamic_fee_tx.go b/core/types/dynamic_fee_tx.go
index 7e288210..c4ec28c5 100644
--- a/core/types/dynamic_fee_tx.go
+++ b/core/types/dynamic_fee_tx.go
@@ -35,8 +35,8 @@
@@ -35,8 +35,8 @@ import (
type DynamicFeeTx struct {
    ChainID      *big.Int
    Nonce        uint64
-    GasTipCap    *big.Int
-    GasFeeCap    *big.Int
+    GasTipCap    *big.Int // a.k.a. maxPriorityFeePerGas
+    GasFeeCap    *big.Int // a.k.a. maxFeePerGas
    Gas          uint64
    To            *common.Address `rlp:"nil"` // nil means contract creation
    Value         *big.Int
diff --git a/core/types/hashing_test.go b/core/types/hashing_test.go
index 60f9da10..004fa888 100644
--- a/core/types/hashing_test.go
+++ b/core/types/hashing_test.go
@@ -34,12 +34,12 @@
@@ -34,12 +34,12 @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/trie"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/trie"
)

func TestDeriveSha(t *testing.T) {
diff --git a/core/types/receipt.go b/core/types/receipt.go
index 176cb109..885722d2 100644
--- a/core/types/receipt.go
+++ b/core/types/receipt.go
@@ -34,11 +34,11 @@
@@ -34,11 +34,11 @@ import (
    "math/big"
    "unsafe"

-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
-    "github.com/ethereum/go-ethereum/crypto"

```

```

        "github.com/ethereum/go-ethereum/rlp"
+       "github.com/flare-foundation/coreth/params"
    )

    //go:generate gencodec -type Receipt -field-override receiptMarshaling -out gen_receipt_json.go
diff --git a/core/types/receipt_test.go b/core/types/receipt_test.go
index d0c1553e..a52f6555 100644
--- a/core/types/receipt_test.go
+++ b/core/types/receipt_test.go
@@ -33,10 +33,10 @@ import (
    "reflect"
    "testing"

-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/params"
)

var (
diff --git a/core/types/transaction_signing.go b/core/types/transaction_signing.go
index a717749b..eb0e1d8d 100644
--- a/core/types/transaction_signing.go
+++ b/core/types/transaction_signing.go
@@ -32,9 +32,9 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/params"
)

var ErrInvalidChainId = errors.New("invalid chain id for signer")
diff --git a/core/vm/access_list_tracer.go b/core/vm/access_list_tracer.go
index fb67897e..c062a406 100644
--- a/core/vm/access_list_tracer.go
+++ b/core/vm/access_list_tracer.go
@@ -30,8 +30,8 @@ import (
    "math/big"
    "time"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

// accessList is an accumulator for the set of accounts and storage slots an EVM
diff --git a/core/vm/analysis.go b/core/vm/analysis.go
index 752e27dd..cfbf0e7f 100644
--- a/core/vm/analysis.go
+++ b/core/vm/analysis.go
@@ -27,12 +27,12 @@ package vm

const (
-    set2BitsMask = uint16(0b1100_0000_0000_0000)
-    set3BitsMask = uint16(0b1110_0000_0000_0000)
-    set4BitsMask = uint16(0b1111_0000_0000_0000)
-    set5BitsMask = uint16(0b1111_1000_0000_0000)
-    set6BitsMask = uint16(0b1111_1100_0000_0000)
-    set7BitsMask = uint16(0b1111_1110_0000_0000)
+    set2BitsMask = uint16(0b11)
+    set3BitsMask = uint16(0b111)
+    set4BitsMask = uint16(0b1111)
+    set5BitsMask = uint16(0b1_1111)
+    set6BitsMask = uint16(0b11_1111)
+    set7BitsMask = uint16(0b111_1111)
)

// bitvec is a bit vector which maps bytes in a program.
@@ -40,32 +40,32 @@ @@ -40,32 +40,26 @@ const (
// it's data (i.e. argument of PUSHxx).
type bitvec []byte

-var lookup = [8]byte{
-    0x80, 0x40, 0x20, 0x10, 0x8, 0x4, 0x2, 0x1,
-}
-
func (bits bitvec) set1(pos uint64) {
-    bits[pos/8] |= lookup[pos%8]
+    bits[pos/8] |= 1 << (pos % 8)
}

func (bits bitvec) setN(flag uint16, pos uint64) {
-    a := flag >> (pos % 8)
-    bits[pos/8] |= byte(a >> 8)
-    if b := byte(a); b != 0 {
-        // If the bit-setting affects the neighbouring byte, we can assign - no need to OR it,
-        // since it's the first write to that byte
+    a := flag << (pos % 8)
+    bits[pos/8] |= byte(a)
+    if b := byte(a >> 8); b != 0 {
+        bits[pos/8+1] = b
    }
}

func (bits bitvec) set8(pos uint64) {
-    a := byte(0xFF >> (pos % 8))
+    a := byte(0xFF << (pos % 8))
+    bits[pos/8] |= a
+    bits[pos/8+1] = ^a
}

func (bits bitvec) set16(pos uint64) {
-    a := byte(0xFF >> (pos % 8))
+    a := byte(0xFF << (pos % 8))
+    bits[pos/8] |= a
+    bits[pos/8+1] = 0xFF
+    bits[pos/8+2] = ^a
@@ -73,7 +67,7 @@ @@ -73,7 +67,7 @@ func (bits bitvec) set16(pos uint64) {

// codeSegment checks if the position is in a code segment.
func (bits *bitvec) codeSegment(pos uint64) bool {
-    return ((*bits)[pos/8] & (0x80 >> (pos % 8))) == 0
+    return ((*bits)[pos/8] >> (pos % 8) & 1) == 0
}

// codeBitmap collects data locations in code.
diff --git a/core/vm/analysis_test.go b/core/vm/analysis_test.go
index 96fa0377..ae6bef9b 100644
--- a/core/vm/analysis_test.go
+++ b/core/vm/analysis_test.go
@@ -27,6 +27,6 @@ package vm

import (
+    "math/bits"
    "testing"

    "github.com/ethereum/go-ethereum/crypto"

```

```

@@ -38,24 +39,27 @@ func TestJumpDestAnalysis(t *testing.T) {
    exp byte
    which int
}

-    {[[]byte{byte(PUSH1), 0x01, 0x01, 0x01}, 0x40, 0},
-    {[[]byte{byte(PUSH1), byte(PUSH1), byte(PUSH1)}, 0x50, 0},
-    {[[]byte{byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), 0x01, 0x01, 0x01}, 0x7F, 0},
-    {[[]byte{byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x80, 1},
-    {[[]byte{0x01, 0x01, 0x01, 0x01, 0x01, byte(PUSH2), byte(PUSH2), byte(PUSH2), 0x01, 0x01, 0x01}, 0x03, 0},
-    {[[]byte{0x01, 0x01, 0x01, 0x01, 0x01, byte(PUSH2), 0x01, 0x01, 0x01, 0x01, 0x01}, 0x00, 1},
-    {[[]byte{byte(PUSH3), 0x01, 0x01, 0x01, byte(PUSH1), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0x74, 0},
-    {[[]byte{byte(PUSH3), 0x01, 0x01, 0x01, byte(PUSH1), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0x00, 1},
-    {[[]byte{0x01, byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0x3F, 0},
-    {[[]byte{0x01, byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0xC0, 1},
-    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0x7F, 0},
-    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0xFF, 1},
-    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0x80, 2},
-    {[[]byte{byte(PUSH8), 0x01, 0x02, 0x03, 0x04, 0x05, 0x06, 0x07, 0x08, byte(PUSH1), 0x01}, 0x7f, 0},
-    {[[]byte{byte(PUSH8), 0x01, 0x02, 0x03, 0x04, 0x05, 0x06, 0x07, 0x08, byte(PUSH1), 0x01}, 0xA0, 1},
-    {[[]byte{byte(PUSH32)}, 0x7F, 0},
-    {[[]byte{byte(PUSH32)}, 0xFF, 1},
-    {[[]byte{byte(PUSH32)}, 0xFF, 2},
+    {[[]byte{byte(PUSH1), 0x01, 0x01, 0x01}, 0b0000_0010, 0},
+    {[[]byte{byte(PUSH1), byte(PUSH1), byte(PUSH1)}, 0b0000_1010, 0},
+    {[[]byte{0x00, byte(PUSH1), 0x00, byte(PUSH1), 0x00, byte(PUSH1)}, 0b0101_0100, 0},
+    {[[]byte{byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), byte(PUSH8), 0x01, 0x01, 0x01}, bits.Reverse8(0x7F), 0},
+    {[[]byte{byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b0000_0001, 1},
+    {[[]byte{0x01, 0x01, 0x01, 0x01, 0x01, byte(PUSH2), byte(PUSH2), 0x01, 0x01, 0x01}, 0b1100_0000, 0},
+    {[[]byte{0x01, 0x01, 0x01, 0x01, byte(PUSH2), 0x01, 0x01, 0x01, 0x01}, 0b0000_0000, 1},
+    {[[]byte{byte(PUSH3), 0x01, 0x01, 0x01, byte(PUSH1), 0x01, 0x01, 0x01, 0x01, 0x01}, 0b0010_1110, 0},
+    {[[]byte{byte(PUSH3), 0x01, 0x01, 0x01, byte(PUSH1), 0x01, 0x01, 0x01, 0x01, 0x01}, 0b0000_0000, 1},
+    {[[]byte{0x01, byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b1111_1100, 0},
+    {[[]byte{0x01, byte(PUSH8), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b0000_0011, 1},
+    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b1111_1110, 0},
+    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b1111_1111, 1},
+    {[[]byte{byte(PUSH16), 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01, 0x01}, 0b0000_0001, 2},
+    {[[]byte{byte(PUSH8), 0x01, 0x02, 0x03, 0x04, 0x05, 0x06, 0x07, 0x08, byte(PUSH1), 0x01}, 0b1111_1110, 0},
+    {[[]byte{byte(PUSH8), 0x01, 0x02, 0x03, 0x04, 0x05, 0x06, 0x07, 0x08, byte(PUSH1), 0x01}, 0b0000_0101, 1},
+    {[[]byte{byte(PUSH32)}, 0b1111_1110, 0},
+    {[[]byte{byte(PUSH32)}, 0b1111_1111, 1},
+    {[[]byte{byte(PUSH32)}, 0b1111_1111, 2},
+    {[[]byte{byte(PUSH32)}, 0b1111_1111, 3},
+    {[[]byte{byte(PUSH32)}, 0b0000_0001, 4},
+    }
    for i, test := range tests {
        ret := codeBitmap(test.code)
diff --git a/core/vm/contract.go b/core/vm/contract.go
index c7e4b53a..78e9fcd8 100644
--- a/core/vm/contract.go
+++ b/core/vm/contract.go
@@ -153,16 +153,11 @@ func (c *Contract) AsDelegate() *Contract {

    // GetOp returns the n'th element in the contract's byte array
    func (c *Contract) GetOp(n uint64) OpCode {
-        return OpCode(c.GetByte(n))
-    }
-
-    // GetByte returns the n'th byte in the contract's byte array
-    func (c *Contract) GetByte(n uint64) byte {
        if n < uint64(len(c.Code)) {
-            return c.Code[n]
+            return OpCode(c.Code[n])
        }

-        return 0
+        return STOP
    }

    // Caller returns the caller of the contract.
diff --git a/core/vm/contracts.go b/core/vm/contracts.go
index cfb26714..40880420 100644
--- a/core/vm/contracts.go
+++ b/core/vm/contracts.go
@@ -32,13 +32,13 @@ import (
    "errors"
    "math/big"

-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/math"
-    "github.com/ethereum/go-ethereum/crypto"
-    "github.com/ethereum/go-ethereum/crypto/blake2b"
-    "github.com/ethereum/go-ethereum/crypto/bls12381"
-    "github.com/ethereum/go-ethereum/crypto/bn256"
+    "github.com/flare-foundation/coreth/params"

    //lint:ignore SA1019 Needed for precompile
    "golang.org/x/crypto/ripemd160"
@@ -101,8 +101,8 @@ var PrecompiledContractsApricotPhase2 = map[common.Address]StatefulPrecompiledCo
common.BytesToAddress([]byte{8}): new(WrappedPrecompiledContract(&bn256PairingIstanbul{})),
common.BytesToAddress([]byte{9}): new(WrappedPrecompiledContract(&blake2F{})),
genesisContractAddr: &deprecatedContract{},
-    nativeAssetBalanceAddr: &nativeAssetBalance{gasCost: params.AssetBalanceApricot},
-    nativeAssetCallAddr: &nativeAssetCall{gasCost: params.AssetCallApricot},
+    NativeAssetBalanceAddr: &nativeAssetBalance{gasCost: params.AssetBalanceApricot},
+    NativeAssetCallAddr: &nativeAssetCall{gasCost: params.AssetCallApricot},
}

var (
diff --git a/core/vm/contracts_stateful.go b/core/vm/contracts_stateful.go
index 981ebd72..bffc33a9 100644
--- a/core/vm/contracts_stateful.go
+++ b/core/vm/contracts_stateful.go
@@ -7,8 +7,8 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/holiman/uint256"
)

@@ -20,8 +20,8 @@ import (

var (
    genesisContractAddr = common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000000")
-    nativeAssetBalanceAddr = common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000001")
-    nativeAssetCallAddr = common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000002")
+    NativeAssetBalanceAddr = common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000001")
+    NativeAssetCallAddr = common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000002")
)

// StatefulPrecompiledContract is the interface for executing a precompiled contract
@@ -54,6 +54,6 @@ type nativeAssetBalance struct {
    gasCost uint64
}

+// PackNativeAssetBalanceInput packs the arguments into the required input data for a transaction to be passed into
+// the native asset balance precompile.
func PackNativeAssetBalanceInput(address common.Address, assetID common.Hash) []byte {
    input := make([]byte, 52)
    copy(input, address.Bytes())
@@ -61,6 +61,6 @@ func PackNativeAssetBalanceInput(address common.Address, assetID common.Hash) []

```

```

    return input
}

+// UnpackNativeAssetBalanceInput attempts to unpack [input] into the arguments to the native asset balance precompile
func UnpackNativeAssetBalanceInput(input []byte) (common.Address, common.Hash, error) {
    if len(input) != 52 {
        return common.Address{}, common.Hash{}, fmt.Errorf("native asset balance input had unexpcted length %d", len(input))
    }
}

-97,6 +100,9 @@ type nativeAssetCall struct {
    gasCost uint64
}

+// PackNativeAssetCallInput packs the arguments into the required input data for a transaction to be passed into
+// the native asset precompile.
+// Assumes that [assetAmount] is non-nil.
func PackNativeAssetCallInput(address common.Address, assetID common.Hash, assetAmount *big.Int, callData []byte) []byte {
    input := make([]byte, 84+len(callData))
    copy(input[0:20], address.Bytes())
}

-106,13 +112,13 @@ func PackNativeAssetCallInput(address common.Address, assetID common.Hash, asset
    return input
}

-func UnpackNativeAssetCallInput(input []byte) (common.Address, *common.Hash, *big.Int, []byte, error) {
+// UnpackNativeAssetCallInput attempts to unpack [input] into the arguments to the native asset call precompile
+func UnpackNativeAssetCallInput(input []byte) (common.Address, common.Hash, *big.Int, []byte, error) {
    if len(input) < 84 {
        return common.Address{}, nil, nil, nil, fmt.Errorf("native asset call input had unexpcted length %d", len(input))
    }
    return common.Address{}, common.Hash{}, nil, nil, fmt.Errorf("native asset call input had unexpected length %d", len(input))
}

-    to := common.BytesToAddress(input[:20])
-    assetID := new(common.Hash)
-    assetID.SetBytes(input[20:52])
+    assetID := common.BytesToHash(input[20:52])
+    assetAmount := new(big.Int).SetBytes(input[52:84])
+    callData := input[84:]
    return to, assetID, assetAmount, callData, nil
}

-135,6 +141,8 @@ func (c *nativeAssetCall) Run(evm *EVM, caller ContractRef, addr common.Address,
    return nil, remainingGas, ErrExecutionReverted
}

+    // Note: it is not possible for a negative assetAmount to be passed in here due to the fact that decoding a
+    // byte slice into a *big.Int type will always return a positive value.
+    if assetAmount.Sign() != 0 && !evm.Context.CanTransferMC(evm.StateDB, caller.Address(), to, assetID, assetAmount) {
        return nil, remainingGas, ErrInsufficientBalance
    }
}

-174,6 +182,6 @@ func (c *nativeAssetCall) Run(evm *EVM, caller ContractRef, addr common.Address,
    type deprecatedContract struct{}

-func (c *deprecatedContract) Run(evm *EVM, caller ContractRef, addr common.Address, input []byte, suppliedGas uint64, readOnly bool) (ret []byte, remainingGas uint64, err error) {
+func (c *deprecatedContract) Run(evm *EVM, caller ContractRef, addr common.Address, input []byte, suppliedGas uint64, readOnly bool) (ret []byte, remainingGas uint64, err error) {
    return nil, suppliedGas, ErrExecutionReverted
}

diff --git a/core/vm/contracts_stateful_test.go b/core/vm/contracts_stateful_test.go
index 6126078f..d2d2df10 100644
--- a/core/vm/contracts_stateful_test.go
+++ b/core/vm/contracts_stateful_test.go
@@ -7,11 +7,11 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/stretchr/testify/assert"
)

@@ -36,16 +36,9 @@ func CanTransfer(db StateDB, addr common.Address, amount *big.Int) bool {
    return db.GetBalance(addr).Cmp(amount) >= 0
}

-func CanTransferMC(db StateDB, addr common.Address, to common.Address, coinID *common.Hash, amount *big.Int) bool {
+func CanTransferMC(db StateDB, addr common.Address, to common.Address, coinID common.Hash, amount *big.Int) bool {
    log.Info("CanTransferMC", "address", addr, "to", to, "coinID", coinID, "amount", amount)
    if coinID == nil {
        return true
    }
    if db.GetBalanceMultiCoin(addr, *coinID).Cmp(amount) >= 0 {
        return true
    }
    // insufficient balance
    return false
}

+    return db.GetBalanceMultiCoin(addr, coinID).Cmp(amount) >= 0
}

// Transfer subtracts amount from sender and adds amount to recipient using the given Db
@@ -55,12 +55,12 @@ func Transfer(db StateDB, sender, recipient common.Address, amount *big.Int) {
}

// Transfer subtracts amount from sender and adds amount to recipient using the given Db
-func TransferMultiCoin(db StateDB, sender, recipient common.Address, coinID *common.Hash, amount *big.Int) {
-    if coinID == nil {
-        return
-    }
-    db.SubBalanceMultiCoin(sender, *coinID, amount)
-    db.AddBalanceMultiCoin(recipient, *coinID, amount)
+func TransferMultiCoin(db StateDB, sender, recipient common.Address, coinID common.Hash, amount *big.Int) {
+    db.SubBalanceMultiCoin(sender, coinID, amount)
+    db.AddBalanceMultiCoin(recipient, coinID, amount)
}

func TestPackNativeAssetCallInput(t *testing.T) {
@@ -74,7 +74,7 @@ func TestPackNativeAssetCallInput(t *testing.T) {
    unpackedAddr, unpackedAssetID, unpackedAssetAmount, unpackedCallData, err := UnpackNativeAssetCallInput(input)
    assert.NoError(t, err)
    assert.Equal(t, addr, unpackedAddr, "address")
    assert.Equal(t, &assetID, unpackedAssetID, "assetID")
+    assert.Equal(t, assetID, unpackedAssetID, "assetID")
    assert.Equal(t, assetAmount, unpackedAssetAmount, "assetAmount")
    assert.Equal(t, callData, unpackedCallData, "callData")
}

@@ -131,7 +131,7 @@ func TestStatefulPrecompile(t *testing.T) {
    return statedb
    },
    from:        userAddr1,
-    precompileAddr: nativeAssetBalanceAddr,
+    precompileAddr: NativeAssetBalanceAddr,
    input:       PackNativeAssetBalanceInput(userAddr1, assetID),
    value:       big0,
    gasInput:    params.AssetBalanceApricot,
}

@@ -157,7 +157,7 @@ func TestStatefulPrecompile(t *testing.T) {
    return statedb
    },
    from:        userAddr1,
-    precompileAddr: nativeAssetBalanceAddr,
+    precompileAddr: NativeAssetBalanceAddr,
    input:       PackNativeAssetBalanceInput(userAddr1, assetID),
    value:       big0,
    gasInput:    params.AssetBalanceApricot,
}

@@ -182,7 +182,7 @@ func TestStatefulPrecompile(t *testing.T) {
}

```

```

        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetBalanceAddr,
+   precompileAddr:      NativeAssetBalanceAddr,
    input:                PackNativeAssetBalanceInput(userAddr1, assetID),
    value:                big0,
    gasInput:             params.AssetBalanceApricot,
@@ -200,7 +190,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetBalanceAddr,
+   precompileAddr:      NativeAssetBalanceAddr,
    input:                nil,
    value:                big0,
    gasInput:             params.AssetBalanceApricot,
@@ -218,7 +208,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetBalanceAddr,
+   precompileAddr:      NativeAssetBalanceAddr,
    input:                PackNativeAssetBalanceInput(userAddr1, assetID),
    value:                big0,
    gasInput:             params.AssetBalanceApricot - 1,
@@ -236,7 +226,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetBalanceAddr,
+   precompileAddr:      NativeAssetBalanceAddr,
    input:                PackNativeAssetBalanceInput(userAddr1, assetID),
    value:                bigHundred,
    gasInput:             params.AssetBalanceApricot,
@@ -257,7 +247,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(50), nil),
    value:                big0,
    gasInput:             params.AssetCallApricot + params.CallNewAccountGas,
@@ -290,7 +280,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(50), nil),
    value:                big.NewInt(49),
    gasInput:             params.AssetCallApricot + params.CallNewAccountGas,
@@ -301,7 +291,7 @@ func TestStatefulPrecompile(t *testing.T) {
    stateDBCheck: func(t *testing.T, stateDB StateDB) {
        user1Balance := stateDB.GetBalance(userAddr1)
        user2Balance := stateDB.GetBalance(userAddr2)
        nativeAssetCallAddrBalance := stateDB.GetBalance(nativeAssetCallAddr)
+       nativeAssetCallAddrBalance := stateDB.GetBalance(NativeAssetCallAddr)
        user1AssetBalance := stateDB.GetBalanceMultiCoin(userAddr1, assetID)
        user2AssetBalance := stateDB.GetBalanceMultiCoin(userAddr2, assetID)
        expectedBalance := big.NewInt(50)
@@ -325,7 +315,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(51), nil),
    value:                big.NewInt(50),
    gasInput:             params.AssetCallApricot,
@@ -357,7 +347,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(50), nil),
    value:                big.NewInt(51),
    gasInput:             params.AssetCallApricot,
@@ -389,7 +379,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(50), nil),
    value:                big.NewInt(50),
    gasInput:             params.AssetCallApricot - 1,
@@ -410,7 +400,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                PackNativeAssetCallInput(userAddr2, assetID, big.NewInt(50), nil),
    value:                big.NewInt(50),
    gasInput:             params.AssetCallApricot + params.CallNewAccountGas - 1,
@@ -442,7 +432,7 @@ func TestStatefulPrecompile(t *testing.T) {
        return statedb
    },
    from:                userAddr1,
    precompileAddr:      nativeAssetCallAddr,
+   precompileAddr:      NativeAssetCallAddr,
    input:                make([]byte, 24),
    value:                big.NewInt(50),
    gasInput:             params.AssetCallApricot + params.CallNewAccountGas,
diff --git a/core/vm/eips.go b/core/vm/eips.go
index e525a73c..8d377dcc 100644
--- a/core/vm/eips.go
+++ b/core/vm/eips.go
@@ -30,7 +30,7 @@ import (
    "fmt"
    "sort"

    "github.com/ava-labs/coreth/params"
+   "github.com/flare-foundation/coreth/params"
    "github.com/holiman/uint256"
)

@@ -166,8 +166,8 @@ func enableAPI(jt *JumpTable) {
}

func enableAP2(jt *JumpTable) {
-   jt[BALANCEMC] = nil
-   jt[CALLEX] = nil
+   jt[BALANCEMC] = &operation{execute: opUndefined, maxStack: maxStack(0, 0)}
+   jt[CALLEX] = &operation{execute: opUndefined, maxStack: maxStack(0, 0)}
}

// enable3198 applies EIP-3198 (BASEFEE Opcode)
diff --git a/core/vm/errors.go b/core/vm/errors.go
index 7f777dc3..5c81bb28 100644
--- a/core/vm/errors.go

```

```

++ b/core/vm/errors.go
@@ -45,6 +45,12 @@ var (
    ErrReturnDataOutOfBounds = errors.New("return data out of bounds")
    ErrGasUintOverflow       = errors.New("gas uint64 overflow")
    ErrInvalidCode           = errors.New("invalid code: must not begin with 0xef")
+   ErrNonceUintOverflow     = errors.New("nonce uint64 overflow")
+   ErrNoSenderBlackhole    = errors.New("blackhole address cannot be used as sender")
+
+   // errStopToken is an internal token indicating interpreter loop termination,
+   // never returned to outside callers.
+   errStopToken = errors.New("stop token")
)

// ErrStackUnderflow wraps an evm error when the items on the stack less
diff --git a/core/vm/evm.go b/core/vm/evm.go
index 823f1013..f5a20bbb 100644
--- a/core/vm/evm.go
+++ b/core/vm/evm.go
@@ -31,9 +31,9 @@ import (
    "sync/atomic"
    "time"

-   "github.com/ava-labs/coreth/params"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/holiman/uint256"
)

@@ -44,10 +44,10 @@ var emptyCodeHash = crypto.Keccak256Hash(nil)
type (
    // CanTransferFunc is the signature of a transfer guard function
    CanTransferFunc func(StateDB, common.Address, *big.Int) bool
-   CanTransferMCFunc func(StateDB, common.Address, common.Address, *common.Hash, *big.Int) bool
+   CanTransferMCFunc func(StateDB, common.Address, common.Address, common.Hash, *big.Int) bool
+   CanTransferMCFunc func(StateDB, common.Address, common.Address, common.Hash, *big.Int) bool
    // TransferFunc is the signature of a transfer function
    TransferFunc func(StateDB, common.Address, common.Address, *big.Int)
-   TransferMCFunc func(StateDB, common.Address, common.Address, *common.Hash, *big.Int)
+   TransferMCFunc func(StateDB, common.Address, common.Address, common.Hash, *big.Int)
+   TransferMCFunc func(StateDB, common.Address, common.Address, common.Hash, *big.Int)
    // GetHashFunc returns the n'th block hash in the blockchain
    // and is used by the BLOCKHASH EVM op code.
    GetHashFunc func(uint64) common.Hash
@@ -182,14 +182,14 @@ func (evm *EVM) Interpreter() *EVMInterpreter {
// the necessary steps to create accounts and reverses the state in case of an
// execution error or failed value transfer.
func (evm *EVM) Call(caller ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {
-   if evm.Config.NoRecursion && evm.depth > 0 {
-       return nil, gas, nil
-   }
+   // Fail if we're trying to execute above the call depth limit
+   if evm.depth > int(params.CallCreateDepth) {
+       return nil, gas, ErrDepth
+   }
+   // Fail if we're trying to transfer more than the available balance
+   // Note: it is not possible for a negative value to be passed in here due to the fact
+   // that [value] will be popped from the stack and decoded to a *big.Int, which will
+   // always yield a positive result.
+   if value.Sign() != 0 && !evm.Context.CanTransfer(evm.StateDB, caller.Address(), value) {
+       return nil, gas, ErrInsufficientBalance
+   }
}

@@ -264,16 +264,16 @@ func (evm *EVM) Call(caller ContractRef, addr common.Address, input []byte, gas
}

// This allows the user transfer balance of a specified coinID in addition to a normal Call().
-func (evm *EVM) CallExpert(caller ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int, coinID *common.Hash, value2 *big.Int) (ret []byte, leftOverGas uint64, err error) {
-   if evm.Config.NoRecursion && evm.depth > 0 {
-       return nil, gas, nil
-   }
+func (evm *EVM) CallExpert(caller ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int, coinID common.Hash, value2 *big.Int) (ret []byte, leftOverGas uint64, err error) {
+   // Fail if we're trying to execute above the call depth limit
+   if evm.depth > int(params.CallCreateDepth) {
+       return nil, gas, ErrDepth
+   }
+   // Fail if we're trying to transfer more than the available balance
+   // Note: it is not possible for a negative value to be passed in here due to the fact
+   // that [value] will be popped from the stack and decoded to a *big.Int, which will
+   // always yield a positive result.
+   if value.Sign() != 0 && !evm.Context.CanTransfer(evm.StateDB, caller.Address(), value) {
+       return nil, gas, ErrInsufficientBalance
+   }
}

@@ -348,9 +348,9 @@ func (evm *EVM) CallExpert(caller ContractRef, addr common.Address, input []byte
// CallCode differs from Call in the sense that it executes the given address'
// code with the caller as context.
func (evm *EVM) CallCode(caller ContractRef, addr common.Address, input []byte, gas uint64, value *big.Int) (ret []byte, leftOverGas uint64, err error) {
-   if evm.Config.NoRecursion && evm.depth > 0 {
-       return nil, gas, nil
-   }
+   // Fail if we're trying to execute above the call depth limit
+   if evm.depth > int(params.CallCreateDepth) {
+       return nil, gas, ErrDepth
+   }
}

@@ -359,6 +359,9 @@ func (evm *EVM) CallCode(caller ContractRef, addr common.Address, input []byte,
// Note although it's noop to transfer X ether to caller itself. But
// if caller doesn't have enough balance, it would be an error to allow
// over-charging itself. So the check here is necessary.
+   // Note: it is not possible for a negative value to be passed in here due to the fact
+   // that [value] will be popped from the stack and decoded to a *big.Int, which will
+   // always yield a positive result.
+   if !evm.Context.CanTransfer(evm.StateDB, caller.Address(), value) {
+       return nil, gas, ErrInsufficientBalance
+   }
}

@@ -399,9 +399,9 @@ func (evm *EVM) CallCode(caller ContractRef, addr common.Address, input []byte,
// DelegateCall differs from CallCode in the sense that it executes the given address'
// code with the caller as context and the caller is set to the caller of the caller.
func (evm *EVM) DelegateCall(caller ContractRef, addr common.Address, input []byte, gas uint64) (ret []byte, leftOverGas uint64, err error) {
-   if evm.Config.NoRecursion && evm.depth > 0 {
-       return nil, gas, nil
-   }
+   // Fail if we're trying to execute above the call depth limit
+   if evm.depth > int(params.CallCreateDepth) {
+       return nil, gas, ErrDepth
+   }
}

@@ -441,9 +441,9 @@ func (evm *EVM) DelegateCall(caller ContractRef, addr common.Address, input []by
// Opcodes that attempt to perform such modifications will result in exceptions
// instead of performing the modifications.
func (evm *EVM) StaticCall(caller ContractRef, addr common.Address, input []byte, gas uint64) (ret []byte, leftOverGas uint64, err error) {
-   if evm.Config.NoRecursion && evm.depth > 0 {
-       return nil, gas, nil
-   }
+   // Fail if we're trying to execute above the call depth limit
+   if evm.depth > int(params.CallCreateDepth) {
+       return nil, gas, ErrDepth
+   }
}

@@ -514,10 +514,10 @@ func (evm *EVM) create(caller ContractRef, codeAndHash *codeAndHash, gas uint64,
if evm.depth > int(params.CallCreateDepth) {
    return nil, common.Address{}, gas, ErrDepth
}
+   // Note: it is not possible for a negative value to be passed in here due to the fact
+   // that [value] will be popped from the stack and decoded to a *big.Int, which will
+   // always yield a positive result.
+   if !evm.Context.CanTransfer(evm.StateDB, caller.Address(), value) {
+       return nil, common.Address{}, gas, ErrInsufficientBalance
+   }
}
+   // If there is any collision with the Blackhole address, return an error instead
+   // of allowing the contract to be created.
+   if address == evm.Context.Coinbase {

```

```

return nil, common.Address{}, gas, ErrNoSenderBlackhole
+ }
+ nonce := evm.StateDB.GetNonce(caller.Address())
+ if nonce+1 < nonce {
+     return nil, common.Address{}, gas, ErrNonceUintOverflow
+ }
+ evm.StateDB.SetNonce(caller.Address(), nonce+1)
+ // We add this to the access list before taking a snapshot. Even if the creation fails,
+ // the access-list change should not be rolled back
@@ -542,10 +547,6 @@ func (evm *EVM) create(caller ContractRef, codeAndHash *codeAndHash, gas uint64,
contract := NewContract(caller, AccountRef(address), value, gas)
contract.SetCodeOptionalHash(&address, codeAndHash)

- if evm.Config.NoRecursion && evm.depth > 0 {
-     return nil, address, gas, nil
- }
-
if evm.Config.Debug {
    if evm.depth == 0 {
        evm.Config.Tracer.CaptureStart(evm, caller.Address(), address, true, codeAndHash.code, gas, value)
@@ -609,7 +610,7 @@ func (evm *EVM) Create(caller ContractRef, code []byte, gas uint64, value *big.I
// Create2 creates a new contract using code as deployment code.
//
-// The different between Create2 with Create is Create2 uses sha3(0xff ++ msg.sender ++ salt ++ sha3(init_code))[12:]
+// The different between Create2 with Create is Create2 uses keccak256(0xff ++ msg.sender ++ salt ++ keccak256(init_code))[12:]
// instead of the usual sender-and-nonce-hash as the address where the contract is initialized at.
func (evm *EVM) Create2(caller ContractRef, code []byte, gas uint64, endowment *big.Int, salt *uint256.Int) (ret []byte, contractAddr common.Address, leftOverGas uint64, err error) {
diff --git a/core/vm/gas_table.go b/core/vm/gas_table.go
index fdbaf1eb..0d80c5c0 100644
--- a/core/vm/gas_table.go
+++ b/core/vm/gas_table.go
@@ -29,9 +29,9 @@ package vm
import (
    "errors"

-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/flare-foundation/coreth/params"
+
)

// memoryGasCost calculates the quadratic gas for memory expansion. It does so
@@ -292,7 +292,7 @@ func makeGasLog(n uint64) gasFunc {
}

-func gasSha3(evm *EVM, contract *Contract, stack *Stack, mem *Memory, memorySize uint64) (uint64, error) {
+func gasKeccak256(evm *EVM, contract *Contract, stack *Stack, mem *Memory, memorySize uint64) (uint64, error) {
    gas, err := memoryGasCost(mem, memorySize)
    if err != nil {
        return 0, err
@@ -301,7 +301,7 @@ func gasSha3(evm *EVM, contract *Contract, stack *Stack, mem *Memory, memorySize
    if overflow {
        return 0, ErrGasUintOverflow
    }
-    if wordGas, overflow = math.SafeMul(toWordSize(wordGas), params.Sha3WordGas); overflow {
+    if wordGas, overflow = math.SafeMul(toWordSize(wordGas), params.Keccak256WordGas); overflow {
        return 0, ErrGasUintOverflow
    }
    if gas, overflow = math.SafeAdd(gas, wordGas); overflow {
@@ -335,7 +335,7 @@ func gasCreate2(evm *EVM, contract *Contract, stack *Stack, mem *Memory, memoryS
    if overflow {
        return 0, ErrGasUintOverflow
    }
-    if wordGas, overflow = math.SafeMul(toWordSize(wordGas), params.Sha3WordGas); overflow {
+    if wordGas, overflow = math.SafeMul(toWordSize(wordGas), params.Keccak256WordGas); overflow {
        return 0, ErrGasUintOverflow
    }
    if gas, overflow = math.SafeAdd(gas, wordGas); overflow {
diff --git a/core/vm/gas_table_test.go b/core/vm/gas_table_test.go
index 92d5d301..fe6a928b 100644
--- a/core/vm/gas_table_test.go
+++ b/core/vm/gas_table_test.go
@@ -31,11 +31,11 @@ import (
    "math/big"
    "testing"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/core/holiman"
+    "golang.org/x/crypto/sha3"
)

func TestMemoryGasCost(t *testing.T) {
diff --git a/core/vm/instructions.go b/core/vm/instructions.go
index 35fb3a86..f30c1d9b 100644
--- a/core/vm/instructions.go
+++ b/core/vm/instructions.go
@@ -28,10 +28,10 @@ package vm
import (
    "errors"
+    "sync/atomic"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/core/types"
    "github.com/flare-foundation/coreth/core/holiman"
    "golang.org/x/crypto/sha3"
)
@@ -243,7 +244,7 @@ func opSAR(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte
return nil, nil
}

-func opSha3(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+func opKeccak256(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
    offset, size := scope.Stack.pop(), scope.Stack.peak()
    data := scope.Memory.GetPtr(int64(offset.Uint64()), int64(size.Uint64()))
@@ -537,6 +538,6 @@ func opSload(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]by
}

func opSstore(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+    if interpreter.ReadOnly {
+        return nil, ErrWriteProtection
+    }
    loc := scope.Stack.pop()
    val := scope.Stack.pop()
    interpreter.evm.StateDB.SetState(scope.Contract.Address(),
@@ -545,23 +549,27 @@ func opSstore(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]b
}

func opJump(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+    if atomic.LoadInt32(&interpreter.evm.abort) != 0 {
+        return nil, errStopToken
+    }
}

```

```

pos := scope.Stack.pop()
if !scope.Contract.validJumpdest(&pos) {
    return nil, ErrInvalidJump
}
- *pc = pos.Uint64()
+ *pc = pos.Uint64() - 1 // pc will be increased by the interpreter loop
return nil, nil
}

func opJumpi(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+ if atomic.LoadInt32(&interpreter.evm.abort) != 0 {
+     return nil, errStopToken
+ }
pos, cond := scope.Stack.pop(), scope.Stack.pop()
if !cond.IsZero() {
    if !scope.Contract.validJumpdest(&pos) {
        return nil, ErrInvalidJump
    }
-     *pc = pos.Uint64()
- } else {
-     *pc++
+ *pc = pos.Uint64() - 1 // pc will be increased by the interpreter loop
}
return nil, nil
}
@@ -586,6 +594,9 @@ func opGas(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte
)

func opCreate(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+ if interpreter.readOnly {
+     return nil, ErrWriteProtection
+ }
var (
    value      = scope.Stack.pop()
    offset, size = scope.Stack.pop(), scope.Stack.pop()
@@ -621,12 +632,17 @@ func opCreate2(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]b
scope.Contract.Gas += returnGas

    if suberr == ErrExecutionReverted {
+         interpreter.returnData = res // set REVERT data to return data buffer
        return res, nil
    }
+    interpreter.returnData = nil // clear dirty return data buffer
    return nil, nil
}

func opCreate2(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+ if interpreter.readOnly {
+     return nil, ErrWriteProtection
+ }
var (
    endowment = scope.Stack.pop()
    offset, size = scope.Stack.pop(), scope.Stack.pop()
@@ -657,8 +673,18 @@ func opCreate2(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]
scope.Contract.Gas += returnGas

    if suberr == ErrExecutionReverted {
+         interpreter.returnData = res // set REVERT data to return data buffer
        return res, nil
    }
+    interpreter.returnData = nil // clear dirty return data buffer
    return nil, nil
}

@@ -674,6 +692,9 @@ func opCall(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byt
// Get the arguments from the memory.
args := scope.Memory.GetPtr(int64(inOffset.Uint64()), int64(inSize.Uint64()))

+ if interpreter.readOnly && !value.IsZero() {
+     return nil, ErrWriteProtection
+ }
var bigVal = big0
//TODO0: use uint256.Int instead of converting with toBig()
// By using big0 here, we save an alloc for the most common case (non-ether-transferring contract calls),
@@ -697,6 +718,7 @@ func opCall(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byt
}
scope.Contract.Gas += returnGas

+ interpreter.returnData = ret
return ret, nil
}

@@ -713,6 +735,9 @@ func opCallExpert(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext)
// Get the arguments from the memory.
args := scope.Memory.GetPtr(int64(inOffset.Uint64()), int64(inSize.Uint64()))

+ if interpreter.readOnly && !value.IsZero() {
+     return nil, ErrWriteProtection
+ }
var bigVal = big0
//TODO0: use uint256.Int instead of converting with toBig()
// By using big0 here, we save an alloc for the most common case (non-ether-transferring contract calls),
@@ -730,7 +755,7 @@ func opCallExpert(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext)
bigVal2 = value2.ToBig()
}

- ret, returnGas, err := interpreter.evm.CallExpert(scope.Contract, toAddr, args, gas, bigVal, &coinID, bigVal2)
+ ret, returnGas, err := interpreter.evm.CallExpert(scope.Contract, toAddr, args, gas, bigVal, coinID, bigVal2)

if err != nil {
    temp.Clear()
@@ -744,6 +769,7 @@ func opCallExpert(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext)
}
scope.Contract.Gas += returnGas

+ interpreter.returnData = ret
return ret, nil
}

func opCallCode(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
@@ -778,6 +804,7 @@ func opCallCode(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) (
}
scope.Contract.Gas += returnGas

+ interpreter.returnData = ret
return ret, nil
}

@@ -806,6 +833,7 @@ func opDelegateCall(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext
}
scope.Contract.Gas += returnGas

+ interpreter.returnData = ret
return ret, nil
}

@@ -834,6 +862,7 @@ func opStaticCall(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext)
}
scope.Contract.Gas += returnGas

+ interpreter.returnData = ret
return ret, nil
}

@@ -841,21 +870,29 @@ func opReturn(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]b

```



```

offset, size := scope.Stack.pop(), scope.Stack.pop()
ret := scope.Memory.GetPtr(int64(offset.Uint64()), int64(size.Uint64()))

-   return ret, nil
+   return ret, errStopToken
}

func opRevert(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
    offset, size := scope.Stack.pop(), scope.Stack.pop()
    ret := scope.Memory.GetPtr(int64(offset.Uint64()), int64(size.Uint64()))

-   return ret, nil
+   interpreter.returnData = ret
+   return ret, ErrExecutionReverted
}

+func opUndefined(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+    return nil, &ErrInvalidOpCode{opcode: OpCode(scope.Contract.Code["pc"])}
+}

func opStop(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
-   return nil, nil
+   return nil, errStopToken
}

-func opSuicide(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+func opSelfdestruct(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+    if interpreter.readOnly {
+        return nil, ErrWriteProtection
+    }
    beneficiary := scope.Stack.pop()
    balance := interpreter.evm.StateDB.GetBalance(scope.Contract.Address())
    interpreter.evm.StateDB.AddBalance(beneficiary.Bytes20(), balance)
@@ -864,7 +901,7 @@ func opSuicide(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]
    interpreter.cfg.Tracer.CaptureEnter(SELFDESTRUCT, scope.Contract.Address(), beneficiary.Bytes20(), []byte{}, 0, balance)
    interpreter.cfg.Tracer.CaptureExit([]byte{}, 0, nil)
-   return nil, nil
+   return nil, errStopToken
}

// following functions are used by the instruction jump table
@@ -872,6 +909,9 @@ func opSuicide(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]
// make log instruction function
func makeLog(size int) executionFunc {
    return func(pc *uint64, interpreter *EVMInterpreter, scope *ScopeContext) ([]byte, error) {
+        if interpreter.readOnly {
+            return nil, ErrWriteProtection
+        }
        topics := make([]common.Hash, size)
        stack := scope.Stack
        mStart, mSize := stack.pop(), stack.pop()
diff --git a/core/vm/instructions_test.go b/core/vm/instructions_test.go
index 1570b36f..Scde24ea 100644
--- a/core/vm/instructions_test.go
+++ b/core/vm/instructions_test.go
@@ -33,9 +33,9 @@ import (
    "io/ioutil"
    "testing"

-   "github.com/ava-labs/coreth/params"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/holiman/uint256"
)

@@ -535,12 +535,14 @@ func TestOpMstore(t *testing.T) {
    mem.Resize(64)
    pc := uint64(0)
    v := "abcdef00000000000000abba00000000dea000000c0de0010000000013700"
-   stack.pushN("new(uint256.Int).SetBytes(common.Hex2Bytes(v)), "new(uint256.Int))
+   stack.push(new(uint256.Int).SetBytes(common.Hex2Bytes(v)))
+   stack.push(new(uint256.Int).SetBytes(common.Hex2Bytes(v)))
+   stack.push(new(uint256.Int))
    opMstore(&pc, evmInterpreter, &ScopeContext{mem, stack, nil})
    if got := common.Bytes2Hex(mem.GetCopy(0, 32)); got != v {
        t.Fatalf("Mstore fail, got %v, expected %v", got, v)
    }
-   stack.pushN("new(uint256.Int).SetUint64(0x1), "new(uint256.Int))
+   stack.push(new(uint256.Int).SetUint64(0x1))
+   stack.push(new(uint256.Int).SetUint64(0x1))
+   stack.push(new(uint256.Int))
    opMstore(&pc, evmInterpreter, &ScopeContext{mem, stack, nil})
    if common.Bytes2Hex(mem.GetCopy(0, 32)) != "0000000000000000000000000000000000000000000000000000000000000001" {
        t.Fatalf("Mstore failed to overwrite previous value")
    }
@@ -563,12 +563,13 @@ func BenchmarkOpMstore(bench *testing.B) {
    bench.ResetTimer()
    for i := 0; i < bench.N; i++ {
-       stack.pushN(*value, *memStart)
+       stack.push(value)
+       stack.push(value)
+       stack.push(memStart)
        opMstore(&pc, evmInterpreter, &ScopeContext{mem, stack, nil})
    }
}

-func BenchmarkOpSHA3(bench *testing.B) {
+func BenchmarkOpKeccak256(bench *testing.B) {
    var (
        env      = NewEVM(BlockContext{}, TxContext{}, nil, params.TestChainConfig, Config{})
        stack    = newstack()
@@ -582,9 +585,9 @@ func BenchmarkOpSHA3(bench *testing.B) {
    bench.ResetTimer()
    for i := 0; i < bench.N; i++ {
-       stack.pushN("uint256.NewInt(32), *start)
+       stack.push("uint256.NewInt(32)", *start)
        opSha3(&pc, evmInterpreter, &ScopeContext{mem, stack, nil})
+       stack.push("uint256.NewInt(32)")
+       stack.push(start)
+       opKeccak256(&pc, evmInterpreter, &ScopeContext{mem, stack, nil})
    }
}

diff --git a/core/vm/interface.go b/core/vm/interface.go
index bde4b08e..e5acbfbb 100644
--- a/core/vm/interface.go
+++ b/core/vm/interface.go
@@ -29,8 +29,8 @@ package vm
import (
    "math/big"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ethereum/go-ethereum/common"
+   "github.com/flare-foundation/coreth/core/types"
+)

// StateDB is an EVM database for full state querying.
diff --git a/core/vm/interpreter.go b/core/vm/interpreter.go
index 805aea61..653d64ca 100644
--- a/core/vm/interpreter.go
+++ b/core/vm/interpreter.go
@@ -28,7 +28,7 @@ package vm
import (
    "hash"

```

```

-     "sync/atomic"
-
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/common/math"
@@ -46,11 +45,10 @@ var (
type Config struct {
    Debug                bool           // Enables debugging
    Tracer               EVMLogger     // Opcode logger
-   NoRecursion          bool           // Disables call, callcode, delegate call and create
-   NoBaseFee            bool           // Forces the EIP-1559 baseFee to 0 (needed for 0 price calls)
    EnablePreimageRecording bool           // Enables recording of SHA3/keccak preimages
-
-   JumpTable [256]*operation // EVM instruction table, automatically populated if unset
+   JumpTable *JumpTable // EVM instruction table, automatically populated if unset
-
-   ExtraEips []int // Additional EIPS that are to be enabled
-
@@ -88,41 +86,39 @@ type EVMInterpreter struct {
// NewEVMInterpreter returns a new instance of the Interpreter.
func NewEVMInterpreter(evm *EVM, cfg Config) *EVMInterpreter {
-   // We use the STOP instruction whether to see
-   // the jump table was initialised. If it was not
-   // we'll set the default jump table.
-   if cfg.JumpTable[STOP] == nil {
-       var jt JumpTable
+   // If jump table was not initialised we set the default one.
+   if cfg.JumpTable == nil {
+       switch {
-   case evm.chainRules.IsApricotPhase3:
-       jt = apricotPhase3InstructionSet
-       cfg.JumpTable = &apricotPhase3InstructionSet
-   case evm.chainRules.IsApricotPhase2:
-       jt = apricotPhase2InstructionSet
-       cfg.JumpTable = &apricotPhase2InstructionSet
-   case evm.chainRules.IsApricotPhase1:
-       jt = apricotPhase1InstructionSet
-       cfg.JumpTable = &apricotPhase1InstructionSet
-   case evm.chainRules.IsIstanbul:
-       jt = istanbulInstructionSet
-       cfg.JumpTable = &istanbulInstructionSet
-   case evm.chainRules.IsConstantinople:
-       jt = constantinopleInstructionSet
-       cfg.JumpTable = &constantinopleInstructionSet
-   case evm.chainRules.IsByzantium:
-       jt = byzantiumInstructionSet
-       cfg.JumpTable = &byzantiumInstructionSet
-   case evm.chainRules.IsEIP150:
-       jt = spuriousDragonInstructionSet
-       cfg.JumpTable = &spuriousDragonInstructionSet
-   case evm.chainRules.IsEIP150:
-       jt = tangerineWhistleInstructionSet
-       cfg.JumpTable = &tangerineWhistleInstructionSet
-   case evm.chainRules.IsHomestead:
-       jt = homesteadInstructionSet
-       cfg.JumpTable = &homesteadInstructionSet
-   default:
-       jt = frontierInstructionSet
-       cfg.JumpTable = &frontierInstructionSet
-   }
-   for i, eip := range cfg.ExtraEips {
-       if err := EnableEIP(eip, &jt); err != nil {
-           copy := *cfg.JumpTable
-           if err := EnableEIP(eip, &copy); err != nil {
-               // Disable it, so caller can check if it's activated or not
-               cfg.ExtraEips = append(cfg.ExtraEips[:i], cfg.ExtraEips[i+1:]...)
-               log.Error("EIP activation failed", "eip", eip, "error", err)
-           }
-           cfg.JumpTable = &copy
-       }
-       cfg.JumpTable = jt
-   }
-   }
-   return &EVMInterpreter{
@@ -138,7 +134,11 @@ func NewEVMInterpreter(evm *EVM, cfg Config) *EVMInterpreter {
// considered a revert-and-consume-all-gas operation except for
// ErrExecutionReverted which means revert-and-keep-gas-left.
func (in *EVMInterpreter) Run(contract *Contract, input []byte, readOnly bool) (ret []byte, err error) {
-   if contract.Address() == BuiltinAddr {
+   // Deprecate special handling of [BuiltinAddr] as of ApricotPhase2.
+   // In ApricotPhase2, the contract deployed in the genesis is overridden by a deprecated precompiled
+   // contract which will return an error immediately if its ever called. Therefore, this function should
+   // never be called after ApricotPhase2 with [BuiltinAddr] as the contract address.
+   if !in.evm.chainRules.IsApricotPhase2 && contract.Address() == BuiltinAddr {
+       self := AccountRef(contract.Caller())
+       if _, ok := contract.caller.(*Contract); ok {
+           contract = contract.AsDelegate()
@@ -201,9 +201,9 @@ func (in *EVMInterpreter) Run(contract *Contract, input []byte, readOnly bool) {
defer func() {
    if err != nil {
        if !logged {
-           in.cfg.Tracer.CaptureState(in.evm, pcCopy, op, gasCopy, cost, callContext, in.returnData, in.evm.depth, err)
+           in.cfg.Tracer.CaptureState(pcCopy, op, gasCopy, cost, callContext, in.returnData, in.evm.depth, err)
        } else {
-           in.cfg.Tracer.CaptureFault(in.evm, pcCopy, op, gasCopy, cost, callContext, in.evm.depth, err)
+           in.cfg.Tracer.CaptureFault(pcCopy, op, gasCopy, cost, callContext, in.evm.depth, err)
        }
    }
}()
@@ -212,101 +212,71 @@ func (in *EVMInterpreter) Run(contract *Contract, input []byte, readOnly bool) {
// explicit STOP, RETURN or SELFDESTRUCT is executed, an error occurred during
// the execution of one of the operations or until the done flag is set by the
// parent context.
-   steps := 0
-   for {
-       steps++
-       if steps%1000 == 0 && atomic.LoadInt32(&in.evm.abort) != 0 {
-           break
-       }
-       if in.cfg.Debug {
-           // Capture pre-execution values for tracing.
-           logged, pcCopy, gasCopy = false, pc, contract.Gas
-       }
-
-       // Get the operation from the jump table and validate the stack to ensure there are
-       // enough stack items available to perform the operation.
-       op = contract.GetOp(pc)
-       operation := in.cfg.JumpTable[op]
-       if operation == nil {
-           return nil, &ErrInvalidOpCode{opcode: op}
-       }
-       cost = operation.constantGas // For tracing
-       // Validate stack
-       if sLen := stack.len(); sLen < operation.minStack {
-           return nil, &ErrStackUnderflow{stackLen: sLen, required: operation.minStack}
-       } else if sLen > operation.maxStack {
-           return nil, &ErrStackOverflow{stackLen: sLen, limit: operation.maxStack}
-       }
-       // If the operation is valid, enforce write restrictions
-       if in.readOnly && in.evm.chainRules.IsByzantium {
-           // If the interpreter is operating in readonly mode, make sure no
-           // state-modifying operation is performed. The 3rd stack item
-           // for a call operation is the value. Transferring value from one
-           // account to the others means the state is modified and should also

```

```

-         // return with an error.
-         if operation.writes || ((op == CALL || op == CALLEX) && stack.Back(2).Sign() != 0) {
-             return nil, ErrWriteProtection
-         }
-     }
-     // Static portion of gas
-     cost = operation.constantGas // For tracing
-     if !contract.UseGas(operation.constantGas) {
+     if !contract.UseGas(cost) {
-         return nil, ErrOutOfGas
-     }
-
-     var memorySize uint64
-     // calculate the new memory size and expand the memory to fit
-     // the operation
-     // Memory check needs to be done prior to evaluating the dynamic gas portion,
-     // to detect calculation overflows
-     if operation.memorySize != nil {
-         memSize, overflow := operation.memorySize(stack)
-         if overflow {
-             return nil, ErrGasUintOverflow
-         }
-         // memory is expanded in words of 32 bytes. Gas
-         // is also calculated in words.
-         if memorySize, overflow = math.SafeMul(toWordSize(memSize), 32); overflow {
-             return nil, ErrGasUintOverflow
-         }
-     }
-     // Dynamic portion of gas
-     // consume the gas and return an error if not enough gas is available.
-     // cost is explicitly set so that the capture state defer method can get the proper cost
-     if operation.dynamicGas != nil {
+         // All ops with a dynamic memory usage also has a dynamic gas cost.
+         var memorySize uint64
+         // calculate the new memory size and expand the memory to fit
+         // the operation
+         // Memory check needs to be done prior to evaluating the dynamic gas portion,
+         // to detect calculation overflows
+         if operation.memorySize != nil {
+             memSize, overflow := operation.memorySize(stack)
+             if overflow {
+                 return nil, ErrGasUintOverflow
+             }
+             // memory is expanded in words of 32 bytes. Gas
+             // is also calculated in words.
+             if memorySize, overflow = math.SafeMul(toWordSize(memSize), 32); overflow {
+                 return nil, ErrGasUintOverflow
+             }
+         }
+         // Consume the gas and return an error if not enough gas is available.
+         // cost is explicitly set so that the capture state defer method can get the proper cost
+         var dynamicCost uint64
+         dynamicCost, err = operation.dynamicGas(in.evm, contract, stack, mem, memorySize)
+         cost += dynamicCost // total cost, for debug tracing
+         cost += dynamicCost // for tracing
+         if err != nil || !contract.UseGas(dynamicCost) {
+             return nil, ErrOutOfGas
+         }
+         if memorySize > 0 {
+             mem.Resize(memorySize)
+         }
+     }
+     if memorySize > 0 {
+         mem.Resize(memorySize)
+     }
-     if in.cfg.Debug {
-         in.cfg.Tracer.CaptureState(in.evm, pc, op, gasCopy, cost, callContext, in.returnData, in.evm.depth, err)
+         in.cfg.Tracer.CaptureState(pc, op, gasCopy, cost, callContext, in.returnData, in.evm.depth, err)
-         logged = true
-     }
-
-     // execute the operation
-     res, err = operation.execute(&pc, in, callContext)
-     // if the operation clears the return data (e.g. it has returning data)
-     // set the last return to the result of the operation.
-     if operation.returns {
-         in.returnData = res
-     }
-
-     switch {
-     case err != nil:
-         return nil, err
-     case operation.reverts:
-         return res, ErrExecutionReverted
-     case operation.halts:
-         return res, nil
-     case !operation.jumps:
-         pc++
+     if err != nil {
+         break
+     }
+     pc++
+ }
+ return nil, nil
+ if err == errStopToken {
+     err = nil // clear stop token error
+ }
+ }
+ return res, err
}

```

diff --git a/core/vm/interpreter_test.go b/core/vm/interpreter_test.go

new file mode 100644

index 00000000..770f3f67

--- /dev/null

+++ b/core/vm/interpreter_test.go

@@ -0,0 +1,87 @@

+/+ (c) 2020-2021, Ava Labs, Inc.

+/+ This file is a derived work, based on the go-ethereum library whose original notices appear below.

+/+ It is distributed under a license compatible with the licensing terms of the original code from which it is derived.

+/+ Much love to the original authors for their work.

+/+ Copyright 2021 The go-ethereum Authors

+/+ This file is part of the go-ethereum library.

+/+ The go-ethereum library is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or

+/+ (at your option) any later version.

+/+ The go-ethereum library is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

+/+ You should have received a copy of the GNU Lesser General Public License along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.

+package vm

```

+
+import (
+    "math/big"
+    "testing"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
+)
+
+var loopInterruptTests = []string{
+    // infinite loop using JUMP: push(2) jumpdest dup1 jump
+    "60025b8056",
+    // infinite loop using JUMPI: push(1) push(4) jumpdest dup2 dup2 jumpi
+    "600160045b818157",
+}
+
+func TestLoopInterrupt(t *testing.T) {
+    address := common.BytesToAddress([]byte("contract"))
+    vmctx := BlockContext{
+        Transfer: func(StateDB, common.Address, common.Address, *big.Int) {},
+    }
+
+    for i, tt := range loopInterruptTests {
+        statedb, _ := state.New(common.Hash{}, state.NewDatabase(rawdb.NewMemoryDatabase()), nil)
+        statedb.CreateAccount(address)
+        statedb.SetCode(address, common.Hex2Bytes(tt))
+        statedb.Finalise(true)
+
+        evm := NewEVM(vmctx, TxContext{}, statedb, params.TestChainConfig, Config{})
+
+        errChannel := make(chan error)
+        timeout := make(chan bool)
+
+        go func(evm *EVM) {
+            _, _, err := evm.Call(AccountRef(common.Address{}), address, nil, math.MaxUint64, new(big.Int))
+            errChannel <- err
+        }(evm)
+
+        go func() {
+            <-time.After(time.Second)
+            timeout <- true
+        }()
+
+        evm.Cancel()
+
+        select {
+        case <-timeout:
+            t.Errorf("test %d timed out", i)
+        case err := <-errChannel:
+            if err != nil {
+                t.Errorf("test %d failure: %v", i, err)
+            }
+        }
+    }
+}
+
diff --git a/core/vm/jump_table.go b/core/vm/jump_table.go
index d833e523..2d250301 100644
--- a/core/vm/jump_table.go
+++ b/core/vm/jump_table.go
@@ -27,7 +27,9 @@
package vm

import (
-    "github.com/ava-labs/coreth/params"
+    "fmt"
+
+    "github.com/flare-foundation/coreth/params"
)

type (
@@ -50,12 +52,6 @@ type operation struct {

    // memorySize returns the memory size required for the operation
    memorySize memorySizeFunc

-    halts    bool // indicates whether the operation should halt further execution
-    jumps    bool // indicates whether the program counter should not increment
-    writes   bool // determines whether this is a state modifying operation
-    reverts  bool // determines whether the operation reverts state (implicitly halts)
-    returns  bool // determines whether the operations sets the return data content
+}

var (
@@ -74,12 +70,30 @@ var (
// JumpTable contains the EVM opcodes supported at a given fork.
type JumpTable [256]*operation

+func validate(jt JumpTable) JumpTable {
+    for i, op := range jt {
+        if op == nil {
+            panic(fmt.Sprintf("op 0x%x is not set", i))
+        }
+
+        // The interpreter has an assumption that if the memorySize function is
+        // set, then the dynamicGas function is also set. This is a somewhat
+        // arbitrary assumption, and can be removed if we need to -- but it
+        // allows us to avoid a condition check. As long as we have that assumption
+        // in there, this little sanity check prevents us from merging in a
+        // change which violates it.
+        if op.memorySize != nil && op.dynamicGas == nil {
+            panic(fmt.Sprintf("op %v has dynamic memory but not dynamic gas", OpCode(i).String()))
+        }
+    }
+    return jt
+}
+
// newApricotPhase3InstructionSet returns the frontier, homestead, byzantium,
// constantinople, istanbul, petersburg, apricotPhase1, 2, and 3 instructions.
func newApricotPhase3InstructionSet() JumpTable {
    instructionSet := newApricotPhase2InstructionSet()
    enable3198(&instructionSet) // Base fee opcode https://eips.ethereum.org/EIPS/eip-3198
-    return instructionSet
+    return validate(instructionSet)
}

// newApricotPhase1InstructionSet returns the frontier,
@@ -91,7 +85,7 @@ func newApricotPhase2InstructionSet() JumpTable {
    enable2929(&instructionSet)
    enableAP2(&instructionSet)

-    return instructionSet
+    return validate(instructionSet)
}

// newApricotPhase1InstructionSet returns the frontier,
@@ -102,7 +96,7 @@ func newApricotPhase1InstructionSet() JumpTable {
    enableAPI(&instructionSet)

-    return instructionSet
+}

```

```

+         return validate(instructionSet)
+     }

// newIstanbulInstructionSet returns the frontier,
@@ -114,7 +128,7 @@ func newIstanbulInstructionSet() JumpTable {
    enable1884(&instructionSet) // Reprice reader opcodes - https://eips.ethereum.org/EIPS/eip-1884
    enable2200(&instructionSet) // Net metered SSTORE - https://eips.ethereum.org/EIPS/eip-2200

-     return instructionSet
+     return validate(instructionSet)
+ }

// newConstantinopleInstructionSet returns the frontier, homestead,
@@ -152,10 +166,8 @@ func newConstantinopleInstructionSet() JumpTable {
    minStack:    minStack(4, 1),
    maxStack:    maxStack(4, 1),
    memorySize:  memoryCreate2,
-     writes:     true,
-     returns:    true,
- }
- return instructionSet
+     return validate(instructionSet)
+ }

// newByzantiumInstructionSet returns the frontier, homestead and
@@ -169,7 +181,6 @@ func newByzantiumInstructionSet() JumpTable {
    minStack:    minStack(6, 1),
    maxStack:    maxStack(6, 1),
    memorySize:  memoryStaticCall,
-     returns:    true,
- }
    instructionSet[RETURNDATASIZE] = &operation{
        execute:    opReturnDataSize,
@@ -191,17 +202,15 @@ func newByzantiumInstructionSet() JumpTable {
    minStack:    minStack(2, 0),
    maxStack:    maxStack(2, 0),
    memorySize:  memoryRevert,
-     reverts:    true,
-     returns:    true,
- }
- return instructionSet
+     return validate(instructionSet)
+ }

// EIP 158 a.k.a Spurious Dragon
func newSpuriousDragonInstructionSet() JumpTable {
    instructionSet := newTangerineWhistleInstructionSet()
    instructionSet[EXP].dynamicGas = gasExpEIP158
-     return instructionSet
+     return validate(instructionSet)
+ }

@@ -216,7 +225,7 @@ func newTangerineWhistleInstructionSet() JumpTable {
    instructionSet[CALLEX].constantGas = params.CallGasEIP150
    instructionSet[CALLCODE].constantGas = params.CallGasEIP150
    instructionSet[DELEGATECALL].constantGas = params.CallGasEIP150
-     return instructionSet
+     return validate(instructionSet)
+ }

// newHomesteadInstructionSet returns the frontier and homestead
@@ -230,21 +239,19 @@ func newHomesteadInstructionSet() JumpTable {
    minStack:    minStack(6, 1),
    maxStack:    maxStack(6, 1),
    memorySize:  memoryDelegateCall,
-     returns:    true,
- }
- return instructionSet
+     return validate(instructionSet)
+ }

// newFrontierInstructionSet returns the frontier instructions
// that can be executed during the frontier phase.
func newFrontierInstructionSet() JumpTable {
-     return JumpTable{
+     tbl := JumpTable{
        STOP: {
            execute:    opStop,
            constantGas: 0,
            minStack:   minStack(0, 0),
            maxStack:   maxStack(0, 0),
-             halts:     true,
-         },
-         ADD: {
            execute:    opAdd,
@@ -378,13 +385,13 @@ func newFrontierInstructionSet() JumpTable {
    minStack:    minStack(2, 1),
    maxStack:    maxStack(2, 1),
    },
    SHA3: {
-         execute:    opSha3,
-         constantGas: params.Sha3Gas,
-         dynamicGas:  gasSha3,
+         execute:    opKeccak256,
+         constantGas: params.Keccak256Gas,
+         dynamicGas:  gasKeccak256,
        },
        KECCAK256: {
            execute:    opKeccak256,
            constantGas: params.Keccak256Gas,
            dynamicGas:  gasKeccak256,
            minStack:   minStack(2, 1),
            maxStack:   maxStack(2, 1),
            memorySize:  memorySha3,
            memorySize:  memoryKeccak256,
        },
        ADDRESS: {
            execute:    opAddress,
@@ -553,21 +560,18 @@ func newFrontierInstructionSet() JumpTable {
    dynamicGas: gasSStore,
    minStack:   minStack(2, 0),
    maxStack:   maxStack(2, 0),
-     writes:    true,
- },
    JUMP: {
        execute:    opJump,
        constantGas: GasMidStep,
        minStack:   minStack(1, 0),
        maxStack:   maxStack(1, 0),
-         jumps:    true,
- },
    JUMPI: {
        execute:    opJumpi,
        constantGas: GasSlowStep,
        minStack:   minStack(2, 0),
        maxStack:   maxStack(2, 0),
-         jumps:    true,
- },
    PC: {
        execute:    opPc,
@@ -983,7 +987,6 @@ func newFrontierInstructionSet() JumpTable {
    minStack:    minStack(2, 0),
    maxStack:    maxStack(2, 0),
    memorySize:  memoryLog,
-     writes:     true,
- },
    LOG1: {

```

```

        execute:    makeLog(1),
@@ -991,7 +994,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(3, 0),
        maxStack:  maxStack(3, 0),
        memorySize: memoryLog,
        writes:    true,
    },
    LOG2: {
        execute:    makeLog(2),
@@ -999,7 +1001,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(4, 0),
        maxStack:  maxStack(4, 0),
        memorySize: memoryLog,
        writes:    true,
    },
    LOG3: {
        execute:    makeLog(3),
@@ -1007,7 +1008,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(5, 0),
        maxStack:  maxStack(5, 0),
        memorySize: memoryLog,
        writes:    true,
    },
    LOG4: {
        execute:    makeLog(4),
@@ -1015,7 +1015,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(6, 0),
        maxStack:  maxStack(6, 0),
        memorySize: memoryLog,
        writes:    true,
    },
    CREATE: {
        execute:    opCreate,
@@ -1024,8 +1023,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(3, 1),
        maxStack:  maxStack(3, 1),
        memorySize: memoryCreate,
        writes:    true,
        returns:   true,
    },
    CALL: {
        execute:    opCall,
@@ -1034,8 +1031,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(7, 1),
        maxStack:  maxStack(7, 1),
        memorySize: memoryCall,
        returns:   true,
    },
    CALLEX: {
        execute:    opCallExpert,
@@ -1043,8 +1039,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(9, 1),
        maxStack:  maxStack(9, 1),
        memorySize: memoryCallExpert,
        returns:   true,
    },
    CALLCODE: {
        execute:    opCallCode,
@@ -1052,8 +1047,6 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(7, 1),
        maxStack:  maxStack(7, 1),
        memorySize: memoryCall,
        returns:   true,
    },
    RETURN: {
        execute:    opReturn,
@@ -1060,15 +1054,21 @@ func newFrontierInstructionSet() JumpTable {
        minStack:  minStack(2, 0),
        maxStack:  maxStack(2, 0),
        memorySize: memoryReturn,
        halts:     true,
    },
    SELFDSTRUCT: {
        execute:    opSuicide,
+        execute:    opSelfdestruct,
        dynamicGas: gasSelfdestruct,
        minStack:  minStack(1, 0),
        maxStack:  maxStack(1, 0),
        halts:     true,
        writes:    true,
    },
}

+
+ // Fill all unassigned slots with opUndefined.
+ for i, entry := range tbl {
+     if entry == nil {
+         tbl[i] = &operation{execute: opUndefined, maxStack: maxStack(0, 0)}
+     }
+ }
+
+ return validate(tbl)
+}

diff --git a/core/vm/logger.go b/core/vm/logger.go
index ff11a6f5..ba962f77 100644
--- a/core/vm/logger.go
+++ b/core/vm/logger.go
@@ -27,87 +27,125 @@ package vm

import (
-     "encoding/hex"
-     "fmt"
-     "io"
-     "math/big"
-     "strings"
-     "time"
-
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/params"
-     "github.com/ethereum/go-ethereum/common"
-     "github.com/ethereum/go-ethereum/common/hexutil"
-     "github.com/ethereum/go-ethereum/common/math"
-     "github.com/holiman/uint256"
)

-// Storage represents a contract's storage.
-type Storage map[common.Hash]common.Hash
-
-// Copy duplicates the current storage.
-func (s Storage) Copy() Storage {
-     cpy := make(Storage)
-     for key, value := range s {
-         cpy[key] = value
-     }
-     return cpy
-}
-
-// LogConfig are the configuration options for structured logger the EVM
-type LogConfig struct {
-     EnableMemory    bool // enable memory capture
-     DisableStack    bool // disable stack capture
-     DisableStorage  bool // disable storage capture
-     EnableReturnData bool // enable return data capture
-     Debug           bool // print output during capture end

```

```

-     Limit      int // maximum length of output, but zero means unlimited
-     // Chain overrides, can be used to execute a trace using future fork rules
-     Overrides *params.ChainConfig `json:"overrides,omitempty"`
- }
-
- //go:generate gencodec -type StructLog -field-override structLogMarshaling -out gen_structlog.go
-
- // StructLog is emitted to the EVM each cycle and lists information about the current internal state
- // prior to the execution of the statement.
- type StructLog struct {
-     Pc          uint64          `json:"pc"`
-     Op          OpCode         `json:"op"`
-     Gas         uint64         `json:"gas"`
-     GasCost     uint64         `json:"gasCost"`
-     Memory      []byte         `json:"memory"`
-     MemorySize  int            `json:"memSize"`
-     Stack       []uint256.Int `json:"stack"`
-     ReturnData  []byte         `json:"returnData"`
-     Storage     map[common.Hash]common.Hash `json:"- "`
-     Depth       int            `json:"depth"`
-     RefundCounter uint64         `json:"refund"`
-     Err         error          `json:"- "`
- }
-
- // overrides for gencodec
- type structLogMarshaling struct {
-     Gas      math.HexOrDecimal64
-     GasCost  math.HexOrDecimal64
-     Memory   hexutil.Bytes
-     ReturnData hexutil.Bytes
-     OpName    string `json:"opName"` // adds call to OpName() in MarshalJSON
-     ErrorString string `json:"error"` // adds call to ErrorString() in MarshalJSON
- }
-
- // OpName formats the operand name in a human-readable format.
- func (s *StructLog) OpName() string {
-     return s.Op.String()
- }
-
- // ErrorString formats the log's error as a string.
- func (s *StructLog) ErrorString() string {
-     if s.Err != nil {
-         return s.Err.Error()
-     }
-     return ""
- }
-
- // EVMLogger is used to collect execution traces from an EVM transaction
- // execution. CaptureState is called for each step of the VM with the
- // current VM state.
- @@ -115,252 +40,9 @@ func (s *StructLog) ErrorString() string {
- // if you need to retain them beyond the current call.
- type EVMLogger interface {
-     CaptureStart(env *EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int)
-     CaptureState(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, rData []byte, depth int, err error)
+     CaptureState(pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, rData []byte, depth int, err error)
+     CaptureEnter(typ OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int)
+     CaptureExit(output []byte, gasUsed uint64, err error)
-     CaptureFault(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, depth int, err error)
+     CaptureFault(pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, depth int, err error)
+     CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error)
- }
-
- // StructLogger is an EVM state logger and implements EVMLogger.
- //
- // StructLogger can capture state based on the given Log configuration and also keeps
- // a track record of modified storage which is used in reporting snapshots of the
- // contract their storage.
- type StructLogger struct {
-     cfg LogConfig
-
-     storage map[common.Address]Storage
-     logs    []StructLog
-     output  []byte
-     err     error
- }
-
- // NewStructLogger returns a new logger
- func NewStructLogger(cfg *LogConfig) *StructLogger {
-     logger := &StructLogger{
-         storage: make(map[common.Address]Storage),
-     }
-     if cfg != nil {
-         logger.cfg = *cfg
-     }
-     return logger
- }
-
- // Reset clears the data held by the logger.
- func (l *StructLogger) Reset() {
-     l.storage = make(map[common.Address]Storage)
-     l.output = make([]byte, 0)
-     l.logs = l.logs[:0]
-     l.err = nil
- }
-
- // CaptureStart implements the EVMLogger interface to initialize the tracing operation.
- func (l *StructLogger) CaptureStart(env *EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
- }
-
- // CaptureState logs a new structured log message and pushes it out to the environment
- //
- // CaptureState also tracks SLOAD/SSTORE ops to track storage change.
- func (l *StructLogger) CaptureState(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, rData []byte, depth int, err error) {
-     memory := scope.Memory
-     stack := scope.Stack
-     contract := scope.Contract
-     // check if already accumulated the specified number of logs
-     if l.cfg.Limit != 0 && l.cfg.Limit <= len(l.logs) {
-         return
-     }
-     // Copy a snapshot of the current memory state to a new buffer
-     var mem []byte
-     if l.cfg.EnableMemory {
-         mem = make([]byte, len(memory.Data()))
-         copy(mem, memory.Data())
-     }
-     // Copy a snapshot of the current stack state to a new buffer
-     var stck []uint256.Int
-     if !l.cfg.DisableStack {
-         stck = make([]uint256.Int, len(stack.Data()))
-         for i, item := range stack.Data() {
-             stck[i] = item
-         }
-     }
-     // Copy a snapshot of the current storage to a new container
-     var storage Storage
-     if !l.cfg.DisableStorage && (op == SLOAD || op == SSTORE) {
-         // initialise new changed values storage container for this contract
-         // if not present.
-         if l.storage[contract.Address()] == nil {
-             l.storage[contract.Address()] = make(Storage)
-         }
-         // capture SLOAD opcodes and record the read entry in the local storage

```

```

-         if op == SLOAD && stack.len() >= 1 {
-             var (
-                 address = common.Hash(stack.data[stack.len()-1].Bytes32())
-                 value   = env.StateDB.GetState(contract.Address(), address)
-             )
-             l.storage[contract.Address()][address] = value
-             storage = l.storage[contract.Address()].Copy()
-         } else if op == SSTORE && stack.len() >= 2 {
-             // capture SSTORE opcodes and record the written entry in the local storage.
-             var (
-                 value   = common.Hash(stack.data[stack.len()-2].Bytes32())
-                 address = common.Hash(stack.data[stack.len()-1].Bytes32())
-             )
-             l.storage[contract.Address()][address] = value
-             storage = l.storage[contract.Address()].Copy()
-         }
-     }
-     var rdata []byte
-     if l.cfg.EnableReturnData {
-         rdata = make([]byte, len(rData))
-         copy(rdata, rData)
-     }
-     // create a new snapshot of the EVM.
-     log := StructLog{pc, op, gas, cost, mem, memory.Len(), stck, rdata, storage, depth, env.StateDB.GetRefund(), err}
-     l.logs = append(l.logs, log)
- }
-
- // CaptureFault implements the EVMLogger interface to trace an execution fault
- // while running an opcode.
- func (l *StructLogger) CaptureFault(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, depth int, err error) {
- }
-
- // CaptureEnd is called after the call finishes to finalize the tracing.
- func (l *StructLogger) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {
-     l.output = output
-     l.err = err
-     if l.cfg.Debug {
-         fmt.Printf("0x%x\n", output)
-         if err != nil {
-             fmt.Printf(" error: %v\n", err)
-         }
-     }
- }
-
- func (l *StructLogger) CaptureEnter(typ OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
- }
-
- func (l *StructLogger) CaptureExit(output []byte, gasUsed uint64, err error) {}
-
- // StructLogs returns the captured log entries.
- func (l *StructLogger) StructLogs() []StructLog { return l.logs }
-
- // Error returns the VM error captured by the trace.
- func (l *StructLogger) Error() error { return l.err }
-
- // Output returns the VM return value captured by the trace.
- func (l *StructLogger) Output() []byte { return l.output }
-
- // WriteTrace writes a formatted trace to the given writer
- func WriteTrace(writer io.Writer, logs []StructLog) {
-     for _, log := range logs {
-         fmt.Fprintf(writer, "%-16spc=%08d gas=%v cost=%v", log.Op, log.Pc, log.Gas, log.GasCost)
-         if log.Err != nil {
-             fmt.Fprintf(writer, " ERROR: %v", log.Err)
-         }
-         fmt.Fprintln(writer)
-
-         if len(log.Stack) > 0 {
-             fmt.Fprintln(writer, "Stack:")
-             for i := len(log.Stack) - 1; i >= 0; i-- {
-                 fmt.Fprintf(writer, "%08d %s\n", len(log.Stack)-i-1, log.Stack[i].Hex())
-             }
-         }
-
-         if len(log.Memory) > 0 {
-             fmt.Fprintln(writer, "Memory:")
-             fmt.Fprint(writer, hex.Dump(log.Memory))
-         }
-
-         if len(log.Storage) > 0 {
-             fmt.Fprintln(writer, "Storage:")
-             for h, item := range log.Storage {
-                 fmt.Fprintf(writer, "%x: %x\n", h, item)
-             }
-         }
-
-         if len(log.ReturnData) > 0 {
-             fmt.Fprintln(writer, "ReturnData:")
-             fmt.Fprint(writer, hex.Dump(log.ReturnData))
-         }
-         fmt.Fprintln(writer)
-     }
- }
-
- // WriteLogs writes vm logs in a readable format to the given writer
- func WriteLogs(writer io.Writer, logs []*types.Log) {
-     for _, log := range logs {
-         fmt.Fprintf(writer, "LOG%d: %x bn=%d txi=%x\n", len(log.Topics), log.Address, log.BlockNumber, log.TxIndex)
-
-         for i, topic := range log.Topics {
-             fmt.Fprintf(writer, "%08d %x\n", i, topic)
-         }
-
-         fmt.Fprint(writer, hex.Dump(log.Data))
-         fmt.Fprintln(writer)
-     }
- }
-
- type mdLogger struct {
-     out io.Writer
-     cfg *LogConfig
- }
-
- // NewMarkdownLogger creates a logger which outputs information in a format adapted
- // for human readability, and is also a valid markdown table
- func NewMarkdownLogger(cfg *LogConfig, writer io.Writer) *mdLogger {
-     l := &mdLogger{writer, cfg}
-     if l.cfg == nil {
-         l.cfg = &LogConfig{}
-     }
-     return l
- }
-
- func (t *mdLogger) CaptureStart(env *EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
-     if !create {
-         fmt.Fprintf(t.out, "From: `%v`\nTo: `%v`\nData: `0x%x`\nGas: `%d`\nValue `%v` wei\n",
-             from.String(), to.String(),
-             input, gas, value)
-     } else {
-         fmt.Fprintf(t.out, "From: `%v`\nCreate at: `%v`\nData: `0x%x`\nGas: `%d`\nValue `%v` wei\n",
-             from.String(), to.String(),
-             input, gas, value)
-     }
-
-     fmt.Fprintf(t.out, `
-| Pc      | Op      | Cost   | Stack  | RStack | Refund |
-|-----|-----|-----|-----|-----|-----|

```



```

-})
-}
-
-// CaptureState also tracks SLOAD/SSTORE ops to track storage change.
-func (t *mdLogger) CaptureState(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, rData []byte, depth int, err error) {
-    stack := scope.Stack
-    fmt.Fprintf(t.out, "| %4d | %10v | %3d |", pc, op, cost)
-
-    if !t.cfg.DisableStack {
-        // format stack
-        var a []string
-        for _, elem := range stack.data {
-            a = append(a, elem.Hex())
-        }
-        b := fmt.Sprintf("%v]", strings.Join(a, ","))
-        fmt.Fprintf(t.out, "%10v |", b)
-    }
-    fmt.Fprintf(t.out, "%10v |", env.StateDB.GetRefund())
-    fmt.Fprintln(t.out, "")
-    if err != nil {
-        fmt.Fprintf(t.out, "Error: %v\n", err)
-    }
-}
-
-func (t *mdLogger) CaptureFault(env *EVM, pc uint64, op OpCode, gas, cost uint64, scope *ScopeContext, depth int, err error) {
-    fmt.Fprintf(t.out, "\nError: at pc=%d, op=%v: %v\n", pc, op, err)
-}
-
-func (t *mdLogger) CaptureEnd(output []byte, gasUsed uint64, tm time.Duration, err error) {
-    fmt.Fprintf(t.out, "\nOutput: `0x%x`\nConsumed gas: `%d`\nError: `%v`\n",
-        output, gasUsed, err)
-}
-
-func (t *mdLogger) CaptureEnter(typ OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
-}
-
-func (t *mdLogger) CaptureExit(output []byte, gasUsed uint64, err error) {}
diff --git a/core/vm/logger_test.go b/core/vm/logger_test.go
index 1d87e46a..4fb66b02 100644
--- a/core/vm/logger_test.go
+++ b/core/vm/logger_test.go
@@ -30,10 +30,12 @@ import (
     "math/big"
     "testing"

-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/holiman/uint256"
+    "github.com/ethereum/go-ethereum/common"
+
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
)

type dummyContractRef struct {
diff --git a/core/vm/memory_table.go b/core/vm/memory_table.go
index 45dd99bb..4af8c93c 100644
--- a/core/vm/memory_table.go
+++ b/core/vm/memory_table.go
@@ -26,7 +26,7 @@
package vm

-func memorySha3(stack *Stack) (uint64, bool) {
+func memoryKeccak256(stack *Stack) (uint64, bool) {
    return calcMemSize64(stack.Back(0), stack.Back(1))
}

diff --git a/core/vm/opcodes.go b/core/vm/opcodes.go
index 55f86ffc..c8739d5b 100644
--- a/core/vm/opcodes.go
+++ b/core/vm/opcodes.go
@@ -42,75 +42,70 @@ func (op OpCode) IsPush() bool {
    return false
}

-// IsStaticJump specifies if an opcode is JUMP.
-func (op OpCode) IsStaticJump() bool {
-    return op == JUMP
-}
-
-// 0x0 range - arithmetic ops.
-const (
-    STOP OpCode = iota
-    ADD
-    MUL
-    SUB
-    DIV
-    SDIV
-    MOD
-    SMOD
-    ADDMOD
-    MULMOD
-    EXP
-    SIGNEXTEND
+    STOP      OpCode = 0x0
+    ADD       OpCode = 0x1
+    MUL       OpCode = 0x2
+    SUB       OpCode = 0x3
+    DIV       OpCode = 0x4
+    SDIV      OpCode = 0x5
+    MOD       OpCode = 0x6
+    SMOD      OpCode = 0x7
+    ADDMOD    OpCode = 0x8
+    MULMOD    OpCode = 0x9
+    EXP       OpCode = 0xa
+    SIGNEXTEND OpCode = 0xb
+
+    LT      OpCode = 0x10
+    GT      OpCode = 0x11
+    SLT     OpCode = 0x12
+    SGT     OpCode = 0x13
+    EQ      OpCode = 0x14
+    ISZERO  OpCode = 0x15
+    AND     OpCode = 0x16

```

```

+         OR      OpCode = 0x17
+         XOR      OpCode = 0x18
+         NOT      OpCode = 0x19
+         BYTE     OpCode = 0x1a
+         SHL      OpCode = 0x1b
+         SHR      OpCode = 0x1c
+         SAR      OpCode = 0x1d
+
+         SHA3 OpCode = 0x20
+         KECCAK256 OpCode = 0x20
+     )
+
+ // 0x30 range - closure state.
+ const (
+     ADDRESS OpCode = 0x30 + iota
+     BALANCE OpCode = 0x31
+     ORIGIN OpCode = 0x32
+     CALLER OpCode = 0x33
+     CALLVALUE OpCode = 0x34
+     CALLDATALOAD OpCode = 0x35
+     CALLDATASIZE OpCode = 0x36
+     CALLDATACOPY OpCode = 0x37
+     CODESIZE OpCode = 0x38
+     CODECOPY OpCode = 0x39
+     GASPRICE OpCode = 0x3a
+     EXTCODESIZE OpCode = 0x3b
+     EXTCODECOPY OpCode = 0x3c
+     RETURNDATASIZE OpCode = 0x3d
+     RETURNDATACOPY OpCode = 0x3e
+     EXTCODEHASH OpCode = 0x3f
+ )
+
+ // 0x40 range - block operations.
+ const (
+     BLOCKHASH OpCode = 0x40 + iota
+     COINBASE OpCode = 0x41
+     TIMESTAMP OpCode = 0x42
+     NUMBER OpCode = 0x43
+     DIFFICULTY OpCode = 0x44
+     GASLIMIT OpCode = 0x45
+     BLOCKHASH OpCode = 0x40
+     COINBASE OpCode = 0x41
+     TIMESTAMP OpCode = 0x42
+     NUMBER OpCode = 0x43
+     DIFFICULTY OpCode = 0x44
+     GASLIMIT OpCode = 0x45
+     CHAINID OpCode = 0x46
+     SELFBALANCE OpCode = 0x47
+     BASEFEE OpCode = 0x48
+
+ @ -132,7 +127,7 @@ const (
+     JUMPDEST OpCode = 0x5b
+ )
+
+ // 0x60 range - pushes.
+ const (
+     PUSH1 OpCode = 0x60 + iota
+     PUSH2
+ @ -166,7 +161,11 @@ const (
+     PUSH30
+     PUSH31
+     PUSH32
+     DUP1
+ )
+
+ // 0x80 range - dups.
+ const (
+     DUP1 = 0x80 + iota
+     DUP2
+     DUP3
+     DUP4
+ @ -182,7 +181,11 @@ const (
+     DUP14
+     DUP15
+     DUP16
+     SWAP1
+ )
+
+ // 0x90 range - swaps.
+ const (
+     SWAP1 = 0x90 + iota
+     SWAP2
+     SWAP3
+     SWAP4
+ @ -209,13 +212,6 @@ const (
+     LOG4
+ )
+
+ // unofficial opcodes used for parsing.
+ const (
+     PUSH OpCode = 0xb0 + iota
+     DUP
+     SWAP
+ )
+
+ const (
+     BALANCE OpCode = 0xcd
+     CALLEX OpCode = 0xcf
+ @ -223,14 +219,16 @@ const (
+
+ // 0xf0 range - closures.
+ const (
+     CREATE OpCode = 0xf0 + iota
+     CALL OpCode = 0xf1
+     CALLCODE OpCode = 0xf2
+     RETURN OpCode = 0xf3
+     DELEGATECALL OpCode = 0xf4
+     CREATE2 OpCode = 0xf5
+
+     STATICCALL OpCode = 0xfa
+     REVERT OpCode = 0xfd
+     INVALID OpCode = 0xfe

```

```

        SELFDestruct OpCode = 0xff
    )

@@ -267,7 +265,7 @@ var opCodeToString = map[OpCode]string{
    MULMOD: "MULMOD",

    // 0x20 range - crypto.
    -   SHA3: "SHA3",
    +   KECCAK256: "KECCAK256",

    // 0x30 range - closure state.
    ADDRESS: "ADDRESS",
@@ -398,11 +396,8 @@ var opCodeToString = map[OpCode]string{
    CREATE2: "CREATE2",
    STATICCALL: "STATICCALL",
    REVERT: "REVERT",
    +   INVALID: "INVALID",
    SELFDestruct: "SELFDestruct",

    -
    -   PUSH: "PUSH",
    -   DUP: "DUP",
    -   SWAP: "SWAP",
    }

func (op OpCode) String() string {
@@ -441,7 +436,7 @@ var stringToOp = map[string]OpCode{
    "SAR": SAR,
    "ADDMOD": ADDMOD,
    "MULMOD": MULMOD,
    -   "SHA3": SHA3,
    +   "KECCAK256": KECCAK256,
    "ADDRESS": ADDRESS,
    "BALANCE": BALANCE,
    "BALANCEMC": BALANCEMC,
@@ -558,6 +553,7 @@ var stringToOp = map[string]OpCode{
    "RETURN": RETURN,
    "CALLCODE": CALLCODE,
    "REVERT": REVERT,
    +   "INVALID": INVALID,
    "SELFDestruct": SELFDestruct,
}

diff --git a/core/vm/operations_acl.go b/core/vm/operations_acl.go
index adfe7729..31c012cb 100644
--- a/core/vm/operations_acl.go
+++ b/core/vm/operations_acl.go
@@ -29,9 +29,9 @@ package vm
import (
    "errors"

    -   "github.com/ava-labs/coreth/params"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/math"
    +   "github.com/flare-foundation/coreth/params"
)

// gasSStoreEIP2929 implements gas cost for SSTORE according to EIP-2929
diff --git a/core/vm/runtime/env.go b/core/vm/runtime/env.go
index 5293c408..00718ebe 100644
--- a/core/vm/runtime/env.go
+++ b/core/vm/runtime/env.go
@@ -27,8 +27,8 @@ package runtime
import (
    -   "github.com/ava-labs/coreth/core"
    -   "github.com/ava-labs/coreth/core/vm"
    +   "github.com/flare-foundation/coreth/core"
    +   "github.com/flare-foundation/coreth/core/vm"
)

func NewEnv(cfg *Config) *vm.EVM {
diff --git a/core/vm/runtime/runtime.go b/core/vm/runtime/runtime.go
index 74d5499f..ac87b5e0 100644
--- a/core/vm/runtime/runtime.go
+++ b/core/vm/runtime/runtime.go
@@ -31,12 +31,12 @@ import (
    "math/big"
    "time"

    -   "github.com/ava-labs/coreth/core/rawdb"
    -   "github.com/ava-labs/coreth/core/state"
    -   "github.com/ava-labs/coreth/core/vm"
    -   "github.com/ava-labs/coreth/params"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
    +   "github.com/flare-foundation/coreth/core/rawdb"
    +   "github.com/flare-foundation/coreth/core/state"
    +   "github.com/flare-foundation/coreth/core/vm"
    +   "github.com/flare-foundation/coreth/params"
)

// Config is a basic type specifying certain configuration flags for running
diff --git a/core/vm/runtime/runtime_example_test.go b/core/vm/runtime/runtime_example_test.go
index 9850e283..7d9c1352 100644
--- a/core/vm/runtime/runtime_example_test.go
+++ b/core/vm/runtime/runtime_example_test.go
@@ -29,8 +29,8 @@ package runtime_test
import (
    "fmt"

    -   "github.com/ava-labs/coreth/core/vm/runtime"
    "github.com/ethereum/go-ethereum/common"
    +   "github.com/flare-foundation/coreth/core/vm/runtime"
)

func ExampleExecute() {
diff --git a/core/vm/runtime/runtime_test.go b/core/vm/runtime/runtime_test.go
index 456ae781..b3271145 100644
--- a/core/vm/runtime/runtime_test.go
+++ b/core/vm/runtime/runtime_test.go
@@ -34,17 +34,21 @@ import (
    "testing"
    "time"

    -   "github.com/ava-labs/coreth/consensus"
    -   "github.com/ava-labs/coreth/core"
    -   "github.com/ava-labs/coreth/core/rawdb"
    -   "github.com/ava-labs/coreth/core/state"
    -   "github.com/ava-labs/coreth/core/types"
    -   "github.com/ava-labs/coreth/core/vm"
    -   "github.com/ava-labs/coreth/eth/tracers"
    -   "github.com/ava-labs/coreth/params"
    -   "github.com/ethereum/go-ethereum/accounts/abi"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/core/asm"
    +   "github.com/flare-foundation/coreth/accounts/abi"
    +   "github.com/flare-foundation/coreth/consensus"
    +   "github.com/flare-foundation/coreth/core"
    +   "github.com/flare-foundation/coreth/core/rawdb"
    +   "github.com/flare-foundation/coreth/core/state"
    +   "github.com/flare-foundation/coreth/core/types"
    +   "github.com/flare-foundation/coreth/core/vm"
    +   "github.com/flare-foundation/coreth/eth/tracers"

```

```

+     "github.com/flare-foundation/coreth/eth/tracers/logger"
+     "github.com/flare-foundation/coreth/params"
+
+     // force-load js tracers to trigger registration
+     _ "github.com/flare-foundation/coreth/eth/tracers/js"
+ }

func TestDefaults(t *testing.T) {
@@ -333,19 +337,19 @@ func TestBlockhash(t *testing.T) {
}

type stepCounter struct {
-     inner *vm.JSONLogger
+     inner *logger.JSONLogger
    steps int
}

func (s *stepCounter) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
}

-func (s *stepCounter) CaptureFault(env *vm.EVM, pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+func (s *stepCounter) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
}

func (s *stepCounter) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {}

-func (s *stepCounter) CaptureState(env *vm.EVM, pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+func (s *stepCounter) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
    s.steps++
    // Enable this for more output
    //s.inner.CaptureState(env, pc, op, gas, cost, memory, stack, rStack, contract, depth, err)
@@ -500,7 +504,7 @@ func BenchmarkSimpleLoop(b *testing.B) {
    byte(vm.JUMP),
}

-     //tracer := vm.NewJSONLogger(nil, os.Stdout)
+     //tracer := logger.NewJSONLogger(nil, os.Stdout)
    //Execute(loopingCode, nil, &Config{
    //     EVMConfig: vm.Config{
    //         Debug: true,
@@ -521,7 +525,7 @@ func BenchmarkSimpleLoop(b *testing.B) {
    // TestEip2929Cases contains various testcases that are used for
    // EIP-2929 about gas repricings
    func TestEip2929Cases(t *testing.T) {
-
+     t.Skip("Test only useful for generating documentation")
        id := 1
        prettyPrint := func(comment string, code []byte) {
@@ -543,7 +547,7 @@ func TestEip2929Cases(t *testing.T) {
            Execute(code, nil, &Config{
                EVMConfig: vm.Config{
                    Debug: true,
                    Tracer: vm.NewMarkdownLogger(nil, os.Stdout),
                    Tracer: logger.NewMarkdownLogger(nil, os.Stdout),
                    ExtraEips: []int{2929},
                },
            },
        })
@@ -693,7 +697,7 @@ func TestColdAccountAccessCost(t *testing.T) {
        want: 7600,
    },
} {
-     tracer := vm.NewStructLogger(nil)
+     tracer := logger.NewStructLogger(nil)
    Execute(tc.code, nil, &Config{
        EVMConfig: vm.Config{
            Debug: true,
diff --git a/core/vm/stack.go b/core/vm/stack.go
index ebfeef1..7ff708c6 100644
--- a/core/vm/stack.go
+++ b/core/vm/stack.go
@@ -64,10 +64,6 @@ func (st *Stack) push(d *uint256.Int) {
    // NOTE push limit (1024) is checked in baseCheck
    st.data = append(st.data, *d)
}

-func (st *Stack) pushN(ds ...uint256.Int) {
-    // FIXME: Is there a way to pass args by pointers.
-    st.data = append(st.data, ds...)
-}

func (st *Stack) pop() (ret uint256.Int) {
    ret = st.data[len(st.data)-1]
diff --git a/core/vm/stack_table.go b/core/vm/stack_table.go
index 487acaef..42197e64 100644
--- a/core/vm/stack_table.go
+++ b/core/vm/stack_table.go
@@ -27,7 +27,7 @@ package vm

import (
-     "github.com/ava-labs/coreth/params"
+     "github.com/flare-foundation/coreth/params"
)

func minSwapStack(n int) int {
diff --git a/eth/api.go b/eth/api.go
index e58e1545..d00efb5a 100644
--- a/eth/api.go
+++ b/eth/api.go
@@ -36,17 +36,17 @@ import (
    "strings"
    "time"

-     "github.com/ava-labs/coreth/core"
-     "github.com/ava-labs/coreth/core/rawdb"
-     "github.com/ava-labs/coreth/core/state"
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/internal/ethapi"
-     "github.com/ava-labs/coreth/rpc"
-     "github.com/ava-labs/coreth/trie"
+     "github.com/flare-foundation/coreth/core"
+     "github.com/flare-foundation/coreth/core/rawdb"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/internal/ethapi"
+     "github.com/flare-foundation/coreth/rpc"
+     "github.com/flare-foundation/coreth/trie"
)

// PublicEthereumAPI provides an API to access Ethereum full node-related
@@ -93,7 +93,7 @@ func (api *PrivateAdminAPI) ExportChain(file string, first *uint64, last *uint64
    last = &head
}

if _, err := os.Stat(file); err == nil {
-     // File already exists. Allowing overwrite could be a DoS vecotor,
+     // File already exists. Allowing overwrite could be a DoS vector,
    // since the 'file' may point to arbitrary paths on the drive
    return false, errors.New("location would overwrite an existing file")
}

diff --git a/eth/api_backend.go b/eth/api_backend.go

```

```

index 3cc182b3..c2189509 100644
--- a/eth/api_backend.go
+++ b/eth/api_backend.go
@@ -32,21 +32,21 @@ import (
    "math/big"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/bloombits"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/eth/gasprice"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/bloombits"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/gasprice"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/rpc"
)

var (
@@ -242,13 +242,6 @@ func (b *EthAPIBackend) GetLogs(ctx context.Context, hash common.Hash) ([][]*typ
    return logs, nil
}

-func (b *EthAPIBackend) GetTd(ctx context.Context, hash common.Hash) *big.Int {
-    if header := b.eth.blockchain.GetHeaderByHash(hash); header != nil {
-        return b.eth.blockchain.GetTd(hash, header.Number.Uint64())
-    }
-    return nil
-}
-
func (b *EthAPIBackend) GetEVM(ctx context.Context, msg core.Message, state *state.StateDB, header *types.Header, vmConfig *vm.Config) (*vm.EVM, func() error, error) {
    vmError := func() error { return nil }
    if vmConfig == nil {
@@ -428,7 +421,7 @@ func (b *EthAPIBackend) GetMaxBlocksPerRequest() int64 {
}

func (b *EthAPIBackend) StateAtBlock(ctx context.Context, block *types.Block, reexec uint64, base *state.StateDB, checkLive bool, preferDisk bool) (*state.StateDB, error) {
-    return b.eth.stateAtBlock(block, reexec, base, checkLive, preferDisk)
+    return b.eth.StateAtBlock(block, reexec, base, checkLive, preferDisk)
}

func (b *EthAPIBackend) StateAtTransaction(ctx context.Context, block *types.Block, txIndex int, reexec uint64) (core.Message, vm.BlockContext, *state.StateDB, error) {
diff --git a/eth/backend.go b/eth/backend.go
index 8b76075f..32cfbd1 100644
--- a/eth/backend.go
+++ b/eth/backend.go
@@ -33,27 +33,30 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/bloombits"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/eth/ethconfig"
-    "github.com/ava-labs/coreth/eth/filters"
-    "github.com/ava-labs/coreth/eth/gasprice"
-    "github.com/ava-labs/coreth/eth/tracers"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/internal/ethapi"
-    "github.com/ava-labs/coreth/miner"
-    "github.com/ava-labs/coreth/node"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/bloombits"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/state/pruner"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/ethconfig"
+    "github.com/flare-foundation/coreth/eth/filters"
+    "github.com/flare-foundation/coreth/eth/gasprice"
+    "github.com/flare-foundation/coreth/eth/tracers"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/internal/ethapi"
+    "github.com/flare-foundation/coreth/internal/shutdowncheck"
+    "github.com/flare-foundation/coreth/miner"
+    "github.com/flare-foundation/coreth/node"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/rpc"
+    "github.com/flare-foundation/flare/utl/timer/mockable"
)

// Config contains the configuration options of the ETH protocol.
@@ -97,16 +100,23 @@ type Ethereum struct {

    lock sync.RWMutex // Protects the variadic fields (e.g. gas price and etherbase)

+    shutdownTracker *shutdowncheck.ShutdownTracker // Tracks if and when the node has shutdown ungracefully
+
+    stackRPCs []rpc.API
+
    settings Settings // Settings for Ethereum API
}

// New creates a new Ethereum object (including the
// initialisation of the common Ethereum object)
-func New(stack *node.Node, config *Config,
+func New(
+    stack *node.Node,
+    config *Config,
+    cb *dummy.ConsensusCallbacks,
+    chainDb ethdb.Database,
+    settings Settings,

```

```

        lastAcceptedHash common.Hash,
        clock *mockable.Clock,
    ) (*Ethereum, error) {
        if chainDb == nil {
            return nil, errors.New("chainDb cannot be nil")
        }
    }
@@ -131,14 +141,19 @@ func New(stack *node.Node, config *Config,
    log.Info("Initialised chain configuration", "config", chainConfig)

- // FIXME RecoverPruning once that package is migrated over
- // if err := pruner.RecoverPruning(stack.ResolvePath(""), chainDb, stack.ResolvePath(config.TrieCleanCacheJournal)); err != nil {
- //     log.Error("Failed to recover state", "error", err)
- // }
+ // Note: RecoverPruning must be called to handle the case that we are midway through offline pruning.
+ // If the data directory is changed in between runs preventing RecoverPruning from performing its job correctly,
+ // it may cause DB corruption.
+ // Since RecoverPruning will only continue a pruning run that already began, we do not need to ensure that
+ // reprocessState has already been called and completed successfully. To ensure this, we must maintain
+ // that Prune is only run after reprocessState has finished successfully.
+ if err := pruner.RecoverPruning(config.OfflinePruningDataDirectory, chainDb); err != nil {
+     log.Error("Failed to recover state", "error", err)
+ }
    eth := &Ethereum{
        config:      config,
        chainDb:     chainDb,
-       eventMux:    stack.EventMux(),
+       eventMux:    new(event.TypeMux),
        accountManager: stack.AccountManager(),
        engine:       dummy.NewDummyEngine(cb),
        closeBloomHandler: make(chan struct{}),
@@ -147,6 +162,7 @@ func New(stack *node.Node, config *Config,
        bloomRequests:  make(chan chan *bloombits.Retrieval),
        bloomIndexer:   core.NewBloomIndexer(chainDb, params.BloomBitsBlocks, params.BloomConfirms),
        settings:       settings,
+       shutdownTracker: shutdowncheck.NewShutdownTracker(chainDb),
    }

    bcVersion := rawdb.ReadDatabaseVersion(chainDb)
@@ -184,16 +200,17 @@ func New(stack *node.Node, config *Config,
    if err != nil {
        return nil, err
    }

+    if err := eth.handleOfflinePruning(cacheConfig, chainConfig, vmConfig, lastAcceptedHash); err != nil {
+        return nil, err
+    }

    eth.bloomIndexer.Start(eth.blockchain)

- // Original code (requires disk):
- // if config.TxPool.Journal != "" {
- //     config.TxPool.Journal = stack.ResolvePath(config.TxPool.Journal)
- // }
+ config.TxPool.Journal = ""
    eth.txPool = core.NewTxPool(config.TxPool, chainConfig, eth.blockchain)

- eth.miner = miner.New(eth, &config.Miner, chainConfig, eth.EventMux(), eth.engine)
+ eth.miner = miner.New(eth, &config.Miner, chainConfig, eth.EventMux(), eth.engine, clock)

    eth.APIBackend = &EthAPIBackend{
        extRPCEnabled: stack.Config().ExtRPCEnabled(),
@@ -213,8 +230,10 @@ func New(stack *node.Node, config *Config,
    // Start the RPC service
    eth.netRPCService = ethapi.NewPublicNetAPI(eth.NetVersion())

- // Register the backend on the node
- stack.RegisterAPIs(eth.APIs())
+ eth.stackRPCs = stack.APIs()

+ // Successful startup; push a marker and check previous unclean shutdowns.
+ eth.shutdownTracker.MarkStartup()

    return eth, nil
}
@@ -227,8 +246,8 @@ func (s *Ethereum) APIs() []rpc.API {
    // Append tracing APIs
    apis = append(apis, tracers.APIs(s.APIBackend)...)

- // Append any APIs exposed explicitly by the consensus engine
- apis = append(apis, s.engine.APIs(s.BlockChain())...)
+ // Add the APIs from the node
+ apis = append(apis, s.stackRPCs...)

    // Append all the local APIs and return
    return append(apis, []rpc.API{
@@ -237,29 +256,35 @@ func (s *Ethereum) APIs() []rpc.API {
        Version: "1.0",
        Service:  NewPublicEthereumAPI(s),
        Public:   true,
        Name:     "public-eth",
    }, {
        Namespace: "eth",
        Version:   "1.0",
        Service:   filters.NewPublicFilterAPI(s.APIBackend, false, 5*time.Minute),
        Public:    true,
        Name:      "public-eth-filter",
    }, {
        Namespace: "admin",
        Version:   "1.0",
        Service:   NewPrivateAdminAPI(s),
        Name:      "private-admin",
    }, {
        Namespace: "debug",
        Version:   "1.0",
        Service:   NewPublicDebugAPI(s),
        Public:    true,
        Name:      "public-debug",
    }, {
        Namespace: "debug",
        Version:   "1.0",
        Service:   NewPrivateDebugAPI(s),
        Name:      "private-debug",
    }, {
        Namespace: "net",
        Version:   "1.0",
        Service:   s.netRPCService,
        Public:    true,
        Name:      "net",
    },
    }...)
}
@@ -314,6 +339,9 @@ func (s *Ethereum) BloomIndexer() *core.ChainIndexer { return s.bloomIndexer }
func (s *Ethereum) Start() {
    // Start the bloom bits servicing goroutines
    s.startBloomHandlers(params.BloomBitsBlocks)

+ // Regularly update shutdown marker
+ s.shutdownTracker.Start()
+ }

    // Stop implements node.Lifecycle, terminating all internal goroutines used by the
@@ -325,6 +353,10 @@ func (s *Ethereum) Stop() error {
    s.txPool.Stop()
    s.blockchain.Stop()
}

```

```

s.engine.Close()
+
+ // Clean shutdown marker as the last thing before closing db
+ s.shutdownTracker.Stop()
+
+ s.chainDb.Close()
+ s.eventMux.Stop()
+ return nil
@@ -333,3 +365,46 @@ func (s *Ethereum) Stop() error {
func (s *Ethereum) LastAcceptedBlock() *types.Block {
return s.blockchain.LastAcceptedBlock()
}
+
+func (s *Ethereum) handleOfflinePruning(cacheConfig *core.CacheConfig, chainConfig *params.ChainConfig, vmConfig vm.Config, lastAcceptedHash common.Hash) error {
+ if !s.config.OfflinePruning {
+ // Delete the offline pruning marker to indicate that the node started with offline pruning disabled.
+ if err := rawdb.DeleteOfflinePruning(s.chainDb); err != nil {
+ return fmt.Errorf("failed to write offline pruning disabled marker: %w", err)
+ }
+ return nil
+ }
+
+ // Perform offline pruning after NewBlockChain has been called to ensure that we have rolled back the chain
+ // to the last accepted block before pruning begins.
+ // If offline pruning marker is on disk, then we force the node to be started with offline pruning disabled
+ // before allowing another run of offline pruning.
+ if _, err := rawdb.ReadOfflinePruning(s.chainDb); err == nil {
+ log.Error("Offline pruning is not meant to be left enabled permanently. Please disable offline pruning and allow your node to start successfully before running offline pruning again.")
+ return errors.New("cannot start chain with offline pruning enabled on consecutive starts")
+ }
+
+ // Clean up middle roots
+ if err := s.blockchain.CleanBlockRootsAboveLastAccepted(); err != nil {
+ return err
+ }
+ targetRoot := s.blockchain.LastAcceptedBlock().Root()
+
+ // Allow the blockchain to be garbage collected immediately, since we will shut down the chain after offline pruning completes.
+ s.blockchain.Stop()
+ s.blockchain = nil
+ log.Info("Starting offline pruning", "dataDir", s.config.OfflinePruningDataDirectory, "bloomFilterSize", s.config.OfflinePruningBloomFilterSize)
+ pruner, err := pruner.NewPruner(s.chainDb, s.config.OfflinePruningDataDirectory, s.config.OfflinePruningBloomFilterSize)
+ if err != nil {
+ return fmt.Errorf("failed to create new pruner with data directory: %s, size: %d, due to: %w", s.config.OfflinePruningDataDirectory, s.config.OfflinePruningBloomFilterSize, err)
+ }
+ if err := pruner.Prune(targetRoot); err != nil {
+ return fmt.Errorf("failed to prune blockchain with target root: %s due to: %w", targetRoot, err)
+ }
+ s.blockchain, err = core.NewBlockChain(s.chainDb, cacheConfig, chainConfig, s.engine, vmConfig, lastAcceptedHash)
+ if err != nil {
+ return fmt.Errorf("failed to re-initialize blockchain after offline pruning: %w", err)
+ }
+
+ return nil
+}
diff --git a/eth/bloombits.go b/eth/bloombits.go
index ecc0aaf1..1a0d5d46 100644
--- a/eth/bloombits.go
+++ b/eth/bloombits.go
@@ -29,8 +29,8 @@ package eth
import (
"time"

- "github.com/ava-labs/coreth/core/rawdb"
+ "github.com/ethereum/go-ethereum/common/bitutil"
+ "github.com/flare-foundation/coreth/core/rawdb"
)

const (
diff --git a/eth/ethconfig/config.go b/eth/ethconfig/config.go
index d2d76292..a0ab38ae 100644
--- a/eth/ethconfig/config.go
+++ b/eth/ethconfig/config.go
@@ -29,15 +29,15 @@ package ethconfig
import (
"time"

- "github.com/ava-labs/coreth/core"
- "github.com/ava-labs/coreth/eth/gasprice"
- "github.com/ava-labs/coreth/miner"
+ "github.com/ethereum/go-ethereum/common"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/flare-foundation/coreth/eth/gasprice"
+ "github.com/flare-foundation/coreth/miner"
)

// DefaultFullGPOConfig contains default gasprice oracle settings for full node.
var DefaultFullGPOConfig = gasprice.Config{
- Blocks: 20,
+ Blocks: 40,
Percentile: 60,
MaxHeaderHistory: 1024,
MaxBlockHistory: 1024,
@@ -51,22 +51,19 @@ var DefaultConfig = NewDefaultConfig()

func NewDefaultConfig() Config {
return Config{
- NetworkId: 1,
- LightPeers: 100,
- UltraLightFraction: 75,
- DatabaseCache: 512,
- TrieCleanCache: 75,
- TrieCleanCacheJournal: "triecache",
- TrieCleanCacheRejournal: 60 * time.Minute,
- TrieDirtyCache: 256,
- TrieTimeout: 60 * time.Minute,
- SnapshotCache: 128,
- Miner: miner.Config{},
- TxPool: core.DefaultTxPoolConfig,
- RPCGasCap: 25000000,
- RPCEVMTimeout: 5 * time.Second,
- GPO: DefaultFullGPOConfig,
- RPCTxFeeCap: 1, // 1 AVAX
+ NetworkId: 1,
+ LightPeers: 100,
+ UltraLightFraction: 75,
+ DatabaseCache: 512,
+ TrieCleanCache: 75,
+ TrieDirtyCache: 256,
+ SnapshotCache: 128,
+ Miner: miner.Config{},
+ TxPool: core.DefaultTxPoolConfig,
+ RPCGasCap: 25000000,
+ RPCEVMTimeout: 5 * time.Second,
+ GPO: DefaultFullGPOConfig,
+ RPCTxFeeCap: 1, // 1 AVAX
}
}

@@ -110,13 +107,10 @@ type Config struct {
DatabaseCache int
// DatabaseFreezer string

- TrieCleanCache int

```

```

- TrieCleanCacheJournal string `toml:",omitempty"` // Disk journal directory for trie cache to survive node restarts
- TrieCleanCacheRejournal time.Duration `toml:",omitempty"` // Time interval to regenerate the journal for clean cache
- TrieDirtyCache int
- TrieTimeout time.Duration
- SnapshotCache int
- Preimages bool
+ TrieCleanCache int
+ TrieDirtyCache int
+ SnapshotCache int
+ Preimages bool

// Mining options
Miner miner.Config
@@ -150,4 +144,10 @@ type Config struct {
// Unprotected transactions are transactions that are signed without EIP-155
// replay protection.
AllowUnprotectedTx bool

+
+ // OfflinePruning enables offline pruning on startup of the node. If a node is started
+ // with this configuration option, it must finish pruning before resuming normal operation.
+ OfflinePruning bool
+ OfflinePruningBloomFilterSize uint64
+ OfflinePruningDataDirectory string
}

diff --git a/eth/filters/api.go b/eth/filters/api.go
index f269156c..9d70fa0b 100644
--- a/eth/filters/api.go
+++ b/eth/filters/api.go
@@ -35,13 +35,12 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ava-labs/coreth/rpc"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/hexutil"
    "github.com/ethereum/go-ethereum/event"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    "github.com/flare-foundation/coreth/rpc"
)

// filter is a helper struct that holds meta information over the filter type
@@ -61,7 +60,6 @@ type PublicFilterAPI struct {
    Backend Backend
    mux      *event.TypeMux
    quit     chan struct{}
-    chainDb  ethdb.Database
    events   *EventSystem
    filtersMu sync.Mutex
    filters  map[rpc.ID]*filter

@@ -72,7 +70,6 @@ type PublicFilterAPI struct {
func NewPublicFilterAPI(backend Backend, lightMode bool, timeout time.Duration) *PublicFilterAPI {
    api := &PublicFilterAPI{
        backend: backend,
-        chainDb: backend.ChainDb(),
        events:  NewEventSystem(backend, lightMode),
        filters: make(map[rpc.ID]*filter),
        timeout: timeout,
    }

diff --git a/eth/filters/api_test.go b/eth/filters/api_test.go
index 6f356395..041bd46b 100644
--- a/eth/filters/api_test.go
+++ b/eth/filters/api_test.go
@@ -21,8 +21,8 @@ import (
    "fmt"
    "testing"

-    "github.com/ava-labs/coreth/rpc"
    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/rpc"
)

func TestUnmarshalJSONNewFilterArgs(t *testing.T) {
diff --git a/eth/filters/filter.go b/eth/filters/filter.go
index 227ccd64..b90bca8b 100644
--- a/eth/filters/filter.go
+++ b/eth/filters/filter.go
@@ -32,15 +32,15 @@ import (
    "fmt"
    "math/big"

-    "github.com/ava-labs/coreth/core/vm"
+    "github.com/flare-foundation/coreth/core/vm"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/bloombits"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/rpc"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/event"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/bloombits"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/rpc"
)

type Backend interface {
diff --git a/eth/filters/filter_system.go b/eth/filters/filter_system.go
index 86f5a0bd..67829974 100644
--- a/eth/filters/filter_system.go
+++ b/eth/filters/filter_system.go
@@ -34,14 +34,14 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ava-labs/coreth/rpc"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/event"
    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/interfaces"
+    "github.com/flare-foundation/coreth/rpc"
)

// Type determines the kind of filter and is used to put the filter in to
diff --git a/eth/gasprice/feehistory.go b/eth/gasprice/feehistory.go
index 081f6c48..59b3b08d 100644
--- a/eth/gasprice/feehistory.go
+++ b/eth/gasprice/feehistory.go
@@ -36,11 +36,11 @@ import (
    "sort"
    "sync/atomic"

-    _ "github.com/ava-labs/coreth/consensus/misc"

```



```

-      "github.com/ava-labs/coreth/core/types"
-      "github.com/ava-labs/coreth/rpc"
-      "github.com/ethereum/go-ethereum/common"
-      "github.com/ethereum/go-ethereum/log"
+      "github.com/flare-foundation/coreth/consensus/misc"
+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/rpc"
)

var (
diff --git a/eth/gasprice/feehistory_test.go b/eth/gasprice/feehistory_test.go
index 146c2814..e60c4d9c 100644
--- a/eth/gasprice/feehistory_test.go
+++ b/eth/gasprice/feehistory_test.go
@@ -32,12 +32,12 @@ import (
    "math/big"
    "testing"

-      "github.com/ava-labs/coreth/core"
-      "github.com/ava-labs/coreth/core/types"
+      "github.com/flare-foundation/coreth/core"
+      "github.com/flare-foundation/coreth/core/types"

-      "github.com/ava-labs/coreth/params"
-      "github.com/ava-labs/coreth/rpc"
-      "github.com/ethereum/go-ethereum/common"
+      "github.com/flare-foundation/coreth/params"
+      "github.com/flare-foundation/coreth/rpc"
)

func TestFeeHistory(t *testing.T) {
diff --git a/eth/gasprice/gasprice.go b/eth/gasprice/gasprice.go
index 210bc701..ce67d395 100644
--- a/eth/gasprice/gasprice.go
+++ b/eth/gasprice/gasprice.go
@@ -32,22 +32,24 @@ import (
    "sort"
    "sync"

-      "github.com/ava-labs/avalanchego/utils/timer/mockable"
-      "github.com/ava-labs/coreth/consensus/dummy"
-      "github.com/ava-labs/coreth/core"
-      "github.com/ava-labs/coreth/core/types"
-      "github.com/ava-labs/coreth/params"
-      "github.com/ava-labs/coreth/rpc"
-      "github.com/ethereum/go-ethereum/common"
+      "github.com/ethereum/go-ethereum/common/math"
+      "github.com/ethereum/go-ethereum/event"
+      "github.com/ethereum/go-ethereum/log"
+      "github.com/flare-foundation/coreth/consensus/dummy"
+      "github.com/flare-foundation/coreth/core"
+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/params"
+      "github.com/flare-foundation/coreth/rpc"
+      "github.com/flare-foundation/flare/utils/timer/mockable"
+      lru "github.com/hashicorp/golang-lru"
)

var (
    DefaultMaxPrice = big.NewInt(150 * params.GWei)
    DefaultMinPrice = big.NewInt(0 * params.GWei)
-    DefaultMinGasUsed = big.NewInt(2_000_000) // block gas limit is 8,000,000
+    DefaultMinBaseFee = big.NewInt(params.ApricotPhase3InitialBaseFee)
+    DefaultMinGasUsed = big.NewInt(12_000_000) // block gas target is 15,000,000
)

type Config struct {
@@ -74,9 +76,10 @@ type OracleBackend interface {
// Oracle recommends gas prices based on the content of recent
// blocks. Suitable for both light and full clients.
type Oracle struct {
-    backend      OracleBackend
-    lastHead     common.Hash
-    lastPrice    *big.Int
+    backend      OracleBackend
+    lastHead     common.Hash
+    lastPrice    *big.Int
+    lastBaseFee  *big.Int
    // [minPrice] ensures we don't get into a positive feedback loop where tips
    // sink to 0 during a period of slow block production, such that nobody's
    // transactions will be included until the full block fee duration has
@@ -109,8 +112,8 @@ func NewOracle(backend OracleBackend, config Config) *Oracle {
    if percent < 0 {
        percent = 0
        log.Warn("Sanitizing invalid gasprice oracle sample percentile", "provided", config.Percentile, "updated", percent)
-    }
-    if percent > 100 {
+    } else if percent > 100 {
        percent = 100
        log.Warn("Sanitizing invalid gasprice oracle sample percentile", "provided", config.Percentile, "updated", percent)
    }
@@ -129,6 +131,6 @@ func NewOracle(backend OracleBackend, config Config) *Oracle {
    minGasUsed = DefaultMinGasUsed
    log.Warn("Sanitizing invalid gasprice oracle min gas used", "provided", config.MinGasUsed, "updated", minGasUsed)

    maxHeaderHistory := config.MaxHeaderHistory
    if maxHeaderHistory < 1 {
        maxHeaderHistory = 1
        log.Warn("Sanitizing invalid gasprice oracle max header history", "provided", config.MaxHeaderHistory, "updated", maxHeaderHistory)
    }
    maxBlockHistory := config.MaxBlockHistory
    if maxBlockHistory < 1 {
        maxBlockHistory = 1
        log.Warn("Sanitizing invalid gasprice oracle max block history", "provided", config.MaxBlockHistory, "updated", maxBlockHistory)
    }
}

cache, _ := lru.New(2048)
headEvent := make(chan core.ChainHeadEvent, 1)
@@ -146,13 +158,14 @@ func NewOracle(backend OracleBackend, config Config) *Oracle {
    return &Oracle{
        backend:      backend,
        lastPrice:    minPrice,
+        lastBaseFee:  DefaultMinBaseFee,
        minPrice:     minPrice,
        maxPrice:     maxPrice,
        minGasUsed:   minGasUsed,
        checkBlocks:  blocks,
        percentile:   percent,
-        maxHeaderHistory: config.MaxHeaderHistory,
-        maxBlockHistory:  config.MaxBlockHistory,
+        maxHeaderHistory: maxHeaderHistory,
+        maxBlockHistory:  maxBlockHistory,
        historyCache: cache,
    }
}

@@ -161,6 +174,34 @@ func NewOracle(backend OracleBackend, config Config) *Oracle {
// produced at the current time. If ApricotPhase3 has not been activated, it may
// return a nil value and a nil error.
func (oracle *Oracle) EstimateBaseFee(ctx context.Context) (*big.Int, error) {
    _, baseFee, err := oracle.suggestDynamicFees(ctx)
    if err != nil {
        return nil, err
    }
}

```

```

+ // We calculate the [nextBaseFee] if a block were to be produced immediately.
+ // If [nextBaseFee] is lower than the estimate from sampling, then we return it
+ // to prevent returning an incorrectly high fee when the network is quiescent.
+ nextBaseFee, err := oracle.estimateNextBaseFee(ctx)
+ if err != nil {
+     log.Warn("failed to estimate next base fee", "err", err)
+     return baseFee, nil
+ }
+ // If base fees have not been enabled, return a nil value.
+ if nextBaseFee == nil {
+     return nil, nil
+ }
+
+ baseFee = math.BigMin(baseFee, nextBaseFee)
+ return baseFee, nil
+}
+
+// estimateNextBaseFee calculates what the base fee should be on the next block if it
+// were produced immediately. If the current time is less than the timestamp of the latest
+// block, this estimate uses the timestamp of the latest block instead.
+// If the latest block has a nil base fee, this function will return nil as the base fee
+// of the next block.
+func (oracle *Oracle) estimateNextBaseFee(ctx context.Context) (*big.Int, error) {
+    // Fetch the most recent block by number
+    block, err := oracle.backend.BlockByNumber(ctx, rpc.LatestBlockNumber)
+    if err != nil {
+@@ -171,37 +212,35 @@ func (oracle *Oracle) EstimateBaseFee(ctx context.Context) (*big.Int, error) {
+        return nil, nil
+    }
+
+    // If the current time is prior to the parent timestamp, then we use the parent
+    // timestamp instead.
+    header := block.Header()
+    timestamp := oracle.clock.Unix()
+    if timestamp < header.Time {
+        timestamp = header.Time
+    }
+    // If the block does have a baseFee, calculate the next base fee
+    // based on the current time and add it to the tip to estimate the
+    // total gas price estimate.
+    _, nextBaseFee, err := dummy.CalcBaseFee(oracle.backend.ChainConfig(), header, timestamp)
+    _, nextBaseFee, err := dummy.EstimateNextBaseFee(oracle.backend.ChainConfig(), block.Header(), oracle.clock.Unix())
+    return nextBaseFee, err
+}
+
+// SuggestPrice returns an estimated price for legacy transactions.
+func (oracle *Oracle) SuggestPrice(ctx context.Context) (*big.Int, error) {
+    // Estimate the effective tip based on recent blocks.
+    tip, err := oracle.suggestTipCap(ctx)
+    tip, baseFee, err := oracle.suggestDynamicFees(ctx)
+    if err != nil {
+        return nil, err
+    }
+    nextBaseFee, err := oracle.EstimateBaseFee(ctx)
+
+    // We calculate the [nextBaseFee] if a block were to be produced immediately.
+    // If [nextBaseFee] is lower than the estimate from sampling, then we return it
+    // to prevent returning an incorrectly high fee when the network is quiescent.
+    nextBaseFee, err := oracle.estimateNextBaseFee(ctx)
+    if err != nil {
+        return nil, err
+    }
+    log.Warn("failed to estimate next base fee", "err", err)
+
+    // If [nextBaseFee] is nil, return [tip] without modification.
+    if nextBaseFee == nil {
+        return tip, nil
+    }
+    // Separately from checking the error value, check that [nextBaseFee] is non-nil
+    // before attempting to take the minimum.
+    if nextBaseFee != nil {
+        baseFee = math.BigMin(baseFee, nextBaseFee)
+    }
+
+    return new(big.Int).Add(tip, nextBaseFee), nil
+    return new(big.Int).Add(tip, baseFee), nil
+}
+
+// SuggestTipCap returns a tip cap so that newly created transaction can have a
+@@ -211,54 +259,43 @@ func (oracle *Oracle) SuggestPrice(ctx context.Context) (*big.Int, error) {
+    // necessary to add the basefee to the returned number to fall back to the legacy
+    // behavior.
+    func (oracle *Oracle) SuggestTipCap(ctx context.Context) (*big.Int, error) {
+        return oracle.suggestTipCap(ctx)
+    tip, _, err := oracle.suggestDynamicFees(ctx)
+    return tip, err
+    }
+
+// suggestTipCap checks the clock to estimate what network rules will be applied to
+// new transactions and then suggests a gas tip cap based on the response.
+func (oracle *Oracle) suggestTipCap(ctx context.Context) (*big.Int, error) {
+    bigTimestamp := big.NewInt(oracle.clock.Time().Unix())
+
+    switch {
+    case oracle.backend.ChainConfig().IsApricotPhase4(bigTimestamp):
+        return oracle.suggestDynamicTipCap(ctx)
+    case oracle.backend.ChainConfig().IsApricotPhase3(bigTimestamp):
+        return new(big.Int).Set(common.Big0), nil
+    case oracle.backend.ChainConfig().IsApricotPhase1(bigTimestamp):
+        return big.NewInt(params.ApricotPhase1MinGasPrice), nil
+    default:
+        return big.NewInt(params.LaunchMinGasPrice), nil
+    }
+
+// suggestDynamicFees estimates the gas tip and base fee based on a simple sampling method
+func (oracle *Oracle) suggestDynamicFees(ctx context.Context) (*big.Int, *big.Int, error) {
+    head, err := oracle.backend.HeaderByNumber(ctx, rpc.LatestBlockNumber)
+    if err != nil {
+        return nil, nil, err
+    }
+
+// suggestDynamicTipCap estimates the gas tip based on a simple sampling method
+func (oracle *Oracle) suggestDynamicTipCap(ctx context.Context) (*big.Int, error) {
+    head, _ := oracle.backend.HeaderByNumber(ctx, rpc.LatestBlockNumber)
+    headHash := head.Hash()
+
+    // If the latest gasprice is still available, return it.
+    oracle.cacheLock.RLock()
+    lastHead, lastPrice := oracle.lastHead, oracle.lastPrice
+    lastHead, lastPrice, lastBaseFee := oracle.lastHead, oracle.lastPrice, oracle.lastBaseFee
+    oracle.cacheLock.RUnlock()
+    if headHash == lastHead {
+        return new(big.Int).Set(lastPrice), nil
+        return new(big.Int).Set(lastPrice), new(big.Int).Set(lastBaseFee), nil
+    }
+
+    oracle.fetchLock.Lock()
+    defer oracle.fetchLock.Unlock()
+
+    // Try checking the cache again, maybe the last fetch fetched what we need
+    oracle.cacheLock.RLock()
+    lastHead, lastPrice = oracle.lastHead, oracle.lastPrice
+    lastHead, lastPrice, lastBaseFee = oracle.lastHead, oracle.lastPrice, oracle.lastBaseFee
+    oracle.cacheLock.RUnlock()
+    if headHash == lastHead {
+        return new(big.Int).Set(lastPrice), nil
+        return new(big.Int).Set(lastPrice), new(big.Int).Set(lastBaseFee), nil
+    }
+}

```

```

var (
    sent, exp int
    number      = head.Number.Uint64()
    result      = make(chan results, oracle.checkBlocks)
    quit        = make(chan struct{})
    results     []*big.Int
+   sent, exp   int
+   number      = head.Number.Uint64()
+   result      = make(chan results, oracle.checkBlocks)
+   quit        = make(chan struct{})
+   tipResults  []*big.Int
+   baseFeeResults []*big.Int
)
for sent < oracle.checkBlocks && number > 0 {
    go oracle.getBlockTips(ctx, number, result, quit)
@@ -270,19 +298,31 @@ func (oracle *Oracle) suggestDynamicTipCap(ctx context.Context) (*big.Int, error
    res := <-result
    if res.err != nil {
        close(quit)
-       return new(big.Int).Set(lastPrice), res.err
+       return new(big.Int).Set(lastPrice), new(big.Int).Set(lastBaseFee), res.err
    }
    exp--
    if res.value != nil {
        results = append(results, res.value)
    }
    if res.tip != nil {
        tipResults = append(tipResults, res.tip)
    } else {
        tipResults = append(tipResults, new(big.Int).Set(common.Big0))
    }
    if res.baseFee != nil {
        baseFeeResults = append(baseFeeResults, res.baseFee)
    } else {
        results = append(results, new(big.Int).Set(common.Big0))
        baseFeeResults = append(baseFeeResults, new(big.Int).Set(common.Big0))
    }
}
price := lastPrice
if len(results) > 0 {
    sort.Sort(bigIntArray(results))
    price = results[(len(results)-1)*oracle.percentile/100]
+   baseFee := lastBaseFee
+   if len(tipResults) > 0 {
+       sort.Sort(bigIntArray(tipResults))
+       price = tipResults[(len(tipResults)-1)*oracle.percentile/100]
+   }
+   if len(baseFeeResults) > 0 {
+       sort.Sort(bigIntArray(baseFeeResults))
+       baseFee = baseFeeResults[(len(baseFeeResults)-1)*oracle.percentile/100]
+   }
+   if price.Cmp(oracle.maxPrice) > 0 {
+       price = new(big.Int).Set(oracle.maxPrice)
@@ -293,14 +333,16 @@ func (oracle *Oracle) suggestDynamicTipCap(ctx context.Context) (*big.Int, error
    oracle.cacheLock.Lock()
    oracle.lastHead = headHash
    oracle.lastPrice = price
+   oracle.lastBaseFee = baseFee
    oracle.cacheLock.Unlock()

    return new(big.Int).Set(price), nil
+   return new(big.Int).Set(price), new(big.Int).Set(baseFee), nil
}

type results struct {
-   value *big.Int
-   err   error
+   tip   *big.Int
+   baseFee *big.Int
+   err   error
}

// getBlockTips calculates the minimum required tip to be included in a given
@@ -309,7 +351,7 @@ func (oracle *Oracle) getBlockTips(ctx context.Context, blockNum uint64, result
    header, err := oracle.backend.HeaderByNumber(ctx, rpc.BlockNumber(blockNum))
    if header == nil {
        select {
-       case result <- results(nil, err):
+       case result <- results(nil, nil, err):
+       case result <- results(nil, nil, err):
        case <-quit:
        }
        return
@@ -319,7 +361,7 @@ func (oracle *Oracle) getBlockTips(ctx context.Context, blockNum uint64, result
    // expedite block production.
    if header.GasUsed < oracle.minGasUsed.Uint64() {
        select {
-       case result <- results(nil, nil):
+       case result <- results(nil, header.BaseFee, nil):
+       case result <- results(nil, header.BaseFee, nil):
        case <-quit:
        }
        return
@@ -335,7 +377,7 @@ func (oracle *Oracle) getBlockTips(ctx context.Context, blockNum uint64, result
    // delay in transaction inclusion.
    minTip, err := oracle.backend.MinRequiredTip(ctx, header)
    select {
-       case result <- results(minTip, err):
+       case result <- results(minTip, header.BaseFee, err):
+       case result <- results(minTip, header.BaseFee, err):
        case <-quit:
        }
    }
}

diff --git a/eth/gasprice/gasprice_test.go b/eth/gasprice/gasprice_test.go
index 4a52849b..fce0cf9a 100644
--- a/eth/gasprice/gasprice_test.go
+++ b/eth/gasprice/gasprice_test.go
@@ -31,17 +31,17 @@ import (
    "math/big"
    "testing"

-   "github.com/ava-labs/coreth/consensus/dummy"
-   "github.com/ava-labs/coreth/core"
-   "github.com/ava-labs/coreth/core/rawdb"
-   "github.com/ava-labs/coreth/core/state"
-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/core/vm"
-   "github.com/ava-labs/coreth/params"
-   "github.com/ava-labs/coreth/rpc"
+   "github.com/ethereum/go-ethereum/common"
+   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/ethereum/go-ethereum/event"
+   "github.com/flare-foundation/coreth/consensus/dummy"
+   "github.com/flare-foundation/coreth/core"
+   "github.com/flare-foundation/coreth/core/rawdb"
+   "github.com/flare-foundation/coreth/core/state"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/core/vm"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/flare-foundation/coreth/rpc"
)

const testHead = 32
@@ -195,37 +195,6 @@ func applyGasPriceTest(t *testing.T, test suggestTipCapTest) {
}
}

```

```

-func TestSuggestTipCapNetworkUpgrades(t *testing.T) {
-     tests := map[string]suggestTipCapTest{
-         "launch": {
-             chainConfig: params.TestLaunchConfig,
-             expectedTip: big.NewInt(params.LaunchMinGasPrice),
-         },
-         "apricot phase 1": {
-             chainConfig: params.TestApricotPhase1Config,
-             expectedTip: big.NewInt(params.ApricotPhase1MinGasPrice),
-         },
-         "apricot phase 2": {
-             chainConfig: params.TestApricotPhase2Config,
-             expectedTip: big.NewInt(params.ApricotPhase1MinGasPrice),
-         },
-         "apricot phase 3": {
-             chainConfig: params.TestApricotPhase3Config,
-             expectedTip: big.NewInt(0),
-         },
-         "apricot phase 4": {
-             chainConfig: params.TestApricotPhase4Config,
-             expectedTip: DefaultMinPrice,
-         },
-     },
-     for name, test := range tests {
-         t.Run(name, func(t *testing.T) {
-             applyGasPriceTest(t, test)
-         })
-     }
-}

func TestSuggestTipCapEmptyExtDataGasUsage(t *testing.T) {
    txTip := big.NewInt(55 * params.GWei)
    applyGasPriceTest(t, suggestTipCapTest{
@@ -255,7 +224,7 @@ func TestSuggestTipCapEmptyExtDataGasUsage(t *testing.T) {
        b.AddTx(tx)
    },
    expectedTip: big.NewInt(2_844_353_281),
+    expectedTip: big.NewInt(11_427_927_927),
})

@@ -288,7 +257,7 @@ func TestSuggestTipCapSimple(t *testing.T) {
    },
    expectedTip: big.NewInt(2_844_353_281),
+    expectedTip: big.NewInt(11_427_927_927),
})

@@ -369,7 +338,7 @@ func TestSuggestTipCapSmallTips(t *testing.T) {
    },
    // NOTE: small tips do not bias estimate
    expectedTip: big.NewInt(2_844_353_281),
+    expectedTip: big.NewInt(11_427_927_927),
})

@@ -402,12 +371,12 @@ func TestSuggestTipCapExtDataUsage(t *testing.T) {
    },
    expectedTip: big.NewInt(2_840_938_303),
+    expectedTip: big.NewInt(11_413_453_299),
})

func TestSuggestTipCapMinGas(t *testing.T) {
-    txTip := big.NewInt(55 * params.GWei)
+    txTip := big.NewInt(500 * params.GWei)
    applyGasPriceTest(t, suggestTipCapTest{
        chainConfig: params.TestChainConfig,
        numBlocks:    3,
@@ -438,3 +407,40 @@ func TestSuggestTipCapMinGas(t *testing.T) {
        expectedTip: big.NewInt(0),
    })
}

+// Regression test to ensure that SuggestPrice does not panic prior to activation of ApricotPhase3
+// Note: support for gas estimation without activated hard forks has been deprecated, but we still
+// ensure that the call does not panic.
+func TestSuggestGasPricePreAP3(t *testing.T) {
+    config := Config{
+        Blocks:    20,
+        Percentile: 60,
+    }
+
+    backend := newTestBackend(t, params.TestApricotPhase2Config, 3, nil, func(i int, b *core.BlockGen) {
+        b.SetCoinbase(common.Address{1})
+
+        signer := types.LatestSigner(params.TestApricotPhase2Config)
+        gasPrice := big.NewInt(params.ApricotPhase1MinGasPrice)
+        for j := 0; j < 50; j++ {
+            tx := types.NewTx(&types.LegacyTx{
+                Nonce:    b.TxNonce(addr),
+                To:      &common.Address{},
+                Gas:      params.TxGas,
+                GasPrice: gasPrice,
+                Data:     []byte{},
+            })
+            tx, err := types.SignTx(tx, signer, key)
+            if err != nil {
+                t.Fatalf("failed to create tx: %s", err)
+            }
+            b.AddTx(tx)
+        }
+    })
+    oracle := NewOracle(backend, config)
+    _, err := oracle.SuggestPrice(context.Background())
+    if err != nil {
+        t.Fatal(err)
+    }
+}

diff --git a/eth/state_accessor.go b/eth/state_accessor.go
index 1ef40509..a1efb07b 100644
--- a/eth/state_accessor.go
+++ b/eth/state_accessor.go
@@ -32,15 +32,20 @@ import (
    "math/big"
    "time"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/trie"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core"

```

```

+         "github.com/flare-foundation/coreth/core/state"
+         "github.com/flare-foundation/coreth/core/types"
+         "github.com/flare-foundation/coreth/core/vm"
+         "github.com/flare-foundation/coreth/trie"
+     )
+
+ // StateAtBlock retrieves the state database associated with a certain block.
+ // If no state is locally available for the given block, a number of blocks
+ // are attempted to be reexecuted to generate the desired state. The optional
+ // base layer statedb can be passed then it's regarded as the statedb of the
+ // parent block.
+ // Parameters:
+ // - block: The block for which we want the state (== state at the stateRoot of the parent)
+ // - reexec: The maximum number of blocks to reprocess trying to obtain the desired state
+@@ -51,7 +56,7 @@ import (
+     // storing trash persistently
+     // - preferDisk: this arg can be used by the caller to signal that even though the 'base' is provided,
+     // it would be preferable to start from a fresh state, if we have it on disk.
+func (eth *Ethereum) stateAtBlock(block *types.Block, reexec uint64, base *state.StateDB, checkLive bool, preferDisk bool) (statedb *state.StateDB, err error) {
+func (eth *Ethereum) StateAtBlock(block *types.Block, reexec uint64, base *state.StateDB, checkLive bool, preferDisk bool) (statedb *state.StateDB, err error) {
+    var (
+        current *types.Block
+        database state.Database
+@@ -178,7 +183,7 @@ func (eth *Ethereum) stateAtTransaction(block *types.Block, txIndex int, reexec
+    }
+    // Lookup the statedb of parent block from the live database,
+    // otherwise regenerate it on the flight.
+    - statedb, err := eth.stateAtBlock(parent, reexec, nil, true, false)
+    + statedb, err := eth.StateAtBlock(parent, reexec, nil, true, false)
+    if err != nil {
+        return nil, vm.BlockContext{}, nil, err
+    }
+diff --git a/eth/tracers/api.go b/eth/tracers/api.go
index fcff10f0..8ada419a 100644
--- a/eth/tracers/api.go
+++ b/eth/tracers/api.go
+@@ -36,19 +36,20 @@ import (
+     "sync"
+     "time"
+
+ -     "github.com/ava-labs/coreth/consensus"
+ -     "github.com/ava-labs/coreth/core"
+ -     "github.com/ava-labs/coreth/core/state"
+ -     "github.com/ava-labs/coreth/core/types"
+ -     "github.com/ava-labs/coreth/core/vm"
+ -     "github.com/ava-labs/coreth/ethdb"
+ -     "github.com/ava-labs/coreth/internal/ethapi"
+ -     "github.com/ava-labs/coreth/params"
+ -     "github.com/ava-labs/coreth/rpc"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/ethereum/go-ethereum/common/hexutil"
+     "github.com/ethereum/go-ethereum/log"
+     "github.com/ethereum/go-ethereum/rpc"
+ +     "github.com/flare-foundation/coreth/consensus"
+ +     "github.com/flare-foundation/coreth/core"
+ +     "github.com/flare-foundation/coreth/core/state"
+ +     "github.com/flare-foundation/coreth/core/types"
+ +     "github.com/flare-foundation/coreth/core/vm"
+ +     "github.com/flare-foundation/coreth/eth/tracers/logger"
+ +     "github.com/flare-foundation/coreth/ethdb"
+ +     "github.com/flare-foundation/coreth/internal/ethapi"
+ +     "github.com/flare-foundation/coreth/params"
+ +     "github.com/flare-foundation/coreth/rpc"
+ )
+
+ const (
+@@ -170,7 +171,7 @@ func (api *API) blockByNumberAndHash(ctx context.Context, number rpc.BlockNumber
+
+ // TraceConfig holds extra parameters to trace functions.
+ type TraceConfig struct {
+ -     *vm.LogConfig
+ +     *logger.Config
+     Tracer *string
+     Timeout *string
+     Reexec *uint64
+@@ -179,7 +180,7 @@ type TraceConfig struct {
+ // TraceCallConfig is the config for traceCall API. It holds one more
+ // field to override the state for tracing.
+ type TraceCallConfig struct {
+ -     *vm.LogConfig
+ +     *logger.Config
+     Tracer *string
+     Timeout *string
+     Reexec *uint64
+@@ -188,7 +189,7 @@ type TraceCallConfig struct {
+
+ // StdTraceConfig holds extra parameters to standard-json trace functions.
+ type StdTraceConfig struct {
+ -     vm.LogConfig
+ +     logger.Config
+     Reexec *uint64
+     TxHash common.Hash
+ }
+@@ -570,12 +571,13 @@ func (api *API) traceBlock(ctx context.Context, block *types.Block, config *Trac
+ if threads > len(txs) {
+     threads = len(txs)
+ }
+ -     blockCtx := core.NewEVMBlockContext(block.Header(), api.chainContext(ctx), nil)
+     blockHash := block.Hash()
+     for th := 0; th < threads; th++ {
+         pend.Add(1)
+         go func() {
+             defer pend.Done()
+
+ +             blockCtx := core.NewEVMBlockContext(block.Header(), api.chainContext(ctx), nil)
+             // Fetch and execute the next transaction trace tasks
+             for task := range jobs {
+                 msg, _ := txs[task.index].AsMessage(signer, block.BaseFee())
+@@ -595,6 +597,7 @@ func (api *API) traceBlock(ctx context.Context, block *types.Block, config *Trac
+ }
+ // Feed the transactions into the tracers and return
+ var failed error
+ +     blockCtx := core.NewEVMBlockContext(block.Header(), api.chainContext(ctx), nil)
+     for i, tx := range txs {
+         // Send the trace task over for execution
+         jobs <- &txTraceTask{statedb: statedb.Copy(), index: i}
+@@ -697,10 +700,10 @@ func (api *API) TraceCall(ctx context.Context, args ethapi.TransactionArgs, bloc
+ var traceConfig *TraceConfig
+ if config != nil {
+     traceConfig = &TraceConfig{
+ -         LogConfig: config.LogConfig,
+ -         Tracer:     config.Tracer,
+ -         Timeout:    config.Timeout,
+ -         Reexec:     config.Reexec,
+ +         Config:     config.Config,
+ +         Tracer:     config.Tracer,
+ +         Timeout:    config.Timeout,
+ +         Reexec:     config.Reexec,
+     }
+ }
+ return api.traceTx(ctx, msg, new(Context), vmctx, statedb, traceConfig)
+@@ -718,7 +721,7 @@ func (api *API) traceTx(ctx context.Context, message core.Message, txctx *Contex
+ )
+ switch {

```

```

case config == nil:
-   tracer = vm.NewStructLogger(nil)
+   tracer = logger.NewStructLogger(nil)
case config.Tracer != nil:
    // Define a meaningful timeout of a single transaction trace
    timeout := defaultTraceTimeout
@@ -742,7 +745,7 @@ func (api *API) traceTx(ctx context.Context, message core.Message, txctx *Context
}

default:
-   tracer = vm.NewStructLogger(config.LogConfig)
+   tracer = logger.NewStructLogger(config.Config)
}
// Run the transaction with tracing enabled.
vmenv := vm.NewEVM(vmctx, txContext, statedb, api.backend.ChainConfig(), vm.Config{Debug: true, Tracer: tracer, NoBaseFee: true})
@@ -757,7 +760,7 @@ func (api *API) traceTx(ctx context.Context, message core.Message, txctx *Context

// Depending on the tracer type, format and return the output.
switch tracer := tracer.(type) {
-   case *vm.StructLogger:
+   case *logger.StructLogger:
        // If the result contains a revert reason, return it.
        returnVal := fmt.Sprintf("%x", result.Return())
        if len(result.Revert()) > 0 {
@@ -787,6 +790,7 @@ func APIS(backed Backend) []rpc.API {
        Version:  "1.0",
        Service:   NewAPI(backed),
        Public:    false,
+       Name:      "debug-tracer",
    },
}

}

diff --git a/eth/tracers/api_test.go b/eth/tracers/api_test.go
index a76dc10d..713e0c70 100644
--- a/eth/tracers/api_test.go
+++ b/eth/tracers/api_test.go
@@ -30,7 +30,6 @@ import (
    "bytes"
    "context"
    "crypto/ecdsa"
-   "encoding/json"
    "errors"
    "fmt"
    "math/big"
@@ -38,20 +37,20 @@ import (
    "sort"
    "testing"

-   "github.com/ava-labs/coreth/consensus"
-   "github.com/ava-labs/coreth/consensus/dummy"
-   "github.com/ava-labs/coreth/core"
-   "github.com/ava-labs/coreth/core/rawdb"
-   "github.com/ava-labs/coreth/core/state"
-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/core/vm"
-   "github.com/ava-labs/coreth/ethdb"
-   "github.com/ava-labs/coreth/internal/ethapi"
-   "github.com/ava-labs/coreth/params"
-   "github.com/ava-labs/coreth/rpc"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/common/hexutil"
-   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/flare-foundation/coreth/consensus"
+   "github.com/flare-foundation/coreth/consensus/dummy"
+   "github.com/flare-foundation/coreth/core"
+   "github.com/flare-foundation/coreth/core/rawdb"
+   "github.com/flare-foundation/coreth/core/state"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/core/vm"
+   "github.com/flare-foundation/coreth/ethdb"
+   "github.com/flare-foundation/coreth/internal/ethapi"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/flare-foundation/coreth/rpc"
)

var (
@@ -324,147 +323,6 @@ func TestTraceCall(t *testing.T) {
}

-func TestOverriddenTraceCall(t *testing.T) {
-   t.Parallel()
-
-   // Initialize test accounts
-   accounts := newAccounts(3)
-   genesis := &core.Genesis{Alloc: core.GenesisAlloc{
-       accounts[0].addr: {Balance: big.NewInt(params.Ether)},
-       accounts[1].addr: {Balance: big.NewInt(params.Ether)},
-       accounts[2].addr: {Balance: big.NewInt(params.Ether)},
-   }}
-   genBlocks := 10
-   signer := types.HomesteadSigner{}
-   api := NewAPI(newTestBackend(t, genBlocks, genesis, func(i int, b *core.BlockGen) {
-       // Transfer from account[0] to account[1]
-       //   value: 1000 wei
-       //   fee: 0 wei
-       tx, _ := types.SignTx(types.NewTransaction(uint64(i), accounts[1].addr, big.NewInt(1000), params.TxGas, new(big.Int).Add(b.BaseFee(), big.NewInt(int64(500*params.GWei))), nil), signer, ac
-       b.AddTx(tx)
-   }))
-   randomAccounts, tracer := newAccounts(3), "callTracerJs"
-
-   var testSuite = []struct {
-       blockNumber rpc.BlockNumber
-       call        ethapi.TransactionArgs
-       config      *TraceCallConfig
-       expectErr   error
-       expect      *callTrace
-   }{
-       // Successful call with state overriding
-       {
-           blockNumber: rpc.PendingBlockNumber,
-           call: ethapi.TransactionArgs{
-               From: &randomAccounts[0].addr,
-               To:   &randomAccounts[1].addr,
-               Value: (*hexutil.Big)(big.NewInt(1000)),
-           },
-           config: &TraceCallConfig{
-               Tracer: &tracer,
-               StateOverrides: &ethapi.StateOverride{
-                   randomAccounts[0].addr: ethapi.OverrideAccount{Balance: newRPCBalance(new(big.Int).Mul(big.NewInt(1), big.NewInt(params.Ether)))},
-               },
-           },
-           expectErr: nil,
-           expect: &callTrace{
-               Type: "CALL",
-               From: randomAccounts[0].addr,
-               To:   randomAccounts[1].addr,
-               Gas:  newRPCUint64(24979000),
-               GasUsed: newRPCUint64(0),
-               Value:  (*hexutil.Big)(big.NewInt(1000)),
-           },
-       },
-       // Invalid call without state overriding
-       {

```

```

blockNumber: rpc.PendingBlockNumber,
call: ethapi.TransactionArgs{
    From: &randomAccounts[0].addr,
    To: &randomAccounts[1].addr,
    Value: (*hexutil.Big)(big.NewInt(1000)),
},
config: &TraceCallConfig{
    Tracer: &tracer,
},
expectErr: core.ErrInsufficientFunds,
expect: nil,
},
// Successful simple contract call
//
// // SPDX-License-Identifier: GPL-3.0
//
// pragma solidity >=0.7.0 <0.8.0;
//
// /**
//  * @title Storage
//  * @dev Store & retrieve value in a variable
//  */
// contract Storage {
//     uint256 public number;
//     constructor() {
//         number = block.number;
//     }
// }
{
    blockNumber: rpc.PendingBlockNumber,
    call: ethapi.TransactionArgs{
        From: &randomAccounts[0].addr,
        To: &randomAccounts[2].addr,
        Data: newRPCBytes(common.Hex2Bytes("8381f58a")), // call number()
    },
    config: &TraceCallConfig{
        Tracer: &tracer,
        StateOverrides: &ethapi.StateOverride{
            randomAccounts[2].addr: ethapi.OverrideAccount{
                Code: newRPCBytes(common.Hex2Bytes("6080604052348015600f57600080fd5b506004361060285760003560e01c80638381f58a14602d575b600080fd5b60336049565b6040518082
                StateDiff: newStates([]common.Hash{}, []common.Hash{common.BigToHash(big.NewInt(123))}),
            },
        },
    },
    expectErr: nil,
    expect: &callTrace{
        Type: "CALL",
        From: randomAccounts[0].addr,
        To: randomAccounts[2].addr,
        Input: hexutil.Bytes(common.Hex2Bytes("8381f58a")),
        Output: hexutil.Bytes(common.BigToHash(big.NewInt(123)).Bytes()),
        Gas: newRPCUint64(24978936),
        GasUsed: newRPCUint64(2283),
        Value: (*hexutil.Big)(big.NewInt(0)),
    },
},
}
for i, testspec := range testSuite {
    result, err := api.TraceCall(context.Background(), testspec.call, rpc.BlockNumberOrHash{BlockNumber: &testspec.blockNumber}, testspec.config)
    if testspec.expectErr != nil {
        if err == nil {
            t.Errorf("test %d: want error %v, have nothing", i, testspec.expectErr)
            continue
        }
        if !errors.Is(err, testspec.expectErr) {
            t.Errorf("test %d: error mismatch, want %v, have %v", i, testspec.expectErr, err)
        }
    } else {
        if err != nil {
            t.Errorf("test %d: want no error, have %v", i, err)
            continue
        }
        ret := new(callTrace)
        if err := json.Unmarshal(result.(json.RawMessage), ret); err != nil {
            t.Fatalf("test %d: failed to unmarshal trace result: %v", i, err)
        }
        if !json.Equal(ret, testspec.expect) {
            // uncomment this for easier debugging
            //have, _ := json.MarshalIndent(ret, "", " ")
            //want, _ := json.MarshalIndent(testspec.expect, "", " ")
            //t.Fatalf("trace mismatch: \nhave %v\nwant %v", string(have), string(want))
            t.Fatalf("trace mismatch: \nhave %v\nwant %v", ret, testspec.expect)
        }
    }
}
}
}

func TestTraceTransaction(t *testing.T) {
    t.Parallel()

@@ -629,29 +487,3 @@ func newAccounts(n int) (accounts Accounts) {
    sort.Sort(accounts)
    return accounts
}

-func newRPCBalance(balance *big.Int) **hexutil.Big {
-    rpcBalance := (*hexutil.Big)(balance)
-    return &rpcBalance
-}

-func newRPCUint64(number uint64) *hexutil.Uint64 {
-    rpcUint64 := hexutil.Uint64(number)
-    return &rpcUint64
-}

-func newRPCBytes(bytes []byte) *hexutil.Bytes {
-    rpcBytes := hexutil.Bytes(bytes)
-    return &rpcBytes
-}

-func newStates(keys []common.Hash, vals []common.Hash) *map[common.Hash]common.Hash {
-    if len(keys) != len(vals) {
-        panic("invalid input")
-    }
-    m := make(map[common.Hash]common.Hash)
-    for i := 0; i < len(keys); i++ {
-        m[keys[i]] = vals[i]
-    }
-    return &m
-}

diff --git a/eth/tracers/internal/tracetest/calltrace_test.go b/eth/tracers/internal/tracetest/calltrace_test.go
new file mode 100644
index 00000000..2e628832
--- /dev/null
+++ b/eth/tracers/internal/tracetest/calltrace_test.go
@@ -0,0 +1,404 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+//

```

```

+// Much love to the original authors for their work.
+// *****
+// Copyright 2021 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package tracetest
+
+import (
+    "encoding/json"
+    "io/ioutil"
+    "math/big"
+    "path/filepath"
+    "reflect"
+    "strings"
+    "testing"
+    "unicode"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/tracers"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/tests"
+
+    // Force-load native and js packages, to trigger registration
+    _ "github.com/flare-foundation/coreth/eth/tracers/js"
+    _ "github.com/flare-foundation/coreth/eth/tracers/native"
+)
+
+// To generate a new callTracer test, copy paste the makeTest method below into
+// a Geth console and call it with a transaction hash you which to export.
+
+/*
+// makeTest generates a callTracer test by running a prestate reassembled and a
+// call trace run, assembling all the gathered information into a test case.
+var makeTest = function(tx, rewind) {
+    // Generate the genesis block from the block, transaction and prestate data
+    var block    = eth.getBlock(eth.getTransaction(tx).blockHash);
+    var genesis = eth.getBlock(block.parentHash);
+
+    delete genesis.gasUsed;
+    delete genesis.logsBloom;
+    delete genesis.parentHash;
+    delete genesis.receiptsRoot;
+    delete genesis.sha3Uncles;
+    delete genesis.size;
+    delete genesis.transactions;
+    delete genesis.transactionsRoot;
+    delete genesis.uncles;
+
+    genesis.gasLimit = genesis.gasLimit.toString();
+    genesis.number   = genesis.number.toString();
+    genesis.timestamp = genesis.timestamp.toString();
+
+    genesis.alloc = debug.traceTransaction(tx, {tracer: "prestateTracer", rewind: rewind});
+    for (var key in genesis.alloc) {
+        genesis.alloc[key].nonce = genesis.alloc[key].nonce.toString();
+    }
+    genesis.config = admin.nodeInfo.protocols.eth.config;
+
+    // Generate the call trace and produce the test input
+    var result = debug.traceTransaction(tx, {tracer: "callTracer", rewind: rewind});
+    delete result.time;
+
+    console.log(JSON.stringify({
+        genesis: genesis,
+        context: {
+            number:    block.number.toString(),
+            difficulty: block.difficulty,
+            timestamp: block.timestamp.toString(),
+            gasLimit:  block.gasLimit.toString(),
+            miner:     block.miner,
+        },
+        input:  eth.getRawTransaction(tx),
+        result: result,
+    }, null, 2));
+}
+*/
+
+type callContext struct {
+    Number    math.HexOrDecimal64 `json:"number"`
+    Difficulty *math.HexOrDecimal256 `json:"difficulty"`
+    Time      math.HexOrDecimal64 `json:"timestamp"`
+    GasLimit  math.HexOrDecimal64 `json:"gasLimit"`
+    Miner     common.Address `json:"miner"`
+}
+
+// callTrace is the result of a callTracer run.
+type callTrace struct {
+    Type      string `json:"type"`
+    From      common.Address `json:"from"`
+    To        common.Address `json:"to"`
+    Input     hexutil.Bytes `json:"input"`
+    Output    hexutil.Bytes `json:"output"`
+    Gas       *hexutil.Uint64 `json:"gas,omitempty"`
+    GasUsed   *hexutil.Uint64 `json:"gasUsed,omitempty"`
+    Value     *hexutil.Big `json:"value,omitempty"`
+    Error     string `json:"error,omitempty"`
+    Calls     []callTrace `json:"calls,omitempty"`
+}
+
+// callTracerTest defines a single test to check the call tracer against.
+type callTracerTest struct {
+    Genesis *core.Genesis `json:"genesis"`
+    Context *callContext `json:"context"`
+    Input   string `json:"input"`
+    Result  *callTrace `json:"result"`
+}
+
+// Iterates over all the input-output datasets in the tracer test harness and
+// runs the JavaScript tracers against them.
+func TestCallTracerLegacy(t *testing.T) {
+    testCallTracer("callTracerLegacy", "call_tracer_legacy", t)
+}
+
+

```



```

+func TestCallTracerJs(t *testing.T) {
+    testCallTracer("callTracerJs", "call_tracer", t)
+}
+
+func TestCallTracerNative(t *testing.T) {
+    testCallTracer("callTracer", "call_tracer", t)
+}
+
+func testCallTracer(tracerName string, dirPath string, t *testing.T) {
+    files, err := ioutil.ReadDir(filepath.Join("testdata", dirPath))
+    if err != nil {
+        t.Fatalf("failed to retrieve tracer test suite: %v", err)
+    }
+    for _, file := range files {
+        if !strings.HasSuffix(file.Name(), ".json") {
+            continue
+        }
+        file := file // capture range variable
+        t.Run(camel(strings.TrimSuffix(file.Name(), ".json")), func(t *testing.T) {
+            t.Parallel()
+
+            var (
+                test = new(callTracerTest)
+                tx   = new(types.Transaction)
+            )
+            // Call tracer test found, read if from disk
+            if blob, err := ioutil.ReadFile(filepath.Join("testdata", dirPath, file.Name())); err != nil {
+                t.Fatalf("failed to read testcase: %v", err)
+            } else if err := json.Unmarshal(blob, test); err != nil {
+                t.Fatalf("failed to parse testcase: %v", err)
+            }
+            if err := rlp.DecodeBytes(common.FromHex(test.Input), tx); err != nil {
+                t.Fatalf("failed to parse testcase input: %v", err)
+            }
+            // Configure a blockchain with the given prestate
+            var (
+                signer    = types.MakeSigner(test.Genesis.Config, new(big.Int).SetUint64(uint64(test.Context.Number)), new(big.Int).SetUint64(uint64(test.Context.Time)))
+                origin, _ = signer.Sender(tx)
+                txContext = vm.TxContext{
+                    Origin:   origin,
+                    GasPrice: tx.GasPrice(),
+                }
+                context = vm.BlockContext{
+                    CanTransfer: core.CanTransfer,
+                    Transfer:    core.Transfer,
+                    Coinbase:    test.Context.Miner,
+                    BlockNumber: new(big.Int).SetUint64(uint64(test.Context.Number)),
+                    Time:        new(big.Int).SetUint64(uint64(test.Context.Time)),
+                    Difficulty:  (*big.Int)(test.Context.Difficulty),
+                    GasLimit:    uint64(test.Context.GasLimit),
+                }
+                _, statedb = tests.MakePreState(rawdb.NewMemoryDatabase(), test.Genesis.Alloc, false)
+            )
+            tracer, err := tracers.New(tracerName, new(tracers.Context))
+            if err != nil {
+                t.Fatalf("failed to create call tracer: %v", err)
+            }
+            evm := vm.NewEVM(context, txContext, statedb, test.Genesis.Config, vm.Config{Debug: true, Tracer: tracer})
+            msg, err := tx.AsMessage(signer, nil)
+            if err != nil {
+                t.Fatalf("failed to prepare transaction for tracing: %v", err)
+            }
+            st := core.NewStateTransition(evm, msg, new(core.GasPool).AddGas(tx.Gas()))
+            if _, err = st.TransitionDb(); err != nil {
+                t.Fatalf("failed to execute transaction: %v", err)
+            }
+            // Retrieve the trace result and compare against the etalon
+            res, err := tracer.GetResult()
+            if err != nil {
+                t.Fatalf("failed to retrieve trace result: %v", err)
+            }
+            ret := new(callTrace)
+            if err := json.Unmarshal(res, ret); err != nil {
+                t.Fatalf("failed to unmarshal trace result: %v", err)
+            }
+
+            if !jsonEqual(ret, test.Result) {
+                // uncomment this for easier debugging
+                //have, _ := json.MarshalIndent(ret, "", " ")
+                //want, _ := json.MarshalIndent(test.Result, "", " ")
+                //t.Fatalf("trace mismatch: \nhave %v\nwant %v", string(have), string(want))
+                t.Fatalf("trace mismatch: \nhave %v\nwant %v", ret, test.Result)
+            }
+        })
+    }
+}
+
+// jsonEqual is similar to reflect.DeepEqual, but does a 'bounce' via json prior to
+// comparison
+func jsonEqual(x, y interface{}) bool {
+    xTrace := new(callTrace)
+    yTrace := new(callTrace)
+    if xj, err := json.Marshal(x); err == nil {
+        json.Unmarshal(xj, xTrace)
+    } else {
+        return false
+    }
+    if yj, err := json.Marshal(y); err == nil {
+        json.Unmarshal(yj, yTrace)
+    } else {
+        return false
+    }
+    return reflect.DeepEqual(xTrace, yTrace)
+}
+
+// camel converts a snake cased input string into a camel cased output.
+func camel(str string) string {
+    pieces := strings.Split(str, "_")
+    for i := 1; i < len(pieces); i++ {
+        pieces[i] = string(unicode.ToUpper(rune(pieces[i][0]))) + pieces[i][1:]
+    }
+    return strings.Join(pieces, "")
+}
+
+func BenchmarkTracers(b *testing.B) {
+    files, err := ioutil.ReadDir(filepath.Join("testdata", "call_tracer"))
+    if err != nil {
+        b.Fatalf("failed to retrieve tracer test suite: %v", err)
+    }
+    for _, file := range files {
+        if !strings.HasSuffix(file.Name(), ".json") {
+            continue
+        }
+        file := file // capture range variable
+        b.Run(camel(strings.TrimSuffix(file.Name(), ".json")), func(b *testing.B) {
+            blob, err := ioutil.ReadFile(filepath.Join("testdata", "call_tracer", file.Name()))
+            if err != nil {
+                b.Fatalf("failed to read testcase: %v", err)
+            }
+            test := new(callTracerTest)
+            if err := json.Unmarshal(blob, test); err != nil {
+                b.Fatalf("failed to parse testcase: %v", err)
+            }
+            benchTracer("callTracerNative", test, b)
+        })
+    }
+}

```

```

+    }
+}
+
+func benchTracer(tracerName string, test *callTracerTest, b *testing.B) {
+    // Configure a blockchain with the given prestate
+    tx := new(types.Transaction)
+    if err := rlp.DecodeBytes(common.FromHex(test.Input), tx); err != nil {
+        b.Fatalf("failed to parse testcase input: %v", err)
+    }
+    signer := types.MakeSigner(test.Genesis.Config, new(big.Int).SetUint64(uint64(test.Context.Number)), new(big.Int).SetUint64(uint64(test.Context.Time)))
+    msg, err := tx.AsMessage(signer, nil)
+    if err != nil {
+        b.Fatalf("failed to prepare transaction for tracing: %v", err)
+    }
+    origin, _ := signer.Sender(tx)
+    txContext := vm.TxContext{
+        Origin:    origin,
+        GasPrice: tx.GasPrice(),
+    }
+    context := vm.BlockContext{
+        CanTransfer: core.CanTransfer,
+        Transfer:    core.Transfer,
+        Coinbase:    test.Context.Miner,
+        BlockNumber: new(big.Int).SetUint64(uint64(test.Context.Number)),
+        Time:        new(big.Int).SetUint64(uint64(test.Context.Time)),
+        Difficulty:   (*big.Int)(test.Context.Difficulty),
+        GasLimit:    uint64(test.Context.GasLimit),
+    }
+    _, statedb := tests.MakePreState(rawdb.NewMemoryDatabase(), test.Genesis.Alloc, false)
+
+    b.ReportAllocs()
+    b.ResetTimer()
+    for i := 0; i < b.N; i++ {
+        tracer, err := tracers.New(tracerName, new(tracers.Context))
+        if err != nil {
+            b.Fatalf("failed to create call tracer: %v", err)
+        }
+        evm := vm.NewEVM(context, txContext, statedb, test.Genesis.Config, vm.Config{Debug: true, Tracer: tracer})
+        snap := statedb.Snapshot()
+        st := core.NewStateTransition(evm, msg, new(core.GasPool).AddGas(tx.Gas()))
+        if _, err := st.TransitionDb(); err != nil {
+            b.Fatalf("failed to execute transaction: %v", err)
+        }
+        if _, err = tracer.GetResult(); err != nil {
+            b.Fatal(err)
+        }
+        statedb.RevertToSnapshot(snap)
+    }
+}
+
+// TestZeroValueToNotExitCall tests the calltracer(s) on the following:
+// Tx to A, A calls B with zero value. B does not already exist.
+// Expected: that enter/exit is invoked and the inner call is shown in the result
+func TestZeroValueToNotExitCall(t *testing.T) {
+    var to = common.HexToAddress("0x0000000000000000000000000000000000000000000000000000000000000000deadbeef")
+    privkey, err := crypto.HexToECDSA("0000000000000000000000000000000000000000000000000000000000000000deadbeef")
+    if err != nil {
+        t.Fatalf("err %v", err)
+    }
+    signer := types.NewEIP155Signer(big.NewInt(1))
+    tx, err := types.SignNewTx(privkey, signer, &types.LegacyTx{
+        GasPrice: big.NewInt(0),
+        Gas:      50000,
+        To:       &to,
+    })
+    if err != nil {
+        t.Fatalf("err %v", err)
+    }
+    origin, _ := signer.Sender(tx)
+    txContext := vm.TxContext{
+        Origin:    origin,
+        GasPrice: big.NewInt(1),
+    }
+    context := vm.BlockContext{
+        CanTransfer: core.CanTransfer,
+        Transfer:    core.Transfer,
+        Coinbase:    common.Address{},
+        BlockNumber: new(big.Int).SetUint64(8000000),
+        Time:        new(big.Int).SetUint64(5),
+        Difficulty:   big.NewInt(0x30000),
+        GasLimit:    uint64(6000000),
+    }
+    var code = []byte{
+        byte(vm.PUSH1), 0x0, byte(vm.DUP1), byte(vm.DUP1), byte(vm.DUP1), // in and outs zero
+        byte(vm.DUP1), byte(vm.PUSH1), 0xff, byte(vm.GAS), // value=0, address=0xff, gas=GAS
+        byte(vm.CALL),
+    }
+    var alloc = core.GenesisAlloc{
+        to: core.GenesisAccount{
+            Nonce: 1,
+            Code:  code,
+        },
+        origin: core.GenesisAccount{
+            Nonce: 0,
+            Balance: big.NewInt(5000000000000000),
+        },
+    }
+    _, statedb := tests.MakePreState(rawdb.NewMemoryDatabase(), alloc, false)
+    // Create the tracer, the EVM environment and run it
+    tracer, err := tracers.New("callTracer", nil)
+    if err != nil {
+        t.Fatalf("failed to create call tracer: %v", err)
+    }
+    evm := vm.NewEVM(context, txContext, statedb, params.AvalancheMainnetChainConfig, vm.Config{Debug: true, Tracer: tracer})
+    msg, err := tx.AsMessage(signer, nil)
+    if err != nil {
+        t.Fatalf("failed to prepare transaction for tracing: %v", err)
+    }
+    st := core.NewStateTransition(evm, msg, new(core.GasPool).AddGas(tx.Gas()))
+    if _, err := st.TransitionDb(); err != nil {
+        t.Fatalf("failed to execute transaction: %v", err)
+    }
+    // Retrieve the trace result and compare against the etalon
+    res, err := tracer.GetResult()
+    if err != nil {
+        t.Fatalf("failed to retrieve trace result: %v", err)
+    }
+    have := new(callTrace)
+    if err := json.Unmarshal(res, have); err != nil {
+        t.Fatalf("failed to unmarshal trace result: %v", err)
+    }
+    wantStr := `{"type": "CALL", "from": "0x682a80a6f560eec50d54e63cbda1c324c5f8d1b", "to": "0x0000000000000000000000000000000000000000000000000000000000000000deadbeef", "value": "0x0", "gas": "0x7148", "gasUsed": "0x2d0", "input": "0x", "output": "0x"}`
+    want := new(callTrace)
+    json.Unmarshal([]byte(wantStr), want)
+    if !json.Equal(have, want) {
+        t.Error("have != want")
+    }
+}
+
+diff --git a/eth/tracers/internal/tracetest/testdata/call_tracer/create.json b/eth/tracers/internal/tracetest/testdata/call_tracer/create.json
+new file mode 100644
+index 00000000..8699bf3e
+--- /dev/null
+++ b/eth/tracers/internal/tracetest/testdata/call_tracer/create.json
@@ -0,0 +1,58 @@

```

[illegible]

[illegible]

[illegible]

```
+         "to": "0x2ccc5e0538493c235d1c5ef6580f77d99e91396",
+         "type": "CALL",
+         "value": "0x0"
+     },
+     {
+         "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+         "gas": "0x17bec",
+         "gasUsed": "0x229",
+         "input": "0x2e94420f",
+         "output": "0x5842545553440000000000000000000000000000000000000000000000000000",
+         "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "type": "CALL",
+         "value": "0x0"
+     },
+     {
+         "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+         "gas": "0x1764e",
+         "gasUsed": "0x45c",
+         "input": "0xf92eb774584254555344000000000000000000000000000000000000000000000000",
+         "output": "0x0000000000000000000000000000000000000000000000000000000000000000",
+         "to": "0xc00ffdf997ad14939736f026006498e3f099baaf",
+         "type": "CALL",
+         "value": "0x0"
+     },
+     {
+         "calls": [
+             {
+                 "from": "0x2a98c5f40bfa3dee83431103c535f6fae9a8ad38",
+                 "gas": "0x108ba",
+                 "gasUsed": "0x24d",
+                 "input": "0x13bc6d4b00000000000000000000000000000000000000000000000000000000",
+                 "output": "0x0000000000000000000000000000000000000000000000000000000000000000",
+                 "to": "0x2ccc5e0538493c235d1c5ef6580f77d99e91396",
+                 "type": "CALL",
+                 "value": "0x0"
+             }
+         ],
+         "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+         "gas": "0x16e62",
+         "gasUsed": "0xebb",
+         "input": "0x645a3b72584254555344000000000000000000000000000000000000000000000000",
+         "output": "0x",
+         "to": "0x2a98c5f40bfa3dee83431103c535f6fae9a8ad38",
+         "type": "CALL",
+         "value": "0x0"
+     },
+     {
+         "from": "0xb4fe7aa695b326c9d219158d2ca50db77b39f99f",
+         "gas": "0x283b9",
+         "gasUsed": "0xc51c",
+         "input": "0x949ae47900000000000000000000000000000000000000000000000000000000",
+         "output": "0x",
+         "to": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+         "type": "CALL",
+         "value": "0x0"
+     },
+     {
+         "from": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "gas": "0x30b4a",
+         "gasUsed": "0xedb7",
+         "input": "0x51a34eb800000000000000000000000000000000000000000000000000000000",
+         "output": "0x",
+         "to": "0xb4fe7aa695b326c9d219158d2ca50db77b39f99f",
+         "type": "CALL",
+         "value": "0x0"
+     },
+ ],
+ "from": "0x70c9217d814985faef62b124420f8dfbddd96433",
+ "gas": "0x37b38",
+ "gasUsed": "0x12bb3",
+ "input": "0x51a34eb800000000000000000000000000000000000000000000000000000000",
+ "output": "0x",
+ "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+ "type": "CALL",
+ "value": "0x0"
+ }
+ )
+ )
+ diff --git a/eth/tracers/internal/tracetest/testdata/call_tracer/delegatecall.json b/eth/tracers/internal/tracetest/testdata/call_tracer/delegatecall.json
+ new file mode 100644
+ index 00000000..f7ad6df5
+ --- /dev/null
+ +++ b/eth/tracers/internal/tracetest/testdata/call_tracer/delegatecall.json
+ @@ -0,0 +1,97 @@
+ +{
+ +  "context": {
+ +    "difficulty": "31927752",
+ +    "gasLimit": "4707788",
+ +    "miner": "0x5659922ce141eedbc2733678f9806c77b4eebee8",
+ +    "number": "11495",
+ +    "timestamp": "1479735917"
+ +  },
+ +  "genesis": {
+ +    "alloc": {
+ +      "0x13204f5d64c28326fd7bd05fd4ea855302d7f2ff": {
+ +        "balance": "0x0",
+ +        "code": "0x606060405236156100825760e060020a60003504630a0313a981146100875780630a3b0a4f146101095780630cd40fea1461021257806329092d0e1461021f5780634cd06a5f146103295780635dbe47e8146103395780637a9e541f",
+ +        "nonce": "1",
+ +        "storage": {
+ +          "0x4d140b25abf3c71052885c66f73ce07cfff141c1afabffdaf5cba04d625b7ebcc": "0x0000000000000000000000000000000000000000000000000000000000000001"
+ +        }
+ +      },
+ +      "0x269296dddce321a6bcbaa2f0181127593d732cba": {
+ +        "balance": "0x0",
+ +        "code": "0x606060405236156101275760e060020a60003504630cd40fea811461012c578063173825d9146101395780631849cb5a146101c7578063285791371461030f5780632a58b3301461033f5780632cb0d48a146103565780632f54bf6f",
+ +        "nonce": "1",
+ +        "storage": {
+ +          "0x0000000000000000000000000000000000000000000000000000000000000001": "0x000113204f5d64c28326fd7bd05fd4ea855302d7f2ff0000000000000000000000000000000000000000000000000000000000000000"
+ +        }
+ +      },
+ +      "0x42b025deeb78f34cd5ac896473b63e6c99a71a2": {
+ +        "balance": "0x0",
+ +        "code": "0x6504032353da7150606060405236156100695760e060020a60003504631bf7509d811461006e57806321ce24d41461008157806333556e84146100ec578063685a1f3c146101035780637d65837a1461011757806389489a8714610",
+ +        "nonce": "1",
+ +        "storage": {}
+ +      }
+ +    },
+ +    "0xa529806c67cc6486d4d62024471772f47f6fd672": {
+ +      "balance": "0x67820e39ac8fe9800",
+ +      "code": "0x",
+ +      "nonce": "68",
+ +      "storage": {}
+ +    }
+ +  },
+ +  "config": {
+ +    "byzantiumBlock": 1700000,
+ +    "chainId": 3,
+ +    "daoForkSupport": true,
+ +    "eip150Block": 0,
+ +    "eip150Hash": "0x41941023680923e0fe4d7a34bdac8141f2540e3ae90623718e47d66d1ca4a2d",
+ +    "eip155Block": 10,
+ +    "eip158Block": 10,
+ +    "ethash": {},
+ +    "homesteadBlock": 0
+ +  },
+ +  "difficulty": "31912170",
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
diff --git a/eth/tracers/internal/tracetest/testdata/call_tracer_legacy/deep_calls.json b/eth/tracers/internal/tracetest/testdata/call_tracer_legacy/deep_calls.json
new file mode 100644
index 00000000..0353d4cf
--- /dev/null
+++ b/eth/tracers/internal/tracetest/testdata/call_tracer_legacy/deep_calls.json
@@ -0,0 +1,415 @@
+{
+  "context": {
+    "difficulty": "117066904",
+    "gasLimit": "4712384",
+    "miner": "0x1977c248e1014cc103929dd7f154199c916e39ec",
+    "number": "25001",
+    "timestamp": "1479891545"
+  },
+  "genesis": {
+    "alloc": {
+      "0x2a98c5f40bfa3dee83431103c535f6fae9a8ad38": {
+        "balance": "0x0",
+        "code": "0x606060405236156100825760e00020a600035046302d05d3f811461008a5780630accce061461009c5780631ab9075a146100c757806331ed274614610102578063645a3b7214610133578063772fdae314610155578063a7f4377f"
+      },
+      "0x0": {
+        "balance": "1",
+        "code": "0x"
+      }
+    }
+  }
+}
```

[illegible]

```

+ "gas": "0x32366",
+ "gasUsed": "0x273",
+ "input": "0x16c6cc60000000000000000000000000c212e03b9e060e36facad5fd8f4435412ca22e6b",
+ "output": "0x0000000000000000000000000000000000000000000000000000000000000001",
+ "to": "0x7986bad81f4cbd9317f5a46861437dae58d69113",
+ "type": "CALL",
+ "value": "0x0"
+ },
+ ],
+ {
+   "from": "0xb4fe7aa695b326c9d219158d2ca50db77b39f99f",
+   "gas": "0x29f35",
+   "gasUsed": "0xf8d",
+   "input": "0x16c6cc60000000000000000000000000c212e03b9e060e36facad5fd8f4435412ca22e6b",
+   "output": "0x0000000000000000000000000000000000000000000000000000000000000001",
+   "to": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+   "type": "CALL",
+   "value": "0x0"
+ },
+ },
+ {
+   "from": "0xb4fe7aa695b326c9d219158d2ca50db77b39f99f",
+   "gas": "0x28a9e",
+   "gasUsed": "0x334",
+   "input": "0xe16c7d98636f6e747261637463746c0000000000000000000000000000000000000000",
+   "output": "0x0000000000000000000000003e9286eafa2db8101246c2131c09b49080d00690",
+   "to": "0x2ccccf5e0538493c235d1c5ef6580f77d99e91396",
+   "type": "CALL",
+   "value": "0x0"
+ },
+ },
+ {
+   "calls": [
+     {
+       "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+       "gas": "0x21d79",
+       "gasUsed": "0x24d",
+       "input": "0x13bc6d4b000000000000000000000000b4fe7aa695b326c9d219158d2ca50db77b39f99f",
+       "output": "0x0000000000000000000000000000000000000000000000000000000000000001",
+       "to": "0x2ccccf5e0538493c235d1c5ef6580f77d99e91396",
+       "type": "CALL",
+       "value": "0x0"
+     },
+     {
+       "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+       "gas": "0x2165b",
+       "gasUsed": "0x334",
+       "input": "0xe16c7d986d61726b6574646200000000000000000000000000000000000000000000",
+       "output": "0x0000000000000000000000000000cf00ffd997ad14939736f026006498e3f099baaf",
+       "to": "0x2ccccf5e0538493c235d1c5ef6580f77d99e91396",
+       "type": "CALL",
+       "value": "0x0"
+     },
+     {
+       "calls": [
+         {
+           "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+           "gas": "0x1a8e8",
+           "gasUsed": "0x24d",
+           "input": "0x13bc6d4b00000000000000000000000000003e9286eafa2db8101246c2131c09b49080d00690",
+           "output": "0x0000000000000000000000000000000000000000000000000000000000000001",
+           "to": "0x2ccccf5e0538493c235d1c5ef6580f77d99e91396",
+           "type": "CALL",
+           "value": "0x0"
+         },
+         {
+           "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+           "gas": "0x1a2c6",
+           "gasUsed": "0x3cb",
+           "input": "0xc9503fe2",
+           "output": "0x000000000000000000000000000000000000000000000000000000008ac7230489e80000",
+           "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+           "type": "CALL",
+           "value": "0x0"
+         },
+       ],
+       {
+         "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+         "gas": "0x19b72",
+         "gasUsed": "0x3cb",
+         "input": "0xc9503fe2",
+         "output": "0x000000000000000000000000000000000000000000000000000000008ac7230489e80000",
+         "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "type": "CALL",
+         "value": "0x0"
+       },
+       {
+         "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+         "gas": "0x19428",
+         "gasUsed": "0x305",
+         "input": "0x6f265b93",
+         "output": "0x00000000000000000000000000000000000000000000000000000000283c7b9181eca20000",
+         "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "type": "CALL",
+         "value": "0x0"
+       },
+       {
+         "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+         "gas": "0x18d45",
+         "gasUsed": "0x229",
+         "input": "0xe94420f",
+         "output": "0x5842545534400000000000000000000000000000000000000000000000000000",
+         "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "type": "CALL",
+         "value": "0x0"
+       },
+       {
+         "from": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+         "gas": "0x1734e",
+         "gasUsed": "0x229",
+         "input": "0xe94420f",
+         "output": "0x5842545534400000000000000000000000000000000000000000000000000000",
+         "to": "0xc212e03b9e060e36facad5fd8f4435412ca22e6b",
+         "type": "CALL",
+         "value": "0x0"
+       }
+     ],
+     {
+       "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+       "gas": "0x20ee1",
+       "gasUsed": "0x5374",
+       "input": "0x581d5fd0000000000000000000000000c212e03b9e060e36facad5fd8f4435412ca22e6b0b00000000000000000000000000280faf689c35ac0000",
+       "output": "0x0",
+       "to": "0xcff0ffd997ad14939736f026006498e3f099baaf",
+       "type": "CALL",
+       "value": "0x0"
+     },
+     {
+       "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",
+       "gas": "0x1b6c1",
+       "gasUsed": "0x334",
+       "input": "0xe16c7d986c6f676d677200000000000000000000000000000000000000000000000000",
+       "output": "0x000000000000000000000000000000000000000000002a98c5f40bf3dee83431103c535f6fae9a8ad38",
+       "to": "0x2ccccf5e0538493c235d1c5ef6580f77d99e91396",
+       "type": "CALL",
+       "value": "0x0"
+     },
+     {
+       "from": "0x3e9286eafa2db8101246c2131c09b49080d00690",

```

```
diff --git a/eth/tracers/internal/tracetest/testdata/call tracer legacy/delegatecall.json b/eth/tracers/internal/tracetest/testdata/call tracer legacy/delegatecall.json
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```

+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package js
+
+// bigIntegerJS is the minified version of https://github.com/peterolson/BigInteger.js.
+const bigIntegerJS = `var bigInt=function(undefiend){"use strict";var BASE=1e7,LOG_BASE=7,MAX_INT=9007199254740992,MAX_INT_ARR=smallToArray(MAX_INT),LOG_MAX_INT=Math.log(MAX_INT);function Integer(v,radi:
diff --git a/eth/tracers/js/internal/tracers/4byte_tracer_legacy.js b/eth/tracers/js/internal/tracers/4byte_tracer_legacy.js
new file mode 100644
index 00000000..c27f9ae0
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/4byte_tracer_legacy.js
@@ -0,0 +1,96 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// 4byteTracer searches for 4byte-identifiers, and collects them for post-processing.
+// It collects the methods identifiers along with the size of the supplied data, so
+// a reversed signature can be matched against the size of the data.
+//
+// Example:
+// > debug.traceTransaction( "0x214e597e35da083692f5386141e69f47e973b2c56e7a8073b1ea08fd7571e9de", {tracer: "4byteTracer"})
+// {
+//   0x27dc297e-128: 1,
+//   0x38cc4831-0: 2,
+//   0x524f3889-96: 1,
+//   0xadf59f99-288: 1,
+//   0xc281d19e-0: 1
+// }
+
+({
+  // ids aggregates the 4byte ids found.
+  ids : {},
+
+  // callType returns 'false' for non-calls, or the peek-index for the first param
+  // after 'value', i.e. meminstart.
+  callType: function(opstr){
+    switch(opstr){
+      case "CALL": case "CALLCODE":
+        // gas, addr, val, memin, meminsz, memout, memoutsz
+        return 3; // stack ptr to memin
+
+      case "DELEGATECALL": case "STATICCALL":
+        // gas, addr, memin, meminsz, memout, memoutsz
+        return 2; // stack ptr to memin
+
+    }
+    return false;
+  },
+
+  // store save the given identifier and datasize.
+  store: function(id, size){
+    var key = "" + toHex(id) + "-" + size;
+    this.ids[key] = this.ids[key] + 1 || 1;
+  },
+
+  // step is invoked for every opcode that the VM executes.
+  step: function(log, db) {
+    // Skip any opcodes that are not internal calls
+    var ct = this.callType(log.op.toString());
+    if (!ct) {
+      return;
+    }
+    // Skip any pre-compile invocations, those are just fancy opcodes
+    if (isPrecompiled(toAddress(log.stack.peek(1).toString(16)))) {
+      return;
+    }
+    // Gather internal call details
+    var inSz = log.stack.peek(ct + 1).value0f();
+    if (inSz >= 4) {
+      var inOff = log.stack.peek(ct).value0f();
+      this.store(log.memory.slice(inOff, inOff + 4), inSz-4);
+    }
+  },
+
+  // fault is invoked when the actual execution of an opcode fails.
+  fault: function(log, db) { },
+
+  // result is invoked when all the opcodes have been iterated over and returns
+  // the final result of the tracing.
+  result: function(ctx) {
+    // Save the outer calldata also
+    if (ctx.input.length >= 4) {
+      this.store(slice(ctx.input, 0, 4), ctx.input.length-4)
+    }
+    return this.ids;
+  },
+})
diff --git a/eth/tracers/js/internal/tracers/assets.go b/eth/tracers/js/internal/tracers/assets.go
new file mode 100644
index 00000000..a2bb69de
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/assets.go
@@ -0,0 +1,181 @@
+// Code generated by go-bindata. DO NOT EDIT.
+// sources:
+// 4byte_tracer_legacy.js (2.933kB)
+// bigram_tracer.js (1.712kB)
+// call_tracer.js.js (3.497kB)
+// call_tracer_legacy.js (8.956kB)
+// evmdis_tracer.js (4.215kB)
+// noop_tracer.js (1.271kB)
+// opcount_tracer.js (1.372kB)
+// prestate_tracer.js (4.287kB)
+// trigram_tracer.js (1.788kB)
+// unigram_tracer.js (1.469kB)
+

```

```
+package tracers
+
+import (
+    "bytes"
+    "compress/gzip"
+    "crypto/sha256"
+    "fmt"
+    "io"
+    "io/ioutil"
+    "os"
+    "path/filepath"
+    "strings"
+    "time"
+)
+
+func bindataRead(data []byte, name string) ([]byte, error) {
+    gz, err := gzip.NewReader(bytes.NewBuffer(data))
+    if err != nil {
+        return nil, fmt.Errorf("read %q: %w", name, err)
+    }
+
+    var buf bytes.Buffer
+    _, err = io.Copy(&buf, gz)
+    clErr := gz.Close()
+
+    if err != nil {
+        return nil, fmt.Errorf("read %q: %w", name, err)
+    }
+    if clErr != nil {
+        return nil, err
+    }
+
+    return buf.Bytes(), nil
+}
+
+type asset struct {
+    bytes []byte
+    info  os.FileInfo
+    digest [sha256.Size]byte
+}
+
+type bindataFileInfo struct {
+    name      string
+    size      int64
+    mode      os.FileMode
+    modTime   time.Time
+}
+
+func (fi bindataFileInfo) Name() string {
+    return fi.name
+}
+
+func (fi bindataFileInfo) Size() int64 {
+    return fi.size
+}
+
+func (fi bindataFileInfo) Mode() os.FileMode {
+    return fi.mode
+}
+
+func (fi bindataFileInfo) ModTime() time.Time {
+    return fi.modTime
+}
+
+func (fi bindataFileInfo) IsDir() bool {
+    return false
+}
+
+func (fi bindataFileInfo) Sys() interface{} {
+    return nil
+}
+
+var __4byte_tracer_legacyJs = []byte("\x1f\x8b\x00\x00\x00\x00\x00\x00\xff\x94\x56\x5b\x6f\xdb\x4a\x0e\x7e\xb6\x7f\x05\xd7\x2f\xb5\x51\x59\x8e\x2f\x89\x2f\xd9\x16\xf0\xe6\xa4\x6d\x80\x9c\x24\x88\xdd"
+
+func _4byte_tracer_legacyJsBytes() ([]byte, error) {
+    return bindataRead(
+        __4byte_tracer_legacyJs,
+        "4byte_tracer_legacy.js",
+    )
+}
+
+func _4byte_tracer_legacyJs() (*asset, error) {
+    bytes, err := _4byte_tracer_legacyJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "4byte_tracer_legacy.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0xb4, 0xc5, 0x48, 0x2d, 0xd9, 0x43, 0x95, 0x93, 0x3b, 0x93, 0x2c, 0x47, 0x8c, 0x84, 0x32, 0x3c, 0x8b, 0x2e, 0xf3, 0x72, 0xc4, 0x57, 0xe6, 0:
+    return a, nil
+}
+
+var _bigram_tracerJs = []byte("\x1f\x8b\x00\x00\x00\x00\x00\x00\xff\x8c\x54\x5b\x6f\xdb\x36\x14\x7e\xf7\xaf\xf8\xde\x92\x20\xae\x4d\x6e\x2f\x83\x33\x0f\xd0\xb2\xa4\x35\x90\xda\x81\xad\xac\x30\x86\x3"
+
+func bigram_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _bigram_tracerJs,
+        "bigram_tracer.js",
+    )
+}
+
+func bigram_tracerJs() (*asset, error) {
+    bytes, err := bigram_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "bigram_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x77, 0x6c, 0xd, 0x24, 0xf2, 0x49, 0xbd, 0x58, 0x8b, 0xb5, 0xd1, 0xc9, 0xcd, 0xcf, 0x5b, 0x3e, 0x5c, 0xfb, 0x14, 0x50, 0xe7, 0xe3, 0xb9, 0x:
+    return a, nil
+}
+
+var _call_tracer_jsJs = []byte("\x1f\x8b\x00\x00\x00\x00\x00\x00\xff\x8c\x56\x5f\x6f\xdb\x38\x0c\x7f\x8e\x3f\x05\xaf\x0f\x4b\x82\x65\x71\xbb\x03\xf6\xd0\x2d\x03\x72\x45\xbb\x05\xe8\xb5\x45\x9a\xde\x"
+
+func call_tracer_jsJsBytes() ([]byte, error) {
+    return bindataRead(
+        _call_tracer_jsJs,
+        "call_tracer_js.js",
+    )
+}
+
+func call_tracer_jsJs() (*asset, error) {
+    bytes, err := call_tracer_jsJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "call_tracer_js.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x42, 0x13, 0x7a, 0x14, 0xbf, 0xa7, 0x49, 0x4f, 0xb4, 0x4f, 0x45, 0x1, 0xbc, 0x9e, 0xd1, 0x8e, 0xc7, 0xee, 0x61, 0xfa, 0x82, 0x52, 0xa4, 0x:
+    return a, nil
+}
+
+var _call_tracer_legacyJs = []byte("\x1f\x8b\x00\x00\x00\x00\x00\x00\xff\x4d\x5a\xdf\x6f\x1b\x37\xf2\x7f\x96\xfe\xa8\x89\x1f\xa6\xa0\x9\x51\x24\x39\xe9\xb7\x5f\xc0\xae\xa7\xa\x50\x1d\x25\x35\xe0\xc6\x81\x"
+
+func call_tracer_legacyJsBytes() ([]byte, error) {
+    return bindataRead(
+        _call_tracer_legacyJs,
+        "call_tracer_legacy.js",
+    )
+}
```

```
+}
+
+func call_tracer_legacyJs() (*asset, error) {
+    bytes, err := call_tracer_legacyJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "call_tracer_legacy.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x46, 0x79, 0xb6, 0xbc, 0xd2, 0xc, 0x25, 0xb1, 0x22, 0x56, 0xef, 0x77, 0xb9, 0x5e, 0x2e, 0xf4, 0xda, 0xb2, 0x2f, 0x53, 0xa4, 0xff, 0xc8, 0x:
+    return a, nil
+}
+
+var _evmdis_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\xac\x97\xf6\xf6\xda\x48\x13\xc0\x5f\xc3\xa7\x18\xe5\x15\xa8\x14\xb0\x31\x04\x9c\xcb\x49\x3c\x29\xe9\xe5\x51\x9a\x44\x40\xee\x54\xa:
+
+func evmdis_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _evmdis_tracerJs,
+        "evmdis_tracer.js",
+    )
+}
+
+func evmdis_tracerJs() (*asset, error) {
+    bytes, err := evmdis_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "evmdis_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x13, 0xeb, 0xca, 0x1f, 0x5f, 0xd3, 0x29, 0x81, 0xbb, 0xd8, 0xc8, 0x4a, 0x3a, 0x38, 0x10, 0xe2, 0xe7, 0xa4, 0xcd, 0xde, 0x78, 0x85, 0xc2, 0:
+    return a, nil
+}
+
+var _noop_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\x8c\x93\x4f\x6f\xdb\x46\x10\xc5\xcf\xe6\xa7\x78\xc7\x04\x50\xc5\xfe\x39\x14\x70\x8a\x02\xac\x61\x27\x2a\x1c\xdb\x90\xe8\x06\x3e\x0e\
+
+func noop_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _noop_tracerJs,
+        "noop_tracer.js",
+    )
+}
+
+func noop_tracerJs() (*asset, error) {
+    bytes, err := noop_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "noop_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0xe3, 0xf, 0x1c, 0x6f, 0x65, 0xaf, 0x90, 0x31, 0xab, 0xf, 0xe0, 0xca, 0x54, 0x7, 0xfd, 0xd3, 0xa1, 0x4a, 0x14, 0x1, 0x2a, 0x9d, 0xdc, 0xb9,
+    return a, nil
+}
+
+var _opcount_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\x8c\x94\xcf\x6e\xdb\x46\x10\x87\xcf\xe2\x53\xfc\x8e\x09\xa2\x92\x69\x7b\x28\xe0\x16\x05\x58\xc3\x4e\x04\xd8\xb2\x21\xd1\x09\x7c\x:
+
+func opcount_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _opcount_tracerJs,
+        "opcount_tracer.js",
+    )
+}
+
+func opcount_tracerJs() (*asset, error) {
+    bytes, err := opcount_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "opcount_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x27, 0xe, 0x97, 0x88, 0x9b, 0x53, 0xbb, 0x20, 0x44, 0xd8, 0xf5, 0xeb, 0x41, 0xd2, 0x7e, 0xd6, 0xda, 0x6b, 0xf5, 0xaf, 0x0, 0x75, 0x9f, 0xd:
+    return a, nil
+}
+
+var _prestate_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\x9c\x57\xdd\xf6\xdb\x38\x12\x7f\xb6\xfe\x8a\x41\x5f\x6c\x5d\x5d\xb9\xcd\x02\x7b\x80\x73\x39\x40\x75\xdd\x36\x40\x36\x09\x6c\xe7\
+
+func prestate_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _prestate_tracerJs,
+        "prestate_tracer.js",
+    )
+}
+
+func prestate_tracerJs() (*asset, error) {
+    bytes, err := prestate_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "prestate_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0xd4, 0x9, 0xf9, 0x44, 0x13, 0x31, 0x89, 0xf7, 0x35, 0x9a, 0xc6, 0xf0, 0x86, 0x9d, 0xb2, 0xe3, 0x57, 0xe2, 0xc0, 0xde, 0xc9, 0x3a, 0x4c, 0x:
+    return a, nil
+}
+
+var _trigram_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\x8c\x94\x41\xf6\xdb\xc6\x13\xc5\xef\xfa\x14\xef\x68\x23\xfa\x8b\xc9\xbf\x97\x42\x69\x0a\xb0\x86\xd\x08\x70\x64\x43\xa2\x1b\x18\x:
+
+func trigram_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _trigram_tracerJs,
+        "trigram_tracer.js",
+    )
+}
+
+func trigram_tracerJs() (*asset, error) {
+    bytes, err := trigram_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "trigram_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0x40, 0x63, 0xe1, 0x42, 0x60, 0x7, 0x1b, 0x79, 0x47, 0x1, 0xa1, 0xbf, 0xc4, 0x66, 0x19, 0x9b, 0x2b, 0x5a, 0x1f, 0x82, 0x3d, 0xcf, 0xee, 0xe:
+    return a, nil
+}
+
+var _unigram_tracerJs = []byte{"\x1f\x8b\x08\x00\x00\x00\x00\x00\xff\x8c\x94\x41\xf6\xdb\xc6\x13\xc5\xef\xfa\x14\xef\x68\x23\xfa\x8b\xc9\xbf\x97\x42\x69\x0a\xb0\x86\xd\x08\x70\x64\x43\xa2\x1b\x18\x:
+
+func unigram_tracerJsBytes() ([]byte, error) {
+    return bindataRead(
+        _unigram_tracerJs,
+        "unigram_tracer.js",
+    )
+}
+
+func unigram_tracerJs() (*asset, error) {
+    bytes, err := unigram_tracerJsBytes()
+    if err != nil {
+        return nil, err
+    }
+
+    info := bindataFileInfo{name: "unigram_tracer.js", size: 0, mode: os.FileMode(0), modTime: time.Unix(0, 0)}
+    a := &asset{bytes: bytes, info: info, digest: [32]uint8{0xc, 0xe6, 0x5c, 0x88, 0x18, 0xa7, 0x85, 0x61, 0x18, 0xc6, 0xec, 0x17, 0xfc, 0xdf, 0x9d, 0xc0, 0x1b, 0x49, 0xf8, 0x8d, 0xf1, 0xeb, 0x35, 0x:
+    return a, nil
+}
+
+// Asset loads and returns the asset for the given name.
```



```

+// It returns an error if the asset could not be found or
+// could not be loaded.
+func Asset(name string) ([]byte, error) {
+    canonicalName := strings.Replace(name, "\\", "/", -1)
+    if f, ok := _bindata[canonicalName]; ok {
+        a, err := f()
+        if err != nil {
+            return nil, fmt.Errorf("Asset %s can't read by error: %v", name, err)
+        }
+        return a.bytes, nil
+    }
+    return nil, fmt.Errorf("Asset %s not found", name)
+}
+
+// AssetString returns the asset contents as a string (instead of a []byte).
+func AssetString(name string) (string, error) {
+    data, err := Asset(name)
+    return string(data), err
+}
+
+// MustAsset is like Asset but panics when Asset would return an error.
+// It simplifies safe initialization of global variables.
+func MustAsset(name string) []byte {
+    a, err := Asset(name)
+    if err != nil {
+        panic("asset: Asset(" + name + "): " + err.Error())
+    }
+    return a
+}
+
+// MustAssetString is like AssetString but panics when Asset would return an
+// error. It simplifies safe initialization of global variables.
+func MustAssetString(name string) string {
+    return string(MustAsset(name))
+}
+
+// AssetInfo loads and returns the asset info for the given name.
+// It returns an error if the asset could not be found or
+// could not be loaded.
+func AssetInfo(name string) (os.FileInfo, error) {
+    canonicalName := strings.Replace(name, "\\", "/", -1)
+    if f, ok := _bindata[canonicalName]; ok {
+        a, err := f()
+        if err != nil {
+            return nil, fmt.Errorf("AssetInfo %s can't read by error: %v", name, err)
+        }
+        return a.info, nil
+    }
+    return nil, fmt.Errorf("AssetInfo %s not found", name)
+}
+
+// AssetDigest returns the digest of the file with the given name. It returns an
+// error if the asset could not be found or the digest could not be loaded.
+func AssetDigest(name string) ([sha256.Size]byte, error) {
+    canonicalName := strings.Replace(name, "\\", "/", -1)
+    if f, ok := _bindata[canonicalName]; ok {
+        a, err := f()
+        if err != nil {
+            return [sha256.Size]byte{}, fmt.Errorf("AssetDigest %s can't read by error: %v", name, err)
+        }
+        return a.digest, nil
+    }
+    return [sha256.Size]byte{}, fmt.Errorf("AssetDigest %s not found", name)
+}
+
+// Digests returns a map of all known files and their checksums.
+func Digests() (map[string][sha256.Size]byte, error) {
+    mp := make(map[string][sha256.Size]byte, len(_bindata))
+    for name := range _bindata {
+        a, err := _bindata[name]()
+        if err != nil {
+            return nil, err
+        }
+        mp[name] = a.digest
+    }
+    return mp, nil
+}
+
+// AssetNames returns the names of the assets.
+func AssetNames() []string {
+    names := make([]string, 0, len(_bindata))
+    for name := range _bindata {
+        names = append(names, name)
+    }
+    return names
+}
+
+// _bindata is a table, holding each asset generator, mapped to its name.
+var _bindata = map[string]func() (*asset, error){
+    "4byte_tracer_legacy.js": _4byte_tracer_legacyJs,
+    "bigram_tracer.js":      bigram_tracerJs,
+    "call_tracer.js.js":     call_tracer_jsJs,
+    "call_tracer_legacy.js": call_tracer_legacyJs,
+    "evmdis_tracer.js":      evmdis_tracerJs,
+    "noop_tracer.js":        noop_tracerJs,
+    "opcount_tracer.js":     opcount_tracerJs,
+    "prestate_tracer.js":    prestate_tracerJs,
+    "trigram_tracer.js":     trigram_tracerJs,
+    "unigram_tracer.js":     unigram_tracerJs,
+}
+
+// AssetDebug is true if the assets were built with the debug flag enabled.
+const AssetDebug = false
+
+// AssetDir returns the file names below a certain
+// directory embedded in the file by go-bindata.
+// For example if you run go-bindata on data/... and data contains the
+// following hierarchy:
+//     data/
+//     foo.txt
+//     img/
+//     a.png
+//     b.png
+// then AssetDir("data") would return []string{"foo.txt", "img"},
+// AssetDir("data/img") would return []string{"a.png", "b.png"},
+// AssetDir("foo.txt") and AssetDir("notexist") would return an error, and
+// AssetDir("") will return []string{"data"}.
+func AssetDir(name string) ([]string, error) {
+    node := _bintree
+    if len(name) != 0 {
+        canonicalName := strings.Replace(name, "\\", "/", -1)
+        pathList := strings.Split(canonicalName, "/")
+        for _, p := range pathList {
+            node = node.Children[p]
+            if node == nil {
+                return nil, fmt.Errorf("Asset %s not found", name)
+            }
+        }
+    }
+    if node.Func != nil {
+        return nil, fmt.Errorf("Asset %s not found", name)
+    }
+    rv := make([]string, 0, len(node.Children))

```

```

+         for childName := range node.Children {
+             rv = append(rv, childName)
+         }
+         return rv, nil
+     }
+ }
+
+type bintree struct {
+     Func      func() (*asset, error)
+     Children  map[string]*bintree
+ }
+
+var _bintree = &bintree(nil, map[string]*bintree{
+     "4byte_tracer_legacy.js": {4byte_tracer_legacyJs, map[string]*bintree{}},
+     "bigram_tracer.js":      {bigram_tracerJs, map[string]*bintree{}},
+     "call_tracer.js":       {call_tracer_jsJs, map[string]*bintree{}},
+     "call_tracer_legacy.js": {call_tracer_legacyJs, map[string]*bintree{}},
+     "evmdis_tracer.js":     {evmdis_tracerJs, map[string]*bintree{}},
+     "noop_tracer.js":       {noop_tracerJs, map[string]*bintree{}},
+     "opcount_tracer.js":    {opcount_tracerJs, map[string]*bintree{}},
+     "prestate_tracer.js":   {prestate_tracerJs, map[string]*bintree{}},
+     "trigram_tracer.js":    {trigram_tracerJs, map[string]*bintree{}},
+     "unigram_tracer.js":    {unigram_tracerJs, map[string]*bintree{}},
+ })
+
+// RestoreAsset restores an asset under the given directory.
+func RestoreAsset(dir, name string) error {
+     data, err := Asset(name)
+     if err != nil {
+         return err
+     }
+     info, err := AssetInfo(name)
+     if err != nil {
+         return err
+     }
+     err = os.MkdirAll(_filePath(dir, filepath.Dir(name)), os.FileMode(0755))
+     if err != nil {
+         return err
+     }
+     err = ioutil.WriteFile(_filePath(dir, name), data, info.Mode())
+     if err != nil {
+         return err
+     }
+     return os.Chtimes(_filePath(dir, name), info.ModTime(), info.ModTime())
+ }
+
+// RestoreAssets restores an asset under the given directory recursively.
+func RestoreAssets(dir, name string) error {
+     children, err := AssetDir(name)
+     // File
+     if err != nil {
+         return RestoreAsset(dir, name)
+     }
+     // Dir
+     for _, child := range children {
+         err = RestoreAssets(dir, filepath.Join(name, child))
+         if err != nil {
+             return err
+         }
+     }
+     return nil
+ }
+
+func _filePath(dir, name string) string {
+     canonicalName := strings.Replace(name, "\\", "/", -1)
+     return filepath.Join(append([]string{dir}, strings.Split(canonicalName, "/")...))
+ }
+
diff --git a/eth/tracers/js/internal/tracers/bigram_tracer.js b/eth/tracers/js/internal/tracers/bigram_tracer.js
new file mode 100644
index 00000000..24f9af95
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/bigram_tracer.js
@@ -0,0 +1,57 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2018 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+{
+     // hist is the counters of opcode bigrams
+     hist: {},
+     // lastOp is last operation
+     lastOp: '',
+     // execution depth of last op
+     lastDepth: 0,
+     // step is invoked for every opcode that the VM executes.
+     step: function(log, db) {
+         var op = log.op.toString();
+         var depth = log.getDepth();
+         if (depth == this.lastDepth){
+             var key = this.lastOp+'-'+op;
+             if (this.hist[key]){
+                 this.hist[key]++;
+             }
+             else {
+                 this.hist[key] = 1;
+             }
+         }
+         this.lastOp = op;
+         this.lastDepth = depth;
+     },
+     // fault is invoked when the actual execution of an opcode fails.
+     fault: function(log, db) {},
+     // result is invoked when all the opcodes have been iterated over and returns
+     // the final result of the tracing.
+     result: function(ctx) {
+         return this.hist;
+     }
+ },
+
diff --git a/eth/tracers/js/internal/tracers/call_tracer.js b/eth/tracers/js/internal/tracers/call_tracer.js
new file mode 100644
index 00000000..31c20a28
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/call_tracer.js

```

```

@@ -0,0 +1,122 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2021 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+
+// callFrameTracer uses the new call frame tracing methods to report useful information
+// about internal messages of a transaction.
+{
+    callstack: [],
+    fault: function(log, db) {},
+    result: function(ctx, db) {
+        // Prepare outer message info
+        var result = {
+            type:      ctx.type,
+            from:      toHex(ctx.from),
+            to:        toHex(ctx.to),
+            value:     '0x' + ctx.value.toString(16),
+            gas:       '0x' + bigInt(ctx.gas).toString(16),
+            gasUsed:   '0x' + bigInt(ctx.gasUsed).toString(16),
+            input:     toHex(ctx.input),
+            output:    toHex(ctx.output),
+        }
+        if (this.callstack[0].calls !== undefined) {
+            result.calls = this.callstack[0].calls
+        }
+        if (this.callstack[0].error !== undefined) {
+            result.error = this.callstack[0].error
+        } else if (ctx.error !== undefined) {
+            result.error = ctx.error
+        }
+        if (result.error !== undefined && (result.error !== "execution reverted" || result.output === "0x")) {
+            delete result.output
+        }
+        return this.finalize(result)
+    },
+    enter: function(frame) {
+        var call = {
+            type: frame.getType(),
+            from: toHex(frame.getFrom()),
+            to: toHex(frame.getTo()),
+            input: toHex(frame.getInput()),
+            gas: '0x' + bigInt(frame.getGas()).toString('16'),
+        }
+        if (frame.getValue() !== undefined){
+            call.value='0x' + bigInt(frame.getValue()).toString(16)
+        }
+        this.callstack.push(call)
+    },
+    exit: function(frameResult) {
+        var len = this.callstack.length
+        if (len > 1) {
+            var call = this.callstack.pop()
+            call.gasUsed = '0x' + bigInt(frameResult.getGasUsed()).toString('16')
+            var error = frameResult.getError()
+            if (error === undefined) {
+                call.output = toHex(frameResult.getOutput())
+            } else {
+                call.error = error
+                if (call.type === 'CREATE' || call.type === 'CREATE2') {
+                    delete call.to
+                }
+            }
+            len -= 1
+            if (this.callstack[len-1].calls === undefined) {
+                this.callstack[len-1].calls = []
+            }
+            this.callstack[len-1].calls.push(call)
+        }
+    },
+    // finalize recreates a call object using the final desired field order for json
+    // serialization. This is a nicety feature to pass meaningfully ordered results
+    // to users who don't interpret it, just display it.
+    finalize: function(call) {
+        var sorted = {
+            type:      call.type,
+            from:      call.from,
+            to:        call.to,
+            value:     call.value,
+            gas:       call.gas,
+            gasUsed:   call.gasUsed,
+            input:     call.input,
+            output:    call.output,
+            error:     call.error,
+            time:      call.time,
+            calls:     call.calls,
+        }
+        for (var key in sorted) {
+            if (sorted[key] === undefined) {
+                delete sorted[key]
+            }
+        }
+        if (sorted.calls !== undefined) {
+            for (var i=0; i<sorted.calls.length; i++) {
+                sorted.calls[i] = this.finalize(sorted.calls[i])
+            }
+        }
+        return sorted
+    }
+}
+
diff --git a/eth/tracers/js/internal/tracers/call_tracer_legacy.js b/eth/tracers/js/internal/tracers/call_tracer_legacy.js
new file mode 100644
index 00000000..7081dffa
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/call_tracer_legacy.js
@@ -0,0 +1,262 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//

```

```

+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// callTracer is a full blown transaction tracer that extracts and reports all
+// the internal calls made by a transaction, along with any useful information.
+{
+    // callstack is the current recursive call stack of the EVM execution.
+    callstack: {},
+
+    // descended tracks whether we've just descended from an outer transaction into
+    // an inner call.
+    descended: false,
+
+    // step is invoked for every opcode that the VM executes.
+    step: function(log, db) {
+        // Capture any errors immediately
+        var error = log.getError();
+        if (error !== undefined) {
+            this.fault(log, db);
+            return;
+        }
+        // We only care about system opcodes, faster if we pre-check once
+        var syscall = (log.op.toNumber() & 0xf0) == 0xf0;
+        if (syscall) {
+            var op = log.op.toString();
+
+            // If a new contract is being created, add to the call stack
+            if (syscall && (op == 'CREATE' || op == "CREATE2")) {
+                var in0ff = log.stack.peek(1).valueOf();
+                var inEnd = in0ff + log.stack.peek(2).valueOf();
+
+                // Assemble the internal call report and store for completion
+                var call = {
+                    type:    op,
+                    from:    toHex(log.contract.getAddress()),
+                    input:    toHex(log.memory.slice(in0ff, inEnd)),
+                    gasIn:    log.getGas(),
+                    gasCost: log.getCost(),
+                    value:    '0x' + log.stack.peek(0).toString(16)
+                };
+                this.callstack.push(call);
+                this.descended = true
+                return;
+            }
+
+            // If a contract is being self destructed, gather that as a subcall too
+            if (syscall && op == 'SELFDESTRUCT') {
+                var left = this.callstack.length;
+                if (this.callstack[left-1].calls === undefined) {
+                    this.callstack[left-1].calls = [];
+                }
+                this.callstack[left-1].calls.push({
+                    type:    op,
+                    from:    toHex(log.contract.getAddress()),
+                    to:      toHex(toAddress(log.stack.peek(0).toString(16))),
+                    gasIn:    log.getGas(),
+                    gasCost: log.getCost(),
+                    value:    '0x' + db.getBalance(log.contract.getAddress()).toString(16)
+                });
+                return
+            }
+
+            // If a new method invocation is being done, add to the call stack
+            if (syscall && (op == 'CALL' || op == 'CALLCODE' || op == 'DELEGATECALL' || op == 'STATICCALL')) {
+                // Skip any pre-compile invocations, those are just fancy opcodes
+                var to = toAddress(log.stack.peek(1).toString(16));
+                if (isPrecompiled(to)) {
+                    return
+                }
+                var off = (op == 'DELEGATECALL' || op == 'STATICCALL' ? 0 : 1);
+
+                var in0ff = log.stack.peek(2 + off).valueOf();
+                var inEnd = in0ff + log.stack.peek(3 + off).valueOf();
+
+                // Assemble the internal call report and store for completion
+                var call = {
+                    type:    op,
+                    from:    toHex(log.contract.getAddress()),
+                    to:      toHex(to),
+                    input:    toHex(log.memory.slice(in0ff, inEnd)),
+                    gasIn:    log.getGas(),
+                    gasCost: log.getCost(),
+                    out0ff:   log.stack.peek(4 + off).valueOf(),
+                    outLen:   log.stack.peek(5 + off).valueOf()
+                };
+                if (op != 'DELEGATECALL' && op != 'STATICCALL') {
+                    call.value = '0x' + log.stack.peek(2).toString(16);
+                }
+                this.callstack.push(call);
+                this.descended = true
+                return;
+            }
+
+            // If we've just descended into an inner call, retrieve it's true allowance. We
+            // need to extract if from within the call as there may be funky gas dynamics
+            // with regard to requested and actually given gas (2300 stipend, 63/64 rule).
+            if (this.descended) {
+                if (log.getDepth() >= this.callstack.length) {
+                    this.callstack[this.callstack.length - 1].gas = log.getGas();
+                } else {
+                    // TODO(karalabe): The call was made to a plain account. We currently don't
+                    // have access to the true gas amount inside the call and so any amount will
+                    // mostly be wrong since it depends on a lot of input args. Skip gas for now.
+                }
+                this.descended = false;
+            }
+
+            // If an existing call is returning, pop off the call stack
+            if (syscall && op == 'REVERT') {
+                this.callstack[this.callstack.length - 1].error = "execution reverted";
+                return;
+            }
+
+            if (log.getDepth() == this.callstack.length - 1) {
+                // Pop off the last call and get the execution results
+                var call = this.callstack.pop();
+
+                if (call.type == 'CREATE' || call.type == "CREATE2") {
+                    // If the call was a CREATE, retrieve the contract address and output code
+                    call.gasUsed = '0x' + bigInt(call.gasIn - call.gasCost - log.getGas()).toString(16);
+                    delete call.gasIn; delete call.gasCost;
+                }
+            }
+        }
+    }
+}

```

```

+         var ret = log.stack.peek(0);
+         if (!ret.equals(0)) {
+             call.to = toHex(toAddress(ret.toString(16)));
+             call.output = toHex(db.getCode(toAddress(ret.toString(16))));
+         } else if (call.error === undefined) {
+             call.error = "internal failure"; // TODO(karalabe): surface these faults somehow
+         }
+     } else {
+         // If the call was a contract call, retrieve the gas usage and output
+         if (call.gas !== undefined) {
+             call.gasUsed = '0x' + bigInt(call.gasIn - call.gasCost + call.gas - log.getGas()).toString(16);
+         }
+         var ret = log.stack.peek(0);
+         if (!ret.equals(0)) {
+             call.output = toHex(log.memory.slice(call.outOff, call.outOff + call.outLen));
+         } else if (call.error === undefined) {
+             call.error = "internal failure"; // TODO(karalabe): surface these faults somehow
+         }
+         delete call.gasIn; delete call.gasCost;
+         delete call.outOff; delete call.outLen;
+     }
+     if (call.gas !== undefined) {
+         call.gas = '0x' + bigInt(call.gas).toString(16);
+     }
+     // Inject the call into the previous one
+     var left = this.callstack.length;
+     if (this.callstack[left-1].calls === undefined) {
+         this.callstack[left-1].calls = [];
+     }
+     this.callstack[left-1].calls.push(call);
+ },
+
+ // fault is invoked when the actual execution of an opcode fails.
+ fault: function(log, db) {
+     // If the topmost call already reverted, don't handle the additional fault again
+     if (this.callstack[this.callstack.length - 1].error !== undefined) {
+         return;
+     }
+     // Pop off the just failed call
+     var call = this.callstack.pop();
+     call.error = log.getError();
+
+     // Consume all available gas and clean any leftovers
+     if (call.gas !== undefined) {
+         call.gas = '0x' + bigInt(call.gas).toString(16);
+         call.gasUsed = call.gas
+     }
+     delete call.gasIn; delete call.gasCost;
+     delete call.outOff; delete call.outLen;
+
+     // Flatten the failed call into its parent
+     var left = this.callstack.length;
+     if (left > 0) {
+         if (this.callstack[left-1].calls === undefined) {
+             this.callstack[left-1].calls = [];
+         }
+         this.callstack[left-1].calls.push(call);
+         return;
+     }
+     // Last call failed too, leave it in the stack
+     this.callstack.push(call);
+ },
+
+ // result is invoked when all the opcodes have been iterated over and returns
+ // the final result of the tracing.
+ result: function(ctx, db) {
+     var result = {
+         type:      ctx.type,
+         from:      toHex(ctx.from),
+         to:        toHex(ctx.to),
+         value:     '0x' + ctx.value.toString(16),
+         gas:       '0x' + bigInt(ctx.gas).toString(16),
+         gasUsed:   '0x' + bigInt(ctx.gasUsed).toString(16),
+         input:     toHex(ctx.input),
+         output:    toHex(ctx.output),
+         time:      ctx.time,
+     };
+     if (this.callstack[0].calls !== undefined) {
+         result.calls = this.callstack[0].calls;
+     }
+     if (this.callstack[0].error !== undefined) {
+         result.error = this.callstack[0].error;
+     } else if (ctx.error !== undefined) {
+         result.error = ctx.error;
+     }
+     if (result.error !== undefined && (result.error !== "execution reverted" || result.output === "0x")) {
+         delete result.output;
+     }
+     return this.finalize(result);
+ },
+
+ // finalize recreates a call object using the final desired field order for json
+ // serialization. This is a nicety feature to pass meaningfully ordered results
+ // to users who don't interpret it, just display it.
+ finalize: function(call) {
+     var sorted = {
+         type:      call.type,
+         from:      call.from,
+         to:        call.to,
+         value:     call.value,
+         gas:       call.gas,
+         gasUsed:   call.gasUsed,
+         input:     call.input,
+         output:    call.output,
+         error:     call.error,
+         time:      call.time,
+         calls:     call.calls,
+     };
+     for (var key in sorted) {
+         if (sorted[key] === undefined) {
+             delete sorted[key];
+         }
+     }
+     if (sorted.calls !== undefined) {
+         for (var i=0; i<sorted.calls.length; i++) {
+             sorted.calls[i] = this.finalize(sorted.calls[i]);
+         }
+     }
+     return sorted;
+ }
+ }
+
+diff --git a/eth/tracers/js/internal/tracers/evmdis_tracer.js b/eth/tracers/js/internal/tracers/evmdis_tracer.js
new file mode 100644
index 00000000..db3422ed
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/evmdis_tracer.js
@@ -0,0 +1,103 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+diff
```

```

+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// evmdisTracer returns sufficient information from a trace to perform evmdis-style
+// disassembly.
+{
+    stack: [{ops: []}],
+
+    npushes: {0: 0, 1: 1, 2: 1, 3: 1, 4: 1, 5: 1, 6: 1, 7: 1, 8: 1, 9: 1, 10: 1, 11: 1, 16: 1, 17: 1, 18: 1, 19: 1, 20: 1, 21: 1, 22: 1, 23: 1, 24: 1, 25: 1, 26: 1, 32: 1, 48: 1, 49: 1, 50: 1, 51: 1,
+
+    // result is invoked when all the opcodes have been iterated over and returns
+    // the final result of the tracing.
+    result: function() { return this.stack[0].ops; },
+
+    // fault is invoked when the actual execution of an opcode fails.
+    fault: function(log, db) { },
+
+    // step is invoked for every opcode that the VM executes.
+    step: function(log, db) {
+        var frame = this.stack[this.stack.length - 1];
+
+        var error = log.getError();
+        if (error) {
+            frame["error"] = error;
+        } else if (log.getDepth() == this.stack.length) {
+            opinfo = {
+                op: log.op.toNumber(),
+                depth: log.getDepth(),
+                result: [],
+            };
+            if (frame.ops.length > 0) {
+                var prevop = frame.ops[frame.ops.length - 1];
+                for(var i = 0; i < this.npushes[prevop.op]; i++)
+                    prevop.result.push(log.stack.peek(i).toString(16));
+            }
+            switch(log.op.toString()) {
+            case "CALL": case "CALLCODE":
+                var instart = log.stack.peek(3).valueOf();
+                var insize = log.stack.peek(4).valueOf();
+                opinfo["gas"] = log.stack.peek(0).valueOf();
+                opinfo["to"] = log.stack.peek(1).toString(16);
+                opinfo["value"] = log.stack.peek(2).toString();
+                opinfo["input"] = log.memory.slice(instart, instart + insize);
+                opinfo["error"] = null;
+                opinfo["return"] = null;
+                opinfo["ops"] = [];
+                this.stack.push(opinfo);
+                break;
+            case "DELEGATECALL": case "STATICCALL":
+                var instart = log.stack.peek(2).valueOf();
+                var insize = log.stack.peek(3).valueOf();
+                opinfo["op"] = log.op.toString();
+                opinfo["gas"] = log.stack.peek(0).valueOf();
+                opinfo["to"] = log.stack.peek(1).toString(16);
+                opinfo["input"] = log.memory.slice(instart, instart + insize);
+                opinfo["error"] = null;
+                opinfo["return"] = null;
+                opinfo["ops"] = [];
+                this.stack.push(opinfo);
+                break;
+            case "RETURN": case "REVERT":
+                var out = log.stack.peek(0).valueOf();
+                var outsize = log.stack.peek(1).valueOf();
+                frame.return = log.memory.slice(out, out + outsize);
+                break;
+            case "STOP": case "SELFDESTRUCT":
+                frame.return = log.memory.slice(0, 0);
+                break;
+            case "JUMPDEST":
+                opinfo["pc"] = log.getPC();
+            }
+            if(log.op.isPush()) {
+                opinfo["len"] = log.op.toNumber() - 0x5e;
+            }
+            frame.ops.push(opinfo);
+        } else {
+            this.stack = this.stack.slice(0, log.getDepth());
+        }
+    }
+}
+
+diff --git a/eth/tracers/js/internal/tracers/noop_tracer.js b/eth/tracers/js/internal/tracers/noop_tracer.js
new file mode 100644
index 00000000..c6881a43
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/noop_tracer.js
@@ -0,0 +1,39 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// noopTracer is just the barebone boilerplate code required from a JavaScript
+// object to be usable as a transaction tracer.
+{
+    // step is invoked for every opcode that the VM executes.

```

```

+         step: function(log, db) { },
+
+         // fault is invoked when the actual execution of an opcode fails.
+         fault: function(log, db) { },
+
+         // result is invoked when all the opcodes have been iterated over and returns
+         // the final result of the tracing.
+         result: function(ctx, db) { return {}; }
+    }
+}
diff --git a/eth/tracers/js/internal/tracers/opcount_tracer.js b/eth/tracers/js/internal/tracers/opcount_tracer.js
new file mode 100644
index 00000000..b0b307fa
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/opcount_tracer.js
@@ -0,0 +1,42 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// opcountTracer is a sample tracer that just counts the number of instructions
+// executed by the EVM before the transaction terminated.
+{
+    // count tracks the number of EVM instructions executed.
+    count: 0,
+
+    // step is invoked for every opcode that the VM executes.
+    step: function(log, db) { this.count++ },
+
+    // fault is invoked when the actual execution of an opcode fails.
+    fault: function(log, db) { },
+
+    // result is invoked when all the opcodes have been iterated over and returns
+    // the final result of the tracing.
+    result: function(ctx, db) { return this.count }
+}
diff --git a/eth/tracers/js/internal/tracers/prestate_tracer.js b/eth/tracers/js/internal/tracers/prestate_tracer.js
new file mode 100644
index 00000000..a264b7f1
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/prestate_tracer.js
@@ -0,0 +1,118 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// prestateTracer outputs sufficient information to create a local execution of
+// the transaction from a custom assembled genesis block.
+{
+    // prestate is the genesis that we're building.
+    prestate: null,
+
+    // lookupAccount injects the specified account into the prestate object.
+    lookupAccount: function(addr, db){
+        var acc = toHex(addr);
+        if (this.prestate[acc] === undefined) {
+            this.prestate[acc] = {
+                balance: '0x' + db.getBalance(addr).toString(16),
+                nonce: db.getNonce(addr),
+                code: toHex(db.getCode(addr)),
+                storage: {}
+            };
+        }
+    },
+
+    // lookupStorage injects the specified storage entry of the given account into
+    // the prestate object.
+    lookupStorage: function(addr, key, db){
+        var acc = toHex(addr);
+        var idx = toHex(key);
+
+        if (this.prestate[acc].storage[idx] === undefined) {
+            this.prestate[acc].storage[idx] = toHex(db.getState(addr, key));
+        }
+    },
+
+    // result is invoked when all the opcodes have been iterated over and returns
+    // the final result of the tracing.
+    result: function(ctx, db) {
+        // At this point, we need to deduct the 'value' from the
+        // outer transaction, and move it back to the origin
+        this.lookupAccount(ctx.from, db);
+
+        var fromBal = bigInt(this.prestate[toHex(ctx.from)].balance.slice(2), 16);
+        var toBal = bigInt(this.prestate[toHex(ctx.to)].balance.slice(2), 16);
+
+        this.prestate[toHex(ctx.to)].balance = '0x'+toBal.subtract(ctx.value).toString(16);
+        this.prestate[toHex(ctx.from)].balance = '0x'+fromBal.add(ctx.value).add((ctx.gasUsed + ctx.intrinsicGas) * ctx.gasPrice).toString(16);
+
+        // Decrement the caller's nonce, and remove empty create targets
+        this.prestate[toHex(ctx.from)].nonce--;
+    }
+}

```

```

+         if (ctx.type == 'CREATE') {
+             // We can blibldy delete the contract prestate, as any existing state would
+             // have caused the transaction to be rejected as invalid in the first place.
+             delete this.prestate[toHex(ctx.to)];
+         }
+         // Return the assembled allocations (prestate)
+         return this.prestate;
+     },
+
+     // step is invoked for every opcode that the VM executes.
+     step: function(log, db) {
+         // Add the current account if we just started tracing
+         if (this.prestate == null){
+             this.prestate = {};
+             // Balance will potentially be wrong here, since this will include the value
+             // sent along with the message. We fix that in 'result()'.
+             this.lookupAccount(log.contract.getAddress(), db);
+         }
+         // Whenever new state is accessed, add it to the prestate
+         switch (log.op.toString()) {
+             case "EXTCODECOPY": case "EXTCODESIZE": case "BALANCE":
+                 this.lookupAccount(toAddress(log.stack.peek(0).toString(16)), db);
+                 break;
+             case "CREATE":
+                 var from = log.contract.getAddress();
+                 this.lookupAccount(toContract(from, db.getNonce(from)), db);
+                 break;
+             case "CREATE2":
+                 var from = log.contract.getAddress();
+                 // stack: salt, size, offset, endowment
+                 var offset = log.stack.peek(1).valueOf()
+                 var size = log.stack.peek(2).valueOf()
+                 var end = offset + size
+                 this.lookupAccount(toContract2(from, log.stack.peek(3).toString(16), log.memory.slice(offset, end)), db);
+                 break;
+             case "CALL": case "CALLCODE": case "DELEGATECALL": case "STATICCALL":
+                 this.lookupAccount(toAddress(log.stack.peek(1).toString(16)), db);
+                 break;
+             case 'SSTORE':case 'SLOAD':
+                 this.lookupStorage(log.contract.getAddress(), toWord(log.stack.peek(0).toString(16)), db);
+                 break;
+         }
+     },
+
+     // fault is invoked when the actual execution of an opcode fails.
+     fault: function(log, db) {}
+ })

```

diff --git a/eth/tracers/js/internal/tracers/tracers.go b/eth/tracers/js/internal/tracers/tracers.go

new file mode 100644

index 00000000..c921db72

--- /dev/null

+++ b/eth/tracers/js/internal/tracers/tracers.go

```

@@ -0,0 +1,31 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+//go:generate go-bindata -nometadata -o assets.go -pkg tracers -ignore tracers.go -ignore assets.go ./...
+//go:generate gofmt -s -w assets.go
+
+// Package tracers contains the actual JavaScript tracer assets.
+package tracers

```

diff --git a/eth/tracers/js/internal/tracers/trigram_tracer.js b/eth/tracers/js/internal/tracers/trigram_tracer.js

new file mode 100644

index 00000000..f43c690b

--- /dev/null

+++ b/eth/tracers/js/internal/tracers/trigram_tracer.js

```

@@ -0,0 +1,59 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2018 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+{
+    // hist is the map of trigram counters
+    hist: {},
+    // lastOp is last operation
+    lastOps: ['', ''],
+    lastDepth: 0,
+
+    // step is invoked for every opcode that the VM executes.
+    step: function(log, db) {
+        var depth = log.getDepth();
+        if (depth != this.lastDepth){
+            this.lastOps = ['', ''];
+            this.lastDepth = depth;
+            return;
+        }
+        var op = log.op.toString();
+        var key = this.lastOps[0]+'-'+this.lastOps[1]+'-'+op;
+        if (this.hist[key]){
+            this.hist[key]++;
+        }
+    }
+}

```



```

+     }
+     else {
+         this.hist[key] = 1;
+     }
+     this.lastOps[0] = this.lastOps[1];
+     this.lastOps[1] = op;
+ },
+ // fault is invoked when the actual execution of an opcode fails.
+ fault: function(log, db) {},
+ // result is invoked when all the opcodes have been iterated over and returns
+ // the final result of the tracing.
+ result: function(ctx) {
+     return this.hist;
+ },
+ },
+}
diff --git a/eth/tracers/js/internal/tracers/unigram_tracer.js b/eth/tracers/js/internal/tracers/unigram_tracer.js
new file mode 100644
index 00000000..8ca82381
--- /dev/null
+++ b/eth/tracers/js/internal/tracers/unigram_tracer.js
@@ -0,0 +1,51 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2018 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+{
+    // hist is the map of opcodes to counters
+    hist: {},
+    // nops counts number of ops
+    nops: 0,
+    // step is invoked for every opcode that the VM executes.
+    step: function(log, db) {
+        var op = log.op.toString();
+        if (this.hist[op]){
+            this.hist[op]++;
+        }
+        else {
+            this.hist[op] = 1;
+        }
+        this.nops++;
+    },
+    // fault is invoked when the actual execution of an opcode fails.
+    fault: function(log, db) {},
+
+    // result is invoked when all the opcodes have been iterated over and returns
+    // the final result of the tracing.
+    result: function(ctx) {
+        return this.hist;
+    },
+},
+}
diff --git a/eth/tracers/js/tracer.go b/eth/tracers/js/tracer.go
new file mode 100644
index 00000000..4fdf7904
--- /dev/null
+++ b/eth/tracers/js/tracer.go
@@ -0,0 +1,985 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2017 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+// package js is a collection of tracers written in javascript.
+package js
+
+import (
+    "encoding/json"
+    "errors"
+    "fmt"
+    "math/big"
+    "strings"
+    "sync/atomic"
+    "time"
+    "unicode"
+    "unsafe"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/vm"
+    tracers2 "github.com/flare-foundation/coreth/eth/tracers"
+    "github.com/flare-foundation/coreth/eth/tracers/js/internal/tracers"
+    "gopkg.in/olebedev/go-duktape.v3"
+)
+
+// camel converts a snake cased input string into a camel cased output.
+func camel(str string) string {
+    pieces := strings.Split(str, "_")
+    for i := 1; i < len(pieces); i++ {
+        pieces[i] = string(unicode.ToUpper(rune(pieces[i][0]))) + pieces[i][1:]
+    }

```

```

+     }
+     return strings.Join(pieces, "")
+}
+
+var assetTracers = make(map[string]string)
+
+// init retrieves the JavaScript transaction tracers included in go-ethereum.
+func init() {
+    for _, file := range tracers.AssetNames() {
+        name := camel(strings.TrimSuffix(file, ".js"))
+        assetTracers[name] = string(tracers.MustAsset(file))
+    }
+    tracers2.RegisterLookup(true, newJsTracer)
+}
+
+// makeSlice convert an unsafe memory pointer with the given type into a Go byte
+// slice.
+//
+// Note, the returned slice uses the same memory area as the input arguments.
+// If those are duktape stack items, popping them off **will** make the slice
+// contents change.
+func makeSlice(ptr unsafe.Pointer, size uint) []byte {
+    var sl = struct {
+        addr uintptr
+        len  int
+        cap  int
+    }(uintptr(ptr), int(size), int(size))
+
+    return *(*[]byte)(unsafe.Pointer(&sl))
+}
+
+// popSlice pops a buffer off the JavaScript stack and returns it as a slice.
+func popSlice(ctx *duktape.Context) []byte {
+    blob := common.CopyBytes(makeSlice(ctx.GetBuffer(-1)))
+    ctx.Pop()
+    return blob
+}
+
+// pushBigInt create a JavaScript BigInteger in the VM.
+func pushBigInt(n *big.Int, ctx *duktape.Context) {
+    ctx.GetGlobalString("bigint")
+    ctx.PushString(n.String())
+    ctx.Call(1)
+}
+
+// opWrapper provides a JavaScript wrapper around OpCode.
+type opWrapper struct {
+    op vm.OpCode
+}
+
+// pushObject assembles a JSVM object wrapping a swappable opcode and pushes it
+// onto the VM stack.
+func (ow *opWrapper) pushObject(vm *duktape.Context) {
+    obj := vm.PushObject()
+
+    vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushInt(int(ow.op)); return 1 })
+    vm.PutPropString(obj, "toNumber")
+
+    vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushString(ow.op.String()); return 1 })
+    vm.PutPropString(obj, "toString")
+
+    vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushBoolean(ow.op.IsPush()); return 1 })
+    vm.PutPropString(obj, "isPush")
+}
+
+// memoryWrapper provides a JavaScript wrapper around vm.Memory.
+type memoryWrapper struct {
+    memory *vm.Memory
+}
+
+// slice returns the requested range of memory as a byte slice.
+func (mw *memoryWrapper) slice(begin, end int64) []byte {
+    if end == begin {
+        return []byte{}
+    }
+    if end < begin || begin < 0 {
+        // TODO(karalabe): We can't js-throw from Go inside duktape inside Go. The Go
+        // runtime goes belly up https://github.com/golang/go/issues/15639.
+        log.Warn("Tracer accessed out of bound memory", "offset", begin, "end", end)
+        return nil
+    }
+    if mw.memory.Len() < int(end) {
+        // TODO(karalabe): We can't js-throw from Go inside duktape inside Go. The Go
+        // runtime goes belly up https://github.com/golang/go/issues/15639.
+        log.Warn("Tracer accessed out of bound memory", "available", mw.memory.Len(), "offset", begin, "size", end-begin)
+        return nil
+    }
+    return mw.memory.GetCopy(begin, end-begin)
+}
+
+// getUint returns the 32 bytes at the specified address interpreted as a uint.
+func (mw *memoryWrapper) getUint(addr int64) *big.Int {
+    if mw.memory.Len() < int(addr)+32 || addr < 0 {
+        // TODO(karalabe): We can't js-throw from Go inside duktape inside Go. The Go
+        // runtime goes belly up https://github.com/golang/go/issues/15639.
+        log.Warn("Tracer accessed out of bound memory", "available", mw.memory.Len(), "offset", addr, "size", 32)
+        return new(big.Int)
+    }
+    return new(big.Int).SetBytes(mw.memory.GetPtr(addr, 32))
+}
+
+// pushObject assembles a JSVM object wrapping a swappable memory and pushes it
+// onto the VM stack.
+func (mw *memoryWrapper) pushObject(vm *duktape.Context) {
+    obj := vm.PushObject()
+
+    // Generate the `slice` method which takes two ints and returns a buffer
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        blob := mw.slice(int64(ctx.GetInt(-2)), int64(ctx.GetInt(-1)))
+        ctx.Pop2()
+
+        ptr := ctx.PushFixedBuffer(len(blob))
+        copy(makeSlice(ptr, uint(len(blob))), blob)
+        return 1
+    })
+    vm.PutPropString(obj, "slice")
+
+    // Generate the `getUint` method which takes an int and returns a bigint
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        offset := int64(ctx.GetInt(-1))
+        ctx.Pop()
+
+        pushBigInt(mw.getUint(offset), ctx)
+        return 1
+    })
+    vm.PutPropString(obj, "getUint")
+}
+
+// stackWrapper provides a JavaScript wrapper around vm.Stack.
+type stackWrapper struct {
+    stack *vm.Stack
+}
+
+// peek returns the nth-from-the-top element of the stack.

```

```

+func (sw *stackWrapper) peek(idx int) *big.Int {
+    if len(sw.stack.Data()) <= idx || idx < 0 {
+        // TODO(karalabe): We can't js-throw from Go inside duktape inside Go. The Go
+        // runtime goes belly up https://github.com/golang/go/issues/15639.
+        log.Warn("Tracer accessed out of bound stack", "size", len(sw.stack.Data()), "index", idx)
+        return new(big.Int)
+    }
+    return sw.stack.Back(idx).ToBig()
+}
+
+// pushObject assembles a JSVM object wrapping a swappable stack and pushes it
+// onto the VM stack.
+func (sw *stackWrapper) pushObject(vm *duktape.Context) {
+    obj := vm.PushObject()
+
+    vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushInt(len(sw.stack.Data())); return 1 })
+    vm.PutPropString(obj, "length")
+
+    // Generate the `peek` method which takes an int and returns a bigint
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        offset := ctx.GetInt(-1)
+        ctx.Pop()
+
+        pushBigInt(sw.peek(offset), ctx)
+        return 1
+    })
+    vm.PutPropString(obj, "peek")
+}
+
+// dbWrapper provides a JavaScript wrapper around vm.Database.
+type dbWrapper struct {
+    db vm.StateDB
+}
+
+// pushObject assembles a JSVM object wrapping a swappable database and pushes it
+// onto the VM stack.
+func (dw *dbWrapper) pushObject(vm *duktape.Context) {
+    obj := vm.PushObject()
+
+    // Push the wrapper for statedb.GetBalance
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        pushBigInt(dw.db.GetBalance(common.BytesToAddress(popSlice(ctx))), ctx)
+        return 1
+    })
+    vm.PutPropString(obj, "getBalance")
+
+    // Push the wrapper for statedb.GetNonce
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        ctx.PushInt(int(dw.db.GetNonce(common.BytesToAddress(popSlice(ctx)))))
+        return 1
+    })
+    vm.PutPropString(obj, "getNonce")
+
+    // Push the wrapper for statedb.GetCode
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        code := dw.db.GetCode(common.BytesToAddress(popSlice(ctx)))
+
+        ptr := ctx.PushFixedBuffer(len(code))
+        copy(makeSlice(ptr, uint(len(code))), code)
+        return 1
+    })
+    vm.PutPropString(obj, "getCode")
+
+    // Push the wrapper for statedb.GetState
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        hash := popSlice(ctx)
+        addr := popSlice(ctx)
+
+        state := dw.db.GetState(common.BytesToAddress(addr), common.BytesToHash(hash))
+
+        ptr := ctx.PushFixedBuffer(len(state))
+        copy(makeSlice(ptr, uint(len(state))), state[:])
+        return 1
+    })
+    vm.PutPropString(obj, "getState")
+
+    // Push the wrapper for statedb.Exists
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        ctx.PushBoolean(dw.db.Exist(common.BytesToAddress(popSlice(ctx))))
+        return 1
+    })
+    vm.PutPropString(obj, "exists")
+}
+
+// contractWrapper provides a JavaScript wrapper around vm.Contract
+type contractWrapper struct {
+    contract *vm.Contract
+}
+
+// pushObject assembles a JSVM object wrapping a swappable contract and pushes it
+// onto the VM stack.
+func (cw *contractWrapper) pushObject(vm *duktape.Context) {
+    obj := vm.PushObject()
+
+    // Push the wrapper for contract.Caller
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        ptr := ctx.PushFixedBuffer(20)
+        copy(makeSlice(ptr, 20), cw.contract.Caller().Bytes())
+        return 1
+    })
+    vm.PutPropString(obj, "getCaller")
+
+    // Push the wrapper for contract.Address
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        ptr := ctx.PushFixedBuffer(20)
+        copy(makeSlice(ptr, 20), cw.contract.Address().Bytes())
+        return 1
+    })
+    vm.PutPropString(obj, "getAddress")
+
+    // Push the wrapper for contract.Value
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        pushBigInt(cw.contract.Value(), ctx)
+        return 1
+    })
+    vm.PutPropString(obj, "getValue")
+
+    // Push the wrapper for contract.Input
+    vm.PushGoFunction(func(ctx *duktape.Context) int {
+        blob := cw.contract.Input
+
+        ptr := ctx.PushFixedBuffer(len(blob))
+        copy(makeSlice(ptr, uint(len(blob))), blob)
+        return 1
+    })
+    vm.PutPropString(obj, "getInput")
+}
+
+type frame struct {
+    typ *string
+    from *common.Address
+    to *common.Address
+    input []byte
+    gas *uint

```

```

+         value *big.Int
+     }
+ }
+
+func newFrame() *frame {
+     return &frame{
+         typ:  new(string),
+         from: new(common.Address),
+         to:   new(common.Address),
+         gas:  new(uint),
+     }
+ }
+
+func (f *frame) pushObject(vm *duktape.Context) {
+     obj := vm.PushObject()
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, *f.typ); return 1 })
+     vm.PutPropString(obj, "getType")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, *f.from); return 1 })
+     vm.PutPropString(obj, "getFrom")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, *f.to); return 1 })
+     vm.PutPropString(obj, "getTo")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, f.input); return 1 })
+     vm.PutPropString(obj, "getInput")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, *f.gas); return 1 })
+     vm.PutPropString(obj, "getGas")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int {
+         if f.value != nil {
+             pushValue(ctx, f.value)
+         } else {
+             ctx.PushUndefined()
+         }
+         return 1
+     })
+     vm.PutPropString(obj, "getValue")
+ }
+
+type frameResult struct {
+     gasUsed    *uint
+     output     []byte
+     errorValue *string
+ }
+
+func newFrameResult() *frameResult {
+     return &frameResult{
+         gasUsed: new(uint),
+     }
+ }
+
+func (r *frameResult) pushObject(vm *duktape.Context) {
+     obj := vm.PushObject()
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, *r.gasUsed); return 1 })
+     vm.PutPropString(obj, "getGasUsed")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int { pushValue(ctx, r.output); return 1 })
+     vm.PutPropString(obj, "getOutput")
+
+     vm.PushGoFunction(func(ctx *duktape.Context) int {
+         if r.errorValue != nil {
+             pushValue(ctx, *r.errorValue)
+         } else {
+             ctx.PushUndefined()
+         }
+         return 1
+     })
+     vm.PutPropString(obj, "getError")
+ }
+
+// jsTracer provides an implementation of Tracer that evaluates a Javascript
+// function for each VM execution step.
+type jsTracer struct {
+     vm      *duktape.Context // Javascript VM instance
+     env     *vm.EVM           // EVM instance executing the code being traced
+
+     tracerObject int // Stack index of the tracer JavaScript object
+     stateObject  int // Stack index of the global state to pull arguments from
+
+     opWrapper    *opWrapper // Wrapper around the VM opcode
+     stackWrapper *stackWrapper // Wrapper around the VM stack
+     memoryWrapper *memoryWrapper // Wrapper around the VM memory
+     contractWrapper *contractWrapper // Wrapper around the contract object
+     dbWrapper     *dbWrapper // Wrapper around the VM environment
+
+     pcValue    *uint // Swappable pc value wrapped by a log accessor
+     gasValue   *uint // Swappable gas value wrapped by a log accessor
+     costValue  *uint // Swappable cost value wrapped by a log accessor
+     depthValue *uint // Swappable depth value wrapped by a log accessor
+     errorValue *string // Swappable error value wrapped by a log accessor
+     refundValue *uint // Swappable refund value wrapped by a log accessor
+
+     frame      *frame // Represents entry into call frame. Fields are swappable
+     frameResult *frameResult // Represents exit from a call frame. Fields are swappable
+
+     ctx map[string]interface{} // Transaction context gathered throughout execution
+     err error // Error, if one has occurred
+
+     interrupt uint32 // Atomic flag to signal execution interruption
+     reason     error // Textual reason for the interruption
+
+     activePrecompiles []common.Address // Updated on CaptureStart based on given rules
+     traceSteps        bool // When true, will invoke step() on each opcode
+     traceCallFrames   bool // When true, will invoke enter() and exit() js funcs
+ }
+
+// New instantiates a new tracer instance. code specifies a Javascript snippet,
+// which must evaluate to an expression returning an object with 'step', 'fault'
+// and 'result' functions.
+func newJsTracer(code string, ctx *tracers2.Context) (tracers2.Tracer, error) {
+     if c, ok := assetTracers[code]; ok {
+         code = c
+     }
+     if ctx == nil {
+         ctx = new(tracers2.Context)
+     }
+     tracer := &jsTracer{
+         vm:      duktape.New(),
+         ctx:     make(map[string]interface{}),
+         opWrapper: new(opWrapper),
+         stackWrapper: new(stackWrapper),
+         memoryWrapper: new(memoryWrapper),
+         contractWrapper: new(contractWrapper),
+         dbWrapper: new(dbWrapper),
+         pcValue: new(uint),
+         gasValue: new(uint),
+         costValue: new(uint),
+         depthValue: new(uint),
+         refundValue: new(uint),
+         frame: newFrame(),
+         frameResult: newFrameResult(),

```

```

+ }
+ if ctx.BlockHash != (common.Hash{}) {
+     tracer.ctx["blockHash"] = ctx.BlockHash
+
+     if ctx.TxHash != (common.Hash{}) {
+         tracer.ctx["txIndex"] = ctx.TxIndex
+         tracer.ctx["txHash"] = ctx.TxHash
+     }
+ }
+
+ // Set up builtins for this environment
+ tracer.vm.PushGlobalGoFunction("toHex", func(ctx *duktape.Context) int {
+     ctx.PushString(hexutil.Encode(popSlice(ctx)))
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("toWord", func(ctx *duktape.Context) int {
+     var word common.Hash
+     if ptr, size := ctx.GetBuffer(-1); ptr != nil {
+         word = common.BytesToHash(makeSlice(ptr, size))
+     } else {
+         word = common.HexToHash(ctx.GetString(-1))
+     }
+     ctx.Pop()
+     copy(makeSlice(ctx.PushFixedBuffer(32), 32), word[:])
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("toAddress", func(ctx *duktape.Context) int {
+     var addr common.Address
+     if ptr, size := ctx.GetBuffer(-1); ptr != nil {
+         addr = common.BytesToAddress(makeSlice(ptr, size))
+     } else {
+         addr = common.HexToAddress(ctx.GetString(-1))
+     }
+     ctx.Pop()
+     copy(makeSlice(ctx.PushFixedBuffer(20), 20), addr[:])
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("toContract", func(ctx *duktape.Context) int {
+     var from common.Address
+     if ptr, size := ctx.GetBuffer(-2); ptr != nil {
+         from = common.BytesToAddress(makeSlice(ptr, size))
+     } else {
+         from = common.HexToAddress(ctx.GetString(-2))
+     }
+     nonce := uint64(ctx.GetInt(-1))
+     ctx.Pop2()
+
+     contract := crypto.CreateAddress(from, nonce)
+     copy(makeSlice(ctx.PushFixedBuffer(20), 20), contract[:])
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("toContract2", func(ctx *duktape.Context) int {
+     var from common.Address
+     if ptr, size := ctx.GetBuffer(-3); ptr != nil {
+         from = common.BytesToAddress(makeSlice(ptr, size))
+     } else {
+         from = common.HexToAddress(ctx.GetString(-3))
+     }
+
+     // Retrieve salt hex string from js stack
+     salt := common.HexToHash(ctx.GetString(-2))
+     // Retrieve code slice from js stack
+     var code []byte
+     if ptr, size := ctx.GetBuffer(-1); ptr != nil {
+         code = common.CopyBytes(makeSlice(ptr, size))
+     } else {
+         code = common.FromHex(ctx.GetString(-1))
+     }
+
+     codeHash := crypto.Keccak256(code)
+     ctx.Pop3()
+     contract := crypto.CreateAddress2(from, salt, codeHash)
+     copy(makeSlice(ctx.PushFixedBuffer(20), 20), contract[:])
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("isPrecompiled", func(ctx *duktape.Context) int {
+     addr := common.BytesToAddress(popSlice(ctx))
+     for _, p := range tracer.activePrecompiles {
+         if p == addr {
+             ctx.PushBoolean(true)
+             return 1
+         }
+     }
+     ctx.PushBoolean(false)
+     return 1
+ })
+
+ tracer.vm.PushGlobalGoFunction("slice", func(ctx *duktape.Context) int {
+     start, end := ctx.GetInt(-2), ctx.GetInt(-1)
+     ctx.Pop2()
+
+     blob := popSlice(ctx)
+     size := end - start
+
+     if start < 0 || start > end || end > len(blob) {
+         // TODO(karalabe): We can't js-throw from Go inside duktape inside Go. The Go
+         // runtime goes belly up https://github.com/golang/go/issues/15639.
+         log.Warn("Tracer accessed out of bound memory", "available", len(blob), "offset", start, "size", size)
+         ctx.PushFixedBuffer(0)
+         return 1
+     }
+
+     copy(makeSlice(ctx.PushFixedBuffer(size), uint(size)), blob[start:end])
+     return 1
+ })
+
+ // Push the JavaScript tracer as object #0 onto the JSVM stack and validate it
+ if err := tracer.vm.PevalString("(" + code + ")"); err != nil {
+     log.Warn("Failed to compile tracer", "err", err)
+     return nil, err
+ }
+
+ tracer.tracerObject = 0 // yeah, nice, eval can't return the index itself
+
+ hasStep := tracer.vm.GetPropString(tracer.tracerObject, "step")
+ tracer.vm.Pop()
+
+ if !tracer.vm.GetPropString(tracer.tracerObject, "fault") {
+     return nil, fmt.Errorf("trace object must expose a function fault()")
+ }
+ tracer.vm.Pop()
+
+ if !tracer.vm.GetPropString(tracer.tracerObject, "result") {
+     return nil, fmt.Errorf("trace object must expose a function result()")
+ }
+ tracer.vm.Pop()
+
+ hasEnter := tracer.vm.GetPropString(tracer.tracerObject, "enter")
+ tracer.vm.Pop()
+ hasExit := tracer.vm.GetPropString(tracer.tracerObject, "exit")
+ tracer.vm.Pop()
+ if hasEnter != hasExit {
+     return nil, fmt.Errorf("trace object must expose either both or none of enter() and exit()")
+ }
+
+ tracer.traceCallFrames = hasEnter && hasExit
+ tracer.traceSteps = hasStep
+
+ // Tracer is valid, inject the big int library to access large numbers
+ tracer.vm.EvalString(bigIntegerJS)
+ tracer.vm.PutGlobalString("bigInt")

```

```

+ // Push the global environment state as object #1 into the JSVM stack
+ tracer.stateObject = tracer.vm.PushObject()
+
+ logObject := tracer.vm.PushObject()
+
+ tracer.opWrapper.pushObject(tracer.vm)
+ tracer.vm.PutPropString(logObject, "op")
+
+ tracer.stackWrapper.pushObject(tracer.vm)
+ tracer.vm.PutPropString(logObject, "stack")
+
+ tracer.memoryWrapper.pushObject(tracer.vm)
+ tracer.vm.PutPropString(logObject, "memory")
+
+ tracer.contractWrapper.pushObject(tracer.vm)
+ tracer.vm.PutPropString(logObject, "contract")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushUint(*tracer.pcValue); return 1 })
+ tracer.vm.PutPropString(logObject, "getPC")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushUint(*tracer.gasValue); return 1 })
+ tracer.vm.PutPropString(logObject, "getGas")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushUint(*tracer.costValue); return 1 })
+ tracer.vm.PutPropString(logObject, "getCost")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushUint(*tracer.depthValue); return 1 })
+ tracer.vm.PutPropString(logObject, "getDepth")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int { ctx.PushUint(*tracer.refundValue); return 1 })
+ tracer.vm.PutPropString(logObject, "getRefund")
+
+ tracer.vm.PushGoFunction(func(ctx *duktape.Context) int {
+     if tracer.errorValue != nil {
+         ctx.PushString(*tracer.errorValue)
+     } else {
+         ctx.PushUndefined()
+     }
+     return 1
+ })
+ tracer.vm.PutPropString(logObject, "getError")
+
+ tracer.vm.PutPropString(tracer.stateObject, "log")
+
+ tracer.frame.pushObject(tracer.vm)
+ tracer.vm.PutPropString(tracer.stateObject, "frame")
+
+ tracer.frameResult.pushObject(tracer.vm)
+ tracer.vm.PutPropString(tracer.stateObject, "frameResult")
+
+ tracer.dbWrapper.pushObject(tracer.vm)
+ tracer.vm.PutPropString(tracer.stateObject, "db")
+
+ return tracer, nil
+}
+
+// Stop terminates execution of the tracer at the first opportune moment.
+func (jst *jsTracer) Stop(err error) {
+    jst.reason = err
+    atomic.StoreUint32(&jst.interrupt, 1)
+}
+
+// call executes a method on a JS object, catching any errors, formatting and
+// returning them as error objects.
+func (jst *jsTracer) call(noret bool, method string, args ...string) (json.RawMessage, error) {
+    // Execute the JavaScript call and return any error
+    jst.vm.PushString(method)
+    for _, arg := range args {
+        jst.vm.GetPropString(jst.stateObject, arg)
+    }
+    code := jst.vm.PcallProp(jst.tracerObject, len(args))
+    defer jst.vm.Pop()
+
+    if code != 0 {
+        err := jst.vm.SafeToString(-1)
+        return nil, errors.New(err)
+    }
+    // No error occurred, extract return value and return
+    if noret {
+        return nil, nil
+    }
+
+    // Push a JSON marshaller onto the stack. We can't marshal from the out-
+    // side because duktape can crash on large nestings and we can't catch
+    // C++ exceptions ourselves from Go. TODO(karalabe): Yuck, why wrap?!
+    jst.vm.PushString("(JSON.stringify)")
+    jst.vm.Eval()
+
+    jst.vm.Swap(-1, -2)
+    if code = jst.vm.Pcall(1); code != 0 {
+        err := jst.vm.SafeToString(-1)
+        return nil, errors.New(err)
+    }
+    return json.RawMessage(jst.vm.SafeToString(-1)), nil
+}
+
+func wrapError(context string, err error) error {
+    return fmt.Errorf("%w in server-side tracer function '%v'", err, context)
+}
+
+// CaptureStart implements the Tracer interface to initialize the tracing operation.
+func (jst *jsTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+    jst.env = env
+    jst.ctx["type"] = "CALL"
+    if create {
+        jst.ctx["type"] = "CREATE"
+    }
+    jst.ctx["from"] = from
+    jst.ctx["to"] = to
+    jst.ctx["input"] = input
+    jst.ctx["gas"] = gas
+    jst.ctx["gasPrice"] = env.TxContext.GasPrice
+    jst.ctx["value"] = value
+
+    // Initialize the context
+    jst.ctx["block"] = env.Context.BlockNumber.Uint64()
+    jst.dbWrapper.db = env.StateDB
+    // Update list of precompiles based on current block
+    rules := env.ChainConfig().AvalancheRules(env.Context.BlockNumber, env.Context.Time)
+    jst.activePrecompiles = vm.ActivePrecompiles(rules)
+
+    // Compute intrinsic gas
+    isHomestead := env.ChainConfig().IsHomestead(env.Context.BlockNumber)
+    isIstanbul := env.ChainConfig().IsIstanbul(env.Context.BlockNumber)
+    intrinsicGas, err := core.IntrinsicGas(input, nil, jst.ctx["type"] == "CREATE", isHomestead, isIstanbul)
+    if err != nil {
+        return
+    }
+    jst.ctx["intrinsicGas"] = intrinsicGas
+}
+
+// CaptureState implements the Tracer interface to trace a single step of VM execution.
+func (jst *jsTracer) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+    if !jst.traceSteps {
+        return
+    }

```

```

+     }
+     if jst.err != nil {
+         return
+     }
+     // If tracing was interrupted, set the error and stop
+     if atomic.LoadUint32(&jst.interrupt) > 0 {
+         jst.err = jst.reason
+         jst.env.Cancel()
+         return
+     }
+     jst.opWrapper.op = op
+     jst.stackWrapper.stack = scope.Stack
+     jst.memoryWrapper.memory = scope.Memory
+     jst.contractWrapper.contract = scope.Contract
+
+     *jst.pcValue = uint(pc)
+     *jst.gasValue = uint(gas)
+     *jst.costValue = uint(cost)
+     *jst.depthValue = uint(depth)
+     *jst.refundValue = uint(jst.env.StateDB.GetRefund())
+
+     jst.errorValue = nil
+     if err != nil {
+         jst.errorValue = new(string)
+         *jst.errorValue = err.Error()
+     }
+
+     if _, err := jst.call(true, "step", "log", "db"); err != nil {
+         jst.err = wrapError("step", err)
+     }
+ }
+}
+
+// CaptureFault implements the Tracer interface to trace an execution fault
+func (jst *jsTracer) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+     if jst.err != nil {
+         return
+     }
+     // Apart from the error, everything matches the previous invocation
+     jst.errorValue = new(string)
+     *jst.errorValue = err.Error()
+
+     if _, err := jst.call(true, "fault", "log", "db"); err != nil {
+         jst.err = wrapError("fault", err)
+     }
+ }
+}
+
+// CaptureEnd is called after the call finishes to finalize the tracing.
+func (jst *jsTracer) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {
+     jst.ctx["output"] = output
+     jst.ctx["time"] = t.String()
+     jst.ctx["gasUsed"] = gasUsed
+
+     if err != nil {
+         jst.ctx["error"] = err.Error()
+     }
+ }
+}
+
+// CaptureEnter is called when EVM enters a new scope (via call, create or selfdestruct).
+func (jst *jsTracer) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+     if !jst.traceCallFrames {
+         return
+     }
+     if jst.err != nil {
+         return
+     }
+     // If tracing was interrupted, set the error and stop
+     if atomic.LoadUint32(&jst.interrupt) > 0 {
+         jst.err = jst.reason
+         return
+     }
+
+     *jst.frame.typ = typ.String()
+     *jst.frame.from = from
+     *jst.frame.to = to
+     jst.frame.input = common.CopyBytes(input)
+     *jst.frame.gas = uint(gas)
+     jst.frame.value = nil
+     if value != nil {
+         jst.frame.value = new(big.Int).SetBytes(value.Bytes())
+     }
+
+     if _, err := jst.call(true, "enter", "frame"); err != nil {
+         jst.err = wrapError("enter", err)
+     }
+ }
+}
+
+// CaptureExit is called when EVM exits a scope, even if the scope didn't
+// execute any code.
+func (jst *jsTracer) CaptureExit(output []byte, gasUsed uint64, err error) {
+     if !jst.traceCallFrames {
+         return
+     }
+     // If tracing was interrupted, set the error and stop
+     if atomic.LoadUint32(&jst.interrupt) > 0 {
+         jst.err = jst.reason
+         return
+     }
+
+     jst.frameResult.output = common.CopyBytes(output)
+     *jst.frameResult.gasUsed = uint(gasUsed)
+     jst.frameResult.errorValue = nil
+     if err != nil {
+         jst.frameResult.errorValue = new(string)
+         *jst.frameResult.errorValue = err.Error()
+     }
+
+     if _, err := jst.call(true, "exit", "frameResult"); err != nil {
+         jst.err = wrapError("exit", err)
+     }
+ }
+}
+
+// GetResult calls the Javascript 'result' function and returns its value, or any accumulated error
+func (jst *jsTracer) GetResult() (json.RawMessage, error) {
+     // Transform the context into a JavaScript object and inject into the state
+     obj := jst.vm.PushObject()
+
+     for key, val := range jst.ctx {
+         jst.addToObj(obj, key, val)
+     }
+     jst.vm.PutPropString(jst.stateObject, "ctx")
+
+     // Finalize the trace and return the results
+     result, err := jst.call(false, "result", "ctx", "db")
+     if err != nil {
+         jst.err = wrapError("result", err)
+     }
+
+     // Clean up the JavaScript environment
+     jst.vm.DestroyHeap()
+     jst.vm.Destroy()
+
+     return result, jst.err
+}
+}
+
+// addToObj pushes a field to a JS object.

```

```

+func (jst *jsTracer) addToObj(obj int, key string, val interface{}) {
+    pushValue(jst.vm, val)
+    jst.vm.PutPropString(obj, key)
+}
+
+func pushValue(ctx *duktape.Context, val interface{}) {
+    switch val := val.(type) {
+    case uint64:
+        ctx.PushUint(uint(val))
+    case string:
+        ctx.PushString(val)
+    case []byte:
+        ptr := ctx.PushFixedBuffer(len(val))
+        copy(makeSlice(ptr, uint(len(val))), val)
+    case common.Address:
+        ptr := ctx.PushFixedBuffer(20)
+        copy(makeSlice(ptr, 20), val[:])
+    case *big.Int:
+        pushBigInt(val, ctx)
+    case int:
+        ctx.PushInt(val)
+    case uint:
+        ctx.PushUint(val)
+    case common.Hash:
+        ptr := ctx.PushFixedBuffer(32)
+        copy(makeSlice(ptr, 32), val[:])
+    default:
+        panic(fmt.Sprintf("unsupported type: %T", val))
+    }
+}
+
diff --git a/eth/tracers/js/tracer_test.go b/eth/tracers/js/tracer_test.go
new file mode 100644
index 00000000..a43eb9b3
--- /dev/null
+++ b/eth/tracers/js/tracer_test.go
@@ -0,0 +1,267 @@
// (c) 2020-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
//
// *****
// Copyright 2017 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package js
+
+import (
+    "encoding/json"
+    "errors"
+    "math/big"
+    "testing"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/tracers"
+    "github.com/flare-foundation/coreth/params"
+)
+
+type account struct{}
+
+func (account) SubBalance(amount *big.Int) {}
+func (account) AddBalance(amount *big.Int) {}
+func (account) SetAddress(common.Address) {}
+func (account) Value() *big.Int { return nil }
+func (account) SetBalance(*big.Int) {}
+func (account) SetNonce(uint64) {}
+func (account) Balance() *big.Int { return nil }
+func (account) Address() common.Address { return common.Address{} }
+func (account) SetCode(common.Hash, []byte) {}
+func (account) ForEachStorage(cb func(key, value common.Hash) bool) {}
+
+type dummyStatedb struct {
+    state.StateDB
+}
+
+func (*dummyStatedb) GetRefund() uint64 { return 1337 }
+func (*dummyStatedb) GetBalance(addr common.Address) *big.Int { return new(big.Int) }
+
+type vmContext struct {
+    blockCtx vm.BlockContext
+    txCtx    vm.TxContext
+}
+
+func testCtx() *vmContext {
+    return &vmContext{blockCtx: vm.BlockContext{BlockNumber: big.NewInt(1)}, txCtx: vm.TxContext{GasPrice: big.NewInt(100000)}}
+}
+
+func runTrace(tracer tracers.Tracer, vmctx *vmContext, chaincfg *params.ChainConfig) (json.RawMessage, error) {
+    var (
+        env      = vm.NewEVM(vmctx.blockCtx, vmctx.txCtx, &dummyStatedb{}, chaincfg, vm.Config{Debug: true, Tracer: tracer})
+        startGas uint64 = 10000
+        value     = big.NewInt(0)
+        contract  = vm.NewContract(account{}, account{}, value, startGas)
+    )
+    contract.Code = []byte{byte(vm.PUSH1), 0x1, byte(vm.PUSH1), 0x1, 0x0}
+
+    tracer.CaptureStart(env, contract.Caller(), contract.Address(), false, []byte{}, startGas, value)
+    ret, err := env.Interpreter().Run(contract, []byte{}, false)
+    tracer.CaptureEnd(ret, startGas-contract.Gas, 1, err)
+    if err != nil {
+        return nil, err
+    }
+    return tracer.GetResult()
+}
+
+func TestTracer(t *testing.T) {
+    execTracer := func(code string) ([]byte, string) {
+        t.Helper()
+        tracer, err := newJsTracer(code, nil)
+        if err != nil {
+            t.Fatal(err)
+        }
+        ret, err := runTrace(tracer, testCtx(), params.TestChainConfig)
+        if err != nil {

```



```

return nil, err.Error() // Stringify to allow comparison without nil checks
}
    return ret, ""
}
for i, tt := range []struct {
    code string
    want string
    fail string
}{
    { // tests that we don't panic on bad arguments to memory access
        code: "{depths: [], step: function(log) { this.depths.push(log.memory.slice(-1,-2)); }, fault: function() {}, result: function() { return this.depths; }}",
        want: "[[],[],[]]",
    }, { // tests that we don't panic on bad arguments to stack peeks
        code: "{depths: [], step: function(log) { this.depths.push(log.stack.peek(-1)); }, fault: function() {}, result: function() { return this.depths; }}",
        want: "[\"0\",\"0\",\"0\"]",
    }, { // tests that we don't panic on bad arguments to memory getUint
        code: "{ depths: [], step: function(log, db) { this.depths.push(log.memory.getUint(-64));}, fault: function() {}, result: function() { return this.depths; }}",
        want: "[\"0\",\"0\",\"0\"]",
    }, { // tests some general counting
        code: "{count: 0, step: function() { this.count += 1; }, fault: function() {}, result: function() { return this.count; }}",
        want: "3",
    }, { // tests that depth is reported correctly
        code: "{depths: [], step: function(log) { this.depths.push(log.stack.length()); }, fault: function() {}, result: function() { return this.depths; }}",
        want: "[0,1,2]",
    }, { // tests to-string of opcodes
        code: "{opcodes: [], step: function(log) { this.opcodes.push(log.op.toString()); }, fault: function() {}, result: function() { return this.opcodes; }}",
        want: "[\"PUSH1\", \"PUSH1\", \"STOP\"]",
    }, { // tests intrinsic gas
        code: "{depths: [], step: function() {}, fault: function() {}, result: function(ctx) { return ctx.gasPrice+'.'+ctx.gasUsed+'.'+ctx.intrinsicGas; }}",
        want: "100000.6.21000",
    }, { // tests too deep object / serialization crash
        code: "(step: function() {}, fault: function() {}, result: function() { var o={}; var x=o; for (var i=0; i<1000; i++){ o.foo={}; o=o.foo; } return x;})",
        fail: "RangeError: json encode recursion limit in server-side tracer function 'result'",
    },
} {
    if have, err := execTracer(tt.code); tt.want != string(have) || tt.fail != err {
        t.Errorf("testcase %d: expected return value to be '%s' got '%s', error to be '%s' got '%s'\n\tcode: %v", i, tt.want, string(have), tt.fail, err, tt.code)
    }
}
}
}

+func TestHalt(t *testing.T) {
+    t.Skip("duktape doesn't support abortion")
+    timeout := errors.New("stahp")
+    tracer, err := newJsTracer("{step: function() { while(1); }, result: function() { return null; }, fault: function({})}", nil)
+    if err != nil {
+        t.Fatal(err)
+    }
+    go func() {
+        time.Sleep(1 * time.Second)
+        tracer.Stop(timeout)
+    }()
+    if _, err = runTrace(tracer, testCtx(), params.TestChainConfig); err.Error() != "stahp in server-side tracer function 'step'" {
+        t.Errorf("Expected timeout error, got %v", err)
+    }
+}

+func TestHaltBetweenSteps(t *testing.T) {
+    tracer, err := newJsTracer("{step: function() {}, fault: function() {}, result: function() { return null; }}", nil)
+    if err != nil {
+        t.Fatal(err)
+    }
+    env := vm.NewEVM(vm.BlockContext{BlockNumber: big.NewInt(1)}, vm.TxContext{GasPrice: big.NewInt(1)}, &dummyStatedb{}, params.TestChainConfig, vm.Config{Debug: true, Tracer: tracer})
+    scope := &vm.ScopeContext{
+        Contract: vm.NewContract(&account{}, &account{}, big.NewInt(0), 0),
+    }
+    tracer.CaptureStart(env, common.Address{}, common.Address{}, false, [byte]{}, 0, big.NewInt(0))
+    tracer.CaptureState(0, 0, 0, 0, scope, nil, 0, nil)
+    timeout := errors.New("stahp")
+    tracer.Stop(timeout)
+    tracer.CaptureState(0, 0, 0, 0, scope, nil, 0, nil)
+
+    if _, err := tracer.GetResult(); err.Error() != timeout.Error() {
+        t.Errorf("Expected timeout error, got %v", err)
+    }
+}

+// TestNoStepExec tests a regular value transfer (no exec), and accessing the statedb
+// in 'result'
+func TestNoStepExec(t *testing.T) {
+    execTracer := func(code string) [byte] {
+        t.Helper()
+        tracer, err := newJsTracer(code, nil)
+        if err != nil {
+            t.Fatal(err)
+        }
+        env := vm.NewEVM(vm.BlockContext{BlockNumber: big.NewInt(1)}, vm.TxContext{GasPrice: big.NewInt(100)}, &dummyStatedb{}, params.TestChainConfig, vm.Config{Debug: true, Tracer: tracer})
+        tracer.CaptureStart(env, common.Address{}, common.Address{}, false, [byte]{}, 1000, big.NewInt(0))
+        tracer.CaptureEnd(nil, 0, 1, nil)
+        ret, err := tracer.GetResult()
+        if err != nil {
+            t.Fatal(err)
+        }
+        return ret
+    }
+    for i, tt := range []struct {
+        code string
+        want string
+    ){
+        { // tests that we don't panic on accessing the db methods
+            code: "{depths: [], step: function() {}, fault: function() {}, result: function(ctx, db){ return db.getBalance(ctx.to); }}",
+            want: "0",
+        },
+    } {
+        if have := execTracer(tt.code); tt.want != string(have) {
+            t.Errorf("testcase %d: expected return value to be %s got %s\n\tcode: %v", i, tt.want, string(have), tt.code)
+        }
+    }
+}

+func TestIsPrecompile(t *testing.T) {
+    chaincfg := &params.ChainConfig{ChainID: big.NewInt(1), HomesteadBlock: big.NewInt(0), DAOForkBlock: nil, DAOForkSupport: false, EIP150Block: big.NewInt(0), EIP150Hash: common.Hash{}, EIP155Block:
+    chaincfg.ByzantiumBlock = big.NewInt(100)
+    chaincfg.IstanbulBlock = big.NewInt(200)
+    chaincfg.ApricotPhase2BlockTimestamp = big.NewInt(300)
+    txCtx := vm.TxContext{GasPrice: big.NewInt(100000)}
+    tracer, err := newJsTracer("{addr: toAddress('0000000000000000000000000000000000000000'), res: null, step: function() { this.res = isPrecompiled(this.addr); }, fault: function() {}, result: functio
+    if err != nil {
+        t.Fatal(err)
+    }
+
+    blockCtx := vm.BlockContext{BlockNumber: big.NewInt(150)}
+    res, err := runTrace(tracer, &vmContext{blockCtx, txCtxt}, chaincfg)
+    if err != nil {
+        t.Error(err)
+    }
+    if string(res) != "false" {
+        t.Errorf("Tracer should not consider blake2f as precompile in byzantium")
+    }
+
+    tracer, _ = newJsTracer("{addr: toAddress('0000000000000000000000000000000000000000'), res: null, step: function() { this.res = isPrecompiled(this.addr); }, fault: function() {}, result: function
+    blockCtx := vm.BlockContext{BlockNumber: big.NewInt(250)}
+    res, err := runTrace(tracer, &vmContext{blockCtx, txCtxt}, chaincfg)
+    if err != nil {

```

```

+         t.Error(err)
+     }
+     if string(res) != "true" {
+         t.Errorf("Tracer should consider blake2f as precompile in istanbul")
+     }
+ }
+}
+
+func TestEnterExit(t *testing.T) {
+    // test that either both or none of enter() and exit() are defined
+    if _, err := newJsTracer("{step: function() {}, fault: function() {}, result: function() { return null; }, enter: function() {}}", new(tracers.Context)); err == nil {
+        t.Fatal("tracer creation should've failed without exit() definition")
+    }
+    if _, err := newJsTracer("{step: function() {}, fault: function() {}, result: function() { return null; }, enter: function() {}, exit: function() {}}", new(tracers.Context)); err != nil {
+        t.Fatal(err)
+    }
+    // test that the enter and exit method are correctly invoked and the values passed
+    tracer, err := newJsTracer("{enters: 0, exits: 0, enterGas: 0, gasUsed: 0, step: function() {}, fault: function() {}, result: function() { return {enters: this.enters, exits: this.exits, enterGas}
+    if err != nil {
+        t.Fatal(err)
+    }
+    scope := &vm.ScopeContext{
+        Contract: vm.NewContract(&account{}, &account{}, big.NewInt(0), 0),
+    }
+    tracer.CaptureEnter(vm.CALL, scope.Contract.Caller(), scope.Contract.Address(), []byte{}, 1000, new(big.Int))
+    tracer.CaptureExit([]byte{}, 400, nil)
+
+    have, err := tracer.GetResult()
+    if err != nil {
+        t.Fatal(err)
+    }
+    want := `{"enters":1,"exits":1,"enterGas":1000,"gasUsed":400}`
+    if string(have) != want {
+        t.Errorf("Number of invocations of enter() and exit() is wrong. Have %s, want %s\n", have, want)
+    }
+}
+}
diff --git a/eth/tracers/logger/access_list_tracer.go b/eth/tracers/logger/access_list_tracer.go
new file mode 100644
index 00000000..72ea7f98
--- /dev/null
+++ b/eth/tracers/logger/access_list_tracer.go
@@ -0,0 +1,185 @@
+// Copyright 2021 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package logger
+
+import (
+    "math/big"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+)
+
+// accessList is an accumulator for the set of accounts and storage slots an EVM
+// contract execution touches.
+type accessList map[common.Address]accessListSlots
+
+// accessListSlots is an accumulator for the set of storage slots within a single
+// contract that an EVM contract execution touches.
+type accessListSlots map[common.Hash]struct{}
+
+// newAccessList creates a new accessList.
+func newAccessList() accessList {
+    return make(map[common.Address]accessListSlots)
+}
+
+// addAddress adds an address to the accessList.
+func (al accessList) addAddress(address common.Address) {
+    // Set address if not previously present
+    if _, present := al[address]; !present {
+        al[address] = make(map[common.Hash]struct{})
+    }
+}
+
+// addSlot adds a storage slot to the accessList.
+func (al accessList) addSlot(address common.Address, slot common.Hash) {
+    // Set address if not previously present
+    al.addAddress(address)
+
+    // Set the slot on the surely existent storage set
+    al[address][slot] = struct{}{}
+}
+
+// equal checks if the content of the current access list is the same as the
+// content of the other one.
+func (al accessList) equal(other accessList) bool {
+    // Cross reference the accounts first
+    if len(al) != len(other) {
+        return false
+    }
+    for addr := range al {
+        if _, ok := other[addr]; !ok {
+            return false
+        }
+    }
+    for addr := range other {
+        if _, ok := al[addr]; !ok {
+            return false
+        }
+    }
+    // Accounts match, cross reference the storage slots too
+    for addr, slots := range al {
+        otherslots := other[addr]
+
+        if len(slots) != len(otherslots) {
+            return false
+        }
+        for hash := range slots {
+            if _, ok := otherslots[hash]; !ok {
+                return false
+            }
+        }
+        for hash := range otherslots {
+            if _, ok := slots[hash]; !ok {
+                return false
+            }
+        }
+    }
+    return true
+}

```

```

+}
+
+// accesslist converts the accesslist to a types.AccessList.
+func (al accessList) accessList() types.AccessList {
+    acl := make(types.AccessList, 0, len(al))
+    for addr, slots := range al {
+        tuple := types.AccessTuple{Address: addr, StorageKeys: []common.Hash{}}
+        for slot := range slots {
+            tuple.StorageKeys = append(tuple.StorageKeys, slot)
+        }
+        acl = append(acl, tuple)
+    }
+    return acl
+}
+
+// AccessListTracer is a tracer that accumulates touched accounts and storage
+// slots into an internal set.
+type AccessListTracer struct {
+    excl map[common.Address]struct{} // Set of account to exclude from the list
+    list accessList                  // Set of accounts and storage slots touched
+}
+
+// NewAccessListTracer creates a new tracer that can generate AccessLists.
+// An optional AccessList can be specified to occupy slots and addresses in
+// the resulting accesslist.
+func NewAccessListTracer(acl types.AccessList, from, to common.Address, precompiles []common.Address) *AccessListTracer {
+    excl := map[common.Address]struct{}{
+        from: {}, to: {},
+    }
+    for _, addr := range precompiles {
+        excl[addr] = struct{}{}
+    }
+    list := newAccessList()
+    for _, al := range acl {
+        if _, ok := excl[al.Address]; !ok {
+            list.AddAddress(al.Address)
+        }
+        for _, slot := range al.StorageKeys {
+            list.AddSlot(al.Address, slot)
+        }
+    }
+    return &AccessListTracer{
+        excl: excl,
+        list: list,
+    }
+}
+
+func (a *AccessListTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+}
+
+// CaptureState captures all opcodes that touch storage or addresses and adds them to the accesslist.
+func (a *AccessListTracer) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+    stack := scope.Stack
+    stackData := stack.Data()
+    stackLen := len(stackData)
+    if (op == vm.SLOAD || op == vm.SSTORE) && stackLen >= 1 {
+        slot := common.Hash(stackData[stackLen-1].Bytes32())
+        a.list.AddSlot(scope.Contract.Address(), slot)
+    }
+    if (op == vm.EXTCODECOPY || op == vm.EXTCODEHASH || op == vm.EXTCODESIZE || op == vm.BALANCE || op == vm.SELFDESTRUCT) && stackLen >= 1 {
+        addr := common.Address(stackData[stackLen-1].Bytes20())
+        if _, ok := a.excl[addr]; !ok {
+            a.list.AddAddress(addr)
+        }
+    }
+    if (op == vm.DELEGATECALL || op == vm.CALL || op == vm.STATICCALL || op == vm.CALLCODE) && stackLen >= 5 {
+        addr := common.Address(stackData[stackLen-2].Bytes20())
+        if _, ok := a.excl[addr]; !ok {
+            a.list.AddAddress(addr)
+        }
+    }
+}
+
+func (*AccessListTracer) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+}
+
+func (*AccessListTracer) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {}
+
+func (*AccessListTracer) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+}
+
+func (*AccessListTracer) CaptureExit(output []byte, gasUsed uint64, err error) {}
+
+// AccessList returns the current accesslist maintained by the tracer.
+func (a *AccessListTracer) AccessList() types.AccessList {
+    return a.list.accessList()
+}
+
+// Equal returns if the content of two access list traces are equal.
+func (a *AccessListTracer) Equal(other *AccessListTracer) bool {
+    return a.list.equal(other.list)
+}
+
+diff --git a/eth/tracers/logger/gen_structlog.go b/eth/tracers/logger/gen_structlog.go
new file mode 100644
index 00000000..20569f62
--- /dev/null
+++ b/eth/tracers/logger/gen_structlog.go
@@ -0,0 +1,118 @@
+// Code generated by github.com/fjl/gencodec. DO NOT EDIT.
+
+package logger
+
+import (
+    "encoding/json"
+
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/holiman/uint256"
+)
+
+var _ = (*structLogMarshaling)(nil)
+
+// MarshalJSON marshals as JSON.
+func (s StructLog) MarshalJSON() ([]byte, error) {
+    type StructLog struct {
+        Pc          uint64          `json:"pc"`
+        Op          vm.OpCode       `json:"op"`
+        Gas         math.HexOrDecimal64 `json:"gas"`
+        GasCost     math.HexOrDecimal64 `json:"gasCost"`
+        Memory      hexutil.Bytes    `json:"memory"`
+        MemorySize  int              `json:"memSize"`
+        Stack       []uint256.Int    `json:"stack"`
+        ReturnData  hexutil.Bytes    `json:"returnData"`
+        Storage     map[common.Hash]common.Hash `json:"-"`
+        Depth       int              `json:"depth"`
+        RefundCounter uint64           `json:"refund"`
+        Err         error            `json:"-"`
+        OpName      string           `json:"opName"`
+        ErrorString string           `json:"error"`
+    }
+    var enc StructLog
+    enc.Pc = s.Pc

```

```

+     enc.Op = s.Op
+     enc.Gas = math.HexOrDecimal64(s.Gas)
+     enc.GasCost = math.HexOrDecimal64(s.GasCost)
+     enc.Memory = s.Memory
+     enc.MemorySize = s.MemorySize
+     enc.Stack = s.Stack
+     enc.ReturnData = s.ReturnData
+     enc.Storage = s.Storage
+     enc.Depth = s.Depth
+     enc.RefundCounter = s.RefundCounter
+     enc.Err = s.Err
+     enc.OpName = s.OpName()
+     enc.ErrorString = s.ErrorString()
+     return json.Marshal(&enc)
+}
+
+// UnmarshalJSON unmarshals from JSON.
+func (s *StructLog) UnmarshalJSON(input []byte) error {
+    type StructLog struct {
+        Pc          *uint64          `json:"pc"`
+        Op          *vm.OpCode       `json:"op"`
+        Gas         *math.HexOrDecimal64
+        GasCost     *math.HexOrDecimal64 `json:"gasCost"`
+        Memory      *hexutil.Bytes    `json:"memory"`
+        MemorySize  *int              `json:"memSize"`
+        Stack       []uint256.Int     `json:"stack"`
+        ReturnData  *hexutil.Bytes    `json:"returnData"`
+        Storage     map[common.Hash]common.Hash `json:"-:"`
+        Depth       *int              `json:"depth"`
+        RefundCounter *uint64          `json:"refund"`
+        Err         error             `json:"-:"`
+    }
+    var dec StructLog
+    if err := json.Unmarshal(input, &dec); err != nil {
+        return err
+    }
+    if dec.Pc != nil {
+        s.Pc = *dec.Pc
+    }
+    if dec.Op != nil {
+        s.Op = *dec.Op
+    }
+    if dec.Gas != nil {
+        s.Gas = uint64(*dec.Gas)
+    }
+    if dec.GasCost != nil {
+        s.GasCost = uint64(*dec.GasCost)
+    }
+    if dec.Memory != nil {
+        s.Memory = *dec.Memory
+    }
+    if dec.MemorySize != nil {
+        s.MemorySize = *dec.MemorySize
+    }
+    if dec.Stack != nil {
+        s.Stack = dec.Stack
+    }
+    if dec.ReturnData != nil {
+        s.ReturnData = *dec.ReturnData
+    }
+    if dec.Storage != nil {
+        s.Storage = dec.Storage
+    }
+    if dec.Depth != nil {
+        s.Depth = *dec.Depth
+    }
+    if dec.RefundCounter != nil {
+        s.RefundCounter = *dec.RefundCounter
+    }
+    if dec.Err != nil {
+        s.Err = dec.Err
+    }
+    return nil
+}
+
+diff --git a/eth/tracers/logger/logger.go b/eth/tracers/logger/logger.go
+new file mode 100644
+index 00000000..5b5f73ec
+--- /dev/null
++++ b/eth/tracers/logger/logger.go
+@@ -0,0 +1,349 @@
+// Copyright 2015 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
++package logger
+
++import (
++    "encoding/hex"
++    "fmt"
++    "io"
++    "math/big"
++    "strings"
++    "time"
+
++    "github.com/ethereum/go-ethereum/common"
++    "github.com/ethereum/go-ethereum/common/hexutil"
++    "github.com/ethereum/go-ethereum/common/math"
++    "github.com/flare-foundation/coreth/core/types"
++    "github.com/flare-foundation/coreth/core/vm"
++    "github.com/flare-foundation/coreth/params"
++    "github.com/holiman/uint256"
++)
++
++// Storage represents a contract's storage.
++type Storage map[common.Hash]common.Hash
++
++// Copy duplicates the current storage.
++func (s Storage) Copy() Storage {
++    cpy := make(Storage)
++    for key, value := range s {
++        cpy[key] = value
++    }
++    return cpy
++}
++
++// Config are the configuration options for structured logger the EVM
++type Config struct {
++    EnableMemory    bool // enable memory capture
++    DisableStack    bool // disable stack capture
++    DisableStorage  bool // disable storage capture
++    EnableReturnData bool // enable return data capture
++    Debug           bool // print output during capture end

```

```

+     Limit      int // maximum length of output, but zero means unlimited
+     // Chain overrides, can be used to execute a trace using future fork rules
+     Overrides *params.ChainConfig `json:"overrides,omitempty"`
+}
+
+//go:generate gencodec -type StructLog -field-override structLogMarshaling -out gen_structlog.go
+
+// StructLog is emitted to the EVM each cycle and lists information about the current internal state
+// prior to the execution of the statement.
+type StructLog struct {
+    Pc          uint64          `json:"pc"`
+    Op          vm.OpCode       `json:"op"`
+    Gas         uint64          `json:"gas"`
+    GasCost     uint64          `json:"gasCost"`
+    Memory      []byte          `json:"memory"`
+    MemorySize  int             `json:"memSize"`
+    Stack       []uint256.Int   `json:"stack"`
+    ReturnData  []byte          `json:"returnData"`
+    Storage     map[common.Hash]common.Hash `json:"-:"`
+    Depth       int             `json:"depth"`
+    RefundCounter uint64          `json:"refund"`
+    Err         error           `json:"-:"`
+}
+
+// overrides for gencodec
+type structLogMarshaling struct {
+    Gas      math.HexOrDecimal64
+    GasCost  math.HexOrDecimal64
+    Memory   hexutil.Bytes
+    ReturnData hexutil.Bytes
+    OpName    string `json:"opName"` // adds call to OpName() in MarshalJSON
+    ErrorString string `json:"error"` // adds call to ErrorString() in MarshalJSON
+}
+
+// OpName formats the operand name in a human-readable format.
+func (s *StructLog) OpName() string {
+    return s.Op.String()
+}
+
+// ErrorString formats the log's error as a string.
+func (s *StructLog) ErrorString() string {
+    if s.Err != nil {
+        return s.Err.Error()
+    }
+    return ""
+}
+
+// StructLogger is an EVM state logger and implements EVMLogger.
+//
+// StructLogger can capture state based on the given Log configuration and also keeps
+// a track record of modified storage which is used in reporting snapshots of the
+// contract their storage.
+type StructLogger struct {
+    cfg Config
+    env *vm.EVM
+
+    storage map[common.Address]Storage
+    logs    []StructLog
+    output  []byte
+    err     error
+}
+
+// NewStructLogger returns a new logger
+func NewStructLogger(cfg *Config) *StructLogger {
+    logger := &StructLogger{
+        storage: make(map[common.Address]Storage),
+    }
+    if cfg != nil {
+        logger.cfg = *cfg
+    }
+    return logger
+}
+
+// Reset clears the data held by the logger.
+func (l *StructLogger) Reset() {
+    l.storage = make(map[common.Address]Storage)
+    l.output = make([]byte, 0)
+    l.logs = l.logs[:0]
+    l.err = nil
+}
+
+// CaptureStart implements the EVMLogger interface to initialize the tracing operation.
+func (l *StructLogger) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+    l.env = env
+}
+
+// CaptureState logs a new structured log message and pushes it out to the environment
+//
+// CaptureState also tracks SLOAD/SSTORE ops to track storage change.
+func (l *StructLogger) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+    memory := scope.Memory
+    stack := scope.Stack
+    contract := scope.Contract
+    // check if already accumulated the specified number of logs
+    if l.cfg.Limit != 0 && l.cfg.Limit <= len(l.logs) {
+        return
+    }
+    // Copy a snapshot of the current memory state to a new buffer
+    var mem []byte
+    if l.cfg.EnableMemory {
+        mem = make([]byte, len(memory.Data()))
+        copy(mem, memory.Data())
+    }
+    // Copy a snapshot of the current stack state to a new buffer
+    var stck []uint256.Int
+    if !l.cfg.DisableStack {
+        stck = make([]uint256.Int, len(stack.Data()))
+        for i, item := range stack.Data() {
+            stck[i] = item
+        }
+    }
+    stackData := stack.Data()
+    stackLen := len(stackData)
+    // Copy a snapshot of the current storage to a new container
+    var storage Storage
+    if !l.cfg.DisableStorage && (op == vm.SLOAD || op == vm.SSTORE) {
+        // initialise new changed values storage container for this contract
+        // if not present.
+        if l.storage[contract.Address()] == nil {
+            l.storage[contract.Address()] = make(Storage)
+        }
+        // capture SLOAD opcodes and record the read entry in the local storage
+        if op == vm.SLOAD && stackLen >= 1 {
+            var (
+                address = common.Hash(stackData[stackLen-1].Bytes32())
+                value    = l.env.StateDB.GetState(contract.Address(), address)
+            )
+            l.storage[contract.Address()][address] = value
+            storage = l.storage[contract.Address()].Copy()
+        } else if op == vm.SSTORE && stackLen >= 2 {
+            // capture SSTORE opcodes and record the written entry in the local storage.
+            var (
+                value    = common.Hash(stackData[stackLen-2].Bytes32())
+                address = common.Hash(stackData[stackLen-1].Bytes32())
+            )

```

```

+         )
+         l.storage[contract.Address()][address] = value
+         storage = l.storage[contract.Address()].Copy()
+     }
+ }
+ var rdata []byte
+ if l.cfg.EnableReturnData {
+     rdata = make([]byte, len(rData))
+     copy(rdata, rData)
+ }
+ // create a new snapshot of the EVM.
+ log := StructLog{pc, op, gas, cost, mem, memory.Len(), stck, rdata, storage, depth, l.env.StateDB.GetRefund(), err}
+ l.logs = append(l.logs, log)
+}
+
+// CaptureFault implements the EVMLogger interface to trace an execution fault
+// while running an opcode.
+func (l *StructLogger) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+}
+
+// CaptureEnd is called after the call finishes to finalize the tracing.
+func (l *StructLogger) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {
+    l.output = output
+    l.err = err
+    if l.cfg.Debug {
+        fmt.Printf("0x%x\n", output)
+        if err != nil {
+            fmt.Printf(" error: %v\n", err)
+        }
+    }
+}
+
+func (l *StructLogger) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+}
+
+func (l *StructLogger) CaptureExit(output []byte, gasUsed uint64, err error) {}
+
+// StructLogs returns the captured log entries.
+func (l *StructLogger) StructLogs() []StructLog { return l.logs }
+
+// Error returns the VM error captured by the trace.
+func (l *StructLogger) Error() error { return l.err }
+
+// Output returns the VM return value captured by the trace.
+func (l *StructLogger) Output() []byte { return l.output }
+
+// WriteTrace writes a formatted trace to the given writer
+func WriteTrace(writer io.Writer, logs []StructLog) {
+    for _, log := range logs {
+        fmt.Fprintf(writer, "%-16spc=%08d gas=%v cost=%v", log.Op, log.Pc, log.Gas, log.GasCost)
+        if log.Err != nil {
+            fmt.Fprintf(writer, " ERROR: %v", log.Err)
+        }
+        fmt.Fprintln(writer)
+
+        if len(log.Stack) > 0 {
+            fmt.Fprintln(writer, "Stack:")
+            for i := len(log.Stack) - 1; i >= 0; i-- {
+                fmt.Fprintf(writer, "%08d %s\n", len(log.Stack)-i-1, log.Stack[i].Hex())
+            }
+        }
+        if len(log.Memory) > 0 {
+            fmt.Fprintln(writer, "Memory:")
+            fmt.Fprint(writer, hex.Dump(log.Memory))
+        }
+        if len(log.Storage) > 0 {
+            fmt.Fprintln(writer, "Storage:")
+            for h, item := range log.Storage {
+                fmt.Fprintf(writer, "%x: %x\n", h, item)
+            }
+        }
+        if len(log.ReturnData) > 0 {
+            fmt.Fprintln(writer, "ReturnData:")
+            fmt.Fprint(writer, hex.Dump(log.ReturnData))
+        }
+        fmt.Fprintln(writer)
+    }
+}
+
+// WriteLogs writes vm logs in a readable format to the given writer
+func WriteLogs(writer io.Writer, logs []*types.Log) {
+    for _, log := range logs {
+        fmt.Fprintf(writer, "LOG%d: %x bn=%d txi=%x\n", len(log.Topics), log.Address, log.BlockNumber, log.TxIndex)
+
+        for i, topic := range log.Topics {
+            fmt.Fprintf(writer, "%08d %x\n", i, topic)
+        }
+
+        fmt.Fprint(writer, hex.Dump(log.Data))
+        fmt.Fprintln(writer)
+    }
+}
+
+type mdLogger struct {
+    out io.Writer
+    cfg *Config
+    env *vm.EVM
+}
+
+// NewMarkdownLogger creates a logger which outputs information in a format adapted
+// for human readability, and is also a valid markdown table
+func NewMarkdownLogger(cfg *Config, writer io.Writer) *mdLogger {
+    l := &mdLogger{out: writer, cfg: cfg}
+    if l.cfg == nil {
+        l.cfg = &Config{}
+    }
+    return l
+}
+
+func (t *mdLogger) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+    t.env = env
+    if !create {
+        fmt.Fprintf(t.out, "From: `%v`\nTo: `%v`\nData: `0x%x`\nGas: `%d`\nValue `%v` wei\n",
+            from.String(), to.String(),
+            input, gas, value)
+    } else {
+        fmt.Fprintf(t.out, "From: `%v`\nCreate at: `%v`\nData: `0x%x`\nGas: `%d`\nValue `%v` wei\n",
+            from.String(), to.String(),
+            input, gas, value)
+    }
+
+    fmt.Fprintf(t.out, `
+| Pc   | Op   | Cost | Stack | RStack | Refund |
+|-----|-----|-----|-----|-----|-----|
+`)
+}
+
+// CaptureState also tracks SLOAD/SSTORE ops to track storage change.
+func (t *mdLogger) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+    stack := scope.Stack
+    fmt.Fprintf(t.out, "| %4d | %10v | %3d |", pc, op, cost)
+
+    if !t.cfg.DisableStack {
+        // format stack

```

```

+         var a []string
+         for _, elem := range stack.Data() {
+             a = append(a, elem.Hex())
+         }
+         b := fmt.Sprintf("%v]", strings.Join(a, ","))
+         fmt.Fprintf(t.out, "%10v |", b)
+     }
+     fmt.Fprintf(t.out, "%10v |", t.env.StateDB.GetRefund())
+     fmt.Fprintln(t.out, "")
+     if err != nil {
+         fmt.Fprintf(t.out, "Error: %v\n", err)
+     }
+ }
+}
+
+func (t *mdLogger) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+     fmt.Fprintf(t.out, "\nError: at pc=%d, op=%v: %v\n", pc, op, err)
+}
+
+func (t *mdLogger) CaptureEnd(output []byte, gasUsed uint64, tm time.Duration, err error) {
+     fmt.Fprintf(t.out, "\nOutput: `0x%x`\nConsumed gas: `%d`\nError: `%v`\n",
+         output, gasUsed, err)
+}
+
+func (t *mdLogger) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+}
+
+func (t *mdLogger) CaptureExit(output []byte, gasUsed uint64, err error) {}
+
+diff --git a/eth/tracers/logger/logger_json.go b/eth/tracers/logger/logger_json.go
new file mode 100644
index 00000000..5d222466
--- /dev/null
+++ b/eth/tracers/logger/logger_json.go
@@ -0,0 +1,100 @@
@@ -0,0 +1,100 @@
// Copyright 2017 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package logger
+
+import (
+     "encoding/json"
+     "io"
+     "math/big"
+     "time"
+
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/ethereum/go-ethereum/common/math"
+     "github.com/flare-foundation/coreth/core/vm"
+)
+
+type JSONLogger struct {
+     encoder *json.Encoder
+     cfg     *Config
+     env     *vm.EVM
+}
+
+// NewJSONLogger creates a new EVM tracer that prints execution steps as JSON objects
+// into the provided stream.
+func NewJSONLogger(cfg *Config, writer io.Writer) *JSONLogger {
+     l := &JSONLogger{encoder: json.NewEncoder(writer), cfg: cfg}
+     if l.cfg == nil {
+         l.cfg = &Config{}
+     }
+     return l
+}
+
+func (l *JSONLogger) CaptureStart(env *vm.EVM, from, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+     l.env = env
+}
+
+func (l *JSONLogger) CaptureFault(pc uint64, op vm.OpCode, gas uint64, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+     // TODO: Add rData to this interface as well
+     l.CaptureState(pc, op, gas, cost, scope, nil, depth, err)
+}
+
+// CaptureState outputs state information on the logger.
+func (l *JSONLogger) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+     memory := scope.Memory
+     stack := scope.Stack
+
+     log := StructLog{
+         Pc:      pc,
+         Op:      op,
+         Gas:     gas,
+         GasCost: cost,
+         MemorySize: memory.Len(),
+         Depth:   depth,
+         RefundCounter: l.env.StateDB.GetRefund(),
+         Err:      err,
+     }
+
+     if l.cfg.EnableMemory {
+         log.Memory = memory.Data()
+     }
+     if !l.cfg.DisableStack {
+         log.Stack = stack.Data()
+     }
+     if l.cfg.EnableReturnData {
+         log.ReturnData = rData
+     }
+     l.encoder.Encode(log)
+}
+
+// CaptureEnd is triggered at end of execution.
+func (l *JSONLogger) CaptureEnd(output []byte, gasUsed uint64, t time.Duration, err error) {
+     type endLog struct {
+         Output string `json:"output"`
+         GasUsed math.HexOrDecimal64 `json:"gasUsed"`
+         Time time.Duration `json:"time"`
+         Err string `json:"error,omitempty"`
+     }
+
+     var errMsg string
+     if err != nil {
+         errMsg = err.Error()
+     }
+     l.encoder.Encode(endLog{common.Bytes2Hex(output), math.HexOrDecimal64(gasUsed), t, errMsg})
+}
+
+func (l *JSONLogger) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+}
+
+func (l *JSONLogger) CaptureExit(output []byte, gasUsed uint64, err error) {}
+
+diff --git a/eth/tracers/logger/logger_test.go b/eth/tracers/logger/logger_test.go

```

```
new file mode 100644
index 00000000..954cc119
--- /dev/null
+++ b/eth/tracers/logger/logger_test.go
@@ -0,0 +1,74 @@
// Copyright 2016 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package logger
+
+import (
+    "math/big"
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
+)
+
+type dummyContractRef struct {
+    calledForEach bool
+}
+
+func (dummyContractRef) Address() common.Address    { return common.Address{} }
+func (dummyContractRef) Value() *big.Int           { return new(big.Int) }
+func (dummyContractRef) SetCode(common.Hash, []byte) {}
+func (d *dummyContractRef) ForEachStorage(callback func(key, value common.Hash) bool) {
+    d.calledForEach = true
+}
+
+func (d *dummyContractRef) SubBalance(amount *big.Int) {}
+func (d *dummyContractRef) AddBalance(amount *big.Int) {}
+func (d *dummyContractRef) SetBalance(*big.Int) {}
+func (d *dummyContractRef) SetNonce(uint64) {}
+func (d *dummyContractRef) Balance() *big.Int    { return new(big.Int) }
+
+type dummyStatedb struct {
+    state.StateDB
+}
+
+func (*dummyStatedb) GetRefund() uint64 { return 1337 }
+func (*dummyStatedb) GetState(_ common.Address, _ common.Hash) common.Hash { return common.Hash{} }
+func (*dummyStatedb) SetState(_ common.Address, _ common.Hash, _ common.Hash) {}
+
+func TestStoreCapture(t *testing.T) {
+    var (
+        logger      = NewStructLogger(nil)
+        env         = vm.NewEVM(vm.BlockContext{}, vm.TxContext{}, &dummyStatedb, params.TestChainConfig, vm.Config{Debug: true, Tracer: logger})
+        contract    = vm.NewContract(&dummyContractRef{}, &dummyContractRef{}, new(big.Int), 100000)
+    )
+    contract.Code = []byte{byte(vm.PUSH1), 0x1, byte(vm.PUSH1), 0x0, byte(vm.SSTORE)}
+    var index common.Hash
+    logger.CaptureStart(env, common.Address{}, contract.Address(), false, nil, 0, nil)
+    _, err := env.Interpreter().Run(contract, []byte{}, false)
+    if err != nil {
+        t.Fatal(err)
+    }
+    if len(logger.storage[contract.Address()]) == 0 {
+        t.Fatalf("expected exactly 1 changed value on address %x, got %d", contract.Address(),
+            len(logger.storage[contract.Address()]))
+    }
+    exp := common.BigToHash(big.NewInt(1))
+    if logger.storage[contract.Address()][index] != exp {
+        t.Errorf("expected %x, got %x", exp, logger.storage[contract.Address()][index])
+    }
+}
diff --git a/eth/tracers/native/4byte.go b/eth/tracers/native/4byte.go
new file mode 100644
index 00000000..473abc02
--- /dev/null
+++ b/eth/tracers/native/4byte.go
@@ -0,0 +1,158 @@
// (c) 2020-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
//
// *****
// Copyright 2021 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package native
+
+import (
+    "encoding/json"
+    "math/big"
+    "strconv"
+    "sync/atomic"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/tracers"
+)
+
+func init() {
+    register("4byteTracer", newFourByteTracer)
+}
+
+// fourByteTracer searches for 4byte-identifiers, and collects them for post-processing.
+// It collects the methods identifiers along with the size of the supplied data, so
+// a reversed signature can be matched against the size of the data.
+//
+// Example:
```



```

+// > debug.traceTransaction( "0x214e597e35da083692f5386141e69f47e973b2c56e7a8073b1ea08fd7571e9de", {tracer: "4byteTracer"})
+// {
+//   0x27dc297e-128: 1,
+//   0x38cc4831-0: 2,
+//   0x524f3889-96: 1,
+//   0xadf59f99-288: 1,
+//   0xc281d19e-0: 1
+// }
+type fourByteTracer struct {
+  env          *vm.EVM
+  ids          map[string]int // ids aggregates the 4byte ids found
+  interrupt     uint32    // Atomic flag to signal execution interruption
+  reason       error     // Textual reason for the interruption
+  activePrecompiles []common.Address // Updated on CaptureStart based on given rules
+}
+
+// newFourByteTracer returns a native go tracer which collects
+// 4 byte-identifiers of a tx, and implements vm.EVMLogger.
+func newFourByteTracer() tracers.Tracer {
+  t := &fourByteTracer{
+    ids: make(map[string]int),
+  }
+  return t
+}
+
+// isPrecompiled returns whether the addr is a precompile. Logic borrowed from newJsTracer in eth/tracers/js/tracer.go
+func (t *fourByteTracer) isPrecompiled(addr common.Address) bool {
+  for _, p := range t.activePrecompiles {
+    if p == addr {
+      return true
+    }
+  }
+  return false
+}
+
+// store saves the given identifier and datasize.
+func (t *fourByteTracer) store(id []byte, size int) {
+  key := bytesToHex(id) + "-" + strconv.Itoa(size)
+  t.ids[key] += 1
+}
+
+// CaptureStart implements the EVMLogger interface to initialize the tracing operation.
+func (t *fourByteTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+  t.env = env
+
+  // Update list of precompiles based on current block
+  rules := env.ChainConfig().AvalancheRules(env.Context.BlockNumber, env.Context.Time)
+  t.activePrecompiles = vm.ActivePrecompiles(rules)
+
+  // Save the outer calldata also
+  if len(input) >= 4 {
+    t.store(input[0:4], len(input)-4)
+  }
+}
+
+// CaptureState implements the EVMLogger interface to trace a single step of VM execution.
+func (t *fourByteTracer) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+}
+
+// CaptureEnter is called when EVM enters a new scope (via call, create or selfdestruct).
+func (t *fourByteTracer) CaptureEnter(op vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
+  // Skip if tracing was interrupted
+  if atomic.LoadUint32(&t.interrupt) > 0 {
+    t.env.Cancel()
+    return
+  }
+  if len(input) < 4 {
+    return
+  }
+  // primarily we want to avoid CREATE/CREATE2/SELFDESTRUCT
+  if op != vm.DELEGATECALL && op != vm.STATICCALL &&
+    op != vm.CALL && op != vm.CALLCODE {
+    return
+  }
+  // Skip any pre-compile invocations, those are just fancy opcodes
+  if t.isPrecompiled(to) {
+    return
+  }
+  t.store(input[0:4], len(input)-4)
+}
+
+// CaptureExit is called when EVM exits a scope, even if the scope didn't
+// execute any code.
+func (t *fourByteTracer) CaptureExit(output []byte, gasUsed uint64, err error) {
+}
+
+// CaptureFault implements the EVMLogger interface to trace an execution fault.
+func (t *fourByteTracer) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, depth int, err error) {
+}
+
+// CaptureEnd is called after the call finishes to finalize the tracing.
+func (t *fourByteTracer) CaptureEnd(output []byte, gasUsed uint64, _ time.Duration, err error) {
+}
+
+// GetResult returns the json-encoded nested list of call traces, and any
+// error arising from the encoding or forceful termination (via `Stop`).
+func (t *fourByteTracer) GetResult() (json.RawMessage, error) {
+  res, err := json.Marshal(t.ids)
+  if err != nil {
+    return nil, err
+  }
+  return res, t.reason
+}
+
+// Stop terminates execution of the tracer at the first opportune moment.
+func (t *fourByteTracer) Stop(err error) {
+  t.reason = err
+  atomic.StoreUint32(&t.interrupt, 1)
+}
+
+diff --git a/eth/tracers/native/call.go b/eth/tracers/native/call.go
index 757a7653..f248b3d5 100644
--- a/eth/tracers/native/call.go
+++ b/eth/tracers/native/call.go
@@ -1,4 +1,4 @@
-// (c) 2019-2020, Ava Labs, Inc.
+// (c) 2020-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
@@ -35,13 +35,13 @@ import (
  "sync/atomic"
  "time"

-  "github.com/ava-labs/coreth/core/vm"
-  "github.com/ava-labs/coreth/eth/tracers"
-  "github.com/ethereum/go-ethereum/common"
+  "github.com/flare-foundation/coreth/core/vm"
+  "github.com/flare-foundation/coreth/eth/tracers"
 )

 func init() {
-  tracers.RegisterNativeTracer("callTracer", NewCallTracer)
+  register("callTracer", newCallTracer)
 }

```

```

type callFrame struct {
@@ -58,21 +58,24 @@ type callFrame struct {
}

type callTracer struct {
+     env      *vm.EVM
    callstack []callFrame
    interrupt uint32 // Atomic flag to signal execution interruption
    reason    error  // Textual reason for the interruption
}

-// NewCallTracer returns a native go tracer which tracks
+// newCallTracer returns a native go tracer which tracks
// call frames of a tx, and implements vm.EVMLogger.
-func NewCallTracer() tracers.Tracer {
+func newCallTracer() tracers.Tracer {
    // First callframe contains tx context info
    // and is populated on start and end.
    t := &callTracer{callstack: make([]callFrame, 1)}
    return t
}

+// CaptureStart implements the EVMLogger interface to initialize the tracing operation.
func (t *callTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
+
    t.env = env
    t.callstack[0] = callFrame{
        Type: "CALL",
        From: addrToHex(from),
@@ -86,6 +89,7 @@ func (t *callTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Ad
    }
}

+// CaptureEnd is called after the call finishes to finalize the tracing.
func (t *callTracer) CaptureEnd(output []byte, gasUsed uint64, _ time.Duration, err error) {
    t.callstack[0].GasUsed = uintToHex(gasUsed)
    if err != nil {
@@ -98,16 +102,19 @@ func (t *callTracer) CaptureEnd(output []byte, gasUsed uint64, _ time.Duration,
    }
}

-func (t *callTracer) CaptureState(env *vm.EVM, pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
+// CaptureState implements the EVMLogger interface to trace a single step of VM execution.
+func (t *callTracer) CaptureState(pc uint64, op vm.OpCode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
}

-func (t *callTracer) CaptureFault(env *vm.EVM, pc uint64, op vm.OpCode, gas, cost uint64, _ *vm.ScopeContext, depth int, err error) {
+// CaptureFault implements the EVMLogger interface to trace an execution fault.
+func (t *callTracer) CaptureFault(pc uint64, op vm.OpCode, gas, cost uint64, _ *vm.ScopeContext, depth int, err error) {
}

+// CaptureEnter is called when EVM enters a new scope (via call, create or selfdestruct).
func (t *callTracer) CaptureEnter(typ vm.OpCode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
    // Skip if tracing was interrupted
    if atomic.LoadUint32(&t.interrupt) > 0 {
-
-     // TODO: env.Cancel()
+
+     t.env.Cancel()
    return
}

@@ -122,6 +129,8 @@ func (t *callTracer) CaptureEnter(typ vm.OpCode, from common.Address, to common.
    t.callstack = append(t.callstack, call)
}

+// CaptureExit is called when EVM exits a scope, even if the scope didn't
+// execute any code.
func (t *callTracer) CaptureExit(output []byte, gasUsed uint64, err error) {
    size := len(t.callstack)
    if size <= 1 {
@@ -144,6 +153,8 @@ func (t *callTracer) CaptureExit(output []byte, gasUsed uint64, err error) {
    t.callstack[size-1].Calls = append(t.callstack[size-1].Calls, call)
}

+// GetResult returns the json-encoded nested list of call traces, and any
+// error arising from the encoding or forceful termination (via `Stop`).
func (t *callTracer) GetResult() (json.RawMessage, error) {
    if len(t.callstack) != 1 {
        return nil, errors.New("incorrect number of top-level calls")
@@ -155,6 +166,7 @@ func (t *callTracer) GetResult() (json.RawMessage, error) {
    return json.RawMessage(res), t.reason
}

+// Stop terminates execution of the tracer at the first opportune moment.
func (t *callTracer) Stop(err error) {
    t.reason = err
    atomic.StoreUint32(&t.interrupt, 1)
}
diff --git a/eth/tracers/native/noop.go b/eth/tracers/native/noop.go
index 2b0eb09c..cddb9f 100644
--- a/eth/tracers/native/noop.go
+++ b/eth/tracers/native/noop.go
@@ -1,3 +1,3 @@
@@ -1,3 +1,3 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2021 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
package native

import (
@@ -5,42 +31,54 @@ import (
    "math/big"
    "time"

-     "github.com/ava-labs/coreth/core/vm"
-     "github.com/ava-labs/coreth/eth/tracers"
-     "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/coreth/core/vm"
+     "github.com/flare-foundation/coreth/eth/tracers"
)

func init() {
-     tracers.RegisterNativeTracer("noopTracerNative", NewNoopTracer)
+     register("noopTracerNative", newNoopTracer)
}

```

```

}

// noopTracer is a go implementation of the Tracer interface which
// performs no action. It's mostly useful for testing purposes.
type noopTracer struct{}

-func NewNoopTracer() tracers.Tracer {
// newNoopTracer returns a new noop tracer.
+func newNoopTracer() tracers.Tracer {
    return &noopTracer{}
}

// CaptureStart implements the EVMLogger interface to initialize the tracing operation.
func (t *noopTracer) CaptureStart(env *vm.EVM, from common.Address, to common.Address, create bool, input []byte, gas uint64, value *big.Int) {
}

// CaptureEnd is called after the call finishes to finalize the tracing.
func (t *noopTracer) CaptureEnd(output []byte, gasUsed uint64, _ time.Duration, err error) {
}

-func (t *noopTracer) CaptureState(env *vm.EVM, pc uint64, op vm.Opcode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
// CaptureState implements the EVMLogger interface to trace a single step of VM execution.
+func (t *noopTracer) CaptureState(pc uint64, op vm.Opcode, gas, cost uint64, scope *vm.ScopeContext, rData []byte, depth int, err error) {
}

-func (t *noopTracer) CaptureFault(env *vm.EVM, pc uint64, op vm.Opcode, gas, cost uint64, _ *vm.ScopeContext, depth int, err error) {
// CaptureFault implements the EVMLogger interface to trace an execution fault.
+func (t *noopTracer) CaptureFault(pc uint64, op vm.Opcode, gas, cost uint64, _ *vm.ScopeContext, depth int, err error) {
}

// CaptureEnter is called when EVM enters a new scope (via call, create or selfdestruct).
func (t *noopTracer) CaptureEnter(typ vm.Opcode, from common.Address, to common.Address, input []byte, gas uint64, value *big.Int) {
}

// CaptureExit is called when EVM exits a scope, even if the scope didn't
// execute any code.
func (t *noopTracer) CaptureExit(output []byte, gasUsed uint64, err error) {
}

// GetResult returns an empty json object.
func (t *noopTracer) GetResult() (json.RawMessage, error) {
    return json.RawMessage(`{}`, nil)
}

// Stop terminates execution of the tracer at the first opportune moment.
func (t *noopTracer) Stop(err error) {
}

diff --git a/eth/tracers/native/tracer.go b/eth/tracers/native/tracer.go
new file mode 100644
index 00000000..b60883ab
--- /dev/null
+++ b/eth/tracers/native/tracer.go
@@ -0,0 +1,89 @@
// (c) 2020-2021, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
//
// *****
// Copyright 2021 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+/*
+Package native is a collection of tracers written in go.
+
+In order to add a native tracer and have it compiled into the binary, a new
+file needs to be added to this folder, containing an implementation of the
+'eth.tracers.Tracer' interface.
+
+Aside from implementing the tracer, it also needs to register itself, using the
+'register' method -- and this needs to be done in the package initialization.
+
+Example:
+
+```golang
+func init() {
+    register("noopTracerNative", newNoopTracer)
+}
+```
+*/
+package native
+
+import (
+    "errors"
+
+    "github.com/flare-foundation/coreth/eth/tracers"
+)
+
+// init registers itself this packages as a lookup for tracers.
+func init() {
+    tracers.RegisterLookup(false, lookup)
+}
+
+/*
+ctors is a map of package-local tracer constructors.
+
+We cannot be certain about the order of init-functions within a package,
+The go spec (https://golang.org/ref/spec#Package_initialization) says
+
+> To ensure reproducible initialization behavior, build systems
+> are encouraged to present multiple files belonging to the same
+> package in lexical file name order to a compiler.
+
+Hence, we cannot make the map in init, but must make it upon first use.
+*/
+var ctors map[string]func() tracers.Tracer
+
+// register is used by native tracers to register their presence.
+func register(name string, ctor func() tracers.Tracer) {
+    if ctors == nil {
+        ctors = make(map[string]func() tracers.Tracer)
+    }
+    ctors[name] = ctor
+}
+
+// lookup returns a tracer, if one can be matched to the given name.
+func lookup(name string, ctx *tracers.Context) (tracers.Tracer, error) {

```

```

+     if ctors == nil {
+         ctors = make(map[string]func() tracers.Tracer)
+     }
+     if ctor, ok := ctors[name]; ok {
+         return ctor(), nil
+     }
+     return nil, errors.New("no tracer found")
+}
diff --git a/eth/tracers/testing/calltrace_test.go b/eth/tracers/testing/calltrace_test.go
index 8cac229d..072f624c 100644
--- a/eth/tracers/testing/calltrace_test.go
+++ b/eth/tracers/testing/calltrace_test.go
@@ -10,19 +10,20 @@ import (
    "testing"
    "unicode"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/eth/tracers"
-    "github.com/ava-labs/coreth/tests"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
-    "github.com/ethereum/go-ethereum/common/math"
-    "github.com/ethereum/go-ethereum/rpc"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/tracers"
+    "github.com/flare-foundation/coreth/tests"
+
+    // Force-load the native, to trigger registration
-    "github.com/ava-labs/coreth/eth/tracers/native"
+    "github.com/flare-foundation/coreth/eth/tracers/native"
)

type callContext struct {
@@ -106,7 +107,7 @@ func testCallTracer(tracerName string, dirPath string, t *testing.T) {
    context = vm.BlockContext{
        CanTransfer: core.CanTransfer,
        Transfer:     core.Transfer,
-       Coinbase:     test.Context.Miner,
+       Coinbase:     common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000000"),
        BlockNumber:  new(big.Int).SetUint64(uint64(test.Context.Number)),
        Time:          new(big.Int).SetUint64(uint64(test.Context.Time)),
        Difficulty:    (*big.Int)(test.Context.Difficulty),
@@ -217,7 +218,7 @@ func benchTracer(tracerName string, test *callTracerTest, b *testing.B) {
    context := vm.BlockContext{
        CanTransfer: core.CanTransfer,
        Transfer:     core.Transfer,
-       Coinbase:     test.Context.Miner,
+       Coinbase:     common.HexToAddress("0x0100000000000000000000000000000000000000000000000000000000000000"),
        BlockNumber:  new(big.Int).SetUint64(uint64(test.Context.Number)),
        Time:          new(big.Int).SetUint64(uint64(test.Context.Time)),
        Difficulty:    (*big.Int)(test.Context.Difficulty),
diff --git a/eth/tracers/tracer.go b/eth/tracers/tracer.go
index 416c085b..8b531633 100644
--- a/eth/tracers/tracer.go
+++ b/eth/tracers/tracer.go
@@ -35,13 +35,15 @@ import (
    "time"
    "unsafe"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/vm"
+    "gopkg.in/olebedev/go-duktape.v3"
+
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/hexutil"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/log"
-    "gopkg.in/olebedev/go-duktape.v3"
+
    "github.com/flare-foundation/coreth/core"
    "github.com/flare-foundation/coreth/core/vm"
)

// bigIntegerJS is the minified version of https://github.com/peterolson/BigInteger.js.
diff --git a/eth/tracers/tracer_test.go b/eth/tracers/tracer_test.go
index 72a62026..10c64b64 100644
--- a/eth/tracers/tracer_test.go
+++ b/eth/tracers/tracer_test.go
@@ -33,10 +33,11 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/params"
)

type account struct{}
diff --git a/eth/tracers/tracers.go b/eth/tracers/tracers.go
index d90cec62..908084c0 100644
--- a/eth/tracers/tracers.go
+++ b/eth/tracers/tracers.go
@@ -1,13 +1,14 @@
-// (c) 2019-2020, Ava Labs, Inc.
-//
-// This file is a derived work, based on the go-ethereum library whose original
-// notices appear below.
-//
-// It is distributed under a license compatible with the licensing terms of the
-// original code from which it is derived.
-//
-// Much love to the original authors for their work.
-// *****
-// Copyright 2017 The go-ethereum Authors
-// This file is part of the go-ethereum library.
-//
@@ -24,19 +24,20 @@
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.

-// Package tracers is a collection of JavaScript transaction tracers.
package tracers

import (
    "encoding/json"

    "strings"
    "unicode"
    "errors"

    "github.com/ava-labs/coreth/core/vm"

```

```

-      "github.com/ava-labs/coreth/eth/tracers/internal/tracers"
+      "github.com/ethereum/go-ethereum/common"
+      "github.com/flare-foundation/coreth/core/vm"
)

+// Context contains some contextual infos for a transaction execution that is not
+// available from within the EVM object.
+type Context struct {
+    BlockHash common.Hash // Hash of the block the tx is contained within (zero if dangling tx or call)
+    TxIndex   int           // Index of the transaction within a block (zero if dangling tx or call)
+    TxHash    common.Hash // Hash of the transaction being traced (zero if dangling call)
+}
+
+// Tracer interface extends vm.EVMLogger and additionally
+// allows collecting the tracing result.
+type Tracer interface {
@@ -46,50 +42,31 @@ type Tracer interface {
    Stop(err error)
}

+type lookupFunc func(string, *Context) (Tracer, error)
+
+var (
-    nativeTracers map[string]func() Tracer = make(map[string]func() Tracer)
-    jsTracers      = make(map[string]string)
+    lookups []lookupFunc
)

-// RegisterNativeTracer makes native tracers which adhere
-// to the `Tracer` interface available to the rest of the codebase.
-// It is typically invoked in the `init()` function, e.g. see the `native/call.go`.
-func RegisterNativeTracer(name string, ctor func() Tracer) {
-    nativeTracers[name] = ctor
-}
-
-// New returns a new instance of a tracer,
-// 1. If 'code' is the name of a registered native tracer, then that tracer
-//    is instantiated and returned
-// 2. If 'code' is the name of a registered js-tracer, then that tracer is
-//    instantiated and returned
-// 3. Otherwise, the code is interpreted as the js code of a js-tracer, and
-//    is evaluated and returned.
-func New(code string, ctx *Context) (Tracer, error) {
-    // Resolve native tracer
-    if fn, ok := nativeTracers[code]; ok {
-        return fn(), nil
+// RegisterLookup registers a method as a lookup for tracers, meaning that
+// users can invoke a named tracer through that lookup. If 'wildcard' is true,
+// then the lookup will be placed last. This is typically meant for interpreted
+// engines (js) which can evaluate dynamic user-supplied code.
+func RegisterLookup(wildcard bool, lookup lookupFunc) {
+    if wildcard {
+        lookups = append(lookups, lookup)
+    } else {
+        lookups = append([]lookupFunc{lookup}, lookups...)
+    }
+    // Resolve js-tracers by name and assemble the tracer object
+    if tracer, ok := jsTracers[code]; ok {
+        code = tracer
+    }
+    return newJsTracer(code, ctx)
}

-// camel converts a snake cased input string into a camel cased output.
-func camel(str string) string {
-    pieces := strings.Split(str, "_")
-    for i := 1; i < len(pieces); i++ {
-        pieces[i] = string(unicode.ToUpper(rune(pieces[i][0]))) + pieces[i][1:]
-    }
-    return strings.Join(pieces, "")
-}
-
-// init retrieves the JavaScript transaction tracers included in go-ethereum.
-func init() {
-    for _, file := range tracers.AssetNames() {
-        name := camel(strings.TrimSuffix(file, ".js"))
-        jsTracers[name] = string(tracers.MustAsset(file))
+// New returns a new instance of a tracer, by iterating through the
+// registered lookups.
+func New(code string, ctx *Context) (Tracer, error) {
+    for _, lookup := range lookups {
+        if tracer, err := lookup(code, ctx); err == nil {
+            return tracer, nil
+        }
+    }
+    return nil, errors.New("tracer not found")
}

diff --git a/eth/tracers/tracers_test.go b/eth/tracers/tracers_test.go
index 3421d7c0..d6a2217d 100644
--- a/eth/tracers/tracers_test.go
+++ b/eth/tracers/tracers_test.go
@@ -27,74 +27,21 @@ package tracers

import (
-    "crypto/ecdsa"
-    "crypto/rand"
-    "encoding/json"
-    "math/big"
-    "reflect"
-    "testing"
-
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/tests"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/common/hexutil"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/eth/tracers/logger"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/tests"
)

-// To generate a new callTracer test, copy paste the makeTest method below into
-// a Geth console and call it with a transaction hash you which to export.
-
-/*
-// makeTest generates a callTracer test by running a prestate reassembled and a
-// call trace run, assembling all the gathered information into a test case.
-var makeTest = function(tx, rewind) {
-    // Generate the genesis block from the block, transaction and prestate data
-    var block = eth.getBlock(eth.getTransaction(tx).blockHash);
-    var genesis = eth.getBlock(block.parentHash);
-
-    delete genesis.gasUsed;
-    delete genesis.logsBloom;
-    delete genesis.parentHash;

```

```

- delete genesis.receiptsRoot;
- delete genesis.sha3Uncles;
- delete genesis.size;
- delete genesis.transactions;
- delete genesis.transactionsRoot;
- delete genesis.uncles;
-
- genesis.gasLimit = genesis.gasLimit.toString();
- genesis.number = genesis.number.toString();
- genesis.timestamp = genesis.timestamp.toString();
-
- genesis.alloc = debug.traceTransaction(tx, {tracer: "prestateTracer", rewind: rewind});
- for (var key in genesis.alloc) {
-     genesis.alloc[key].nonce = genesis.alloc[key].nonce.toString();
- }
- genesis.config = admin.nodeInfo.protocols.eth.config;
-
- // Generate the call trace and produce the test input
- var result = debug.traceTransaction(tx, {tracer: "callTracer", rewind: rewind});
- delete result.time;
-
- console.log(JSON.stringify({
-     genesis: genesis,
-     context: {
-         number: block.number.toString(),
-         difficulty: block.difficulty,
-         timestamp: block.timestamp.toString(),
-         gasLimit: block.gasLimit.toString(),
-         miner: block.miner,
-     },
-     input: eth.getRawTransaction(tx),
-     result: result,
- }, null, 2));
-}
-*/
-
-// callTrace is the result of a callTracer run.
-type callTrace struct {
-    Type string `json:"type"`
-}
-@ -109,184 +56,6 @@ type callTrace struct {
-    Calls []callTrace `json:"calls,omitempty"`
-}
-
-// TestZeroValueToNotExitCall tests the calltracer(s) on the following:
-// Tx to A, A calls B with zero value. B does not already exist.
-// Expected: that enter/exit is invoked and the inner call is shown in the result
-+func TestZeroValueToNotExitCall(t *testing.T) {
-+    var to = common.HexToAddress("0x0000000000000000000000000000000000000000000000000000000000000000deadbeef")
-+    privkey, err := crypto.HexToECDSA("0000000000000000000000000000000000000000000000000000000000000000deadbeef")
-+    if err != nil {
-+        t.Fatalf("err %v", err)
-+    }
-+    signer := types.NewEIP155Signer(big.NewInt(1))
-+    tx, err := types.SignNewTx(privkey, signer, &types.LegacyTx{
-+        GasPrice: big.NewInt(0),
-+        Gas:      50000,
-+        To:       &to,
-+    })
-+    if err != nil {
-+        t.Fatalf("err %v", err)
-+    }
-+    origin, _ := signer.Sender(tx)
-+    txContext := vm.TxContext{
-+        Origin: origin,
-+        GasPrice: big.NewInt(1),
-+    }
-+    context := vm.BlockContext{
-+        CanTransfer: core.CanTransfer,
-+        Transfer:    core.Transfer,
-+        Coinbase:    common.Address{},
-+        BlockNumber: new(big.Int).SetUint64(8000000),
-+        Time:        new(big.Int).SetUint64(5),
-+        Difficulty:   big.NewInt(0x30000),
-+        GasLimit:     uint64(6000000),
-+    }
-+    var code = []byte{
-+        byte(vm.PUSH1), 0x0, byte(vm.DUP1), byte(vm.DUP1), byte(vm.DUP1), // in and outs zero
-+        byte(vm.DUP1), byte(vm.PUSH1), 0xff, byte(vm.GAS), // value=0, address=0xff, gas=GAS
-+        byte(vm.CALL),
-+    }
-+    var alloc = core.GenesisAlloc{
-+        to: core.GenesisAccount{
-+            Nonce: 1,
-+            Code: code,
-+        },
-+        origin: core.GenesisAccount{
-+            Nonce: 0,
-+            Balance: big.NewInt(500000000000000),
-+        },
-+    }
-+    , statedb := tests.MakePreState(rawdb.NewMemoryDatabase(), alloc, false)
-+    // Create the tracer, the EVM environment and run it
-+    tracer, err := New("callTracerJs", new(Context))
-+    if err != nil {
-+        t.Fatalf("failed to create call tracer: %v", err)
-+    }
-+    evm := vm.NewEVM(context, txContext, statedb, params.AvalancheMainnetChainConfig, vm.Config{Debug: true, Tracer: tracer})
-+    msg, err := tx.AsMessage(signer, nil)
-+    if err != nil {
-+        t.Fatalf("failed to prepare transaction for tracing: %v", err)
-+    }
-+    st := core.NewStateTransition(evm, msg, new(core.GasPool).AddGas(tx.Gas()))
-+    if _, err = st.TransitionDb(); err != nil {
-+        t.Fatalf("failed to execute transaction: %v", err)
-+    }
-+    // Retrieve the trace result and compare against the etalon
-+    res, err := tracer.GetResult()
-+    if err != nil {
-+        t.Fatalf("failed to retrieve trace result: %v", err)
-+    }
-+    have := new(callTrace)
-+    if err := json.Unmarshal(res, have); err != nil {
-+        t.Fatalf("failed to unmarshal trace result: %v", err)
-+    }
-+    wantStr := `{"type":"CALL","from":"0x682a80a6f560eec50d54e63cbda1c324c5f8d1b","to":"0x0000000000000000000000000000000000000000000000000000000000000000deadbeef","value":"0x0","gas":"0x7148","gasUsed":"0x2d0","input":"0x","output":null}`
-+    want := new(callTrace)
-+    json.Unmarshal([]byte(wantStr), want)
-+    if !json.Equal(have, want) {
-+        t.Error("have != want")
-+    }
-+}
-
-+func TestPrestateTracerCreate2(t *testing.T) {
-+    unsignedTx := types.NewTransaction(1, common.HexToAddress("0x0000000000000000000000000000000000000000000000000000000000000000deadbeef"),
-+        new(big.Int), 5000000, big.NewInt(1), []byte{})
-+
-+    privateKeyECDSA, err := ecdsa.GenerateKey(crypto.S256(), rand.Reader)
-+    if err != nil {
-+        t.Fatalf("err %v", err)
-+    }
-+    signer := types.NewEIP155Signer(big.NewInt(1))
-+    tx, err := types.SignTx(unsignedTx, signer, privateKeyECDSA)
-+    if err != nil {
-+        t.Fatalf("err %v", err)
-+    }
-
-}

```

```

-     }
-     /**
-         This comes from one of the test-vectors on the Skinny Create2 - EIP
-
-         address 0x00000000000000000000000000000000deadbeef
-         salt 0x0000000000000000000000000000000000000000000000000000000000000000cafebabe
-         init_code 0xdeadbeef
-         gas (assuming no mem expansion): 32006
-         result: 0x60f3f640a8508fc6a86d45df051962668e1e8ac7
-     */
-     origin, _ := signer.Sender(tx)
-     txContext := vm.TxContext{
-         Origin:    origin,
-         GasPrice:   big.NewInt(1),
-     }
-     context := vm.BlockContext{
-         CanTransfer: core.CanTransfer,
-         Transfer:    core.Transfer,
-         Coinbase:    common.Address{},
-         BlockNumber: new(big.Int).SetUint64(8000000),
-         Time:        new(big.Int).SetUint64(5),
-         Difficulty:   big.NewInt(0x30000),
-         GasLimit:     uint64(6000000),
-     }
-     alloc := core.GenesisAlloc{
-
-         // The code pushes 'deadbeef' into memory, then the other params, and calls CREATE2, then returns
-         // the address
-         alloc[common.HexToAddress("0x00000000000000000000000000000000deadbeef")] = core.GenesisAccount{
-             Nonce:    1,
-             Code:     hexutil.MustDecode("0x63deadbeef60005263cafebabe6004601c6000f560005260206000f3"),
-             Balance:   big.NewInt(1),
-         }
-         alloc[origin] = core.GenesisAccount{
-             Nonce:    1,
-             Code:     []byte{},
-             Balance:   big.NewInt(5000000000000000),
-         }
-     }
-     _, statedb := tests.MakePreState(rawdb.NewMemoryDatabase(), alloc, false)
-
-     // Create the tracer, the EVM environment and run it
-     tracer, err := New("prestateTracer", new(Context))
-     if err != nil {
-         t.Fatalf("failed to create call tracer: %v", err)
-     }
-     evm := vm.NewEVM(context, txContext, statedb, params.AvalancheMainnetChainConfig, vm.Config{Debug: true, Tracer: tracer})
-
-     msg, err := tx.AsMessage(signer, nil)
-     if err != nil {
-         t.Fatalf("failed to prepare transaction for tracing: %v", err)
-     }
-     st := core.NewStateTransition(evm, msg, new(core.GasPool).AddGas(tx.Gas()))
-     if _, err = st.TransitionDb(); err != nil {
-         t.Fatalf("failed to execute transaction: %v", err)
-     }
-     // Retrieve the trace result and compare against the etalon
-     res, err := tracer.GetResult()
-     if err != nil {
-         t.Fatalf("failed to retrieve trace result: %v", err)
-     }
-     ret := make(map[string]interface{})
-     if err := json.Unmarshal(res, &ret); err != nil {
-         t.Fatalf("failed to unmarshal trace result: %v", err)
-     }
-     if _, has := ret["0x60f3f640a8508fc6a86d45df051962668e1e8ac7"]; !has {
-         t.Fatalf("Expected 0x60f3f640a8508fc6a86d45df051962668e1e8ac7 in result")
-     }
- }
-
- // jsonEqual is similar to reflect.DeepEqual, but does a 'bounce' via json prior to
- // comparison
- func jsonEqual(x, y interface{}) bool {
-     xTrace := new(callTrace)
-     yTrace := new(callTrace)
-     if xj, err := json.Marshal(x); err == nil {
-         json.Unmarshal(xj, xTrace)
-     } else {
-         return false
-     }
-     if yj, err := json.Marshal(y); err == nil {
-         json.Unmarshal(yj, yTrace)
-     } else {
-         return false
-     }
-     return reflect.DeepEqual(xTrace, yTrace)
- }
-
- func BenchmarkTransactionTrace(b *testing.B) {
-     key, _ := crypto.HexToECDSA("b71c71a67e1177ad4e901695elb4b9ee17ae16c6668d313eac2f96dbcca3f291")
-     from := crypto.PublicKeyToAddress(key.PublicKey)
@@ -337,7 +106,7 @@ func BenchmarkTransactionTrace(b *testing.B) {
     }, statedb := tests.MakePreState(rawdb.NewMemoryDatabase(), alloc, false)
     // Create the tracer, the EVM environment and run it
-     tracer := vm.NewStructLogger(&vm.LogConfig{
+     tracer := logger.NewStructLogger(&logger.Config{
         Debug: false,
         //DisableStorage: true,
         //EnableMemory: false,
diff --git a/ethclient/client_interface_test.go b/ethclient/client_interface_test.go
new file mode 100644
index 00000000..80783bcc
--- /dev/null
+++ b/ethclient/client_interface_test.go
@@ -0,0 +1,17 @@
+package ethclient
+
+import (
+    "reflect"
+    "testing"
+)
+
+func TestInterfaceStructOneToOne(t *testing.T) {
+    // checks struct provides at least the methods signatures in the interface
+    var _ Client = (*client)(nil)
+    // checks interface and struct have the same number of methods
+    clientType := reflect.TypeOf(&client{})
+    ClientType := reflect.TypeOf((*Client)(nil)).Elem()
+    if clientType.NumMethod() != ClientType.NumMethod() {
+        t.Fatalf("no 1 to 1 compliance between struct methods (%v) and interface methods (%v)", clientType.NumMethod(), ClientType.NumMethod())
+    }
+}
diff --git a/ethclient/coreethclient/coreethclient.go b/ethclient/coreethclient/coreethclient.go
index 8d7654fc..a0881dd5 100644
--- a/ethclient/coreethclient/coreethclient.go
+++ b/ethclient/coreethclient/coreethclient.go
@@ -33,12 +33,12 @@ import (
     "runtime"
     "runtime/debug"
 
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/ethclient"
-    "github.com/ava-labs/coreth/interfaces"
-    "github.com/ava-labs/coreth/rpc"

```

```

        "github.com/ethereum/go-ethereum/common"
        "github.com/ethereum/go-ethereum/common/hexutil"
+       "github.com/flare-foundation/coreth/core/types"
+       "github.com/flare-foundation/coreth/ethclient"
+       "github.com/flare-foundation/coreth/interfaces"
+       "github.com/flare-foundation/coreth/rpc"
    )

    // Client is a wrapper around rpc.Client that implements geth-specific functionality.
diff --git a/ethclient/ethclient.go b/ethclient/ethclient.go
index 81d988e2..b0d6e182 100644
--- a/ethclient/ethclient.go
+++ b/ethclient/ethclient.go
@@ -34,46 +34,84 @@ import (
    "fmt"
    "math/big"

-   "github.com/ava-labs/avalanchego/ids"
-   "github.com/ava-labs/coreth/accounts/abi/bind"
-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/interfaces"
-   "github.com/ava-labs/coreth/rpc"
+   "github.com/ethereum/go-ethereum/common"
+   "github.com/ethereum/go-ethereum/common/hexutil"
+   "github.com/flare-foundation/coreth/accounts/abi/bind"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/interfaces"
+   "github.com/flare-foundation/coreth/rpc"
+   "github.com/flare-foundation/flare/ids"
)

// Verify that Client implements required interfaces
var (
-   _ bind.AcceptedContractCaller = (*Client)(nil)
-   _ bind.ContractBackend        = (*Client)(nil)
-   _ bind.ContractFilterer       = (*Client)(nil)
-   _ bind.ContractTransactor     = (*Client)(nil)
-   _ bind.DeployBackend          = (*Client)(nil)
-
-   _ interfaces.ChainReader      = (*Client)(nil)
-   _ interfaces.ChainStateReader = (*Client)(nil)
-   _ interfaces.TransactionReader = (*Client)(nil)
-   _ interfaces.TransactionSender = (*Client)(nil)
-   _ interfaces.ContractCaller   = (*Client)(nil)
-   _ interfaces.GasEstimator     = (*Client)(nil)
-   _ interfaces.GasPricer        = (*Client)(nil)
-   _ interfaces.LogFilterer      = (*Client)(nil)
-   _ interfaces.AcceptedStateReader = (*Client)(nil)
-   _ interfaces.AcceptedContractCaller = (*Client)(nil)
+   _ bind.AcceptedContractCaller = (*client)(nil)
+   _ bind.ContractBackend        = (*client)(nil)
+   _ bind.ContractFilterer       = (*client)(nil)
+   _ bind.ContractTransactor     = (*client)(nil)
+   _ bind.DeployBackend          = (*client)(nil)
+
+   _ interfaces.ChainReader      = (*client)(nil)
+   _ interfaces.ChainStateReader = (*client)(nil)
+   _ interfaces.TransactionReader = (*client)(nil)
+   _ interfaces.TransactionSender = (*client)(nil)
+   _ interfaces.ContractCaller   = (*client)(nil)
+   _ interfaces.GasEstimator     = (*client)(nil)
+   _ interfaces.GasPricer        = (*client)(nil)
+   _ interfaces.LogFilterer      = (*client)(nil)
+   _ interfaces.AcceptedStateReader = (*client)(nil)
+   _ interfaces.AcceptedContractCaller = (*client)(nil)
+
+   _ Client = (*client)(nil)
)

-// Client defines typed wrappers for the Ethereum RPC API.
-type Client struct {
+// Client defines interface for typed wrappers for the Ethereum RPC API.
+type Client interface {
+   Close()
+   ChainID(context.Context) (*big.Int, error)
+   BlockByHash(context.Context, common.Hash) (*types.Block, error)
+   BlockByNumber(context.Context, *big.Int) (*types.Block, error)
+   BlockNumber(context.Context) (uint64, error)
+   HeaderByHash(context.Context, common.Hash) (*types.Header, error)
+   HeaderByNumber(context.Context, *big.Int) (*types.Header, error)
+   TransactionByHash(context.Context, common.Hash) (tx *types.Transaction, isPending bool, err error)
+   TransactionSender(context.Context, *types.Transaction, common.Hash, uint) (common.Address, error)
+   TransactionCount(context.Context, common.Hash) (uint, error)
+   TransactionInBlock(context.Context, common.Hash, uint) (*types.Transaction, error)
+   TransactionReceipt(context.Context, common.Hash) (*types.Receipt, error)
+   SubscribeNewAcceptedTransactions(context.Context, chan<- *common.Hash) (interfaces.Subscription, error)
+   SubscribeNewPendingTransactions(context.Context, chan<- *common.Hash) (interfaces.Subscription, error)
+   SubscribeNewHead(context.Context, chan<- *types.Header) (interfaces.Subscription, error)
+   NetworkID(context.Context) (*big.Int, error)
+   BalanceAt(context.Context, common.Address, *big.Int) (*big.Int, error)
+   AssetBalanceAt(context.Context, common.Address, ids.ID, *big.Int) (*big.Int, error)
+   StorageAt(context.Context, common.Address, common.Hash, *big.Int) ([]byte, error)
+   CodeAt(context.Context, common.Address, *big.Int) ([]byte, error)
+   NonceAt(context.Context, common.Address, *big.Int) (uint64, error)
+   FilterLogs(context.Context, interfaces.FilterQuery) ([]types.Log, error)
+   SubscribeFilterLogs(context.Context, interfaces.FilterQuery, chan<- types.Log) (interfaces.Subscription, error)
+   AcceptedCodeAt(context.Context, common.Address) ([]byte, error)
+   AcceptedNonceAt(context.Context, common.Address) (uint64, error)
+   AcceptedCallContract(context.Context, interfaces.CallMsg) ([]byte, error)
+   CallContract(context.Context, interfaces.CallMsg, *big.Int) ([]byte, error)
+   SuggestGasPrice(context.Context) (*big.Int, error)
+   SuggestGasTipCap(context.Context) (*big.Int, error)
+   EstimateGas(context.Context, interfaces.CallMsg) (uint64, error)
+   EstimateBaseFee(context.Context) (*big.Int, error)
+   SendTransaction(context.Context, *types.Transaction) error
+}
+
+// client defines implementation for typed wrappers for the Ethereum RPC API.
+type client struct {
+   c *rpc.Client
+}

    // Dial connects a client to the given URL.
-   -func Dial(rawurl string) (*Client, error) {
+   +func Dial(rawurl string) (Client, error) {
        return DialContext(context.Background(), rawurl)
    }

-   -func DialContext(ctx context.Context, rawurl string) (*Client, error) {
+   +func DialContext(ctx context.Context, rawurl string) (Client, error) {
        c, err := rpc.DialContext(ctx, rawurl)
        if err != nil {
            return nil, err
        }
@@ -82,18 +120,18 @@ func DialContext(ctx context.Context, rawurl string) (*Client, error) {
    }

    // NewClient creates a client that uses the given RPC client.
-   -func NewClient(c *rpc.Client) *Client {
-       return &Client{c}
+   +func NewClient(c *rpc.Client) Client {
+       return &client{c}
    }

-   -func (ec *Client) Close() {

```



```

+func (ec *Client) Close() {
    ec.c.Close()
}

// Blockchain Access

// ChainID retrieves the current chain ID for transaction replay protection.
-func (ec *Client) ChainID(ctx context.Context) (*big.Int, error) {
+func (ec *Client) ChainID(ctx context.Context) (*big.Int, error) {
    var result hexutil.Big
    err := ec.c.CallContext(ctx, &result, "eth_chainId")
    if err != nil {
@@ -106,7 +144,7 @@ func (ec *Client) ChainID(ctx context.Context) (*big.Int, error) {
    //
    // Note that loading full blocks requires two requests. Use HeaderByHash
    // if you don't need all transactions or uncle headers.
-func (ec *Client) BlockByHash(ctx context.Context, hash common.Hash) (*types.Block, error) {
+func (ec *Client) BlockByHash(ctx context.Context, hash common.Hash) (*types.Block, error) {
    return ec.getBlock(ctx, "eth_getBlockByHash", hash, true)
}

@@ -115,12 +153,12 @@ func (ec *Client) BlockByHash(ctx context.Context, hash common.Hash) (*types.Blo
//
// Note that loading full blocks requires two requests. Use HeaderByNumber
// if you don't need all transactions or uncle headers.
-func (ec *Client) BlockByNumber(ctx context.Context, number *big.Int) (*types.Block, error) {
+func (ec *Client) BlockByNumber(ctx context.Context, number *big.Int) (*types.Block, error) {
    return ec.getBlock(ctx, "eth_getBlockByNumber", ToBlockNumArg(number), true)
}

// BlockNumber returns the most recent block number
-func (ec *Client) BlockNumber(ctx context.Context) (uint64, error) {
+func (ec *Client) BlockNumber(ctx context.Context) (uint64, error) {
    var result hexutil.Uint64
    err := ec.c.CallContext(ctx, &result, "eth_blockNumber")
    return uint64(result), err
@@ -134,7 +172,7 @@ type rpcBlock struct {
    BlockExtraData *hexutil.Bytes `json:"blockExtraData"`
}

-func (ec *Client) getBlock(ctx context.Context, method string, args ...interface{}) (*types.Block, error) {
+func (ec *Client) getBlock(ctx context.Context, method string, args ...interface{}) (*types.Block, error) {
    var raw json.RawMessage
    err := ec.c.CallContext(ctx, &raw, method, args...)
    if err != nil {
@@ -200,7 +238,7 @@ func (ec *Client) getBlock(ctx context.Context, method string, args ...interface
}

// HeaderByHash returns the block header with the given hash.
-func (ec *Client) HeaderByHash(ctx context.Context, hash common.Hash) (*types.Header, error) {
+func (ec *Client) HeaderByHash(ctx context.Context, hash common.Hash) (*types.Header, error) {
    var head *types.Header
    err := ec.c.CallContext(ctx, &head, "eth_getBlockByHash", hash, false)
    if err == nil && head == nil {
@@ -211,7 +249,7 @@ func (ec *Client) HeaderByHash(ctx context.Context, hash common.Hash) (*types.He
//
// HeaderByNumber returns a block header from the current canonical chain. If number is
// nil, the latest known header is returned.
-func (ec *Client) HeaderByNumber(ctx context.Context, number *big.Int) (*types.Header, error) {
+func (ec *Client) HeaderByNumber(ctx context.Context, number *big.Int) (*types.Header, error) {
    var head *types.Header
    err := ec.c.CallContext(ctx, &head, "eth_getBlockByNumber", ToBlockNumArg(number), false)
    if err == nil && head == nil {
@@ -239,7 +277,7 @@ func (tx *rpcTransaction) UnmarshalJSON(msg []byte) error {
}

// TransactionByHash returns the transaction with the given hash.
-func (ec *Client) TransactionByHash(ctx context.Context, hash common.Hash) (tx *types.Transaction, isPending bool, err error) {
+func (ec *Client) TransactionByHash(ctx context.Context, hash common.Hash) (tx *types.Transaction, isPending bool, err error) {
    var json *rpcTransaction
    err = ec.c.CallContext(ctx, &json, "eth_getTransactionByHash", hash)
    if err != nil {
@@ -261,12 +299,14 @@ func (ec *Client) TransactionByHash(ctx context.Context, hash common.Hash) (tx *
//
// There is a fast-path for transactions retrieved by TransactionByHash and
// TransactionInBlock. Getting their sender address can be done without an RPC interaction.
-func (ec *Client) TransactionSender(ctx context.Context, tx *types.Transaction, block common.Hash, index uint) (common.Address, error) {
+func (ec *Client) TransactionSender(ctx context.Context, tx *types.Transaction, block common.Hash, index uint) (common.Address, error) {
    // Try to load the address from the cache.
    sender, err := types.Sender(&senderFromServer{blockhash: block}, tx)
    if err == nil {
        return sender, nil
    }
+
+    // It was not found in cache, ask the server.
    var meta struct {
        Hash common.Hash
        From common.Address
@@ -281,14 +321,14 @@ func (ec *Client) TransactionSender(ctx context.Context, tx *types.Transaction,
}

// TransactionCount returns the total number of transactions in the given block.
-func (ec *Client) TransactionCount(ctx context.Context, blockHash common.Hash) (uint, error) {
+func (ec *Client) TransactionCount(ctx context.Context, blockHash common.Hash) (uint, error) {
    var num hexutil.Uint
    err := ec.c.CallContext(ctx, &num, "eth_getBlockTransactionCountByHash", blockHash)
    return uint(num), err
}

// TransactionInBlock returns a single transaction at index in the given block.
-func (ec *Client) TransactionInBlock(ctx context.Context, blockHash common.Hash, index uint) (*types.Transaction, error) {
+func (ec *Client) TransactionInBlock(ctx context.Context, blockHash common.Hash, index uint) (*types.Transaction, error) {
    var json *rpcTransaction
    err := ec.c.CallContext(ctx, &json, "eth_getTransactionByBlockHashAndIndex", blockHash, hexutil.Uint64(index))
    if err != nil {
@@ -307,7 +347,7 @@ func (ec *Client) TransactionInBlock(ctx context.Context, blockHash common.Hash,
//
// TransactionReceipt returns the receipt of a transaction by transaction hash.
// Note that the receipt is not available for pending transactions.
-func (ec *Client) TransactionReceipt(ctx context.Context, txHash common.Hash) (*types.Receipt, error) {
+func (ec *Client) TransactionReceipt(ctx context.Context, txHash common.Hash) (*types.Receipt, error) {
    var r *types.Receipt
    err := ec.c.CallContext(ctx, &r, "eth_getTransactionReceipt", txHash)
    if err == nil {
@@ -319,25 +359,25 @@ func (ec *Client) TransactionReceipt(ctx context.Context, txHash common.Hash) (*
}

// SubscribeNewAcceptedTransactions subscribes to notifications about the accepted transaction hashes on the given channel.
-func (ec *Client) SubscribeNewAcceptedTransactions(ctx context.Context, ch chan<- *common.Hash) (interfaces.Subscription, error) {
+func (ec *Client) SubscribeNewAcceptedTransactions(ctx context.Context, ch chan<- *common.Hash) (interfaces.Subscription, error) {
    return ec.c.EthSubscribe(ctx, ch, "newAcceptedTransactions")
}

// SubscribeNewAcceptedTransactions subscribes to notifications about the accepted transaction hashes on the given channel.
-func (ec *Client) SubscribeNewPendingTransactions(ctx context.Context, ch chan<- *common.Hash) (interfaces.Subscription, error) {
+func (ec *Client) SubscribeNewPendingTransactions(ctx context.Context, ch chan<- *common.Hash) (interfaces.Subscription, error) {
    return ec.c.EthSubscribe(ctx, ch, "newPendingTransactions")
}

// SubscribeNewHead subscribes to notifications about the current blockchain head
// on the given channel.
-func (ec *Client) SubscribeNewHead(ctx context.Context, ch chan<- *types.Header) (interfaces.Subscription, error) {
+func (ec *Client) SubscribeNewHead(ctx context.Context, ch chan<- *types.Header) (interfaces.Subscription, error) {
    return ec.c.EthSubscribe(ctx, ch, "newHeads")
}

```

```

}

// State Access

// NetworkID returns the network ID (also known as the chain ID) for this chain.
-func (ec *Client) NetworkID(ctx context.Context) (*big.Int, error) {
+func (ec *Client) NetworkID(ctx context.Context) (*big.Int, error) {
    version := new(big.Int)
    var ver string
    if err := ec.c.CallContext(ctx, &ver, "net_version"); err != nil {
@@ -351,7 +391,7 @@ func (ec *Client) NetworkID(ctx context.Context) (*big.Int, error) {

// BalanceAt returns the wei balance of the given account.
// The block number can be nil, in which case the balance is taken from the latest known block.
-func (ec *Client) BalanceAt(ctx context.Context, account common.Address, blockNumber *big.Int) (*big.Int, error) {
+func (ec *Client) BalanceAt(ctx context.Context, account common.Address, blockNumber *big.Int) (*big.Int, error) {
    var result hexutil.Big
    err := ec.c.CallContext(ctx, &result, "eth_getBalance", account, ToBlockNumArg(blockNumber))
    return (*big.Int)(&result), err
@@ -359,7 +399,7 @@ func (ec *Client) BalanceAt(ctx context.Context, account common.Address, blockNu

// AssetBalanceAt returns the [assetID] balance of the given account
// The block number can be nil, in which case the balance is taken from the latest known block.
-func (ec *Client) AssetBalanceAt(ctx context.Context, account common.Address, assetID ids.ID, blockNumber *big.Int) (*big.Int, error) {
+func (ec *Client) AssetBalanceAt(ctx context.Context, account common.Address, assetID ids.ID, blockNumber *big.Int) (*big.Int, error) {
    var result hexutil.Big
    err := ec.c.CallContext(ctx, &result, "eth_getAssetBalance", account, ToBlockNumArg(blockNumber), assetID)
    return (*big.Int)(&result), err
@@ -367,7 +407,7 @@ func (ec *Client) AssetBalanceAt(ctx context.Context, account common.Address, as

// StorageAt returns the value of key in the contract storage of the given account.
// The block number can be nil, in which case the value is taken from the latest known block.
-func (ec *Client) StorageAt(ctx context.Context, account common.Address, key common.Hash, blockNumber *big.Int) ([]byte, error) {
+func (ec *Client) StorageAt(ctx context.Context, account common.Address, key common.Hash, blockNumber *big.Int) ([]byte, error) {
    var result hexutil.Bytes
    err := ec.c.CallContext(ctx, &result, "eth_getStorageAt", account, key, ToBlockNumArg(blockNumber))
    return result, err
@@ -375,7 +415,7 @@ func (ec *Client) StorageAt(ctx context.Context, account common.Address, key com

// CodeAt returns the contract code of the given account.
// The block number can be nil, in which case the code is taken from the latest known block.
-func (ec *Client) CodeAt(ctx context.Context, account common.Address, blockNumber *big.Int) ([]byte, error) {
+func (ec *Client) CodeAt(ctx context.Context, account common.Address, blockNumber *big.Int) ([]byte, error) {
    var result hexutil.Bytes
    err := ec.c.CallContext(ctx, &result, "eth_getCode", account, ToBlockNumArg(blockNumber))
    return result, err
@@ -383,7 +423,7 @@ func (ec *Client) CodeAt(ctx context.Context, account common.Address, blockNumbe

// NonceAt returns the account nonce of the given account.
// The block number can be nil, in which case the nonce is taken from the latest known block.
-func (ec *Client) NonceAt(ctx context.Context, account common.Address, blockNumber *big.Int) (uint64, error) {
+func (ec *Client) NonceAt(ctx context.Context, account common.Address, blockNumber *big.Int) (uint64, error) {
    var result hexutil.Uint64
    err := ec.c.CallContext(ctx, &result, "eth_getTransactionCount", account, ToBlockNumArg(blockNumber))
    return uint64(result), err
@@ -392,7 +432,7 @@ func (ec *Client) NonceAt(ctx context.Context, account common.Address, blockNumb
// Filters

// FilterLogs executes a filter query.
-func (ec *Client) FilterLogs(ctx context.Context, q interfaces.FilterQuery) ([]types.Log, error) {
+func (ec *Client) FilterLogs(ctx context.Context, q interfaces.FilterQuery) ([]types.Log, error) {
    var result []types.Log
    arg, err := toFilterArg(q)
    if err != nil {
@@ -400,7 +440,7 @@ func (ec *Client) FilterLogs(ctx context.Context, q interfaces.FilterQuery) ([]t
    }

// SubscribeFilterLogs subscribes to the results of a streaming filter query.
-func (ec *Client) SubscribeFilterLogs(ctx context.Context, q interfaces.FilterQuery, ch chan<- types.Log) (interfaces.Subscription, error) {
+func (ec *Client) SubscribeFilterLogs(ctx context.Context, q interfaces.FilterQuery, ch chan<- types.Log) (interfaces.Subscription, error) {
    arg, err := toFilterArg(q)
    if err != nil {
        return nil, err
@@ -433,19 +473,19 @@ func toFilterArg(q interfaces.FilterQuery) (interface{}, error) {
    }

// AcceptedCodeAt returns the contract code of the given account in the accepted state.
-func (ec *Client) AcceptedCodeAt(ctx context.Context, account common.Address) ([]byte, error) {
+func (ec *Client) AcceptedCodeAt(ctx context.Context, account common.Address) ([]byte, error) {
    return ec.CodeAt(ctx, account, nil)
}

// AcceptedNonceAt returns the account nonce of the given account in the accepted state.
// This is the nonce that should be used for the next transaction.
-func (ec *Client) AcceptedNonceAt(ctx context.Context, account common.Address) (uint64, error) {
+func (ec *Client) AcceptedNonceAt(ctx context.Context, account common.Address) (uint64, error) {
    return ec.NonceAt(ctx, account, nil)
}

// AcceptedCallContract executes a message call transaction in the accepted
// state.
-func (ec *Client) AcceptedCallContract(ctx context.Context, msg interfaces.CallMsg) ([]byte, error) {
+func (ec *Client) AcceptedCallContract(ctx context.Context, msg interfaces.CallMsg) ([]byte, error) {
    return ec.CallContract(ctx, msg, nil)
}

@@ -457,7 +497,7 @@ func (ec *Client) AcceptedCallContract(ctx context.Context, msg interfaces.CallM
// blockNumber selects the block height at which the call runs. It can be nil, in which
// case the code is taken from the latest known block. Note that state from very old
// blocks might not be available.
-func (ec *Client) CallContract(ctx context.Context, msg interfaces.CallMsg, blockNumber *big.Int) ([]byte, error) {
+func (ec *Client) CallContract(ctx context.Context, msg interfaces.CallMsg, blockNumber *big.Int) ([]byte, error) {
    var hex hexutil.Bytes
    err := ec.c.CallContext(ctx, &hex, "eth_call", toCallArg(msg), ToBlockNumArg(blockNumber))
    if err != nil {
@@ -468,7 +508,7 @@ func (ec *Client) CallContract(ctx context.Context, msg interfaces.CallMsg, bloc

// SuggestGasPrice retrieves the currently suggested gas price to allow a timely
// execution of a transaction.
-func (ec *Client) SuggestGasPrice(ctx context.Context) (*big.Int, error) {
+func (ec *Client) SuggestGasPrice(ctx context.Context) (*big.Int, error) {
    var hex hexutil.Big
    if err := ec.c.CallContext(ctx, &hex, "eth_gasPrice"); err != nil {
        return nil, err
@@ -478,7 +518,7 @@ func (ec *Client) SuggestGasPrice(ctx context.Context) (*big.Int, error) {

// SuggestGasTipCap retrieves the currently suggested gas tip cap after 1559 to
// allow a timely execution of a transaction.
-func (ec *Client) SuggestGasTipCap(ctx context.Context) (*big.Int, error) {
+func (ec *Client) SuggestGasTipCap(ctx context.Context) (*big.Int, error) {
    var hex hexutil.Big
    if err := ec.c.CallContext(ctx, &hex, "eth_maxPriorityFeePerGas"); err != nil {
        return nil, err
@@ -490,7 +530,7 @@ func (ec *Client) SuggestGasTipCap(ctx context.Context) (*big.Int, error) {
// the current pending state of the backend blockchain. There is no guarantee that this is
// the true gas limit requirement as other transactions may be added or removed by miners,
// but it should provide a basis for setting a reasonable default.
-func (ec *Client) EstimateGas(ctx context.Context, msg interfaces.CallMsg) (uint64, error) {
+func (ec *Client) EstimateGas(ctx context.Context, msg interfaces.CallMsg) (uint64, error) {
    var hex hexutil.Uint64
    err := ec.c.CallContext(ctx, &hex, "eth_estimateGas", toCallArg(msg))
    if err != nil {
@@ -502,7 +542,7 @@ func (ec *Client) EstimateGas(ctx context.Context, msg interfaces.CallMsg) (uint
// EstimateBaseFee tries to estimate the base fee for the next block if it were created

```

```

// immediately. There is no guarantee that this will be the base fee used in the next block
// or that the next base fee will be higher or lower than the returned value.
-func (ec *Client) EstimateBaseFee(ctx context.Context) (*big.Int, error) {
+func (ec *client) EstimateBaseFee(ctx context.Context) (*big.Int, error) {
    var hex hexutil.Big
    err := ec.c.CallContext(ctx, &hex, "eth_baseFee")
    if err != nil {
@@ -515,7 +555,7 @@ func (ec *Client) EstimateBaseFee(ctx context.Context) (*big.Int, error) {
    //
    // If the transaction was a contract creation use the TransactionReceipt method to get the
    // contract address after the transaction has been mined.
-func (ec *Client) SendTransaction(ctx context.Context, tx *types.Transaction) error {
+func (ec *client) SendTransaction(ctx context.Context, tx *types.Transaction) error {
    data, err := tx.MarshalBinary()
    if err != nil {
        return err
diff --git a/ethclient/signer.go b/ethclient/signer.go
index dafa943b..43baf1bc 100644
--- a/ethclient/signer.go
+++ b/ethclient/signer.go
@@ -30,8 +30,8 @@ import (
    "errors"
    "math/big"

-    "github.com/ava-labs/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

// senderFromServer is a types.Signer that remembers the sender address returned by the RPC
@@ -55,7 +55,7 @@ func (s *senderFromServer) Equal(other types.Signer) bool {
}

func (s *senderFromServer) Sender(tx *types.Transaction) (common.Address, error) {
-    if s.blockhash == (common.Hash{}) {
+    if s.addr == (common.Address{}) {
        return common.Address{}, errNotCached
    }
    return s.addr, nil
diff --git a/ethdb/dbtest/testsuite.go b/ethdb/dbtest/testsuite.go
index 90c92ee3..2209bf79 100644
--- a/ethdb/dbtest/testsuite.go
+++ b/ethdb/dbtest/testsuite.go
@@ -32,7 +32,7 @@ import (
    "sort"
    "testing"

-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

// TestDatabaseSuite runs a suite of tests against a KeyValueStore database
diff --git a/ethdb/leveldb/leveldb.go b/ethdb/leveldb/leveldb.go
new file mode 100644
index 00000000..1bb02f19
--- /dev/null
+++ b/ethdb/leveldb/leveldb.go
@@ -0,0 +1,531 @@
+// (c) 2021-2022, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+//
+// *****
+// Copyright 2018 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+//go:build !js
+// +build !js
+
+// Package leveldb implements the key-value database layer based on LevelDB.
+package leveldb
+
+import (
+    "fmt"
+    "strconv"
+    "strings"
+    "sync"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/metrics"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/syndtr/goleveldb/leveldb"
+    "github.com/syndtr/goleveldb/leveldb/errors"
+    "github.com/syndtr/goleveldb/leveldb/filter"
+    "github.com/syndtr/goleveldb/leveldb/opt"
+    "github.com/syndtr/goleveldb/leveldb/util"
+)
+
+const (
+    // degradationWarnInterval specifies how often warning should be printed if the
+    // leveldb database cannot keep up with requested writes.
+    degradationWarnInterval = time.Minute
+
+    // minCache is the minimum amount of memory in megabytes to allocate to leveldb
+    // read and write caching, split half and half.
+    minCache = 16
+
+    // minHandles is the minimum number of files handles to allocate to the open
+    // database files.
+    minHandles = 16
+
+    // metricsGatheringInterval specifies the interval to retrieve leveldb database
+    // compaction, io and pause stats to report to the user.
+    metricsGatheringInterval = 3 * time.Second
+)
+
+// Database is a persistent key-value store. Apart from basic data storage
+// functionality it also supports batch writes and iterating over the keyspace in
+// binary-alphabetical order.
+type Database struct {
+    fn string // filename for reporting
+    db *leveldb.DB // LevelDB instance
+
+    compTimeMeter    metrics.Meter // Meter for measuring the total time spent in database compaction

```

```

+     compReadMeter      metrics.Meter // Meter for measuring the data read during compaction
+     compWriteMeter     metrics.Meter // Meter for measuring the data written during compaction
+     writeDelayMeter    metrics.Meter // Meter for measuring the write delay number due to database compaction
+     writeDelayMeter    metrics.Meter // Meter for measuring the write delay duration due to database compaction
+     diskSizeGauge      metrics.Gauge // Gauge for tracking the size of all the levels in the database
+     diskReadMeter      metrics.Meter // Meter for measuring the effective amount of data read
+     diskWriteMeter     metrics.Meter // Meter for measuring the effective amount of data written
+     memCompGauge       metrics.Gauge // Gauge for tracking the number of memory compaction
+     level0CompGauge    metrics.Gauge // Gauge for tracking the number of table compaction in level0
+     nonlevel0CompGauge metrics.Gauge // Gauge for tracking the number of table compaction in non0 level
+     seekCompGauge      metrics.Gauge // Gauge for tracking the number of table compaction caused by read opt
+
+
+     quitLock sync.Mutex // Mutex protecting the quit channel access
+     quitChan chan chan error // Quit channel to stop the metrics collection before closing the database
+
+     log log.Logger // Contextual logger tracking the database path
+}
+
+// New returns a wrapped LevelDB object. The namespace is the prefix that the
+// metrics reporting should use for surfacing internal stats.
+func New(file string, cache int, handles int, namespace string, readonly bool) (*Database, error) {
+    return NewCustom(file, namespace, func(options *opt.Options) {
+        // Ensure we have some minimal caching and file guarantees
+        if cache < minCache {
+            cache = minCache
+        }
+        if handles < minHandles {
+            handles = minHandles
+        }
+        // Set default options
+        options.OpenFilesCacheCapacity = handles
+        options.BlockCacheCapacity = cache / 2 * opt.MiB
+        options.WriteBuffer = cache / 4 * opt.MiB // Two of these are used internally
+        if readonly {
+            options.ReadOnly = true
+        }
+    })
+}
+
+// NewCustom returns a wrapped LevelDB object. The namespace is the prefix that the
+// metrics reporting should use for surfacing internal stats.
+// The customize function allows the caller to modify the leveldb options.
+func NewCustom(file string, namespace string, customize func(options *opt.Options)) (*Database, error) {
+    options := configureOptions(customize)
+    logger := log.New("database", file)
+    usedCache := options.GetBlockCacheCapacity() + options.GetWriteBuffer()*2
+    logCtx := []interface{}{"cache", common.StorageSize(usedCache), "handles", options.GetOpenFilesCacheCapacity()}
+    if options.ReadOnly {
+        logCtx = append(logCtx, "readonly", "true")
+    }
+    logger.Info("Allocated cache and file handles", logCtx...)
+
+    // Open the db and recover any potential corruptions
+    db, err := leveldb.OpenFile(file, options)
+    if _, corrupted := err.(*errors.ErrCorrupted); corrupted {
+        db, err = leveldb.RecoverFile(file, nil)
+    }
+    if err != nil {
+        return nil, err
+    }
+    // Assemble the wrapper with all the registered metrics
+    ldb := &Database{
+        fn:      file,
+        db:      db,
+        log:      logger,
+        quitChan: make(chan chan error),
+    }
+    ldb.compTimeMeter = metrics.NewRegisteredMeter(namespace+"compact/time", nil)
+    ldb.compReadMeter = metrics.NewRegisteredMeter(namespace+"compact/input", nil)
+    ldb.compWriteMeter = metrics.NewRegisteredMeter(namespace+"compact/output", nil)
+    ldb.diskSizeGauge = metrics.NewRegisteredGauge(namespace+"disk/size", nil)
+    ldb.diskReadMeter = metrics.NewRegisteredMeter(namespace+"disk/read", nil)
+    ldb.diskWriteMeter = metrics.NewRegisteredMeter(namespace+"disk/write", nil)
+    ldb.writeDelayMeter = metrics.NewRegisteredMeter(namespace+"compact/writedelay/duration", nil)
+    ldb.writeDelayNMeter = metrics.NewRegisteredMeter(namespace+"compact/writedelay/counter", nil)
+    ldb.memCompGauge = metrics.NewRegisteredGauge(namespace+"compact/memory", nil)
+    ldb.level0CompGauge = metrics.NewRegisteredGauge(namespace+"compact/level0", nil)
+    ldb.nonlevel0CompGauge = metrics.NewRegisteredGauge(namespace+"compact/nonlevel0", nil)
+    ldb.seekCompGauge = metrics.NewRegisteredGauge(namespace+"compact/seek", nil)
+
+    // Start up the metrics gathering and return
+    go ldb.meter(metricsGatheringInterval)
+    return ldb, nil
+}
+
+// configureOptions sets some default options, then runs the provided setter.
+func configureOptions(customizeFn func(*opt.Options)) *opt.Options {
+    // Set default options
+    options := &opt.Options{
+        Filter:          filter.NewBloomFilter(10),
+        DisableSeeksCompaction: true,
+    }
+    // Allow caller to make custom modifications to the options
+    if customizeFn != nil {
+        customizeFn(options)
+    }
+    return options
+}
+
+// Close stops the metrics collection, flushes any pending data to disk and closes
+// all io accesses to the underlying key-value store.
+func (db *Database) Close() error {
+    db.quitLock.Lock()
+    defer db.quitLock.Unlock()
+
+    if db.quitChan != nil {
+        errc := make(chan error)
+        db.quitChan <- errc
+        if err := <-errc; err != nil {
+            db.log.Error("Metrics collection failed", "err", err)
+        }
+        db.quitChan = nil
+    }
+    return db.db.Close()
+}
+
+// Has retrieves if a key is present in the key-value store.
+func (db *Database) Has(key []byte) (bool, error) {
+    return db.db.Has(key, nil)
+}
+
+// Get retrieves the given key if it's present in the key-value store.
+func (db *Database) Get(key []byte) ([]byte, error) {
+    dat, err := db.db.Get(key, nil)
+    if err != nil {
+        return nil, err
+    }
+    return dat, nil
+}
+
+// Put inserts the given value into the key-value store.
+func (db *Database) Put(key []byte, value []byte) error {
+    return db.db.Put(key, value, nil)
+}

```

```

+
+// Delete removes the key from the key-value store.
+func (db *Database) Delete(key []byte) error {
+    return db.db.Delete(key, nil)
+}
+
+// NewBatch creates a write-only key-value store that buffers changes to its host
+// database until a final write is called.
+func (db *Database) NewBatch() ethdb.Batch {
+    return &batch{
+        db: db.db,
+        b:  new(leveldb.Batch),
+    }
+}
+
+// NewIterator creates a binary-alphabetical iterator over a subset
+// of database content with a particular key prefix, starting at a particular
+// initial key (or after, if it does not exist).
+func (db *Database) NewIterator(prefix []byte, start []byte) ethdb.Iterator {
+    return db.db.NewIterator(bytesPrefixRange(prefix, start), nil)
+}
+
+// Stat returns a particular internal stat of the database.
+func (db *Database) Stat(property string) (string, error) {
+    return db.db.GetProperty(property)
+}
+
+// Compact flattens the underlying data store for the given key range. In essence,
+// deleted and overwritten versions are discarded, and the data is rearranged to
+// reduce the cost of operations needed to access them.
+//
+// A nil start is treated as a key before all keys in the data store; a nil limit
+// is treated as a key after all keys in the data store. If both is nil then it
+// will compact entire data store.
+func (db *Database) Compact(start []byte, limit []byte) error {
+    return db.db.CompactRange(util.Range{Start: start, Limit: limit})
+}
+
+// Path returns the path to the database directory.
+func (db *Database) Path() string {
+    return db.fn
+}
+
+// meter periodically retrieves internal leveldb counters and reports them to
+// the metrics subsystem.
+//
+// This is how a LevelDB stats table looks like (currently):
+// Compactions
+// Level | Tables | Size(MB) | Time(sec) | Read(MB) | Write(MB)
+// -----|-----|-----|-----|-----|-----
+// 0      |      0 |    0.00000 |    1.27969 |    0.00000 |    12.31098
+// 1      |     85 |   109.27913 |   28.09293 |   213.92493 |   214.26294
+// 2      |    523 |  1000.37159 |   7.26059  |   66.86342  |   66.77884
+// 3      |    570 |  1113.18458 |   0.00000  |   0.00000  |   0.00000
+//
+// This is how the write delay look like (currently):
+// DelayN:5 Delay:406.604657ms Paused: false
+//
+// This is how the iostats look like (currently):
+// Read(MB):3895.04860 Write(MB):3654.64712
+func (db *Database) meter(refresh time.Duration) {
+    // Create the counters to store current and previous compaction values
+    compactions := make([][]float64, 2)
+    for i := 0; i < 2; i++ {
+        compactions[i] = make([]float64, 4)
+    }
+    // Create storage for iostats.
+    var iostats [2]float64
+
+    // Create storage and warning log tracer for write delay.
+    var (
+        delaystats [2]int64
+        lastWritePaused time.Time
+    )
+
+    var (
+        errc chan error
+        merr error
+    )
+
+    timer := time.NewTimer(refresh)
+    defer timer.Stop()
+
+    // Iterate ad infinitum and collect the stats
+    for i := 1; errc == nil && merr == nil; i++ {
+        // Retrieve the database stats
+        stats, err := db.db.GetProperty("leveldb.stats")
+        if err != nil {
+            db.log.Error("Failed to read database stats", "err", err)
+            merr = err
+            continue
+        }
+        // Find the compaction table, skip the header
+        lines := strings.Split(stats, "\n")
+        for len(lines) > 0 && strings.TrimSpace(lines[0]) != "Compactions" {
+            lines = lines[1:]
+        }
+        if len(lines) <= 3 {
+            db.log.Error("Compaction leveldbTable not found")
+            merr = errors.New("compaction leveldbTable not found")
+            continue
+        }
+        lines = lines[3:]
+
+        // Iterate over all the leveldbTable rows, and accumulate the entries
+        for j := 0; j < len(compactions[i%2]); j++ {
+            compactions[i%2][j] = 0
+        }
+        for _, line := range lines {
+            parts := strings.Split(line, "|")
+            if len(parts) != 6 {
+                break
+            }
+            for idx, counter := range parts[2:] {
+                value, err := strconv.ParseFloat(strings.TrimSpace(counter), 64)
+                if err != nil {
+                    db.log.Error("Compaction entry parsing failed", "err", err)
+                    merr = err
+                    continue
+                }
+                compactions[i%2][idx] += value
+            }
+        }
+        // Update all the requested meters
+        if db.diskSizeGauge != nil {
+            db.diskSizeGauge.Update(int64(compactions[i%2][0] * 1024 * 1024))
+        }
+        if db.compTimeMeter != nil {
+            db.compTimeMeter.Mark(int64((compactions[i%2][1] - compactions[(i-1)%2][1]) * 1000 * 1000 * 1000))
+        }
+        if db.compReadMeter != nil {
+            db.compReadMeter.Mark(int64((compactions[i%2][2] - compactions[(i-1)%2][2]) * 1024 * 1024))
+        }
+        if db.compWriteMeter != nil {

```

```

+         db.compWriteMeter.Mark(int64((compactions[i%2][3] - compactions[(i-1)%2][3]) * 1024 * 1024))
+     }
+     // Retrieve the write delay statistic
+     writedelay, err := db.db.GetProperty("leveldb.writedelay")
+     if err != nil {
+         db.log.Error("Failed to read database write delay statistic", "err", err)
+         merr = err
+         continue
+     }
+     var (
+         delayN      int64
+         delayDuration string
+         duration     time.Duration
+         paused       bool
+     )
+     if n, err := fmt.Sscanf(writedelay, "DelayN:%d Delay:%s Paused:%t", &delayN, &delayDuration, &paused); n != 3 || err != nil {
+         db.log.Error("Write delay statistic not found")
+         merr = err
+         continue
+     }
+     duration, err = time.ParseDuration(delayDuration)
+     if err != nil {
+         db.log.Error("Failed to parse delay duration", "err", err)
+         merr = err
+         continue
+     }
+     if db.writeDelayNMeter != nil {
+         db.writeDelayNMeter.Mark(delayN - delaystats[0])
+     }
+     if db.writeDelayMeter != nil {
+         db.writeDelayMeter.Mark(duration.Nanoseconds() - delaystats[1])
+     }
+     // If a warning that db is performing compaction has been displayed, any subsequent
+     // warnings will be withheld for one minute not to overwhelm the user.
+     if paused && delayN-delaystats[0] == 0 && duration.Nanoseconds()-delaystats[1] == 0 &&
+         time.Now().After(lastWritePaused.Add(degradationWarnInterval)) {
+         db.log.Warn("Database compacting, degraded performance")
+         lastWritePaused = time.Now()
+     }
+     delaystats[0], delaystats[1] = delayN, duration.Nanoseconds()
+
+     // Retrieve the database iostats.
+     ioStats, err := db.db.GetProperty("leveldb.iostats")
+     if err != nil {
+         db.log.Error("Failed to read database iostats", "err", err)
+         merr = err
+         continue
+     }
+     var nRead, nWrite float64
+     parts := strings.Split(ioStats, " ")
+     if len(parts) < 2 {
+         db.log.Error("Bad syntax of ioStats", "ioStats", ioStats)
+         merr = fmt.Errorf("bad syntax of ioStats %s", ioStats)
+         continue
+     }
+     if n, err := fmt.Sscanf(parts[0], "Read(MB):%f", &nRead); n != 1 || err != nil {
+         db.log.Error("Bad syntax of read entry", "entry", parts[0])
+         merr = err
+         continue
+     }
+     if n, err := fmt.Sscanf(parts[1], "Write(MB):%f", &nWrite); n != 1 || err != nil {
+         db.log.Error("Bad syntax of write entry", "entry", parts[1])
+         merr = err
+         continue
+     }
+     if db.diskReadMeter != nil {
+         db.diskReadMeter.Mark(int64((nRead - iostats[0]) * 1024 * 1024))
+     }
+     if db.diskWriteMeter != nil {
+         db.diskWriteMeter.Mark(int64((nWrite - iostats[1]) * 1024 * 1024))
+     }
+     iostats[0], iostats[1] = nRead, nWrite
+
+     compCount, err := db.db.GetProperty("leveldb.compcount")
+     if err != nil {
+         db.log.Error("Failed to read database iostats", "err", err)
+         merr = err
+         continue
+     }
+
+     var (
+         memComp      uint32
+         level0Comp   uint32
+         nonLevel0Comp uint32
+         seekComp     uint32
+     )
+     if n, err := fmt.Sscanf(compCount, "MemComp:%d Level0Comp:%d NonLevel0Comp:%d SeekComp:%d", &memComp, &level0Comp, &nonLevel0Comp, &seekComp); n != 4 || err != nil {
+         db.log.Error("Compaction count statistic not found")
+         merr = err
+         continue
+     }
+     db.memCompGauge.Update(int64(memComp))
+     db.level0CompGauge.Update(int64(level0Comp))
+     db.nonLevel0CompGauge.Update(int64(nonLevel0Comp))
+     db.seekCompGauge.Update(int64(seekComp))
+
+     // Sleep a bit, then repeat the stats collection
+     select {
+     case errc = <-db.quitChan:
+         // Quit requesting, stop hammering the database
+     case <-timer.C:
+         timer.Reset(refresh)
+         // Timeout, gather a new set of stats
+     }
+
+     if errc == nil {
+         errc = <-db.quitChan
+     }
+     errc <- merr
+ }
+
+ // batch is a write-only leveldb batch that commits changes to its host database
+ // when Write is called. A batch cannot be used concurrently.
+ type batch struct {
+     db      *leveldb.DB
+     b       *leveldb.Batch
+     size int
+ }
+
+ // Put inserts the given value into the batch for later committing.
+ func (b *batch) Put(key, value []byte) error {
+     b.b.Put(key, value)
+     b.size += len(value)
+     return nil
+ }
+
+ // Delete inserts the a key removal into the batch for later committing.
+ func (b *batch) Delete(key []byte) error {
+     b.b.Delete(key)
+     b.size += len(key)
+     return nil
+ }
+
+

```

```

// ValueSize retrieves the amount of data queued up for writing.
+func (b *batch) ValueSize() int {
+    return b.size
+}
+
+// Write flushes any accumulated data to disk.
+func (b *batch) Write() error {
+    return b.db.Write(b.b, nil)
+}
+
+// Reset resets the batch for reuse.
+func (b *batch) Reset() {
+    b.b.Reset()
+    b.size = 0
+}
+
+// Replay replays the batch contents.
+func (b *batch) Replay(w ethdb.KeyValueWriter) error {
+    return b.b.Replay(&replayer{writer: w})
+}
+
+// replayer is a small wrapper to implement the correct replay methods.
+type replayer struct {
+    writer ethdb.KeyValueWriter
+    failure error
+}
+
+// Put inserts the given value into the key-value data store.
+func (r *replayer) Put(key, value []byte) {
+    // If the replay already failed, stop executing ops
+    if r.failure != nil {
+        return
+    }
+    r.failure = r.writer.Put(key, value)
+}
+
+// Delete removes the key from the key-value data store.
+func (r *replayer) Delete(key []byte) {
+    // If the replay already failed, stop executing ops
+    if r.failure != nil {
+        return
+    }
+    r.failure = r.writer.Delete(key)
+}
+
+// bytesPrefixRange returns key range that satisfy
+// - the given prefix, and
+// - the given seek position
+func bytesPrefixRange(prefix, start []byte) *util.Range {
+    r := util.BytesPrefix(prefix)
+    r.Start = append(r.Start, start...)
+    return r
+}
diff --git a/ethdb/leveldb/leveldb_test.go b/ethdb/leveldb/leveldb_test.go
new file mode 100644
index 00000000..fff2a97d
--- /dev/null
+++ b/ethdb/leveldb/leveldb_test.go
@@ -0,0 +1,50 @@
+// (c) 2021-2022, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2019 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package leveldb
+
+import (
+    "testing"
+
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb/dbtest"
+    "github.com/syndtr/goleveldb/leveldb"
+    "github.com/syndtr/goleveldb/leveldb/storage"
+)
+
+func TestLevelDB(t *testing.T) {
+    t.Run("DatabaseSuite", func(t *testing.T) {
+        dbtest.TestDatabaseSuite(t, func() ethdb.KeyValueStore {
+            db, err := leveldb.Open(storage.NewMemStorage(), nil)
+            if err != nil {
+                t.Fatal(err)
+            }
+            return &Database{
+                db: db,
+            }
+        })
+    })
+}
diff --git a/ethdb/memorydb/memorydb.go b/ethdb/memorydb/memorydb.go
index 769581ff..6b56519d 100644
--- a/ethdb/memorydb/memorydb.go
+++ b/ethdb/memorydb/memorydb.go
@@ -33,8 +33,8 @@ import (
     "strings"
     "sync"
 
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/ethdb"
 )
 
 var (
diff --git a/ethdb/memorydb/memorydb_test.go b/ethdb/memorydb/memorydb_test.go
index 34361e9f..92df1f67 100644
--- a/ethdb/memorydb/memorydb_test.go
+++ b/ethdb/memorydb/memorydb_test.go
@@ -29,8 +29,8 @@ package memorydb
 import (
     "testing"
 
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/ethdb/dbtest"
+    "github.com/flare-foundation/coreth/ethdb"
 )

```

```
+      "github.com/flare-foundation/coreth/ethdb/dbtest"
+    )

    func TestMemoryDB(t *testing.T) {
diff --git a/go.mod b/go.mod
index f607b56d..1262c316 100644
--- a/go.mod
+++ b/go.mod
@@ -1,28 +1,30 @@
-module github.com/ava-labs/coreth
+module github.com/flare-foundation/coreth

go 1.16

require (
    github.com/VictoriaMetrics/fastcache v1.6.0
-    github.com/ava-labs/avalanchego v1.6.4
    github.com/btcsuite/btcd v0.21.0-beta // indirect
    github.com/cespare/cp v0.1.0
    github.com/davecgh/go-spew v1.1.1
    github.com/deckarep/golang-set v1.7.1
    github.com/ethereum/go-ethereum v1.10.12
    github.com/fjl/memsize v0.0.0-20190710130421-bcb5799ab5e5
+    github.com/flare-foundation/flare v0.5.0
    github.com/gballet/go-libpcsclite v0.0.0-20191108122812-4678299bea08
    github.com/google/uuid v1.1.5
    github.com/gorilla/rpc v1.2.0
    github.com/gorilla/websocket v1.4.2
    github.com/hashicorp/go-bexpr v0.1.10
-    github.com/hashicorp/go-plugin v1.3.0
+    github.com/hashicorp/go-plugin v1.4.3
    github.com/hashicorp/golang-lru v0.5.5-0.20210104140557-80c98217689d
    github.com/holiman/bloomfilter/v2 v2.0.3
    github.com/holiman/uint256 v1.2.0
    github.com/mattn/go-colorable v0.1.8
    github.com/mattn/go-isatty v0.0.12
    github.com/olekukonko/tablewriter v0.0.5
+    github.com/prometheus/client_golang v1.7.1
+    github.com/prometheus/client_model v0.2.0
    github.com/prometheus/tsdb v0.10.0 // indirect
    github.com/rjeczalik/notify v0.9.2
    github.com/spf13/cast v1.3.1
@@ -30,8 +32,10 @@ require (
    github.com/spf13/viper v1.7.1
    github.com/status-im/keycard-go v0.0.0-20200402102358-957c09536969
    github.com/stretchr/testify v1.7.0
+    github.com/syndtr/goleveldb v0.1.1-0.20210819022825-2ae1ddf74ef7
    github.com/tyler-smith/go-bip39 v1.0.2
    golang.org/x/crypto v0.0.0-20210322153248-0c34fe9e7dc2
+    golang.org/x/sync v0.0.0-20210220032951-036812b2e83c
    golang.org/x/text v0.3.6
    golang.org/x/time v0.0.0-20210723032227-1f47c861a9ac
    gopkg.in/olebedev/go-duktape.v3 v3.0.0-20200619000410-60c24ae608a6

diff --git a/go.sum b/go.sum
index 5b333235..93546743 100644
--- a/go.sum
+++ b/go.sum
@@ -1,7 +1,7 @@
-module github.com
+++ b/go.sum
@@ -40,7 +40,6 @@
github.com/BurntSushi/toml v0.3.1/go.mod h1:xHWCNGjB5oqIDr8zfn03MMHue2Ht5sIBksp03
github.com/BurntSushi/xgb v0.0.0-20160522181843-27f122750802/go.mod h1:VnqG0Eym/WL8OVXweHU+Q+/VP0lqqI8lqeDX9IjBq=
github.com/DATA-DOG/go-sqlmock v1.3.3/go.mod h1:f/Ixk793poVm4qj/V1dPUg2JEAkC73Q5eFN3EC/SaM=
github.com/Microsoft/go-winio v0.4.14/go.mod h1:QXK3Q3Ka7+6tgxaGTie4Kpcdsi+P8jBhyozl1bpyYA=
-github.com/NYTimes/gziphandler v1.1.1 h1:ZUDjp0ae29j0ryrS0u/B8HZfJ8tBQHjqw2rQ2cqUQ31=
github.com/NYTimes/gziphandler v1.1.1/go.mod h1:n/CVRwUEdQIxrgPvAQhUUr9oeUtrvhMomdKFjzJNB0c=
github.com/OneOfOne/xxhash v1.2.2/go.mod h1:H5dpLMjK5SmBqAxs5vPj2TmRDMfkzw+cTzAELWljhcU=
github.com/StackExchange/wmi v0.0.0-20180116203802-5d049714c4a6 h1:fljPD/aNc3UIOAE6tdi6QkUempXK3P9BI7mr2hd6gx8=
@@ -60,8 +59,6 @@
github.com/armon/circbuf v0.0.0-20150827004946-bbba097214e/go.mod h1:3U/XgcO3hCbZ8TKRvWD2dDTcfh9M9ya+I9JpbB708=
github.com/armon/go-metrics v0.0.0-20180917152333-f0300d1749da/go.mod h1:U73ZrmVTwzksrR9V5SSury031EELLFMUz1kky1939pY=
github.com/armon/go-radix v0.0.0-20180808171621-7fddfc383310/go.mod h1:uFluZ+zHj4x4tNLv4JWepY2hxWSp5RywhrRMgIH9cCH8=
-github.com/ava-labs/avalanchego v1.6.4 h1:EbJGqYU9MqpsRVC9nmNmUNlFM8aiTTKNmAwuiHdmPs=
-github.com/ava-labs/avalanchego v1.6.4/go.mod h1:DzxlGkF8hj3GiwRo0fHq9GjYt4EyeIAsdfB08g6mI=
github.com/aws/aws-sdk-go v2 v1.2.0/go.mod h1:ZE0s0Y2RbmW1dJK0PoJv3ygDYOFFre1ejLJWL8FwAuQo=
github.com/aws/aws-sdk-go v2/config v1.1.1/go.mod h1:0XsVy9lBI/BCxm+2TuvT39YmdHwS5undQmxZ0Ye8F5Y=
github.com/aws/aws-sdk-go v2/credentials v1.1.1/go.mod h1:mM2iIjwL7LULWtS6J3CACYInboHirisUuDK8PoTHMM0U=
@@ -157,6 +154,8 @@
github.com/fatih/color v1.7.0 h1:DKwD4oS2D8LGgTVIwJXJSL5Vp2ffCQg58nFV38Ys=
github.com/fatih/color v1.7.0/go.mod h1:Zm6kSWBoL9eyXnKykthP6abPY2pDugNf5KwzbycvMj4=
github.com/fjl/memsize v0.0.0-20190710130421-bcb5799ab5e5 h1:FtmdgXiULNeRsonMF1KLDt+5+6hbjVMEW6RG07aUf7c=
github.com/fjl/memsize v0.0.0-20190710130421-bcb5799ab5e5/go.mod h1:VvhXp0YNQvB+uIk2RvXzuaQtKQJzIx6LSBe1xv7hi0=
+github.com/flare-foundation/flare v0.5.0 h1:pQJhItsDwEC8U7Q0lFCsBBrkMcD03KR2/0v7S5fbH2u=
+github.com/flare-foundation/flare v0.5.0/go.mod h1:ai++PWjLrL0ZYheY7qtU3bReW5JSe5X047QcIQ0W0U=
github.com/fogleman/gg v1.2.1-0.20190220221249-0403632d5b90/go.mod h1:R/bRT+9gY/C5z7JzPU0zXsXHKM4/aya+zqcVNZzPa1k=
github.com/fogleman/gg v1.3.0/go.mod h1:R/bRT+9gY/C5z7JzPU0zXsXHKM4/aya+zqcVNZzPa1k=
github.com/fortytwo/leaktest v1.3.0 h1:u8491cBMT08ft8aeV+adlcytMZylmA5nnwmKRZjI8vw=
@@ -256,9 +255,7 @@
github.com/googoleapis/gax-go v2 v2.0.5/go.mod h1:DWXyrwAJ9X0FpwwEdw+IPEYBICEFu5mhpdKc/us6b0k=
github.com/gopherjs/gopherjs v0.0.0-20181017120253-0766667cb4d1 h1:EGx4pi6eqNxGaHF6gqu48+N2wcFQ5qg5FXgd0dqsJ5d8=
github.com/gopherjs/gopherjs v0.0.0-20181017120253-0766667cb4d1/go.mod h1:wjF0RRmW1u3UXTncJ5qlYoELFm8eSnnE06hX4iZ3EWY=
-github.com/gorilla/handlers v1.4.2 h1:00niY0USkh01RGclFkxeNHK9bkDHGRYGMDFBCS+YARg=
github.com/gorilla/handlers v1.4.2/go.mod h1:Qkdc/uu4t4H4gmTK6auzZ766c4CAGn8+o/OAinr0iO=
-github.com/gorilla/mux v1.8.0 h1:140aqfKr1h25LN9hojwV5ZA91wcXF0vkdNieDF05koi=
github.com/gorilla/mux v1.8.0/go.mod h1:DVbg23sWSpFRCP05f1EN6jmj59UnW/n46H5rLB71So=
github.com/gorilla/rpc v1.2.0 h1:WvvdC2lNeT1SP32zrIcs10ECBfAlmrmsSbusC57wfk=
github.com/gorilla/rpc v1.2.0/go.mod h1:V4h9r+4sF5HnzqbwIez0fKSpANP0zlyd3qr7p36jktQ=
@@ -281,8 +278,9 @@
github.com/hashicorp/go-hclog v0.14.1/go.mod h1:whpDnt7S5deAju8AMKIwsul05p54N/39
github.com/hashicorp/go-immutable-radix v1.0.0/go.mod h1:0y9vanUI8NX6FsYo03zmjHv/C5i9g4Q3DwcSNZ4P60=
github.com/hashicorp/go-msgpack v0.5.3/go.mod h1:LdHT0lPgSu+cZNNakkCN/P3h0UDHfCYQXV3XU0M6sGGrk=
github.com/hashicorp/go-multierror v1.0.0/go.mod h1:dHt0lPgSu+cZNNakkCN/P3h0UDHfCYQXV3XU0M6sGGrk=
-github.com/hashicorp/go-plugin v1.3.0 h1:4d/WJoJzVHVI4i/rRjVaeuyxWrlzDE1mCydD8fXS8=
github.com/hashicorp/go-plugin v1.3.0/go.mod h1:F9eh4LrE/ZsRdbwhfjs9kHoDuWAhNytXdgmf1AVNs0=
+github.com/hashicorp/go-plugin v1.4.3 h1:DXmviwbWD5qd18ts9TpBC7BYL1Aia5sxbRq0B+v6Uzm=
+github.com/hashicorp/go-plugin v1.4.3/go.mod h1:5fGEH170VWtTcR0zV7yhdPLLMfX9YSZ38b18Udy6vYQ=
github.com/hashicorp/go-rootcerts v1.0.0/go.mod h1:K6zTfqRlCUIjKwsM4Z+hi5fzStQa6eBIzfwkVWnHU=
github.com/hashicorp/go-sockaddr v1.0.0/go.mod h1:7Xibr9yA9jJqL1pNB2Vw7kxv8xerXegt+ozgdvDeDU=
github.com/hashicorp/go-syslog v1.0.0/go.mod h1:qPfrKkXGihmCbJM2mZgKGvGkl1dFdvslp1gctolz4=
diff --git a/interfaces/interfaces.go b/interfaces/interfaces.go
index 100e658c..4e75ef41 100644
--- a/interfaces/interfaces.go
+++ b/interfaces/interfaces.go
@@ -1,7 +1,7 @@
-module github.com
+++ b/interfaces/interfaces.go
@@ -32,8 +32,8 @@
import (
    "errors"
    "math/big"

-    "github.com/ava-labs/coreth/core/types"
    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/core/types"
)

// NotFound is returned by API methods if the requested item does not exist.
diff --git a/internal/ethapi/api.go b/internal/ethapi/api.go
index 60e0242c..7ed7af3b 100644
--- a/internal/ethapi/api.go
+++ b/internal/ethapi/api.go
@@ -1,7 +1,7 @@
-module github.com
+++ b/internal/ethapi/api.go
@@ -34,6 +34,6 @@
import (
    "strings"
    "time"

-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/accounts/keystore"
-    "github.com/ava-labs/coreth/accounts/scwallet"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/vm"

```



```

- "github.com/ava-labs/coreth/params"
- "github.com/ava-labs/coreth/rpc"
- "github.com/davecgh/go-spew/spew"
- "github.com/ethereum/go-ethereum/accounts/abi"
- "github.com/ethereum/go-ethereum/common"
@@ -52,6 +42,17 @@ import (
- "github.com/ethereum/go-ethereum/crypto"
- "github.com/ethereum/go-ethereum/log"
- "github.com/ethereum/go-ethereum/rlp"
+ "github.com/flare-foundation/coreth/accounts"
+ "github.com/flare-foundation/coreth/accounts/keystore"
+ "github.com/flare-foundation/coreth/accounts/scwallet"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/flare-foundation/coreth/core/state"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+ "github.com/flare-foundation/coreth/eth/tracers/logger"
+ "github.com/flare-foundation/coreth/params"
+ "github.com/flare-foundation/coreth/rpc"
+ "github.com/flare-foundation/flare/ids"
+ "github.com/tyler-smith/go-bip39"
)

@@ -599,9 +600,9 @@ func NewPublicBlockchainAPI(b Backend) *PublicBlockchainAPI {
}

// ChainId is the EIP-155 replay-protection chain id for the current ethereum chain config.
-func (api *PublicBlockchainAPI) ChainId() (*hexutil.Big, error) {
+func (s *PublicBlockchainAPI) ChainId() (*hexutil.Big, error) {
// if current block is at or past the EIP-155 replay-protection fork block, return chainID from config
- if config := api.b.ChainConfig(); config.IsEIP155(api.b.CurrentBlock().Number()) {
+ if config := s.b.ChainConfig(); config.IsEIP155(s.b.CurrentBlock().Number()) {
return (*hexutil.Big)(config.ChainID), nil
}
return nil, fmt.Errorf("chain not synced beyond EIP-155 replay-protection fork block")
@@ -900,7 +901,11 @@ func DoCall(ctx context.Context, b Backend, args TransactionArgs, blockNrOrHash
if blkNumber, isNum := blockNrOrHash.Number(); isNum && blkNumber == rpc.PendingBlockNumber {
// Override header with a copy to ensure the original header is not modified
header = types.CopyHeader(header)
+ // Grab the hash of the unmodified header, so that the modified header can point to the
+ // prior block as its parent.
parentHash := header.Hash()
header.Time = uint64(time.Now().Unix())
+ header.ParentHash = parentHash
header.Number = new(big.Int).Add(header.Number, big.NewInt(1))
estimatedBaseFee, err := b.EstimateBaseFee(ctx)
if err != nil {
@@ -1157,7 +1162,7 @@ type StructLogRes struct {
}

// FormatLogs formats EVM returned structured logs for json output
-func FormatLogs(logs []vm.StructLog) []StructLogRes {
+func FormatLogs(logs []logger.StructLog) []StructLogRes {
formatted := make([]StructLogRes, len(logs))
for index, trace := range logs {
formatted[index] = StructLogRes{
@@ -1270,7 +1275,9 @@ func RPCMarshalBlock(block *types.Block, inclTx bool, fullTx bool, config *param
// a 'PublicBlockchainAPI'.
func (s *PublicBlockchainAPI) rpcMarshalHeader(ctx context.Context, header *types.Header) map[string]interface{} {
fields := RPCMarshalHeader(header)
- fields["totalDifficulty"] = (*hexutil.Big)(s.b.GetTd(ctx, header.Hash()))
+ // Note: Coreth enforces that the difficulty of a block is always 1, such that the total difficulty of a block
+ // will be equivalent to its height.
fields["totalDifficulty"] = (*hexutil.Big)(header.Number)
return fields
}

@@ -1282,7 +1289,9 @@ func (s *PublicBlockchainAPI) rpcMarshalBlock(ctx context.Context, b *types.Bloc
return nil, err
}
if inclTx {
- fields["totalDifficulty"] = (*hexutil.Big)(s.b.GetTd(ctx, b.Hash()))
+ // Note: Coreth enforces that the difficulty of a block is always 1, such that the total difficulty of a block
+ // will be equivalent to its height.
fields["totalDifficulty"] = (*hexutil.Big)(b.Number())
}
return fields, err
}

@@ -1459,9 +1468,9 @@ func AccessList(ctx context.Context, b Backend, blockNrOrHash rpc.BlockNumberOrH
precompiles := vm.ActivePrecompiles(b.ChainConfig().AvalancheRules(header.Number, new(big.Int).SetUint64(header.Time)))

// Create an initial tracer
- prevTracer := vm.NewAccessListTracer(nil, args.from(), to, precompiles)
+ prevTracer := logger.NewAccessListTracer(nil, args.from(), to, precompiles)
if args.AccessList != nil {
- prevTracer = vm.NewAccessListTracer(*args.AccessList, args.from(), to, precompiles)
+ prevTracer = logger.NewAccessListTracer(*args.AccessList, args.from(), to, precompiles)
}
for {
// Retrieve the current access list to expand
@@ -1488,7 +1497,7 @@ func AccessList(ctx context.Context, b Backend, blockNrOrHash rpc.BlockNumberOrH
}

// Apply the transaction with the access list tracer
- tracer := vm.NewAccessListTracer(accessList, args.from(), to, precompiles)
+ tracer := logger.NewAccessListTracer(accessList, args.from(), to, precompiles)
config := vm.Config{Tracer: tracer, Debug: true, NoBaseFee: true}
vmenv, _, err := b.GetEVM(ctx, msg, statedb, header, &config)
if err != nil {
diff --git a/internal/ethapi/backend.go b/internal/ethapi/backend.go
index 0f0795fa..f964de6 100644
--- a/internal/ethapi/backend.go
+++ b/internal/ethapi/backend.go
@@ -32,18 +32,18 @@ import (
- "math/big"
- "time"
-
- "github.com/ava-labs/coreth/accounts"
- "github.com/ava-labs/coreth/consensus"
- "github.com/ava-labs/coreth/core"
- "github.com/ava-labs/coreth/core/bloombits"
- "github.com/ava-labs/coreth/core/state"
- "github.com/ava-labs/coreth/core/types"
- "github.com/ava-labs/coreth/core/vm"
- "github.com/ava-labs/coreth/ethdb"
- "github.com/ava-labs/coreth/params"
- "github.com/ava-labs/coreth/rpc"
- "github.com/ethereum/go-ethereum/common"
- "github.com/ethereum/go-ethereum/event"
+ "github.com/flare-foundation/coreth/accounts"
+ "github.com/flare-foundation/coreth/consensus"
+ "github.com/flare-foundation/coreth/core"
+ "github.com/flare-foundation/coreth/core/bloombits"
+ "github.com/flare-foundation/coreth/core/state"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/core/vm"
+ "github.com/flare-foundation/coreth/ethdb"
+ "github.com/flare-foundation/coreth/params"
+ "github.com/flare-foundation/coreth/rpc"
)

// Backend interface provides the common API services (that are provided by
@@ -74,7 +74,6 @@ type Backend interface {
StateAndHeaderByNumber(ctx context.Context, number rpc.BlockNumber) (*state.StateDB, *types.Header, error)

```

```

StateAndHeaderByNumberOrHash(ctx context.Context, blockNrOrHash rpc.BlockNumberOrHash) (*state.StateDB, *types.Header, error)
GetReceipts(ctx context.Context, hash common.Hash) (types.Receipts, error)
- GetTd(ctx context.Context, hash common.Hash) *big.Int
GetEVM(ctx context.Context, msg core.Message, state *state.StateDB, header *types.Header, vmConfig *vm.Config) (*vm.EVM, func() error, error)
SubscribeChainEvent(ch chan<- core.ChainEvent) event.Subscription
SubscribeChainHeadEvent(ch chan<- core.ChainHeadEvent) event.Subscription
@@ -113,48 +112,48 @@ func GetAPIs(apiBackend Backend) []rpc.API {
    Version:      "1.0",
    Service:      NewPublicEthereumAPI(apiBackend),
    Public:       true,
+   Name:         "internal-public-eth",
+ }, {
    Namespace:    "eth",
    Version:      "1.0",
    Service:      NewPublicBlockChainAPI(apiBackend),
    Public:       true,
+   Name:         "internal-public-blockchain",
+ }, {
    Namespace:    "eth",
    Version:      "1.0",
    Service:      NewPublicTransactionPoolAPI(apiBackend, nonceLock),
    Public:       true,
+   Name:         "internal-public-transaction-pool",
+ }, {
    Namespace:    "txpool",
    Version:      "1.0",
    Service:      NewPublicTxPoolAPI(apiBackend),
    Public:       true,
+   Name:         "internal-public-tx-pool",
+ }, {
    Namespace:    "debug",
    Version:      "1.0",
    Service:      NewPublicDebugAPI(apiBackend),
    Public:       true,
+   Name:         "internal-public-debug",
+ }, {
    Namespace:    "debug",
    Version:      "1.0",
    Service:      NewPrivateDebugAPI(apiBackend),
    Name:         "internal-private-debug",
+ }, {
    Namespace:    "eth",
    Version:      "1.0",
    Service:      NewPublicAccountAPI(apiBackend.AccountManager()),
    Public:       true,
+   Name:         "internal-public-account",
+ }, {
    Namespace:    "personal",
    Version:      "1.0",
    Service:      NewPrivateAccountAPI(apiBackend, nonceLock),
    Public:       false,
+   Name:         "internal-private-personal",
+ },
    },
}

diff --git a/internal/ethapi/transaction_args.go b/internal/ethapi/transaction_args.go
index 0048b910..8feab39c 100644
--- a/internal/ethapi/transaction_args.go
+++ b/internal/ethapi/transaction_args.go
@@ -33,12 +33,12 @@ import (
    "fmt"
    "math/big"

-   "github.com/ava-labs/coreth/core/types"
-   "github.com/ava-labs/coreth/rpc"
+   "github.com/ethereum/go-ethereum/common"
+   "github.com/ethereum/go-ethereum/common/hexutil"
+   "github.com/ethereum/go-ethereum/common/math"
+   "github.com/ethereum/go-ethereum/log"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/rpc"
)

// TransactionArgs represents the arguments to construct a new transaction
diff --git a/internal/shutdowncheck/shutdown_tracker.go b/internal/shutdowncheck/shutdown_tracker.go
new file mode 100644
index 00000000..bfb2fbde
--- /dev/null
+++ b/internal/shutdowncheck/shutdown_tracker.go
@@ -0,0 +1,95 @@
@@ -0,0 +1,95 @@
+// (c) 2020-2021, Ava Labs, Inc.
+//
+// This file is a derived work, based on the go-ethereum library whose original
+// notices appear below.
+//
+// It is distributed under a license compatible with the licensing terms of the
+// original code from which it is derived.
+//
+// Much love to the original authors for their work.
+// *****
+// Copyright 2021 The go-ethereum Authors
+// This file is part of the go-ethereum library.
+//
+// The go-ethereum library is free software: you can redistribute it and/or modify
+// it under the terms of the GNU Lesser General Public License as published by
+// the Free Software Foundation, either version 3 of the License, or
+// (at your option) any later version.
+//
+// The go-ethereum library is distributed in the hope that it will be useful,
+// but WITHOUT ANY WARRANTY; without even the implied warranty of
+// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
+// GNU Lesser General Public License for more details.
+//
+// You should have received a copy of the GNU Lesser General Public License
+// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package shutdowncheck
+
+import (
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
+)
+
+// ShutdownTracker is a service that reports previous unclean shutdowns
+// upon start. It needs to be started after a successful start-up and stopped
+// after a successful shutdown, just before the db is closed.
+type ShutdownTracker struct {
+    db      ethdb.Database
+    stopCh chan struct{}
+}
+
+// NewShutdownTracker creates a new ShutdownTracker instance and has
+// no other side-effect.
+func NewShutdownTracker(db ethdb.Database) *ShutdownTracker {
+    return &ShutdownTracker{
+        db:      db,
+        stopCh: make(chan struct{}),
+    }
+}

```

```
// MarkStartup is to be called in the beginning when the node starts. It will:
// - Push a new startup marker to the db
// - Report previous unclean shutdowns
+func (t *ShutdownTracker) MarkStartup() {
+    if uncleanShutdowns, discards, err := rawdb.PushUncleanShutdownMarker(t.db); err != nil {
+        log.Error("Could not update unclean-shutdown-marker list", "error", err)
+    } else {
+        if discards > 0 {
+            log.Warn("Old unclean shutdowns found", "count", discards)
+        }
+        for _, tstamp := range uncleanShutdowns {
+            t := time.Unix(int64(tstamp), 0)
+            log.Warn("Unclean shutdown detected", "booted", t,
+                "age", common.PrettyAge(t))
+        }
+    }
+}
+
+// Start runs an event loop that updates the current marker's timestamp every 5 minutes.
+func (t *ShutdownTracker) Start() {
+    go func() {
+        ticker := time.NewTicker(5 * time.Minute)
+        defer ticker.Stop()
+        for {
+            select {
+            case <-ticker.C:
+                rawdb.UpdateUncleanShutdownMarker(t.db)
+            case <-t.stopCh:
+                return
+            }
+        }
+    }()
+}
+
+// Stop will stop the update loop and clear the current marker.
+func (t *ShutdownTracker) Stop() {
+    // Stop update loop.
+    t.stopCh <- struct{}{}
+    // Clear last marker.
+    rawdb.PopUncleanShutdownMarker(t.db)
+}
+
diff --git a/metrics/prometheus/prometheus.go b/metrics/prometheus/prometheus.go
new file mode 100644
index 00000000..83b92a7e
--- /dev/null
+++ b/metrics/prometheus/prometheus.go
@@ -0,0 +1,185 @@
+// (c) 2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package prometheus
+
+import (
+    "sort"
+    "strings"
+
+    "github.com/ethereum/go-ethereum/metrics"
+    "github.com/prometheus/client_golang/prometheus"
+    dto "github.com/prometheus/client_model/go"
+
+    pv          = []float64{.5, .75, .95, .99, .999, .9999}
+    pvShortPercent = []float64{50, 95, 99}
+    pvShort      = []float64{.50, .95, .99}
+
+type gatherer struct {
+    reg metrics.Registry
+}
+
+func (g gatherer) Gather() ([]*dto.MetricFamily, error) {
+    // Gather and pre-sort the metrics to avoid random listings
+    var names []string
+    g.reg.Each(func(name string, i interface{}) {
+        names = append(names, name)
+    })
+    sort.Strings(names)
+
+    mfs := make([]*dto.MetricFamily, 0, len(names))
+    for _, name := range names {
+        mIntf := g.reg.Get(name)
+        name := strings.Replace(name, "/", "_", -1)
+
+        switch m := mIntf.(type) {
+        case metrics.Counter:
+            val := m.Snapshot().Count()
+            valFloat := float64(val)
+            mfs = append(mfs, &dto.MetricFamily{
+                Name: &name,
+                Type: dto.MetricType_COUNTER.Enum(),
+                Metric: []*dto.Metric{{
+                    Counter: &dto.Counter{
+                        Value: &valFloat,
+                    },
+                }},
+            })
+        case metrics.Gauge:
+            val := m.Snapshot().Value()
+            valFloat := float64(val)
+            mfs = append(mfs, &dto.MetricFamily{
+                Name: &name,
+                Type: dto.MetricType_GAUGE.Enum(),
+                Metric: []*dto.Metric{{
+                    Gauge: &dto.Gauge{
+                        Value: &valFloat,
+                    },
+                }},
+            })
+        case metrics.GaugeFloat64:
+            val := m.Snapshot().Value()
+            mfs = append(mfs, &dto.MetricFamily{
+                Name: &name,
+                Type: dto.MetricType_GAUGE.Enum(),
+                Metric: []*dto.Metric{{
+                    Gauge: &dto.Gauge{
+                        Value: &val,
+                    },
+                }},
+            })
+        case metrics.Histogram:
+            snapshot := m.Snapshot()
+            count := snapshot.Count()
+            countUint := uint64(count)
+            sum := snapshot.Sum()
+            sumFloat := float64(sum)
+
+            ps := m.Percentiles(pv)
+            qs := make([]*dto.Quantile, len(pv))
+            for i := range ps {
+                v := pv[i]
+                s := ps[i]
```

```

+         qs[i] = &dto.Quantile{
+             Quantile: &v,
+             Value:     &s,
+         }
+     }
+
+     mfs = append(mfs, &dto.MetricFamily{
+         Name: &name,
+         Type:  dto.MetricType_SUMMARY.Enum(),
+         Metric: []*dto.Metric{{
+             Summary: &dto.Summary{
+                 SampleCount: &countUint,
+                 SampleSum:    &sumFloat,
+                 Quantile:     qs,
+             },
+         }},
+     })
+
+ case metrics.Meter:
+     val := m.Snapshot().Count()
+     valFloat := float64(val)
+     mfs = append(mfs, &dto.MetricFamily{
+         Name: &name,
+         Type:  dto.MetricType_GAUGE.Enum(),
+         Metric: []*dto.Metric{{
+             Gauge: &dto.Gauge{
+                 Value: &valFloat,
+             },
+         }},
+     })
+
+ case metrics.Timer:
+     snapshot := m.Snapshot()
+     count := snapshot.Count()
+     countUint := uint64(count)
+     sum := snapshot.Sum()
+     sumFloat := float64(sum)
+
+     ps := m.Percentiles(pv)
+     qs := make([]*dto.Quantile, len(pv))
+     for i := range ps {
+         v := pv[i]
+         s := ps[i]
+         qs[i] = &dto.Quantile{
+             Quantile: &v,
+             Value:     &s,
+         }
+     }
+
+     mfs = append(mfs, &dto.MetricFamily{
+         Name: &name,
+         Type:  dto.MetricType_SUMMARY.Enum(),
+         Metric: []*dto.Metric{{
+             Summary: &dto.Summary{
+                 SampleCount: &countUint,
+                 SampleSum:    &sumFloat,
+                 Quantile:     qs,
+             },
+         }},
+     })
+
+ case metrics.ResettingTimer:
+     snapshot := m.Snapshot()
+
+     vals := snapshot.Values()
+     count := uint64(len(vals))
+     if count == 0 {
+         continue
+     }
+
+     ps := m.Percentiles(pvShortPercent)
+     qs := make([]*dto.Quantile, len(pv))
+     for i := range pvShort {
+         v := pv[i]
+         s := float64(ps[i])
+         qs[i] = &dto.Quantile{
+             Quantile: &v,
+             Value:     &s,
+         }
+     }
+
+     mfs = append(mfs, &dto.MetricFamily{
+         Name: &name,
+         Type:  dto.MetricType_SUMMARY.Enum(),
+         Metric: []*dto.Metric{{
+             Summary: &dto.Summary{
+                 SampleCount: &count,
+                 // TODO: do we need to specify SampleSum here? and if so
+                 // what should that be?
+                 Quantile: qs,
+             },
+         }},
+     })
+ }
+
+ return mfs, nil
+}
+
+func Gatherer(reg metrics.Registry) prometheus.Gatherer {
+    return gatherer(reg: reg)
+}
+
+diff --git a/miner/miner.go b/miner/miner.go
index 263439c8..f05ae5c2 100644
--- a/miner/miner.go
+++ b/miner/miner.go
@@ -28,12 +28,13 @@
package miner

import (
-     "github.com/ava-labs/coreth/consensus"
-     "github.com/ava-labs/coreth/core"
-     "github.com/ava-labs/coreth/core/types"
-     "github.com/ava-labs/coreth/params"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/ethereum/go-ethereum/event"
+     "github.com/flare-foundation/coreth/consensus"
+     "github.com/flare-foundation/coreth/core"
+     "github.com/flare-foundation/coreth/core/types"
+     "github.com/flare-foundation/coreth/params"
+     "github.com/flare-foundation/flare/utls/timer/mockable"
)

// Backend wraps all methods required for mining.
@@ -51,9 +52,9 @@
type Miner struct {
    worker *worker
}

-func New(eth Backend, config *Config, chainConfig *params.ChainConfig, mux *event.TypeMux, engine consensus.Engine) *Miner {
+func New(eth Backend, config *Config, chainConfig *params.ChainConfig, mux *event.TypeMux, engine consensus.Engine, clock *mockable.Clock) *Miner {
    return &Miner{
        worker: newWorker(config, chainConfig, engine, eth, mux),
        worker: newWorker(config, chainConfig, engine, eth, mux, clock),
    }
}

diff --git a/miner/worker.go b/miner/worker.go
index 1e6c3df8..40702504 100644

```

```

--- a/miner/worker.go
+++ b/miner/worker.go
@@ -36,16 +36,17 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/coreth/consensus"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/consensus/misc"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/event"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/consensus"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/consensus/misc"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/flare/utlis/timer/mockable"
)

// environment is the worker's current environment and holds all of the current state information.
@@ -81,9 +82,10 @@ type worker struct {
    mux      *event.TypeMux // TODO replace
    mu       sync.RWMutex   // The lock used to protect the coinbase and extra fields
    coinbase common.Address
+    clock    *mockable.Clock // Allows us mock the clock for testing
}

-func newWorker(config *Config, chainConfig *params.ChainConfig, engine consensus.Engine, eth Backend, mux *event.TypeMux) *worker {
+func newWorker(config *Config, chainConfig *params.ChainConfig, engine consensus.Engine, eth Backend, mux *event.TypeMux, clock *mockable.Clock) *worker {
    worker := &worker{
        config:      config,
        chainConfig: chainConfig,
@@ -91,6 +93,7 @@ func newWorker(config *Config, chainConfig *params.ChainConfig, engine consensus
    eth:         eth,
    mux:         mux,
    chain:       eth.BlockChain(),
+    clock:       clock,
    }

    return worker
@@ -108,7 +111,7 @@ func (w *worker) commitNewWork() (*types.Block, error) {
    w.mu.RLock()
    defer w.mu.RUnlock()

-    tstart := time.Now()
+    tstart := w.clock.Time()
    timestamp := tstart.Unix()
    parent := w.chain.CurrentBlock()
    // Note: in order to support asynchronous block production, blocks are allowed to have
@@ -119,11 +122,11 @@ func (w *worker) commitNewWork() (*types.Block, error) {
    }

    var gasLimit uint64
-    if w.chainConfig.IsApricotPhase1(big.NewInt(timestamp)) {
+    if w.chainConfig.IsApricotPhase5(big.NewInt(timestamp)) {
+        gasLimit = params.ApricotPhase5GasLimit
    } else if w.chainConfig.IsApricotPhase1(big.NewInt(timestamp)) {
        gasLimit = params.ApricotPhase1GasLimit
    } else {
        // The gas limit is set in phase1 to ApricotPhase1GasLimit because the ceiling and floor were set to the same value
        // such that the gas limit converged to it. Since this is hardbaked now, we remove the ability to configure it.
        gasLimit = core.CalcGasLimit(parent.GasUsed(), parent.GasLimit(), params.ApricotPhase1GasLimit, params.ApricotPhase1GasLimit)
    }

    num := parent.Number()
    diff --git a/node/api.go b/node/api.go
index c38b2df2..9f5ce07a 100644
--- a/node/api.go
+++ b/node/api.go
@@ -27,24 +27,28 @@ package node

import (
-    "github.com/ava-labs/coreth/internal/debug"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/internal/debug"
+    "github.com/flare-foundation/coreth/rpc"
)

// apis returns the collection of built-in RPC APIs.
func (n *Node) apis() []rpc.API {
-    return []rpc.API{{
-        Namespace: "debug",
-        Version:   "1.0",
-        Service:   debug.Handler,
-    }, {
-        Namespace: "web3",
-        Version:   "1.0",
-        Service:   &publicWeb3API{n},
-        Public:    true,
-    }},
+    return []rpc.API{
+        {
+            Namespace: "debug",
+            Version:   "1.0",
+            Service:   debug.Handler,
+            Name:      "debug-handler",
+        },
+        {
+            Namespace: "web3",
+            Version:   "1.0",
+            Service:   &publicWeb3API{n},
+            Public:    true,
+            Name:      "web3",
+        },
    }
}

diff --git a/node/config.go b/node/config.go
index 5b7be4e3..ab20ff59 100644
--- a/node/config.go
+++ b/node/config.go
@@ -32,11 +32,11 @@ import (
    "os"
    "path/filepath"

-    "github.com/ava-labs/coreth/accounts"
-    "github.com/ava-labs/coreth/accounts/external"
-    "github.com/ava-labs/coreth/accounts/keystore"
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/accounts"
+    "github.com/flare-foundation/coreth/accounts/external"
+    "github.com/flare-foundation/coreth/accounts/keystore"
+    "github.com/flare-foundation/coreth/rpc"
)

```

```

// Config represents a small collection of configuration values to fine tune the
diff --git a/node/defaults.go b/node/defaults.go
index e4c826b9..1d8728fb 100644
--- a/node/defaults.go
+++ b/node/defaults.go
@@ -27,7 +27,7 @@
package node

import (
-     "github.com/ava-labs/coreth/rpc"
+     "github.com/flare-foundation/coreth/rpc"
)

const (
diff --git a/node/node.go b/node/node.go
index 8d5cc504..819f32a5 100644
--- a/node/node.go
+++ b/node/node.go
@@ -27,59 +27,18 @@
package node

import (
-     "sync"
-
-     "github.com/ava-labs/coreth/accounts"
-     "github.com/ava-labs/coreth/rpc"
-     "github.com/ethereum/go-ethereum/event"
+     "github.com/flare-foundation/coreth/accounts"
+     "github.com/flare-foundation/coreth/rpc"
)

// Node is a container on which services can be registered.
type Node struct {
-     eventmux *event.TypeMux
-     config    *Config
-     accman    *accounts.Manager
-     // log      log.Logger
-     // ephemKeystore string // if non-empty, the key directory that will be removed by Stop
-     // dirLock    fileutil.Releaser // prevents concurrent use of instance directory
-     // stop       chan struct{} // Channel to wait for termination notifications
-     // server     *p2p.Server // Currently running P2P networking layer
-     // startStopLock sync.Mutex // Start/Stop are protected by an additional lock
-     // state int // Tracks state of node lifecycle
-
-     lock sync.Mutex
-     rpcAPIs []rpc.API // List of APIs currently provided by the node
-     // inprocHandler *rpc.Server // In-process RPC request handler to process the API requests
-
-     // databases map[*closeTrackingDB]struct{} // All open databases
+     config *Config
+     accman *accounts.Manager
+
+     corethVersion string
}

-// const (
-//     initializingState = iota
-//     closedState
-// )
-
-// func (n *Node) openDataDir() error {
-//     if n.config.DataDir == "" {
-//         return nil // ephemeral
-//     }
-//
-//     instdir := filepath.Join(n.config.DataDir, n.config.name())
-//     if err := os.MkdirAll(instdir, 0700); err != nil {
-//         return err
-//     }
-//     // Lock the instance directory to prevent concurrent use by another instance as well as
-//     // accidental use of the instance directory as a database.
-//     release, _, err := fileutil.Flock(filepath.Join(instdir, "LOCK"))
-//     if err != nil {
-//         return convertFileLockError(err)
-//     }
-//     n.dirLock = release
-//     return nil
-// }
-
-// New creates a new P2P node, ready for protocol registration.
func New(conf *Config) (*Node, error) {
    // Copy config and resolve the datadir so future changes to the current
@@ -87,13 +46,7 @@
func New(conf *Config) (*Node, error) {
    confCopy := *conf
    conf = &confCopy

    node := &Node{
        config:    conf,
        eventmux:  new(event.TypeMux),
    }

    // Register built-in APIs.
    node.rpcAPIs = append(node.rpcAPIs, node.apis()...)
    node := &Node{config: conf}

    // Ensure that the AccountManager method works before the node has started. We rely on
    // this in cmd/geth.
@@ -103,8 +56,6 @@
func New(conf *Config) (*Node, error) {
    }
    node.accman = am

    // Configure RPC servers.

    return node, nil
}

@@ -118,16 +69,7 @@
func (n *Node) AccountManager() *accounts.Manager {
    return n.accman
}

// EventMux retrieves the event multiplexer used by all the network services in
// the current protocol stack.
-func (n *Node) EventMux() *event.TypeMux {
-    return n.eventmux
-}
-
// RegisterAPIs registers the APIs a service provides on the node.
-func (n *Node) RegisterAPIs(apis []rpc.API) {
-    n.lock.Lock()
-    defer n.lock.Unlock()
-
-    n.rpcAPIs = append(n.rpcAPIs, apis...)
+func (n *Node) APIs() []rpc.API {
+    return n.apis()
}

diff --git a/params/avalanche_params.go b/params/avalanche_params.go
index a6124df6..df871bf0 100644
--- a/params/avalanche_params.go
+++ b/params/avalanche_params.go
@@ -4,11 +4,13 @@
package params

import (

```

```

-      "github.com/ava-labs/avalanchego/units/units"
-      "math/big"
+
+      "github.com/flare-foundation/flare/units/units"
+    )
+
+    // Minimum Gas Price
+    -var (
+const (
+    // MinGasPrice is the number of nAVAX required per gas unit for a
+    // transaction to be valid, measured in wei
+    LaunchMinGasPrice      uint64 = 470_000_000_000
+
+@@ -17,11 +19,30 @@ var (
+    AvalancheAtomicTxFee = units.MilliAvax
+
+    ApricotPhase1GasLimit uint64 = 8_000_000
+    ApricotPhase5GasLimit uint64 = 30_000_000
+
+    +
+    ApricotPhase3ExtraDataSize      = 80
+    ApricotPhase3MinBaseFee          int64 = 75_000_000_000
+    ApricotPhase3MaxBaseFee          int64 = 225_000_000_000
+    ApricotPhase3InitialBaseFee      int64 = 225_000_000_000
+    ApricotPhase3TargetGas           uint64 = 10_000_000
+    ApricotPhase4MinBaseFee          int64 = 25_000_000_000
+    ApricotPhase4MaxBaseFee          int64 = 1_000_000_000_000
+    ApricotPhase4BaseFeeChangeDenominator uint64 = 12
+    ApricotPhase5TargetGas           uint64 = 150_000_000
+    ApricotPhase5BaseFeeChangeDenominator uint64 = 36
+
+    -    ApricotPhase3ExtraDataSize      = 80
+    -    ApricotPhase3MinBaseFee          int64 = 75_000_000_000
+    -    ApricotPhase3MaxBaseFee          int64 = 225_000_000_000
+    -    ApricotPhase3InitialBaseFee      int64 = 225_000_000_000
+    -    ApricotPhase4MinBaseFee          int64 = 25_000_000_000
+    -    ApricotPhase4MaxBaseFee          int64 = 1_000_000_000_000
+    +    // The base cost to charge per atomic transaction. Added in Apricot Phase 5.
+    +    AtomicTxBaseCost uint64 = 10_000
+    +)
+
+    +var (
+    +    // The atomic gas limit specifies the maximum amount of gas that can be consumed by the atomic
+    +    // transactions included in a block and is enforced as of ApricotPhase5. Prior to ApricotPhase5,
+    +    // a block included a single atomic transaction. As of ApricotPhase5, each block can include a set
+    +    // of atomic transactions where the cumulative atomic gas consumed is capped by the atomic gas limit,
+    +    // similar to the block gas limit.
+    +    //
+    +    // This value must always remain <= MaxUint64.
+    +    AtomicGasLimit *big.Int = big.NewInt(100_000)
+    +)
+
+diff --git a/params/config.go b/params/config.go
index 5bb4a55c..596f0cd4 100644
--- a/params/config.go
+++ b/params/config.go
+@@ -37,20 +37,22 @@ import (
+
+ // Avalanche ChainIDs
+ var (
+     -    // AvalancheMainnetChainID ...
+     -    AvalancheMainnetChainID = big.NewInt(43114)
+     -    // AvalancheFujiChainID ...
+     -    AvalancheFujiChainID = big.NewInt(43113)
+     -    // AvalancheLocalChainID ...
+     -    AvalancheLocalChainID = big.NewInt(43112)
+     +    // FlareChainID ...
+     +    FlareChainID = big.NewInt(14)
+     +    // SongbirdChainID ...
+     +    SongbirdChainID = big.NewInt(19)
+     +    // CostonChainID ...
+     +    CostonChainID = big.NewInt(16)
+     +    // LocalChainID ...
+     +    LocalChainID = big.NewInt(4294967295)
+
+     errNonGenesisForkByHeight = errors.New("coreth only supports forking by height at the genesis block")
+ )
+
+ var (
+     -    // AvalancheMainnetChainConfig is the configuration for Avalanche Main Network
+     -    AvalancheMainnetChainConfig = &ChainConfig{
+     -        ChainID:      AvalancheMainnetChainID,
+     +    // FlareChainConfig is the configuration for Flare main network.
+     +    FlareChainConfig = &ChainConfig{
+     +        ChainID:      FlareChainID,
+     +        HomesteadBlock: big.NewInt(0),
+     +        DAOForkBlock:    big.NewInt(0),
+     +        DAOForkSupport:    true,
+@@ -63,15 +65,16 @@ var (
+     +        PetersburgBlock: big.NewInt(0),
+     +        IstanbulBlock:    big.NewInt(0),
+     +        MuirGlacierBlock:  big.NewInt(0),
+     -        ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2021, time.March, 31, 14, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2021, time.May, 10, 11, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2021, time.August, 24, 14, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2021, time.September, 22, 21, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase5BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     }
+
+     -    // AvalancheFujiChainConfig is the configuration for the Fuji Test Network
+     -    AvalancheFujiChainConfig = &ChainConfig{
+     -        ChainID:      AvalancheFujiChainID,
+     +    // SongbirdChainConfig is the configuration for the Songbird canary network.
+     +    SongbirdChainConfig = &ChainConfig{
+     +        ChainID:      SongbirdChainID,
+     +        HomesteadBlock: big.NewInt(0),
+     +        DAOForkBlock:    big.NewInt(0),
+     +        DAOForkSupport:    true,
+@@ -84,15 +87,16 @@ var (
+     +        PetersburgBlock: big.NewInt(0),
+     +        IstanbulBlock:    big.NewInt(0),
+     +        MuirGlacierBlock:  big.NewInt(0),
+     -        ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2021, time.March, 26, 14, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2021, time.May, 5, 14, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2021, time.August, 16, 19, 0, 0, 0, time.UTC).Unix()),
+     -        ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2021, time.September, 16, 21, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     +        ApricotPhase5BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, 0, time.UTC).Unix()),
+     }
+
+     -    // AvalancheLocalChainConfig is the configuration for the Avalanche Local Network
+     -    AvalancheLocalChainConfig = &ChainConfig{
+     -        ChainID:      AvalancheLocalChainID,
+     +    // CostonChainConfig is the configuration for the Coston test network.
+     +    CostonChainConfig = &ChainConfig{
+     +        ChainID:      CostonChainID,
+     +        HomesteadBlock: big.NewInt(0),
+     +        DAOForkBlock:    big.NewInt(0),
+     +        DAOForkSupport:    true,
+@@ -105,18 +109,42 @@ var (

```

```

PetersburgBlock:      big.NewInt(0),
IstanbulBlock:        big.NewInt(0),
MuirGlacierBlock:     big.NewInt(0),
ApricotPhase1BlockTimestamp: big.NewInt(0),
ApricotPhase2BlockTimestamp: big.NewInt(0),
ApricotPhase3BlockTimestamp: big.NewInt(0),
ApricotPhase4BlockTimestamp: big.NewInt(0),
}

TestChainConfig = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestLaunchConfig = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase1Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase2Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase3Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase4Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase5BlockTimestamp: big.NewInt(time.Date(2100, time.January, 1, 0, 0, 0, time.UTC).Unix()),
}

// LocalChainConfig is the configuration for the local network.
LocalChainConfig = &ChainConfig{
ChainID:      LocalChainID,
HomesteadBlock:      big.NewInt(0),
DAOForkBlock:        big.NewInt(0),
DAOForkSupport:      true,
EIP150Block:         big.NewInt(0),
EIP150Hash:          common.HexToHash("0x2086799aeebeae135c246c5021c82b4e15a2c451340993aacfd2751886514f0"),
EIP155Block:         big.NewInt(0),
EIP158Block:         big.NewInt(0),
ByzantiumBlock:      big.NewInt(0),
ConstantinopleBlock: big.NewInt(0),
PetersburgBlock:     big.NewInt(0),
IstanbulBlock:       big.NewInt(0),
MuirGlacierBlock:    big.NewInt(0),
ApricotPhase1BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase2BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase3BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase4BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
ApricotPhase5BlockTimestamp: big.NewInt(time.Date(2000, time.January, 1, 0, 0, 0, time.UTC).Unix()),
}

TestChainConfig = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestLaunchConfig = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase1Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase2Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase3Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase4Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestApricotPhase5Config = &ChainConfig{big.NewInt(1), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}, big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), big.NewInt(0), nil, false, big.NewInt(0), common.Hash{}},
TestRules = TestChainConfig.AvalancheRules(new(big.Int), new(big.Int))
}

@@ -155,11 +183,13 @@ type ChainConfig struct {
ApricotPhase3BlockTimestamp *big.Int `json:"apricotPhase3BlockTimestamp,omitempty"`
// Apricot Phase 4 introduces the notion of a block fee to the dynamic fee algorithm (nil = no fork, 0 = already activated)
ApricotPhase4BlockTimestamp *big.Int `json:"apricotPhase4BlockTimestamp,omitempty"`
// Apricot Phase 5 introduces a batch of atomic transactions with a maximum atomic gas limit per block. (nil = no fork, 0 = already activated)
ApricotPhase5BlockTimestamp *big.Int `json:"apricotPhase5BlockTimestamp,omitempty"`
}

// String implements the fmt.Stringer interface.
func (c *ChainConfig) String() string {
return fmt.Sprintf("{ChainID: %v Homestead: %v DAO: %v DAOSupport: %v EIP150: %v EIP155: %v EIP158: %v Byzantium: %v Constantinople: %v Petersburg: %v Istanbul: %v, Muir Glacier: %v, Apricot Phase1: %v, Apricot Phase2: %v, Apricot Phase3: %v, Apricot Phase4: %v, Apricot Phase5: %v}",
c.ChainID,
c.HomesteadBlock,
c.DAOForkBlock,
c.DAOForkSupport,
c.EIP150Block,
c.EIP150Hash,
c.EIP155Block,
c.EIP158Block,
c.ByzantiumBlock,
c.ConstantinopleBlock,
c.PetersburgBlock,
c.IstanbulBlock,
c.MuirGlacierBlock,
c.ApricotPhase1BlockTimestamp,
c.ApricotPhase2BlockTimestamp,
c.ApricotPhase3BlockTimestamp,
c.ApricotPhase4BlockTimestamp,
c.ApricotPhase5BlockTimestamp,
}

@@ -176,6 +206,7 @@ func (c *ChainConfig) String() string {
c.ChainID,
c.HomesteadBlock,
c.DAOForkBlock,
c.DAOForkSupport,
c.EIP150Block,
c.EIP150Hash,
c.EIP155Block,
c.EIP158Block,
c.ByzantiumBlock,
c.ConstantinopleBlock,
c.PetersburgBlock,
c.IstanbulBlock,
c.MuirGlacierBlock,
c.ApricotPhase1BlockTimestamp,
c.ApricotPhase2BlockTimestamp,
c.ApricotPhase3BlockTimestamp,
c.ApricotPhase4BlockTimestamp,
c.ApricotPhase5BlockTimestamp,
}

@@ -257,15 +288,22 @@ func (c *ChainConfig) IsApricotPhase4(blockTimestamp *big.Int) bool {
return isForked(c.ApricotPhase4BlockTimestamp, blockTimestamp)
}

// IsApricotPhase5 returns whether [blockTimestamp] represents a block
// with a timestamp after the Apricot Phase 5 upgrade time.
+func (c *ChainConfig) IsApricotPhase5(blockTimestamp *big.Int) bool {
return isForked(c.ApricotPhase5BlockTimestamp, blockTimestamp)
}

// CheckCompatible checks whether scheduled fork transitions have been imported
// with a mismatching chain configuration.
-func (c *ChainConfig) CheckCompatible(newcfg *ChainConfig, height uint64) *ConfigCompatError {
+func (c *ChainConfig) CheckCompatible(newcfg *ChainConfig, height uint64, timestamp uint64) *ConfigCompatError {
bhead := new(big.Int).SetUint64(height)
bheadTimestamp := new(big.Int).SetUint64(timestamp)

// Iterate checkCompatible to find the lowest conflict.
var lasterr *ConfigCompatError
for {
err := c.checkCompatible(newcfg, bhead)
err := c.checkCompatible(newcfg, bhead, bheadTimestamp)
if err == nil || (lasterr != nil && err.RewindTo == lasterr.RewindTo) {
break
}
}

@@ -356,48 +394,63 @@ func (c *ChainConfig) CheckConfigForkOrder() error {
return nil
}

-func (c *ChainConfig) checkCompatible(newcfg *ChainConfig, head *big.Int) *ConfigCompatError {
if isForkIncompatible(c.HomesteadBlock, newcfg.HomesteadBlock, head) {
return newCompatError("Homestead fork block", c.HomesteadBlock, newcfg.HomesteadBlock)
}
if isForkIncompatible(c.DAOForkBlock, newcfg.DAOForkBlock, head) {
return newCompatError("DAO fork block", c.DAOForkBlock, newcfg.DAOForkBlock)
}
if c.IsDAOFork(head) && c.DAOForkSupport != newcfg.DAOForkSupport {
return newCompatError("DAO fork support flag", c.DAOForkBlock, newcfg.DAOForkBlock)
}
if isForkIncompatible(c.EIP150Block, newcfg.EIP150Block, head) {
return newCompatError("EIP150 fork block", c.EIP150Block, newcfg.EIP150Block)
}
if isForkIncompatible(c.EIP155Block, newcfg.EIP155Block, head) {
return newCompatError("EIP155 fork block", c.EIP155Block, newcfg.EIP155Block)
}
if isForkIncompatible(c.EIP158Block, newcfg.EIP158Block, head) {
return newCompatError("EIP158 fork block", c.EIP158Block, newcfg.EIP158Block)
}
if c.IsEIP158(head) && !configNumEqual(c.ChainID, newcfg.ChainID) {
if c.IsEIP158(headHeight) && !configNumEqual(c.ChainID, newcfg.ChainID) {
return newCompatError("EIP158 fork block", c.EIP158Block, newcfg.EIP158Block)
}
}
}

```



```

        return newCompatError("EIP158 chain ID", c.EIP158Block, newcfg.EIP158Block)
    }
-   if isForkIncompatible(c.ByzantiumBlock, newcfg.ByzantiumBlock, head) {
+   if isForkIncompatible(c.ByzantiumBlock, newcfg.ByzantiumBlock, headHeight) {
        return newCompatError("Byzantium fork block", c.ByzantiumBlock, newcfg.ByzantiumBlock)
    }
-   if isForkIncompatible(c.ConstantinopleBlock, newcfg.ConstantinopleBlock, head) {
+   if isForkIncompatible(c.ConstantinopleBlock, newcfg.ConstantinopleBlock, headHeight) {
        return newCompatError("Constantinople fork block", c.ConstantinopleBlock, newcfg.ConstantinopleBlock)
    }
-   if isForkIncompatible(c.PetersburgBlock, newcfg.PetersburgBlock, head) {
+   if isForkIncompatible(c.PetersburgBlock, newcfg.PetersburgBlock, headHeight) {
        // the only case where we allow Petersburg to be set in the past is if it is equal to Constantinople
        // mainly to satisfy fork ordering requirements which state that Petersburg fork be set if Constantinople fork is set
-   if isForkIncompatible(c.ConstantinopleBlock, newcfg.PetersburgBlock, head) {
+   if isForkIncompatible(c.ConstantinopleBlock, newcfg.PetersburgBlock, headHeight) {
        return newCompatError("Petersburg fork block", c.PetersburgBlock, newcfg.PetersburgBlock)
    }
}
-   if isForkIncompatible(c.IstanbulBlock, newcfg.IstanbulBlock, head) {
+   if isForkIncompatible(c.IstanbulBlock, newcfg.IstanbulBlock, headHeight) {
        return newCompatError("Istanbul fork block", c.IstanbulBlock, newcfg.IstanbulBlock)
    }
-   if isForkIncompatible(c.MuirGlacierBlock, newcfg.MuirGlacierBlock, head) {
+   if isForkIncompatible(c.MuirGlacierBlock, newcfg.MuirGlacierBlock, headHeight) {
        return newCompatError("Muir Glacier fork block", c.MuirGlacierBlock, newcfg.MuirGlacierBlock)
    }
-   // TODO(aaronbuchwald) ensure that Avalanche Blocktimestamps are not modified
+   if isForkIncompatible(c.ApricotPhase1BlockTimestamp, newcfg.ApricotPhase1BlockTimestamp, headTimestamp) {
+       return newCompatError("ApricotPhase1 fork block timestamp", c.ApricotPhase1BlockTimestamp, newcfg.ApricotPhase1BlockTimestamp)
+   }
+   if isForkIncompatible(c.ApricotPhase2BlockTimestamp, newcfg.ApricotPhase2BlockTimestamp, headTimestamp) {
+       return newCompatError("ApricotPhase2 fork block timestamp", c.ApricotPhase2BlockTimestamp, newcfg.ApricotPhase2BlockTimestamp)
+   }
+   if isForkIncompatible(c.ApricotPhase3BlockTimestamp, newcfg.ApricotPhase3BlockTimestamp, headTimestamp) {
+       return newCompatError("ApricotPhase3 fork block timestamp", c.ApricotPhase3BlockTimestamp, newcfg.ApricotPhase3BlockTimestamp)
+   }
+   if isForkIncompatible(c.ApricotPhase4BlockTimestamp, newcfg.ApricotPhase4BlockTimestamp, headTimestamp) {
+       return newCompatError("ApricotPhase4 fork block timestamp", c.ApricotPhase4BlockTimestamp, newcfg.ApricotPhase4BlockTimestamp)
+   }
+   if isForkIncompatible(c.ApricotPhase5BlockTimestamp, newcfg.ApricotPhase5BlockTimestamp, headTimestamp) {
+       return newCompatError("ApricotPhase5 fork block timestamp", c.ApricotPhase5BlockTimestamp, newcfg.ApricotPhase5BlockTimestamp)
+   }
}
return nil
}

@@ -467,10 +520,7 @@ type Rules struct {
    IsByzantium, IsConstantinople, IsPetersburg, IsIstanbul bool

    // Rules for Avalanche releases
-   IsApricotPhase1 bool
-   IsApricotPhase2 bool
-   IsApricotPhase3 bool
-   IsApricotPhase4 bool
+   IsApricotPhase1, IsApricotPhase2, IsApricotPhase3, IsApricotPhase4, IsApricotPhase5 bool
}

// Rules ensures c's ChainID is not nil.
@@ -501,5 +551,6 @@ func (c *ChainConfig) AvalancheRules(blockNum, blockTimestamp *big.Int) Rules {
    rules.IsApricotPhase2 = c.IsApricotPhase2(blockTimestamp)
    rules.IsApricotPhase3 = c.IsApricotPhase3(blockTimestamp)
    rules.IsApricotPhase4 = c.IsApricotPhase4(blockTimestamp)
+   rules.IsApricotPhase5 = c.IsApricotPhase5(blockTimestamp)
    return rules
}

diff --git a/params/config_test.go b/params/config_test.go
new file mode 100644
index 00000000..62fea4be
--- /dev/null
+++ b/params/config_test.go
@@ -0,0 +1,138 @@
// (c) 2019-2020, Ava Labs, Inc.
//
// This file is a derived work, based on the go-ethereum library whose original
// notices appear below.
//
// It is distributed under a license compatible with the licensing terms of the
// original code from which it is derived.
//
// Much love to the original authors for their work.
// *****
// Copyright 2017 The go-ethereum Authors
// This file is part of the go-ethereum library.
//
// The go-ethereum library is free software: you can redistribute it and/or modify
// it under the terms of the GNU Lesser General Public License as published by
// the Free Software Foundation, either version 3 of the License, or
// (at your option) any later version.
//
// The go-ethereum library is distributed in the hope that it will be useful,
// but WITHOUT ANY WARRANTY; without even the implied warranty of
// MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
// GNU Lesser General Public License for more details.
//
// You should have received a copy of the GNU Lesser General Public License
// along with the go-ethereum library. If not, see <http://www.gnu.org/licenses/>.
+
+package params
+
+import (
+   "math/big"
+   "reflect"
+   "testing"
+)
+
+func TestCheckCompatible(t *testing.T) {
+   type test struct {
+       stored, new          *ChainConfig
+       headHeight, headTimestamp uint64
+       wantErr              *ConfigCompatError
+   }
+   tests := []test{
+       {stored: TestChainConfig, new: TestChainConfig, headHeight: 0, headTimestamp: 0, wantErr: nil},
+       {stored: TestChainConfig, new: TestChainConfig, headHeight: 100, headTimestamp: 1000, wantErr: nil},
+       {
+           stored:      &ChainConfig{EIP150Block: big.NewInt(10)},
+           new:         &ChainConfig{EIP150Block: big.NewInt(20)},
+           headHeight:  9,
+           headTimestamp: 90,
+           wantErr:     nil,
+       },
+       {
+           stored:      TestChainConfig,
+           new:         &ChainConfig{HomesteadBlock: nil},
+           headHeight:  3,
+           headTimestamp: 30,
+           wantErr:     &ConfigCompatError{
+               What:      "Homestead fork block",
+               StoredConfig: big.NewInt(0),
+               NewConfig:  nil,
+               RewindTo:   0,
+           },
+       },
+   },
+}

```

```

+         {
+             stored:      TestChainConfig,
+             new:          &ChainConfig{HomesteadBlock: big.NewInt(1)},
+             headHeight:   3,
+             headTimestamp: 30,
+             wantErr: &ConfigCompatError{
+                 What:      "Homestead fork block",
+                 StoredConfig: big.NewInt(0),
+                 NewConfig:   big.NewInt(1),
+                 RewindTo:    0,
+             },
+         },
+         {
+             stored:      &ChainConfig{HomesteadBlock: big.NewInt(30), EIP150Block: big.NewInt(10)},
+             new:          &ChainConfig{HomesteadBlock: big.NewInt(25), EIP150Block: big.NewInt(20)},
+             headHeight:   25,
+             headTimestamp: 250,
+             wantErr: &ConfigCompatError{
+                 What:      "EIP150 fork block",
+                 StoredConfig: big.NewInt(10),
+                 NewConfig:   big.NewInt(20),
+                 RewindTo:    9,
+             },
+         },
+         {
+             stored:      &ChainConfig{ConstantinopleBlock: big.NewInt(30)},
+             new:          &ChainConfig{ConstantinopleBlock: big.NewInt(30), PetersburgBlock: big.NewInt(30)},
+             headHeight:   40,
+             headTimestamp: 400,
+             wantErr:      nil,
+         },
+         {
+             stored:      &ChainConfig{ConstantinopleBlock: big.NewInt(30)},
+             new:          &ChainConfig{ConstantinopleBlock: big.NewInt(30), PetersburgBlock: big.NewInt(31)},
+             headHeight:   40,
+             headTimestamp: 400,
+             wantErr: &ConfigCompatError{
+                 What:      "Petersburg fork block",
+                 StoredConfig: nil,
+                 NewConfig:   big.NewInt(31),
+                 RewindTo:    30,
+             },
+         },
+         {
+             stored:      TestChainConfig,
+             new:          TestApricotPhase4Config,
+             headHeight:   0,
+             headTimestamp: 0,
+             wantErr: &ConfigCompatError{
+                 What:      "ApricotPhase5 fork block timestamp",
+                 StoredConfig: big.NewInt(0),
+                 NewConfig:   nil,
+                 RewindTo:    0,
+             },
+         },
+         {
+             stored:      TestChainConfig,
+             new:          TestApricotPhase4Config,
+             headHeight:   10,
+             headTimestamp: 100,
+             wantErr: &ConfigCompatError{
+                 What:      "ApricotPhase5 fork block timestamp",
+                 StoredConfig: big.NewInt(0),
+                 NewConfig:   nil,
+                 RewindTo:    0,
+             },
+         },
+     },
+ }
+
+ for _, test := range tests {
+     err := test.stored.CheckCompatible(test.new, test.headHeight, test.headTimestamp)
+     if !reflect.DeepEqual(err, test.wantErr) {
+         t.Errorf("error mismatch:\nstored: %v\nnew: %v\nheadHeight: %v\nerr: %v\nwant: %v", test.stored, test.new, test.headHeight, err, test.wantErr)
+     }
+ }
+
+diff --git a/params/protocol_params.go b/params/protocol_params.go
index 00db2bfd..2137387a 100644
--- a/params/protocol_params.go
+++ b/params/protocol_params.go
@@ -29,9 +29,10 @@ package params
 import "math/big"

 const (
-    GasLimitBoundDivisor uint64 = 1024    // The bound divisor of the gas limit, used in update calculations.
-    MinGasLimit           uint64 = 5000    // Minimum the gas limit may ever be.
-    GenesisGasLimit        uint64 = 4712388 // Gas limit of the Genesis block.
+    GasLimitBoundDivisor uint64 = 1024    // The bound divisor of the gas limit, used in update calculations.
+    MinGasLimit           uint64 = 5000    // Minimum the gas limit may ever be.
+    MaxGasLimit           uint64 = 0x7fffffffffffffff // Maximum the gas limit (2^63-1).
+    GenesisGasLimit        uint64 = 4712388 // Gas limit of the Genesis block.

    // Note: MaximumExtraDataSize has been reduced to 32 in Geth, but is kept the same in Coreth for
    // backwards compatibility.
@@ -47,8 +47,8 @@ const (
 LogDataGas      uint64 = 8    // Per byte in a LOG* operation's data.
 CallStipend      uint64 = 2300 // Free gas given at beginning of call.

-    Sha3Gas      uint64 = 30 // Once per SHA3 operation.
-    Sha3WordGas  uint64 = 6  // Once per word of the SHA3 operation's data.
+    Keccak256Gas  uint64 = 30 // Once per KECCAK256 operation.
+    Keccak256WordGas uint64 = 6 // Once per word of the KECCAK256 operation's data.

    SstoreSetGas  uint64 = 20000 // Once per SSTORE operation.
    SstoreResetGas uint64 = 5000 // Once per SSTORE operation if the zeroness changes from zero.
@@ -130,8 +130,8 @@ const (
 // Introduced in Tangerine Whistle (Eip 150)
 CreateBySelfdestructGas uint64 = 25000

-    BaseFeeChangeDenominator = 12 // Bounds the amount the base fee can change between blocks.
-
    MaxCodeSize = 24576 // Maximum bytecode to permit for a contract

    // Precompiled contract gas prices
diff --git a/params/version.go b/params/version.go
index 22334a6a..a6df0da8 100644
--- a/params/version.go
+++ b/params/version.go
@@ -33,7 +33,7 @@ import (
 const (
     VersionMajor = 1    // Major version component of the current release
     VersionMinor = 10   // Minor version component of the current release
-    VersionPatch = 12   // Patch version component of the current release
+    VersionPatch = 15   // Patch version component of the current release
     VersionMeta  = "stable" // Version metadata to append to the version string
 )

diff --git a/peer/client.go b/peer/client.go
new file mode 100644
index 00000000..6b46f7ca
--- /dev/null
+++ b/peer/client.go
@@ -0,0 +1,52 @@
+// (c) 2019-2022, Ava Labs, Inc. All rights reserved.

```

```

+// See the file LICENSE for licensing terms.
+
+package peer
+
+import (
+    "github.com/flare-foundation/flare/version"
+)
+
+var _ Client = &client{}
+
+// Client defines ability to send request / response through the Network
+type Client interface {
+    // RequestAny synchronously sends request to the first connected peer that matches the specified minVersion in
+    // random order.
+    // A peer is considered a match if its version is greater than or equal to the specified minVersion
+    // Returns errNoPeersMatchingVersion if no peer could be found matching specified version
+    RequestAny(minVersion version.Application, request []byte) ([]byte, bool, error)
+
+    // Gossip sends given gossip message to peers
+    Gossip(gossip []byte) error
+}
+
+// client implements Client interface
+// provides ability to send request / responses through the Network
+type client struct {
+    network Network
+}
+
+// RequestAny synchronously sends request to the first connected peer that matches the specified minVersion in
+// random order.
+// Returns response bytes, whether the request failed and optional error
+// Returns errNoPeersMatchingVersion if no peer could be found matching specified version
+// This function is blocks until a response is received from the peer
+func (c *client) RequestAny(minVersion version.Application, request []byte) ([]byte, bool, error) {
+    waitingHandler := newWaitingResponseHandler()
+    if err := c.network.RequestAny(minVersion, request, waitingHandler); err != nil {
+        return nil, true, err
+    }
+    return <-waitingHandler.responseChan, waitingHandler.failed, nil
+}
+
+func (c *client) Gossip(gossip []byte) error {
+    return c.network.Gossip(gossip)
+}
+
+// NewClient returns Client for a given network
+func NewClient(network Network) Client {
+    return &client{
+        network: network,
+    }
+}
+
+diff --git a/peer/network.go b/peer/network.go
+new file mode 100644
+index 00000000..700f4bb2
+--- /dev/null
++++ b/peer/network.go
+@@ -0,0 +1,343 @@
+// (c) 2019-2022, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package peer
+
+import (
+    "context"
+    "errors"
+    "fmt"
+    "sync"
+    "time"
+
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+
+    "github.com/flare-foundation/flare/snow/validators"
+
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/snow/engine/common"
+    "github.com/flare-foundation/flare/version"
+    "golang.org/x/sync/semaphore"
+)
+
+// Minimum amount of time to handle a request
+const minRequestHandlingDuration = 100 * time.Millisecond
+
+var (
+    errAcquiringSemaphore = errors.New("error acquiring semaphore")
+    _ Network              = &network{}
+    _ validators.Connector = &network{}
+    _ common.AppHandler    = &network{}
+)
+
+type Network interface {
+    validators.Connector
+    common.AppHandler
+
+    // RequestAny synchronously sends request to the first connected peer that matches the specified minVersion in
+    // random order.
+    // A peer is considered a match if its version is greater than or equal to the specified minVersion
+    // Returns errNoPeersMatchingVersion if no peer could be found matching specified version
+    RequestAny(minVersion version.Application, message []byte, handler message.ResponseHandler) error
+
+    // Gossip sends given gossip message to peers
+    Gossip(gossip []byte) error
+
+    // Shutdown stops all peer channel listeners and marks the node to have stopped
+    // n.Start() can be called again but the peers will have to be reconnected
+    // by calling OnPeerConnected for each peer
+    Shutdown()
+
+    // SetGossipHandler sets the provided gossip handler as the gossip handler
+    SetGossipHandler(handler message.GossipHandler)
+
+    // SetRequestHandler sets the provided request handler as the request handler
+    SetRequestHandler(handler message.RequestHandler)
+
+    // Size returns the size of the network in number of connected peers
+    Size() uint32
+}
+
+// network is an implementation of Network that processes message requests for
+// each peer in linear fashion
+type network struct {
+    lock sync.RWMutex // lock for mutating state of this Network struct
+    self ids.ShortID  // NodeID of this node
+    requestIDGen uint32 // requestID counter used to track outbound requests
+    outstandingResponseHandlerMap map[uint32]message.ResponseHandler // maps avalanche requestID => response handler
+    activeRequests *semaphore.Weighted // controls maximum number of active outbound requests
+    appSender common.AppSender // avalanche AppSender for sending messages
+    codec codec.Manager // Codec used for parsing messages
+    requestHandler message.RequestHandler // maps request type => handler
+    gossipHandler message.GossipHandler // maps gossip type => handler
+    peers map[ids.ShortID]version.Application // maps nodeID => version.Version
+}

```

```

+func NewNetwork(appSender common.AppSender, codec codec.Manager, self ids.ShortID, maxActiveRequests int64) Network {
+    return &network{
+        appSender:      appSender,
+        codec:           codec,
+        self:            self,
+        outstandingResponseHandlerMap: make(map[uint32]message.ResponseHandler),
+        peers:           make(map[ids.ShortID]version.Application),
+        activeRequests:   semaphore.NewWeighted(maxActiveRequests),
+    }
+}
+
+// RequestAny sends given request to the first connected peer that matches the specified minVersion
+// A peer is considered a match if its version is greater than or equal to the specified minVersion
+// If minVersion is nil, then the request will be sent to any peer regardless of their version
+// Returns a non-nil error if we were not able to send a request to a peer with >= [minVersion]
+// or we fail to send a request to the selected peer.
+func (n *network) RequestAny(minVersion version.Application, request []byte, handler message.ResponseHandler) error {
+    // Take a slot from total [activeRequests] and block until a slot becomes available.
+    if err := n.activeRequests.Acquire(context.Background(), 1); err != nil {
+        return errAcquiringSemaphore
+    }
+
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    for nodeID, nodeVersion := range n.peers {
+        // map iteration is sufficiently random to avoid hitting same peer so here
+        // we get a random peerID key that we compare minimum version that
+        // we have
+        if minVersion == nil || nodeVersion.Compare(minVersion) >= 0 {
+            return n.request(nodeID, request, handler)
+        }
+    }
+
+    n.activeRequests.Release(1)
+    return fmt.Errorf("no peers found matching version %s out of %d peers", minVersion, len(n.peers))
+}
+
+// Request sends request message bytes to specified nodeID and adds [responseHandler] to [outstandingResponseHandlerMap]
+// so that it can be invoked when the network receives either a response or failure message.
+// Assumes [nodeID] is never [self] since we guarantee [self] will not be added to the [peers] map.
+// Returns an error if [appSender] is unable to make the request.
+// Assumes write lock is held
+func (n *network) request(nodeID ids.ShortID, request []byte, responseHandler message.ResponseHandler) error {
+    log.Debug("sending request to peer", "nodeID", nodeID, "requestLen", len(request))
+
+    // generate requestID
+    requestID := n.requestIDGen
+    n.requestIDGen++
+
+    n.outstandingResponseHandlerMap[requestID] = responseHandler
+
+    nodeIDs := ids.NewShortSet(1)
+    nodeIDs.Add(nodeID)
+
+    // send app request to the peer
+    // on failure: release the activeRequests slot, mark message as processed and return fatal error
+    // Send app request to [nodeID].
+    // On failure, release the slot from active requests and [outstandingResponseHandlerMap].
+    if err := n.appSender.SendAppRequest(nodeIDs, requestID, request); err != nil {
+        n.activeRequests.Release(1)
+        delete(n.outstandingResponseHandlerMap, requestID)
+        log.Error("could not send app message", "err", err, "nodeID", nodeID, "requestID", requestID)
+        return err
+    }
+
+    log.Debug("sent request message to peer", "nodeID", nodeID, "requestID", requestID)
+    return nil
+}
+
+// AppRequest is called by avalancheGo -> VM when there is an incoming AppRequest from a peer
+// error returned by this function is expected to be treated as fatal by the engine
+// returns error if the requestHandler returns an error
+// sends a response back to the sender if length of response returned by the handler is >0
+// expects the deadline to not have been passed
+func (n *network) AppRequest(nodeID ids.ShortID, requestID uint32, deadline time.Time, request []byte) error {
+    n.lock.RLock()
+    defer n.lock.RUnlock()
+
+    log.Debug("received AppRequest from node", "nodeID", nodeID, "requestID", requestID, "requestLen", len(request))
+
+    var req message.Request
+    if _, err := n.codec.Unmarshal(request, &req); err != nil {
+        log.Debug("failed to unmarshal app request", "nodeID", nodeID, "requestID", requestID, "requestLen", len(request), "err", err)
+        return nil
+    }
+
+    // calculate how much time is left until the deadline
+    timeTillDeadline := time.Until(deadline)
+
+    // bufferedDeadline is half the time till actual deadline so that the message has a reasonable chance
+    // of completing its processing and sending the response to the peer.
+    timeTillDeadline = time.Duration(timeTillDeadline.Nanoseconds() / 2)
+    bufferedDeadline := time.Now().Add(timeTillDeadline)
+
+    // check if we have enough time to handle this request
+    if time.Until(bufferedDeadline) < minRequestHandlingDuration {
+        // Drop the request if we already missed the deadline to respond.
+        log.Debug("deadline to process AppRequest has expired, skipping", "nodeID", nodeID, "requestID", requestID, "type", req.Type())
+        return nil
+    }
+
+    log.Debug("processing incoming request", "nodeID", nodeID, "requestID", requestID, "type", req.Type())
+    ctx, cancel := context.WithDeadline(context.Background(), bufferedDeadline)
+    defer cancel()
+
+    responseBytes, err := req.Handle(ctx, nodeID, requestID, n.requestHandler)
+    switch {
+    case err != nil && err != context.DeadlineExceeded:
+        return err // Return a fatal error
+    case responseBytes != nil:
+        return n.appSender.SendAppResponse(nodeID, requestID, responseBytes) // Propagate fatal error
+    default:
+        return nil
+    }
+}
+
+// AppResponse is invoked when there is a response received from a peer regarding a request
+// Error returned by this function is expected to be treated as fatal by the engine
+// If [requestID] is not known, this function will emit a log and return a nil error.
+// If the response handler returns an error it is propagated as a fatal error.
+func (n *network) AppResponse(nodeID ids.ShortID, requestID uint32, response []byte) error {
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    log.Debug("received AppResponse from peer", "nodeID", nodeID, "requestID", requestID)
+
+    handler, exists := n.getRequestHandler(requestID)
+    if !exists {
+        // Should never happen since the engine should be managing outstanding requests
+        log.Error("received response to unknown request", "nodeID", nodeID, "requestID", requestID, "responseLen", len(response))
+        return nil
+    }
+}

```

```

+         return handler.OnResponse(nodeID, requestID, response)
+}
+
+// AppRequestFailed can be called by the avalanche -> VM in following cases:
+// - node is benched
+// - failed to send message to [nodeID] due to a network issue
+// - timeout
+// error returned by this function is expected to be treated as fatal by the engine
+// returns error only when the response handler returns an error
+func (n *network) AppRequestFailed(nodeID ids.ShortID, requestID uint32) error {
+    n.lock.Lock()
+    defer n.lock.Unlock()
+    log.Debug("received AppRequestFailed from peer", "nodeID", nodeID, "requestID", requestID)
+
+    handler, exists := n.getRequestHandler(requestID)
+    if !exists {
+        // Should never happen since the engine should be managing outstanding requests
+        log.Error("received request failed to unknown request", "nodeID", nodeID, "requestID", requestID)
+        return nil
+    }
+
+    return handler.OnFailure(nodeID, requestID)
+}
+
+// getHandler fetches the handler for [requestID] and marks the request with [requestID] as having been fulfilled.
+// This is called by either [AppResponse] or [AppRequestFailed].
+// assumes that the write lock is held.
+func (n *network) getHandler(requestID uint32) (message.ResponseHandler, bool) {
+    handler, exists := n.outstandingResponseHandlerMap[requestID]
+    if !exists {
+        return nil, false
+    }
+    // mark message as processed, release activeRequests slot
+    delete(n.outstandingResponseHandlerMap, requestID)
+    n.activeRequests.Release(1)
+    return handler, true
+}
+
+// Gossip sends given gossip message to peers
+func (n *network) Gossip(gossip []byte) error {
+    return n.appSender.SendAppGossip(gossip)
+}
+
+// AppGossip is called by avalanche -> VM when there is an incoming AppGossip from a peer
+// error returned by this function is expected to be treated as fatal by the engine
+// returns error if request could not be parsed as message.Request or when the requestHandler returns an error
+func (n *network) AppGossip(nodeID ids.ShortID, gossipBytes []byte) error {
+    var gossipMsg message.Message
+    if _, err := n.codec.Unmarshal(gossipBytes, &gossipMsg); err != nil {
+        log.Debug("could not parse app gossip", "nodeID", nodeID, "gossipLen", len(gossipBytes), "err", err)
+        return nil
+    }
+
+    log.Debug("processing AppGossip from node", "nodeID", nodeID, "type", gossipMsg.Type(), "gossipLen", len(gossipBytes))
+    return gossipMsg.Handle(n.gossipHandler, nodeID)
+}
+
+// Connected adds the given nodeID to the peer list so that it can receive messages
+func (n *network) Connected(nodeID ids.ShortID, nodeVersion version.Application) error {
+    log.Debug("adding new peer", "nodeID", nodeID)
+
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    if nodeID == n.self {
+        log.Debug("skipping registering self as peer")
+        return nil
+    }
+
+    if storedVersion, exists := n.peers[nodeID]; exists {
+        // Peer is already connected, update the version if it has changed.
+        // Log a warning message since the consensus engine should never call Connected on a peer
+        // that we have already marked as Connected.
+        if nodeVersion.Compare(storedVersion) != 0 {
+            n.peers[nodeID] = nodeVersion
+            log.Warn("received Connected message for already connected peer, updating node version", "nodeID", nodeID, "storedVersion", storedVersion, "nodeVersion", nodeVersion)
+        } else {
+            log.Warn("ignoring peer connected event for already connected peer with identical version", "nodeID", nodeID)
+        }
+        return nil
+    }
+
+    n.peers[nodeID] = nodeVersion
+    return nil
+}
+
+// Disconnected removes given [nodeID] from the peer list
+func (n *network) Disconnected(nodeID ids.ShortID) error {
+    log.Debug("disconnecting peer", "nodeID", nodeID)
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    // if this peer already exists, log a warning and ignore the request
+    if _, exists := n.peers[nodeID]; !exists {
+        // we're not connected to this peer, nothing to do here
+        log.Warn("received peer disconnect request to unconnected peer", "nodeID", nodeID)
+        return nil
+    }
+
+    delete(n.peers, nodeID)
+    return nil
+}
+
+// Shutdown disconnects all peers
+func (n *network) Shutdown() {
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    // reset peers map
+    n.peers = make(map[ids.ShortID]version.Application)
+}
+
+func (n *network) SetGossipHandler(handler message.GossipHandler) {
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    n.gossipHandler = handler
+}
+
+func (n *network) SetRequestHandler(handler message.RequestHandler) {
+    n.lock.Lock()
+    defer n.lock.Unlock()
+
+    n.requestHandler = handler
+}
+
+func (n *network) Size() uint32 {
+    n.lock.RLock()
+    defer n.lock.RUnlock()
+
+    return uint32(len(n.peers))
+}
+
diff --git a/peer/network_test.go b/peer/network_test.go
new file mode 100644

```

```

index 00000000..d6f8a072
--- /dev/null
+++ b/peer/network_test.go
@@ -0,0 +1,484 @@
+// (c) 2019-2022, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package peer
+
+import (
+    "context"
+    "errors"
+    "fmt"
+    "sync"
+    "sync/atomic"
+    "testing"
+    "time"
+
+    "github.com/flare-foundation/flare/snow/engine/common"
+
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/codec/linearcodec"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/version"
+    "github.com/stretchr/testify/assert"
+)
+
+var (
+    defaultPeerVersion = version.NewDefaultApplication("corethtest", 1, 0, 0)
+
+    _ message.Request = &HelloRequest{}
+    _                 = &HelloResponse{}
+    _                 = &GreetingRequest{}
+    _                 = &GreetingResponse{}
+    _                 = &TestMessage{}
+
+    _ message.RequestHandler = &HelloGreetingRequestHandler{}
+    _ message.RequestHandler = &testRequestHandler{}
+
+    _ common.AppSender    = testAppSender{}
+    _ message.Message     = HelloGossip{}
+    _ message.GossipHandler = &testGossipHandler{}
+)
+
+func TestNetworkDoesNotConnectToItself(t *testing.T) {
+    selfNodeID := ids.GenerateTestShortID()
+    n := NewNetwork(nil, nil, selfNodeID, 1)
+    assert.NoError(t, n.Connected(selfNodeID, version.NewDefaultApplication("avalanchego", 1, 0, 0)))
+    assert.EqualValues(t, 0, n.Size())
+}
+
+func TestRequestsRoutingAndResponse(t *testing.T) {
+    callNum := uint32(0)
+    senderWg := &sync.WaitGroup{}
+    var net Network
+    sender := testAppSender{
+        sendAppRequestFn: func(nodes ids.ShortSet, requestID uint32, requestBytes []byte) error {
+            nodeID, _ := nodes.Pop()
+            senderWg.Add(1)
+            go func() {
+                defer senderWg.Done()
+                if err := net.AppRequest(nodeID, requestID, time.Now().Add(5*time.Second), requestBytes); err != nil {
+                    panic(err)
+                }
+            }()
+            return nil
+        },
+        sendAppResponseFn: func(nodeID ids.ShortID, requestID uint32, responseBytes []byte) error {
+            senderWg.Add(1)
+            go func() {
+                defer senderWg.Done()
+                if err := net.AppResponse(nodeID, requestID, responseBytes); err != nil {
+                    panic(err)
+                }
+            }()
+            atomic.AddUint32(&callNum, 1)
+            return nil
+        },
+    },
+}
+
+    codecManager := buildCodec(t, HelloRequest{}, HelloResponse{})
+    net = NewNetwork(sender, codecManager, ids.ShortEmpty, 16)
+    net.SetRequestHandler(&HelloGreetingRequestHandler{codec: codecManager})
+    client := NewClient(net)
+    nodeID := ids.GenerateTestShortID()
+    assert.NoError(t, net.Connected(nodeID, defaultPeerVersion))
+
+    requestMessage := HelloRequest{Message: "this is a request"}
+
+    defer net.Shutdown()
+    assert.NoError(t, net.Connected(nodeID, defaultPeerVersion))
+
+    totalRequests := 5000
+    numCallsPerRequest := 1 // on sending response
+    totalCalls := totalRequests * numCallsPerRequest
+
+    requestWg := &sync.WaitGroup{}
+    requestWg.Add(totalCalls)
+    for i := 0; i < totalCalls; i++ {
+        go func(wg *sync.WaitGroup) {
+            defer wg.Done()
+            requestBytes, err := message.RequestToBytes(codecManager, requestMessage)
+            assert.NoError(t, err)
+            responseBytes, failed, err := client.RequestAny(defaultPeerVersion, requestBytes)
+            assert.NoError(t, err)
+            assert.False(t, failed)
+            assert.NotNil(t, responseBytes)
+
+            var response TestMessage
+            if _, err = codecManager.Unmarshal(responseBytes, &response); err != nil {
+                panic(fmt.Errorf("unexpected error during unmarshal: %w", err))
+            }
+            assert.Equal(t, "Hi", response.Message)
+        }(requestWg)
+    }
+
+    requestWg.Wait()
+    senderWg.Wait()
+    assert.Equal(t, totalCalls, int(atomic.LoadUint32(&callNum)))
+}
+
+func TestRequestMinVersion(t *testing.T) {
+    callNum := uint32(0)
+    nodeID := ids.GenerateTestShortID()
+    codecManager := buildCodec(t, TestMessage{})
+
+    var net Network
+    sender := testAppSender{
+        sendAppRequestFn: func(nodes ids.ShortSet, reqID uint32, messageBytes []byte) error {
+            atomic.AddUint32(&callNum, 1)
+            assert.True(t, nodes.Contains(nodeID), "request nodes should contain expected nodeID")
+            assert.Len(t, nodes, 1, "request nodes should contain exactly one node")

```

```

+
+         go func() {
+             time.Sleep(200 * time.Millisecond)
+             atomic.AddUint32(&callNum, 1)
+             responseBytes, err := codecManager.Marshal(message.Version, TestMessage{Message: "this is a response"})
+             if err != nil {
+                 panic(err)
+             }
+             err = net.AppResponse(nodeID, reqID, responseBytes)
+             assert.NoError(t, err)
+         }()
+         return nil
+     },
+ }
+
+ // passing nil as codec works because the net.AppRequest is never called
+ net = NewNetwork(sender, codecManager, ids.ShortEmpty, 1)
+ client := NewClient(net)
+ requestMessage := TestMessage{Message: "this is a request"}
+ requestBytes, err := message.RequestToBytes(codecManager, requestMessage)
+ assert.NoError(t, err)
+ assert.NoError(t, net.Connected(nodeID, version.NewDefaultApplication("corethtest", 1, 7, 1)))
+
+ // ensure version does not match
+ responseBytes, failed, err := client.RequestAny(version.NewDefaultApplication("corethtest", 2, 0, 0), requestBytes)
+ assert.Equal(t, err.Error(), "no peers found matching version corethtest/2.0.0 out of 1 peers")
+ assert.True(t, failed)
+ assert.Nil(t, responseBytes)
+
+ // ensure version matches and the request goes through
+ responseBytes, failed, err = client.RequestAny(version.NewDefaultApplication("corethtest", 1, 0, 0), requestBytes)
+ assert.NoError(t, err)
+ assert.False(t, failed)
+
+ var response TestMessage
+ if _, err = codecManager.Unmarshal(responseBytes, &response); err != nil {
+     t.Fatal("unexpected error during unmarshal", err)
+ }
+ assert.Equal(t, "this is a response", response.Message)
+}
+
+func TestOnRequestHonoursDeadline(t *testing.T) {
+    var net Network
+    responded := false
+    sender := testAppSender{
+        sendAppRequestFn: func(nodes ids.ShortSet, reqID uint32, message []byte) error {
+            return nil
+        },
+        sendAppResponseFn: func(nodeID ids.ShortID, reqID uint32, message []byte) error {
+            responded = true
+            return nil
+        },
+    }
+
+    codecManager := buildCodec(t, TestMessage{})
+
+    requestBytes, err := marshalStruct(codecManager, TestMessage{Message: "hello there"})
+    assert.NoError(t, err)
+
+    requestHandler := &testRequestHandler{
+        processingDuration: 500 * time.Millisecond,
+    }
+    net = NewNetwork(sender, codecManager, ids.ShortEmpty, 1)
+    net.SetRequestHandler(requestHandler)
+    nodeID := ids.GenerateTestShortID()
+
+    requestHandler.response, err = marshalStruct(codecManager, TestMessage{Message: "hi there"})
+    assert.NoError(t, err)
+    err = net.AppRequest(nodeID, 1, time.Now().Add(1*time.Millisecond), requestBytes)
+    assert.NoError(t, err)
+    // ensure the handler didn't get called (as peer.Network would've dropped the request)
+    assert.EqualValues(t, requestHandler.calls, 0)
+
+    requestHandler.processingDuration = 0
+    err = net.AppRequest(nodeID, 2, time.Now().Add(250*time.Millisecond), requestBytes)
+    assert.NoError(t, err)
+    assert.True(t, responded)
+    assert.EqualValues(t, requestHandler.calls, 1)
+}
+
+func TestGossip(t *testing.T) {
+    codecManager := buildCodec(t, HelloGossip{})
+
+    nodeID := ids.GenerateTestShortID()
+    var clientNetwork Network
+    wg := &sync.WaitGroup{}
+    sentGossip := false
+    wg.Add(1)
+    sender := testAppSender{
+        sendAppGossipFn: func(msg []byte) error {
+            go func() {
+                defer wg.Done()
+                err := clientNetwork.AppGossip(nodeID, msg)
+                assert.NoError(t, err)
+            }()
+            sentGossip = true
+            return nil
+        },
+    }
+
+    gossipHandler := &testGossipHandler{}
+    clientNetwork = NewNetwork(sender, codecManager, ids.ShortEmpty, 1)
+    clientNetwork.SetGossipHandler(gossipHandler)
+
+    assert.NoError(t, clientNetwork.Connected(nodeID, defaultPeerVersion))
+
+    client := NewClient(clientNetwork)
+    defer clientNetwork.Shutdown()
+
+    b, err := buildGossip(codecManager, HelloGossip{Msg: "hello there!"})
+    assert.NoError(t, err)
+
+    err = client.Gossip(b)
+    assert.NoError(t, err)
+
+    wg.Wait()
+    assert.True(t, sentGossip)
+    assert.True(t, gossipHandler.received)
+}
+
+func TestHandleInvalidMessages(t *testing.T) {
+    codecManager := buildCodec(t, HelloGossip{}, TestMessage{})
+
+    nodeID := ids.GenerateTestShortID()
+    requestID := uint32(1)
+    sender := testAppSender{
+
+    clientNetwork := NewNetwork(sender, codecManager, ids.ShortEmpty, 1)
+    clientNetwork.SetGossipHandler(message.NoopMempoolGossipHandler{})
+    clientNetwork.SetRequestHandler(&testRequestHandler{})
+
+    assert.NoError(t, clientNetwork.Connected(nodeID, defaultPeerVersion))
+
+    defer clientNetwork.Shutdown()

```

```

+
+ // Ensure a valid gossip message sent as any App specific message type does not trigger a fatal error
+ gossipMsg, err := buildGossip(codecManager, HelloGossip{Msg: "hello there!"})
+ assert.NoError(t, err)
+
+ // Ensure a valid request message sent as any App specific message type does not trigger a fatal error
+ requestMessage, err := marshalStruct(codecManager, TestMessage{Message: "Hello"})
+ assert.NoError(t, err)
+
+ // Ensure a random message sent as any App specific message type does not trigger a fatal error
+ garbageResponse := make([]byte, 10)
+ // Ensure a zero-length message sent as any App specific message type does not trigger a fatal error
+ emptyResponse := make([]byte, 0)
+ // Ensure a nil byte slice sent as any App specific message type does not trigger a fatal error
+ var nilResponse []byte
+
+ // Check for edge cases
+ assert.NoError(t, clientNetwork.AppGossip(nodeID, gossipMsg))
+ assert.NoError(t, clientNetwork.AppGossip(nodeID, requestMessage))
+ assert.NoError(t, clientNetwork.AppGossip(nodeID, garbageResponse))
+ assert.NoError(t, clientNetwork.AppGossip(nodeID, emptyResponse))
+ assert.NoError(t, clientNetwork.AppGossip(nodeID, nilResponse))
+ assert.NoError(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), gossipMsg))
+ assert.NoError(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), requestMessage))
+ assert.NoError(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), garbageResponse))
+ assert.NoError(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), emptyResponse))
+ assert.NoError(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), nilResponse))
+ assert.NoError(t, clientNetwork.AppResponse(nodeID, requestID, gossipMsg))
+ assert.NoError(t, clientNetwork.AppResponse(nodeID, requestID, requestMessage))
+ assert.NoError(t, clientNetwork.AppResponse(nodeID, requestID, garbageResponse))
+ assert.NoError(t, clientNetwork.AppResponse(nodeID, requestID, emptyResponse))
+ assert.NoError(t, clientNetwork.AppResponse(nodeID, requestID, nilResponse))
+ assert.NoError(t, clientNetwork.AppRequestFailed(nodeID, requestID))
+}
+
+func TestNetworkPropagatesRequestHandlerError(t *testing.T) {
+    codecManager := buildCodec(t, TestMessage{})
+
+    nodeID := ids.GenerateTestShortID()
+    requestID := uint32(1)
+    sender := testAppSender{}
+
+    clientNetwork := NewNetwork(sender, codecManager, ids.ShortEmpty, 1)
+    clientNetwork.SetGossipHandler(message.NoopMemPoolGossipHandler{})
+    clientNetwork.SetRequestHandler(&testRequestHandler{err: errors.New("fail")}) // Return an error from the request handler
+
+    assert.NoError(t, clientNetwork.Connected(nodeID, defaultPeerVersion))
+
+    defer clientNetwork.Shutdown()
+
+    // Ensure a valid request message sent as any App specific message type does not trigger a fatal error
+    requestMessage, err := marshalStruct(codecManager, TestMessage{Message: "Hello"})
+    assert.NoError(t, err)
+
+    // Check that if the request handler returns an error, it is propagated as a fatal error.
+    assert.Error(t, clientNetwork.AppRequest(nodeID, requestID, time.Now().Add(time.Second), requestMessage))
+}
+
+func buildCodec(t *testing.T, types ...interface{}) codec.Manager {
+    codecManager := codec.NewDefaultManager()
+    c := linearcodec.NewDefault()
+    for _, typ := range types {
+        assert.NoError(t, c.RegisterType(typ))
+    }
+    assert.NoError(t, codecManager.RegisterCodec(message.Version, c))
+    return codecManager
+}
+
+// marshalStruct is a helper method used to marshal an object as `interface{}`
+// so that the codec is able to include the TypeID in the resulting bytes
+func marshalStruct(codec codec.Manager, obj interface{}) ([]byte, error) {
+    return codec.Marshal(message.Version, &obj)
+}
+
+func buildGossip(codec codec.Manager, msg message.Message) ([]byte, error) {
+    return codec.Marshal(message.Version, &msg)
+}
+
+type testAppSender struct {
+    sendAppRequestFn func(ids.ShortSet, uint32, []byte) error
+    sendAppResponseFn func(ids.ShortID, uint32, []byte) error
+    sendAppGossipFn func([]byte) error
+}
+
+func (t testAppSender) SendAppGossipSpecific(ids.ShortSet, []byte) error {
+    panic("not implemented")
+}
+
+func (t testAppSender) SendAppRequest(nodeIDs ids.ShortSet, requestID uint32, message []byte) error {
+    return t.sendAppRequestFn(nodeIDs, requestID, message)
+}
+
+func (t testAppSender) SendAppResponse(nodeID ids.ShortID, requestID uint32, message []byte) error {
+    return t.sendAppResponseFn(nodeID, requestID, message)
+}
+
+func (t testAppSender) SendAppGossip(message []byte) error {
+    return t.sendAppGossipFn(message)
+}
+
+type HelloRequest struct {
+    Message string `serialize:"true"`
+}
+
+func (h HelloRequest) Handle(ctx context.Context, nodeID ids.ShortID, requestID uint32, handler message.RequestHandler) ([]byte, error) {
+    // casting is only necessary for test since RequestHandler does not implement anything at the moment
+    return handler.(TestRequestHandler).HandleHelloRequest(ctx, nodeID, requestID, &h)
+}
+
+func (h HelloRequest) Type() string {
+    return "hello-request"
+}
+
+type GreetingRequest struct {
+    Greeting string `serialize:"true"`
+}
+
+func (g GreetingRequest) Handle(ctx context.Context, nodeID ids.ShortID, requestID uint32, handler message.RequestHandler) ([]byte, error) {
+    // casting is only necessary for test since RequestHandler does not implement anything at the moment
+    return handler.(TestRequestHandler).HandleGreetingRequest(ctx, nodeID, requestID, &g)
+}
+
+func (g GreetingRequest) Type() string {
+    return "greeting-request"
+}
+
+type HelloResponse struct {
+    Response string `serialize:"true"`
+}
+
+type GreetingResponse struct {
+    Greet string `serialize:"true"`
+}
+

```



```

+type TestRequestHandler interface {
+    HandleHelloRequest(ctx context.Context, nodeID ids.ShortID, requestID uint32, request *HelloRequest) ([]byte, error)
+    HandleGreetingRequest(ctx context.Context, nodeID ids.ShortID, requestID uint32, request *GreetingRequest) ([]byte, error)
+}
+
+type HelloGreetingRequestHandler struct {
+    codec codec.Manager
+}
+
+func (h *HelloGreetingRequestHandler) HandleHelloRequest(ctx context.Context, nodeID ids.ShortID, requestID uint32, request *HelloRequest) ([]byte, error) {
+    return h.codec.Marshal(message.Version, HelloResponse{Response: "Hi"})
+}
+
+func (h *HelloGreetingRequestHandler) HandleGreetingRequest(ctx context.Context, nodeID ids.ShortID, requestID uint32, request *GreetingRequest) ([]byte, error) {
+    return h.codec.Marshal(message.Version, GreetingResponse{Greet: "Hey there"})
+}
+
+type TestMessage struct {
+    Message string `serialize:"true"`
+}
+
+func (t TestMessage) Handle(ctx context.Context, nodeID ids.ShortID, requestID uint32, handler message.RequestHandler) ([]byte, error) {
+    return handler.(*testRequestHandler).handleTestRequest(ctx, nodeID, requestID, &t)
+}
+
+func (t TestMessage) Type() string {
+    return "test-message"
+}
+
+type HelloGossip struct {
+    message.Message
+    Msg string `serialize:"true"`
+}
+
+func (h HelloGossip) Handle(handler message.GossipHandler, nodeID ids.ShortID) error {
+    return handler.HandleEthTx(nodeID, nil)
+}
+
+func (h HelloGossip) Type() string {
+    return "hello-gossip"
+}
+
+func (h HelloGossip) initialize(_ []byte) {
+    // no op
+}
+
+func (h HelloGossip) Bytes() []byte {
+    // no op
+    return nil
+}
+
+type testGossipHandler struct {
+    received bool
+    nodeID    ids.ShortID
+    msg       []byte
+}
+
+func (t *testGossipHandler) HandleAtomicTx(nodeID ids.ShortID, _ *message.AtomicTx) error {
+    t.received = true
+    t.nodeID = nodeID
+    return nil
+}
+
+func (t *testGossipHandler) HandleEthTx(nodeID ids.ShortID, _ *message.EthTx) error {
+    t.received = true
+    t.nodeID = nodeID
+    return nil
+}
+
+type testRequestHandler struct {
+    calls      uint32
+    processingDuration time.Duration
+    response    []byte
+    err         error
+}
+
+func (r *testRequestHandler) handleTestRequest(ctx context.Context, _ ids.ShortID, _ uint32, _ *TestMessage) ([]byte, error) {
+    r.calls++
+    select {
+    case <-time.After(r.processingDuration):
+        break
+    case <-ctx.Done():
+        return nil, ctx.Err()
+    }
+    return r.response, r.err
+}
diff --git a/peer/waiting_handler.go b/peer/waiting_handler.go
new file mode 100644
index 00000000..b656dff4
--- /dev/null
+++ b/peer/waiting_handler.go
@@ -0,0 +1,39 @@
+// (c) 2019-2022, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package peer
+
+import (
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+    "github.com/flare-foundation/flare/ids"
+)
+
+var _ message.ResponseHandler = &waitingResponseHandler{}
+
+// waitingResponseHandler implements the ResponseHandler interface
+// Internally used to wait for response after making a request synchronously
+// responseChan may contain response bytes if the original request has not failed
+// responseChan is closed in either fail or success scenario
+type waitingResponseHandler struct {
+    responseChan chan []byte // blocking channel with response bytes
+    failed       bool      // whether the original request is failed
+}
+
+// OnResponse passes the response bytes to the responseChan and closes the channel
+func (w *waitingResponseHandler) OnResponse(_ ids.ShortID, _ uint32, response []byte) error {
+    w.responseChan <- response
+    close(w.responseChan)
+    return nil
+}
+
+// OnFailure sets the failed flag to true and closes the channel
+func (w *waitingResponseHandler) OnFailure(ids.ShortID, uint32) error {
+    w.failed = true
+    close(w.responseChan)
+    return nil
+}
+
+// newWaitingResponseHandler returns new instance of the waitingResponseHandler
+func newWaitingResponseHandler() *waitingResponseHandler {
+    return &waitingResponseHandler{responseChan: make(chan []byte)}
+}
diff --git a/plugin/evm/admin.go b/plugin/evm/admin.go
index 0cdeb2ad..b9418795 100644
--- a/plugin/evm/admin.go

```

```

+++ b/plugin/evm/admin.go
@@ -7,9 +7,9 @@ import (
    "fmt"
    "net/http"

-    "github.com/ava-labs/avalanchego/api"
-    "github.com/ava-labs/avalanchego/utls/profiler"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/flare/api"
+    "github.com/flare-foundation/flare/utls/profiler"
)

// Admin is the API service for admin API calls
diff --git a/plugin/evm/atomic_trie.go b/plugin/evm/atomic_trie.go
new file mode 100644
index 00000000..44c2531c
--- /dev/null
+++ b/plugin/evm/atomic_trie.go
@@ -0,0 +1,437 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "encoding/binary"
+    "fmt"
+    "time"
+
+    "github.com/flare-foundation/flare/database"
+    "github.com/flare-foundation/flare/database/prefixdb"
+    "github.com/flare-foundation/flare/database/versiondb"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/trie"
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utls/wrappers"
+)
+
+const (
+    commitHeightInterval = uint64(4096)
+    progressLogUpdate     = 30 * time.Second
+)
+
+var (
+    lastCommittedKey = []byte("atomicTrieLastCommittedBlock")
+)
+
+// AtomicTrie maintains an index of atomic operations by blockchainIDs for every block
+// height containing atomic transactions. The backing data structure for this index is
+// a Trie. The keys of the trie are block heights and the values (leaf nodes)
+// are the atomic operations applied to shared memory while processing the block accepted
+// at the corresponding height.
+type AtomicTrie interface {
+    // Index indexes the given atomicOps at the specified block height
+    // Returns an optional root hash
+    // A non-empty root hash is returned when the atomic trie has been committed
+    // Atomic trie is committed if the block height is multiple of commit interval
+    Index(height uint64, atomicOps map[ids.ID]*atomic.Requests) error
+
+    // Iterator returns an AtomicTrieIterator to iterate the trie at the given
+    // root hash
+    Iterator(hash common.Hash, startHeight uint64) (AtomicTrieIterator, error)
+
+    // LastCommitted returns the last committed hash and corresponding block height
+    LastCommitted() (common.Hash, uint64)
+
+    // TrieDB returns the underlying trie database
+    TrieDB() *trie.Database
+
+    // Root returns hash if it exists at specified height
+    // if trie was not committed at provided height, it returns
+    // common.Hash{} instead
+    Root(height uint64) (common.Hash, error)
+}
+
+// AtomicTrieIterator is a stateful iterator that iterates the leafs of an AtomicTrie
+type AtomicTrieIterator interface {
+    // Next advances the iterator to the next node in the atomic trie and
+    // returns true if there are more nodes to iterate
+    Next() bool
+
+    // BlockNumber returns the current block number
+    BlockNumber() uint64
+
+    // BlockchainID returns the current blockchain ID at the current block number
+    BlockchainID() ids.ID
+
+    // AtomicOps returns a map of blockchainIDs to the set of atomic requests
+    // for that blockchainID at the current block number
+    AtomicOps() *atomic.Requests
+
+    // Error returns error, if any encountered during this iteration
+    Error() error
+}
+
+// atomicTrie implements the AtomicTrie interface
+// using the eth trie.Trie implementation
+type atomicTrie struct {
+    commitHeightInterval uint64 // commit interval, same as commitHeightInterval by default
+    db                   *versiondb.Database // Underlying database
+    bonusBlocks          map[uint64]ids.ID    // Map of height to blockID for blocks to skip indexing
+    metadataDB           database.Database     // Underlying database containing the atomic trie metadata
+    atomicTrieDB          database.Database     // Underlying database containing the atomic trie
+    trieDB               *trie.Database       // Trie database
+    trie                 *trie.Trie           // Atomic trie.Trie mapping key (height+blockchainID) and value (codec serialized atomic.Requests)
+    repo                 AtomicTxRepository
+    lastCommittedHash     common.Hash // trie root hash of the most recent commit
+    lastCommittedHeight   uint64      // index height of the most recent commit
+    codec                 codec.Manager
+    log                   log.Logger // struct logger
+}
+
+var _ AtomicTrie = &atomicTrie{}
+
+// NewAtomicTrie returns a new instance of a atomicTrie with the default commitHeightInterval.
+// Initializes the trie before returning it.
+func NewAtomicTrie(db *versiondb.Database, bonusBlocks map[uint64]ids.ID, repo AtomicTxRepository, codec codec.Manager, lastAcceptedHeight uint64) (AtomicTrie, error) {
+    return newAtomicTrie(db, bonusBlocks, repo, codec, lastAcceptedHeight, commitHeightInterval)
+}
+
+// newAtomicTrie returns a new instance of a atomicTrie with a configurable commitHeightInterval, used in testing.
+// Initializes the trie before returning it.
+func newAtomicTrie(
+    db *versiondb.Database, bonusBlocks map[uint64]ids.ID, repo AtomicTxRepository, codec codec.Manager, lastAcceptedHeight uint64, commitHeightInterval uint64,
+) (*atomicTrie, error) {
+    atomicTrieDB := prefixdb.New(atomicTrieDBPrefix, db)
+    metadataDB := prefixdb.New(atomicTrieMetaDBPrefix, db)
+    root, height, err := lastCommittedRootIfExists(metadataDB)
+    if err != nil {
+        return nil, err
+    }

```

```

+     }
+
+     triedb := trie.NewDatabaseWithConfig(
+         Database{atomicTrieDB},
+         &trie.Config{
+             Cache:      10, // Allocate 10MB of memory for clean cache
+             Preimages: false, // Keys are not hashed, so there is no need for preimages
+         },
+     )
+     t, err := trie.New(root, triedb)
+     if err != nil {
+         return nil, err
+     }
+
+     atomicTrie := &atomicTrie{
+         commitHeightInterval: commitHeightInterval,
+         db:                    db,
+         bonusBlocks:           bonusBlocks,
+         atomicTrieDB:          atomicTrieDB,
+         metadataDB:            metadataDB,
+         trieDB:                triedb,
+         trie:                  t,
+         repo:                  repo,
+         codec:                 codec,
+         lastCommittedHash:     root,
+         lastCommittedHeight:   height,
+         log:                   log.New("c", "atomicTrie"),
+     }
+     return atomicTrie, atomicTrie.initialize(lastAcceptedHeight)
+ }
+
+ // lastCommittedRootIfExists returns the last committed trie root and height if it exists
+ // else returns empty common.Hash{} and 0
+ // returns error only if there are issues with the underlying data store
+ // or if values present in the database are not as expected
+ func lastCommittedRootIfExists(db database.Database) (common.Hash, uint64, error) {
+     // read the last committed entry if it exists and set the root hash
+     lastCommittedHeightBytes, err := db.Get(lastCommittedKey)
+     switch {
+     case err == database.ErrNotFound:
+         return common.Hash{}, 0, nil
+     case err != nil:
+         return common.Hash{}, 0, err
+     case len(lastCommittedHeightBytes) != wrappers.LongLen:
+         return common.Hash{}, 0, fmt.Errorf("expected value of lastCommittedKey to be %d but was %d", wrappers.LongLen, len(lastCommittedHeightBytes))
+     }
+     height := binary.BigEndian.Uint64(lastCommittedHeightBytes)
+     hash, err := db.Get(lastCommittedHeightBytes)
+     if err != nil {
+         return common.Hash{}, 0, fmt.Errorf("committed hash does not exist for committed height: %d: %w", height, err)
+     }
+     return common.BytesToHash(hash), height, nil
+ }
+
+ // nearestCommitHeight returns the nearest multiple of commitInterval less than or equal to blockNumber
+ func nearestCommitHeight(blockNumber uint64, commitInterval uint64) uint64 {
+     return blockNumber - (blockNumber % commitInterval)
+ }
+
+ // initializes the atomic trie using the atomic repository height index.
+ // Iterating from the last indexed height to lastAcceptedBlockNumber, making a single commit at the
+ // most recent height divisible by the commitInterval.
+ // Subsequent updates to this trie are made using the Index call as blocks are accepted.
+ func (a *atomicTrie) initialize(lastAcceptedBlockNumber uint64) error {
+     start := time.Now()
+     a.log.Info("initializing atomic trie", "lastAcceptedBlockNumber", lastAcceptedBlockNumber)
+     // finalCommitHeight is the highest block that can be committed i.e. is divisible by b.commitHeightInterval
+     // Txs from heights greater than commitHeight are to be included in the trie corresponding to the block at
+     // finalCommitHeight+b.commitHeightInterval, which has not been accepted yet.
+     finalCommitHeight := nearestCommitHeight(lastAcceptedBlockNumber, a.commitHeightInterval)
+     uncommittedOpsMap := make(map[uint64]map[ids.ID]*atomic.Requests, lastAcceptedBlockNumber-finalCommitHeight)
+
+     heightBytes := make([]byte, wrappers.LongLen)
+     binary.BigEndian.PutUint64(heightBytes, a.lastCommittedHeight)
+     // iterate by height, from lastCommittedHeight to the lastAcceptedBlockNumber
+     iter := a.repo.IterateByHeight(heightBytes)
+     defer iter.Release()
+
+     preCommitBlockIndexed := 0
+     postCommitTxIndexed := 0
+     lastUpdate := time.Now()
+
+     // keep track of the latest generated trie's root and height.
+     lastHash := common.Hash{}
+     lastHeight := a.lastCommittedHeight
+     for iter.Next() {
+         // Get the height and transactions for this iteration (from the key and value, respectively)
+         // iterate over the transactions, indexing them if the height is < commit height
+         // otherwise, add the atomic operations from the transaction to the uncommittedOpsMap
+         height := binary.BigEndian.Uint64(iter.Key())
+         txs, err := ExtractAtomicTxs(iter.Value(), true, a.codec)
+         if err != nil {
+             return err
+         }
+
+         // combine atomic operations from all transactions at this block height
+         combinedOps, err := mergeAtomicOps(txs)
+         if err != nil {
+             return err
+         }
+
+         if _, skipBonusBlock := a.bonusBlocks[height]; skipBonusBlock {
+             // If [height] is a bonus block, do not index the atomic operations into the trie
+         } else if height > finalCommitHeight {
+             // if height is greater than commit height, add it to the map so that we can write it later
+             // this is to ensure we have all the data before the commit height so that we can commit the
+             // trie
+             uncommittedOpsMap[height] = combinedOps
+         } else {
+             if err := a.updateTrie(height, combinedOps); err != nil {
+                 return err
+             }
+             preCommitBlockIndexed++
+         }
+
+         if time.Since(lastUpdate) > progressLogUpdate {
+             a.log.Info("imported entries into atomic trie pre-commit", "heightsIndexed", preCommitBlockIndexed)
+             lastUpdate = time.Now()
+         }
+
+         // if height has reached or skipped over the next commit interval,
+         // keep track of progress and keep commit size under commitSizeCap
+         commitHeight := nearestCommitHeight(height, a.commitHeightInterval)
+         if lastHeight < commitHeight {
+             hash, _, err := a.trie.Commit(nil)
+             if err != nil {
+                 return err
+             }
+             // Dereference lastHash to avoid writing more intermediary
+             // trie nodes than needed to disk, while keeping the commit
+             // size under commitSizeCap (approximately).
+             // Check [lastHash != hash] here to avoid dereferencing the
+             // trie root in case there were no atomic txs since the
+             // last commit.

```

```

+         if (lastHash != common.Hash{} && lastHash != hash) {
+             a.trieDB.Dereference(lastHash)
+         }
+         storage, _ := a.trieDB.Size()
+         if storage > commitSizeCap {
+             a.log.Info("committing atomic trie progress", "storage", storage)
+             a.commit(commitHeight)
+             // Flush any remaining changes that have not been committed yet in the versiondb.
+             if err := a.db.Commit(); err != nil {
+                 return err
+             }
+         }
+         lastHash = hash
+         lastHeight = commitHeight
+     }
+ }
+ if err := iter.Error(); err != nil {
+     return err
+ }
+
+ // Note: we should never create a commit at the genesis block (should not contain any atomic txs)
+ if lastAcceptedBlockNumber == 0 {
+     return nil
+ }
+ // now that all heights < finalCommitHeight have been processed
+ // commit the trie
+ if err := a.commit(finalCommitHeight); err != nil {
+     return err
+ }
+ // Flush any remaining changes that have not been committed yet in the versiondb.
+ if err := a.db.Commit(); err != nil {
+     return err
+ }
+
+ // process uncommitted ops for heights > finalCommitHeight
+ for height, ops := range uncommittedOpsMap {
+     if err := a.updateTrie(height, ops); err != nil {
+         return fmt.Errorf("failed to update trie at height %d: %w", height, err)
+     }
+ }
+
+ postCommitTxIndexed++
+ if time.Since(lastUpdate) > progressLogUpdate {
+     a.log.Info("imported entries into atomic trie post-commit", "entriesIndexed", postCommitTxIndexed)
+     lastUpdate = time.Now()
+ }
+ }
+
+ a.log.Info(
+     "finished initializing atomic trie",
+     "lastAcceptedBlockNumber", lastAcceptedBlockNumber,
+     "preCommitEntriesIndexed", preCommitBlockIndexed, "postCommitEntriesIndexed", postCommitTxIndexed,
+     "time", time.Since(start),
+ )
+ return nil
+}
+
+// Index updates the trie with entries in atomicOps
+// This function updates the following:
+// - heightBytes => trie root hash (if the trie was committed)
+// - lastCommittedBlock => height (if the trie was committed)
+func (a *atomicTrie) Index(height uint64, atomicOps map[ids.ID]*atomic.Requests) error {
+    if err := a.validateIndexHeight(height); err != nil {
+        return err
+    }
+
+    if err := a.updateTrie(height, atomicOps); err != nil {
+        return err
+    }
+
+    if height%a.commitHeightInterval == 0 {
+        return a.commit(height)
+    }
+
+    return nil
+}
+
+// validateIndexHeight returns an error if [height] is not currently valid to be indexed.
+func (a *atomicTrie) validateIndexHeight(height uint64) error {
+    // Do not allow a height that we have already passed to be indexed
+    if height < a.lastCommittedHeight {
+        return fmt.Errorf("height %d must be after last committed height %d", height, a.lastCommittedHeight)
+    }
+
+    // Do not allow a height that is more than a commit interval ahead
+    // of the current index
+    nextCommitHeight := a.lastCommittedHeight + a.commitHeightInterval
+    if height > nextCommitHeight {
+        return fmt.Errorf("height %d not within the next commit height %d", height, nextCommitHeight)
+    }
+
+    return nil
+}
+
+// commit calls commit on the trie to generate a root, commits the underlying trieDB, and updates the
+// metadata pointers.
+// assumes that the caller is aware of the commit rules i.e. the height being within commitInterval.
+// returns the trie root from the commit
+func (a *atomicTrie) commit(height uint64) error {
+    hash, _, err := a.trie.Commit(nil)
+    if err != nil {
+        return err
+    }
+
+    a.log.Info("committed atomic trie", "hash", hash.String(), "height", height)
+    if err := a.trieDB.Commit(hash, false, nil); err != nil {
+        return err
+    }
+
+    // all good here, update the heightBytes
+    heightBytes := make([]byte, wrappers.LongLen)
+    binary.BigEndian.PutUint64(heightBytes, height)
+
+    // now save the trie hash against the height it was committed at
+    if err := a.metadataDB.Put(heightBytes, hash[:]); err != nil {
+        return err
+    }
+
+    // update lastCommittedKey with the current height
+    if err := a.metadataDB.Put(lastCommittedKey, heightBytes); err != nil {
+        return err
+    }
+
+    a.lastCommittedHash = hash
+    a.lastCommittedHeight = height
+    return nil
+}
+
+func (a *atomicTrie) updateTrie(height uint64, atomicOps map[ids.ID]*atomic.Requests) error {
+    for blockchainID, requests := range atomicOps {
+        valueBytes, err := a.codec.Marshal(codecVersion, requests)
+        if err != nil {
+            // highly unlikely but possible if atomic.Element
+            // has a change that is unsupported by the codec
+            return err
+        }
+    }
+}

```

```

+     }
+
+     // key is [height]+[blockchainID]
+     keyPacker := wrappers.Packer{Bytes: make([]byte, wrappers.LongLen+common.HashLength)}
+     keyPacker.PackLong(height)
+     keyPacker.PackFixedBytes(blockchainID[:])
+     if err := a.trie.TryUpdate(keyPacker.Bytes, valueBytes); err != nil {
+         return err
+     }
+ }
+
+ return nil
+}
+
+// LastCommitted returns the last committed trie hash and last committed height
+func (a *atomicTrie) LastCommitted() (common.Hash, uint64) {
+    return a.lastCommittedHash, a.lastCommittedHeight
+}
+
+// Iterator returns a types.AtomicTrieIterator that iterates the trie from the given
+// atomic trie root, starting at the specified height
+func (a *atomicTrie) Iterator(root common.Hash, startHeight uint64) (AtomicTrieIterator, error) {
+    startKey := make([]byte, wrappers.LongLen)
+    binary.BigEndian.PutUint64(startKey, startHeight)
+
+    t, err := trie.New(root, a.trieDB)
+    if err != nil {
+        return nil, err
+    }
+
+    iter := trie.NewIterator(t.NodeIterator(startKey))
+    return NewAtomicTrieIterator(iter, a.codec, iter.Err)
+}
+
+func (a *atomicTrie) TrieDB() *trie.Database {
+    return a.trieDB
+}
+
+// Root returns hash if it exists at specified height
+// if trie was not committed at provided height, it returns
+// common.Hash{} instead
+func (a *atomicTrie) Root(height uint64) (common.Hash, error) {
+    heightBytes := make([]byte, wrappers.LongLen)
+    binary.BigEndian.PutUint64(heightBytes, height)
+
+    hash, err := a.metadataDB.Get(heightBytes)
+    switch {
+    case err == database.ErrNotFound:
+        return common.Hash{}, nil
+    case err != nil:
+        return common.Hash{}, err
+    }
+    return common.BytesToHash(hash), nil
+}
+
diff --git a/plugin/evm/atomic_trie_iterator.go b/plugin/evm/atomic_trie_iterator.go
new file mode 100644
index 00000000..2c64c44f
--- /dev/null
+++ b/plugin/evm/atomic_trie_iterator.go
@@ -0,0 +1,110 @@
+// (c) 2019-2020, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "encoding/binary"
+    "fmt"
+
+    "github.com/flare-foundation/flare/codec"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/trie"
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utils/wrappers"
+)
+
+const atomicTrieKeyLen = wrappers.LongLen + common.HashLength
+
+// atomicTrieIterator is an implementation of types.AtomicTrieIterator that serves
+// parsed data with each iteration
+type atomicTrieIterator struct {
+    trieIterator *trie.Iterator // underlying trie.Iterator
+    codec        codec.Manager
+    atomicOps    *atomic.Requests // atomic operation entries at this iteration
+    blockchainID ids.ID           // blockchain ID
+    blockNumber  uint64          // block number at this iteration
+    err          error            // error if any has occurred
+}
+
+func NewAtomicTrieIterator(trieIterator *trie.Iterator, codec codec.Manager) AtomicTrieIterator {
+    return &atomicTrieIterator{trieIterator: trieIterator, codec: codec}
+}
+
+// Error returns error, if any encountered during this iteration
+func (a *atomicTrieIterator) Error() error {
+    return a.err
+}
+
+// Next returns whether there are more nodes to iterate over
+// On success, this function sets the blockNumber and atomicOps fields
+// In case of an error during this iteration, it sets the error value and resets the above fields.
+// It is the responsibility of the caller to check the result of Error() after an iterator reports
+// having no more elements to iterate.
+func (a *atomicTrieIterator) Next() bool {
+    hasNext := a.trieIterator.Next()
+
+    if a.trieIterator.Err != nil {
+        a.resetFields(a.trieIterator.Err)
+        return false
+    }
+
+    if !hasNext {
+        a.resetFields(nil)
+        return false
+    }
+
+    // if the underlying iterator has data to iterate over, parse and set the fields
+    // key is [blockNumberBytes]+[blockchainIDBytes] = 8+32=40 bytes
+    keyLen := len(a.trieIterator.Key)
+    // If the key has an unexpected length, set the error and stop the iteration since the data is
+    // no longer reliable.
+    if keyLen != atomicTrieKeyLen {
+        a.resetFields(fmt.Errorf("expected atomic trie key length to be %d but was %d", atomicTrieKeyLen, keyLen))
+        return false
+    }
+
+    blockNumber := binary.BigEndian.Uint64(a.trieIterator.Key[:wrappers.LongLen])
+    blockchainID, err := ids.ToID(a.trieIterator.Key[wrappers.LongLen:])
+    if err != nil {
+        a.resetFields(err)
+        return false
+    }
+}

```

```

+ // The value in the iterator should be the atomic requests serialized the the codec.
+ requests := new(atomic.Requests)
+ if _, err = a.codec.Unmarshal(a.trieIterator.Value, requests); err != nil {
+     a.resetFields(err)
+     return false
+ }
+
+ // Success, update the struct fields
+ a.blockNumber = blockNumber
+ a.blockchainID = blockchainID
+ a.atomicOps = requests
+ return true
+}
+
+// resetFields resets the value fields of the iterator to their nil values and sets the error value to [err].
+func (a *atomicTrieIterator) resetFields(err error) {
+    a.err = err
+    a.blockNumber = 0
+    a.blockchainID = ids.ID{}
+    a.atomicOps = nil
+}
+
+// BlockNumber returns the current block number
+func (a *atomicTrieIterator) BlockNumber() uint64 {
+    return a.blockNumber
+}
+
+// BlockchainID returns the current blockchain ID at the current block number
+func (a *atomicTrieIterator) BlockchainID() ids.ID {
+    return a.blockchainID
+}
+
+// AtomicOps returns atomic requests for the blockchainID at the current block number
+func (a *atomicTrieIterator) AtomicOps() *atomic.Requests {
+    return a.atomicOps
+}
+
diff --git a/plugin/evm/atomic_trie_iterator_test.go b/plugin/evm/atomic_trie_iterator_test.go
new file mode 100644
index 00000000..ca43ebd8
--- /dev/null
+++ b/plugin/evm/atomic_trie_iterator_test.go
@@ -0,0 +1,90 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/database/memdb"
+    "github.com/flare-foundation/flare/database/versiondb"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utls"
+    "github.com/stretchr/testify/assert"
+)
+
+func TestIteratorCanIterate(t *testing.T) {
+    lastAcceptedHeight := uint64(1000)
+    db := versiondb.New(memdb.New())
+    codec := testTxCodec()
+    repo, err := NewAtomicTxRepository(db, codec, lastAcceptedHeight)
+    assert.NoError(t, err)
+
+    // create state with multiple transactions
+    // since each test transaction generates random ID for blockchainID we should get
+    // multiple blockchain IDs per block in the overall combined atomic operation map
+    operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+    writeTx(t, repo, 0, lastAcceptedHeight, constTxPerHeight(3), nil, operationsMap)
+
+    // create an atomic trie
+    // on create it will initialize all the transactions from the above atomic repository
+    atomicTrie1, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 100)
+    assert.NoError(t, err)
+
+    lastCommittedHash1, lastCommittedHeight1 := atomicTrie1.LastCommitted()
+    assert.NoError(t, err)
+    assert.NotEqual(t, common.Hash{}, lastCommittedHash1)
+    assert.EqualValues(t, 1000, lastCommittedHeight1)
+
+    verifyOperations(t, atomicTrie1, codec, lastCommittedHash1, 0, 1000, operationsMap)
+
+    // iterate on a new atomic trie to make sure there is no resident state affecting the data and the
+    // iterator
+    atomicTrie2, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 100)
+    assert.NoError(t, err)
+    lastCommittedHash2, lastCommittedHeight2 := atomicTrie2.LastCommitted()
+    assert.NoError(t, err)
+    assert.NotEqual(t, common.Hash{}, lastCommittedHash2)
+    assert.EqualValues(t, 1000, lastCommittedHeight2)
+
+    verifyOperations(t, atomicTrie2, codec, lastCommittedHash1, 0, 1000, operationsMap)
+}
+
+func TestIteratorHandlesInvalidData(t *testing.T) {
+    lastAcceptedHeight := uint64(1000)
+    db := versiondb.New(memdb.New())
+    codec := testTxCodec()
+    repo, err := NewAtomicTxRepository(db, codec, lastAcceptedHeight)
+    assert.NoError(t, err)
+
+    // create state with multiple transactions
+    // since each test transaction generates random ID for blockchainID we should get
+    // multiple blockchain IDs per block in the overall combined atomic operation map
+    operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+    writeTx(t, repo, 0, lastAcceptedHeight, constTxPerHeight(3), nil, operationsMap)
+
+    // create an atomic trie
+    // on create it will initialize all the transactions from the above atomic repository
+    atomicTrie, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 100)
+    assert.NoError(t, err)
+
+    lastCommittedHash, lastCommittedHeight := atomicTrie.LastCommitted()
+    assert.NoError(t, err)
+    assert.NotEqual(t, common.Hash{}, lastCommittedHash)
+    assert.EqualValues(t, 1000, lastCommittedHeight)
+
+    verifyOperations(t, atomicTrie, codec, lastCommittedHash, 0, 1000, operationsMap)
+
+    // Add a random key-value pair to the atomic trie in order to test that the iterator correctly
+    // handles an error when it runs into an unexpected key-value pair in the trie.
+    assert.NoError(t, atomicTrie.trie.TryUpdate(utls.RandomBytes(50), utls.RandomBytes(50)))
+    assert.NoError(t, atomicTrie.commit(lastCommittedHeight+1))
+    corruptedHash, _ := atomicTrie.LastCommitted()
+    iter, err := atomicTrie.Iterator(corruptedHash, 0)
+    assert.NoError(t, err)
+    for iter.Next() {
+    }
+    assert.Error(t, iter.Error())
+}
diff --git a/plugin/evm/atomic_trie_test.go b/plugin/evm/atomic_trie_test.go
new file mode 100644

```

```

index 00000000..ba3fb68c
--- /dev/null
+++ b/plugin/evm/atomic_trie_test.go
@@ -0,0 +1,401 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "testing"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/database/memdb"
+    "github.com/flare-foundation/flare/database/versiondb"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/stretchr/testify/assert"
+)
+
+const testCommitInterval = 100
+
+func (tx *Tx) mustAtomicOps() map[ids.ID]*atomic.Requests {
+    id, reqs, err := tx.AtomicOps()
+    if err != nil {
+        panic(err)
+    }
+    return map[ids.ID]*atomic.Requests{id: reqs}
+}
+
+func TestNearestCommitHeight(t *testing.T) {
+    type test struct {
+        height, commitInterval, expectedCommitHeight uint64
+    }
+
+    for _, test := range []test{
+        {
+            height:          4500,
+            commitInterval:  4096,
+            expectedCommitHeight: 4096,
+        },
+        {
+            height:          8500,
+            commitInterval:  4096,
+            expectedCommitHeight: 8192,
+        },
+        {
+            height:          950,
+            commitInterval:  100,
+            expectedCommitHeight: 900,
+        },
+    } {
+        commitHeight := nearestCommitHeight(test.height, test.commitInterval)
+        assert.Equal(t, commitHeight, test.expectedCommitHeight)
+    }
+}
+
+func TestAtomicTrieInitialize(t *testing.T) {
+    type test struct {
+        commitInterval, lastAcceptedHeight, expectedCommitHeight uint64
+        numTxPerBlock func(uint64) int
+    }
+
+    for name, test := range map[string]test{
+        "genesis": {
+            commitInterval: 10,
+            lastAcceptedHeight: 0,
+            expectedCommitHeight: 0,
+            numTxPerBlock: constTxPerHeight(0),
+        },
+        "before first commit": {
+            commitInterval: 10,
+            lastAcceptedHeight: 5,
+            expectedCommitHeight: 0,
+            numTxPerBlock: constTxPerHeight(3),
+        },
+        "first commit": {
+            commitInterval: 10,
+            lastAcceptedHeight: 10,
+            expectedCommitHeight: 10,
+            numTxPerBlock: constTxPerHeight(3),
+        },
+        "past first commit": {
+            commitInterval: 10,
+            lastAcceptedHeight: 15,
+            expectedCommitHeight: 10,
+            numTxPerBlock: constTxPerHeight(3),
+        },
+        "many existing commits": {
+            commitInterval: 10,
+            lastAcceptedHeight: 1000,
+            expectedCommitHeight: 1000,
+            numTxPerBlock: constTxPerHeight(3),
+        },
+        "many existing commits plus 1": {
+            commitInterval: 10,
+            lastAcceptedHeight: 1001,
+            expectedCommitHeight: 1000,
+            numTxPerBlock: constTxPerHeight(3),
+        },
+        "some blocks without atomic tx": {
+            commitInterval: 10,
+            lastAcceptedHeight: 101,
+            expectedCommitHeight: 100,
+            numTxPerBlock: func(height uint64) int {
+                if height <= 50 || height == 101 {
+                    return 1
+                }
+                return 0
+            },
+        },
+    } {
+        t.Run(name, func(t *testing.T) {
+            db := versiondb.New(memdb.New())
+            codec := testTxCodec()
+            repo, err := NewAtomicTxRepository(db, codec, test.lastAcceptedHeight)
+            if err != nil {
+                t.Fatal(err)
+            }
+            operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+            writeTx(t, repo, 0, test.lastAcceptedHeight+1, test.numTxPerBlock, nil, operationsMap)
+
+            // Construct the atomic trie for the first time
+            atomicTrie1, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, test.lastAcceptedHeight, test.commitInterval)
+            if err != nil {
+                t.Fatal(err)
+            }
+            rootHash1, commitHeight1 := atomicTrie1.LastCommitted()
+            assert.EqualValues(t, test.expectedCommitHeight, commitHeight1)
+            if test.expectedCommitHeight != 0 {
+                assert.NotEqual(t, common.Hash{}, rootHash1)
+            }
+
+            // Verify the operations up to the expected commit height
+            verifyOperations(t, atomicTrie1, codec, rootHash1, 0, test.expectedCommitHeight, operationsMap)
+        })
+    }
+}

```

```

+
+         // Construct the atomic trie a second time and ensure that it produces the same hash
+         atomicTrie2, err := newAtomicTrie(versiondb.New(memdb.New()), make(map[uint64]ids.ID), repo, codec, test.lastAcceptedHeight, test.commitInterval)
+         if err != nil {
+             t.Fatal(err)
+         }
+         rootHash2, commitHeight2 := atomicTrie2.LastCommitted()
+         assert.EqualValues(t, commitHeight1, commitHeight2)
+         assert.EqualValues(t, rootHash1, rootHash2)
+
+         // We now index additional operations up the next commit interval in order to confirm that nothing
+         // during the initialization phase will cause an invalid root when indexing continues.
+         nextCommitHeight := nearestCommitHeight(test.lastAcceptedHeight+test.commitInterval, test.commitInterval)
+         for i := test.lastAcceptedHeight + 1; i <= nextCommitHeight; i++ {
+             txs := newTestTxs(test.numTxsPerBlock(i))
+             if err := repo.Write(i, txs); err != nil {
+                 t.Fatal(err)
+             }
+
+             atomicOps, err := mergeAtomicOps(txs)
+             if err != nil {
+                 t.Fatal(err)
+             }
+             if err := atomicTrie1.Index(i, atomicOps); err != nil {
+                 t.Fatal(err)
+             }
+             operationsMap[i] = atomicOps
+         }
+
+         updatedRoot, updatedLastCommitHeight := atomicTrie1.LastCommitted()
+         assert.EqualValues(t, nextCommitHeight, updatedLastCommitHeight)
+         assert.NotEqual(t, common.Hash{}, updatedRoot)
+
+         // Verify the operations up to the new expected commit height
+         verifyOperations(t, atomicTrie1, codec, updatedRoot, 0, updatedLastCommitHeight, operationsMap)
+
+         // Generate a new atomic trie to compare the root against.
+         atomicTrie3, err := newAtomicTrie(versiondb.New(memdb.New()), make(map[uint64]ids.ID), repo, codec, nextCommitHeight, test.commitInterval)
+         if err != nil {
+             t.Fatal(err)
+         }
+         rootHash3, commitHeight3 := atomicTrie3.LastCommitted()
+         assert.EqualValues(t, rootHash3, updatedRoot)
+         assert.EqualValues(t, updatedLastCommitHeight, commitHeight3)
+     })
+ }
+
+}
+
+func TestIndexerInitializesOnlyOnce(t *testing.T) {
+     lastAcceptedHeight := uint64(25)
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+     repo, err := NewAtomicTxRepository(db, codec, lastAcceptedHeight)
+     assert.NoError(t, err)
+     operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+     writeTxs(t, repo, 0, lastAcceptedHeight+1, constTxsPerHeight(2), nil, operationsMap)
+
+     // Initialize atomic repository
+     atomicTrie, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 10 /*commitHeightInterval*/)
+     assert.NoError(t, err)
+
+     hash, height := atomicTrie.LastCommitted()
+     assert.NotEqual(t, common.Hash{}, hash)
+     assert.Equal(t, uint64(20), height)
+
+     // We write another tx at a height below the last committed height in the repo and then
+     // re-initialize the atomic trie since initialize is not supposed to run again the height
+     // at the trie should still be the old height with the old commit hash without any changes.
+     // This scenario is not realistic, but is used to test potential double initialization behavior.
+     err = repo.Write(15, []*Tx{testDataExportTx()})
+     assert.NoError(t, err)
+
+     // Re-initialize the atomic trie
+     atomicTrie, err = newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 10 /*commitHeightInterval*/)
+     assert.NoError(t, err)
+
+     newHash, newHeight := atomicTrie.LastCommitted()
+     assert.Equal(t, height, newHeight, "height should not have changed")
+     assert.Equal(t, hash, newHash, "hash should be the same")
+ }
+
+func newTestAtomicTrieIndexer(t *testing.T) AtomicTrie {
+     db := versiondb.New(memdb.New())
+     repo, err := NewAtomicTxRepository(db, testTxCodec(), 0)
+     assert.NoError(t, err)
+     indexer, err := newAtomicTrie(db, make(map[uint64]ids.ID), repo, testTxCodec(), 0, testCommitInterval)
+     assert.NoError(t, err)
+     assert.NotNil(t, indexer)
+     return indexer
+ }
+
+func TestIndexerWriteAndRead(t *testing.T) {
+     atomicTrie := newTestAtomicTrieIndexer(t)
+
+     blockRootMap := make(map[uint64]common.Hash)
+     lastCommittedBlockHeight := uint64(0)
+     var lastCommittedBlockHash common.Hash
+
+     // process 205 blocks so that we get three commits (0, 100, 200)
+     for height := uint64(0); height <= testCommitInterval*2+5; /*=205*/ height++ {
+         atomicRequests := testDataImportTx().mustAtomicOps()
+         err := atomicTrie.Index(height, atomicRequests)
+         assert.NoError(t, err)
+         if height%testCommitInterval == 0 {
+             lastCommittedBlockHash, lastCommittedBlockHeight = atomicTrie.LastCommitted()
+             assert.NoError(t, err)
+             assert.NotEqual(t, common.Hash{}, lastCommittedBlockHash)
+             blockRootMap[lastCommittedBlockHeight] = lastCommittedBlockHash
+         }
+     }
+
+     // ensure we have 3 roots
+     assert.Len(t, blockRootMap, 3)
+
+     hash, height := atomicTrie.LastCommitted()
+     assert.EqualValues(t, lastCommittedBlockHeight, height, "expected %d was %d", 200, lastCommittedBlockHeight)
+     assert.Equal(t, lastCommittedBlockHash, hash)
+
+     // Verify that [atomicTrie] can access each of the expected roots
+     for height, hash := range blockRootMap {
+         root, err := atomicTrie.Root(height)
+         assert.NoError(t, err)
+         assert.Equal(t, hash, root)
+     }
+
+     // Ensure that Index refuses to accept blocks older than the last committed height
+     err := atomicTrie.Index(10, testDataExportTx().mustAtomicOps())
+     assert.Error(t, err)
+     assert.Equal(t, "height 10 must be after last committed height 200", err.Error())
+
+     // Ensure Index does not accept blocks beyond the next commit interval
+     nextCommitHeight := lastCommittedBlockHeight + testCommitInterval + 1 // =301
+     err = atomicTrie.Index(nextCommitHeight, testDataExportTx().mustAtomicOps())
+     assert.Error(t, err)
+ }

```



```

+         assert.Equal(t, "height 301 not within the next commit height 300", err.Error())
+     }
+ }
+
+func TestAtomicOpsAreNotTxOrderDependent(t *testing.T) {
+     atomicTrie1 := newTestAtomicTrieIndexer(t)
+     atomicTrie2 := newTestAtomicTrieIndexer(t)
+
+     for height := uint64(0); height <= testCommitInterval; /*=205*/ height++ {
+         tx1 := testDataImportTx()
+         tx2 := testDataImportTx()
+         atomicRequests1, err := mergeAtomicOps([]*Tx{tx1, tx2})
+         assert.NoError(t, err)
+         atomicRequests2, err := mergeAtomicOps([]*Tx{tx2, tx1})
+         assert.NoError(t, err)
+
+         err = atomicTrie1.Index(height, atomicRequests1)
+         assert.NoError(t, err)
+         err = atomicTrie2.Index(height, atomicRequests2)
+         assert.NoError(t, err)
+     }
+     root1, height1 := atomicTrie1.LastCommitted()
+     root2, height2 := atomicTrie2.LastCommitted()
+     assert.NotEqual(t, common.Hash{}, root1)
+     assert.Equal(t, uint64(testCommitInterval), height1)
+     assert.Equal(t, uint64(testCommitInterval), height2)
+     assert.Equal(t, root1, root2)
+ }
+
+func TestAtomicTrieSkipsBonusBlocks(t *testing.T) {
+     lastAcceptedHeight := uint64(100)
+     numTxsPerBlock := 3
+     commitInterval := uint64(10)
+     expectedCommitHeight := uint64(100)
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+     repo, err := NewAtomicTxRepository(db, codec, lastAcceptedHeight)
+     if err != nil {
+         t.Fatal(err)
+     }
+     operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+     writeTxs(t, repo, 0, lastAcceptedHeight, constTxsPerHeight(numTxsPerBlock), nil, operationsMap)
+
+     bonusBlocks := map[uint64]ids.ID{
+         10: {},
+         13: {},
+         14: {},
+     }
+
+     // Construct the atomic trie for the first time
+     atomicTrie, err := newAtomicTrie(db, bonusBlocks, repo, codec, lastAcceptedHeight, commitInterval)
+     if err != nil {
+         t.Fatal(err)
+     }
+     rootHash, commitHeight := atomicTrie.LastCommitted()
+     assert.EqualValues(t, expectedCommitHeight, commitHeight)
+     assert.NotEqual(t, common.Hash{}, rootHash)
+
+     // Verify the operations are as expected with the bonus block heights removed from the operations map
+     for height := range bonusBlocks {
+         delete(operationsMap, height)
+     }
+     verifyOperations(t, atomicTrie, codec, rootHash, 0, expectedCommitHeight, operationsMap)
+ }
+
+func TestIndexingNilShouldNotImpactTrie(t *testing.T) {
+     // operations to index
+     ops := make([]map[ids.ID]*atomic.Requests, 0)
+     for i := 0; i <= testCommitInterval; i++ {
+         ops = append(ops, testDataImportTx().mustAtomicOps())
+     }
+
+     // without nils
+     a1 := newTestAtomicTrieIndexer(t)
+     for i := uint64(0); i <= testCommitInterval; i++ {
+         if i%2 == 0 {
+             if err := a1.Index(i, ops[i]); err != nil {
+                 t.Fatal(err)
+             }
+         } else {
+             // do nothing
+         }
+     }
+
+     root1, height1 := a1.LastCommitted()
+     assert.NotEqual(t, common.Hash{}, root1)
+     assert.Equal(t, uint64(testCommitInterval), height1)
+
+     // with nils
+     a2 := newTestAtomicTrieIndexer(t)
+     for i := uint64(0); i <= testCommitInterval; i++ {
+         if i%2 == 0 {
+             if err := a2.Index(i, ops[i]); err != nil {
+                 t.Fatal(err)
+             }
+         } else {
+             if err := a2.Index(i, nil); err != nil {
+                 t.Fatal(err)
+             }
+         }
+     }
+
+     root2, height2 := a2.LastCommitted()
+     assert.NotEqual(t, common.Hash{}, root2)
+     assert.Equal(t, uint64(testCommitInterval), height2)
+
+     // key assertion of the test
+     assert.Equal(t, root1, root2)
+ }
+
+func BenchmarkAtomicTrieInit(b *testing.B) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+
+     operationsMap := make(map[uint64]map[ids.ID]*atomic.Requests)
+
+     lastAcceptedHeight := uint64(25000)
+     // add 25000 * 3 = 75000 transactions
+     repo, err := NewAtomicTxRepository(db, codec, lastAcceptedHeight)
+     assert.NoError(b, err)
+     writeTxs(b, repo, 0, 25000, constTxsPerHeight(3), nil, operationsMap)
+
+     var atomicTrie AtomicTrie
+     var hash common.Hash
+     var height uint64
+     b.ReportAllocs()
+     b.ResetTimer()
+     for i := 0; i < b.N; i++ {
+         atomicTrie, err = newAtomicTrie(db, make(map[uint64]ids.ID), repo, codec, lastAcceptedHeight, 5000)
+         assert.NoError(b, err)
+
+         hash, height = atomicTrie.LastCommitted()
+         assert.Equal(b, lastAcceptedHeight, height)
+         assert.NotEqual(b, common.Hash{}, hash)
+     }
+     b.StopTimer()
+ }

```

```

+ // Verify operations
+ verifyOperations(b, atomicTrie, codec, hash, 0, lastAcceptedHeight, operationsMap)
+}
diff --git a/plugin/evm/atomic_tx_repository.go b/plugin/evm/atomic_tx_repository.go
new file mode 100644
index 00000000..e98e2f4d
--- /dev/null
+++ b/plugin/evm/atomic_tx_repository.go
@@ -0,0 +1,363 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "encoding/binary"
+    "fmt"
+    "sort"
+    "time"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/database"
+    "github.com/flare-foundation/flare/database/prefixdb"
+    "github.com/flare-foundation/flare/database/versiondb"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/units"
+    "github.com/flare-foundation/flare/units/wrappers"
+)
+
+const (
+    commitSizeCap = 10 * units.MiB
+)
+
+var (
+    atomicTxIDDBPrefix      = []byte("atomicTxDB")
+    atomicHeightTxDBPrefix  = []byte("atomicHeightTxDB")
+    atomicRepoMetadataDBPrefix = []byte("atomicRepoMetadataDB")
+    maxIndexedHeightKey     = []byte("maxIndexedAtomicTxHeight")
+    bonusBlocksRepairedKey  = []byte("bonusBlocksRepaired")
+)
+
+// AtomicTxRepository defines an entity that manages storage and indexing of
+// atomic transactions
+type AtomicTxRepository interface {
+    GetIndexHeight() (uint64, error)
+    GetByTxID(txID ids.ID) (*Tx, uint64, error)
+    GetByHeight(height uint64) ([]*Tx, error)
+    Write(height uint64, txs []*Tx) error
+    WriteBonus(height uint64, txs []*Tx) error
+
+    IterateByHeight([]byte) database.Iterator
+
+    IsBonusBlocksRepaired() (bool, error)
+    MarkBonusBlocksRepaired(repairedEntries uint64) error
+}
+
+// atomicTxRepository is a prefixdb implementation of the AtomicTxRepository interface
+type atomicTxRepository struct {
+    // [acceptedAtomicTxDB] maintains an index of [txID] => [height]+[atomic tx] for all accepted atomic txs.
+    acceptedAtomicTxDB database.Database
+
+    // [acceptedAtomicTxByHeightDB] maintains an index of [height] => [atomic txs] for all accepted block heights.
+    acceptedAtomicTxByHeightDB database.Database
+
+    // [atomicRepoMetadataDB] maintains a single key-value pair which tracks the height up to which the atomic repository
+    // has indexed.
+    atomicRepoMetadataDB database.Database
+
+    // This db is used to store [maxIndexedHeightKey] to avoid interfering with the iterators over the atomic transaction DBs.
+    db *versiondb.Database
+
+    // Use this codec for serializing
+    codec codec.Manager
+}
+
+func NewAtomicTxRepository(db *versiondb.Database, codec codec.Manager, lastAcceptedHeight uint64) (AtomicTxRepository, error) {
+    repo := &atomicTxRepository{
+        acceptedAtomicTxDB:      prefixdb.New(atomicTxIDDBPrefix, db),
+        acceptedAtomicTxByHeightDB: prefixdb.New(atomicHeightTxDBPrefix, db),
+        atomicRepoMetadataDB:    prefixdb.New(atomicRepoMetadataDBPrefix, db),
+        codec:                   codec,
+        db:                      db,
+    }
+    return repo, repo.initializeHeightIndex(lastAcceptedHeight)
+}
+
+// initializeHeightIndex initializes the atomic repository and takes care of any required migration from the previous database
+// format which did not have a height -> txs index.
+func (a *atomicTxRepository) initializeHeightIndex(lastAcceptedHeight uint64) error {
+    startTime := time.Now()
+    lastLogTime := startTime
+
+    // [lastTxID] will be initialized to the last transaction that we indexed
+    // if we are part way through a migration.
+    var lastTxID ids.ID
+    indexHeightBytes, err := a.atomicRepoMetadataDB.Get(maxIndexedHeightKey)
+    switch err {
+    case nil:
+        break
+    case database.ErrNotFound:
+        break
+    default: // unexpected value in the database
+        return fmt.Errorf("found invalid value at max indexed height: %v", indexHeightBytes)
+    }
+
+    switch len(indexHeightBytes) {
+    case 0:
+        log.Info("Initializing atomic transaction repository from scratch")
+    case common.HashLength: // partially initialized
+        lastTxID, err = ids.ToID(indexHeightBytes)
+        if err != nil {
+            return err
+        }
+        log.Info("Initializing atomic transaction repository from txID", "lastTxID", lastTxID)
+    case wrappers.LongLen: // already initialized
+        return nil
+    default: // unexpected value in the database
+        return fmt.Errorf("found invalid value at max indexed height: %v", indexHeightBytes)
+    }
+
+    // Iterate from [lastTxID] to complete the re-index -> generating an index
+    // from height to a slice of transactions accepted at that height
+    iter := a.acceptedAtomicTxDB.NewIteratorWithStart(lastTxID[:])
+    defer iter.Release()
+
+    indexedTxs := 0
+
+    // Keep track of the size of the currently pending writes
+    pendingBytesApproximation := 0
+    for iter.Next() {
+        // iter.Value() consists of [height packed as uint64] + [tx serialized as packed []byte]

```

```

+         iterValue := iter.Value()
+         if len(iterValue) < wrappers.LongLen {
+             return fmt.Errorf("atomic tx DB iterator value had invalid length (%d) < (%d)", len(iterValue), wrappers.LongLen)
+         }
+         heightBytes := iterValue[:wrappers.LongLen]
+
+         // Get the tx iter is pointing to, len(txs) == 1 is expected here.
+         txBytes := iterValue[wrappers.LongLen+wrappers.IntLen:]
+         tx, err := ExtractAtomicTx(txBytes, a.codec)
+         if err != nil {
+             return err
+         }
+
+         // Check if there are already transactions at [height], to ensure that we
+         // add [txs] to the already indexed transactions at [height] instead of
+         // overwriting them.
+         if err := a.appendTxToHeightIndex(heightBytes, tx); err != nil {
+             return err
+         }
+         lastTxID = tx.ID()
+         pendingBytesApproximation += len(txBytes)
+
+         // call commitFn to write to underlying DB if we have reached
+         // [commitSizeCap]
+         if pendingBytesApproximation > commitSizeCap {
+             if err := a.atomicRepoMetadataDB.Put(maxIndexedHeightKey, lastTxID[:]); err != nil {
+                 return err
+             }
+             if err := a.db.Commit(); err != nil {
+                 return err
+             }
+             log.Info("Committing work initializing the atomic repository", "lastTxID", lastTxID, "pendingBytesApprox", pendingBytesApproximation)
+             pendingBytesApproximation = 0
+         }
+         indexedTxs++
+         // Periodically log progress
+         if time.Since(lastLogTime) > 15*time.Second {
+             lastLogTime = time.Now()
+             log.Info("Atomic repository initialization", "indexedTxs", indexedTxs)
+         }
+     }
+     if err := iter.Error(); err != nil {
+         return fmt.Errorf("atomic tx DB iterator errored while initializing atomic trie: %w", err)
+     }
+
+     // Updated the value stored [maxIndexedHeightKey] to be the lastAcceptedHeight
+     indexedHeight := make([]byte, wrappers.LongLen)
+     binary.BigEndian.PutUint64(indexedHeight, lastAcceptedHeight)
+     if err := a.atomicRepoMetadataDB.Put(maxIndexedHeightKey, indexedHeight); err != nil {
+         return err
+     }
+
+     log.Info("Completed atomic transaction repository migration", "lastAcceptedHeight", lastAcceptedHeight, "duration", time.Since(startTime))
+     return a.db.Commit()
+ }
+
+ // GetIndexHeight returns the last height that was indexed by the atomic repository
+ func (a *atomicTxRepository) GetIndexHeight() (uint64, error) {
+     indexHeightBytes, err := a.atomicRepoMetadataDB.Get(maxIndexedHeightKey)
+     if err != nil {
+         return 0, err
+     }
+
+     if len(indexHeightBytes) != wrappers.LongLen {
+         return 0, fmt.Errorf("unexpected length for indexHeightBytes %d", len(indexHeightBytes))
+     }
+     indexHeight := binary.BigEndian.Uint64(indexHeightBytes)
+     return indexHeight, nil
+ }
+
+ // GetByTxID queries [acceptedAtomicTxDB] for the [txID], parses a [*Tx] object
+ // if an entry is found, and returns it with the block height the atomic tx it
+ // represents was accepted on, along with an optional error.
+ func (a *atomicTxRepository) GetByTxID(txID ids.ID) (*Tx, uint64, error) {
+     indexedTxBytes, err := a.acceptedAtomicTxDB.Get(txID[:])
+     if err != nil {
+         return nil, 0, err
+     }
+
+     if len(indexedTxBytes) < wrappers.LongLen {
+         return nil, 0, fmt.Errorf("acceptedAtomicTxDB entry too short: %d", len(indexedTxBytes))
+     }
+
+     // value is stored as [height]+[tx bytes], decompose with a packer.
+     packer := wrappers.Packer{Bytes: indexedTxBytes}
+     height := packer.UnpackLong()
+     txBytes := packer.UnpackBytes()
+     tx, err := ExtractAtomicTx(txBytes, a.codec)
+     if err != nil {
+         return nil, 0, err
+     }
+
+     return tx, height, nil
+ }
+
+ // GetByHeight returns all atomic txs processed on block at [height].
+ // Returns [database.ErrNotFound] if there are no atomic transactions indexed at [height].
+ // Note: if [height] is below the last accepted height, then this means that there were
+ // no atomic transactions in the block accepted at [height].
+ // If [height] is greater than the last accepted height, then this will always return
+ // [database.ErrNotFound]
+ func (a *atomicTxRepository) GetByHeight(height uint64) ([]*Tx, error) {
+     heightBytes := make([]byte, wrappers.LongLen)
+     binary.BigEndian.PutUint64(heightBytes, height)
+
+     return a.getByHeightBytes(heightBytes)
+ }
+
+ func (a *atomicTxRepository) getByHeightBytes(heightBytes []byte) ([]*Tx, error) {
+     txsBytes, err := a.acceptedAtomicTxByHeightDB.Get(heightBytes)
+     if err != nil {
+         return nil, err
+     }
+
+     return ExtractAtomicTxsBatch(txsBytes, a.codec)
+ }
+
+ // Write updates indexes maintained on atomic txs, so they can be queried
+ // by txID or height. This method must be called only once per height,
+ // and [txs] must include all atomic txs for the block accepted at the
+ // corresponding height.
+ func (a *atomicTxRepository) Write(height uint64, txs []*Tx) error {
+     return a.write(height, txs, false)
+ }
+
+ // WriteBonus is similar to Write, except the [txID] => [height] is not
+ // overwritten if already exists.
+ func (a *atomicTxRepository) WriteBonus(height uint64, txs []*Tx) error {
+     return a.write(height, txs, true)
+ }
+
+ func (a *atomicTxRepository) write(height uint64, txs []*Tx, bonus bool) error {
+     if len(txs) > 1 {
+         // txs should be stored in order of txID to ensure consistency
+         // with txs initialized from the txID index.

```

```

+         copyTx := make([]*Tx, len(txs))
+         copy(copyTx, txs)
+         sort.Slice(copyTx, func(i, j int) bool { return copyTx[i].ID().Hex() < copyTx[j].ID().Hex() })
+         txs = copyTx
+     }
+     heightBytes := make([]byte, wrappers.LongLen)
+     binary.BigEndian.PutUint64(heightBytes, height)
+     // Skip adding an entry to the height index if [txs] is empty.
+     if len(txs) > 0 {
+         for _, tx := range txs {
+             if bonus {
+                 switch _, err := a.GetByTxID(tx.ID()); err {
+                     case nil:
+                         // avoid overwriting existing value if [bonus] is true
+                         continue
+                     case database.ErrNotFound:
+                         // no existing value to overwrite, proceed as normal
+                     default:
+                         // unexpected error
+                         return err
+                     }
+                 }
+                 if err := a.indexTxByID(heightBytes, tx); err != nil {
+                     return err
+                 }
+             }
+             if err := a.indexTxAtHeight(heightBytes, tx); err != nil {
+                 return err
+             }
+         }
+     }
+
+     // Update the index height regardless of if any atomic transactions
+     // were present at [height].
+     return a.atomicRepoMetadataDB.Put(maxIndexedHeightKey, heightBytes)
+}
+
+// indexTxByID writes [tx] into the [acceptedAtomicTxDB] stored as
+// [height] + [tx bytes]
+func (a *atomicTxRepository) indexTxByID(heightBytes []byte, tx *Tx) error {
+     txBytes, err := a.codec.Marshal(codecVersion, tx)
+     if err != nil {
+         return err
+     }
+
+     // map txID => [height]+[tx bytes]
+     heightTxPacker := wrappers.Packer[Bytes: make([]byte, wrappers.LongLen+wrappers.IntLen+len(txBytes))]
+     heightTxPacker.PackFixedBytes(heightBytes)
+     heightTxPacker.PackBytes(txBytes)
+     txID := tx.ID()
+
+     if err := a.acceptedAtomicTxDB.Put(txID[:], heightTxPacker.Bytes); err != nil {
+         return err
+     }
+
+     return nil
+}
+
+// indexTxAtHeight adds [height] -> [txs] to the [acceptedAtomicTxByHeightDB]
+func (a *atomicTxRepository) indexTxAtHeight(heightBytes []byte, txs []*Tx) error {
+     txsBytes, err := a.codec.Marshal(codecVersion, txs)
+     if err != nil {
+         return err
+     }
+     if err := a.acceptedAtomicTxByHeightDB.Put(heightBytes, txsBytes); err != nil {
+         return err
+     }
+
+     return nil
+}
+
+// appendTxToHeightIndex retrieves the transactions stored at [heightBytes] and appends
+// [tx] to the slice of transactions stored there.
+// This function is used while initializing the atomic repository to re-index the atomic transactions
+// by txID into the height -> txs index.
+func (a *atomicTxRepository) appendTxToHeightIndex(heightBytes []byte, tx *Tx) error {
+     txs, err := a.getByHeightBytes(heightBytes)
+     if err != nil && err != database.ErrNotFound {
+         return err
+     }
+
+     // Iterate over the existing transactions to ensure we do not add a
+     // duplicate to the index.
+     for _, existingTx := range txs {
+         if existingTx.ID() == tx.ID() {
+             return nil
+         }
+     }
+
+     txs = append(txs, tx)
+     return a.indexTxAtHeight(heightBytes, txs)
+}
+
+func (a *atomicTxRepository) IterateByHeight(heightBytes []byte) database.Iterator {
+     return a.acceptedAtomicTxByHeightDB.NewIteratorWithStart(heightBytes)
+}
+
+func (a *atomicTxRepository) IsBonusBlocksRepaired() (bool, error) {
+     return a.atomicRepoMetadataDB.Has(bonusBlocksRepairedKey)
+}
+
+func (a *atomicTxRepository) MarkBonusBlocksRepaired(repairedEntries uint64) error {
+     val := make([]byte, wrappers.LongLen)
+     binary.BigEndian.PutUint64(val, repairedEntries)
+     return a.atomicRepoMetadataDB.Put(bonusBlocksRepairedKey, val)
+}
+
+diff --git a/plugin/evm/atomic_tx_repository_test.go b/plugin/evm/atomic_tx_repository_test.go
new file mode 100644
index 00000000..56b0c2dc
--- /dev/null
+++ b/plugin/evm/atomic_tx_repository_test.go
+@ -0,0 +1,287 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+     "sort"
+     "testing"
+
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/flare-foundation/flare/chains/atomic"
+     "github.com/flare-foundation/flare/database"
+     "github.com/flare-foundation/flare/database/prefixdb"
+     "github.com/flare-foundation/flare/database/versiondb"
+
+     "github.com/flare-foundation/flare/codec"
+     "github.com/flare-foundation/flare/utls/wrappers"
+
+     "github.com/stretchr/testify/assert"
+
+     "github.com/flare-foundation/flare/database/memdb"
+     "github.com/flare-foundation/flare/ids"
+
+

```

```

+// addTx writes [txsPerHeight] txs for heights ranging in [fromHeight, toHeight) directly to [acceptedAtomicTxDB],
+// storing the resulting transactions in [txMap] if non-nil and the resulting atomic operations in [operationsMap]
+// if non-nil.
+func addTx(t testing.TB, codec codec.Manager, acceptedAtomicTxDB database.Database, fromHeight uint64, toHeight uint64, txsPerHeight int, txMap map[uint64][]Tx, operationsMap map[uint64]map[ids.ID]*atomic.Requests) {
+    for height := fromHeight; height < toHeight; height++ {
+        txs := make([]Tx, 0, txsPerHeight)
+        for i := 0; i < txsPerHeight; i++ {
+            tx := newTestTx()
+            txs = append(txs, tx)
+            txBytes, err := codec.Marshal(codecVersion, tx)
+            assert.NoError(t, err)
+
+            // Write atomic transactions to the [acceptedAtomicTxDB]
+            // in the format handled prior to the migration to the atomic
+            // tx repository.
+            packer := wrappers.Packer{Bytes: make([]byte, 1), MaxSize: 1024 * 1024}
+            packer.PackLong(height)
+            packer.PackBytes(txBytes)
+            txID := tx.ID()
+            err = acceptedAtomicTxDB.Put(txID[:], packer.Bytes)
+            assert.NoError(t, err)
+        }
+        // save this to the map (if non-nil) for verifying expected results in verifyTx
+        if txMap != nil {
+            txMap[height] = txs
+        }
+        if operationsMap != nil {
+            atomicRequests, err := mergeAtomicOps(txs)
+            if err != nil {
+                t.Fatal(err)
+            }
+            operationsMap[height] = atomicRequests
+        }
+    }
+}
+
+// constTxPerHeight returns a function for passing to [writeTx], which will return a constant number
+// as the number of atomic txs per height to create.
+func constTxPerHeight(txCount int) func(uint64) int {
+    return func(uint64) int { return txCount }
+}
+
+// writeTx writes [txsPerHeight] txs for heights ranging in [fromHeight, toHeight) through the Write call on [repo],
+// storing the resulting transactions in [txMap] if non-nil and the resulting atomic operations in [operationsMap]
+// if non-nil.
+func writeTx(t testing.TB, repo AtomicTxRepository, fromHeight uint64, toHeight uint64, txsPerHeight func(height uint64) int, txMap map[uint64][]Tx, operationsMap map[uint64]map[ids.ID]*atomic.Requests) {
+    for height := fromHeight; height < toHeight; height++ {
+        txs := newTestTx(txsPerHeight(height))
+        if err := repo.Write(height, txs); err != nil {
+            t.Fatal(err)
+        }
+        // save this to the map (if non-nil) for verifying expected results in verifyTx
+        if txMap != nil {
+            txMap[height] = txs
+        }
+        if operationsMap != nil {
+            atomicRequests, err := mergeAtomicOps(txs)
+            if err != nil {
+                t.Fatal(err)
+            }
+            if len(atomicRequests) == 0 {
+                continue
+            }
+            operationsMap[height] = atomicRequests
+        }
+    }
+}
+
+// verifyTx asserts [repo] can find all txs in [txMap] by height and txID
+func verifyTx(t testing.TB, repo AtomicTxRepository, txMap map[uint64][]Tx) {
+    // We should be able to fetch indexed txs by height:
+    getComparator := func(txs []Tx) func(int, int) bool {
+        return func(i, j int) bool {
+            return txs[i].ID().Hex() < txs[j].ID().Hex()
+        }
+    }
+    for height, expectedTx := range txMap {
+        txs, err := repo.GetByHeight(height)
+        assert.NoError(t, err, "unexpected error on GetByHeight at height=%d", height)
+        assert.Len(t, txs, len(expectedTx), "wrong len of txs at height=%d", height)
+        // txs should be stored in order of txID
+        sort.Slice(expectedTx, getComparator(expectedTx))
+
+        txIDs := ids.Set{}
+        for i := 0; i < len(txs); i++ {
+            assert.Equalf(t, expectedTx[i].ID().Hex(), txs[i].ID().Hex(), "wrong txID at height=%d idx=%d", height, i)
+            txIDs.Add(txs[i].ID())
+        }
+        assert.Equalf(t, len(txs), txIDs.Len(), "incorrect number of unique transactions in slice at height %d, expected %d, found %d", height, len(txs), txIDs.Len())
+    }
+}
+
+// verifyOperations creates an iterator over the atomicTrie at [rootHash] and verifies that the all of the operations in the trie in the interval [from, to] are identical to
+// the atomic operations contained in [operationsMap] on the same interval.
+func verifyOperations(t testing.TB, atomicTrie AtomicTrie, codec codec.Manager, rootHash common.Hash, from, to uint64, operationsMap map[uint64]map[ids.ID]*atomic.Requests) {
+    // Start the iterator at [from]
+    iter, err := atomicTrie.Iterator(rootHash, from)
+    if err != nil {
+        t.Fatal(err)
+    }
+
+    // Generate map of the marshalled atomic operations on the interval [from, to]
+    // based on [operationsMap].
+    marshalledOperationsMap := make(map[uint64]map[ids.ID][]byte)
+    for height, blockRequests := range operationsMap {
+        if height < from || height > to {
+            continue
+        }
+        for blockchainID, atomicRequests := range blockRequests {
+            b, err := codec.Marshal(0, atomicRequests)
+            if err != nil {
+                t.Fatal(err)
+            }
+            if requestsMap, exists := marshalledOperationsMap[height]; exists {
+                requestsMap[blockchainID] = b
+            } else {
+                requestsMap = make(map[ids.ID][]byte)
+                requestsMap[blockchainID] = b
+                marshalledOperationsMap[height] = requestsMap
+            }
+        }
+    }
+
+    // Generate map of marshalled atomic operations on the interval [from, to]
+    // based on the contents of the trie.
+    iteratorMarshalledOperationsMap := make(map[uint64]map[ids.ID][]byte)
+    for iter.Next() {
+        height := iter.BlockNumber()
+        if height < from {
+            t.Fatalf("Iterator starting at (%d) found value at block height (%d)", from, height)
+        }
+        if height > to {

```

```

+         continue
+     }
+
+     blockchainID := iter.BlockchainID()
+     b, err := codec.Marshal(0, iter.AtomicOps())
+     if err != nil {
+         t.Fatal(err)
+     }
+     if requestsMap, exists := iteratorMarshaledOperationsMap[height]; exists {
+         requestsMap[blockchainID] = b
+     } else {
+         requestsMap = make(map[ids.ID][]byte)
+         requestsMap[blockchainID] = b
+         iteratorMarshaledOperationsMap[height] = requestsMap
+     }
+ }
+ if err := iter.Error(); err != nil {
+     t.Fatal(err)
+ }
+
+ assert.Equal(t, marshalledOperationsMap, iteratorMarshaledOperationsMap)
+}
+
+func TestAtomicRepositoryReadWriteSingleTx(t *testing.T) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+     repo, err := NewAtomicTxRepository(db, codec, 0)
+     if err != nil {
+         t.Fatal(err)
+     }
+     txMap := make(map[uint64][]Tx)
+
+     writeTxs(t, repo, 0, 100, constTxsPerHeight(1), txMap, nil)
+     verifyTxs(t, repo, txMap)
+}
+
+func TestAtomicRepositoryReadWriteMultipleTxs(t *testing.T) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+     repo, err := NewAtomicTxRepository(db, codec, 0)
+     if err != nil {
+         t.Fatal(err)
+     }
+     txMap := make(map[uint64][]Tx)
+
+     writeTxs(t, repo, 0, 100, constTxsPerHeight(10), txMap, nil)
+     verifyTxs(t, repo, txMap)
+}
+
+func TestAtomicRepositoryPreAP5Migration(t *testing.T) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+
+     acceptedAtomicTxDB := prefixdb.New(atomicTxIDDBPrefix, db)
+     txMap := make(map[uint64][]Tx)
+     addTxs(t, codec, acceptedAtomicTxDB, 0, 100, 1, txMap, nil)
+     if err := db.Commit(); err != nil {
+         t.Fatal(err)
+     }
+
+     // Ensure the atomic repository can correctly migrate the transactions
+     // from the old accepted atomic tx DB to add the height index.
+     repo, err := NewAtomicTxRepository(db, codec, 100)
+     if err != nil {
+         t.Fatal(err)
+     }
+     assert.NoError(t, err)
+     verifyTxs(t, repo, txMap)
+
+     writeTxs(t, repo, 100, 150, constTxsPerHeight(1), txMap, nil)
+     writeTxs(t, repo, 150, 200, constTxsPerHeight(10), txMap, nil)
+     verifyTxs(t, repo, txMap)
+}
+
+func TestAtomicRepositoryPostAP5Migration(t *testing.T) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+
+     acceptedAtomicTxDB := prefixdb.New(atomicTxIDDBPrefix, db)
+     txMap := make(map[uint64][]Tx)
+     addTxs(t, codec, acceptedAtomicTxDB, 0, 100, 1, txMap, nil)
+     addTxs(t, codec, acceptedAtomicTxDB, 100, 200, 10, txMap, nil)
+     if err := db.Commit(); err != nil {
+         t.Fatal(err)
+     }
+
+     // Ensure the atomic repository can correctly migrate the transactions
+     // from the old accepted atomic tx DB to add the height index.
+     repo, err := NewAtomicTxRepository(db, codec, 200)
+     if err != nil {
+         t.Fatal(err)
+     }
+     assert.NoError(t, err)
+     verifyTxs(t, repo, txMap)
+
+     writeTxs(t, repo, 200, 300, constTxsPerHeight(10), txMap, nil)
+     verifyTxs(t, repo, txMap)
+}
+
+func benchAtomicRepositoryIndex10_000(b *testing.B, maxHeight uint64, txsPerHeight int) {
+     db := versiondb.New(memdb.New())
+     codec := testTxCodec()
+
+     acceptedAtomicTxDB := prefixdb.New(atomicTxIDDBPrefix, db)
+     txMap := make(map[uint64][]Tx)
+
+     addTxs(b, codec, acceptedAtomicTxDB, 0, maxHeight, txsPerHeight, txMap, nil)
+     if err := db.Commit(); err != nil {
+         b.Fatal(err)
+     }
+
+     repo, err := NewAtomicTxRepository(db, codec, maxHeight)
+     if err != nil {
+         b.Fatal(err)
+     }
+     assert.NoError(b, err)
+     verifyTxs(b, repo, txMap)
+}
+
+func BenchmarkAtomicRepositoryIndex_10kBlocks_1Tx(b *testing.B) {
+     for n := 0; n < b.N; n++ {
+         benchAtomicRepositoryIndex10_000(b, 10_000, 1)
+     }
+}
+
+func BenchmarkAtomicRepositoryIndex_10kBlocks_10Tx(b *testing.B) {
+     for n := 0; n < b.N; n++ {
+         benchAtomicRepositoryIndex10_000(b, 10_000, 10)
+     }
+}
+
diff --git a/plugin/evm/block.go b/plugin/evm/block.go
index 8789c2a7..53e19576 100644
--- a/plugin/evm/block.go
+++ b/plugin/evm/block.go
@@ -12,90 +12,106 @@ import (

```

```
"github.com/ethereum/go-ethereum/log"
"github.com/ethereum/go-ethereum/rlp"

- "github.com/ava-labs/coreth/core/types"
- "github.com/ava-labs/coreth/params"
+ "github.com/flare-foundation/coreth/core/types"
+ "github.com/flare-foundation/coreth/params"

- "github.com/ava-labs/avalanchego/ids"
- "github.com/ava-labs/avalanchego/snow/choices"
+ "github.com/flare-foundation/flare/chains/atomic"
+ "github.com/flare-foundation/flare/ids"
+ "github.com/flare-foundation/flare/snow/choices"
)

-var bonusBlocks = ids.Set{}
+var (
+    bonusBlocks                = ids.Set{}
+    bonusBlockMainnetHeights = make(map[uint64]ids.ID)
+    // first height that processed a TX included on a
+    // bonus block is the canonical height for that TX.
+    canonicalBonusBlocks = []uint64{
+        102928, 103035, 103038, 103114, 103193,
+        103234, 103338, 103444, 103480, 103491,
+        103513, 103533, 103535, 103538, 103541,
+        103546, 103571, 103572, 103619,
+        103287, 103624, 103591,
+    }
+)

func init() {
-    blockIDStrings := []string{
-        "XMoEsew2DhSgQaydcJFJUQAQYP88TNTYbEJZvtbrV20sX7iE3",
-        "20iHZwLhQ3xLuyyfcdo5yCfuoSqdVdR2ox5ECU19HIsWfRoCGp",
-        "tLLijh7oKfvWt1Yk9zRv4F0vuQSDAiuvb5kHCNN9zh4mqkFMG",
-        "2db2wMbVAoCc5EUJrsBYWwNZDekyY8uNpaaVapdBaQZ5oRaou",
-        "2rAsBj3emqQa13CV8r5fttHogs4sXnjvbbXvzCKPi3WmzhpK9D",
-        "amgH2C1s9H3Av7vSW4y7n7TXb9tKYKHENvrdXutgN6nsejgc",
-        "dWBSRYRwFrcy13DPdLoHsL670kZ5h86hwtVPf94ZBaY18Ekmf",
-        "PgaRk1UAoUvRybhnxSrlQ5t6imWhEa6ksNjbN6hWgs4qPr5zm",
-        "b7XFDDLgwb12DfL7UTWZoxwBpkLPL5mdHtXngD94Y2RoeWXSh",
-        "2i2FP6nJyvhX9FR15qN2D9AvOkSXKgBD212A07FoSpfowxvQDX",
-        "2J8z7HNv4nwh82wqRGyEHqQeuw4J6mCDCSVUgusBu35asnshK",
-        "2cUPPHy1hspr2nAKpQrrAEisLKkaWS59iF2wJfNyFRs8vnSkKk",
-        "2gTygYckZgFzFN5QOWPaPB03nabqjdV55mwy1x1Nd4JmJAwam",
-        "5Mpt5dP6dBMP5wK9GJjeVe39deZJTRh9i82cgN1bjedffrrTf",
-        "2v3smb35s4GLACsK4Zkd2RcLbLDWA4huqrvq8Y3VP4Cve8kFTM",
-        "7KCZKBpxovtX9opb7rMRie9WmW5YbZ8A4HwBBokJ9eShpZPqx",
-        "2oueNTj4dUE2FFtGyPpawmCCsy6EUQeVHWLZy8HMeQmkAcIP4",
-        "Nzs93kFTvcXanFlup8Y8BQKYNzm8BxykVNF3TkdyaEeuxwbP",
-        "2YH2KymFj1BhpKzgt6HXJhLSt5SV9U04tJuUNjfn1nQ0dm5zz",
-        "Qv5v5Ru8ArfnWKB1w6s4G5EYPh7TybHJtF6UsVwAkvZFoqmj",
-        "z3BgePPpCqqlmRBRvUj28rYYxnEtJizkUEHnDBrcZeV7MFVvk",
-        "Ry2sJfFfGEnJxRKUGFSyZn7GR3m4aKaF1scdW2uXSNQ8568Y",
-        "2YgxGHns7Z2hMMHJsPCgVXuJaL7x1b3gnHbmSCfCdyAcYGr6mx",
-        "cwJusfmn98TW3DjAbfLRN9utYR24KAQ82qpAXmVsvjHyJzUM2",
-        "2JbuExUGKw5mYz5KfXATwq11bRD1mgks9wEdYGN5C6Ttey1R4U",
-        "21Jys8UNURmtckKSV8952hntEwYmJsZrLQbdLaNcbXcxDAsQ5a",
-        "MjExz2z1qhwugc1tAy1GxRsCq4GvJwKfyyS29nr4tRVB8oo1c",
-        "9oZha4yBCcVwSgYDoUzRAuausvPNJ3xH6nopK56bwYzMFLoQ2",
-        "uK5FF9iBfDtRepVv9NgCQ1STD1nzLJG3yrf1bHG4mGvmybw6f",
-        "22ckZ27cC38hmBFX2v3jMwXun8eD8psNa1cFyeokS67DxwmPTx",
-        "2AfTQ2FXNj9bkSUQnud9pFXULx6EbF7cbbw6i3ayvc2QNhgxTF",
-        "pTf7gfklksj7bqMrLyMCj1f8FBKtH1uRqQrtfYkMfExhX5xnrl",
-        "2AXxT3PSEnaYHntBTnYrVTf24tKdWjky9sqoFEhydrGXE9iKH",
-        "PJTkrRvKZ1m4AQdPND1MBpUXpCrGN4D0mXmJQA1UrsxPoLQX",
-        "fv8K1U8oQDmfVwK66AwN73a5sU1Whm8quNvNmKsZnBvcY2W",
-        "sg6wAwFBsPQ155Yfyh41cVCKRQbrrXsXmeNyQ1xkunfZsdyv",
-        "soPweZ8DGaoUjrnzjH3V2bypa7ZvvfqBan4UCsMuxMP759gw",
-        "2DnKpQF4mooveyUDfBYQTbFsGDV4wkncQPpEw4kHkfst5T5o5x",
-        "63YldYXfXc5tY3mmwLaDSbXzQHymwVxMP7HKBrd4u3C2iM1",
-        "2tCe88ur6MLQcVgwE5XoaHlTgtSrtHwKN35DbHE4kwiQ7MSTV",
-        "2nG4exd9eUoAGzELfksmBR8XDCKhohY1uDKRFzEXJG4M8p3qa7",
-        "2F5t5QbdTfhZxvkvZqdFp7KR3FrJPKEsDLQK7ktPhNxj1EZAh4",
-        "21o2fVTnzmtgXqkV1yuQeze7YEQhR5JB31jVVD9oVUnaaV8qm",
-        "2p5jfo7rkFcfZ2CqAxqfW8vqM2CU2nVlHrfZ3rwxz43gkVuGo",
-        "2QBMMMFJmhVHaGF45GAPszKyj1gK6ToBERRxYvXtM7yfrdUGPK",
-        "2ez4CA7w4HHR85SobHQUAwFgj2giRNjNFUZK9JvrZfAlAuRj6X",
-        "2DpCuBaH94zKKFNY2XTs4GeJcwsEv6qZTDHC5958tdg97GZpcJ",
-        "i1HoerJlaxognkUKKL58FvF9aLrbZKtvt7TdKlKt5kgzoeU1vB",
-        "2S15ziHHqPjb1qkw7CdGyupokiYpd2b7mMqR1yszurctcA5AKr",
-        "esx5J962LTYm2aSrskplai5e4CMmsA1dsu9iulGJ3KwGsu2M",
-        "2czmtnBS44VCMNRFUM89h4Fe9m3ZeZVYyh7Pe3FhNqjRNgPhXz",
-        "DK9NqAJGry1wAo767uuyC1dYXajUhzwka6vi8d9tNheqzGUTd",
-        "pE93VXY3N5QKfwsEFcM9159UpPFeZ8nxpJNaGaDQyDgsscNf",
-        "AfWvJH3rB2fdHuPW0p6qYNCFTV29MooQPRigD88rKwUDEDhq",
-        "2KPW9G5t1NF14tZNfG4SgHu0rtUYVZyxuo0f3aZ7AntKrQdsHn",
-        "BYqLB6xpqy7HSAgP2XNfGE8Ubg1uEzse5mBPTSJH9258spvMa",
-        "Njm9TCLUXRoJZk8YhEM6ksvfiPdC1TME4zJvGaDXgzMcyB6oB",
+    mainnetBonusBlocks := map[uint64]string{
+        102972: "Njm9TCLUXRoJZk8YhEM6ksvfiPdC1TME4zJvGaDXgzMcyB6oB",
+        103105: "BYqLB6xpqy7HSAgP2XNfGE8Ubg1uEzse5mBPTSJH9258spvMa",
+        103143: "AfWvJH3rB2fdHuPW0p6qYNCFTV29MooQPRigD88rKwUDEDhq",
+        103183: "2KPW9G5t1NF14tZNfG4SgHu0rtUYVZyxuo0f3aZ7AntKrQdsHn",
+        103197: "pE93VXY3N5QKfwsEFcM9159UpPFeZ8nxpJNaGaDQyDgsscNf",
+        103203: "2czmtnBS44VCMNRFUM89h4Fe9m3ZeZVYyh7Pe3FhNqjRNgPhXz",
+        103208: "esx5J962LTYm2aSrskplai5e4CMmsA1dsu9iulGJ3KwGsu2M",
+        103209: "DK9NqAJGry1wAo767uuyC1dYXajUhzwka6vi8d9tNheqzGUTd",
+        103259: "i1HoerJlaxognkUKKL58FvF9aLrbZKtvt7TdKlKt5kgzoeU1vB",
+        103261: "2DpCuBaH94zKKFNY2XTs4GeJcwsEv6qZTDHC5958tdg97GZpcJ",
+        103266: "2ez4CA7w4HHR85SobHQUAwFgj2giRNjNFUZK9JvrZfAlAuRj6X",
+        103287: "2QBMMMFJmhVHaGF45GAPszKyj1gK6ToBERRxYvXtM7yfrdUGPK",
+        103339: "2p5jfo7rkFcfZ2CqAxqfW8vqM2CU2nVlHrfZ3rwxz43gkVuGo",
+        103346: "2S15ziHHqPjb1qkw7CdGyupokiYpd2b7mMqR1yszurctcA5AKr",
+        103350: "2F5t5QbdTfhZxvkvZqdFp7KR3FrJPKEsDLQK7ktPhNxj1EZAh4",
+        103358: "2tCe88ur6MLQcVgwE5XoaHlTgtSrtHwKN35DbHE4kwiQ7MSTV",
+        103437: "21o2fVTnzmtgXqkV1yuQeze7YEQhR5JB31jVVD9oVUnaaV8qm",
+        103472: "2nG4exd9eUoAGzELfksmBR8XDCKhohY1uDKRFzEXJG4M8p3qa7",
+        103478: "63YldYXfXc5tY3mmwLaDSbXzQHymwVxMP7HKBrd4u3C2iM1",
+        103493: "soPweZ8DGaoUjrnzjH3V2bypa7ZvvfqBan4UCsMuxMP759gw",
+        103514: "2DnKpQF4mooveyUDfBYQTbFsGDV4wkncQPpEw4kHkfst5T5o5x",
+        103536: "PJTkrRvKZ1m4AQdPND1MBpUXpCrGN4D0mXmJQA1UrsxPoLQX",
+        103545: "22ckZ27cC38hmBFX2v3jMwXun8eD8psNa1cFyeokS67DxwmPTx",
+        103547: "pTf7gfklksj7bqMrLyMCj1f8FBKtH1uRqQrtfYkMfExhX5xnrl",
+        103554: "9oZha4yBCcVwSgYDoUzRAuausvPNJ3xH6nopK56bwYzMFLoQ2",
+        103555: "MjExz2z1qhwugc1tAy1GxRsCq4GvJwKfyyS29nr4tRVB8oo1c",
+        103559: "cwJusfmn98TW3DjAbfLRN9utYR24KAQ82qpAXmVsvjHyJzUM2",
+        103561: "2YgxGHns7Z2hMMHJsPCgVXuJaL7x1b3gnHbmSCfCdyAcYGr6mx",
+        103563: "2AXxT3PSEnaYHntBTnYrVTf24tKdWjky9sqoFEhydrGXE9iKH",
+        103564: "Ry2sJfFfGEnJxRKUGFSyZn7GR3m4aKaF1scdW2uXSNQ8568Y",
+        103569: "21Jys8UNURmtckKSV8952hntEwYmJsZrLQbdLaNcbXcxDAsQ5a",
+        103570: "sg6wAwFBsPQ155Yfyh41cVCKRQbrrXsXmeNyQ1xkunfZsdyv",
+        103575: "z3BgePPpCqqlmRBRvUj28rYYxnEtJizkUEHnDBrcZeV7MFVvk",
+        103577: "uK5FF9iBfDtRepVv9NgCQ1STD1nzLJG3yrf1bHG4mGvmybw6f",
+        103578: "Qv5v5Ru8ArfnWKB1w6s4G5EYPh7TybHJtF6UsVwAkvZFoqmj",
+        103582: "7KCZKBpxovtX9opb7rMRie9WmW5YbZ8A4HwBBokJ9eShpZPqx",
+        103587: "2AfTQ2FXNj9bkSUQnud9pFXULx6EbF7cbbw6i3ayvc2QNhgxTF",
+        103590: "2gTygYckZgFzFN5QOWPaPB03nabqjdV55mwy1x1Nd4JmJAwam",
+        103591: "2cUPPHy1hspr2nAKpQrrAEisLKkaWS59iF2wJfNyFRs8vnSkKk",
+        103594: "5Mpt5dP6dBMP5wK9GJjeVe39deZJTRh9i82cgN1bjedffrrTf",
+        103597: "2J8z7HNv4nwh82wqRGyEHqQeuw4J6mCDCSVUgusBu35asnshK",
+        103598: "2i2FP6nJyvhX9FR15qN2D9AvOkSXKgBD212A07FoSpfowxvQDX",
+        103603: "2v3smb35s4GLACsK4Zkd2RcLbLDWA4huqrvq8Y3VP4Cve8kFTM",
+        103604: "b7XFDDLgwb12DfL7UTWZoxwBpkLPL5mdHtXngD94Y2RoeWXSh",
    }
```

```

+         103607: "PgaRk1UAoUvRybhXsrLq5t6imWhEa6ksNjbN6hWgs4qPrSz",
+         103612: "ZoueNTJ4dUE2FFtGyPpawnmCCsy6EUQevHVLZy8NHeQmkAc1P4",
+         103614: "2YH21KymFjiBhpXzgt6HXJhLst5SV9UQ4tJuUNjfnInQ0dm5zz",
+         103617: "amgH2C1s9H3Av7vSW4y7n7TXb9tKyKHENvrDXutgNN6nsejgc",
+         103618: "fV8k1U8oQDmfVwK66kAwN73a5sW1whm8quNpVnKmSznBycV2w",
+         103621: "Nzs93kFTvcXanFUp9Y8VQkKynzmH8xykxVNFJTkdyaEeuxWbP",
+         103623: "2rAsBj3emqQa13CV8r5fTtHogs4sXnjvbbXVzcKPi3WmzhpK9D",
+         103624: "2JbuExUGKw5mYz5KfXATwq1lbRDimgks9wEdYGN5C6Ttey1R4U",
+         103627: "tLLij7oKfvWT1yk9zRv4FQvuQ5DAiuvb5kHCNN9zh4mqkFMG",
+         103628: "dWBsRYRwFrcy13DPdLoHsL67QkZ5h86hwtVfP94ZBaY18EkMF",
+         103629: "XMoEsew2Dh5gQaydcJfJUQAQYP8BTNTYbEJZvtbrV2QsX7iE3",
+         103630: "2db2wMbVaoCc5EUJrsBYWvNZDeqyY8uNpaaVapdBAQZ5oRaou",
+         103633: "2QiHZwLhQ3Luyyfcdo5yCUfoSqWdVRZox5ECU19HiswfrC6p",
+     }
+ }
+
+ for _, blkIDStr := range blockIDStrs {
+
+     for height, blkIDStr := range mainnetBonusBlocks {
+         blkID, err := ids.FromString(blkIDStr)
+         if err != nil {
+             panic(err)
+         }
+         bonusBlocks.Add(blkID)
+         bonusBlockMainnetHeights[height] = blkID
+     }
+ }
+
+ // Block implements the snowman.Block interface
+ type Block struct {
+     - id      ids.ID
+     - ethBlock *types.Block
+     - vm      *VM
+     - status  choices.Status
+     + id      ids.ID
+     + ethBlock *types.Block
+     + vm      *VM
+     + status  choices.Status
+     + atomicTxs []*Tx
+ }
+
+ // ID implements the snowman.Block interface
+ @@ -118,23 +134,29 @@ func (b *Block) Accept() error {
+     return fmt.Errorf("failed to put %s as the last accepted block: %w", b.ID(), err)
+ }
+
+ - tx, err := vm.extractAtomicTx(b.ethBlock)
+ + if len(b.atomicTxs) == 0 {
+ +     if err := b.vm.atomicTrie.Index(b.Height(), nil); err != nil {
+ +         return err
+ +     }
+ +     return vm.db.Commit()
+ + }
+
+ batchChainsAndInputs, err := mergeAtomicOps(b.atomicTxs)
+ if err != nil {
+     return err
+ }
+
+ - if tx == nil {
+ -     return vm.db.Commit()
+ + for _, tx := range b.atomicTxs {
+ +     // Remove the accepted transaction from the mempool
+ +     vm.mempool.RemoveTx(tx.ID())
+ + }
+
+ // Remove the accepted transaction from the mempool
+ vm.mempool.RemoveTx(tx.ID())
+
+ // Save the accepted atomic transaction
+ if err := vm.writeAtomicTx(b, tx); err != nil {
+     isBonus := bonusBlocks.Contains(b.id)
+     if err := b.indexAtomics(vm, b.Height(), b.atomicTxs, batchChainsAndInputs, isBonus); err != nil {
+         return err
+     }
+ }
+
+ - if bonusBlocks.Contains(b.id) {
+ + // If [b] is a bonus block, then we commit the database without applying the requests from
+ + // the atomic transactions to shared memory.
+ + if isBonus {
+ +     log.Info("skipping atomic tx acceptance on bonus block", "block", b.id)
+ +     return vm.db.Commit()
+ + }
+
+ @@ -143,8 +165,22 @@ func (b *Block) Accept() error {
+     if err != nil {
+         return fmt.Errorf("failed to create commit batch due to: %w", err)
+     }
+     return vm.ctx.SharedMemory.Apply(batchChainsAndInputs, batch)
+ }
+
+ - return tx.UnsignedAtomicTx.Accept(vm.ctx, batch)
+ +// indexAtomics writes given list of atomic transactions and atomic operations to atomic repository
+ +// and atomic trie respectively
+ +func (b *Block) indexAtomics(vm *VM, height uint64, atomicTxs []*Tx, batchChainsAndInputs map[ids.ID]*atomic.Requests, isBonus bool) error {
+ +     if isBonus {
+ +         // avoid indexing atomic operations of txs on bonus blocks in the trie
+ +         // so we do not re-execute them the second time that they appear
+ +         return vm.atomicTxRepository.WriteBonus(height, atomicTxs)
+ +     }
+
+     if err := vm.atomicTxRepository.Write(height, atomicTxs); err != nil {
+         return err
+     }
+
+     return b.vm.atomicTrie.Index(height, batchChainsAndInputs)
+ }
+
+ // Reject implements the snowman.Block interface
+ @@ -152,14 +188,12 @@ func (b *Block) Accept() error {
+     func (b *Block) Reject() error {
+         b.status = choices.Rejected
+         log.Debug(fmt.Sprintf("Rejecting block %s (%s) at height %d", b.ID().Hex(), b.ID(), b.Height()))
+
+         - tx, _ := b.vm.extractAtomicTx(b.ethBlock)
+         - if tx != nil {
+ +         for _, tx := range b.atomicTxs {
+ +             b.vm.mempool.RemoveTx(tx.ID())
+ +             if err := b.vm.issueTx(tx, false /* set local to false when re-issuing */); err != nil {
+ +                 log.Debug("Failed to re-issue transaction in rejected block", "txID", tx.ID(), "err", err)
+ +             }
+         }
+
+         return b.vm.chain.Reject(b.ethBlock)
+     }
+ }
+
+ @@ -179,7 +213,7 @@ func (b *Block) Parent() ids.ID {
+
+ // Height implements the snowman.Block interface
+ func (b *Block) Height() uint64 {
+     - return b.ethBlock.Number().Uint64()
+ +     return b.ethBlock.NumberU64()
+ }
+
+ // Timestamp implements the snowman.Block interface
+ @@ -209,31 +243,33 @@ func (b *Block) verify(writes bool) error {
+     return fmt.Errorf("syntactic block verification failed: %w", err)
+ }
+
+ - vm := b.vm

```



```

+         if err := b.verifyAtomicTxs(rules); err != nil {
+             return err
+         }
+     }
+     return b.vm.chain.BlockChain().InsertBlockManual(b.ethBlock, writes)
+}

+func (b *Block) verifyAtomicTxs(rules params.Rules) error {
+    // Ensure that the parent was verified and inserted correctly.
+    ancestorID := b.Parent()
+    ancestorHash := common.Hash(ancestorID)
+    if !vm.chain.BlockChain().HasBlock(ancestorHash, b.Height()-1) {
+    - if !b.vm.chain.BlockChain().HasBlock(ancestorHash, b.Height()-1) {
+        return errRejectedParent
+    }

+    // If the tx is an atomic tx, ensure that it doesn't conflict with any of
+    // its processing ancestry.
+    atomicTx, err := vm.extractAtomicTx(b.ethBlock)
+    if err != nil {
+    - return err
+    }
+    if atomicTx != nil {
+    + inputs := &ids.Set{}
+    + for _, atomicTx := range b.atomicTxs {
+        // If the ancestor is unknown, then the parent failed verification when
+        // it was called.
+        // If the ancestor is rejected, then this block shouldn't be inserted
+        // into the canonical chain because the parent is will be missing.
+    - ancestorInf, err := vm.GetBlockInternal(ancestorID)
+    + // into the canonical chain because the parent will be missing.
+    + ancestorInf, err := b.vm.GetBlockInternal(ancestorID)
+        if err != nil {
+            return errRejectedParent
+        }

+        if blkStatus := ancestorInf.Status(); blkStatus == choices.Unknown || blkStatus == choices.Rejected {
+            return errRejectedParent
+        }
+    }

@@ -241,19 +277,22 @@ func (b *Block) verify(writes bool) error {
+    if !ok {
+        return fmt.Errorf("expected %, parent of %, to be *Block but is %T", ancestor.ID(), b.ID(), ancestorInf)
+    }

+    if bonusBlocks.Contains(b.id) {
+        log.Info("Skipping atomic tx verification on bonus block", "block", b.id)
+    } else {
+        utx := atomicTx.UnsignedAtomicTx
+    - if err := utx.SemanticVerify(vm, atomicTx, ancestor, b.ethBlock.BaseFee(), rules); err != nil {
+    + if err := utx.SemanticVerify(b.vm, atomicTx, ancestor, b.ethBlock.BaseFee(), rules); err != nil {
+            return fmt.Errorf("invalid block due to failed semantic verify: %w at height %d", err, b.Height())
+        }
+        txInputs := utx.InputUTXOs()
+        if inputs.Overlaps(txInputs) {
+            return errConflictingAtomicInputs
+        }
+        inputs.Union(txInputs)
+    }
+    }

+    bc := vm.chain.BlockChain()
+    return bc.InsertBlockManual(b.ethBlock, writes)
+    return nil
+}

// Bytes implements the snowman.Block interface
diff --git a/plugin/evm/block_builder.go b/plugin/evm/block_builder.go
index d86bccd4..49b32b78 100644
--- a/plugin/evm/block_builder.go
+++ b/plugin/evm/block_builder.go
@@ -8,13 +8,13 @@ import (
     "sync"
     "time"

-    coreth "github.com/ava-labs/coreth/chain"
-    "github.com/ava-labs/coreth/params"
+    coreth "github.com/flare-foundation/coreth/chain"
+    "github.com/flare-foundation/coreth/params"

-    "github.com/ava-labs/avalanchego/snow"
-    commonEng "github.com/ava-labs/avalanchego/snow/engine/common"
-    "github.com/ava-labs/avalanchego/utlis/timer"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/flare/snow"
+    commonEng "github.com/flare-foundation/flare/snow/engine/common"
+    "github.com/flare-foundation/flare/utlis/timer"
 )

// buildingBlkStatus denotes the current status of the VM in block production.
@@ -52,9 +52,9 @@ type blockBuilder struct {
    ctx          *snow.Context
    chainConfig  *params.ChainConfig

-    chain    *coreth.ETHChain
-    mempool *Mempool
-    network Network
+    chain    *coreth.ETHChain
+    mempool *Mempool
+    gossip   Gossiper

    shutdownChan <-chan struct{}
    shutdownWg   *sync.WaitGroup
}

@@ -90,7 +90,7 @@ func (vm *VM) NewBlockBuilder(notifyBuildBlockChan chan<- commonEng.Message) *bl
    chainConfig:    vm.chainConfig,
    chain:          vm.chain,
    mempool:        vm.mempool,
-    network:        vm.network,
+    gossip:         vm.gossiper,
    shutdownChan:   vm.shutdownChan,
    shutdownWg:     &vm.shutdownWg,
    notifyBuildBlockChan: notifyBuildBlockChan,
}

@@ -279,13 +279,13 @@ func (b *blockBuilder) awaitSubmittedTxs() {
    b.signalTxsReady()

    // We only attempt to invoke [GossipEthTxs] once AP4 is activated
+    if b.isAP4 && b.network != nil && len(ethTxsEvent.Txs) > 0 {
+    if b.isAP4 && b.gossiper != nil && len(ethTxsEvent.Txs) > 0 {
+        // Give time for this node to build a block before attempting to
+        // gossip
+        time.Sleep(waitBlockTime)
+        // [GossipEthTxs] will block unless [pushNetwork.ethTxsToGossipChan] (an
+        // [GossipEthTxs] will block unless [gossiper.ethTxsToGossipChan] (an
+        // unbuffered channel) is listened on
+        if err := b.network.GossipEthTxs(ethTxsEvent.Txs); err != nil {
+        if err := b.gossiper.GossipEthTxs(ethTxsEvent.Txs); err != nil {
+            log.Warn(
+                "failed to gossip new eth transactions",
+                "err", err,
+            )
+        }

@@ -298,11 +298,11 @@ func (b *blockBuilder) awaitSubmittedTxs() {

    // We only attempt to invoke [GossipAtomicTxs] once AP4 is activated
    newTxs := b.mempool.GetNewTxs()
+    if b.isAP4 && b.network != nil && len(newTxs) > 0 {
+    if b.isAP4 && b.gossiper != nil && len(newTxs) > 0 {

```

```

// Give time for this node to build a block before attempting to
// gossip
time.Sleep(waitBlockTime)
- if err := b.network.GossipAtomicTxs(newTxs); err != nil {
+ if err := b.gossiper.GossipAtomicTxs(newTxs); err != nil {
    log.Warn(
        "failed to gossip new atomic transactions",
        "err", err,
    )
}

diff --git a/plugin/evm/block_builder_test.go b/plugin/evm/block_builder_test.go
index 95b3b18f..e6e54ecd 100644
--- a/plugin/evm/block_builder_test.go
+++ b/plugin/evm/block_builder_test.go
@@ -9,9 +9,9 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"

-    "github.com/ava-labs/avalanchego/snow"
+    "github.com/flare-foundation/flare/snow"
)

func TestBlockBuilderShutsDown(t *testing.T) {
diff --git a/plugin/evm/block_verification.go b/plugin/evm/block_verification.go
index 6833f68c..5a0360da 100644
--- a/plugin/evm/block_verification.go
+++ b/plugin/evm/block_verification.go
@@ -7,11 +7,15 @@ import (
    "fmt"
    "math/big"

-    coreth "github.com/ava-labs/coreth/chain"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/trie"
+    "github.com/ethereum/go-ethereum/common"
+
+    safemath "github.com/flare-foundation/flare/utills/math"
+
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/trie"
+
+    coreth "github.com/flare-foundation/coreth/chain"
)

var (
@@ -21,6 +25,7 @@ var (
    apricotPhase1MinGasPrice = big.NewInt(params.ApricotPhase1MinGasPrice)
    phase3BlockValidator      = blockValidatorPhase3{}
    phase4BlockValidator      = blockValidatorPhase4{}
+    phase5BlockValidator      = blockValidatorPhase5{}
)

type BlockValidator interface {
@@ -125,14 +130,8 @@ func (v blockValidatorPhase0) SyntacticVerify(b *Block) error {
    return errUnclesUnsupported
}

// Block must not be empty
//
// Note: extractAtomicTx also asserts a maximum size
atomicTx, err := b.vm.extractAtomicTx(b.ethBlock)
if err != nil {
    return err
}

txs := b.ethBlock.Transactions()
if len(txs) == 0 && atomicTx == nil {
+ if len(txs) == 0 && len(b.atomicTxs) == 0 {
    return errEmptyBlock
}

@@ -230,14 +229,8 @@ func (blockValidatorPhase1) SyntacticVerify(b *Block) error {
    return errUnclesUnsupported
}

// Block must not be empty
//
// Note: extractAtomicTx also asserts a maximum size
atomicTx, err := b.vm.extractAtomicTx(b.ethBlock)
if err != nil {
    return err
}

txs := b.ethBlock.Transactions()
if len(txs) == 0 && atomicTx == nil {
+ if len(txs) == 0 && len(b.atomicTxs) == 0 {
    return errEmptyBlock
}

@@ -340,14 +333,8 @@ func (blockValidatorPhase3) SyntacticVerify(b *Block) error {
    return errUnclesUnsupported
}

// Block must not be empty
//
// Note: extractAtomicTx also asserts a maximum size
atomicTx, err := b.vm.extractAtomicTx(b.ethBlock)
if err != nil {
    return err
}

txs := b.ethBlock.Transactions()
if len(txs) == 0 && atomicTx == nil {
+ if len(txs) == 0 && len(b.atomicTxs) == 0 {
    return errEmptyBlock
}

@@ -443,14 +430,8 @@ func (blockValidatorPhase4) SyntacticVerify(b *Block) error {
    return errUnclesUnsupported
}

// Block must not be empty
//
// Note: extractAtomicTx also asserts a maximum size
atomicTx, err := b.vm.extractAtomicTx(b.ethBlock)
if err != nil {
    return err
}

txs := b.ethBlock.Transactions()
if len(txs) == 0 && atomicTx == nil {
+ if len(txs) == 0 && len(b.atomicTxs) == 0 {
    return errEmptyBlock
}

@@ -465,28 +446,163 @@ func (blockValidatorPhase4) SyntacticVerify(b *Block) error {
    return errNilExtDataGasUsedApricotPhase4
}

if !ethHeader.ExtDataGasUsed.IsUint64() {
- return fmt.Errorf("too large extDataGasUsed : bitlen %d", ethHeader.ExtDataGasUsed.BitLen())
+ return fmt.Errorf("too large extDataGasUsed: %d", ethHeader.ExtDataGasUsed)
}

if atomicTx != nil {
+ var totalGasUsed uint64
+ for _, atomicTx := range b.atomicTxs {
+     // We perform this check manually here to avoid the overhead of having to
+     // reparse the atomicTx in `CalcExtDataGasUsed`.
+     gasUsed, err := atomicTx.GasUsed()
+     gasUsed, err := atomicTx.GasUsed(false)
}

```

```

        if err != nil {
            return err
        }
        if ethHeader.ExtDataGasUsed.Cmp(new(big.Int).SetUint64(gasUsed)) != 0 {
            return fmt.Errorf("invalid extDataGasUsed: have %d, want %d", ethHeader.ExtDataGasUsed, gasUsed)
        }
        totalGasUsed, err = safemath.Add64(totalGasUsed, gasUsed)
        if err != nil {
            return err
        }
    }
}

switch {
case ethHeader.ExtDataGasUsed.Cmp(new(big.Int).SetUint64(totalGasUsed)) != 0:
    return fmt.Errorf("invalid extDataGasUsed: have %d, want %d", ethHeader.ExtDataGasUsed, totalGasUsed)
}

// Make sure BlockGasCost is not nil
// NOTE: ethHeader.BlockGasCost correctness is checked in header verification
if ethHeader.BlockGasCost == nil {
case ethHeader.BlockGasCost == nil:
    return errNilBlockGasCostApricotPhase4
case !ethHeader.BlockGasCost.IsUint64():
    return fmt.Errorf("too large blockGasCost: %d", ethHeader.BlockGasCost)
}
if !ethHeader.BlockGasCost.IsUint64() {
    return fmt.Errorf("too large blockGasCost: bitlen %d", ethHeader.BlockGasCost.BitLen())
}
return nil
})
}

+type blockValidatorPhase5 struct{}
+
+func (blockValidatorPhase5) SyntacticVerify(b *Block) error {
+    if b == nil || b.ethBlock == nil {
+        return errInvalidBlock
+    }
+
+    // Skip verification of the genesis block since it
+    // should already be marked as accepted
+    if b.ethBlock.Hash() == b.vm.genesisHash {
+        return nil
+    }
+
+    // Perform block and header sanity checks
+    ethHeader := b.ethBlock.Header()
+    if ethHeader.Number == nil || !ethHeader.Number.IsUint64() {
+        return errInvalidBlock
+    }
+    if ethHeader.Difficulty == nil || !ethHeader.Difficulty.IsUint64() ||
+        ethHeader.Difficulty.Uint64() != 1 {
+        return fmt.Errorf(
+            "expected difficulty to be 1 but got %v: %w",
+            ethHeader.Difficulty, errInvalidDifficulty,
+        )
+    }
+    if ethHeader.Nonce.Uint64() != 0 {
+        return fmt.Errorf(
+            "expected nonce to be 0 but got %d: %w",
+            ethHeader.Nonce.Uint64(), errInvalidNonce,
+        )
+    }
+    if ethHeader.GasLimit != params.ApricotPhase5GasLimit {
+        return fmt.Errorf(
+            "expected gas limit to be %d in apricot phase 5 but got %d",
+            params.ApricotPhase5GasLimit, ethHeader.GasLimit,
+        )
+    }
+    if ethHeader.MixDigest != (common.Hash{}) {
+        return fmt.Errorf(
+            "expected MixDigest to be empty but got %x: %w",
+            ethHeader.MixDigest, errInvalidMixDigest,
+        )
+    }
+    if hash := types.CalcExtDataHash(b.ethBlock.ExtData()); ethHeader.ExtDataHash != hash {
+        return fmt.Errorf("extra data hash mismatch: have %x, want %x", ethHeader.ExtDataHash, hash)
+    }
+    if headerExtraDataSize := len(ethHeader.Extra); headerExtraDataSize != params.ApricotPhase3ExtraDataSize {
+        return fmt.Errorf(
+            "expected header ExtraData to be %d but got %d: %w",
+            params.ApricotPhase3ExtraDataSize, headerExtraDataSize, errHeaderExtraDataTooBig,
+        )
+    }
+    if ethHeader.BaseFee == nil {
+        return errNilBaseFeeApricotPhase3
+    }
+    if bfLen := ethHeader.BaseFee.BitLen(); bfLen > 256 {
+        return fmt.Errorf("too large base fee: bitlen %d", bfLen)
+    }
+    if b.ethBlock.Version() != 0 {
+        return fmt.Errorf(
+            "expected block version to be 0 but got %d: %w",
+            b.ethBlock.Version(), errInvalidBlockVersion,
+        )
+    }
+
+    // Check that the tx hash in the header matches the body
+    txsHash := types.DeriveSha(b.ethBlock.Transactions(), new(trie.Trie))
+    if txsHash != ethHeader.TxHash {
+        return errTxHashMismatch
+    }
+    // Check that the uncle hash in the header matches the body
+    unclesHash := types.CalcUncleHash(b.ethBlock.Uncles())
+    if unclesHash != ethHeader.UncleHash {
+        return errUncleHashMismatch
+    }
+    // Coinbase must be zero on C-Chain
+    if b.ethBlock.Coinbase() != coreth.BlackholeAddr {
+        return errInvalidBlock
+    }
+    // Block must not have any uncles
+    if len(b.ethBlock.Uncles()) > 0 {
+        return errUnclesUnsupported
+    }
+    // Block must not be empty
+    txs := b.ethBlock.Transactions()
+    if len(txs) == 0 && len(b.atomicTxs) == 0 {
+        return errEmptyBlock
+    }
+
+    // Make sure the block isn't too far in the future
+    blockTimestamp := b.ethBlock.Time()
+    if maxBlockTime := uint64(b.vm.clock.Time()).Add(maxFutureBlockTime).Unix(); blockTimestamp > maxBlockTime {
+        return fmt.Errorf("block timestamp is too far in the future: %d > allowed %d", blockTimestamp, maxBlockTime)
+    }
+
+    // Make sure ExtDataGasUsed is not nil and correct
+    if ethHeader.ExtDataGasUsed == nil {
+        return errNilExtDataGasUsedApricotPhase4
+    }
+    if ethHeader.ExtDataGasUsed.Cmp(params.AtomicGasLimit) == 1 {
+        return fmt.Errorf("too large extDataGasUsed: %d", ethHeader.ExtDataGasUsed)
+    }
+
+    var totalGasUsed uint64
+    for _, atomicTx := range b.atomicTxs {
+        // We perform this check manually here to avoid the overhead of having to

```

```

+         // reparse the atomicTx in `CalcExtDataGasUsed`.
+         gasUsed, err := atomicTx.GasUsed(true)
+         if err != nil {
+             return err
+         }
+         totalGasUsed, err = safemath.Add64(totalGasUsed, gasUsed)
+         if err != nil {
+             return err
+         }
+     }
+
+     switch {
+     case ethHeader.ExtDataGasUsed.Cmp(new(big.Int).SetUint64(totalGasUsed)) != 0:
+         return fmt.Errorf("invalid extDataGasUsed: have %d, want %d", ethHeader.ExtDataGasUsed, totalGasUsed)
+
+         // Make sure BlockGasCost is not nil
+         // NOTE: ethHeader.BlockGasCost correctness is checked in header verification
+     case ethHeader.BlockGasCost == nil:
+         return errNilBlockGasCostApricotPhase4
+     case !ethHeader.BlockGasCost.IsUint64():
+         return fmt.Errorf("too large blockGasCost: %d", ethHeader.BlockGasCost)
+     }
+     return nil
+ }
+ }
+
diff --git a/plugin/evm/client.go b/plugin/evm/client.go
index 4a33485d..3a37c053 100644
--- a/plugin/evm/client.go
+++ b/plugin/evm/client.go
@@ -4,44 +4,66 @@
package evm

import (
+     "context"
+     "fmt"
-     "time"

-     "github.com/ava-labs/avalanchego/api"
-     "github.com/ava-labs/avalanchego/ids"
-     "github.com/ava-labs/avalanchego/utils/formatting"
-     cjson "github.com/ava-labs/avalanchego/utils/json"
-     "github.com/ava-labs/avalanchego/utils/rpc"
-     "github.com/ethereum/go-ethereum/log"
+     "github.com/flare-foundation/flare/api"
+     "github.com/flare-foundation/flare/ids"
+     "github.com/flare-foundation/flare/utils/formatting"
+     cjson "github.com/flare-foundation/flare/utils/json"
+     "github.com/flare-foundation/flare/utils/rpc"
)

// Client ...
-type Client struct {
+// Interface compliance
+var _ Client = (*client)(nil)
+
+// Client interface for interacting with EVM [chain]
+type Client interface {
+     IssueTx(ctx context.Context, txBytes []byte) (ids.ID, error)
+     GetAtomicTxStatus(ctx context.Context, txID ids.ID) (Status, error)
+     GetAtomicTx(ctx context.Context, txID ids.ID) ([]byte, error)
+     GetAtomicUTXOs(ctx context.Context, addrs []string, sourceChain string, limit uint32, startAddress, startUTXOID string) ([][]byte, api.Index, error)
+     ListAddresses(ctx context.Context, userPass api.UserPass) ([]string, error)
+     ExportKey(ctx context.Context, userPass api.UserPass, addr string) (string, string, error)
+     ImportKey(ctx context.Context, userPass api.UserPass, privateKey string) (string, error)
+     Import(ctx context.Context, userPass api.UserPass, to string, sourceChain string) (ids.ID, error)
+     ExportAVAX(ctx context.Context, userPass api.UserPass, amount uint64, to string) (ids.ID, error)
+     Export(ctx context.Context, userPass api.UserPass, amount uint64, to string, assetID string) (ids.ID, error)
+     StartCPUProfiler(ctx context.Context) (bool, error)
+     StopCPUProfiler(ctx context.Context) (bool, error)
+     MemoryProfile(ctx context.Context) (bool, error)
+     LockProfile(ctx context.Context) (bool, error)
+     SetLogLevel(ctx context.Context, level log.Lvl) (bool, error)
+ }
+
+// Client implementation for interacting with EVM [chain]
+type client struct {
+     requester      rpc.EndpointRequester
+     adminRequester rpc.EndpointRequester
+ }

// NewClient returns a Client for interacting with EVM [chain]
-func NewClient(uri, chain string, requestTimeout time.Duration) *Client {
-     return &Client{
-         requester:      rpc.NewEndpointRequester(uri, fmt.Sprintf("/ext/bc/%s/avax", chain), "avax", requestTimeout),
-         adminRequester: rpc.NewEndpointRequester(uri, fmt.Sprintf("/ext/bc/%s/admin", chain), "admin", requestTimeout),
-     }
+func NewClient(uri, chain string) Client {
+     return &client{
+         requester:      rpc.NewEndpointRequester(uri, fmt.Sprintf("/ext/bc/%s/avax", chain), "avax"),
+         adminRequester: rpc.NewEndpointRequester(uri, fmt.Sprintf("/ext/bc/%s/admin", chain), "admin"),
+     }
+ }

// NewCChainClient returns a Client for interacting with the C Chain
-func NewCChainClient(uri string, requestTimeout time.Duration) *Client {
-     return NewClient(uri, "C", requestTimeout)
+func NewCChainClient(uri string) Client {
+     return NewClient(uri, "C")
+ }

// IssueTx issues a transaction to a node and returns the TxID
-func (c *Client) IssueTx(txBytes []byte) (ids.ID, error) {
+func (c *client) IssueTx(txBytes []byte) (ids.ID, error) {
+     res := &api.JSONTxID{}
+     txStr, err := formatting.EncodeWithChecksum(formatting.Hex, txBytes)
+     if err != nil {
+         return res.TxID, fmt.Errorf("problem hex encoding bytes: %w", err)
+     }
+     err = c.requester.SendRequest("issueTx", &api.FormattedTx{
+         err = c.requester.SendRequest(ctx, "issueTx", &api.FormattedTx{
+             Tx:      txStr,
+             Encoding: formatting.Hex,
+         }, res)
@@ -49,18 +71,18 @@ func (c *Client) IssueTx(txBytes []byte) (ids.ID, error) {
}

// GetAtomicTxStatus returns the status of [txID]
-func (c *Client) GetAtomicTxStatus(txID ids.ID) (Status, error) {
+func (c *client) GetAtomicTxStatus(ctx context.Context, txID ids.ID) (Status, error) {
+     res := &GetAtomicTxStatusReply{}
+     err = c.requester.SendRequest("getAtomicTxStatus", &api.JSONTxID{
+         err = c.requester.SendRequest(ctx, "getAtomicTxStatus", &api.JSONTxID{
+             TxID: txID,
+         }, res)
+     return res.Status, err
+ }

// GetAtomicTx returns the byte representation of [txID]
-func (c *Client) GetAtomicTx(txID ids.ID) ([]byte, error) {
+func (c *client) GetAtomicTx(ctx context.Context, txID ids.ID) ([]byte, error) {
+     res := &api.FormattedTx{}
+     err = c.requester.SendRequest("getAtomicTx", &api.GetTxArgs{
+         err = c.requester.SendRequest(ctx, "getAtomicTx", &api.GetTxArgs{
+             TxID:      txID,
+             Encoding: formatting.Hex,
+         }, res)

```

```

@@ -73,9 +95,9 @@ func (c *Client) GetAtomicTx(txID ids.ID) ([]byte, error) {

// GetAtomicUTXOs returns the byte representation of the atomic UTXOs controlled by [addresses]
// from [sourceChain]
-func (c *Client) GetAtomicUTXOs(addr []string, sourceChain string, limit uint32, startAddress, startUTXOID string) ([]byte, api.Index, error) {
+func (c *Client) GetAtomicUTXOs(ctx context.Context, addrs []string, sourceChain string, limit uint32, startAddress, startUTXOID string) ([]byte, api.Index, error) {
    res := &api.GetUTXOsReply{}
    err := c.requester.SendRequest("getUTXOs", &api.GetUTXOsArgs{
+    err := c.requester.SendRequest(ctx, "getUTXOs", &api.GetUTXOsArgs{
        Addresses:  addrs,
        SourceChain: sourceChain,
        Limit:        cjson.Uint32(limit),
@@ -101,17 +123,17 @@ func (c *Client) GetAtomicUTXOs(addr []string, sourceChain string, limit uint32
    }

// ListAddresses returns all addresses on this chain controlled by [user]
-func (c *Client) ListAddresses(user api.UserPass) ([]string, error) {
+func (c *Client) ListAddresses(ctx context.Context, user api.UserPass) ([]string, error) {
    res := &api.JSONAddresses{}
    err := c.requester.SendRequest("listAddresses", &user, res)
+    err := c.requester.SendRequest(ctx, "listAddresses", &user, res)
    return res.Addresses, err
}

// ExportKey returns the private key corresponding to [addr] controlled by [user]
// in both Avalanche standard format and hex format
-func (c *Client) ExportKey(user api.UserPass, addr string) (string, string, error) {
+func (c *Client) ExportKey(ctx context.Context, user api.UserPass, addr string) (string, string, error) {
    res := &ExportKeyReply{}
    err := c.requester.SendRequest("exportKey", &ExportKeyArgs{
+    err := c.requester.SendRequest(ctx, "exportKey", &ExportKeyArgs{
        UserPass: user,
        Address:  addr,
    }, res)
@@ -119,9 +141,9 @@ func (c *Client) ExportKey(user api.UserPass, addr string) (string, string, erro
    }

// ImportKey imports [privateKey] to [user]
-func (c *Client) ImportKey(user api.UserPass, privateKey string) (string, error) {
+func (c *Client) ImportKey(ctx context.Context, user api.UserPass, privateKey string) (string, error) {
    res := &api.JSONAddress{}
    err := c.requester.SendRequest("importKey", &ImportKeyArgs{
+    err := c.requester.SendRequest(ctx, "importKey", &ImportKeyArgs{
        UserPass:  user,
        PrivateKey: privateKey,
    }, res)
@@ -130,9 +152,9 @@ func (c *Client) ImportKey(user api.UserPass, privateKey string) (string, error)

// Import sends an import transaction to import funds from [sourceChain] and
// returns the ID of the newly created transaction
-func (c *Client) Import(user api.UserPass, to, sourceChain string) (ids.ID, error) {
+func (c *Client) Import(ctx context.Context, user api.UserPass, to, sourceChain string) (ids.ID, error) {
    res := &api.JSONTxID{}
    err := c.requester.SendRequest("import", &ImportArgs{
+    err := c.requester.SendRequest(ctx, "import", &ImportArgs{
        UserPass:  user,
        To:         to,
        SourceChain: sourceChain,
@@ -142,25 +164,27 @@ func (c *Client) Import(user api.UserPass, to, sourceChain string) (ids.ID, erro

// ExportAVAX sends AVAX from this chain to the address specified by [to].
// Returns the ID of the newly created atomic transaction
-func (c *Client) ExportAVAX(
+func (c *Client) ExportAVAX(
+    ctx context.Context,
    user api.UserPass,
    amount uint64,
    to string,
) (ids.ID, error) {
-    return c.Export(user, amount, to, "AVAX")
+    return c.Export(ctx, user, amount, to, "AVAX")
}

// Export sends an asset from this chain to the P/C-Chain.
// After this tx is accepted, the AVAX must be imported to the P/C-chain with an importTx.
// Returns the ID of the newly created atomic transaction
-func (c *Client) Export(
+func (c *Client) Export(
+    ctx context.Context,
    user api.UserPass,
    amount uint64,
    to string,
    assetID string,
) (ids.ID, error) {
    res := &api.JSONTxID{}
    err := c.requester.SendRequest("export", &ExportArgs{
+    err := c.requester.SendRequest(ctx, "export", &ExportArgs{
        ExportAVAXArgs: ExportAVAXArgs{
            UserPass: user,
            Amount:   cjson.Uint64(amount),
@@ -171,34 +193,34 @@ func (c *Client) Export(
    return res.TxID, err
}

-func (c *Client) StartCPUProfiler() (bool, error) {
+func (c *Client) StartCPUProfiler(ctx context.Context) (bool, error) {
    res := &api.SuccessResponse{}
    err := c.adminRequester.SendRequest("startCPUProfiler", struct{}{}, res)
+    err := c.adminRequester.SendRequest(ctx, "startCPUProfiler", struct{}{}, res)
    return res.Success, err
}

-func (c *Client) StopCPUProfiler() (bool, error) {
+func (c *Client) StopCPUProfiler(ctx context.Context) (bool, error) {
    res := &api.SuccessResponse{}
    err := c.adminRequester.SendRequest("stopCPUProfiler", struct{}{}, res)
+    err := c.adminRequester.SendRequest(ctx, "stopCPUProfiler", struct{}{}, res)
    return res.Success, err
}

-func (c *Client) MemoryProfile() (bool, error) {
+func (c *Client) MemoryProfile(ctx context.Context) (bool, error) {
    res := &api.SuccessResponse{}
    err := c.adminRequester.SendRequest("memoryProfile", struct{}{}, res)
+    err := c.adminRequester.SendRequest(ctx, "memoryProfile", struct{}{}, res)
    return res.Success, err
}

-func (c *Client) LockProfile() (bool, error) {
+func (c *Client) LockProfile(ctx context.Context) (bool, error) {
    res := &api.SuccessResponse{}
    err := c.adminRequester.SendRequest("lockProfile", struct{}{}, res)
+    err := c.adminRequester.SendRequest(ctx, "lockProfile", struct{}{}, res)
    return res.Success, err
}

// SetLogLevel dynamically sets the log level for the C Chain
-func (c *Client) SetLogLevel(level log.Lvl) (bool, error) {
+func (c *Client) SetLogLevel(ctx context.Context, level log.Lvl) (bool, error) {
    res := &api.SuccessResponse{}
    err := c.adminRequester.SendRequest("setLogLevel", &SetLogLevelArgs{
+    err := c.adminRequester.SendRequest(ctx, "setLogLevel", &SetLogLevelArgs{
        Level: level.String(),
    }, res)
}

```

```

        return res.Success, err
diff --git a/plugin/evm/client_interface_test.go b/plugin/evm/client_interface_test.go
new file mode 100644
index 00000000..d88c4926
--- /dev/null
+++ b/plugin/evm/client_interface_test.go
@@ -0,0 +1,17 @@
+package evm
+
+import (
+    "reflect"
+    "testing"
+)
+
+func TestInterfaceStructOneToOne(t *testing.T) {
+    // checks struct provides at least the methods signatures in the interface
+    var _ Client = (*client)(nil)
+    // checks interface and struct have the same number of methods
+    clientType := reflect.TypeOf(&client{})
+    ClientType := reflect.TypeOf((*Client)(nil)).Elem()
+    if clientType.NumMethod() != ClientType.NumMethod() {
+        t.Fatalf("no 1 to 1 compliance between struct methods (%v) and interface methods (%v)", clientType.NumMethod(), ClientType.NumMethod())
+    }
+}
diff --git a/plugin/evm/codec.go b/plugin/evm/codec.go
index 3a49eb8e..cbc5365c 100644
--- a/plugin/evm/codec.go
+++ b/plugin/evm/codec.go
@@ -4,10 +4,12 @@
package evm

import (
-    "github.com/ava-labs/avalanchego/codec"
-    "github.com/ava-labs/avalanchego/codec/linearcodec"
-    "github.com/ava-labs/avalanchego/stdlib/wrappers"
-    "github.com/ava-labs/avalanchego/vms/secp256k1fx"
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/codec/linearcodec"
+    "github.com/flare-foundation/flare/stdlib/wrappers"
+    "github.com/flare-foundation/flare/vms/secp256k1fx"
+    "fmt"
)

// Codec does serialization and deserialization
@@ -38,3 +40,60 @@ func init() {
    panic(errs.Err)
}

+
+// extractAtomicTxs returns the atomic transactions in [atomicTxBytes] if
+// they exist.
+// if [batch] is true, it attempts to unmarshal [atomicTxBytes] as a slice of
+// transactions (post-ApricotPhase5), and if it is false, then it unmarshals
+// it as a single atomic transaction.
+func ExtractAtomicTxs(atomicTxBytes []byte, batch bool, codec codec.Manager) ([]*Tx, error) {
+    if len(atomicTxBytes) == 0 {
+        return nil, nil
+    }
+
+    if !batch {
+        tx, err := ExtractAtomicTx(atomicTxBytes, codec)
+        if err != nil {
+            return nil, err
+        }
+        return []*Tx{tx}, err
+    }
+    return ExtractAtomicTxsBatch(atomicTxBytes, codec)
+}
+
+// [ExtractAtomicTx] extracts a singular atomic transaction from [atomicTxBytes]
+// and returns a slice of atomic transactions for compatibility with the type returned post
+// ApricotPhase5.
+// Note: this function assumes [atomicTxBytes] is non-empty.
+func ExtractAtomicTx(atomicTxBytes []byte, codec codec.Manager) (*Tx, error) {
+    atomicTx := new(Tx)
+    if _, err := codec.Unmarshal(atomicTxBytes, atomicTx); err != nil {
+        return nil, fmt.Errorf("failed to unmarshal atomic transaction (pre-AP5): %w", err)
+    }
+    if err := atomicTx.Sign(codec, nil); err != nil {
+        return nil, fmt.Errorf("failed to initialize singleton atomic tx due to: %w", err)
+    }
+    return atomicTx, nil
+}
+
+// [ExtractAtomicTxsBatch] extracts a slice of atomic transactions from [atomicTxBytes].
+// Note: this function assumes [atomicTxBytes] is non-empty.
+func ExtractAtomicTxsBatch(atomicTxBytes []byte, codec codec.Manager) ([]*Tx, error) {
+    var atomicTxs []*Tx
+    if _, err := codec.Unmarshal(atomicTxBytes, &atomicTxs); err != nil {
+        return nil, fmt.Errorf("failed to unmarshal atomic tx (AP5) due to %w", err)
+    }
+
+    // Do not allow non-empty extra data field to contain zero atomic transactions. This would allow
+    // people to construct a block that contains useless data.
+    if len(atomicTxs) == 0 {
+        return nil, errs.MissingAtomicTxs
+    }
+
+    for index, atx := range atomicTxs {
+        if err := atx.Sign(codec, nil); err != nil {
+            return nil, fmt.Errorf("failed to initialize atomic tx at index %d: %w", index, err)
+        }
+    }
+    return atomicTxs, nil
+}
diff --git a/plugin/evm/config.go b/plugin/evm/config.go
index 9e3eea4c..4bb98ff8 100644
--- a/plugin/evm/config.go
+++ b/plugin/evm/config.go
@@ -7,28 +7,40 @@
@@ -7,28 +7,40 @@
import (
    "encoding/json"
    "time"

    "github.com/ava-labs/coreth/eth"
    "github.com/flare-foundation/coreth/eth"
    "github.com/spf13/cast"
)

const (
-    defaultEthApiEnabled           = true
-    defaultNetApiEnabled           = true
-    defaultWeb3ApiEnabled          = true
-    defaultPruningEnabled          = true
-    defaultSnapshotAsync           = true
-    defaultRpcGasCap               = 2500000000 // 25000000 X 100
-    defaultRpcTxFeeCap             = 100       // 100 AVAX
-    defaultApiMaxDuration          = 0         // Default to no maximum API call duration
-    defaultWsCpuRefillRate         = 0         // Default to no maximum WS CPU usage
-    defaultWsCpuMaxStored          = 0         // Default to no maximum WS CPU usage
-    defaultMaxBlocksPerRequest     = 0         // Default to no maximum on the number of blocks per getLogs request
-    defaultContinuousProfilerFrequency = 15 * time.Minute
-    defaultContinuousProfilerMaxFiles = 5
-    defaultTxRegossipFrequency     = 1 * time.Minute
+
+    defaultEthApiEnabled           = true
+    defaultNetApiEnabled           = true
+    defaultWeb3ApiEnabled          = true
+    defaultPruningEnabled          = true
+    defaultSnapshotAsync           = true
+    defaultRpcGasCap               = 2500000000 // 25000000 X 100
+    defaultRpcTxFeeCap             = 100       // 100 AVAX
+    defaultApiMaxDuration          = 0         // Default to no maximum API call duration
+    defaultWsCpuRefillRate         = 0         // Default to no maximum WS CPU usage
+    defaultWsCpuMaxStored          = 0         // Default to no maximum WS CPU usage
+    defaultMaxBlocksPerRequest     = 0         // Default to no maximum on the number of blocks per getLogs request
+    defaultContinuousProfilerFrequency = 15 * time.Minute
+    defaultContinuousProfilerMaxFiles = 5
+    defaultTxRegossipFrequency     = 1 * time.Minute
)

```

```

-     defaultTxRegossipMaxSize           = 15
+     defaultPruningEnabled               = true
+     defaultSnapshotAsync                = true
+     defaultRpcGasCap                    = 120_000_000 // Default to 120M Gas Limit
+     defaultRpcTxFeeCap                  = 100         // 100 AVAX
+     defaultMetricsEnabled               = true
+     defaultMetricsExpensiveEnabled      = false
+     defaultApiMaxDuration               = 0 // Default to no maximum API call duration
+     defaultWsCpuRefillRate              = 0 // Default to no maximum WS CPU usage
+     defaultWsCpuMaxStored               = 0 // Default to no maximum WS CPU usage
+     defaultMaxBlocksPerRequest          = 0 // Default to no maximum on the number of blocks per getLogs request
+     defaultContinuousProfilerFrequency  = 15 * time.Minute
+     defaultContinuousProfilerMaxFiles   = 5
+     defaultTxRegossipFrequency          = 1 * time.Minute
+     defaultTxRegossipMaxSize           = 15
+     defaultOfflinePruningBloomFilterSize uint64 = 512 // Default size (MB) for the offline pruner to use
+     defaultLogLevel                     = "info"
+     defaultMaxOutboundActiveRequests    = 8
+ }
+
+var defaultEnabledAPIs = []string{
+    "public-eth",
+    "public-eth-filter",
+    "net",
+    "web3",
+    "internal-public-eth",
+    "internal-public-blockchain",
+    "internal-public-transaction-pool",
+}
+
+type Duration struct {
+    time.Duration
+}
+
+@@ -36,9 +48,13 @@ type Duration struct {
+ // Config ...
+ type Config struct {
+     // Coreth APIs
+     - SnowmanAPIEnabled bool `json:"snowman-api-enabled"`
+     - CorethAdminAPIEnabled bool `json:"coreth-admin-api-enabled"`
+     - NetAPIEnabled bool `json:"net-api-enabled"`
+     + SnowmanAPIEnabled bool `json:"snowman-api-enabled"`
+     + CorethAdminAPIEnabled bool `json:"coreth-admin-api-enabled"`
+     + CorethAdminAPIDir string `json:"coreth-admin-api-dir"`
+
+     // EnabledEthAPIs is a list of Ethereum services that should be enabled
+     // If none is specified, then we use the default list [defaultEnabledAPIs]
+     EnabledEthAPIs []string `json:"eth-apis"`
+
+     // Continuous Profiler
+     ContinuousProfilerDir string `json:"continuous-profiler-dir"` // If set to non-empty string creates a continuous profiler
+
+@@ -49,19 +65,16 @@ type Config struct {
+     RPCGasCap uint64 `json:"rpc-gas-cap"`
+     RPCTxFeeCap float64 `json:"rpc-tx-fee-cap"`
+
+     // Eth APIs
+     - EthAPIEnabled bool `json:"eth-api-enabled"`
+     - PersonalAPIEnabled bool `json:"personal-api-enabled"`
+     - TxPoolAPIEnabled bool `json:"tx-pool-api-enabled"`
+     - DebugAPIEnabled bool `json:"debug-api-enabled"`
+     - Web3APIEnabled bool `json:"web3-api-enabled"`
+
+     // Eth Settings
+     Preimages bool `json:"preimages-enabled"`
+     Pruning bool `json:"pruning-enabled"`
+     SnapshotAsync bool `json:"snapshot-async"`
+     SnapshotVerify bool `json:"snapshot-verification-enabled"`
+
+     // Metric Settings
+     + MetricsEnabled bool `json:"metrics-enabled"`
+     + MetricsExpensiveEnabled bool `json:"metrics-expensive-enabled"`
+
+     // API Settings
+     LocalTxsEnabled bool `json:"local-txs-enabled"`
+     APIMaxDuration Duration `json:"api-max-duration"`
+
+@@ -83,26 +96,23 @@ type Config struct {
+
+     // Log level
+     LogLevel string `json:"log-level"`
+
+     // Offline Pruning Settings
+     + OfflinePruning bool `json:"offline-pruning-enabled"`
+     + OfflinePruningBloomFilterSize uint64 `json:"offline-pruning-bloom-filter-size"`
+     + OfflinePruningDataDirectory string `json:"offline-pruning-data-directory"`
+
+     // VM2VM network
+     + MaxOutboundActiveRequests int64 `json:"max-outbound-active-requests"`
+ }
+
+ // EthAPIs returns an array of strings representing the Eth APIs that should be enabled
+ func (c Config) EthAPIs() []string {
+     - ethAPIs := make([]string, 0)
+
+     - if c.EthAPIEnabled {
+     -     ethAPIs = append(ethAPIs, "eth")
+     - }
+     - if c.PersonalAPIEnabled {
+     -     ethAPIs = append(ethAPIs, "personal")
+     - }
+     - if c.TxPoolAPIEnabled {
+     -     ethAPIs = append(ethAPIs, "txpool")
+     - }
+     - if c.DebugAPIEnabled {
+     -     ethAPIs = append(ethAPIs, "debug")
+     - }
+
+     - return ethAPIs
+     + return c.EnabledEthAPIs
+ }
+
+ func (c Config) EthBackendSettings() eth.Settings {
+@@ -110,11 +116,11 @@ func (c Config) EthBackendSettings() eth.Settings {
+ }
+
+ func (c *Config) SetDefaults() {
+     - c.EthAPIEnabled = defaultEthApiEnabled
+     - c.NetAPIEnabled = defaultNetApiEnabled
+     - c.Web3APIEnabled = defaultWeb3ApiEnabled
+     + c.EnabledEthAPIs = defaultEnabledAPIs
+     c.RPCGasCap = defaultRpcGasCap
+     c.RPCTxFeeCap = defaultRpcTxFeeCap
+     + c.MetricsEnabled = defaultMetricsEnabled
+     + c.MetricsExpensiveEnabled = defaultMetricsExpensiveEnabled
+     c.APIMaxDuration.Duration = defaultApiMaxDuration
+     c.WSCPURefillRate.Duration = defaultWsCpuRefillRate
+     c.WSCPUMaxStored.Duration = defaultWsCpuMaxStored
+@@ -125,6 +131,9 @@ func (c *Config) SetDefaults() {
+     c.SnapshotAsync = defaultSnapshotAsync
+     c.TxRegossipFrequency.Duration = defaultTxRegossipFrequency
+     c.TxRegossipMaxSize = defaultTxRegossipMaxSize
+     + c.OfflinePruningBloomFilterSize = defaultOfflinePruningBloomFilterSize
+     + c.LogLevel = defaultLogLevel
+     + c.MaxOutboundActiveRequests = defaultMaxOutboundActiveRequests
+ }

```

```

func (d *Duration) UnmarshalJSON(data []byte) (err error) {
diff --git a/plugin/evm/database.go b/plugin/evm/database.go
index c86ae5ee..53441e06 100644
--- a/plugin/evm/database.go
+++ b/plugin/evm/database.go
@@ -4,9 +4,9 @@
package evm

import (
-     "github.com/ava-labs/coreth/ethdb"
+     "github.com/flare-foundation/coreth/ethdb"

-     "github.com/ava-labs/avalanchego/database"
+     "github.com/flare-foundation/flare/database"
)

// Database implements ethdb.Database
diff --git a/plugin/evm/export_tx.go b/plugin/evm/export_tx.go
index fcc8c38f..7b241f7b 100644
--- a/plugin/evm/export_tx.go
+++ b/plugin/evm/export_tx.go
@@ -7,20 +7,21 @@
import (
    "fmt"
    "math/big"

-     "github.com/ava-labs/coreth/core/state"
-     "github.com/ava-labs/coreth/params"
+     "github.com/flare-foundation/coreth/core/state"
+     "github.com/flare-foundation/coreth/params"

-     "github.com/ava-labs/avalanchego/chains/atomic"
-     "github.com/ava-labs/avalanchego/database"
-     "github.com/ava-labs/avalanchego/ids"
-     "github.com/ava-labs/avalanchego/snow"
-     "github.com/ava-labs/avalanchego/utills/crypto"
-     "github.com/ava-labs/avalanchego/utills/math"
-     "github.com/ava-labs/avalanchego/utills/wrappers"
-     "github.com/ava-labs/avalanchego/vms/components/avax"
-     "github.com/ava-labs/avalanchego/vms/secp256k1fx"
+     "github.com/ethereum/go-ethereum/common"
+     "github.com/ethereum/go-ethereum/log"
+     "github.com/flare-foundation/flare/chains/atomic"
+     "github.com/flare-foundation/flare/ids"
+     "github.com/flare-foundation/flare/snow"
+     "github.com/flare-foundation/flare/utills/constants"
+     "github.com/flare-foundation/flare/utills/crypto"
+     "github.com/flare-foundation/flare/utills/math"
+     "github.com/flare-foundation/flare/utills/wrappers"
+     "github.com/flare-foundation/flare/vms/components/avax"
+     "github.com/flare-foundation/flare/vms/components/verify"
+     "github.com/flare-foundation/flare/vms/secp256k1fx"
)

// UnsignedExportTx is an unsigned ExportTx
@@ -40,94 +41,94 @@ type UnsignedExportTx struct {

// InputUTXOs returns a set of all the hash(address:nonce) exporting funds.
func (tx *UnsignedExportTx) InputUTXOs() ids.Set {
-     set := ids.NewSet(len(tx.Ins))
-     for _, in := range tx.Ins {
-         // Total populated bytes is 20 (Address) + 8 (Nonce), however, we allocate
-         // 32 bytes to make ids.ID casting easier.
-         var rawID [32]byte
-         packer := wrappers.Packer{Bytes: rawID[:]}
-         packer.PackLong(in.Nonce)
-         packer.PackBytes(in.Address.Bytes())
-         set.Add(ids.ID(rawID))
-     }
-     return set
+     return ids.Set{}
}

// Verify this transaction is well-formed
func (tx *UnsignedExportTx) Verify(
-     xChainID ids.ID,
-     ctx *snow.Context,
-     rules params.Rules,
) error {
-     switch {
-     case tx == nil:
-         return errNilTx
-     case tx.DestinationChain != xChainID:
-         return errWrongChainID
-     case len(tx.ExportedOutputs) == 0:
-         return errNoExportOutputs
-     case tx.NetworkID != ctx.NetworkID:
-         return errWrongNetworkID
-     case ctx.ChainID != tx.BlockchainID:
-         return errWrongBlockchainID
-     }

-     for _, in := range tx.Ins {
-         if err := in.Verify(); err != nil {
-             return err
-         }
-     }

-     for _, out := range tx.ExportedOutputs {
-         if err := out.Verify(); err != nil {
-             return err
-         }
-     }

-     if !avax.IsSortedTransferableOutputs(tx.ExportedOutputs, Codec) {
-         return errOutputsNotSorted
-     }

-     if rules.IsApricotPhase1 && !IsSortedAndUniqueEVMInputs(tx.Ins) {
-         return errInputsNotSortedUnique
-     }

-     return nil
+     return errExportTxsDisabled
}

-func (tx *UnsignedExportTx) GasUsed() (uint64, error) {
-     byteCost := calcBytesCost(len(tx.UnsignedBytes()))
-     numSigs := uint64(len(tx.Ins))
-     sigCost, err := math.Mul64(numSigs, secp256k1fx.CostPerSignature)
-     if err != nil {
-         return 0, err
-     }
-     return math.Add64(byteCost, sigCost)
+func (tx *UnsignedExportTx) GasUsed(fixedFee bool) (uint64, error) {
+     return 0, errExportTxsDisabled
}

// Amount of [assetID] burned by this transaction
func (tx *UnsignedExportTx) Burned(assetID ids.ID) (uint64, error) {
-     var (
-         spent uint64
-         input uint64
-         err error
-     )
-     for _, out := range tx.ExportedOutputs {
-         if out.AssetId() == assetID {

```



```

-         spent, err = math.Add64(spent, out.Output().Amount())
-         if err != nil {
-             return 0, err
-         }
-     }
- }
- for _, in := range tx.Ins {
-     if in.AssetId == assetID {
-         input, err = math.Add64(input, in.Amount)
-         if err != nil {
-             return 0, err
-         }
-     }
- }
- return math.Sub64(input, spent)
+ return 0, errExportTxsDisabled
}

// SemanticVerify this transaction is valid.
@@ -138,105 +69,12 @@ func (tx *UnsignedExportTx) SemanticVerify(
    baseFee *big.Int,
    rules params.Rules,
) error {
-     if err := tx.Verify(vm.ctx.XChainID, vm.ctx, rules); err != nil {
-         return err
-     }
-
-     // Check the transaction consumes and produces the right amounts
-     fc := avax.NewFlowChecker()
-     switch {
-     // Apply dynamic fees to export transactions as of Apricot Phase 3
-     case rules.IsApricotPhase3:
-         gasUsed, err := stx.GasUsed()
-         if err != nil {
-             return err
-         }
-         txFee, err := calculateDynamicFee(gasUsed, baseFee)
-         if err != nil {
-             return err
-         }
-         fc.Produce(vm.ctx.AVAXAssetID, txFee)
-
-     // Apply fees to export transactions before Apricot Phase 3
-     default:
-         fc.Produce(vm.ctx.AVAXAssetID, params.AvalancheAtomicTxFee)
-     }
-     for _, out := range tx.ExportedOutputs {
-         fc.Produce(out.AssetId(), out.Output().Amount())
-     }
-     for _, in := range tx.Ins {
-         fc.Consume(in.AssetId, in.Amount)
-     }
-
-     if err := fc.Verify(); err != nil {
-         return fmt.Errorf("export tx flow check failed due to: %w", err)
-     }
-
-     if len(tx.Ins) != len(stx.Creds) {
-         return fmt.Errorf("export tx contained mismatched number of inputs/credentials (%d vs. %d)", len(tx.Ins), len(stx.Creds))
-     }
-
-     for i, input := range tx.Ins {
-         cred, ok := stx.Creds[i].(*secp256k1fx.Credential)
-         if !ok {
-             return fmt.Errorf("expected *secp256k1fx.Credential but got %T", cred)
-         }
-         if err := cred.Verify(); err != nil {
-             return err
-         }
-
-         if len(cred.Sigs) != 1 {
-             return fmt.Errorf("expected one signature for EVM Input Credential, but found: %d", len(cred.Sigs))
-         }
-         pubKeyIntf, err := vm.secpFactory.RecoverPublicKey(tx.UnsignedBytes(), cred.Sigs[0][:])
-         if err != nil {
-             return err
-         }
-         pubKey, ok := pubKeyIntf.(*crypto.PublicKeySECP256K1R)
-         if !ok {
-             // This should never happen
-             return fmt.Errorf("expected *crypto.PublicKeySECP256K1R but got %T", pubKeyIntf)
-         }
-         if input.Address != PublicKeyToEthAddress(pubKey) {
-             return errPublicKeySignatureMismatch
-         }
-     }
-
-     return nil
+ return errExportTxsDisabled
}

// Accept this transaction.
-func (tx *UnsignedExportTx) Accept(ctx *snow.Context, batch database.Batch) error {
-     txID := tx.ID()
-
-     elems := make([]*atomic.Element, len(tx.ExportedOutputs))
-     for i, out := range tx.ExportedOutputs {
-         utxo := &avax.UTXO{
-             UTXOID: avax.UTXOID{
-                 TxID:      txID,
-                 OutputIndex: uint32(i),
-             },
-             Asset: avax.Asset{ID: out.AssetId()},
-             Out:    out.Out,
-         }
-
-         utxoBytes, err := Codec.Marshal(codecVersion, utxo)
-         if err != nil {
-             return err
-         }
-
-         utxoID := utxo.InputID()
-         elem := &atomic.Element{
-             Key:    utxoID[:],
-             Value: utxoBytes,
-         }
-
-         if out, ok := utxo.Out.(avax.Addressable); ok {
-             elem.Traits = out.Addresses()
-         }
-
-         elems[i] = elem
-     }
-
-     return ctx.SharedMemory.Apply(map[ids.ID]*atomic.Requests{tx.DestinationChain: {PutRequests: elems}}, batch)
+func (tx *UnsignedExportTx) Accept() (ids.ID, *atomic.Requests, error) {
+     return ids.ID{}, nil, errExportTxsDisabled
+ }

// newExportTx returns a new ExportTx
@@ -248,122 +86,10 @@ func (vm *VM) newExportTx(
    baseFee *big.Int, // fee to use post-AP3
    keys []*crypto.PrivateKeySECP256K1R, // Pay the fee and provide the tokens
) (*Tx, error) {
-     if vm.ctx.XChainID != chainID {

```

```

-         return nil, errWrongChainID
-     }
-
-     outs := []*avax.TransferableOutput{ // Exported to X-Chain
-         Asset: avax.Asset{ID: assetID},
-         Out: &secp256k1fx.TransferOutput{
-             Amt: amount,
-             OutputOwners: secp256k1fx.OutputOwners{
-                 Locktime: 0,
-                 Threshold: 1,
-                 Addrs: []ids.ShortID{to},
-             },
-         },
-     }
- }
-
- var (
-     avaxNeeded      uint64 = 0
-     ins, avaxIns     []EVMInput
-     signers, avaxSigners [][]*crypto.PrivateKeySECP256K1R
-     err              error
- )
-
- // consume non-AVAX
- if assetID != vm.ctx.AVAXAssetID {
-     ins, signers, err = vm.GetSpendableFunds(keys, assetID, amount)
-     if err != nil {
-         return nil, fmt.Errorf("couldn't generate tx inputs/signers: %w", err)
-     }
- } else {
-     avaxNeeded = amount
- }
-
- rules := vm.currentRules()
- switch {
- case rules.IsApricotPhase3:
-     utx := &UnsignedExportTx{
-         NetworkID:      vm.ctx.NetworkID,
-         BlockchainID:    vm.ctx.ChainID,
-         DestinationChain: chainID,
-         Ins:             ins,
-         ExportedOutputs: outs,
-     }
-
-     tx := &Tx{UnsignedAtomicTx: utx}
-     if err := tx.Sign(vm.codec, nil); err != nil {
-         return nil, err
-     }
-
-     var cost uint64
-     cost, err = tx.GasUsed()
-     if err != nil {
-         return nil, err
-     }
-
-     avaxIns, avaxSigners, err = vm.GetSpendableAVAXWithFee(keys, avaxNeeded, cost, baseFee)
-
- default:
-     var newAvaxNeeded uint64
-     newAvaxNeeded, err = math.Add64(avaxNeeded, params.AvalancheAtomicTxFee)
-     if err != nil {
-         return nil, errOverflowExport
-     }
-     avaxIns, avaxSigners, err = vm.GetSpendableFunds(keys, vm.ctx.AVAXAssetID, newAvaxNeeded)
- }
-
- if err != nil {
-     return nil, fmt.Errorf("couldn't generate tx inputs/signers: %w", err)
- }
-
- ins = append(ins, avaxIns...)
- signers = append(signers, avaxSigners...)
-
- avax.SortTransferableOutputs(outs, vm.codec)
- SortEVMInputsAndSigners(ins, signers)
-
- // Create the transaction
- utx := &UnsignedExportTx{
-     NetworkID:      vm.ctx.NetworkID,
-     BlockchainID:    vm.ctx.ChainID,
-     DestinationChain: chainID,
-     Ins:             ins,
-     ExportedOutputs: outs,
- }
-
- tx := &Tx{UnsignedAtomicTx: utx}
- if err := tx.Sign(vm.codec, signers); err != nil {
-     return nil, err
- }
-
- return tx, utx.Verify(vm.ctx.XChainID, vm.ctx, vm.currentRules())
+ return nil, errExportTxDisabled
- }
-
- // EVMStateTransfer executes the state update from the atomic export transaction
- func (tx *UnsignedExportTx) EVMStateTransfer(ctx *snow.Context, state *state.StateDB) error {
-     addrs := map[[20]byte]uint64{}
-     for _, from := range tx.Ins {
-         if from.AssetId == ctx.AVAXAssetID {
-             log.Debug("crosschain C->X", "addr", from.Address, "amount", from.Amount, "assetID", "AVAX")
-             // We multiply the input amount by x2cRate to convert AVAX back to the appropriate
-             // denomination before export.
-             amount := new(big.Int).Mul(
-                 new(big.Int).SetUint64(from.Amount), x2cRate)
-             if state.GetBalance(from.Address).Cmp(amount) < 0 {
-                 return errInsufficientFunds
-             }
-             state.SubBalance(from.Address, amount)
-         } else {
-             log.Debug("crosschain C->X", "addr", from.Address, "amount", from.Amount, "assetID", from.AssetId)
-             amount := new(big.Int).SetUint64(from.Amount)
-             if state.GetBalanceMultiCoin(from.Address, common.Hash(from.AssetId)).Cmp(amount) < 0 {
-                 return errInsufficientFunds
-             }
-             state.SubBalanceMultiCoin(from.Address, common.Hash(from.AssetId), amount)
-         }
-         if state.GetNonce(from.Address) != from.Nonce {
-             return errInvalidNonce
-         }
-         addrs[from.Address] = from.Nonce
-     }
-     for addr, nonce := range addrs {
-         state.SetNonce(addr, nonce+1)
-     }
-     return nil
+ return errExportTxDisabled
- }
-
- diff --git a/plugin/evm/export_tx_test.go b/plugin/evm/export_tx_test.go
- deleted file mode 100644
- index 2861af00..00000000
- --- a/plugin/evm/export_tx_test.go
- +++ /dev/null
- @@ -1,1685 +0,0 @@
- -// (c) 2019-2020, Ava Labs, Inc. All rights reserved.
- -// See the file LICENSE for licensing terms.
- -
- -package evm
- -
- -import (
- -    "bytes"
- -    "math/big"

```

```

-     "testing"
-
-     "github.com/ava-labs/avalanchego/chains/atomic"
-     "github.com/ava-labs/avalanchego/ids"
-     engCommon "github.com/ava-labs/avalanchego/snow/engine/common"
-     "github.com/ava-labs/avalanchego/utls/crypto"
-     "github.com/ava-labs/avalanchego/utls/units"
-     "github.com/ava-labs/avalanchego/vms/components/avax"
-     "github.com/ava-labs/avalanchego/vms/secp256k1fx"
-     "github.com/ava-labs/coreth/params"
-     "github.com/ethereum/go-ethereum/common"
- )
-
- // createExportTxOptions adds funds to shared memory, imports them, and returns a list of export transactions
- // that attempt to send the funds to each of the test keys (list of length 3).
- func createExportTxOptions(t *testing.T, vm *VM, issuer chan engCommon.Message, sharedMemory *atomic.Memory) []*Tx {
-     // Add a UTXO to shared memory
-     utxo := &avax.UTXO{
-         UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-         Asset:  avax.Asset{ID: vm.ctx.AVAXAssetID},
-         Out:    &secp256k1fx.TransferOutput{
-             Amt: uint64(500000000),
-             OutputOwners: secp256k1fx.OutputOwners{
-                 Threshold: 1,
-                 Addrs:     []ids.ShortID{testKeys[0].PublicKey().Address()},
-             },
-         },
-     },
-
-     }
-
-     utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
-     inputID := utxo.InputID()
-     if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
-         Key:   inputID[:],
-         Value: utxoBytes,
-         Traits: [][]byte{
-             testKeys[0].PublicKey().Address().Bytes(),
-         },
-     }}}}); err != nil {
-         t.Fatal(err)
-     }
-
-     // Import the funds
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     // Use the funds to create 3 conflicting export transactions sending the funds to each of the test addresses
-     exportTxs := make([]*Tx, 0, 3)
-     for _, addr := range testShortIDAddrs {
-         exportTx, err := vm.newExportTx(vm.ctx.AVAXAssetID, uint64(500000000), vm.ctx.XChainID, addr, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-         if err != nil {
-             t.Fatal(err)
-         }
-
-         exportTxs = append(exportTxs, exportTx)
-     }
-
-     return exportTxs
- }
-
- func TestExportTxEVMStateTransfer(t *testing.T) {
-     key := testKeys[0]
-     addr := key.PublicKey().Address()
-     ethAddr := GetEthAddress(key)
-
-     avaxAmount := 50 * units.MilliAvax
-     avaxUTXOID := avax.UTXOID{
-         OutputIndex: 0,
-     }
-     avaxInputID := avaxUTXOID.InputID()
-
-     customAmount := uint64(100)
-     customAssetID := ids.ID{1, 2, 3, 4, 5, 7}
-     customUTXOID := avax.UTXOID{
-         OutputIndex: 1,
-     }
-     customInputID := customUTXOID.InputID()
-
-     customUTXO := &avax.UTXO{
-         UTXOID: customUTXOID,
-         Asset:  avax.Asset{ID: customAssetID},
-         Out:    &secp256k1fx.TransferOutput{
-             Amt: customAmount,
-             OutputOwners: secp256k1fx.OutputOwners{
-                 Threshold: 1,
-                 Addrs:     []ids.ShortID{addr},
-             },
-         },
-     },
-
-     }
-
-     tests := []struct {
-         name      string
-         tx        []EVMInput
-         avaxBalance *big.Int
-         balances   map[ids.ID]*big.Int
-         expectedNonce uint64
-         shouldErr  bool
-     }{
-         {
-             name:      "no transfers",
-             tx:        nil,
-             avaxBalance: big.NewInt(int64(avaxAmount) * x2cRateInt64),
-             balances: map[ids.ID]*big.Int{
-                 customAssetID: big.NewInt(int64(customAmount)),
-             },
-             expectedNonce: 0,
-             shouldErr:     false,
-         },
-     }

```

```

    },
    {
        name: "spend half AVAX",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: avaxAmount / 2,
                AssetID: testAvaxAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(int64(avaxAmount/2) * x2cRateInt64),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(int64(customAmount)),
        },
        expectedNonce: 1,
        shouldErr: false,
    },
    {
        name: "spend all AVAX",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: avaxAmount,
                AssetID: testAvaxAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(0),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(int64(customAmount)),
        },
        expectedNonce: 1,
        shouldErr: false,
    },
    {
        name: "spend too much AVAX",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: avaxAmount + 1,
                AssetID: testAvaxAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(0),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(int64(customAmount)),
        },
        expectedNonce: 1,
        shouldErr: true,
    },
    {
        name: "spend half custom",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: customAmount / 2,
                AssetID: customAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(int64(avaxAmount) * x2cRateInt64),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(int64(customAmount / 2)),
        },
        expectedNonce: 1,
        shouldErr: false,
    },
    {
        name: "spend all custom",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: customAmount,
                AssetID: customAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(int64(avaxAmount) * x2cRateInt64),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(0),
        },
        expectedNonce: 1,
        shouldErr: false,
    },
    {
        name: "spend too much custom",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: customAmount + 1,
                AssetID: customAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(int64(avaxAmount) * x2cRateInt64),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(0),
        },
        expectedNonce: 1,
        shouldErr: true,
    },
    {
        name: "spend everything",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: customAmount,
                AssetID: customAssetID,
                Nonce: 0,
            },
            {
                Address: ethAddr,
                Amount: avaxAmount,
                AssetID: testAvaxAssetID,
                Nonce: 0,
            },
        },
        avaxBalance: big.NewInt(0),
        balances: map[ids.ID]*big.Int{
            customAssetID: big.NewInt(0),
        },
        expectedNonce: 1,
        shouldErr: false,
    },
    {
        name: "spend everything wrong nonce",
        tx: []EVMInput{
            {
                Address: ethAddr,
                Amount: customAmount,
                AssetID: customAssetID,
                Nonce: 1,
            },
        },
    },

```

```

    },
    {
        Address: ethAddr,
        Amount: avaxAmount,
        AssetID: testAvaxAssetID,
        Nonce: 1,
    },
},
avaxBalance: big.NewInt(0),
balances: map[ids.ID]*big.Int{
    customAssetID: big.NewInt(0),
},
expectedNonce: 1,
shouldErr: true,
},
{
    name: "spend everything changing nonces",
    tx: []EVMInput{
        {
            Address: ethAddr,
            Amount: customAmount,
            AssetID: customAssetID,
            Nonce: 0,
        },
        {
            Address: ethAddr,
            Amount: avaxAmount,
            AssetID: testAvaxAssetID,
            Nonce: 1,
        },
    },
    avaxBalance: big.NewInt(0),
    balances: map[ids.ID]*big.Int{
        customAssetID: big.NewInt(0),
    },
    expectedNonce: 1,
    shouldErr: true,
},
}
for _, test := range tests {
    t.Run(test.name, func(t *testing.T) {
        issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJSONApricotPhase0, "", "")
        defer func() {
            if err := vm.Shutdown(); err != nil {
                t.Fatal(err)
            }
        }()

        avaxUTXO := &avax.UTXO{
            UTXOID: avaxUTXOID,
            Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
            Out: &secp256k1fx.TransferOutput{
                Amt: avaxAmount,
                OutputOwners: secp256k1fx.OutputOwners{
                    Threshold: 1,
                    Addrs: []ids.ShortID{addr},
                },
            },
        },
    }

    avaxUTXOBytes, err := vm.codec.Marshal(codecVersion, avaxUTXO)
    if err != nil {
        t.Fatal(err)
    }

    customUTXOBytes, err := vm.codec.Marshal(codecVersion, customUTXO)
    if err != nil {
        t.Fatal(err)
    }

    xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
    if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []atomic.Element{
        {
            Key: avaxInputID[:],
            Value: avaxUTXOBytes,
            Traits: [][]byte{
                addr.Bytes(),
            },
        },
        {
            Key: customInputID[:],
            Value: customUTXOBytes,
            Traits: [][]byte{
                addr.Bytes(),
            },
        },
    }}}); err != nil {
        t.Fatal(err)
    }

    tx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
    if err != nil {
        t.Fatal(err)
    }

    if err := vm.issueTx(tx, true /*=local*/); err != nil {
        t.Fatal(err)
    }

    <-issuer

    blk, err := vm.BuildBlock()
    if err != nil {
        t.Fatal(err)
    }

    if err := blk.Verify(); err != nil {
        t.Fatal(err)
    }

    if err := vm.SetPreference(blk.ID()); err != nil {
        t.Fatal(err)
    }

    if err := blk.Accept(); err != nil {
        t.Fatal(err)
    }

    newTx := UnsignedExportTx{
        Ins: test.tx,
    }

    stateDB, err := vm.chain.CurrentState()
    if err != nil {
        t.Fatal(err)
    }

    err = newTx.EVMStateTransfer(vm.ctx, stateDB)
    if test.shouldErr {
        if err == nil {
            t.Fatal("expected EVMStateTransfer to fail")
        }
        return
    }
}

```

```

-         if err != nil {
-             t.Fatal(err)
-         }
-
-         avaxBalance := stateDB.GetBalance(ethAddr)
-         if avaxBalance.Cmp(test.avaxBalance) != 0 {
-             t.Fatalf("address balance %s equal %s not %s", addr.String(), avaxBalance, test.avaxBalance)
-         }
-
-         for assetID, expectedBalance := range test.balances {
-             balance := stateDB.GetBalanceMultiCoin(ethAddr, common.Hash(assetID))
-             if avaxBalance.Cmp(test.avaxBalance) != 0 {
-                 t.Fatalf("%s address balance %s equal %s not %s", assetID, addr.String(), balance, expectedBalance)
-             }
-         }
-
-         if stateDB.GetNonce(ethAddr) != test.expectedNonce {
-             t.Fatalf("failed to set nonce to %d", test.expectedNonce)
-         }
-     })
- }
- })
-}
-
- func TestExportTxSemanticVerify(t *testing.T) {
-     _, vm, _, _ := GenesisVM(t, true, genesisJ50NApricotPhase0, "", "")
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     parent := vm.LastAcceptedBlockInternal().(*Block)
-
-     key := testKeys[0]
-     addr := key.PublicKey().Address()
-     ethAddr := testEthAdrs[0]
-
-     var (
-         avaxBalance      = 10 * units.Avox
-         custom0Balance uint64 = 100
-         custom0AssetID   = ids.ID(1, 2, 3, 4, 5)
-         custom1Balance uint64 = 1000
-         custom1AssetID   = ids.ID(1, 2, 3, 4, 5, 6)
-     )
-
-     validExportTx := &UnsignedExportTx{
-         NetworkID:      vm.ctx.NetworkID,
-         BlockchainID:    vm.ctx.ChainID,
-         DestinationChain: vm.ctx.XChainID,
-         Ins: []EVMInput{
-             {
-                 Address: ethAddr,
-                 Amount:  avaxBalance,
-                 AssetID:  vm.ctx.AVAXAssetID,
-                 Nonce:    0,
-             },
-             {
-                 Address: ethAddr,
-                 Amount:  custom0Balance,
-                 AssetID:  custom0AssetID,
-                 Nonce:    0,
-             },
-             {
-                 Address: ethAddr,
-                 Amount:  custom1Balance,
-                 AssetID:  custom1AssetID,
-                 Nonce:    0,
-             },
-         },
-         ExportedOutputs: []*avax.TransferableOutput{
-             {
-                 Asset: avax.Asset{ID: custom0AssetID},
-                 Out: &secp256k1fx.TransferOutput{
-                     Amt: custom0Balance,
-                     OutputOwners: secp256k1fx.OutputOwners{
-                         Threshold: 1,
-                         Addrs:      []ids.ShortID{addr},
-                     },
-                 },
-             },
-         },
-     }
-
-     tests := []struct {
-         name      string
-         tx         *Tx
-         signers    [][]*crypto.PrivateKeySECP256K1R
-         baseFee    *big.Int
-         rules      params.Rules
-         shouldErr bool
-     }{
-         {
-             name: "valid",
-             tx:   &Tx{UnsignedAtomicTx: validExportTx},
-             signers: [][]*crypto.PrivateKeySECP256K1R{
-                 {key},
-                 {key},
-                 {key},
-             },
-             baseFee:  initialBaseFee,
-             rules:    apricotRulesPhase3,
-             shouldErr: false,
-         },
-         {
-             name: "wrong destination",
-             tx: func() *Tx {
-                 validExportTx := *validExportTx
-                 validExportTx.DestinationChain = ids.GenerateTestID()
-                 return &Tx{UnsignedAtomicTx: &validExportTx}
-             }(),
-             signers: [][]*crypto.PrivateKeySECP256K1R{
-                 {key},
-                 {key},
-                 {key},
-             },
-             baseFee:  initialBaseFee,
-             rules:    apricotRulesPhase3,
-             shouldErr: true,
-         },
-         {
-             name: "no outputs",
-             tx: func() *Tx {
-                 validExportTx := *validExportTx
-                 validExportTx.ExportedOutputs = nil
-                 return &Tx{UnsignedAtomicTx: &validExportTx}
-             }(),
-             signers: [][]*crypto.PrivateKeySECP256K1R{
-                 {key},
-                 {key},
-                 {key},
-             },
-             baseFee:  initialBaseFee,
-             rules:    apricotRulesPhase3,
-         },
-     }

```

```

    },
    {
      name: "wrong networkID",
      tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.NetworkID++
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
    {
      name: "wrong chainID",
      tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.BlockchainID = ids.GenerateTestID()
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
    {
      name: "invalid input",
      tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.Ins = append([]EVMInput{}, validExportTx.Ins...)
        validExportTx.Ins[2].Amount = 0
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
    {
      name: "invalid output",
      tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.ExportedOutputs = []*avax.TransferableOutput{
          {
            Asset: avax.Asset{ID: custom0AssetID},
            Out: &secp256k1fx.TransferOutput{
              Amt: custom0Balance,
              OutputOwners: secp256k1fx.OutputOwners{
                Threshold: 0,
                Addrs:      []ids.ShortID{addr},
              },
            },
          },
        }
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
    {
      name: "unsorted outputs",
      tx: func() *Tx {
        validExportTx := *validExportTx
        exportOutputs := []*avax.TransferableOutput{
          {
            Asset: avax.Asset{ID: custom0AssetID},
            Out: &secp256k1fx.TransferOutput{
              Amt: custom0Balance/2 + 1,
              OutputOwners: secp256k1fx.OutputOwners{
                Threshold: 1,
                Addrs:      []ids.ShortID{addr},
              },
            },
          },
          {
            Asset: avax.Asset{ID: custom0AssetID},
            Out: &secp256k1fx.TransferOutput{
              Amt: custom0Balance/2 - 1,
              OutputOwners: secp256k1fx.OutputOwners{
                Threshold: 1,
                Addrs:      []ids.ShortID{addr},
              },
            },
          },
        }
        // Sort the outputs and then swap the ordering to ensure that they are ordered incorrectly
        avax.SortTransferableOutputs(exportOutputs, Codec)
        exportOutputs[0], exportOutputs[1] = exportOutputs[1], exportOutputs[0]
        validExportTx.ExportedOutputs = exportOutputs
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
    {
      name: "not unique inputs",
      tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.Ins = append([]EVMInput{}, validExportTx.Ins...)
        validExportTx.Ins[2] = validExportTx.Ins[1]
        return &Tx{UnsignedAtomicTx: &validExportTx}
      }(),
      signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
      },
      baseFee: initialBaseFee,
      rules:   apricotRulesPhase3,
      shouldErr: true,
    },
  },
  {
    name: "not unique inputs",
    tx: func() *Tx {
      validExportTx := *validExportTx
      validExportTx.Ins = append([]EVMInput{}, validExportTx.Ins...)
      validExportTx.Ins[2] = validExportTx.Ins[1]
      return &Tx{UnsignedAtomicTx: &validExportTx}
    }(),
    signers: [][]*crypto.PrivateKeySECP256K1R{
      {key},
      {key},
      {key},
    },
    baseFee: initialBaseFee,
    rules:   apricotRulesPhase3,
    shouldErr: true,
  },
}

```

```

{
    name: "custom asset insufficient funds",
    tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.ExportedOutputs = []*avax.TransferableOutput{
            {
                Asset: avax.Asset{ID: custom0AssetID},
                Out: &secp256k1fx.TransferOutput{
                    Amt: custom0Balance + 1,
                    OutputOwners: secp256k1fx.OutputOwners{
                        Threshold: 1,
                        Addrs: []ids.ShortID{addr},
                    },
                },
            },
        },
        return &Tx{UnsignedAtomicTx: &validExportTx}
    }(),
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "avax insufficient funds",
    tx: func() *Tx {
        validExportTx := *validExportTx
        validExportTx.ExportedOutputs = []*avax.TransferableOutput{
            {
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                Out: &secp256k1fx.TransferOutput{
                    Amt: avaxBalance, // after fees this should be too much
                    OutputOwners: secp256k1fx.OutputOwners{
                        Threshold: 1,
                        Addrs: []ids.ShortID{addr},
                    },
                },
            },
        },
        return &Tx{UnsignedAtomicTx: &validExportTx}
    }(),
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "too many signatures",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "too few signatures",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "too many signatures on credential",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {key, testKeys[1]},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "too few signatures on credential",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "wrong signature on credential",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {testKeys[1]},
        {key},
        {key},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
{
    name: "no signatures",
    tx: &Tx{UnsignedAtomicTx: validExportTx},
    signers: [][]*crypto.PrivateKeySECP256K1R{
        {},
    },
    baseFee: initialBaseFee,
    rules: apricotRulesPhase3,
    shouldErr: true,
},
}

for _, test := range tests {
    if err := test.tx.Sign(vm.codec, test.signers); err != nil {
        t.Fatal(err)
    }

    t.Run(test.name, func(t *testing.T) {
        tx := test.tx
        exportTx := tx.UnsignedAtomicTx

        err := exportTx.SemanticVerify(vm, tx, parent, test.baseFee, test.rules)
        if test.shouldErr && err == nil {

```



```

-         t.Fatalf("should have errored but returned valid")
-     }
-     if !test.shouldErr && err != nil {
-         t.Fatalf("shouldn't have errored but returned %s", err)
-     }
- })
-}
-
-func TestExportTxAccept(t *testing.T) {
-    _, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJSONApricotPhase0, "", "")
-
-    xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
-
-    defer func() {
-        if err := vm.Shutdown(); err != nil {
-            t.Fatal(err)
-        }
-    }()
-
-    key := testKeys[0]
-    addr := key.PublicKey().Address()
-    ethAddr := testEthAddrs[0]
-
-    var (
-        avaxBalance      = 10 * units.Avox
-        custom0Balance uint64 = 100
-        custom0AssetID    = ids.ID{1, 2, 3, 4, 5}
-    )
-
-    exportTx := &UnsignedExportTx{
-        NetworkID:      vm.ctx.NetworkID,
-        BlockchainID:    vm.ctx.ChainID,
-        DestinationChain: vm.ctx.XChainID,
-        Ins: []EVMInput{
-            {
-                Address: ethAddr,
-                Amount: avaxBalance,
-                AssetID: vm.ctx.AVAXAssetID,
-                Nonce:    0,
-            },
-            {
-                Address: ethAddr,
-                Amount: custom0Balance,
-                AssetID: custom0AssetID,
-                Nonce:    0,
-            },
-        },
-        ExportedOutputs: []*avax.TransferableOutput{
-            {
-                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
-                Out: &secp256k1fx.TransferOutput{
-                    Amt: avaxBalance,
-                    OutputOwners: secp256k1fx.OutputOwners{
-                        Threshold: 1,
-                        Addrs:      []ids.ShortID{addr},
-                    },
-                },
-            },
-            {
-                Asset: avax.Asset{ID: custom0AssetID},
-                Out: &secp256k1fx.TransferOutput{
-                    Amt: custom0Balance,
-                    OutputOwners: secp256k1fx.OutputOwners{
-                        Threshold: 1,
-                        Addrs:      []ids.ShortID{addr},
-                    },
-                },
-            },
-        },
-    }
-
-    tx := &Tx{UnsignedAtomicTx: exportTx}
-
-    signers := [][]*crypto.PrivateKeySECP256K1R{
-        {key},
-        {key},
-        {key},
-    }
-
-    if err := tx.Sign(vm.codec, signers); err != nil {
-        t.Fatal(err)
-    }
-
-    commitBatch, err := vm.db.CommitBatch()
-    if err != nil {
-        t.Fatalf("Failed to create commit batch for VM due to %s", err)
-    }
-    if err := tx.Accept(vm.ctx, commitBatch); err != nil {
-        t.Fatalf("Failed to accept export transaction due to: %s", err)
-    }
-
-    indexedValues, _, _, err := xChainSharedMemory.Indexed(vm.ctx.ChainID, [][]byte{addr.Bytes()}, nil, nil, 3)
-    if err != nil {
-        t.Fatal(err)
-    }
-
-    if len(indexedValues) != 2 {
-        t.Fatalf("expected 2 values but got %d", len(indexedValues))
-    }
-
-    avaxUTXOID := avax.UTXOID{
-        TxID:      tx.ID(),
-        OutputIndex: 0,
-    }
-    avaxInputID := avaxUTXOID.InputID()
-
-    customUTXOID := avax.UTXOID{
-        TxID:      tx.ID(),
-        OutputIndex: 1,
-    }
-    customInputID := customUTXOID.InputID()
-
-    fetchedValues, err := xChainSharedMemory.Get(vm.ctx.ChainID, [][]byte{
-        customInputID[:],
-        avaxInputID[:],
-    })
-    if err != nil {
-        t.Fatal(err)
-    }
-
-    if !bytes.Equal(fetchedValues[0], indexedValues[0]) {
-        t.Fatalf("inconsistent values returned fetched %x indexed %x", fetchedValues[0], indexedValues[0])
-    }
-    if !bytes.Equal(fetchedValues[1], indexedValues[1]) {
-        t.Fatalf("inconsistent values returned fetched %x indexed %x", fetchedValues[1], indexedValues[1])
-    }
-
-    customUTXOBytes, err := Codec.Marshal(codecVersion, &avax.UTXO{
-        UTXOID: customUTXOID,
-        Asset:  avax.Asset{ID: custom0AssetID},
-        Out:    exportTx.ExportedOutputs[1].Out,
-    })
-    if err != nil {

```

```

-         t.Fatal(err)
-     }
-
-    avaxUTXOBytes, err := Codec.Marshal(codecVersion, &avax.UTXO{
-         UTXOID: avax.UTXOID,
-         Asset:   avax.Asset{ID: vm.ctx.AVAXAssetID},
-         Out:     exportTx.ExportedOutputs[0].Out,
-     })
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if !bytes.Equal(fetchedValues[0], customUTXOBytes) {
-         t.Fatalf("incorrect values returned expected %x got %x", customUTXOBytes, fetchedValues[0])
-     }
-
-     if !bytes.Equal(fetchedValues[1], avaxUTXOBytes) {
-         t.Fatalf("incorrect values returned expected %x got %x", avaxUTXOBytes, fetchedValues[1])
-     }
- }
-
-}
-
-// TestExportTxVerifyNil tests that Verify returns nil for a nil transaction
-// nolint:funlen
-func TestExportTxVerifyNil(t *testing.T) {
-    var exportTx *UnsignedExportTx
-    if err := exportTx.Verify(testXChainID, NewContext(), apricotRulesPhase0); err == nil {
-        t.Fatal("Verify should have failed due to nil transaction")
-    }
-}
-
-// TestExportTxVerify tests that Verify returns an error for a non-nil transaction
-// nolint:funlen
-func TestExportTxVerify(t *testing.T) {
-    var exportAmount uint64 = 10000000
-    exportTx := &UnsignedExportTx{
-        NetworkID:      testNetworkID,
-        BlockchainID:    testCChainID,
-        DestinationChain: testXChainID,
-        Ins: []EVMInput{
-            {
-                Address: testEthAddrs[0],
-                Amount: exportAmount,
-                AssetID: testAvaxAssetID,
-                Nonce:   0,
-            },
-            {
-                Address: testEthAddrs[2],
-                Amount: exportAmount,
-                AssetID: testAvaxAssetID,
-                Nonce:   0,
-            },
-        },
-        ExportedOutputs: []*avax.TransferableOutput{
-            {
-                Asset: avax.Asset{ID: testAvaxAssetID},
-                Out: &secp256k1fx.TransferOutput{
-                    Amt: exportAmount,
-                    OutputOwners: secp256k1fx.OutputOwners{
-                        Locktime: 0,
-                        Threshold: 1,
-                        Addrs:   []ids.ShortID{testShortIDAddrs[0]},
-                    },
-                },
-            },
-            {
-                Asset: avax.Asset{ID: testAvaxAssetID},
-                Out: &secp256k1fx.TransferOutput{
-                    Amt: exportAmount,
-                    OutputOwners: secp256k1fx.OutputOwners{
-                        Locktime: 0,
-                        Threshold: 1,
-                        Addrs:   []ids.ShortID{testShortIDAddrs[1]},
-                    },
-                },
-            },
-        },
-    }
-
-    // Sort the inputs and outputs to ensure the transaction is canonical
-    avax.SortTransferableOutputs(exportTx.ExportedOutputs, Codec)
-    // Pass in a list of signers here with the appropriate length
-    // to avoid causing a nil-pointer error in the helper method
-    emptySigners := make([][]*crypto.PrivateKeySECP256K1R, 2)
-    SortEVMInputsAndSigners(exportTx.Ins, emptySigners)
-
-    ctx := NewContext()
-
-    // Test Valid Export Tx
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err != nil {
-        t.Fatalf("Failed to verify valid ExportTx: %s", err)
-    }
-
-    exportTx.NetworkID = testNetworkID + 1
-
-    // Test Incorrect Network ID Errors
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to incorrect network ID")
-    }
-
-    exportTx.NetworkID = testNetworkID
-    exportTx.BlockchainID = nonExistentID
-    // Test Incorrect Blockchain ID Errors
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to incorrect blockchain ID")
-    }
-
-    exportTx.BlockchainID = testCChainID
-    exportTx.DestinationChain = nonExistentID
-    // Test Incorrect Destination Chain ID Errors
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to incorrect destination chain")
-    }
-
-    exportTx.DestinationChain = testXChainID
-    exportedOuts := exportTx.ExportedOutputs
-    exportTx.ExportedOutputs = nil
-    evmInputs := exportTx.Ins
-    // Test No Exported Outputs Errors
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to no exported outputs")
-    }
-
-    exportTx.ExportedOutputs = []*avax.TransferableOutput{exportedOuts[1], exportedOuts[0]}
-    // Test Unsorted outputs Errors
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to no unsorted exported outputs")
-    }
-
-    exportTx.ExportedOutputs = []*avax.TransferableOutput{exportedOuts[0], nil}
-    // Test invalid exported output
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-        t.Fatal("ExportTx should have failed verification due to invalid output")
-    }
-
-    exportTx.ExportedOutputs = []*avax.TransferableOutput{exportedOuts[0], exportedOuts[1]}
-    exportTx.Ins = []EVMInput{evmInputs[1], evmInputs[0]}
-    // Test unsorted EVM Inputs passes before AP1
-    if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase0); err != nil {
-        t.Fatalf("ExportTx should have passed verification before AP1, but failed due to %s", err)
-    }
-}

```

```

-     }
-     // Test unsorted EVM Inputs fails after AP1
-     if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-         t.Fatal("ExportTx should have failed verification due to unsorted EVM Inputs")
-     }
-     exportTx.Ins = []EVMInput{
-         {
-             Address: testEthAddrs[0],
-             Amount:  0,
-             AssetID: testAvaxAssetID,
-             Nonce:   0,
-         },
-     }
-     // Test ExportTx with invalid EVM Input amount 0 fails verification
-     if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-         t.Fatal("ExportTx should have failed verification due to 0 value amount")
-     }
-     exportTx.Ins = []EVMInput{evmInputs[0], evmInputs[0]}
-     // Test non-unique EVM Inputs passes verification before AP1
-     if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase0); err != nil {
-         t.Fatalf("ExportTx with non-unique EVM Inputs should have passed verification prior to AP1, but failed due to %s", err)
-     }
-     exportTx.Ins = []EVMInput{evmInputs[0], evmInputs[0]}
-     // Test non-unique EVM Inputs fails verification after AP1
-     if err := exportTx.Verify(testXChainID, ctx, apricotRulesPhase1); err == nil {
-         t.Fatal("ExportTx should have failed verification due to non-unique inputs")
-     }
- }
-
- // Note: this is a brittle test to ensure that the gas cost of a transaction does
- // not change
- func TestExportTxGasCost(t *testing.T) {
-     avaxAssetID := ids.GenerateTestID()
-     chainID := ids.GenerateTestID()
-     xChainID := ids.GenerateTestID()
-     networkID := uint32(5)
-     exportAmount := uint64(5000000)
-
-     tests := map[string]struct {
-         UnsignedExportTx *UnsignedExportTx
-         Keys              [][]*crypto.PrivateKeySECP256K1R
-
-         BaseFee      *big.Int
-         ExpectedGasUsed uint64
-         ExpectedFee   uint64
-     }{
-         "simple export 1wei BaseFee": {
-             UnsignedExportTx: &UnsignedExportTx{
-                 NetworkID:    networkID,
-                 BlockchainID: chainID,
-                 DestinationChain: xChainID,
-                 Ins: []EVMInput{
-                     {
-                         Address: testEthAddrs[0],
-                         Amount: exportAmount,
-                         AssetID: avaxAssetID,
-                         Nonce:   0,
-                     },
-                 },
-                 ExportedOutputs: []*avax.TransferableOutput{
-                     {
-                         Asset: avax.Asset{ID: avaxAssetID},
-                         Out: &secp256k1fx.TransferOutput{
-                             Amt: exportAmount,
-                             OutputOwners: secp256k1fx.OutputOwners{
-                                 Locktime: 0,
-                                 Threshold: 1,
-                                 Addrs: []ids.ShortID{testShortIDAddrs[0]},
-                             },
-                         },
-                     },
-                 },
-             },
-             Keys:      [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}},
-             ExpectedGasUsed: 1230,
-             ExpectedFee: 1,
-             BaseFee: big.NewInt(1),
-         },
-         "simple export 25Gwei BaseFee": {
-             UnsignedExportTx: &UnsignedExportTx{
-                 NetworkID:    networkID,
-                 BlockchainID: chainID,
-                 DestinationChain: xChainID,
-                 Ins: []EVMInput{
-                     {
-                         Address: testEthAddrs[0],
-                         Amount: exportAmount,
-                         AssetID: avaxAssetID,
-                         Nonce:   0,
-                     },
-                 },
-                 ExportedOutputs: []*avax.TransferableOutput{
-                     {
-                         Asset: avax.Asset{ID: avaxAssetID},
-                         Out: &secp256k1fx.TransferOutput{
-                             Amt: exportAmount,
-                             OutputOwners: secp256k1fx.OutputOwners{
-                                 Locktime: 0,
-                                 Threshold: 1,
-                                 Addrs: []ids.ShortID{testShortIDAddrs[0]},
-                             },
-                         },
-                     },
-                 },
-             },
-             Keys:      [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}},
-             ExpectedGasUsed: 1230,
-             ExpectedFee: 30750,
-             BaseFee: big.NewInt(25 * params.GWei),
-         },
-         "simple export 225Gwei BaseFee": {
-             UnsignedExportTx: &UnsignedExportTx{
-                 NetworkID:    networkID,
-                 BlockchainID: chainID,
-                 DestinationChain: xChainID,
-                 Ins: []EVMInput{
-                     {
-                         Address: testEthAddrs[0],
-                         Amount: exportAmount,
-                         AssetID: avaxAssetID,
-                         Nonce:   0,
-                     },
-                 },
-                 ExportedOutputs: []*avax.TransferableOutput{
-                     {
-                         Asset: avax.Asset{ID: avaxAssetID},
-                         Out: &secp256k1fx.TransferOutput{
-                             Amt: exportAmount,
-                             OutputOwners: secp256k1fx.OutputOwners{
-                                 Locktime: 0,
-                                 Threshold: 1,
-                                 Addrs: []ids.ShortID{testShortIDAddrs[0]},
-                             },
-                         },
-                     },
-                 },
-             },
-         },
-     }
- }

```

```

    },
    Keys:      [][]*crypto.PrivateKeySECP256K1R({testKeys[0]}),
    ExpectedGasUsed: 1230,
    ExpectedFee: 276750,
    BaseFee:     big.NewInt(225 * params.GWei),
},
"complex export 25Gwei BaseFee": {
    UnsignedExportTx: &UnsignedExportTx{
        NetworkID:     networkID,
        BlockchainID:   chainID,
        DestinationChain: xChainID,
        Ins: []EVMInput{
            {
                Address: testEthAdrrs[0],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
            {
                Address: testEthAdrrs[1],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
            {
                Address: testEthAdrrs[2],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
        },
        ExportedOutputs: []*avax.TransferableOutput{
            {
                Asset: avax.Asset{ID: avaxAssetID},
                Out: &secp256k1fx.TransferOutput{
                    Amt: exportAmount * 3,
                    OutputOwners: secp256k1fx.OutputOwners{
                        Locktime: 0,
                        Threshold: 1,
                        Addrs: []ids.ShortID{testShortIDAdrrs[0]},
                    },
                },
            },
        },
    },
    Keys:      [][]*crypto.PrivateKeySECP256K1R({testKeys[0], testKeys[0], testKeys[0]}),
    ExpectedGasUsed: 3366,
    ExpectedFee: 84150,
    BaseFee:     big.NewInt(25 * params.GWei),
},
"complex export 225Gwei BaseFee": {
    UnsignedExportTx: &UnsignedExportTx{
        NetworkID:     networkID,
        BlockchainID:   chainID,
        DestinationChain: xChainID,
        Ins: []EVMInput{
            {
                Address: testEthAdrrs[0],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
            {
                Address: testEthAdrrs[1],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
            {
                Address: testEthAdrrs[2],
                Amount:  exportAmount,
                AssetID: avaxAssetID,
                Nonce:   0,
            },
        },
        ExportedOutputs: []*avax.TransferableOutput{
            {
                Asset: avax.Asset{ID: avaxAssetID},
                Out: &secp256k1fx.TransferOutput{
                    Amt: exportAmount * 3,
                    OutputOwners: secp256k1fx.OutputOwners{
                        Locktime: 0,
                        Threshold: 1,
                        Addrs: []ids.ShortID{testShortIDAdrrs[0]},
                    },
                },
            },
        },
    },
    Keys:      [][]*crypto.PrivateKeySECP256K1R({testKeys[0], testKeys[0], testKeys[0]}),
    ExpectedGasUsed: 3366,
    ExpectedFee: 757350,
    BaseFee:     big.NewInt(225 * params.GWei),
},
}

for name, test := range tests {
    t.Run(name, func(t *testing.T) {
        tx := &Tx{UnsignedAtomicTx: test.UnsignedExportTx}

        // Sign with the correct key
        if err := tx.Sign(Codec, test.Keys); err != nil {
            t.Fatal(err)
        }

        gasUsed, err := tx.GasUsed()
        if err != nil {
            t.Fatal(err)
        }
        if gasUsed != test.ExpectedGasUsed {
            t.Fatalf("Expected gasUsed to be %d, but found %d", test.ExpectedGasUsed, gasUsed)
        }

        fee, err := calculateDynamicFee(gasUsed, test.BaseFee)
        if err != nil {
            t.Fatal(err)
        }
        if fee != test.ExpectedFee {
            t.Fatalf("Expected fee to be %d, but found %d", test.ExpectedFee, fee)
        }
    })
}

-func TestNewExportTx(t *testing.T) {
    tests := []struct {
        name      string
        genesis   string
        rules     params.Rules
        bal       uint64
        expectedBurnedAVAX uint64
    }{

```

```

    {
        name:            "apricot phase 0",
        genesis:          genesisJSONApricotPhase0,
        rules:            apricotRulesPhase0,
        bal:              44000000,
        expectedBurnedAVAX: 1000000,
    },
    {
        name:            "apricot phase 1",
        genesis:          genesisJSONApricotPhase1,
        rules:            apricotRulesPhase1,
        bal:              44000000,
        expectedBurnedAVAX: 1000000,
    },
    {
        name:            "apricot phase 2",
        genesis:          genesisJSONApricotPhase2,
        rules:            apricotRulesPhase2,
        bal:              43000000,
        expectedBurnedAVAX: 1000000,
    },
    {
        name:            "apricot phase 3",
        genesis:          genesisJSONApricotPhase3,
        rules:            apricotRulesPhase3,
        bal:              44446500,
        expectedBurnedAVAX: 276750,
    },
    {
        name:            "apricot phase 4",
        genesis:          genesisJSONApricotPhase4,
        rules:            apricotRulesPhase4,
        bal:              44446500,
        expectedBurnedAVAX: 276750,
    },
}

for _, test := range tests {
    t.Run(test.name, func(t *testing.T) {
        issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, test.genesis, "", "")

        defer func() {
            if err := vm.Shutdown(); err != nil {
                t.Fatal(err)
            }
        }()

        parent := vm.LastAcceptedBlockInternal().(*Block)
        importAmount := uint64(50000000)
        utxoID := avax.UTXOID{TxID: ids.GenerateTestID()}

        utxo := &avax.UTXO{
            UTXOID: utxoID,
            Asset:  avax.Asset{ID: vm.ctx.AVAXAssetID},
            Out:    &secp256k1fx.TransferOutput{
                Amt: importAmount,
                OutputOwners: secp256k1fx.OutputOwners{
                    Threshold: 1,
                    Addrs:      []ids.ShortID{testKeys[0].PublicKey().Address()},
                },
            },
        },
        utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
        if err != nil {
            t.Fatal(err)
        }

        xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
        inputID := utxo.InputID()
        if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
            Key:  inputID[:],
            Value: utxoBytes,
            Traits: [][]byte{
                testKeys[0].PublicKey().Address().Bytes(),
            },
        }}}}); err != nil {
            t.Fatal(err)
        }

        tx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
        if err != nil {
            t.Fatal(err)
        }

        if err := vm.issueTx(tx, true /*=local*/); err != nil {
            t.Fatal(err)
        }

        <-issuer

        blk, err := vm.BuildBlock()
        if err != nil {
            t.Fatal(err)
        }

        if err := blk.Verify(); err != nil {
            t.Fatal(err)
        }

        if err := vm.SetPreference(blk.ID()); err != nil {
            t.Fatal(err)
        }

        if err := blk.Accept(); err != nil {
            t.Fatal(err)
        }

        parent = vm.LastAcceptedBlockInternal().(*Block)
        exportAmount := uint64(50000000)

        tx, err = vm.newExportTx(vm.ctx.AVAXAssetID, exportAmount, vm.ctx.XChainID, testShortIDAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
        if err != nil {
            t.Fatal(err)
        }

        exportTx := tx.UnsignedAtomicTx

        if err := exportTx.SemanticVerify(vm, tx, parent, parent.ethBlock.BaseFee(), test.rules); err != nil {
            t.Fatal("newExportTx created an invalid transaction", err)
        }

        burnedAVAX, err := exportTx.Burned(vm.ctx.AVAXAssetID)
        if err != nil {
            t.Fatal(err)
        }
        if burnedAVAX != test.expectedBurnedAVAX {
            t.Fatalf("burned wrong amount of AVAX - expected %d burned %d", test.expectedBurnedAVAX, burnedAVAX)
        }

        commitBatch, err := vm.db.CommitBatch()
        if err != nil {
            t.Fatalf("Failed to create commit batch for VM due to %s", err)
        }
        if err := exportTx.Accept(vm.ctx, commitBatch); err != nil {
            t.Fatalf("Failed to accept export transaction due to: %s", err)
        }
    })
}

```

```

    }

    sdb, err := vm.chain.CurrentState()
    if err != nil {
        t.Fatal(err)
    }
    err = exportTx.EVMStateTransfer(vm.ctx, sdb)
    if err != nil {
        t.Fatal(err)
    }

    addr := GetEthAddress(testKeys[0])
    if sdb.GetBalance(addr).Cmp(new(big.Int).SetUint64(test.bal*units.Avax)) != 0 {
        t.Fatalf("address balance %s equal %s not %s", addr.String(), sdb.GetBalance(addr), new(big.Int).SetUint64(test.bal*units.Avax))
    }
}

})

-func TestNewExportTxMulticoin(t *testing.T) {
    tests := []struct {
        name      string
        genesis    string
        rules      params.Rules
        bal        uint64
        balmc      uint64
    }{
        {
            name:      "apricot phase 0",
            genesis:  genesisJSONApricotPhase0,
            rules:    apricotRulesPhase0,
            bal:      490000000,
            balmc:    250000000,
        },
        {
            name:      "apricot phase 1",
            genesis:  genesisJSONApricotPhase1,
            rules:    apricotRulesPhase1,
            bal:      490000000,
            balmc:    250000000,
        },
        {
            name:      "apricot phase 2",
            genesis:  genesisJSONApricotPhase2,
            rules:    apricotRulesPhase2,
            bal:      480000000,
            balmc:    250000000,
        },
        {
            name:      "apricot phase 3",
            genesis:  genesisJSONApricotPhase3,
            rules:    apricotRulesPhase3,
            bal:      48947900,
            balmc:    250000000,
        },
    }
    for _, test := range tests {
        t.Run(test.name, func(t *testing.T) {
            issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, test.genesis, "", "")

            defer func() {
                if err := vm.Shutdown(); err != nil {
                    t.Fatal(err)
                }
            }()

            parent := vm.LastAcceptedBlockInternal().(*Block)
            importAmount := uint64(500000000)
            utxoID := avax.UTXOID{TxID: ids.GenerateTestID()}

            utxo := &avax.UTXO{
                UTXOID: utxoID,
                Asset:   avax.Asset{ID: vm.ctx.AVAXAssetID},
                Out: &secp256k1fx.TransferOutput{
                    Amt: importAmount,
                    OutputOwners: secp256k1fx.OutputOwners{
                        Threshold: 1,
                        Addrs:     []ids.ShortID{testKeys[0].PublicKey().Address()},
                    },
                },
            }

            utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
            if err != nil {
                t.Fatal(err)
            }

            inputID := utxo.InputID()

            tid := ids.GenerateTestID()
            importAmount2 := uint64(300000000)
            utxoID2 := avax.UTXOID{TxID: ids.GenerateTestID()}
            utxo2 := &avax.UTXO{
                UTXOID: utxoID2,
                Asset:   avax.Asset{ID: tid},
                Out: &secp256k1fx.TransferOutput{
                    Amt: importAmount2,
                    OutputOwners: secp256k1fx.OutputOwners{
                        Threshold: 1,
                        Addrs:     []ids.ShortID{testKeys[0].PublicKey().Address()},
                    },
                },
            }

            utxoBytes2, err := vm.codec.Marshal(codecVersion, utxo2)
            if err != nil {
                t.Fatal(err)
            }

            xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
            inputID2 := utxo2.InputID()
            if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{
                {
                    Key:   inputID[:],
                    Value: utxoBytes,
                    Traits: [][]byte{
                        testKeys[0].PublicKey().Address().Bytes(),
                    },
                },
                {
                    Key:   inputID2[:],
                    Value: utxoBytes2,
                    Traits: [][]byte{
                        testKeys[0].PublicKey().Address().Bytes(),
                    },
                },
            }}}); err != nil {
                t.Fatal(err)
            }

            tx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
            if err != nil {
                t.Fatal(err)
            }

            if err := vm.issueTx(tx, false); err != nil {

```

```

-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     parent = vm.LastAcceptedBlockInternal().(*Block)
-     exportAmount := uint64(5000000)
-
-     testKeys0Addr := GetEthAddress(testKeys[0])
-     exportId, err := ids.ToShortID(testKeys0Addr[:])
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     tx, err = vm.NewExportTx(tid, exportAmount, vm.ctx.XChainID, exportId, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     exportTx := tx.UnsignedAtomicTx
-
-     if err := exportTx.SemanticVerify(vm, tx, parent, parent.ethBlock.BaseFee(), test.rules); err != nil {
-         t.Fatal("newExportTx created an invalid transaction", err)
-     }
-
-     commitBatch, err := vm.db.CommitBatch()
-     if err != nil {
-         t.Fatalf("Failed to create commit batch for VM due to %s", err)
-     }
-     if err := exportTx.Accept(vm.ctx, commitBatch); err != nil {
-         t.Fatalf("Failed to accept export transaction due to: %s", err)
-     }
-
-     stdb, err := vm.chain.CurrentState()
-     if err != nil {
-         t.Fatal(err)
-     }
-     err = exportTx.EVMStateTransfer(vm.ctx, stdb)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     addr := GetEthAddress(testKeys[0])
-     if stdb.GetBalance(addr).Cmp(new(big.Int).SetUint64(test.bal*units.Avax)) != 0 {
-         t.Fatalf("address balance %s equal %s not %s", addr.String(), stdb.GetBalance(addr), new(big.Int).SetUint64(test.bal*units.Avax))
-     }
-     if stdb.GetBalanceMultiCoin(addr, common.BytesToHash(tid[:])).Cmp(new(big.Int).SetUint64(test.balmc)) != 0 {
-         t.Fatalf("address balance multicoin %s equal %s not %s", addr.String(), stdb.GetBalanceMultiCoin(addr, common.BytesToHash(tid[:])), new(big.Int).SetUint64(test.balmc))
-     }
- })
- }
-}
diff --git a/plugin/evm/ext_data_hashes.go b/plugin/evm/ext_data_hashes.go
index 7648c3bc..1c231aaa 100644
--- a/plugin/evm/ext_data_hashes.go
+++ b/plugin/evm/ext_data_hashes.go
@@ -8,22 +8,22 @@ import (
 )

 var (
- //go:embed fuji_ext_data_hashes.json
- rawFujiExtDataHashes []byte
- fujiExtDataHashes    map[common.Hash]common.Hash
+ //go:embed songbird_ext_data_hashes.json
+ rawSongbirdExtDataHashes []byte
+ songbirdExtDataHashes    map[common.Hash]common.Hash

- //go:embed mainnet_ext_data_hashes.json
- rawMainnetExtDataHashes []byte
- mainnetExtDataHashes    map[common.Hash]common.Hash
+ //go:embed flare_ext_data_hashes.json
+ rawFlareExtDataHashes []byte
+ flareExtDataHashes    map[common.Hash]common.Hash
 )

 func init() {
- if err := json.Unmarshal(rawFujiExtDataHashes, &fujiExtDataHashes); err != nil {
+ if err := json.Unmarshal(rawSongbirdExtDataHashes, &songbirdExtDataHashes); err != nil {
     panic(err)
 }
- rawFujiExtDataHashes = nil
- if err := json.Unmarshal(rawMainnetExtDataHashes, &mainnetExtDataHashes); err != nil {
+ rawSongbirdExtDataHashes = nil
+ if err := json.Unmarshal(rawFlareExtDataHashes, &flareExtDataHashes); err != nil {
     panic(err)
 }
- rawMainnetExtDataHashes = nil
+ rawFlareExtDataHashes = nil
 }
diff --git a/plugin/evm/factory.go b/plugin/evm/factory.go
index 34ecc893..79388ffc 100644
--- a/plugin/evm/factory.go
+++ b/plugin/evm/factory.go
@@ -4,9 +4,9 @@ package evm

 import (
-     "github.com/ava-labs/avalanchego/ids"
-     "github.com/ava-labs/avalanchego/snow"
-     "github.com/ava-labs/avalanchego/vms"
+     "github.com/flare-foundation/flare/ids"
+     "github.com/flare-foundation/flare/snow"
+     "github.com/flare-foundation/flare/vms"
 )

 var (
diff --git a/plugin/evm/mainnet_ext_data_hashes.json b/plugin/evm/flare_ext_data_hashes.json
similarity index 100%
rename from plugin/evm/mainnet_ext_data_hashes.json
rename to plugin/evm/flare_ext_data_hashes.json
diff --git a/plugin/evm/formatting.go b/plugin/evm/formatting.go
index d2194ea3..cf6c6682 100644
--- a/plugin/evm/formatting.go
+++ b/plugin/evm/formatting.go
@@ -6,12 +6,12 @@ package evm

 import (
-     "fmt"

```

```

-      "github.com/ava-labs/avalanchego/ids"
-      "github.com/ava-labs/avalanchego/utils/constants"
-      "github.com/ava-labs/avalanchego/utils/crypto"
-      "github.com/ava-labs/avalanchego/utils/formatting"
-      "github.com/ethereum/go-ethereum/common"
-      ethcrypto "github.com/ethereum/go-ethereum/crypto"
+      "github.com/flare-foundation/flare/ids"
+      "github.com/flare-foundation/flare/utils/constants"
+      "github.com/flare-foundation/flare/utils/crypto"
+      "github.com/flare-foundation/flare/utils/formatting"
+      "github.com/flare-foundation/flare/ids"
)

// ParseLocalAddress takes in an address for this chain and produces the ID
diff --git a/plugin/evm/gasprice_update.go b/plugin/evm/gasprice_update.go
index 71a4ea1a..177be08a 100644
--- a/plugin/evm/gasprice_update.go
+++ b/plugin/evm/gasprice_update.go
@@ -8,7 +8,7 @@ import (
     "sync"
     "time"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"
)

type gasPriceUpdater struct {
diff --git a/plugin/evm/gasprice_update_test.go b/plugin/evm/gasprice_update_test.go
index 24d8337f..5a720fb0 100644
--- a/plugin/evm/gasprice_update_test.go
+++ b/plugin/evm/gasprice_update_test.go
@@ -9,7 +9,7 @@ import (
     "testing"
     "time"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"
)

type mockGasPriceSetter struct {
diff --git a/plugin/evm/gossiper.go b/plugin/evm/gossiper.go
new file mode 100644
index 00000000..2a73873c
--- /dev/null
+++ b/plugin/evm/gossiper.go
@@ -0,0 +1,516 @@
+// (c) 2019-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "container/heap"
+    "math/big"
+    "sync"
+    "time"
+
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/coreth/peer"
+
+    "github.com/flare-foundation/flare/cache"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/snow"
+    "github.com/flare-foundation/flare/utils/wrappers"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+)
+
+const (
+    // We allow [recentCacheSize] to be fairly large because we only store hashes
+    // in the cache, not entire transactions.
+    recentCacheSize = 512
+
+    // [ethTxGossipInterval] is how often we attempt to gossip newly seen
+    // transactions to other nodes.
+    ethTxGossipInterval = 500 * time.Millisecond
+)
+
+// Gossiper handles outgoing gossip of transactions
+type Gossiper interface {
+    // GossipAtomicTx sends AppGossip message containing the given [txs]
+    // error
+    GossipAtomicTx(txs []*Tx) error
+    // GossipEthTx sends AppGossip message containing the given [txs]
+    // error
+    GossipEthTx(txs []*types.Transaction) error
+}
+
+// pushGossiper is used to gossip transactions to the network
+type pushGossiper struct {
+    ctx                *snow.Context
+    gossipActivationTime time.Time
+    config             Config
+
+    client      peer.Client
+    blockchain  *core.BlockChain
+    txPool      *core.TxPool
+    atomicMempool *Mempool
+
+    // We attempt to batch transactions we need to gossip to avoid runaway
+    // amplification of mempool chatter.
+    ethTxToGossipChan chan []*types.Transaction
+    ethTxToGossip    map[common.Hash]*types.Transaction
+    lastGossiped      time.Time
+    shutdownChan      chan struct{}
+    shutdownWg        *sync.WaitGroup
+
+    // [recentAtomicTx] and [recentEthTx] prevent us from over-gossiping the
+    // same transaction in a short period of time.
+    recentAtomicTx *cache.LRU
+    recentEthTx    *cache.LRU
+
+    codec codec.Manager
+}
+
+// newPushGossiper constructs and returns a pushGossiper
+// assumes vm.chainConfig.ApricotPhase4BlockTimestamp is set
+func (vm *VM) newPushGossiper() Gossiper {
+    net := &pushGossiper{
+        ctx:                vm.ctx,
+        gossipActivationTime: time.Unix(vm.chainConfig.ApricotPhase4BlockTimestamp.Int64(), 0),
+        config:             vm.config,
+        client:             vm.client,
+        blockchain:         vm.chain.BlockChain(),
+        txPool:             vm.chain.GetTxPool(),
+        atomicMempool:      vm.mempool,
+        ethTxToGossipChan:  make(chan []*types.Transaction),
+        ethTxToGossip:      make(map[common.Hash]*types.Transaction),

```



```

+         shutdownChan:      vm.shutdownChan,
+         shutdownWg:        &vm.shutdownWg,
+         recentAtomicTxs:   &cache.LRU{Size: recentCacheSize},
+         recentEthTxs:      &cache.LRU{Size: recentCacheSize},
+         codec:             vm.networkCodec,
+     }
+     net.awaitEthTxGossip()
+     return net
+}
+
+// queueExecutableTxs attempts to select up to [maxTxs] from the tx pool for
+// regossiping.
+//
+// We assume that [txs] contains an array of nonce-ordered transactions for a given
+// account. This array of transactions can have gaps and start at a nonce lower
+// than the current state of an account.
+func (n *pushGossiper) queueExecutableTxs(state *state.StateDB, baseFee *big.Int, txs map[common.Address]types.Transactions, maxTxs int) types.Transactions {
+    // Setup heap for transactions
+    heads := make(types.TxBByPriceAndTime, 0, len(txs))
+    for addr, accountTxs := range txs {
+        // Short-circuit here to avoid performing an unnecessary state lookup
+        if len(accountTxs) == 0 {
+            continue
+        }
+
+        // Ensure any transactions regossiped are immediately executable
+        var (
+            currentNonce = state.GetNonce(addr)
+            tx          *types.Transaction
+        )
+        for _, accountTx := range accountTxs {
+            // The tx pool may be out of sync with current state, so we iterate
+            // through the account transactions until we get to one that is
+            // executable.
+            if accountTx.Nonce() == currentNonce {
+                tx = accountTx
+                break
+            }
+            // There may be gaps in the tx pool and we could jump past the nonce we'd
+            // like to execute.
+            if accountTx.Nonce() > currentNonce {
+                break
+            }
+        }
+        if tx == nil {
+            continue
+        }
+
+        // Don't try to regossip a transaction too frequently
+        if time.Since(tx.FirstSeen()) < n.config.TxRegossipFrequency.Duration {
+            continue
+        }
+
+        // Ensure the fee the transaction pays is valid at tip
+        wrapped, err := types.NewTxWithMinerFee(tx, baseFee)
+        if err != nil {
+            log.Debug(
+                "not queuing tx for regossip",
+                "tx", tx.Hash(),
+                "err", err,
+            )
+            continue
+        }
+
+        heads = append(heads, wrapped)
+    }
+    heap.Init(&heads)
+
+    // Add up to [maxTxs] transactions to be gossiped
+    queued := make([]*types.Transaction, 0, maxTxs)
+    for len(heads) > 0 && len(queued) < maxTxs {
+        tx := heads[0].Tx
+        queued = append(queued, tx)
+        heap.Pop(&heads)
+    }
+
+    return queued
+}
+
+// queueRegossipTxs finds the best transactions in the mempool and adds up to
+// [TxRegossipMaxSize] of them to [ethTxsToGossip].
+func (n *pushGossiper) queueRegossipTxs() types.Transactions {
+    // Fetch all pending transactions
+    pending := n.txPool.Pending(true)
+
+    // Split the pending transactions into locals and remotes
+    localTxs := make(map[common.Address]types.Transactions)
+    remoteTxs := pending
+    for _, account := range n.txPool.Locals() {
+        if txs := remoteTxs[account]; len(txs) > 0 {
+            delete(remoteTxs, account)
+            localTxs[account] = txs
+        }
+    }
+
+    // Add best transactions to be gossiped (preferring local txs)
+    tip := n.blockchain.CurrentBlock()
+    state, err := n.blockchain.StateAt(tip.Root())
+    if err != nil || state == nil {
+        log.Debug(
+            "could not get state at tip",
+            "tip", tip.Hash(),
+            "err", err,
+        )
+        return nil
+    }
+
+    localQueued := n.queueExecutableTxs(state, tip.BaseFee(), localTxs, n.config.TxRegossipMaxSize)
+    localCount := len(localQueued)
+    if localCount >= n.config.TxRegossipMaxSize {
+        return localQueued
+    }
+
+    remoteQueued := n.queueExecutableTxs(state, tip.BaseFee(), remoteTxs, n.config.TxRegossipMaxSize-localCount)
+    return append(localQueued, remoteQueued...)
+}
+
+// awaitEthTxGossip periodically gossips transactions that have been queued for
+// gossip at least once every [ethTxsGossipInterval].
+func (n *pushGossiper) awaitEthTxGossip() {
+    n.shutdownWg.Add(1)
+    go n.ctx.Log.RecoverAndPanic(func() {
+        defer n.shutdownWg.Done()
+
+        var (
+            gossipTicker   = time.NewTicker(ethTxsGossipInterval)
+            regossipTicker = time.NewTicker(n.config.TxRegossipFrequency.Duration)
+        )
+
+        for {
+            select {
+            case <-gossipTicker.C:
+                if attempted, err := n.gossipEthTxs(false); err != nil {
+                    log.Warn(
+                        "failed to send eth transactions",
+                        "len(txs)", attempted,
+                    )
+                }
+            }
+        }
+    })
+}

```

```

+         "err", err,
+     )
+ }
+ case <-regossipTicker.C:
+     for _, tx := range n.queueRegossipTxs() {
+         n.ethTxsToGossip[tx.Hash()] = tx
+     }
+     if attempted, err := n.gossipEthTxs(true); err != nil {
+         log.Warn(
+             "failed to send eth transactions",
+             "len(txs)", attempted,
+             "err", err,
+         )
+     }
+ case txs := <-n.ethTxsToGossipChan:
+     for _, tx := range txs {
+         n.ethTxsToGossip[tx.Hash()] = tx
+     }
+     if attempted, err := n.gossipEthTxs(false); err != nil {
+         log.Warn(
+             "failed to send eth transactions",
+             "len(txs)", attempted,
+             "err", err,
+         )
+     }
+ case <-n.shutdownChan:
+     return
+ }
+ })
+ })
+}
+
+func (n *pushGossiper) GossipAtomicTxs(txs []*Tx) error {
+    if time.Now().Before(n.gossipActivationTime) {
+        log.Trace(
+            "not gossiping atomic tx before the gossiping activation time",
+            "txs", txs,
+        )
+        return nil
+    }
+    errs := wrappers.Errors{}
+    for _, tx := range txs {
+        errs.Add(n.gossipAtomicTx(tx))
+    }
+    return errs.Err
+}
+
+func (n *pushGossiper) gossipAtomicTx(tx *Tx) error {
+    txID := tx.ID()
+    // Don't gossip transaction if it has been recently gossiped.
+    if _, has := n.recentAtomicTxs.Get(txID); has {
+        return nil
+    }
+    // If the transaction is not pending according to the mempool
+    // then there is no need to gossip it further.
+    if _, pending := n.atomicMempool.GetPendingTx(txID); !pending {
+        return nil
+    }
+    n.recentAtomicTxs.Put(txID, nil)
+
+    msg := message.AtomicTx{
+        Tx: tx.Bytes(),
+    }
+    msgBytes, err := message.BuildMessage(n.codec, &msg)
+    if err != nil {
+        return err
+    }
+
+    log.Trace(
+        "gossiping atomic tx",
+        "txID", txID,
+    )
+    return n.client.Gossip(msgBytes)
+}
+
+func (n *pushGossiper) sendEthTxs(txs []*types.Transaction) error {
+    if len(txs) == 0 {
+        return nil
+    }
+
+    txBytes, err := rlp.EncodeToBytes(txs)
+    if err != nil {
+        return err
+    }
+    msg := message.EthTxs{
+        Txs: txBytes,
+    }
+    msgBytes, err := message.BuildMessage(n.codec, &msg)
+    if err != nil {
+        return err
+    }
+
+    log.Trace(
+        "gossiping eth txs",
+        "len(txs)", len(txs),
+        "size(txs)", len(msg.Txs),
+    )
+    return n.client.Gossip(msgBytes)
+}
+
+func (n *pushGossiper) gossipEthTxs(force bool) (int, error) {
+    if (!force && time.Since(n.lastGossiped) < ethTxsGossipInterval) || len(n.ethTxsToGossip) == 0 {
+        return 0, nil
+    }
+    n.lastGossiped = time.Now()
+    txs := make([]*types.Transaction, 0, len(n.ethTxsToGossip))
+    for _, tx := range n.ethTxsToGossip {
+        txs = append(txs, tx)
+        delete(n.ethTxsToGossip, tx.Hash())
+    }
+
+    selectedTxs := make([]*types.Transaction, 0)
+    for _, tx := range txs {
+        txHash := tx.Hash()
+        txStatus := n.txPool.Status([]*common.Hash{txHash})[0]
+        if txStatus != core.TxStatusPending {
+            continue
+        }
+
+        if n.config.RemoteTxGossipOnlyEnabled && n.txPool.HasLocal(txHash) {
+            continue
+        }
+
+        // We check [force] outside of the if statement to avoid an unnecessary
+        // cache lookup.
+        if !force {
+            if _, has := n.recentEthTxs.Get(txHash); has {
+                continue
+            }
+        }
+        n.recentEthTxs.Put(txHash, nil)
+
+        selectedTxs = append(selectedTxs, tx)
+    }

```

```

+     }
+
+     if len(selectedTxs) == 0 {
+         return 0, nil
+     }
+
+     // Attempt to gossip [selectedTxs]
+     msgTxs := make([]*types.Transaction, 0)
+     msgTxsSize := common.StorageSize(0)
+     for _, tx := range selectedTxs {
+         size := tx.Size()
+         if msgTxsSize+size > message.EthMsgSoftCapSize {
+             if err := n.sendEthTxs(msgTxs); err != nil {
+                 return len(selectedTxs), err
+             }
+             msgTxs = msgTxs[:0]
+             msgTxsSize = 0
+         }
+         msgTxs = append(msgTxs, tx)
+         msgTxsSize += size
+     }
+
+     // Send any remaining [msgTxs]
+     return len(selectedTxs), n.sendEthTxs(msgTxs)
+}
+
+// GossipEthTxs enqueues the provided [txs] for gossiping. At some point, the
+// [pushGossiper] will attempt to gossip the provided txs to other nodes
+// (usually right away if not under load).
+//
+// NOTE: We never return a non-nil error from this function but retain the
+// option to do so in case it becomes useful.
+func (n *pushGossiper) GossipEthTxs(txs []*types.Transaction) error {
+    if time.Now().Before(n.gossipActivationTime) {
+        log.Trace(
+            "not gossiping eth txs before the gossiping activation time",
+            "len(txs)", len(txs),
+        )
+        return nil
+    }
+
+    select {
+    case n.ethTxsToGossipChan <- txs:
+    case <-n.shutdownChan:
+    }
+    return nil
+}
+
+// GossipHandler handles incoming gossip messages
+type GossipHandler struct {
+    vm          *VM
+    atomicMempool *Mempool
+    txPool      *core.TxPool
+}
+
+func NewGossipHandler(vm *VM) *GossipHandler {
+    return &GossipHandler{
+        vm:          vm,
+        atomicMempool: vm.mempool,
+        txPool:      vm.chain.GetTxPool(),
+    }
+}
+
+func (h *GossipHandler) HandleAtomicTx(nodeID ids.ShortID, msg *message.AtomicTx) error {
+    log.Trace(
+        "AppGossip called with AtomicTx",
+        "peerID", nodeID,
+    )
+
+    if len(msg.Tx) == 0 {
+        log.Trace(
+            "AppGossip received empty AtomicTx Message",
+            "peerID", nodeID,
+        )
+        return nil
+    }
+
+    // In the case that the gossip message contains a transaction,
+    // attempt to parse it and add it as a remote.
+    tx := Tx{}
+    if _, err := Codec.Unmarshal(msg.Tx, &tx); err != nil {
+        log.Trace(
+            "AppGossip provided invalid tx",
+            "err", err,
+        )
+        return nil
+    }
+    unsignedBytes, err := Codec.Marshal(codecVersion, &tx.UnsignedAtomicTx)
+    if err != nil {
+        log.Trace(
+            "AppGossip failed to marshal unsigned tx",
+            "err", err,
+        )
+        return nil
+    }
+    tx.Initialize(unsignedBytes, msg.Tx)
+
+    txID := tx.ID()
+    if _, dropped, found := h.atomicMempool.GetTx(txID); found || dropped {
+        return nil
+    }
+
+    if err := h.vm.issueTx(&tx, false /*=local*/); err != nil {
+        log.Trace(
+            "AppGossip provided invalid transaction",
+            "peerID", nodeID,
+            "err", err,
+        )
+    }
+
+    return nil
+}
+
+func (h *GossipHandler) HandleEthTxs(nodeID ids.ShortID, msg *message.EthTxs) error {
+    log.Trace(
+        "AppGossip called with EthTxs",
+        "peerID", nodeID,
+        "size(txs)", len(msg.Txs),
+    )
+
+    if len(msg.Txs) == 0 {
+        log.Trace(
+            "AppGossip received empty EthTxs Message",
+            "peerID", nodeID,
+        )
+        return nil
+    }
+
+    // The maximum size of this encoded object is enforced by the codec.
+    txs := make([]*types.Transaction, 0)
+    if err := rlp.DecodeBytes(msg.Txs, &txs); err != nil {
+        log.Trace(
+            "AppGossip provided invalid txs",
+            "peerID", nodeID,

```

```

+         "err", err,
+     )
+     return nil
+ }
+ errs := h.txPool.AddRemotes(txs)
+ for i, err := range errs {
+     if err != nil {
+         log.Trace(
+             "AppGossip failed to add to mempool",
+             "err", err,
+             "tx", txs[i].Hash(),
+         )
+     }
+ }
+ return nil
+}
+
+// noopGossiper should be used when gossip communication is not supported
+type noopGossiper struct{}
+
+func (n *noopGossiper) GossipAtomicTxs([]*Tx) error {
+    return nil
+}
+
+func (n *noopGossiper) GossipEthTxs([]*types.Transaction) error {
+    return nil
+}
+
diff --git a/plugin/evm/network_atomic_gossiping_test.go b/plugin/evm/gossiper_atomic_gossiping_test.go
similarity index 91%
rename from plugin/evm/network_atomic_gossiping_test.go
rename to plugin/evm/gossiper_atomic_gossiping_test.go
index 501bc5c8..15bc3322 100644
--- a/plugin/evm/network_atomic_gossiping_test.go
+++ b/plugin/evm/gossiper_atomic_gossiping_test.go
@@ -8,11 +8,11 @@ import (
     "testing"
     "time"

-    "github.com/ava-labs/avalanchego/ids"
+    "github.com/flare-foundation/flare/ids"

-    "github.com/ava-labs/coreth/plugin/evm/message"
+    "github.com/flare-foundation/coreth/plugin/evm/message"
 )

// locally issued txs should be gossiped
@@ -35,7 +35,7 @@ func TestMempoolAtmTxsIssueTxAndGossiping(t *testing.T) {
    gossipedLock.Lock()
    defer gossipedLock.Unlock()

-    notifyMsgIntf, err := message.Parse(gossipedBytes)
+    notifyMsgIntf, err := message.ParseMessage(vm.networkCodec, gossipedBytes)
    assert.NoError(err)

    requestMsg, ok := notifyMsgIntf.(*message.AtomicTx)
@@ -61,7 +61,7 @@ func TestMempoolAtmTxsIssueTxAndGossiping(t *testing.T) {
    gossipedLock.Unlock()

    // Test hash on retry
-    assert.NoError(vm.network.GossipAtomicTxs([]*Tx{tx}))
+    assert.NoError(vm.gossiper.GossipAtomicTxs([]*Tx{tx}))
    gossipedLock.Lock()
    assert.Equal(1, gossiped)
    gossipedLock.Unlock()
@@ -111,7 +111,7 @@ func TestMempoolAtmTxsAppGossipHandling(t *testing.T) {
    msg := message.AtomicTx{
        Tx: tx.Bytes(),
    }

-    msgBytes, err := message.Build(&msg)
+    msgBytes, err := message.BuildMessage(vm.networkCodec, &msg)
    assert.NoError(err)

    // show that no txID is requested
@@ -134,7 +134,7 @@ func TestMempoolAtmTxsAppGossipHandling(t *testing.T) {
    msg = message.AtomicTx{
        Tx: conflictingTx.Bytes(),
    }

-    msgBytes, err = message.Build(&msg)
+    msgBytes, err = message.BuildMessage(vm.networkCodec, &msg)
    assert.NoError(err)
    assert.NoError(vm.AppGossip(nodeID, msgBytes))
    assert.False(txRequested, "tx should not have been requested")
@@ -179,7 +179,7 @@ func TestMempoolAtmTxsAppGossipHandlingDiscardedTx(t *testing.T) {

    mempool.AddTx(tx)
    mempool.NextTx()
-    mempool.DiscardCurrentTx()
+    mempool.DiscardCurrentTx(txID)

    // Check the mempool does not contain the discarded transaction
    assert.False(mempool.has(txID))
@@ -190,7 +190,7 @@ func TestMempoolAtmTxsAppGossipHandlingDiscardedTx(t *testing.T) {
    msg := message.AtomicTx{
        Tx: tx.Bytes(),
    }

-    msgBytes, err := message.Build(&msg)
+    msgBytes, err := message.BuildMessage(vm.networkCodec, &msg)
    assert.NoError(err)

    assert.NoError(vm.AppGossip(nodeID, msgBytes))
@@ -208,7 +208,7 @@ func TestMempoolAtmTxsAppGossipHandlingDiscardedTx(t *testing.T) {
    msg = message.AtomicTx{
        Tx: conflictingTx.Bytes(),
    }

-    msgBytes, err = message.Build(&msg)
+    msgBytes, err = message.BuildMessage(vm.networkCodec, &msg)
    assert.NoError(err)

    assert.NoError(vm.AppGossip(nodeID, msgBytes))
diff --git a/plugin/evm/gossiper_eth_gossiping_test.go b/plugin/evm/gossiper_eth_gossiping_test.go
new file mode 100644
index 00000000..af38466e
--- /dev/null
+++ b/plugin/evm/gossiper_eth_gossiping_test.go
@@ -0,0 +1,378 @@
+// (c) 2019-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "crypto/ecdsa"
+    "encoding/json"
+    "math/big"
+    "strings"
+    "sync"
+    "testing"
+    "time"
+
+    "github.com/flare-foundation/flare/ids"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"

```

```

+      "github.com/ethereum/go-ethereum/rlp"
+
+      "github.com/stretchr/testify/assert"
+
+      "github.com/flare-foundation/coreth/core"
+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/params"
+      "github.com/flare-foundation/coreth/plugin/evm/message"
+)
+
+func fundAddressByGenesis(addr []common.Address) (string, error) {
+    balance := big.NewInt(0xfffffffffff)
+    genesis := &core.Genesis{
+        Difficulty: common.Big0,
+        GasLimit:   uint64(5000000),
+    }
+    funds := make(map[common.Address]core.GenesisAccount)
+    for _, addr := range addr {
+        funds[addr] = core.GenesisAccount{
+            Balance: balance,
+        }
+    }
+    genesis.Alloc = funds
+
+    genesis.Config = &params.ChainConfig{
+        ChainID:           params.AvalancheLocalChainID,
+        ApricotPhase1BlockTimestamp: big.NewInt(0),
+        ApricotPhase2BlockTimestamp: big.NewInt(0),
+        ApricotPhase3BlockTimestamp: big.NewInt(0),
+        ApricotPhase4BlockTimestamp: big.NewInt(0),
+    }
+
+    bytes, err := json.Marshal(genesis)
+    return string(bytes), err
+}
+
+func getValidEthTxs(key *ecdsa.PrivateKey, count int, gasPrice *big.Int) []*types.Transaction {
+    res := make([]*types.Transaction, count)
+
+    to := common.Address{}
+    amount := big.NewInt(10000)
+    gasLimit := uint64(100000)
+
+    for i := 0; i < count; i++ {
+        tx, _ := types.SignTx(
+            types.NewTransaction(
+                uint64(i),
+                to,
+                amount,
+                gasLimit,
+                gasPrice,
+                []byte(strings.Repeat("aaaaaaaa", 100))),
+            types.HomesteadSigner{}, key)
+        tx.SetFirstSeen(time.Now().Add(-1 * time.Minute))
+        res[i] = tx
+    }
+    return res
+}
+
+// show that locally issued eth txs are gossiped
+// Note: channel through which coreth mempool push txs to vm is injected here
+// to ease up UT, which target only VM behaviors in response to coreth mempool
+// signals
+func TestMempoolEthTxsAddedTxsGossipedAfterActivation(t *testing.T) {
+    assert := assert.New(t)
+
+    key, err := crypto.GenerateKey()
+    assert.NoError(err)
+
+    addr := crypto.PubkeyToAddress(key.PublicKey)
+
+    cfgJson, err := fundAddressByGenesis([]common.Address{addr})
+    assert.NoError(err)
+
+    _, vm, _, sender := GenesisVM(t, true, cfgJson, "", "")
+    defer func() {
+        err := vm.Shutdown()
+        assert.NoError(err)
+    }()
+    vm.chain.GetTxPool().SetGasPrice(common.Big1)
+    vm.chain.GetTxPool().SetMinFee(common.Big0)
+
+    // create eth txes
+    ethTxs := getValidEthTxs(key, 3, common.Big1)
+
+    var wg sync.WaitGroup
+    wg.Add(2)
+    sender.CantSendAppGossip = false
+    signal1 := make(chan struct{})
+    seen := 0
+    sender.SendAppGossipF = func(gossipedBytes []byte) error {
+        if seen == 0 {
+            notifyMsgIntf, err := message.ParseMessage(vm.networkCodec, gossipedBytes)
+            assert.NoError(err)
+
+            requestMsg, ok := notifyMsgIntf.(*message.EthTxs)
+            assert.True(ok)
+            assert.NotEmpty(requestMsg.Txs)
+
+            txs := make([]*types.Transaction, 0)
+            assert.NoError(rlp.DecodeBytes(requestMsg.Txs, &txs))
+            assert.Len(txs, 2)
+            assert.ElementsMatch(
+                []common.Hash{ethTxs[0].Hash(), ethTxs[1].Hash()},
+                []common.Hash{txs[0].Hash(), txs[1].Hash()},
+            )
+            seen++
+            close(signal1)
+        } else if seen == 1 {
+            notifyMsgIntf, err := message.ParseMessage(vm.networkCodec, gossipedBytes)
+            assert.NoError(err)
+
+            requestMsg, ok := notifyMsgIntf.(*message.EthTxs)
+            assert.True(ok)
+            assert.NotEmpty(requestMsg.Txs)
+
+            txs := make([]*types.Transaction, 0)
+            assert.NoError(rlp.DecodeBytes(requestMsg.Txs, &txs))
+            assert.Len(txs, 1)
+            assert.Equal(ethTxs[2].Hash(), txs[0].Hash())
+
+            seen++
+        } else {
+            t.Fatal("should not be seen 3 times")
+        }
+    }
+    wg.Done()
+    return nil
+}
+
+// Notify VM about eth txs
+errs := vm.chain.GetTxPool().AddRemotesSync(ethTxs[:2])
+for _, err := range errs {
+    assert.NoError(err, "failed adding coreth tx to mempool")
+}

```

```

+ // Gossip txs again (shouldn't gossip hashes)
+ <-signal1 // wait until reorg processed
+ assert.NoError(vm.gossiper.GossipEthTxs(ethTxs[:2]))
+
+ errs = vm.chain.GetTxPool().AddRemotesSync(ethTxs)
+ assert.Contains(errs[0].Error(), "already known")
+ assert.Contains(errs[1].Error(), "already known")
+ assert.NoError(errs[2], "failed adding coreth tx to mempool")
+
+ attemptAwait(t, &wg, 5*time.Second)
+}
+
+// show that locally issued eth txs are chunked correctly
+func TestMempoolEthTxsAddedTxsGossipedAfterActivationChunking(t *testing.T) {
+    assert := assert.New(t)
+
+    key, err := crypto.GenerateKey()
+    assert.NoError(err)
+
+    addr := crypto.PubkeyToAddress(key.PublicKey)
+
+    cfgJson, err := fundAddressByGenesis([]common.Address{addr})
+    assert.NoError(err)
+
+    _, vm, _, sender := GenesisVM(t, true, cfgJson, "", "")
+    defer func() {
+        err := vm.Shutdown()
+        assert.NoError(err)
+    }()
+    vm.chain.GetTxPool().SetGasPrice(common.Big1)
+    vm.chain.GetTxPool().SetMinFee(common.Big0)
+
+    // create eth txes
+    ethTxs := getValidEthTxs(key, 100, common.Big1)
+
+    var wg sync.WaitGroup
+    wg.Add(2)
+    sender.CantSendAppGossip = false
+    seen := map[common.Hash]struct{}{}
+    sender.SendAppGossipF = func(gossipedBytes []byte) error {
+        notifyMsgIntf, err := message.ParseMessage(vm.networkCodec, gossipedBytes)
+        assert.NoError(err)
+
+        requestMsg, ok := notifyMsgIntf.(*message.EthTxs)
+        assert.True(ok)
+        assert.NotEmpty(requestMsg.Txs)
+
+        txs := make([]*types.Transaction, 0)
+        assert.NoError(rlp.DecodeBytes(requestMsg.Txs, &txs))
+        for _, tx := range txs {
+            seen[tx.Hash()] = struct{}{}
+        }
+        wg.Done()
+        return nil
+    }
+
+    // Notify VM about eth txs
+    errs := vm.chain.GetTxPool().AddRemotesSync(ethTxs)
+    for _, err := range errs {
+        assert.NoError(err, "failed adding coreth tx to mempool")
+    }
+
+    attemptAwait(t, &wg, 5*time.Second)
+
+    for _, tx := range ethTxs {
+        _, ok := seen[tx.Hash()]
+        assert.True(ok, "missing hash: %v", tx.Hash())
+    }
+}
+
+// show that a geth tx discovered from gossip is requested to the same node that
+// gossiped it
+func TestMempoolEthTxsAppGossipHandling(t *testing.T) {
+    assert := assert.New(t)
+
+    key, err := crypto.GenerateKey()
+    assert.NoError(err)
+
+    addr := crypto.PubkeyToAddress(key.PublicKey)
+
+    cfgJson, err := fundAddressByGenesis([]common.Address{addr})
+    assert.NoError(err)
+
+    _, vm, _, sender := GenesisVM(t, true, cfgJson, "", "")
+    defer func() {
+        err := vm.Shutdown()
+        assert.NoError(err)
+    }()
+    vm.chain.GetTxPool().SetGasPrice(common.Big1)
+    vm.chain.GetTxPool().SetMinFee(common.Big0)
+
+    var (
+        wg          sync.WaitGroup
+        txRequested bool
+    )
+    sender.CantSendAppGossip = false
+    sender.SendAppRequestF = func(_ ids.ShortSet, _ uint32, _ []byte) error {
+        txRequested = true
+        return nil
+    }
+
+    wg.Add(1)
+    sender.SendAppGossipF = func(_ []byte) error {
+        wg.Done()
+        return nil
+    }
+
+    // prepare a tx
+    tx := getValidEthTxs(key, 1, common.Big1)[0]
+
+    // show that unknown coreth hashes is requested
+    txBytes, err := rlp.EncodeToBytes([]*types.Transaction{tx})
+    assert.NoError(err)
+    msg := message.EthTxs{
+        Txs: txBytes,
+    }
+    msgBytes, err := message.BuildMessage(vm.networkCodec, &msg)
+    assert.NoError(err)
+
+    nodeID := ids.GenerateTestShortID()
+    err = vm.AppGossip(nodeID, msgBytes)
+    assert.NoError(err)
+    assert.False(txRequested, "tx should not be requested")
+
+    // wait for transaction to be re-gossiped
+    attemptAwait(t, &wg, 5*time.Second)
+}
+
+func TestMempoolEthTxsRegossipSingleAccount(t *testing.T) {
+    assert := assert.New(t)
+
+    key, err := crypto.GenerateKey()
+    assert.NoError(err)

```

[illegible]

```

-         return set
+         return ids.Set{}
    }

    // Verify this transaction is well-formed
    func (tx *UnsignedImportTx) Verify(
-         xChainID ids.ID,
-         ctx *snow.Context,
-         rules params.Rules,
    ) error {
-         switch {
-         case tx == nil:
-             return errNilTx
-         case tx.SourceChain != xChainID:
-             return errWrongChainID
-         case len(tx.ImportedInputs) == 0:
-             return errNoImportInputs
-         case tx.NetworkID != ctx.NetworkID:
-             return errWrongNetworkID
-         case ctx.ChainID != tx.BlockchainID:
-             return errWrongBlockchainID
-         case rules.IsApricotPhase3 && len(tx.Outs) == 0:
-             return errNoEVMOutputs
-         }

-         for _, out := range tx.Outs {
-             if err := out.Verify(); err != nil {
-                 return fmt.Errorf("EVM Output failed verification: %w", err)
-             }
-         }

-         for _, in := range tx.ImportedInputs {
-             if err := in.Verify(); err != nil {
-                 return fmt.Errorf("atomic input failed verification: %w", err)
-             }
-         }

-         if !avax.IsSortedAndUniqueTransferableInputs(tx.ImportedInputs) {
-             return errInputsNotSortedUnique
-         }

-         if rules.IsApricotPhase2 {
-             if !IsSortedAndUniqueEVMOutputs(tx.Outs) {
-                 return errOutputsNotSortedUnique
-             }
-         } else if rules.IsApricotPhase1 {
-             if !IsSortedEVMOutputs(tx.Outs) {
-                 return errOutputsNotSorted
-             }
-         }

-         return nil
+         return errImportTxDisabled
    }

-func (tx *UnsignedImportTx) GasUsed() (uint64, error) {
-    cost := calcBytesCost(len(tx.UnsignedBytes()))
-    for _, in := range tx.ImportedInputs {
-        inCost, err := in.In.Cost()
-        if err != nil {
-            return 0, err
-        }
-        cost, err = math.Add64(cost, inCost)
-        if err != nil {
-            return 0, err
-        }
-    }
-    return cost, nil
+func (tx *UnsignedImportTx) GasUsed(fixedFee bool) (uint64, error) {
+    return 0, errImportTxDisabled
}

// Amount of [assetID] burned by this transaction
func (tx *UnsignedImportTx) Burned(assetID ids.ID) (uint64, error) {
-    var (
-        spent uint64
-        input uint64
-        err error
-    )
-    for _, out := range tx.Outs {
-        if out.AssetId == assetID {
-            spent, err = math.Add64(spent, out.Amount)
-            if err != nil {
-                return 0, err
-            }
-        }
-    }
-    for _, in := range tx.ImportedInputs {
-        if in.AssetId() == assetID {
-            input, err = math.Add64(input, in.Input().Amount())
-            if err != nil {
-                return 0, err
-            }
-        }
-    }
-    return math.Sub64(input, spent)
+    return 0, errImportTxDisabled
}

// SemanticVerify this transaction is valid.
@@ -145,82 +67,7 @@ func (tx *UnsignedImportTx) SemanticVerify(
    baseFee *big.Int,
    rules params.Rules,
) error {
-    if err := tx.Verify(vm.ctx.XChainID, vm.ctx, rules); err != nil {
-        return err
-    }

-    // Check the transaction consumes and produces the right amounts
-    fc := avax.NewFlowChecker()
-    switch {
-    // Apply dynamic fees to import transactions as of Apricot Phase 3
-    case rules.IsApricotPhase3:
-        gasUsed, err := stx.GasUsed()
-        if err != nil {
-            return err
-        }
-        txFee, err := calculateDynamicFee(gasUsed, baseFee)
-        if err != nil {
-            return err
-        }
-        fc.Produce(vm.ctx.AVAXAssetID, txFee)

-    // Apply fees to import transactions as of Apricot Phase 2
-    case rules.IsApricotPhase2:
-        fc.Produce(vm.ctx.AVAXAssetID, params.AvalancheAtomicTxFee)
-    }
-    for _, out := range tx.Outs {
-        fc.Produce(out.AssetId, out.Amount)
-    }
-    for _, in := range tx.ImportedInputs {
-        fc.Consume(in.AssetId(), in.Input().Amount())
-    }
}

```



```

-     if err := fc.Verify(); err != nil {
-         return fmt.Errorf("import tx flow check failed due to: %w", err)
-     }
-
-     if len(stx.Creds) != len(tx.ImportedInputs) {
-         return fmt.Errorf("import tx contained mismatched number of inputs/credentials (%d vs. %d)", len(tx.ImportedInputs), len(stx.Creds))
-     }
-
-     if !vm.ctx.IsBootstrapped() {
-         // Allow for force committing during bootstrapping
-         return nil
-     }
-
-     utxoIDs := make([][]byte, len(tx.ImportedInputs))
-     for i, in := range tx.ImportedInputs {
-         inputID := in.UTXOID.InputID()
-         utxoIDs[i] = inputID[:]
-     }
-
-     // allUTXOBytes is guaranteed to be the same length as utxoIDs
-     allUTXOBytes, err := vm.ctx.SharedMemory.Get(tx.SourceChain, utxoIDs)
-     if err != nil {
-         return fmt.Errorf("failed to fetch import UTXOs from %s due to: %w", tx.SourceChain, err)
-     }
-
-     for i, in := range tx.ImportedInputs {
-         utxoBytes := allUTXOBytes[i]
-
-         utxo := &avax.UTXO{}
-         if _, err := vm.codec.Unmarshal(utxoBytes, utxo); err != nil {
-             return fmt.Errorf("failed to unmarshal UTXO: %w", err)
-         }
-
-         cred := stx.Creds[i]
-
-         utxoAssetID := utxo.AssetID()
-         inAssetID := in.AssetID()
-         if utxoAssetID != inAssetID {
-             return errAssetIDMismatch
-         }
-
-         if err := vm.fx.VerifyTransfer(tx, in.In, cred, utxo.Out); err != nil {
-             return fmt.Errorf("import tx transfer failed verification: %w", err)
-         }
-     }
-
-     return vm.conflicts(tx.InputUTXOs(), parent)
+ return errImportTxDisabled
}

// Accept this transaction and spend imported inputs
@@ -228,13 +75,8 @@ func (tx *UnsignedImportTx) SemanticVerify(
// we don't want to remove an imported UTXO in semanticVerify
// only to have the transaction not be Accepted. This would be inconsistent.
// Recall that imported UTXOs are not kept in a versionDB.
-func (tx *UnsignedImportTx) Accept(ctx *snow.Context, batch database.Batch) error {
-     utxoIDs := make([][]byte, len(tx.ImportedInputs))
-     for i, in := range tx.ImportedInputs {
-         inputID := in.InputID()
-         utxoIDs[i] = inputID[:]
-     }
-     return ctx.SharedMemory.Apply(map[ids.ID]*atomic.Requests{tx.SourceChain: {RemoveRequests: utxoIDs}}, batch)
+func (tx *UnsignedImportTx) Accept() (ids.ID, *atomic.Requests, error) {
+ return ids.ID{}, nil, errImportTxDisabled
}

// newImportTx returns a new ImportTx
@@ -244,160 +86,22 @@ func (vm *VM) newImportTx(
baseFee *big.Int, // fee to use post-AP3
keys []crypto.PrivateKeySECP256K1R, // Keys to import the funds
) (*Tx, error) {
-     if vm.ctx.XChainID != chainID {
-         return nil, errWrongChainID
-     }
-
-     kc := secp256k1fx.NewKeychain()
-     for _, key := range keys {
-         kc.Add(key)
-     }
-
-     atomicUTXOs, _, err := vm.GetAtomicUTXOs(chainID, kc.Addresses(), ids.ShortEmpty, ids.Empty, -1)
-     if err != nil {
-         return nil, fmt.Errorf("problem retrieving atomic UTXOs: %w", err)
-     }
-
-     importedInputs := []*avax.TransferableInput{}
-     signers := [][]*crypto.PrivateKeySECP256K1R{}
-
-     importedAmount := make(map[ids.ID]uint64)
-     now := vm.clock.Unix()
-     for _, utxo := range atomicUTXOs {
-         inputIntf, utxoSigners, err := kc.Spend(utxo.Out, now)
-         if err != nil {
-             continue
-         }
-         input, ok := inputIntf.(avax.TransferableIn)
-         if !ok {
-             continue
-         }
-         aid := utxo.AssetID()
-         importedAmount[aid], err = math.Add64(importedAmount[aid], input.Amount())
-         if err != nil {
-             return nil, err
-         }
-         importedInputs = append(importedInputs, &avax.TransferableInput{
-             UTXOID: utxo.UTXOID,
-             Asset:  utxo.Asset,
-             In:     input,
-         })
-         signers = append(signers, utxoSigners)
-     }
-     avax.SortTransferableInputsWithSigners(importedInputs, signers)
-     importedAVAXAmount := importedAmount[vm.ctx.AVAXAssetID]
-
-     outs := make([]EVMOutput, 0, len(importedAmount))
-     // This will create unique outputs (in the context of sorting)
-     // since each output will have a unique assetID
-     for assetID, amount := range importedAmount {
-         // Skip the AVAX amount since it is included separately to account for
-         // the fee
-         if assetID == vm.ctx.AVAXAssetID || amount == 0 {
-             continue
-         }
-         outs = append(outs, EVMOutput{
-             Address: to,
-             Amount:  amount,
-             AssetID: assetID,
-         })
-     }
-
-     rules := vm.currentRules()
-
-     var (
-         txFeeWithoutChange uint64
-         txFeeWithChange    uint64

```

```

)
switch {
case rules.IsApricotPhase3:
    if baseFee == nil {
        return nil, errNilBaseFeeApricotPhase3
    }
    utx := &UnsignedImportTx{
        NetworkID:    vm.ctx.NetworkID,
        BlockchainID:  vm.ctx.ChainID,
        Outs:          outs,
        ImportedInputs: importedInputs,
        SourceChain:   chainID,
    }
    tx := &Tx{UnsignedAtomicTx: utx}
    if err := tx.Sign(vm.codec, nil); err != nil {
        return nil, err
    }

    gasUsedWithoutChange, err := tx.GasUsed()
    if err != nil {
        return nil, err
    }
    gasUsedWithChange := gasUsedWithoutChange + EVMOutputGas

    txFeeWithoutChange, err = calculateDynamicFee(gasUsedWithoutChange, baseFee)
    if err != nil {
        return nil, err
    }
    txFeeWithChange, err = calculateDynamicFee(gasUsedWithChange, baseFee)
    if err != nil {
        return nil, err
    }
}
case rules.IsApricotPhase2:
    txFeeWithoutChange = params.AvalancheAtomicTxFee
    txFeeWithChange = params.AvalancheAtomicTxFee
}

// AVAX output
if importedAVAXAmount < txFeeWithoutChange { // imported amount goes toward paying tx fee
    return nil, errInsufficientFundsForFee
}

if importedAVAXAmount > txFeeWithChange {
    outs = append(outs, EVMOutput{
        Address: to,
        Amount:  importedAVAXAmount - txFeeWithChange,
        AssetID: vm.ctx.AVAXAssetID,
    })
}

// If no outputs are produced, return an error.
// Note: this can happen if there is exactly enough AVAX to pay the
// transaction fee, but no other funds to be imported.
if len(outs) == 0 {
    return nil, errNoEVMOutputs
}

SortEVMOutputs(outs)
return nil, errImportTxsDisabled
+}

// Create the transaction
utx := &UnsignedImportTx{
    NetworkID:    vm.ctx.NetworkID,
    BlockchainID:  vm.ctx.ChainID,
    Outs:          outs,
    ImportedInputs: importedInputs,
    SourceChain:   chainID,
}
tx := &Tx{UnsignedAtomicTx: utx}
if err := tx.Sign(vm.codec, signers); err != nil {
    return nil, err
}
return tx, utx.Verify(vm.ctx.XChainID, vm.ctx, vm.currentRules())
+// newImportTx returns a new ImportTx
+func (vm *VM) newImportTxWithUTXOs(
+    chainID ids.ID, // chain to import from
+    to common.Address, // Address of recipient
+    baseFee *big.Int, // fee to use post-AP3
+    kc *secp256k1fx.Keychain, // Keychain to use for signing the atomic UTXOs
+    atomicUTXOs []*avax.UTXO, // UTXOs to spend
+) (*Tx, error) {
+    return nil, errImportTxsDisabled
+}

// EVMStateTransfer performs the state transfer to increase the balances of
// accounts accordingly with the imported EVMOutputs
func (tx *UnsignedImportTx) EVMStateTransfer(ctx *snow.Context, state *state.StateDB) error {
    for _, to := range tx.Outs {
        if to.AssetID == ctx.AVAXAssetID {
            log.Debug("crosschain X->C", "addr", to.Address, "amount", to.Amount, "assetID", "AVAX")
            // If the asset is AVAX, convert the input amount in nAVAX to gWei by
            // multiplying by the x2c rate.
            amount := new(big.Int).Mul(
                new(big.Int).SetUint64(to.Amount), x2cRate)
            state.AddBalance(to.Address, amount)
        } else {
            log.Debug("crosschain X->C", "addr", to.Address, "amount", to.Amount, "assetID", to.AssetID)
            amount := new(big.Int).SetUint64(to.Amount)
            state.AddBalanceMultiCoin(to.Address, common.Hash(to.AssetID), amount)
        }
    }
    return nil
+    return errImportTxsDisabled
+}

diff --git a/plugin/evm/import_tx_test.go b/plugin/evm/import_tx_test.go
deleted file mode 100644
index b3604e1a..00000000
--- a/plugin/evm/import_tx_test.go
+++ /dev/null
@@ -1,1153 +0,0 @@
-// (c) 2019-2020, Ava Labs, Inc. All rights reserved.
-// See the file LICENSE for licensing terms.
-
-package evm
-
-import (
-    "math/big"
-    "testing"
-
-    "github.com/ava-labs/coreth/params"
-    "github.com/ethereum/go-ethereum/common"
-
-    "github.com/ava-labs/avalanchego/chains/atomic"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/utls/crypto"
-    "github.com/ava-labs/avalanchego/vms/components/avax"
-    "github.com/ava-labs/avalanchego/vms/secp256k1fx"
-)
-
-// createImportTxOptions adds a UTXO to shared memory and generates a list of import transactions sending this UTXO
-// to each of the three test keys (conflicting transactions)
-func createImportTxOptions(t *testing.T, vm *VM, sharedMemory *atomic.Memory) []*Tx {
-    utxo := &avax.UTXO{

```

```

-         UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-         Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
-         Out: &secp256k1fx.TransferOutput{
-             Amt: uint64(500000000),
-             OutputOwners: secp256k1fx.OutputOwners{
-                 Threshold: 1,
-                 Addrs:      []ids.ShortID{testKeys[0].PublicKey().Address()},
-             },
-         },
-     },
- }
- utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
- if err != nil {
-     t.Fatal(err)
- }
-
- xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
- inputID := utxo.InputID()
- if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
-     Key:   inputID[:],
-     Value: utxoBytes,
-     Traits: [][]byte{
-         testKeys[0].PublicKey().Address().Bytes(),
-     },
- }}}}); err != nil {
-     t.Fatal(err)
- }
-
- importTxs := make([]*Tx, 0, 3)
- for _, ethAddr := range testEthAddrs {
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, ethAddr, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-     importTxs = append(importTxs, importTx)
- }
-
- return importTxs
-}
-
-func TestImportTxVerify(t *testing.T) {
-    ctx := NewContext()
-
-    var importAmount uint64 = 10000000
-    txID := ids.GenerateTestID()
-    importTx := &UnsignedImportTx{
-        NetworkID:   ctx.NetworkID,
-        BlockchainID: ctx.ChainID,
-        SourceChain:  ctx.XChainID,
-        ImportedInputs: []*avax.TransferableInput{
-            {
-                UTXOID: avax.UTXOID{
-                    TxID:   txID,
-                    OutputIndex: uint32(0),
-                },
-                Asset: avax.Asset{ID: ctx.AVAXAssetID},
-                In: &secp256k1fx.TransferInput{
-                    Amt: importAmount,
-                    Input: secp256k1fx.Input{
-                        SigIndices: []uint32{0},
-                    },
-                },
-            },
-        },
-        {
-            UTXOID: avax.UTXOID{
-                TxID:   txID,
-                OutputIndex: uint32(1),
-            },
-            Asset: avax.Asset{ID: ctx.AVAXAssetID},
-            In: &secp256k1fx.TransferInput{
-                Amt: importAmount,
-                Input: secp256k1fx.Input{
-                    SigIndices: []uint32{0},
-                },
-            },
-        },
-    },
-    Outs: []EVMOutput{
-        {
-            Address: testEthAddrs[0],
-            Amount: importAmount - params.AvalancheAtomicTxFee,
-            AssetID: ctx.AVAXAssetID,
-        },
-        {
-            Address: testEthAddrs[1],
-            Amount: importAmount,
-            AssetID: ctx.AVAXAssetID,
-        },
-    },
- }
-
- // // Sort the inputs and outputs to ensure the transaction is canonical
- avax.SortTransferableInputs(importTx.ImportedInputs)
- SortEVMOutputs(importTx.Outs)
-
- tests := map[string]atomicTxVerifyTest{
-     "nil tx": {
-         generate: func(t *testing.T) UnsignedAtomicTx {
-             var importTx *UnsignedImportTx
-             return importTx
-         },
-         ctx:      ctx,
-         rules:    apricotRulesPhase0,
-         expectedErr: errNilTx.Error(),
-     },
-     "valid import tx": {
-         generate: func(t *testing.T) UnsignedAtomicTx {
-             return importTx
-         },
-         ctx:      ctx,
-         rules:    apricotRulesPhase0,
-         expectedErr: "", // Expect this transaction to be valid
-     },
-     "invalid network ID": {
-         generate: func(t *testing.T) UnsignedAtomicTx {
-             tx := *importTx
-             tx.NetworkID++
-             return &tx
-         },
-         ctx:      ctx,
-         rules:    apricotRulesPhase0,
-         expectedErr: errWrongNetworkID.Error(),
-     },
-     "invalid blockchain ID": {
-         generate: func(t *testing.T) UnsignedAtomicTx {
-             tx := *importTx
-             tx.BlockchainID = ids.GenerateTestID()
-             return &tx
-         },
-         ctx:      ctx,
-         rules:    apricotRulesPhase0,
-         expectedErr: errWrongBlockchainID.Error(),
-     },
-     "invalid source chain ID": {
-         generate: func(t *testing.T) UnsignedAtomicTx {

```

```

        tx := *importTx
        tx.SourceChain = ids.GenerateTestID()
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: errWrongChainID.Error(),
},
"no inputs": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.ImportedInputs = nil
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: errNoImportInputs.Error(),
},
"inputs sorted incorrectly": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.ImportedInputs = []*avax.TransferableInput{
            tx.ImportedInputs[1],
            tx.ImportedInputs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: errInputsNotSortedUnique.Error(),
},
"invalid input": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.ImportedInputs = []*avax.TransferableInput{
            tx.ImportedInputs[0],
            nil,
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: "atomic input failed verification",
},
"unsorted outputs phase 0 passes verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[1],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: "",
},
"non-unique outputs phase 0 passes verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[0],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase0,
    expectedErr: "",
},
"unsorted outputs phase 1 fails verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[1],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase1,
    expectedErr: errOutputsNotSorted.Error(),
},
"non-unique outputs phase 1 passes verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[0],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase1,
    expectedErr: "",
},
"outputs not sorted and unique phase 2 fails verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[0],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase2,
    expectedErr: errOutputsNotSortedUnique.Error(),
},
"outputs not sorted phase 2 fails verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            tx.Outs[1],
            tx.Outs[0],
        }
        return &tx
    },
    ctx:      ctx,
    rules:    apricotRulesPhase2,
    expectedErr: errOutputsNotSortedUnique.Error(),
},
"invalid EVMOutput fails verification": {
    generate: func(t *testing.T) UnsignedAtomicTx {
        tx := *importTx
        tx.Outs = []EVMOutput{
            {
                Address: testEthAddrs[0],
                Amount: 0,
                AssetID: testAvaxAssetID,
            },
        }
        return &tx
    },
},

```



```

-    avaxAssetID := ids.GenerateTestID()
-     antAssetID := ids.GenerateTestID()
-     chainID := ids.GenerateTestID()
-     xChainID := ids.GenerateTestID()
-     networkID := uint32(5)
-     importAmount := uint64(5000000)
-
-     tests := map[string]struct {
-         UnsignedImportTx *UnsignedImportTx
-         Keys              [][]*crypto.PrivateKeySECP256K1R
-
-         ExpectedGasUsed uint64
-         ExpectedFee      uint64
-         BaseFee          *big.Int
-     }{
-         "simple import": {
-             UnsignedImportTx: &UnsignedImportTx{
-                 NetworkID:    networkID,
-                 BlockchainID: chainID,
-                 SourceChain:   xChainID,
-                 ImportedInputs: []*avax.TransferableInput{
-                     {
-                         UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                         Asset:  avax.Asset{ID: avaxAssetID},
-                         In:    &secp256k1fx.TransferInput{
-                             Amt:    importAmount,
-                             Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                         },
-                     },
-                 },
-                 Outs: []EVMOutput{
-                     {
-                         Address: testEthAdrs[0],
-                         Amount:  importAmount,
-                         AssetID: avaxAssetID,
-                     },
-                 },
-                 Keys:      [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}},
-                 ExpectedGasUsed: 1230,
-                 ExpectedFee:    30750,
-                 BaseFee:       big.NewInt(25 * params.GWei),
-             },
-             "simple import lwei": {
-                 UnsignedImportTx: &UnsignedImportTx{
-                     NetworkID:    networkID,
-                     BlockchainID: chainID,
-                     SourceChain:   xChainID,
-                     ImportedInputs: []*avax.TransferableInput{
-                         {
-                             UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                             Asset:  avax.Asset{ID: avaxAssetID},
-                             In:    &secp256k1fx.TransferInput{
-                                 Amt:    importAmount,
-                                 Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                             },
-                         },
-                     },
-                     Outs: []EVMOutput{
-                         {
-                             Address: testEthAdrs[0],
-                             Amount:  importAmount,
-                             AssetID: avaxAssetID,
-                         },
-                     },
-                     Keys:      [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}},
-                     ExpectedGasUsed: 1230,
-                     ExpectedFee:    1,
-                     BaseFee:       big.NewInt(1),
-                 },
-                 "simple ANT import": {
-                     UnsignedImportTx: &UnsignedImportTx{
-                         NetworkID:    networkID,
-                         BlockchainID: chainID,
-                         SourceChain:   xChainID,
-                         ImportedInputs: []*avax.TransferableInput{
-                             {
-                                 UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                                 Asset:  avax.Asset{ID: avaxAssetID},
-                                 In:    &secp256k1fx.TransferInput{
-                                     Amt:    importAmount,
-                                     Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                                 },
-                             },
-                             {
-                                 UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                                 Asset:  avax.Asset{ID: antAssetID},
-                                 In:    &secp256k1fx.TransferInput{
-                                     Amt:    importAmount,
-                                     Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                                 },
-                             },
-                         },
-                         Outs: []EVMOutput{
-                             {
-                                 Address: testEthAdrs[0],
-                                 Amount:  importAmount,
-                                 AssetID: antAssetID,
-                             },
-                         },
-                         Keys:      [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}, {testKeys[0]}},
-                         ExpectedGasUsed: 2318,
-                         ExpectedFee:    57950,
-                         BaseFee:       big.NewInt(25 * params.GWei),
-                     },
-                     "complex ANT import": {
-                         UnsignedImportTx: &UnsignedImportTx{
-                             NetworkID:    networkID,
-                             BlockchainID: chainID,
-                             SourceChain:   xChainID,
-                             ImportedInputs: []*avax.TransferableInput{
-                                 {
-                                     UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                                     Asset:  avax.Asset{ID: avaxAssetID},
-                                     In:    &secp256k1fx.TransferInput{
-                                         Amt:    importAmount,
-                                         Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                                     },
-                                 },
-                                 {
-                                     UTXOID: avax.UTXOID{TxID: ids.GenerateTestID()},
-                                     Asset:  avax.Asset{ID: antAssetID},
-                                     In:    &secp256k1fx.TransferInput{
-                                         Amt:    importAmount,
-                                         Input: secp256k1fx.Input{SigIndices: []uint32{0}},
-                                     },
-                                 },
-                             },
-                             Outs: []EVMOutput{
-                                 {
-                                     Address: testEthAdrs[0],
-                                     Amount:  importAmount,
-                                     AssetID: avaxAssetID,
-                                 },
-                                 {
-                                     Address: testEthAdrs[0],
-                                     Amount:  importAmount,
-                                     AssetID: antAssetID,
-                                 },
-                             },
-                         },
-                     },
-                 },
-             },
-         },
-     }

```

[illegible]

```

        {testKeys[0]},
        {testKeys[0]}},
    },
    ExpectedGasUsed: 11022,
    ExpectedFee: 275550,
    BaseFee: big.NewInt(25 * params.GWei),
},
}

for name, test := range tests {
    t.Run(name, func(t *testing.T) {
        tx := &Tx{UnsignedAtomicTx: test.UnsignedImportTx}

        // Sign with the correct key
        if err := tx.Sign(Codec, test.Keys); err != nil {
            t.Fatal(err)
        }

        gasUsed, err := tx.GasUsed()
        if err != nil {
            t.Fatal(err)
        }
        if gasUsed != test.ExpectedGasUsed {
            t.Fatalf("Expected gasUsed to be %d, but found %d", test.ExpectedGasUsed, gasUsed)
        }

        fee, err := calculateDynamicFee(gasUsed, test.BaseFee)
        if err != nil {
            t.Fatal(err)
        }
        if fee != test.ExpectedFee {
            t.Fatalf("Expected fee to be %d, but found %d", test.ExpectedFee, fee)
        }
    })
}

-}

-func TestImportTxSemanticVerify(t *testing.T) {
    tests := map[string]atomicTxTest{
        "UTXO not present during bootstrapping": {
            setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
                tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
                    NetworkID: vm.ctx.NetworkID,
                    BlockchainID: vm.ctx.ChainID,
                    SourceChain: vm.ctx.XChainID,
                    ImportedInputs: []*avax.TransferableInput{{
                        UTXOID: avax.UTXOID{
                            TxID: ids.GenerateTestID(),
                        },
                        Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                        In: &secp256k1fx.TransferInput{
                            Amt: 1,
                            Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                        },
                    }},
                    Outs: []EVMOutput{{
                        Address: testEthAdrs[0],
                        Amount: 1,
                        AssetID: vm.ctx.AVAXAssetID,
                    }},
                }}
                if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R({testKeys[0]})}); err != nil {
                    t.Fatal(err)
                }
                return tx
            },
            bootstrapping: true,
        },
        "UTXO not present": {
            setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
                tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
                    NetworkID: vm.ctx.NetworkID,
                    BlockchainID: vm.ctx.ChainID,
                    SourceChain: vm.ctx.XChainID,
                    ImportedInputs: []*avax.TransferableInput{{
                        UTXOID: avax.UTXOID{
                            TxID: ids.GenerateTestID(),
                        },
                        Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                        In: &secp256k1fx.TransferInput{
                            Amt: 1,
                            Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                        },
                    }},
                    Outs: []EVMOutput{{
                        Address: testEthAdrs[0],
                        Amount: 1,
                        AssetID: vm.ctx.AVAXAssetID,
                    }},
                }}
                if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R({testKeys[0]})}); err != nil {
                    t.Fatal(err)
                }
                return tx
            },
            semanticVerifyErr: "failed to fetch import UTXOs from",
        },
        "garbage UTXO": {
            setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
                utxoID := avax.UTXOID{TxID: ids.GenerateTestID()}
                xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
                inputID := utxoID.InputID()
                if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
                    Key: inputID[:],
                    Value: []byte("hey there"),
                    Traits: [][32]byte{
                        testShortIDAdrs[0].Bytes(),
                    }},
                }}}); err != nil {
                    t.Fatal(err)
                }

                tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
                    NetworkID: vm.ctx.NetworkID,
                    BlockchainID: vm.ctx.ChainID,
                    SourceChain: vm.ctx.XChainID,
                    ImportedInputs: []*avax.TransferableInput{{
                        UTXOID: utxoID,
                        Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                        In: &secp256k1fx.TransferInput{
                            Amt: 1,
                            Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                        },
                    }},
                    Outs: []EVMOutput{{
                        Address: testEthAdrs[0],
                        Amount: 1,
                        AssetID: vm.ctx.AVAXAssetID,
                    }},
                }}
                if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R({testKeys[0]})}); err != nil {
                    t.Fatal(err)
                }
                return tx
            },
        },
    }
}

```



```

    },
    semanticVerifyErr: "failed to unmarshal UTXO",
},
"UTXO AssetID mismatch": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        expectedAssetID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, expectedAssetID, 1, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID}, // Use a different assetID then the actual UTXO
                In: &secp256k1fx.TransferInput{
                    Amt: 1,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{{
                Address: testEthAddrs[0],
                Amount: 1,
                AssetID: vm.ctx.AVAXAssetID,
            }},
        }}
        if err := tx.Sign(vm.codec, [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}}); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    semanticVerifyErr: errAssetIDMismatch.Error(),
},
"insufficient AVAX funds": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, 1, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                In: &secp256k1fx.TransferInput{
                    Amt: 1,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{{
                Address: testEthAddrs[0],
                Amount: 2, // Produce more output than is consumed by the transaction
                AssetID: vm.ctx.AVAXAssetID,
            }},
        }}
        if err := tx.Sign(vm.codec, [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}}); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    semanticVerifyErr: "import tx flow check failed due to",
},
"insufficient non-AVAX funds": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        assetID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, assetID, 1, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: assetID},
                In: &secp256k1fx.TransferInput{
                    Amt: 1,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{{
                Address: testEthAddrs[0],
                Amount: 2, // Produce more output than is consumed by the transaction
                AssetID: assetID,
            }},
        }}
        if err := tx.Sign(vm.codec, [][]*crypto.PrivateKeySECP256K1R{{testKeys[0]}}); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    semanticVerifyErr: "import tx flow check failed due to",
},
"no signatures": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, 1, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                In: &secp256k1fx.TransferInput{
                    Amt: 1,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{{
                Address: testEthAddrs[0],
                Amount: 1,
                AssetID: vm.ctx.AVAXAssetID,
            }},
        }}
        if err := tx.Sign(vm.codec, nil); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    semanticVerifyErr: "import tx flow check failed due to",
},

```

```

    },
    semanticVerifyErr: "import tx contained mismatched number of inputs/credentials",
},
"incorrect signature": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, 1, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                In: &secp256k1fx.TransferInput{
                    Amt: 1,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{{
                Address: testEthAddrs[0],
                Amount: 1,
                AssetID: vm.ctx.AVAXAssetID,
            }},
        }}
        // Sign the transaction with the incorrect key
        if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R{testKeys[1]}}); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    semanticVerifyErr: "import tx transfer failed verification",
},
"non-unique EVM Outputs": {
    setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
        txID := ids.GenerateTestID()
        utxo, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, 2, testShortIDAddrs[0])
        if err != nil {
            t.Fatal(err)
        }

        tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
            NetworkID:    vm.ctx.NetworkID,
            BlockchainID: vm.ctx.ChainID,
            SourceChain:   vm.ctx.XChainID,
            ImportedInputs: []*avax.TransferableInput{{
                UTXOID: utxo.UTXOID,
                Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                In: &secp256k1fx.TransferInput{
                    Amt: 2,
                    Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                },
            }},
            Outs: []EVMOutput{
                {
                    Address: testEthAddrs[0],
                    Amount: 1,
                    AssetID: vm.ctx.AVAXAssetID,
                },
                {
                    Address: testEthAddrs[0],
                    Amount: 1,
                    AssetID: vm.ctx.AVAXAssetID,
                },
            },
        }}
        if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R{testKeys[0]}}); err != nil {
            t.Fatal(err)
        }
        return tx
    },
    genesisJSON:    genesisJSONApricotPhase3,
    semanticVerifyErr: errOutputsNotSortedUnique.Error(),
},
}

for name, test := range tests {
    t.Run(name, func(t *testing.T) {
        executeTxTest(t, test)
    })
}

-}

-func TestImportTxEVMStateTransfer(t *testing.T) {
    assetID := ids.GenerateTestID()
    tests := map[string]atomicTxTest{
        "AVAX UTXO": {
            setup: func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx {
                txID := ids.GenerateTestID()
                utxo, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, 1, testShortIDAddrs[0])
                if err != nil {
                    t.Fatal(err)
                }

                tx := &Tx{UnsignedAtomicTx: &UnsignedImportTx{
                    NetworkID:    vm.ctx.NetworkID,
                    BlockchainID: vm.ctx.ChainID,
                    SourceChain:   vm.ctx.XChainID,
                    ImportedInputs: []*avax.TransferableInput{{
                        UTXOID: utxo.UTXOID,
                        Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
                        In: &secp256k1fx.TransferInput{
                            Amt: 1,
                            Input: secp256k1fx.Input{SigIndices: []uint32{0}},
                        },
                    }},
                    Outs: []EVMOutput{{
                        Address: testEthAddrs[0],
                        Amount: 1,
                        AssetID: vm.ctx.AVAXAssetID,
                    }},
                }}
                if err := tx.Sign(vm.codec, [][32]byte{crypto.PrivateKeySECP256K1R{testKeys[0]}}); err != nil {
                    t.Fatal(err)
                }
                return tx
            },
            checkState: func(t *testing.T, vm *VM) {
                lastAcceptedBlock := vm.LastAcceptedBlockInternal().(*Block)

                sdb, err := vm.chain.BlockState(lastAcceptedBlock.ethBlock)
                if err != nil {
                    t.Fatal(err)
                }

                avaxBalance := sdb.GetBalance(testEthAddrs[0])
                if avaxBalance.Cmp(x2cRate) != 0 {
                    t.Fatalf("Expected AVAX balance to be %d, found balance: %d", x2cRate, avaxBalance)
                }
            },
        },
    },

```



```

    }
    if tx, exists := m.discardedTxs.Get(txID); exists {
        return tx.(*Tx), true, true
    }
@@ -242,13 +243,13 @@ func (m *Mempool) GetTx(txID ids.ID) (*Tx, bool, bool) {
}

// IssueCurrentTx marks [currentTx] as issued if there is one
-func (m *Mempool) IssueCurrentTx() {
+func (m *Mempool) IssueCurrentTxs() {
    m.lock.Lock()
    defer m.lock.Unlock()

-    if m.currentTx != nil {
-        m.issuedTxs[m.currentTx.ID()] = m.currentTx
-        m.currentTx = nil
+    for txID := range m.currentTxs {
+        m.issuedTxs[txID] = m.currentTxs[txID]
+        delete(m.currentTxs, txID)
    }

    // If there are more transactions to be issued, add an item
@@ -258,30 +259,32 @@ func (m *Mempool) IssueCurrentTx() {
}

}

-// CancelCurrentTx marks the attempt to issue [currentTx]
+// CancelCurrentTx marks the attempt to issue [txID]
+// as being aborted. This should be called after NextTx returns [txID]
+// and the transaction [txID] cannot be included in the block, but should
+// not be discarded. For example, CancelCurrentTx should be called if including
+// the transaction will put the block above the atomic tx gas limit.
+func (m *Mempool) CancelCurrentTx(txID ids.ID) {
+    m.lock.Lock()
+    defer m.lock.Unlock()
+
+    if tx, ok := m.currentTxs[txID]; ok {
+        m.cancelTx(tx)
+    }
+}
+
+// [CancelCurrentTxs] marks the attempt to issue [currentTxs]
+// as being aborted. If this is called after a buildBlock error
+// caused by the atomic transaction, then DiscardCurrentTx should have been called
+// such that this call will have no effect and should not re-issue the invalid tx.
-func (m *Mempool) CancelCurrentTx() {
+func (m *Mempool) CancelCurrentTxs() {
    m.lock.Lock()
    defer m.lock.Unlock()

    // If building a block failed, put the currentTx back in [txs]
    // if it exists.
-    if m.currentTx != nil {
-        // Add tx to heap sorted by gasPrice
-        tx := m.currentTx
-        gasPrice, err := m.atomicTxGasPrice(tx)
-        if err == nil {
-            m.txHeap.Push(tx, gasPrice)
-        } else {
-            log.Error("failed to calculate atomic tx gas price while canceling current tx", "err", err)
-            m.utxoSet.Remove(tx.InputUTXOs().List()...)
-            m.discardedTxs.Put(tx.ID(), tx)
-        }
-        // If the err is not nil, we simply discard the transaction because it is
-        // invalid. This should never happen but we guard against the case it does.
-        m.currentTx = nil
+    for _, tx := range m.currentTxs {
+        m.cancelTx(tx)
    }

    // If there are more transactions to be issued, add an item
@@ -291,20 +292,22 @@ func (m *Mempool) CancelCurrentTx() {
}

}

-// DiscardCurrentTx marks [currentTx] as invalid and aborts the attempt
+// cancelTx removes [tx] from current transactions and moves it back into the
+// tx heap.
+// assumes the lock is held.
+func (m *Mempool) cancelTx(tx *Tx) {
+    // Add tx to heap sorted by gasPrice
+    gasPrice, err := m.atomicTxGasPrice(tx)
+    if err == nil {
+        m.txHeap.Push(tx, gasPrice)
+    } else {
+        // If the err is not nil, we simply discard the transaction because it is
+        // invalid. This should never happen but we guard against the case it does.
+        log.Error("failed to calculate atomic tx gas price while canceling current tx", "err", err)
+        m.utxoSet.Remove(tx.InputUTXOs().List()...)
+        m.discardedTxs.Put(tx.ID(), tx)
+    }
+    delete(m.currentTxs, tx.ID())
+}
+
+// DiscardCurrentTx marks a [tx] in the [currentTxs] map as invalid and aborts the attempt
+// to issue it since it failed verification.
+// Adding to Pending should be handled by CancelCurrentTx in this case.
-func (m *Mempool) DiscardCurrentTx() {
+func (m *Mempool) DiscardCurrentTx(txID ids.ID) {
    m.lock.Lock()
    defer m.lock.Unlock()

    if tx, ok := m.currentTxs[txID]; ok {
        m.discardCurrentTx(tx)
    }
+}
+
+// DiscardCurrentTxs marks all txs in [currentTxs] as discarded.
+func (m *Mempool) DiscardCurrentTxs() {
    m.lock.Lock()
    defer m.lock.Unlock()

-    if m.currentTx == nil {
-        return
+    for _, tx := range m.currentTxs {
+        m.discardCurrentTx(tx)
    }
+}

-    m.utxoSet.Remove(m.currentTx.InputUTXOs().List()...)
-    m.discardedTxs.Put(m.currentTx.ID(), m.currentTx)
-    m.currentTx = nil
+// discardCurrentTx discards [tx] from the set of current transactions.
+// Assumes the lock is held.
+func (m *Mempool) discardCurrentTx(tx *Tx) {
+    m.utxoSet.Remove(tx.InputUTXOs().List()...)
+    m.discardedTxs.Put(tx.ID(), tx)
+    delete(m.currentTxs, tx.ID())
+}

    // RemoveTx removes [txID] from the mempool completely.
@@ -313,9 +314,11 @@ func (m *Mempool) RemoveTx(txID ids.ID) {
    defer m.lock.Unlock()
}

```

```

        var removedTx *Tx
        if m.currentTx != nil && m.currentTx.ID() == txID {
            removedTx = m.currentTx
            m.currentTx = nil
        }
        if tx, ok := m.currentTxs[txID]; ok {
            removedTx = tx
            delete(m.currentTxs, txID)
        }
        if tx, ok := m.txHeap.Get(txID); ok {
            removedTx = tx
        }
diff --git a/plugin/evm/mempool_atomic_gossiping_test.go b/plugin/evm/mempool_atomic_gossiping_test.go
index 43c0f9b7..392f64b9 100644
--- a/plugin/evm/mempool_atomic_gossiping_test.go
+++ b/plugin/evm/mempool_atomic_gossiping_test.go
@@ -6,13 +6,13 @@ package evm
import (
    "testing"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"

-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/utls/crypto"
-    "github.com/ava-labs/avalanchego/vms/components/avax"
-    "github.com/ava-labs/avalanchego/vms/components/chain"
-    "github.com/ava-labs/avalanchego/vms/secp256k1fx"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utls/crypto"
+    "github.com/flare-foundation/flare/vms/components/avax"
+    "github.com/flare-foundation/flare/vms/components/chain"
+    "github.com/flare-foundation/flare/vms/secp256k1fx"

    "github.com/stretchr/testify/assert"
)
@@ -71,9 +71,7 @@ func TestMempoolAddLocallyCreateAtomicTx(t *testing.T) {
    evmBlk, ok := blk.(*chain.BlockWrapper).Block.(*Block)
    assert.True(ok, "unknown block type")

-    retrievedTx, err := vm.extractAtomicTx(evmBlk.ethBlock)
-    assert.NoError(err, "could not extract atomic tx")
-    assert.Equal(txID, retrievedTx.ID(), "block does not include expected transaction")
+    assert.Equal(txID, evmBlk.atomicTxs[0].ID(), "block does not include expected transaction")

    has = mempool.has(txID)
    assert.True(has, "tx should stay in mempool until block is accepted")
diff --git a/plugin/evm/message/codec.go b/plugin/evm/message/codec.go
index 8c172740..0e3f395d 100644
--- a/plugin/evm/message/codec.go
+++ b/plugin/evm/message/codec.go
@@ -1,14 +1,19 @@
<div data-bbox="21 373 103 380" data-label="Text">
<div data-bbox="21 380 276 390" data-label="Text">
<div data-bbox="21 390 276 400" data-label="Text">
<div data-bbox="21 400 233 419" data-label="Text">
<div data-bbox="21 426 97 433" data-label="Text">
<div data-bbox="21 440 63 447" data-label="Text">
<div data-bbox="21 447 304 484" data-label="Text">
<div data-bbox="21 484 313 534" data-label="Text">
<div data-bbox="21 541 172 578" data-label="Text">
<div data-bbox="21 585 309 622" data-label="Text">
<div data-bbox="21 622 197 643" data-label="Text">
<div data-bbox="21 643 285 734" data-label="Text">
<div data-bbox="21 734 371 744" data-label="Text">
<div data-bbox="21 744 168 754" data-label="Text">
<div data-bbox="21 754 187 764" data-label="Text">
<div data-bbox="21 764 168 774" data-label="Text">
<div data-bbox="21 774 233 793" data-label="Text">
<div data-bbox="21 793 243 825" data-label="Text">
<div data-bbox="21 832 489 869" data-label="Text">
<div data-bbox="21 876 243 889" data-label="Text">
<div data-bbox="21 896 253 906" data-label="Text">
<div data-bbox="21 906 319 949" data-label="Text">
<div data-bbox="21 956 438 986" data-label="Text">

```

```

+         return nil
+}
+
+func (NoopMempoolGossipHandler) HandleEthTxs(nodeID ids.ShortID, _ *EthTxs) error {
+    log.Debug("dropping unexpected EthTxs message", "peerID", nodeID)
+    return nil
+}
+
+// RequestHandler interface handles incoming requests from peers
+// Must have methods in format of handleType(context.Context, ids.ShortID, uint32, request Type) error
+// so that the Request object of relevant Type can invoke its respective handle method
+// on this struct.
+// Also see GossipHandler for implementation style.
+type RequestHandler interface{}
+
+// ResponseHandler handles response for a sent request
+// Only one of OnResponse or OnFailure is called for a given requestID, not both
+type ResponseHandler interface {
+    // OnResponse is invoked when the peer responded to a request
+    OnResponse(nodeID ids.ShortID, requestID uint32, response []byte) error
+    // OnFailure is invoked when there was a failure in processing a request
+    // The FailureReason outlines the underlying cause.
+    OnFailure(nodeID ids.ShortID, requestID uint32) error
+}
+
+>>>>>> upstream-v0.8.5-rc.2
diff --git a/plugin/evm/message/handler_test.go b/plugin/evm/message/handler_test.go
index 179ac56e..097f037e 100644
--- a/plugin/evm/message/handler_test.go
+++ b/plugin/evm/message/handler_test.go
@@ -6,21 +6,35 @@ package message
import (
    "testing"

-    "github.com/ava-labs/avalanchego/ids"
+<<<<<< HEAD
+    "github.com/stretchr/testify/assert"
+
+    "github.com/flare-foundation/flare/ids"
+=====
+    "github.com/flare-foundation/flare/ids"

    "github.com/stretchr/testify/assert"
>>>>>> upstream-v0.8.5-rc.2
)

type CounterHandler struct {
    AtomicTx, EthTxs int
}

+<<<<<< HEAD
+func (h *CounterHandler) HandleAtomicTx(ids.ShortID, uint32, *AtomicTx) error {
+=====
+func (h *CounterHandler) HandleAtomicTx(ids.ShortID, *AtomicTx) error {
+>>>>>> upstream-v0.8.5-rc.2
+    h.AtomicTx++
+    return nil
+}

+<<<<<< HEAD
+func (h *CounterHandler) HandleEthTxs(ids.ShortID, uint32, *EthTxs) error {
+=====
+func (h *CounterHandler) HandleEthTxs(ids.ShortID, *EthTxs) error {
+>>>>>> upstream-v0.8.5-rc.2
+    h.EthTxs++
+    return nil
+}

@@ -31,7 +45,11 @@ func TestHandleAtomicTx(t *testing.T) {
    handler := CounterHandler{}
    msg := AtomicTx{}

+<<<<<< HEAD
+    err := msg.Handle(&handler, ids.ShortEmpty, 0)
+=====
+    err := msg.Handle(&handler, ids.ShortEmpty)
+>>>>>> upstream-v0.8.5-rc.2
+    assert.NoError(err)
+    assert.Equal(1, handler.AtomicTx)
+    assert.Zero(handler.EthTxs)
@@ -43,7 +61,11 @@ func TestHandleEthTxs(t *testing.T) {
    handler := CounterHandler{}
    msg := EthTxs{}

+<<<<<< HEAD
+    err := msg.Handle(&handler, ids.ShortEmpty, 0)
+=====
+    err := msg.Handle(&handler, ids.ShortEmpty)
+>>>>>> upstream-v0.8.5-rc.2
+    assert.NoError(err)
+    assert.Zero(handler.AtomicTx)
+    assert.Equal(1, handler.EthTxs)
@@ -52,11 +74,20 @@ func TestHandleEthTxs(t *testing.T) {
    func TestNoopHandler(t *testing.T) {
        assert := assert.New(t)

+<<<<<< HEAD
+        handler := NoopHandler{}
+
+        err := handler.HandleAtomicTx(ids.ShortEmpty, 0, nil)
+        assert.NoError(err)
+
+        err = handler.HandleEthTxs(ids.ShortEmpty, 0, nil)
+=====
+        handler := NoopMempoolGossipHandler{}
+
+        err := handler.HandleAtomicTx(ids.ShortEmpty, nil)
+        assert.NoError(err)
+
+        err = handler.HandleEthTxs(ids.ShortEmpty, nil)
+>>>>>> upstream-v0.8.5-rc.2
+        assert.NoError(err)
+    }

diff --git a/plugin/evm/message/message.go b/plugin/evm/message/message.go
index 082adbe5..4286bd03 100644
--- a/plugin/evm/message/message.go
+++ b/plugin/evm/message/message.go
@@ -6,10 +6,19 @@ package message
import (
    "errors"

+<<<<<< HEAD
+    "github.com/ethereum/go-ethereum/common"
+
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/utils/units"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utils/units"
+=====
+    "github.com/flare-foundation/flare/codec"
+
+    "github.com/ethereum/go-ethereum/common"
+
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utils/units"
+>>>>>> upstream-v0.8.5-rc.2

```

```

)

const (
@@ -18,6 +27,11 @@ const (
    // this size, however. Max inbound message size is enforced by the codec
    // (512KB).
    EthMsgSoftCapSize = common.StorageSize(64 * units.KiB)
+<<<<<< HEAD
+=====
+    atomicTxType      = "atomic-tx"
+    ethTxType         = "eth-txs"
+>>>>>> upstream-v0.8.5-rc.2
)

var (
@@ -29,7 +43,11 @@ var (

type Message interface {
    // Handle this message with the correct message handler
+<<<<<< HEAD
    Handle(handler Handler, nodeID ids.ShortID, requestID uint32) error
+=====
+    Handle(handler GossipHandler, nodeID ids.ShortID) error
+>>>>>> upstream-v0.8.5-rc.2

    // initialize should be called whenever a message is built or parsed
    initialize([]byte)
@@ -38,6 +56,12 @@ type Message interface {
    //
    // Bytes should only be called after being initialized
    Bytes() []byte
+<<<<<< HEAD
+=====
+
+    // Type returns user-friendly name for this object that can be used for logging
+    Type() string
+>>>>>> upstream-v0.8.5-rc.2
}

type message []byte
@@ -51,8 +75,17 @@ type AtomicTx struct {
    Tx []byte `serialize:"true"`
}

+<<<<<< HEAD
+func (msg *AtomicTx) Handle(handler Handler, nodeID ids.ShortID, requestID uint32) error {
+    return handler.HandleAtomicTx(nodeID, requestID, msg)
+=====
+func (msg *AtomicTx) Handle(handler GossipHandler, nodeID ids.ShortID) error {
+    return handler.HandleAtomicTx(nodeID, msg)
+}
+
+func (msg *AtomicTx) Type() string {
+    return atomicTxType
+>>>>>> upstream-v0.8.5-rc.2
}

type EthTx struct {
@@ -61,6 +94,7 @@ type EthTx struct {
    Tx []byte `serialize:"true"`
}

+<<<<<< HEAD
+func (msg *EthTx) Handle(handler Handler, nodeID ids.ShortID, requestID uint32) error {
+    return handler.HandleEthTx(nodeID, requestID, msg)
+}
@@ -72,14 +106,36 @@ func Parse(bytes []byte) (Message, error) {
    return nil, err
}

    if version != codecVersion {
+=====
+func (msg *EthTx) Handle(handler GossipHandler, nodeID ids.ShortID) error {
+    return handler.HandleEthTx(nodeID, msg)
+}
+
+func (msg *EthTx) Type() string {
+    return ethTxType
+}
+
+func ParseMessage(codec codec.Manager, bytes []byte) (Message, error) {
+    var msg Message
+    version, err := codec.Unmarshal(bytes, &msg)
+    if err != nil {
+        return nil, err
+    }
+    if version != Version {
+>>>>>> upstream-v0.8.5-rc.2
+        return nil, errUnexpectedCodecVersion
+    }
+    msg.initialize(bytes)
+    return msg, nil
+}

+<<<<<< HEAD
+func Build(msg Message) ([]byte, error) {
+    bytes, err := c.Marshal(codecVersion, &msg)
+=====
+func BuildMessage(codec codec.Manager, msg Message) ([]byte, error) {
+    bytes, err := codec.Marshal(Version, &msg)
+>>>>>> upstream-v0.8.5-rc.2
+    msg.initialize(bytes)
+    return bytes, err
+}

diff --git a/plugin/evm/message/message_test.go b/plugin/evm/message/message_test.go
index 5575e253..0c09bab0 100644
--- a/plugin/evm/message/message_test.go
+++ b/plugin/evm/message/message_test.go
@@ -6,10 +6,17 @@ package message
 import (
     "testing"

-    "github.com/ava-labs/avalanchego/units"
-    "github.com/ava-labs/avalanchego/units/units"
+<<<<<< HEAD
+    "github.com/stretchchr/testify/assert"
+
+    "github.com/flare-foundation/flare/units"
+    "github.com/flare-foundation/flare/units/units"
+=====
+    "github.com/flare-foundation/flare/units"
+    "github.com/flare-foundation/flare/units/units"
+
+    "github.com/stretchchr/testify/assert"
+>>>>>> upstream-v0.8.5-rc.2
)

func TestAtomicTx(t *testing.T) {
@@ -19,11 +26,21 @@ func TestAtomicTx(t *testing.T) {
    builtMsg := AtomicTx{
        Tx: msg,
    }
+<<<<<< HEAD
    builtMsgBytes, err := Build(&builtMsg)
    assert.NoError(err)

```

```

    assert.Equal(builtMsgBytes, builtMsg.Bytes())
    parsedMsgIntf, err := Parse(builtMsgBytes)
+=====
+   codec, err := BuildCodec()
+   assert.NoError(err)
+   builtMsgBytes, err := BuildMessage(codec, &builtMsg)
+   assert.NoError(err)
+   assert.Equal(builtMsgBytes, builtMsg.Bytes())
+
+   parsedMsgIntf, err := ParseMessage(codec, builtMsgBytes)
+>>>>>> upstream-v0.8.5-rc.2
+   assert.NoError(err)
+   assert.Equal(builtMsgBytes, parsedMsgIntf.Bytes())

@@ -40,11 +57,21 @@ func TestEthTxs(t *testing.T) {
    builtMsg := EthTxs{
        Txs: msg,
    }
+<<<<<< HEAD
+   builtMsgBytes, err := Build(&builtMsg)
+   assert.NoError(err)
+   assert.Equal(builtMsgBytes, builtMsg.Bytes())
+
+   parsedMsgIntf, err := Parse(builtMsgBytes)
+=====
+   codec, err := BuildCodec()
+   assert.NoError(err)
+   builtMsgBytes, err := BuildMessage(codec, &builtMsg)
+   assert.NoError(err)
+   assert.Equal(builtMsgBytes, builtMsg.Bytes())
+
+   parsedMsgIntf, err := ParseMessage(codec, builtMsgBytes)
+>>>>>> upstream-v0.8.5-rc.2
+   assert.NoError(err)
+   assert.Equal(builtMsgBytes, parsedMsgIntf.Bytes())

@@ -60,14 +87,27 @@ func TestEthTxsTooLarge(t *testing.T) {
    builtMsg := EthTxs{
        Txs: utils.RandomBytes(1024 * units.KiB),
    }
+<<<<<< HEAD
+   _, err := Build(&builtMsg)
+=====
+   codec, err := BuildCodec()
+   assert.NoError(err)
+   _, err = BuildMessage(codec, &builtMsg)
+>>>>>> upstream-v0.8.5-rc.2
+   assert.Error(err)
+
+   func TestParseGibberish(t *testing.T) {
+       assert := assert.New(t)
+
+<<<<<< HEAD
+   randomBytes := utils.RandomBytes(256 * units.KiB)
+   _, err := Parse(randomBytes)
+=====
+   codec, err := BuildCodec()
+   assert.NoError(err)
+   randomBytes := utils.RandomBytes(256 * units.KiB)
+   _, err = ParseMessage(codec, randomBytes)
+>>>>>> upstream-v0.8.5-rc.2
+   assert.Error(err)
+
+   }
diff --git a/plugin/evm/message/request.go b/plugin/evm/message/request.go
new file mode 100644
index 00000000..929e38eb
--- /dev/null
+++ b/plugin/evm/message/request.go
@@ -0,0 +1,36 @@
+// (c) 2019-2022, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package message
+
+import (
+    "context"
+
+    "github.com/flare-foundation/flare/codec"
+
+    "github.com/flare-foundation/flare/ids"
+)
+
+// Request represents a Network request type
+type Request interface {
+    // Handle allows 'Request' to call respective methods on handler to handle
+    // this particular request type
+    Handle(ctx context.Context, nodeID ids.ShortID, requestID uint32, handler RequestHandler) ([]byte, error)
+
+    // Type returns user-friendly name for this object that can be used for logging
+    Type() string
+}
+
+// BytesToRequest unmarshals the given requestBytes into Request object
+func BytesToRequest(codec codec.Manager, requestBytes []byte) (Request, error) {
+    var request Request
+    if _, err := codec.Unmarshal(requestBytes, &request); err != nil {
+        return nil, err
+    }
+    return request, nil
+}
+
+// RequestToBytes marshals the given request object into bytes
+func RequestToBytes(codec codec.Manager, request Request) ([]byte, error) {
+    return codec.Marshal(Version, &request)
+}
diff --git a/plugin/evm/network.go b/plugin/evm/network.go
index e3690d62..d5c9d60a 100644
--- a/plugin/evm/network.go
+++ b/plugin/evm/network.go
@@ -9,23 +9,20 @@ import (
    "sync"
    "time"

-    "github.com/ava-labs/avalanchego/cache"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/snow"
-    "github.com/ava-labs/avalanchego/utils/wrappers"
-
-    commonEng "github.com/ava-labs/avalanchego/snow/engine/common"
-
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/log"
-    "github.com/ethereum/go-ethereum/rlp"
-
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/plugin/evm/message"
-
-    coreth "github.com/ava-labs/coreth/chain"
+    coreth "github.com/ava-labs/coreth/chain"
+    coreth "github.com/flare-foundation/coreth/chain"
+    "github.com/flare-foundation/coreth/core"

```



```

+      "github.com/flare-foundation/coreth/core/state"
+      "github.com/flare-foundation/coreth/core/types"
+      "github.com/flare-foundation/coreth/plugin/evm/message"
+      "github.com/flare-foundation/flare/cache"
+      "github.com/flare-foundation/flare/ids"
+      "github.com/flare-foundation/flare/snow"
+      commonEng "github.com/flare-foundation/flare/snow/engine/common"
+      "github.com/flare-foundation/flare/utls/wrappers"
+    )
+
+    const (
diff --git a/plugin/evm/network_eth_gossiping_test.go b/plugin/evm/network_eth_gossiping_test.go
index 54a2011e..0f1559dc 100644
--- a/plugin/evm/network_eth_gossiping_test.go
+++ b/plugin/evm/network_eth_gossiping_test.go
@@ -12,18 +12,17 @@ @@ -12,18 +12,17 @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/avalanchego/ids"
+    "github.com/stretchchr/testify/assert"

    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/crypto"
    "github.com/ethereum/go-ethereum/rlp"

-    "github.com/stretchchr/testify/assert"
-
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/params"
-    "github.com/ava-labs/coreth/plugin/evm/message"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+    "github.com/flare-foundation/flare/ids"
+  )

  func fundAddressByGenesis(addr []common.Address) (string, error) {
@@ -41,11 +41,11 @@ @@ -41,11 +40,12 @@ func fundAddressByGenesis(addr []common.Address) (string, error) {
    genesis.Alloc = funds

    genesis.Config = &params.ChainConfig{
-      ChainID:      params.AvalancheLocalChainID,
+      ChainID:      big.NewInt(31337),
      ApricotPhase1BlockTimestamp: big.NewInt(0),
      ApricotPhase2BlockTimestamp: big.NewInt(0),
      ApricotPhase3BlockTimestamp: big.NewInt(0),
      ApricotPhase4BlockTimestamp: big.NewInt(0),
+      ApricotPhase5BlockTimestamp: big.NewInt(0),
    }

    bytes, err := json.Marshal(genesis)
diff --git a/plugin/evm/service.go b/plugin/evm/service.go
index 0dccc7336..435dec91 100644
--- a/plugin/evm/service.go
+++ b/plugin/evm/service.go
@@ -11,17 +11,16 @@ @@ -11,17 +11,16 @@ import (
    "net/http"
    "strings"

-    "github.com/ava-labs/avalanchego/api"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/utls/constants"
-    "github.com/ava-labs/avalanchego/utls/crypto"
-    "github.com/ava-labs/avalanchego/utls/formatting"
-    "github.com/ava-labs/avalanchego/utls/json"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
-    ethcrypto "github.com/ethereum/go-ethereum/crypto"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/flare/api"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utls/constants"
+    "github.com/flare-foundation/flare/utls/crypto"
+    "github.com/flare-foundation/flare/utls/formatting"
+    "github.com/flare-foundation/flare/utls/json"
+  )

  // test constants
@@ -41,27 +41,26 @@ @@ -41,27 +40,6 @@ var (
    initialBaseFee = big.NewInt(params.ApricotPhase3InitialBaseFee)
  )

-// NetAPI offers network related API methods
-type NetAPI struct{ vm *VM }
-
-// Listening returns an indication if the node is listening for network connections.
-func (s *NetAPI) Listening() bool { return true } // always listening
-
-// PeerCount returns the number of connected peers
-func (s *NetAPI) PeerCount() hexutil.Uint { return hexutil.Uint(0) }
-
-// Version returns the current ethereum protocol version.
-func (s *NetAPI) Version() string { return fmt.Sprintf("%d", s.vm.networkID) }
-
-// Web3API offers helper API methods
-type Web3API struct{}
-
-// ClientVersion returns the version of the vm running
-func (s *Web3API) ClientVersion() string { return Version }
-
-// Sha3 returns the bytes returned by hashing [input] with Keccak256
-func (s *Web3API) Sha3(input hexutil.Bytes) hexutil.Bytes { return ethcrypto.Keccak256(input) }
-
-// SnowmanAPI introduces snowman specific functionality to the evm
-type SnowmanAPI struct{ vm *VM }

diff --git a/plugin/evm/fuji_ext_data_hashes.json b/plugin/evm/songbird_ext_data_hashes.json
similarity index 100%
rename from plugin/evm/fuji_ext_data_hashes.json
rename to plugin/evm/songbird_ext_data_hashes.json
diff --git a/plugin/evm/static_service.go b/plugin/evm/static_service.go
index 9c592251..e1db747e 100644
--- a/plugin/evm/static_service.go
+++ b/plugin/evm/static_service.go
@@ -7,8 +7,8 @@ @@ -7,8 +7,8 @@ import (
    "context"
    "encoding/json"

-    "github.com/ava-labs/avalanchego/utls/formatting"
-    "github.com/ava-labs/coreth/core"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/flare/utls/formatting"
+  )

  // StaticService defines the static API services exposed by the evm
diff --git a/plugin/evm/test_tx.go b/plugin/evm/test_tx.go
new file mode 100644
index 00000000..878610d4
--- /dev/null

```

```

+++ b/plugin/evm/test_tx.go
@@ -0,0 +1,152 @@
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.
+// See the file LICENSE for licensing terms.
+
+package evm
+
+import (
+    "math/big"
+    "math/rand"
+
+    "github.com/flare-foundation/flare/utls"
+
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/codec/linearcodec"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/snow"
+    "github.com/flare-foundation/flare/utls/wrappers"
+)
+
+type TestTx struct {
+    GasUsedV          uint64          `serialize:"true"`
+    AcceptRequestsBlockchainIDV ids.ID      `serialize:"true"`
+    AcceptRequestsV    *atomic.Requests `serialize:"true"`
+    VerifyV            error
+    IDV                ids.ID      `serialize:"true" json:"id"`
+    BurnedV            uint64      `serialize:"true"`
+    UnsignedBytesV     []byte
+    BytesV             []byte
+    InputUTXOsV        ids.Set
+    SemanticVerifyV    error
+    EVMStateTransferV  error
+}
+
+var _ UnsignedAtomicTx = &TestTx{}
+
+// GasUsed implements the UnsignedAtomicTx interface
+func (t *TestTx) GasUsed(fixedFee bool) (uint64, error) { return t.GasUsedV, nil }
+
+// Verify implements the UnsignedAtomicTx interface
+func (t *TestTx) Verify(ctx *snow.Context, rules params.Rules) error { return t.VerifyV }
+
+// AtomicOps implements the UnsignedAtomicTx interface
+func (t *TestTx) AtomicOps() (ids.ID, *atomic.Requests, error) {
+    return t.AcceptRequestsBlockchainIDV, t.AcceptRequestsV, nil
+}
+
+// Initialize implements the UnsignedAtomicTx interface
+func (t *TestTx) Initialize(unsignedBytes, signedBytes []byte) {}
+
+// ID implements the UnsignedAtomicTx interface
+func (t *TestTx) ID() ids.ID { return t.IDV }
+
+// Burned implements the UnsignedAtomicTx interface
+func (t *TestTx) Burned(assetID ids.ID) (uint64, error) { return t.BurnedV, nil }
+
+// UnsignedBytes implements the UnsignedAtomicTx interface
+func (t *TestTx) UnsignedBytes() []byte { return t.UnsignedBytesV }
+
+// Bytes implements the UnsignedAtomicTx interface
+func (t *TestTx) Bytes() []byte { return t.BytesV }
+
+// InputUTXOs implements the UnsignedAtomicTx interface
+func (t *TestTx) InputUTXOs() ids.Set { return t.InputUTXOsV }
+
+// SemanticVerify implements the UnsignedAtomicTx interface
+func (t *TestTx) SemanticVerify(vm *VM, stx *Tx, parent *Block, baseFee *big.Int, rules params.Rules) error {
+    return t.SemanticVerifyV
+}
+
+// EVMStateTransfer implements the UnsignedAtomicTx interface
+func (t *TestTx) EVMStateTransfer(ctx *snow.Context, state *state.StateDB) error {
+    return t.EVMStateTransferV
+}
+
+func testTxCodec() codec.Manager {
+    codec := codec.NewDefaultManager()
+    c := linearcodec.NewDefault()
+
+    errs := wrappers.Errs{}
+    errs.Add(
+        c.RegisterType(&TestTx{}),
+        c.RegisterType(&atomic.Element{}),
+        c.RegisterType(&atomic.Requests{}),
+        codec.RegisterCodec(codecVersion, c),
+    )
+
+    if errs.Errorred() {
+        panic(errs.Err)
+    }
+
+    return codec
+}
+
+var blockChainID = ids.GenerateTestID()
+
+func testDataImportTx() *Tx {
+    return &Tx{
+        UnsignedAtomicTx: &TestTx{
+            IDV:                ids.GenerateTestID(),
+            AcceptRequestsBlockchainIDV: blockChainID,
+            AcceptRequestsV: &atomic.Requests{
+                RemoveRequests: [][]byte{
+                    utls.RandomBytes(32),
+                    utls.RandomBytes(32),
+                },
+            },
+        },
+    }
+}
+
+func testDataExportTx() *Tx {
+    return &Tx{
+        UnsignedAtomicTx: &TestTx{
+            IDV:                ids.GenerateTestID(),
+            AcceptRequestsBlockchainIDV: blockChainID,
+            AcceptRequestsV: &atomic.Requests{
+                PutRequests: []*atomic.Element{
+                    {
+                        Key:    utls.RandomBytes(16),
+                        Value: utls.RandomBytes(24),
+                        Traits: [][]byte{
+                            utls.RandomBytes(32),
+                            utls.RandomBytes(32),
+                        },
+                    },
+                },
+            },
+        },
+    }
+}

```

```

+func newTestTx() *Tx {
+    txType := rand.Intn(2)
+    switch txType {
+    case 0:
+        return testDataImportTx()
+    case 1:
+        return testDataExportTx()
+    default:
+        panic("rng generated unexpected value for tx type")
+    }
+}
+
+func newTestTxs(numTxs int) []*Tx {
+    txs := make([]*Tx, 0, numTxs)
+    for i := 0; i < numTxs; i++ {
+        txs = append(txs, newTestTx())
+    }
+
+    return txs
+}
diff --git a/plugin/evm/tx.go b/plugin/evm/tx.go
index 3f0ad923..4f080f57 100644
--- a/plugin/evm/tx.go
+++ b/plugin/evm/tx.go
@@ -12,19 +12,19 @@ import (

    "github.com/ethereum/go-ethereum/common"

-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/params"
-
-    "github.com/ava-labs/avalanchego/codec"
-    "github.com/ava-labs/avalanchego/database"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/snow"
-    "github.com/ava-labs/avalanchego/utills"
-    "github.com/ava-labs/avalanchego/utills/crypto"
-    "github.com/ava-labs/avalanchego/utills/hashing"
-    "github.com/ava-labs/avalanchego/utills/wrappers"
-    "github.com/ava-labs/avalanchego/vms/components/verify"
-    "github.com/ava-labs/avalanchego/vms/secp256k1fx"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/params"
+
+    "github.com/flare-foundation/flare/chains/atomic"
+    "github.com/flare-foundation/flare/codec"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/snow"
+    "github.com/flare-foundation/flare/utills"
+    "github.com/flare-foundation/flare/utills/crypto"
+    "github.com/flare-foundation/flare/utills/hashing"
+    "github.com/flare-foundation/flare/utills/wrappers"
+    "github.com/flare-foundation/flare/vms/components/verify"
+    "github.com/flare-foundation/flare/vms/secp256k1fx"
)

var (
@@ -92,7 +92,7 @@ func (in *EVMInput) Verify() error {
type UnsignedTx interface {
    Initialize(unsignedBytes, signedBytes []byte)
    ID() ids.ID
-    GasUsed() (uint64, error)
+    GasUsed(fixedFee bool) (uint64, error)
    Burned(assetID ids.ID) (uint64, error)
    UnsignedBytes() []byte
    Bytes() []byte
@@ -102,16 +102,16 @@ type UnsignedAtomicTx interface {
type UnsignedTx interface {

-    // UTXOs this tx consumes
+    // InputUTXOs returns the UTXOs this tx consumes
    InputUTXOs() ids.Set
    // Verify attempts to verify that the transaction is well formed
-    // TODO: remove [xChainID] parameter since this is provided on [ctx]
-    Verify(xChainID ids.ID, ctx *snow.Context, rules params.Rules) error
+    Verify(ctx *snow.Context, rules params.Rules) error
    // Attempts to verify this transaction with the provided state.
    SemanticVerify(vm *VM, stx *Tx, parent *Block, baseFee *big.Int, rules params.Rules) error

-    // Accept this transaction with the additionally provided state transitions.
-    Accept(ctx *snow.Context, batch database.Batch) error
+    // AtomicOps returns the blockchainID and set of atomic requests that
+    // must be applied to shared memory for this transaction to be accepted.
+    // The set of atomic requests must be returned in a consistent order.
    AtomicOps() (ids.ID, *atomic.Requests, error)

    EVMStateTransfer(ctx *snow.Context, state *state.StateDB) error
}
@@ -160,14 +160,14 @@ func (tx *Tx) Sign(c codec.Manager, signers [][]*crypto.PrivateKeySECP256K1R) er
// for via this transaction denominated in [avaxAssetID] with [baseFee] used to calculate the
// cost of this transaction. This function also returns the [gasUsed] by the
// transaction for inclusion in the [baseFee] algorithm.
-func (tx *Tx) BlockFeeContribution(avaxAssetID ids.ID, baseFee *big.Int) (*big.Int, *big.Int, error) {
+func (tx *Tx) BlockFeeContribution(fixedFee bool, avaxAssetID ids.ID, baseFee *big.Int) (*big.Int, *big.Int, error) {
    if baseFee == nil {
        return nil, nil, errNilBaseFee
    }
    if baseFee.Cmp(common.Big0) <= 0 {
        return nil, nil, fmt.Errorf("cannot calculate tip with base fee %d <= 0", baseFee)
    }
-    gasUsed, err := tx.GasUsed()
+    gasUsed, err := tx.GasUsed(fixedFee)
    if err != nil {
        return nil, nil, err
    }
@@ -279,3 +279,30 @@ func calculateDynamicFee(cost uint64, baseFee *big.Int) (uint64, error) {
    func calcBytesCost(len int) uint64 {
        return uint64(len) * TxBytesGas
    }
+
+    // mergeAtomicOps merges atomic requests represented by [txs]
+    // to the [output] map, depending on whether [chainID] is present in the map.
+    func mergeAtomicOps(txs []*Tx) (map[ids.ID]*atomic.Requests, error) {
+        if len(txs) > 1 {
+            // txs should be stored in order of txID to ensure consistency
+            // with txs initialized from the txID index.
+            copyTxs := make([]*Tx, len(txs))
+            copy(copyTxs, txs)
+            sort.Slice(copyTxs, func(i, j int) bool { return copyTxs[i].ID().Hex() < copyTxs[j].ID().Hex() })
+            txs = copyTxs
+        }
+        output := make(map[ids.ID]*atomic.Requests)
+        for _, tx := range txs {
+            chainID, txRequest, err := tx.UnsignedAtomicTx.AtomicOps()
+            if err != nil {
+                return nil, err
+            }
+            if request, exists := output[chainID]; exists {
+                request.PutRequests = append(request.PutRequests, txRequest.PutRequests...)
+                request.RemoveRequests = append(request.RemoveRequests, txRequest.RemoveRequests...)
+            } else {
+                output[chainID] = txRequest
+            }
+        }
    }
}

```

```
}  
    return output, nil  
+}  
+}  
diff --git a/plugin/evm/tx_heap.go b/plugin/evm/tx_heap.go  
index 11c1319c..d5731d1b 100644  
--- a/plugin/evm/tx_heap.go  
+++ b/plugin/evm/tx_heap.go  
@@ -1,9 +1,12 @@  
+// (c) 2020-2021, Ava Labs, Inc. All rights reserved.  
+// See the file LICENSE for licensing terms.  
+  
package evm  
  
import (  
    "container/heap"  
  
    "github.com/ava-labs/avalanchego/ids"  
    "github.com/flare-foundation/flare/ids"  
)  
  
// txEntry is used to track the [gasPrice] transactions pay to be included in  
diff --git a/plugin/evm/tx_test.go b/plugin/evm/tx_test.go  
index e0482e2d..6833e567 100644  
--- a/plugin/evm/tx_test.go  
+++ b/plugin/evm/tx_test.go  
@@ -8,9 +8,9 @@ import (  
    "strings"  
    "testing"  
  
    "github.com/ava-labs/avalanchego/chains/atomic"  
    "github.com/ava-labs/avalanchego/snow"  
    "github.com/ava-labs/coreth/params"  
    "github.com/flare-foundation/coreth/params"  
    "github.com/flare-foundation/flare/chains/atomic"  
    "github.com/flare-foundation/flare/snow"  
)  
  
func TestCalculateDynamicFee(t *testing.T) {  
@@ -49,131 +49,3 @@ func TestCalculateDynamicFee(t *testing.T) {  
    }  
}  
-  
-type atomicTxVerifyTest struct {  
-    ctx          *snow.Context  
-    generate      func(t *testing.T) UnsignedAtomicTx  
-    rules         params.Rules  
-    expectedErr   string  
-}  
-  
-// executeTxVerifyTest tests  
-func executeTxVerifyTest(t *testing.T, test atomicTxVerifyTest) {  
-    atomicTx := test.generate(t)  
-    err := atomicTx.Verify(test.ctx.XChainID, test.ctx, test.rules)  
-    if len(test.expectedErr) == 0 {  
-        if err != nil {  
-            t.Fatalf("Atomic tx failed unexpectedly due to: %s", err)  
-        }  
-    } else {  
-        if err == nil {  
-            t.Fatalf("Expected atomic tx test to fail due to: %s, but passed verification", test.expectedErr)  
-        }  
-        if !strings.Contains(err.Error(), test.expectedErr) {  
-            t.Fatalf("Expected Verify to fail due to %s, but failed with: %s", test.expectedErr, err)  
-        }  
-    }  
-}  
-  
-type atomicTxTest struct {  
-    // setup returns the atomic transaction for the test  
-    setup func(t *testing.T, vm *VM, sharedMemory *atomic.Memory) *Tx  
-    // define a string that should be contained in the error message if the tx fails verification  
-    // at some point. If the strings are empty, then the tx should pass verification at the  
-    // respective step.  
-    semanticVerifyErr, evmStateTransferErr, acceptErr string  
-    // checkState is called iff building and verifying a block containing the transaction is successful. Verifies  
-    // the state of the VM following the block's acceptance.  
-    checkState func(t *testing.T, vm *VM)  
-  
-    // Whether or not the VM should be considered to still be bootstrapping  
-    bootstrapping bool  
-    // genesisJSON to use for the VM genesis (also defines the rule set that will be used in verification)  
-    // If this is left empty, [genesisJSONApricotPhase0], will be used  
-    genesisJSON string  
-  
-    // passed directly into GenesisVM  
-    configJSON, upgradeJSON string  
-}  
-  
-func executeTxTest(t *testing.T, test atomicTxTest) {  
-    genesisJSON := test.genesisJSON  
-    if len(genesisJSON) == 0 {  
-        genesisJSON = genesisJSONApricotPhase0  
-    }  
-    issuer, vm, _, sharedMemory, _ := GenesisVM(t, !test.bootstrapping, genesisJSON, test.configJSON, test.upgradeJSON)  
-    rules := vm.currentRules()  
-  
-    tx := test.setup(t, vm, sharedMemory)  
-  
-    var baseFee *big.Int  
-    // If ApricotPhase3 is active, use the initial base fee for the atomic transaction  
-    switch {  
-    case rules.IsApricotPhase3:  
-        baseFee = initialBaseFee  
-    }  
-  
-    lastAcceptedBlock := vm.LastAcceptedBlockInternal().(*Block)  
-    if err := tx.UnsignedAtomicTx.SemanticVerify(vm, tx, lastAcceptedBlock, baseFee, rules); len(test.semanticVerifyErr) == 0 && err != nil {  
-        t.Fatalf("SemanticVerify failed unexpectedly due to: %s", err)  
-    } else if len(test.semanticVerifyErr) != 0 {  
-        if err == nil {  
-            t.Fatalf("SemanticVerify unexpectedly returned a nil error. Expected err: %s", test.semanticVerifyErr)  
-        }  
-        if !strings.Contains(err.Error(), test.semanticVerifyErr) {  
-            t.Fatalf("Expected SemanticVerify to fail due to %s, but failed with: %s", test.semanticVerifyErr, err)  
-        }  
-    }  
-    // If SemanticVerify failed for the expected reason, return early  
-    return  
-}  
-  
-// Retrieve dummy state to test that EVMStateTransfer works correctly  
-sdb, err := vm.chain.BlockState(lastAcceptedBlock.ethBlock)  
if err != nil {  
    t.Fatal(err)  
}  
if err := tx.UnsignedAtomicTx.EVMStateTransfer(vm.ctx, sdb); len(test.evmStateTransferErr) == 0 && err != nil {  
    t.Fatalf("EVMStateTransfer failed unexpectedly due to: %s", err)  
} else if len(test.evmStateTransferErr) != 0 {  
    if err == nil {  
        t.Fatalf("EVMStateTransfer unexpectedly returned a nil error. Expected err: %s", test.evmStateTransferErr)  
    }  
    if !strings.Contains(err.Error(), test.evmStateTransferErr) {  
        t.Fatalf("Expected SemanticVerify to fail due to %s, but failed with: %s", test.evmStateTransferErr, err)
```

```

-         // If EVMStateTransfer failed for the expected reason, return early
-         return
-     }
-
-     if err := vm.issueTx(tx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-     <-issuer
-
-     // If we've reached this point, we expect to be able to build and verify the block without any errors
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); len(test.acceptErr) == 0 && err != nil {
-         t.Fatalf("Accept failed unexpectedly due to: %s", err)
-     } else if len(test.acceptErr) != 0 {
-         if err == nil {
-             t.Fatalf("Accept unexpectedly returned a nil error. Expected err: %s", test.acceptErr)
-         }
-         if !strings.Contains(err.Error(), test.acceptErr) {
-             t.Fatalf("Expected Accept to fail due to %s, but failed with: %s", test.acceptErr, err)
-         }
-     }
-     // If Accept failed for the expected reason, return early
-     return
- }
-
- if test.checkState != nil {
-     test.checkState(t, vm)
- }
- }
-}
diff --git a/plugin/evm/user.go b/plugin/evm/user.go
index e1902598..fealbe4e 100644
--- a/plugin/evm/user.go
+++ b/plugin/evm/user.go
@@ -7,10 +7,10 @@ import (
    "errors"
    "fmt"

-    "github.com/ava-labs/avalanchego/database/encdb"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/utlis/crypto"
+    "github.com/flare-foundation/flare/database/encdb"
+    "github.com/flare-foundation/flare/ids"
+    "github.com/flare-foundation/flare/utlis/crypto"
)

// Key in the database whose corresponding value is the list of
diff --git a/plugin/evm/version.go b/plugin/evm/version.go
index 9028c4de..541dac76 100644
--- a/plugin/evm/version.go
+++ b/plugin/evm/version.go
@@ -11,7 +11,7 @@ var (
    // GitCommit is set by the build script
    GitCommit string
    // Version is the version of Coreth
-    Version string
+    Version string = "v0.5.1"
)

func init() {
diff --git a/plugin/evm/vm.go b/plugin/evm/vm.go
index 66b98eb3..67396386 100644
--- a/plugin/evm/vm.go
+++ b/plugin/evm/vm.go
@@ -12,53 +12,53 @@ import (
    "math/big"
    "os"
    "path/filepath"
+    "sort"
    "strings"
    "sync"
    "time"

-    "github.com/ava-labs/avalanchego/database/versiondb"
-    coreth "github.com/ava-labs/coreth/chain"
-    "github.com/ava-labs/coreth/consensus/dummy"
-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/eth/ethconfig"
-    "github.com/ava-labs/coreth/node"
-    "github.com/ava-labs/coreth/params"
-
-    "github.com/ava-labs/coreth/rpc"
+    "github.com/flare-foundation/coreth/plugin/evm/message"
+
+    coreth "github.com/flare-foundation/coreth/chain"
+    "github.com/flare-foundation/coreth/consensus/dummy"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/eth/ethconfig"
+    "github.com/flare-foundation/coreth/metrics/prometheus"
+    "github.com/flare-foundation/coreth/node"
+    "github.com/flare-foundation/coreth/params"
+    "github.com/flare-foundation/coreth/peer"
+    "github.com/flare-foundation/coreth/rpc"
+
+    // Force-load tracer engine to trigger registration
+    //
+    // We must import this package (not referenced elsewhere) so that the native "callTracer"
+    // is added to a map of client-accessible tracers. In geth, this is done
+    // inside of cmd/geth.
+    _ "github.com/flare-foundation/coreth/eth/tracers/js"
+    _ "github.com/flare-foundation/coreth/eth/tracers/native"
+
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/metrics"
+    "github.com/ethereum/go-ethereum/rpc"
)

avalancheRPC "github.com/gorilla/rpc/v2"

-    "github.com/ava-labs/avalanchego/cache"
-    "github.com/ava-labs/avalanchego/codec"
-    "github.com/ava-labs/avalanchego/codec/linearcodec"
-    "github.com/ava-labs/avalanchego/database"
-    "github.com/ava-labs/avalanchego/database/manager"
-    "github.com/ava-labs/avalanchego/database/prefixdb"
-    "github.com/ava-labs/avalanchego/ids"
-    "github.com/ava-labs/avalanchego/snow"
-    "github.com/ava-labs/avalanchego/snow/choices"
-    "github.com/ava-labs/avalanchego/snow/consensus/snowman"
-    "github.com/ava-labs/avalanchego/snow/engine/snowman/block"
-    "github.com/ava-labs/avalanchego/utlis/constants"
-    "github.com/ava-labs/avalanchego/utlis/crypto"
-    "github.com/ava-labs/avalanchego/utlis/formatting"

```

```

- "github.com/ava-labs/avalanchego/utls/logging"
- "github.com/ava-labs/avalanchego/utls/math"
- "github.com/ava-labs/avalanchego/utls/profiler"
- "github.com/ava-labs/avalanchego/utls/timer/mockable"
- "github.com/ava-labs/avalanchego/utls/wrappers"
- "github.com/ava-labs/avalanchego/vms/components/avax"
- "github.com/ava-labs/avalanchego/vms/components/chain"
- "github.com/ava-labs/avalanchego/vms/secp256k1fx"
-
- commonEng "github.com/ava-labs/avalanchego/snow/engine/common"
-
- avalancheJSON "github.com/ava-labs/avalanchego/utls/json"
+ "github.com/flare-foundation/flare/cache"
+ "github.com/flare-foundation/flare/codec"
+ "github.com/flare-foundation/flare/codec/linearcodec"
+ "github.com/flare-foundation/flare/database"
+ "github.com/flare-foundation/flare/database/manager"
+ "github.com/flare-foundation/flare/database/prefixdb"
+ "github.com/flare-foundation/flare/database/versiondb"
+ "github.com/flare-foundation/flare/ids"
+ "github.com/flare-foundation/flare/snow"
+ "github.com/flare-foundation/flare/snow/choices"
+ "github.com/flare-foundation/flare/snow/consensus/snowman"
+ "github.com/flare-foundation/flare/snow/engine/snowman/block"
+ "github.com/flare-foundation/flare/utls/constants"
+ "github.com/flare-foundation/flare/utls/crypto"
+ "github.com/flare-foundation/flare/utls/formatting"
+ "github.com/flare-foundation/flare/utls/logging"
+ "github.com/flare-foundation/flare/utls/math"
+ "github.com/flare-foundation/flare/utls/perms"
+ "github.com/flare-foundation/flare/utls/profiler"
+ "github.com/flare-foundation/flare/utls/timer/mockable"
+ "github.com/flare-foundation/flare/vms/components/avax"
+ "github.com/flare-foundation/flare/vms/components/chain"
+ "github.com/flare-foundation/flare/vms/secp256k1fx"
+
+ commonEng "github.com/flare-foundation/flare/snow/engine/common"
+
+ avalancheJSON "github.com/flare-foundation/flare/utls/json"
)

const (
@@ -74,7 +88,8 @@ var (
    x2cRate      = big.NewInt(x2cRateInt64)
    x2cRateMinus1 = big.NewInt(x2cRateMinus1Int64)

-    _ block.ChainVM = &VM{}
+    _ block.ChainVM      = &VM{}
+    _ block.HeightIndexedChainVM = &VM{}
)

const (
@@ -101,10 +116,14 @@ const (

var (
    // Set last accepted key to be longer than the keys used to store accepted block IDs.
-    lastAcceptedKey      = []byte("last_accepted_key")
-    acceptedPrefix       = []byte("snowman_accepted")
-    ethDBPrefix          = []byte("ethdb")
-    atomicTxPrefix       = []byte("atomicTxDB")
+    lastAcceptedKey = []byte("last_accepted_key")
+    acceptedPrefix  = []byte("snowman_accepted")
+    ethDBPrefix     = []byte("ethdb")
+
+    // Prefixes for atomic trie
+    atomicTrieDBPrefix      = []byte("atomicTrieDB")
+    atomicTrieMetaDBPrefix = []byte("atomicTrieMetaDB")
+
+    pruneRejectedBlocksKey = []byte("pruned_rejected_blocks")
)

@@ -142,7 +161,7 @@ var (
    errNilBlockGasCostApricotPhase4 = errors.New("nil blockGasCost is invalid after apricotPhase4")
    errConflictingAtomicTx           = errors.New("conflicting atomic tx present")
    errTooManyAtomicTx               = errors.New("too many atomic tx")
-    defaultLogLevel                 = log.LvlDebug
+    errMissingAtomicTxs             = errors.New("cannot build a block with non-empty extra data and zero atomic transactions")
)

var originalStderr *os.File
@@ -176,12 +195,18 @@ type VM struct {
    // [acceptedBlockDB] is the database to store the last accepted
    // block.
    acceptedBlockDB database.Database
    // [acceptedAtomicTxDB] maintains an index of accepted atomic txs.
-    acceptedAtomicTxDB database.Database
+
+    // [atomicTxRepository] maintains two indexes on accepted atomic txs.
+    // - txID to accepted atomic tx
+    // - block height to list of atomic txs accepted on block at that height
+    atomicTxRepository AtomicTxRepository
+    // [atomicTrie] maintains a merkle forest of [height]=>[atomic txs].
+    // Used to state sync clients.
+    atomicTrie AtomicTrie

    builder *blockBuilder

-    network Network
+    gossip Gossiper

    baseCodec codec.Registry
    codec     codec.Manager
@@ -196,14 +221,12 @@ type VM struct {

    // Continuous Profiler
    profiler profiler.ContinuousProfiler
-}

-func (vm *VM) Connected(nodeID ids.ShortID) error {
-    return nil // noop
-}
+    peer.Network
+    client      peer.Client
+    networkCodec codec.Manager

-func (vm *VM) Disconnected(nodeID ids.ShortID) error {
-    return nil // noop
+    bootstrapped bool
+}

// Codec implements the secp256k1fx interface
@@ -218,9 +241,14 @@ func (vm *VM) Clock() *mockable.Clock { return &vm.clock }
// Logger implements the secp256k1fx interface
func (vm *VM) Logger() logging.Logger { return vm.ctx.Log }

-// SetLogLevel sets the log level with the original [os.Stderr] interface
+// setLogLevel sets the log level with the original [os.Stderr] interface along
+// with the context logger.
func (vm *VM) setLogLevel(logLevel log.Lvl) {
-    log.Root().SetHandler(log.LvlFilterHandler(logLevel, log.StreamHandler(originalStderr, log.TerminalFormat(false))))
+    format := log.TerminalFormat(false)
+    log.Root().SetHandler(log.LvlFilterHandler(logLevel, log.MultiHandler(
+        log.StreamHandler(originalStderr, format),
+        log.StreamHandler(vm.ctx.Log, format),
    )

```

```

+     )))
+ }

/*
@@ -235,7 +263,6 @@ func (vm *VM) GetActivationTime() time.Time {
+ }

// Initialize implements the snowman.ChainVM interface
-
- func (vm *VM) Initialize(
+     ctx *snow.Context,
+     dbManager manager.Manager,
@@ -263,6 +290,9 @@ func (vm *VM) Initialize(
+     return errUnsupportedFXs
+ }

+     metrics.Enabled = vm.config.MetricsEnabled
+     metrics.EnabledExpensive = vm.config.MetricsExpensiveEnabled
+
+     vm.shutdownChan = make(chan struct{}, 1)
+     vm.ctx = ctx
+     baseDB := dbManager.Current().Database
@@ -271,42 +301,40 @@ func (vm *VM) Initialize(
+     vm.chaindb = Database(prefixdb.NewNested(ethDBPrefix, baseDB))
+     vm.db = versiondb.New(baseDB)
+     vm.acceptedBlockDB = prefixdb.New(acceptedPrefix, vm.db)
-     vm.acceptedAtomicTxDB = prefixdb.New(atomicTxPrefix, vm.db)
+     g := new(core.Genesis)
+     if err := json.Unmarshal(genesisBytes, g); err != nil {
+         return err
+     }

-     // Set the chain config for mainnet/fuji chain IDs
+     // Set the hard-coded chain config for reference network chain IDs
+     switch {
-     case g.Config.ChainID.Cmp(params.AvalancheMainnetChainID) == 0:
+         g.Config = params.AvalancheMainnetChainConfig
+         phase0BlockValidator.extDataHashes = mainnetExtDataHashes
-     case g.Config.ChainID.Cmp(params.AvalancheFujiChainID) == 0:
+         g.Config = params.AvalancheFujiChainConfig
+         phase0BlockValidator.extDataHashes = fujiExtDataHashes
-     case g.Config.ChainID.Cmp(params.AvalancheLocalChainID) == 0:
+         g.Config = params.AvalancheLocalChainConfig
-     }

-     // Free the memory of the extDataHash map that is not used (i.e. if mainnet
-     // config, free fuji)
-     fujiExtDataHashes = nil
-     mainnetExtDataHashes = nil
+     case g.Config.ChainID.Cmp(params.FlareChainID) == 0:
+         g.Config = params.FlareChainConfig
+         phase0BlockValidator.extDataHashes = flareExtDataHashes
+     case g.Config.ChainID.Cmp(params.CostonChainID) == 0:
+         g.Config = params.CostonChainConfig
+     case g.Config.ChainID.Cmp(params.SongbirdChainID) == 0:
+         g.Config = params.SongbirdChainConfig
+         phase0BlockValidator.extDataHashes = songbirdExtDataHashes
+     case g.Config.ChainID.Cmp(params.LocalChainID) == 0:
+         g.Config = params.LocalChainConfig
+     }

+     // Free the memory of the extDataHash map that is not used (i.e. if flare
+     // config, free songbird)
+     songbirdExtDataHashes = nil
+     flareExtDataHashes = nil

+     vm.chainID = g.Config.ChainID

+     ethConfig := ethconfig.NewDefaultConfig()
+     ethConfig.Genesis = g
+     ethConfig.NetworkID = vm.chainID.Uint64()

+     // Set log level
+     logLevel := defaultLogLevel
+     if vm.config.LogLevel != "" {
+         configLogLevel, err := log.LvlFromString(vm.config.LogLevel)
+         if err != nil {
+             return fmt.Errorf("failed to initialize logger due to: %w ", err)
+         }
+         logLevel = configLogLevel
+     }
+     logLevel, err := log.LvlFromString(vm.config.LogLevel)
+     if err != nil {
+         return fmt.Errorf("failed to initialize logger due to: %w ", err)
+     }

+     vm.setLogLevel(logLevel)
@@ -323,6 +351,16 @@ func (vm *VM) Initialize(
+     ethConfig.Pruning = vm.config.Pruning
+     ethConfig.SnapshotAsync = vm.config.SnapshotAsync
+     ethConfig.SnapshotVerify = vm.config.SnapshotVerify
+     ethConfig.OfflinePruning = vm.config.OfflinePruning
+     ethConfig.OfflinePruningBloomFilterSize = vm.config.OfflinePruningBloomFilterSize
+     ethConfig.OfflinePruningDataDirectory = vm.config.OfflinePruningDataDirectory
+
+     if len(ethConfig.OfflinePruningDataDirectory) != 0 {
+         if err := os.MkdirAll(ethConfig.OfflinePruningDataDirectory, perms.ReadWriteExecute); err != nil {
+             log.Error("failed to create offline pruning data directory", "error", err)
+             return err
+         }
+     }

+     vm.chainConfig = g.Config
+     vm.networkID = ethConfig.NetworkID
@@ -346,7 +384,7 @@ func (vm *VM) Initialize(
+     var lastAcceptedHash common.Hash
+     switch {
-     case lastAcceptedErr == database.ErrNotFound:
+         // // Set [lastAcceptedHash] to the genesis block hash.
+         // Set [lastAcceptedHash] to the genesis block hash.
+         lastAcceptedHash = ethConfig.Genesis.ToBlock(nil).Hash()
-     case lastAcceptedErr != nil:
+         return fmt.Errorf("failed to get last accepted block ID due to: %w", lastAcceptedErr)
@@ -355,20 +393,42 @@ func (vm *VM) Initialize(
+     default:
+         lastAcceptedHash = common.BytesToHash(lastAcceptedBytes)
+     }
+     ethChain, err := coreth.NewETHChain(&ethConfig, &nodecfg, vm.chaindb, vm.config.EthBackendSettings(), vm.createConsensusCallbacks(), lastAcceptedHash)
+     ethChain, err := coreth.NewETHChain(&ethConfig, &nodecfg, vm.chaindb, vm.config.EthBackendSettings(), vm.createConsensusCallbacks(), lastAcceptedHash, &vm.clock)
+     if err != nil {
+         return err
+     }
+     vm.chain = ethChain
+     lastAccepted := vm.chain.LastAcceptedBlock()

+     vm.atomicTxRepository, err = NewAtomicTxRepository(vm.db, vm.codec, lastAccepted.NumberU64())
+     if err != nil {
+         return fmt.Errorf("failed to create atomic repository: %w", err)
+     }

+     bonusBlockHeights := make(map[uint64]ids.ID)
+     if vm.chainID.Cmp(params.AvalancheMainnetChainID) == 0 {
+         bonusBlockHeights = bonusBlockMainnetHeights
+     }
+     if err := vm.repairAtomicRepositoryForBonusBlockTxs(getAtomicRepositoryRepairHeights(vm.chainID), vm.getAtomicTxFromPreApricot5BlockByHeight); err != nil {

```

```

+         return fmt.Errorf("failed to repair atomic repository: %w", err)
+     }
+     vm.atomicTrie, err = NewAtomicTrie(vm.db, bonusBlockHeights, vm.atomicTxRepository, vm.codec, lastAccepted.NumberU64())
+     if err != nil {
+         return fmt.Errorf("failed to create atomic trie: %w", err)
+     }
+
+     // start goroutines to update the tx pool gas minimum gas price when upgrades go into effect
+     vm.handleGasPriceUpdates()
+
-     // initialize new gossip network
-     //
-     // NOTE: This network must be initialized after the atomic mempool.
-     vm.network = vm.NewNetwork(appSender)
+     vm.networkCodec, err = message.BuildCodec()
+     if err != nil {
+         return err
+     }
+
+     // initialize peer network
+     vm.Network = peer.NewNetwork(appSender, vm.networkCodec, ctx.NodeID, vm.config.MaxOutboundActiveRequests)
+     vm.client = peer.NewClient(vm.Network)
+     vm.initGossipHandling()
+
+     // start goroutines to manage block building
+     //
@@ -381,17 +441,23 @@ func (vm *VM) Initialize(
    vm.genesisHash = vm.chain.GetGenesisBlock().Hash()
    log.Info(fmt.Sprintf("lastAccepted = %s", lastAccepted.Hash().Hex()))
+
+    isApricotPhase5 := vm.chainConfig.IsApricotPhase5(new(big.Int).SetUint64(lastAccepted.Time()))
+    atomicTxs, err := ExtractAtomicTxs(lastAccepted.ExtData(), isApricotPhase5, vm.codec)
+    if err != nil {
+        return err
+    }
+    vm.State = chain.NewState(&chain.Config{
+        DecidedCacheSize:    decidedCacheSize,
+        MissingCacheSize:    missingCacheSize,
+        UnverifiedCacheSize: unverifiedCacheSize,
+        LastAcceptedBlock: &Block{
-             id:      ids.ID(lastAccepted.Hash()),
-             ethBlock: lastAccepted,
-             vm:      vm,
-             status:  choices.Accepted,
+             id:      ids.ID(lastAccepted.Hash()),
+             ethBlock: lastAccepted,
+             vm:      vm,
+             status:  choices.Accepted,
+             atomicTxs: atomicTxs,
+        },
-         GetBlockIDAtHeight: vm.getBlockIDAtHeight,
+         GetBlockIDAtHeight: vm.GetBlockIDAtHeight,
+         GetBlock:           vm.getBlock,
+         UnmarshalBlock:     vm.parseBlock,
+         BuildBlock:         vm.buildBlock,
@@ -414,9 +480,27 @@ func (vm *VM) Initialize(
    // return err
    // }
+
+    // Only provide metrics if they are being populated.
+    if metrics.Enabled {
+        gatherer := prometheus.Gatherer(metrics.DefaultRegistry)
+        if err := ctx.Metrics.Register(gatherer); err != nil {
+            return err
+        }
+    }
+
+    return vm.fx.Initialize(vm)
+
+}
+
+func (vm *VM) initGossipHandling() {
+    if vm.chainConfig.ApricotPhase4BlockTimestamp != nil {
+        vm.gossiper = vm.newPushGossiper()
+        vm.Network.SetGossipHandler(NewGossipHandler(vm))
+    } else {
+        vm.gossiper = &noopGossiper{}
+        vm.Network.SetGossipHandler(message.NoopMempoolGossipHandler{})
+    }
+}
+
+func (vm *VM) createConsensusCallbacks() *dummy.ConsensusCallbacks {
+    return &dummy.ConsensusCallbacks{
+        OnFinalizeAndAssemble: vm.onFinalizeAndAssemble,
@@ -424,7 +508,7 @@ func (vm *VM) createConsensusCallbacks() *dummy.ConsensusCallbacks {
    }
+}
+
- func (vm *VM) onFinalizeAndAssemble(header *types.Header, state *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+ func (vm *VM) preBatchOnFinalizeAndAssemble(header *types.Header, state *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+     for {
+         tx, exists := vm.mempool.NextTx()
+         if !exists {
@@ -438,7 +522,7 @@ func (vm *VM) onFinalizeAndAssemble(header *types.Header, state *state.StateDB,
            rules := vm.chainConfig.AvalancheRules(header.Number, new(big.Int).SetUint64(header.Time))
            if err := vm.verifyTx(tx, header.ParentHash, header.BaseFee, state, rules); err != nil {
                // Discard the transaction from the mempool on failed verification.
-                vm.mempool.DiscardCurrentTx()
+                vm.mempool.DiscardCurrentTx(tx.ID())
                state.RevertToSnapshot(snapshot)
                continue
            }
@@ -447,12 +531,12 @@ func (vm *VM) onFinalizeAndAssemble(header *types.Header, state *state.StateDB,
            if err != nil {
                // Discard the transaction from the mempool and error if the transaction
                // cannot be marshalled. This should never happen.
-                vm.mempool.DiscardCurrentTx()
+                vm.mempool.DiscardCurrentTx(tx.ID())
                return nil, nil, nil, fmt.Errorf("failed to marshal atomic transaction %s due to %w", tx.ID(), err)
            }
            var contribution, gasUsed *big.Int
            if rules.IsApricotPhase4 {
-                contribution, gasUsed, err = tx.BlockFeeContribution(vm.ctx.AVAXAssetID, header.BaseFee)
+                contribution, gasUsed, err = tx.BlockFeeContribution(rules.IsApricotPhase5, vm.ctx.AVAXAssetID, header.BaseFee)
            }
            if err != nil {
                return nil, nil, nil, err
            }
@@ -468,28 +552,28 @@ func (vm *VM) onFinalizeAndAssemble(header *types.Header, state *state.StateDB,
            return nil, nil, nil, nil
        }
+    }
+
+    // assumes that we are in at least Apricot Phase 5.
+    func (vm *VM) postBatchOnFinalizeAndAssemble(header *types.Header, state *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+        var (
+            batchAtomicTxs    []*Tx
+            batchAtomicUTXOs  ids.Set
+            batchContribution  *big.Int = new(big.Int).Set(common.Big0)
+            batchGasUsed      *big.Int = new(big.Int).Set(common.Big0)
+            rules             = vm.chainConfig.AvalancheRules(header.Number, new(big.Int).SetUint64(header.Time))
+        )
+
+        for {
+            tx, exists := vm.mempool.NextTx()
+            if !exists {
+                break

```



```

+     }
+
+     var (
+         txGasUsed, txContribution *big.Int
+         err                      error
+     )
+
+     // Note: we do not need to check if we are in at least ApricotPhase4 here because
+     // we assume that this function will only be called when the block is in at least
+     // ApricotPhase5.
+     txContribution, txGasUsed, err = tx.BlockFeeContribution(true, vm.ctx.AVAXAssetID, header.BaseFee)
+     if err != nil {
+         return nil, nil, nil, err
+     }
+     // ensure [gasUsed] + [batchGasUsed] doesn't exceed the [atomicGasLimit]
+     if totalGasUsed := new(big.Int).Add(batchGasUsed, txGasUsed); totalGasUsed.Cmp(params.AtomicGasLimit) > 0 {
+         // Send [tx] back to the mempool's tx heap.
+         vm.mempool.CancelCurrentTx(tx.ID())
+         break
+     }
+
+     if batchAtomicUTXOs.Overlaps(tx.InputUTXOs()) {
+         // Discard the transaction from the mempool since it will fail verification
+         // after this block has been accepted.
+         // Note: if the proposed block is not accepted, the transaction may still be
+         // valid, but we discard it early here based on the assumption that the proposed
+         // block will most likely be accepted.
+         // Discard the transaction from the mempool on failed verification.
+         vm.mempool.DiscardCurrentTx(tx.ID())
+         continue
+     }
+
+     snapshot := state.Snapshot()
+     if err := vm.verifyTx(tx, header.ParentHash, header.BaseFee, state, rules); err != nil {
+         // Discard the transaction from the mempool and reset the state to [snapshot]
+         // if it fails verification here.
+         // Note: prior to this point, we have not modified [state] so there is no need to
+         // revert to a snapshot if we discard the transaction prior to this point.
+         vm.mempool.DiscardCurrentTx(tx.ID())
+         state.RevertToSnapshot(snapshot)
+         continue
+     }
+
+     batchAtomicTxs = append(batchAtomicTxs, tx)
+     batchAtomicUTXOs.Union(tx.InputUTXOs())
+     // Add the [txGasUsed] to the [batchGasUsed] when the [tx] has passed verification
+     batchGasUsed.Add(batchGasUsed, txGasUsed)
+     batchContribution.Add(batchContribution, txContribution)
+ }
+
+ // If there is a non-zero number of transactions, marshal them and return the byte slice
+ // for the block's extra data along with the contribution and gas used.
+ if len(batchAtomicTxs) > 0 {
+     atomicTxBytes, err := vm.codec.Marshal(codecVersion, batchAtomicTxs)
+     if err != nil {
+         // If we fail to marshal the batch of atomic transactions for any reason,
+         // discard the entire set of current transactions.
+         vm.mempool.DiscardCurrentTxs()
+         return nil, nil, nil, fmt.Errorf("failed to marshal batch of atomic transactions due to %w", err)
+     }
+     return atomicTxBytes, batchContribution, batchGasUsed, nil
+ }
+
+ // If there are no regular transactions and there were also no atomic transactions to be included,
+ // then the block is empty and should be considered invalid.
+ if len(txs) == 0 {
+     // this could happen due to the async logic of geth tx pool
+     return nil, nil, nil, errEmptyBlock
+ }
+
+ // If there are no atomic transactions, but there is a non-zero number of regular transactions, then
+ // we return a nil slice with no contribution from the atomic transactions and a nil error.
+ return nil, nil, nil, nil
+}
+
+func (vm *VM) onFinalizeAndAssemble(header *types.Header, state *state.StateDB, txs []*types.Transaction) ([]byte, *big.Int, *big.Int, error) {
+    if !vm.chainConfig.IsApricotPhase5(new(big.Int).SetUint64(header.Time)) {
+        return vm.preBatchOnFinalizeAndAssemble(header, state, txs)
+    }
+    return vm.postBatchOnFinalizeAndAssemble(header, state, txs)
+}
+
+func (vm *VM) onExtraStateChange(block *types.Block, state *state.StateDB) (*big.Int, *big.Int, error) {
+    tx, err := vm.extractAtomicTx(block)
+    var (
+        batchContribution *big.Int = big.NewInt(0)
+        batchGasUsed      *big.Int = big.NewInt(0)
+        timestamp         = new(big.Int).SetUint64(block.Time())
+        isApricotPhase4   = vm.chainConfig.IsApricotPhase4(timestamp)
+        isApricotPhase5   = vm.chainConfig.IsApricotPhase5(timestamp)
+    )
+
+    txs, err := ExtractAtomicTxs(block.ExtData(), isApricotPhase5, vm.codec)
+    if err != nil {
+        return nil, nil, err
+    }
+    // If [tx] is nil, we can return nil for the extra state contribution instead of allocating
+    // a big Int for 0.
+    if tx == nil {
+
+        // If there are no transactions, we can return early
+        if len(txs) == 0 {
+            return nil, nil, nil
+        }
+
+        if err := tx.UnsignedAtomicTx.EVMStateTransfer(vm.ctx, state); err != nil {
+            return nil, nil, err
+        }
+
+        switch {
+            // If ApricotPhase4 is enabled, calculate the block fee contribution
+            case vm.chainConfig.IsApricotPhase4(new(big.Int).SetUint64(block.Time())):
+                return tx.BlockFeeContribution(vm.ctx.AVAXAssetID, block.BaseFee())
+            default:
+                // Otherwise, there is no contribution
+                return nil, nil, nil
+        }
+
+        for _, tx := range txs {
+            if err := tx.UnsignedAtomicTx.EVMStateTransfer(vm.ctx, state); err != nil {
+                return nil, nil, err
+            }
+
+            // If ApricotPhase4 is enabled, calculate the block fee contribution
+            if isApricotPhase4 {
+                contribution, gasUsed, err := tx.BlockFeeContribution(isApricotPhase5, vm.ctx.AVAXAssetID, block.BaseFee())
+                if err != nil {
+                    return nil, nil, err
+                }
+
+                batchContribution.Add(batchContribution, contribution)
+                batchGasUsed.Add(batchGasUsed, gasUsed)
+            }
+
+            // If ApricotPhase5 is enabled, enforce that the atomic gas used does not exceed the
+            // atomic gas limit.
+            if vm.chainConfig.IsApricotPhase5(timestamp) {
+                // Ensure that [tx] does not push [block] above the atomic gas limit.

```

```

+         if batchGasUsed.Cmp(params.AtomicGasLimit) == 1 {
+             return nil, nil, fmt.Errorf("atomic gas used (%d) by block (%s), exceeds atomic gas limit (%d)", batchGasUsed, block.Hash().Hex(), params.AtomicGasLimit)
+         }
+     }
+     return batchContribution, batchGasUsed, nil
+ }
+
+ func (vm *VM) pruneChain() error {
@@ -513,19 +715,20 @@ func (vm *VM) pruneChain() error {
+     if err := vm.db.Put(pruneRejectedBlocksKey, heightBytes); err != nil {
+         return err
+     }
+     return vm.db.Commit()
+ }
+
+ // Bootstrapping notifies this VM that the consensus engine is performing
+ // bootstrapping
+ func (vm *VM) Bootstrapping() error { return vm.fx.Bootstrapping() }
+
+ // Bootstrapped notifies this VM that the consensus engine has finished
+ // bootstrapping
+ func (vm *VM) Bootstrapped() error {
+     vm.ctx.Bootstrapped()
+     return vm.fx.Bootstrapped()
+ }
+ func (vm *VM) SetState(state snow.State) error {
+     switch state {
+     case snow.Bootstrapping:
+         vm.bootstrapped = false
+         return vm.fx.Bootstrapping()
+     case snow.NormalOp:
+         vm.bootstrapped = true
+         return vm.fx.Bootstrapped()
+     default:
+         return snow.ErrUnknownState
+     }
+ }
+
+ // Shutdown implements the snowman.ChainVM interface
@@ -545,15 +748,22 @@ func (vm *VM) buildBlock() (snowman.Block, error) {
+     block, err := vm.chain.GenerateBlock()
+     vm.builder.handleGenerateBlock()
+     if err != nil {
+         vm.mempool.CancelCurrentTx()
+         vm.mempool.CancelCurrentTxs()
+         return nil, err
+     }
+
+     isApricotPhase5 := vm.chainConfig.IsApricotPhase5(new(big.Int).SetUint64(block.Time()))
+     atomicTxs, err := ExtractAtomicTxs(block.ExtData(), isApricotPhase5, vm.codec)
+     if err != nil {
+         vm.mempool.DiscardCurrentTxs()
+         return nil, err
+     }
+     // Note: the status of block is set by ChainState
+     blk := &Block{
+         id:      ids.ID(block.Hash()),
+         ethBlock: block,
+         vm:      vm,
+         id:      ids.ID(block.Hash()),
+         ethBlock: block,
+         vm:      vm,
+         atomicTxs: atomicTxs,
+     }
+
+     // Verify is called on a non-wrapped block here, such that this
@@ -569,14 +779,14 @@ func (vm *VM) buildBlock() (snowman.Block, error) {
+     // to the blk state root in the triedb when we are going to call verify
+     // again from the consensus engine with writes enabled.
+     if err := blk.verify(false /*=writes*/); err != nil {
+         vm.mempool.CancelCurrentTx()
+         vm.mempool.CancelCurrentTxs()
+         return nil, fmt.Errorf("block failed verification due to: %w", err)
+     }
+
+     log.Debug(fmt.Sprintf("Built block %s", blk.ID()))
+     // Marks the current tx from the mempool as being successfully issued
+     // Marks the current transactions from the mempool as being successfully issued
+     // into a block.
+     vm.mempool.IssueCurrentTx()
+     vm.mempool.IssueCurrentTxs()
+     return blk, nil
+ }
+
@@ -586,11 +796,11 @@ func (vm *VM) parseBlock(b []byte) (snowman.Block, error) {
+     if err := rlp.DecodeBytes(b, ethBlock); err != nil {
+         return nil, err
+     }
+
+     isApricotPhase5 := vm.chainConfig.IsApricotPhase5(new(big.Int).SetUint64(ethBlock.Time()))
+     atomicTxs, err := ExtractAtomicTxs(ethBlock.ExtData(), isApricotPhase5, vm.codec)
+     if err != nil {
+         return nil, err
+     }
+     // Note: the status of block is set by ChainState
+     block := &Block{
+         id:      ids.ID(ethBlock.Hash()),
+         ethBlock: ethBlock,
+         vm:      vm,
+         id:      ids.ID(ethBlock.Hash()),
+         ethBlock: ethBlock,
+         vm:      vm,
+         atomicTxs: atomicTxs,
+     }
+
+     // Performing syntactic verification in ParseBlock allows for
+     // short-circuiting bad blocks before they are processed by the VM.
@@ -609,11 +819,11 @@ func (vm *VM) getBlock(id ids.ID) (snowman.Block, error) {
+     if ethBlock == nil {
+         return nil, database.ErrNotFound
+     }
+
+     isApricotPhase5 := vm.chainConfig.IsApricotPhase5(new(big.Int).SetUint64(ethBlock.Time()))
+     atomicTxs, err := ExtractAtomicTxs(ethBlock.ExtData(), isApricotPhase5, vm.codec)
+     if err != nil {
+         return nil, err
+     }
+     // Note: the status of block is set by ChainState
+     blk := &Block{
+         id:      ids.ID(ethBlock.Hash()),
+         ethBlock: ethBlock,
+         vm:      vm,
+         id:      ids.ID(ethBlock.Hash()),
+         ethBlock: ethBlock,
+         vm:      vm,
+         atomicTxs: atomicTxs,
+     }
+     return blk, nil
+ }
+
@@ -631,10 +841,10 @@ func (vm *VM) SetPreference(blkID ids.ID) error {
+     return vm.chain.SetPreference(block.(*Block).ethBlock)
+ }
+
+ // getBlockIDAthHeight retrieves the blkID of the canonical block at [blkHeight]
+ func (vm *VM) VerifyHeightIndex() error {

```

```
// our index is vm.chain.GetBlockByNumber
+ return nil
+
+
+// GetBlockIDAtHeight retrieves the blkID of the canonical block at [blkHeight]
+// if [blkHeight] is less than the height of the last accepted block, this will return
+// a canonical block. Otherwise, it may return a blkID that has not yet been accepted.
+func (vm *VM) GetBlockIDAtHeight(blkHeight uint64) (ids.ID, error) {
+func (vm *VM) GetBlockIDAtHeight(blkHeight uint64) (ids.ID, error) {
+    ethBlock := vm.chain.GetBlockByNumber(blkHeight)
+    if ethBlock == nil {
+        return ids.ID{}, fmt.Errorf("could not find block at height: %d", blkHeight)
+    }
+    @ -673,7 +901,9 @ func (vm *VM) CreateHandlers() (map[string]*commonEng.HTTPHandler, error) {
+    func (vm *VM) CreateHandlers() (map[string]*commonEng.HTTPHandler, error) {
+        handler := vm.chain.NewAPCHandler(vm.config.APIMaxDuration.Duration)
+        enabledAPIs := vm.config.EthAPIs()
+        vm.chain.AttachEthService(handler, enabledAPIs)
+        if err := vm.chain.AttachEthService(handler, enabledAPIs); err != nil {
+            return nil, err
+        }
+
+        primaryAlias, err := vm.ctx.BCLookup.PrimaryAlias(vm.ctx.ChainID)
+        if err != nil {
+    @ -688,7 +918,7 @ func (vm *VM) CreateHandlers() (map[string]*commonEng.HTTPHandler, error) {
+        apis[avaxEndpoint] = avaxAPI
+
+        if vm.config.CorethAdminAPIEnabled {
+            adminAPI, err := newHandler("admin", NewAdminService(vm, fmt.Sprintf("coreth_performance_%s", primaryAlias)))
+            adminAPI, err := newHandler("admin", NewAdminService(vm, os.ExpandEnv(fmt.Sprintf("%s_coreth_performance_%s", vm.config.CorethAdminAPIDir, primaryAlias))))
+            if err != nil {
+                return nil, fmt.Errorf("failed to register service for admin API due to %w", err)
+            }
+        }
+        @ -696,22 +926,12 @ func (vm *VM) CreateHandlers() (map[string]*commonEng.HTTPHandler, error) {
+        enabledAPIs = append(enabledAPIs, "coreth-admin")
+    }
+
+    - errors := wrappers.Errors{}
+    if vm.config.SnowmanAPIEnabled {
+        - errors.Add(handler.RegisterName("snowman", &SnowmanAPI{vm}))
+        + if err := handler.RegisterName("snowman", &SnowmanAPI{vm}); err != nil {
+        +         return nil, err
+        +     }
+        enabledAPIs = append(enabledAPIs, "snowman")
+    }
+
+    - if vm.config.NetAPIEnabled {
+    -     errors.Add(handler.RegisterName("net", &NetAPI{vm}))
+    -     enabledAPIs = append(enabledAPIs, "net")
+    - }
+    - if vm.config.Web3APIEnabled {
+    -     errors.Add(handler.RegisterName("web3", &Web3API{}))
+    -     enabledAPIs = append(enabledAPIs, "web3")
+    - }
+    - if errors.Errorred() {
+    -     return nil, errors.Err
+    - }
+
+    log.Info(fmt.Sprintf("Enabled APIs: %s", strings.Join(enabledAPIs, ", ")))
+    apis[ethRPCEndpoint] = &commonEng.HTTPHandler{
+    @ -740,7 +960,6 @ func (vm *VM) CreateStaticHandlers() (map[string]*commonEng.HTTPHandler, error)
+
+    return map[string]*commonEng.HTTPHandler{
+        "/rpc": {LockOptions: commonEng.NoLock, Handler: handler},
+        - "/ws": {LockOptions: commonEng.NoLock, Handler: handler.WebsocketHandler([]string{""})},
+        }, nil
+    }
+
+    @ -749,35 +968,17 @ func (vm *VM) CreateStaticHandlers() (map[string]*commonEng.HTTPHandler, error)
+    ***** Helpers *****
+    */
+    -// extractAtomicTx returns the atomic transaction in [block] if
+    -// one exists.
+    -func (vm *VM) extractAtomicTx(block *types.Block) (*Tx, error) {
+    -    extdata := block.ExtData()
+    -    if len(extdata) == 0 {
+    -        return nil, nil
+    -    }
+    -    atx := new(Tx)
+    -    if _, err := vm.codec.Unmarshal(extdata, atx); err != nil {
+    -        return nil, fmt.Errorf("failed to unmarshal atomic tx due to %w", err)
+    -    }
+    -    if err := atx.Sign(vm.codec, nil); err != nil {
+    -        return nil, fmt.Errorf("failed to initialize atomic tx in block %s", block.Hash().Hex())
+    -    }
+    -    return atx, nil
+    -}
+
+    // conflicts returns an error if [inputs] conflicts with any of the atomic inputs contained in [ancestor]
+    // or any of its ancestor blocks going back to the last accepted block in its ancestry. If [ancestor] is
+    // accepted, then nil will be returned immediately.
+    // If the ancestry of [ancestor] cannot be fetched, then [errRejectedParent] may be returned.
+    func (vm *VM) conflicts(inputs ids.Set, ancestor *Block) error {
+        for ancestor.Status() != choices.Accepted {
+            - atx, err := vm.extractAtomicTx(ancestor.ethBlock)
+            - if err != nil {
+            -         return fmt.Errorf("problem parsing atomic tx of ancestor block %s: %w", ancestor.ID(), err)
+            -     }
+            -     // If the ancestor isn't an atomic block, it can't conflict with
+            -     // the import tx.
+            -     if atx != nil {
+            -         ancestorInputs := atx.UnsignedAtomicTx.InputUTXOs()
+            -         if inputs.Overlaps(ancestorInputs) {
+            -             // If any of the atomic transactions in the ancestor conflict with [inputs]
+            -             // return an error.
+            -             for _, atomicTx := range ancestor.atomicTxes {
+            -                 if inputs.Overlaps(atomicTx.InputUTXOs()) {
+            -                     return errConflictingAtomicInputs
+            -                 }
+            -             }
+            -         }
+            -     }
+            -     return nil
+            - }
+
+            -// getAcceptedAtomicTx attempts to get [txID] from the database.
+            -func (vm *VM) getAcceptedAtomicTx(txID ids.ID) (*Tx, uint64, error) {
+            -    indexedTxBytes, err := vm.acceptedAtomicTxDB.Get(txID[:])
+            -    if err != nil {
+            -        return nil, 0, err
+            -    }
+            -    packer := wrappers.Packer(Bytes: indexedTxBytes)
+            -    height := packer.UnpackLong()
+            -    txBytes := packer.UnpackBytes()
+            -    tx := &Tx{}
+            -    if _, err := vm.codec.Unmarshal(txBytes, tx); err != nil {
+            -        return nil, 0, fmt.Errorf("problem parsing atomic transaction from db: %w", err)
+            -    }
+            -    if err := tx.Sign(vm.codec, nil); err != nil {
+            -        return nil, 0, fmt.Errorf("problem initializing atomic transaction from db: %w", err)
+            -    }
+            -    return tx, height, nil
+        }
+    }
+}
```

```

-}
-
// getAtomicTx returns the requested transaction, status, and height.
// If the status is Unknown, then the returned transaction will be nil.
func (vm *VM) getAtomicTx(txID ids.ID) (*Tx, Status, uint64, error) {
-   if tx, height, err := vm.getAcceptedAtomicTx(txID); err == nil {
+   if tx, height, err := vm.atomicTxRepository.GetByTxID(txID); err == nil {
+       if tx, height, err := vm.atomicTxRepository.GetByTxID(txID); err == nil {
+           return tx, Accepted, height, nil
+       } else if err != database.ErrNotFound {
+           return nil, Unknown, 0, err
+       }
+
+       tx, dropped, found := vm.mempool.GetTx(txID)
+       switch {
+       case found && dropped:
@@ -851,20 +1029,6 @@ func (vm *VM) getAtomicTx(txID ids.ID) (*Tx, Status, uint64, error) {
    }
}

-// writeAtomicTx writes indexes [tx] in [blk]
-func (vm *VM) writeAtomicTx(blk *Block, tx *Tx) error {
-   // 8 bytes
-   height := blk.ethBlock.NumberU64()
-   // 4 + len(txBytes)
-   txBytes := tx.Bytes()
-   packer := wrappers.Packer{Bytes: make([]byte, 12+len(txBytes))}
-   packer.PackLong(height)
-   packer.PackBytes(txBytes)
-   txID := tx.ID()
-
-   return vm.acceptedAtomicTxDB.Put(txID[:], packer.Bytes)
-}
-
// ParseAddress takes in an address and produces the ID of the chain it's for
// the ID of the address
func (vm *VM) ParseAddress(addrStr string) (ids.ID, ids.ShortID, error) {
@@ -924,7 +1088,6 @@ func (vm *VM) issueTx(tx *Tx, local bool) error {
    }
    return err
}

-   // NOTE: Gossiping of the issued [Tx] is handled in [AddTx]
-   return nil
}
@@ -939,10 +1102,10 @@ func (vm *VM) verifyTxAtTip(tx *Tx) error {
    rules := vm.currentRules()
    parentHeader := preferredBlock.Header()
    var nextBaseFee *big.Int
-   timestamp := time.Now().Unix()
+   timestamp := vm.clock.Time().Unix()
    bigTimestamp := big.NewInt(timestamp)
    if vm.chainConfig.IsApricotPhase3(bigTimestamp) {
-       _, nextBaseFee, err = dummy.CalcBaseFee(vm.chainConfig, parentHeader, uint64(timestamp))
+       _, nextBaseFee, err = dummy.EstimateNextBaseFee(vm.chainConfig, parentHeader, uint64(timestamp))
+       if err != nil {
+           // Return extremely detailed error since CalcBaseFee should never encounter an issue here
+           return fmt.Errorf("failed to calculate base fee with parent timestamp (%d), parent ExtraData: (0x%x), and current timestamp (%d): %w", parentHeader.Time, parentHeader.Extra, times
@@ -1201,6 +1364,8 @@ func (vm *VM) currentRules() params.Rules {
    // follows the ruleset defined by [rules]
    func (vm *VM) getBlockValidator(rules params.Rules) BlockValidator {
        switch {
+       case rules.IsApricotPhase5:
+           return phase5BlockValidator
+       case rules.IsApricotPhase4:
+           return phase4BlockValidator
+       case rules.IsApricotPhase3:
@@ -1255,3 +1420,84 @@ func (vm *VM) estimateBaseFee(ctx context.Context) (*big.Int, error) {
    return baseFee, nil
}

+func getAtomicRepositoryRepairHeights(chainID *big.Int) []uint64 {
+   if chainID.Cmp(params.AvalancheMainnetChainID) != 0 {
+       return nil
+   }
+   repairHeights := make([]uint64, 0, len(bonusBlockMainnetHeights)+len(canonicalBonusBlocks))
+   for height := range bonusBlockMainnetHeights {
+       repairHeights = append(repairHeights, height)
+   }
+   for _, height := range canonicalBonusBlocks {
+       if _, exists := bonusBlockMainnetHeights[height]; !exists {
+           repairHeights = append(repairHeights, height)
+       }
+   }
+   sort.Slice(repairHeights, func(i, j int) bool { return repairHeights[i] < repairHeights[j] })
+   return repairHeights
+}
+
+func (vm *VM) getAtomicTxFromPreApricot5BlockByHeight(height uint64) (*Tx, error) {
+   blk := vm.chain.GetBlockByNumber(height)
+   if blk == nil {
+       return nil, nil
+   }
+   return ExtractAtomicTx(blk.ExtData(), vm.codec)
+}
+
+// repairAtomicRepositoryForBonusBlockTxns ensures that atomic txs that were processed
+// on more than one block (canonical block + a number of bonus blocks) are indexed to
+// the first height they were processed on (canonical block).
+// [sortedHeights] should include all canonical block + bonus block heights in ascending
+// order, and will only be passed as non-empty on mainnet.
+func (vm *VM) repairAtomicRepositoryForBonusBlockTxns(
+   sortedHeights []uint64, getAtomicTxFromBlockByHeight func(height uint64) (*Tx, error),
+) error {
+   done, err := vm.atomicTxRepository.IsBonusBlocksRepaired()
+   if err != nil {
+       return err
+   }
+   if done {
+       return nil
+   }
+   repairedEntries := uint64(0)
+   seenTxns := make(map[ids.ID][]uint64)
+   for _, height := range sortedHeights {
+       // get atomic tx from block
+       tx, err := getAtomicTxFromBlockByHeight(height)
+       if err != nil {
+           return err
+       }
+       if tx == nil {
+           continue
+       }
+
+       // get the tx by txID and update it, the first time we encounter
+       // a given [txID], overwrite the previous [txID] => [height]
+       // mapping. This provides a canonical mapping across nodes.
+       heights, seen := seenTxns[tx.ID()]
+       _, foundHeight, err := vm.atomicTxRepository.GetByTxID(tx.ID())
+       if err != nil && !errors.Is(err, database.ErrNotFound) {
+           return err
+       }
+       if !seen {
+           if err := vm.atomicTxRepository.Write(height, []*Tx{tx}); err != nil {
+               return err
+           }

```

```
+     }
+     } else {
+         if err := vm.atomicTxRepository.WriteBonus(height, []Tx{tx}); err != nil {
+             return err
+         }
+     }
+     if foundHeight != height && !seen {
+         repairedEntries++
+     }
+     seenTxs[tx.ID()] = append(heights, height)
+ }
+ if err := vm.atomicTxRepository.MarkBonusBlocksRepaired(repairedEntries); err != nil {
+     return err
+ }
+ log.Info("repairAtomicRepositoryForBonusBlockTxs complete", "repairedEntries", repairedEntries)
+ return vm.db.Commit()
+})
diff --git a/plugin/evm/vm_test.go b/plugin/evm/vm_test.go
index 2050e189..76f4d1e0 100644
--- a/plugin/evm/vm_test.go
+++ b/plugin/evm/vm_test.go
@@ -12,43 +12,48 @@ import (
    "math/big"
    "os"
    "path/filepath"
+   "sort"
    "strings"
    "testing"
    "time"

-   "github.com/ava-labs/coreth/trie"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/log"
+   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/trie"

-   "github.com/stretchr/testify/assert"
+
+   "github.com/ava-labs/avalanchego/api/keystore"
+   "github.com/ava-labs/avalanchego/chains/atomic"
+   "github.com/ava-labs/avalanchego/database/manager"
+   "github.com/ava-labs/avalanchego/database/prefixdb"
+   "github.com/ava-labs/avalanchego/ids"
+   "github.com/ava-labs/avalanchego/snow"
+   "github.com/ava-labs/avalanchego/snow/choices"
+   "github.com/ava-labs/avalanchego/utls/crypto"
+   "github.com/ava-labs/avalanchego/utls/formatting"
+   "github.com/ava-labs/avalanchego/utls/hashing"
+   "github.com/ava-labs/avalanchego/utls/logging"
+   "github.com/ava-labs/avalanchego/utls/units"
+   "github.com/ava-labs/avalanchego/version"
+   "github.com/ava-labs/avalanchego/vms/components/avax"
+   "github.com/ava-labs/avalanchego/vms/components/chain"
+   "github.com/ava-labs/avalanchego/vms/secp256k1fx"

-   engCommon "github.com/ava-labs/avalanchego/snow/engine/common"
+
+   "github.com/ava-labs/coreth/consensus/dummy"
+   "github.com/ava-labs/coreth/core"
+   "github.com/ava-labs/coreth/core/types"
+   "github.com/ava-labs/coreth/eth"
+   "github.com/ava-labs/coreth/params"
+   "github.com/ava-labs/coreth/rpc"

-   accountKeystore "github.com/ava-labs/coreth/accounts/keystore"
+   "github.com/flare-foundation/flare/api/keystore"
+   "github.com/flare-foundation/flare/chains/atomic"
+   "github.com/flare-foundation/flare/database/manager"
+   "github.com/flare-foundation/flare/database/memdb"
+   "github.com/flare-foundation/flare/database/prefixdb"
+   "github.com/flare-foundation/flare/database/versiondb"
+   "github.com/flare-foundation/flare/ids"
+   "github.com/flare-foundation/flare/snow"
+   "github.com/flare-foundation/flare/snow/choices"
+   "github.com/flare-foundation/flare/utls/constants"
+   "github.com/flare-foundation/flare/utls/crypto"
+   "github.com/flare-foundation/flare/utls/formatting"
+   "github.com/flare-foundation/flare/utls/hashing"
+   "github.com/flare-foundation/flare/utls/logging"
+   "github.com/flare-foundation/flare/utls/units"
+   "github.com/flare-foundation/flare/version"
+   "github.com/flare-foundation/flare/vms/components/avax"
+   "github.com/flare-foundation/flare/vms/components/chain"
+   "github.com/flare-foundation/flare/vms/secp256k1fx"

-   engCommon "github.com/flare-foundation/flare/snow/engine/common"
+
+   "github.com/flare-foundation/coreth/consensus/dummy"
+   "github.com/flare-foundation/coreth/core"
+   "github.com/flare-foundation/coreth/core/types"
+   "github.com/flare-foundation/coreth/eth"
+   "github.com/flare-foundation/coreth/params"
+   "github.com/flare-foundation/coreth/rpc"

-   accountKeystore "github.com/flare-foundation/coreth/accounts/keystore"
+
+   )

var (
@@ -62,19 +62,21 @@ func TestAvaxAssetID (
    testAvaxAssetID = ids.ID{1, 2, 3}
    username         = "Johns"
    password         = "CjasdjhjPeirbSenfeI13" // #nosec G101
-   // Use chainId: 43111, so that it does not overlap with any Avalanche ChainIDs, which may have their
+   // Use chainId: 31337, so that it does not overlap with any Avalanche ChainIDs, which may have their
+   // config overridden in vm.Initialize.
-   genesisJ50NApricotPhase0 = "{\n\"config\":{\n\"chainId\":43111,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
-   genesisJ50NApricotPhase1 = "{\n\"config\":{\n\"chainId\":43111,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
-   genesisJ50NApricotPhase2 = "{\n\"config\":{\n\"chainId\":43111,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
-   genesisJ50NApricotPhase3 = "{\n\"config\":{\n\"chainId\":43111,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
-   genesisJ50NApricotPhase4 = "{\n\"config\":{\n\"chainId\":43111,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase0 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase1 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase2 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase3 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase4 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",
+   genesisJ50NApricotPhase5 = "{\n\"config\":{\n\"chainId\":31337,\n\"homesteadBlock\":0,\n\"daoForkBlock\":0,\n\"daoForkSupport\":true,\n\"eip150Block\":0,\n\"eip150Hash\":\n\"0x2086799aeebeae135c246c65021c82b4e15\",

    apricotRulesPhase0 = params.Rules{}
    apricotRulesPhase1 = params.Rules{IsApricotPhase1: true}
    apricotRulesPhase2 = params.Rules{IsApricotPhase1: true, IsApricotPhase2: true}
    apricotRulesPhase3 = params.Rules{IsApricotPhase1: true, IsApricotPhase2: true, IsApricotPhase3: true}
    apricotRulesPhase4 = params.Rules{IsApricotPhase1: true, IsApricotPhase2: true, IsApricotPhase3: true, IsApricotPhase4: true}
+   apricotRulesPhase5 = params.Rules{IsApricotPhase1: true, IsApricotPhase2: true, IsApricotPhase3: true, IsApricotPhase4: true, IsApricotPhase5: true}

)

func init() {
@@ -123,16 +123,18 @@ func NewContext() *snow.Context {
    ctx.NetworkID = testNetworkID
    ctx.ChainID = testChainID
    ctx.AVAXAssetID = testAvaxAssetID
-   ctx.XChainID = ids.Empty.Prefix(0)
+   ctx.XChainID = testXChainID
+   aliaser := ctx.BCLookup.(ids.Aliaser)
+   _ = aliaser.Alias(testChainID, "C")
}
```

```

- = aliaser.Alias(testCChainID, testCChainID.String())
- = aliaser.Alias(testXChainID, "X")
- = aliaser.Alias(testXChainID, testXChainID.String())

+ ctx.SNLookup = &snLookup{
+   chainsToSubnet: map[ids.ID]ids.ID{
+       constants.PlatformChainID: constants.PrimaryNetworkID,
+       testXChainID:               constants.PrimaryNetworkID,
+       testCChainID:               constants.PrimaryNetworkID,
+   },
+ }
+ return ctx
}

+type snLookup struct {
+   chainsToSubnet map[ids.ID]ids.ID
+}
+
+func (sn *snLookup) SubnetID(chainID ids.ID) (ids.ID, error) {
+   subnetID, ok := sn.chainsToSubnet[chainID]
+   if !ok {
+       return ids.ID{}, errors.New("unknown chain")
+   }
+   return subnetID, nil
+}
+
+func setupGenesis(t *testing.T,
+   genesisJSON string,
+   ) (*snow.Context,
+   @@ -194,17 +219,18 @@ func GenesisVM(t *testing.T,
+   ) {
+   if finishBootstrapping {
+       assert.NoError(t, vm.Bootstrapping())
+       assert.NoError(t, vm.Bootstrapped())
+       assert.NoError(t, vm.SetState(snow.Bootstrapping))
+       assert.NoError(t, vm.SetState(snow.NormalOp))
+   }
+
+   return issuer, vm, dbManager, m, appSender
+}

- func addUTXO(sharedMemory *atomic.Memory, ctx *snow.Context, txID ids.ID, assetID ids.ID, amount uint64, addr ids.ShortID) (*avax.UTXO, error) {
+func addUTXO(sharedMemory *atomic.Memory, ctx *snow.Context, txID ids.ID, index uint32, assetID ids.ID, amount uint64, addr ids.ShortID) (*avax.UTXO, error) {
+   utxo := &avax.UTXO{
+       UTXOID: avax.UTXOID{
+           TxID: txID,
+           TxID: txID,
+           OutputIndex: index,
+       },
+       Asset: avax.Asset{ID: assetID},
+       Out: &secp256k1fx.TransferOutput{
@@ -244,7 +270,7 @@ func GenesisVMWithUTXOs(t *testing.T, finishBootstrapping bool, genesisJSON stri
+   if err != nil {
+       t.Fatalf("Failed to generate txID from addr: %s", err)
+   }
+   if _, err := addUTXO(sharedMemory, vm.ctx, txID, vm.ctx.AVAXAssetID, avaxAmount, addr); err != nil {
+   if _, err := addUTXO(sharedMemory, vm.ctx, txID, 0, vm.ctx.AVAXAssetID, avaxAmount, addr); err != nil {
+       t.Fatalf("Failed to add UTXO to shared memory: %s", err)
+   }
+   }
+}

@@ -254,24 +280,24 @@ func GenesisVMWithUTXOs(t *testing.T, finishBootstrapping bool, genesisJSON stri

func TestVMConfig(t *testing.T) {
+   txFeeCap := float64(11)
-   netAPIEnabled := true
-   configJSON := fmt.Sprintf("{\"rpc-tx-fee-cap\": %g, \"net-api-enabled\": %t}", txFeeCap, netAPIEnabled)
+   enabledEthAPIs := []string{"internal-private-debug"}
+   configJSON := fmt.Sprintf("{\"rpc-tx-fee-cap\": %g, \"eth-apis\": %s}", txFeeCap, fmt.Sprintf("[%q]", enabledEthAPIs[0]))
+   vm, _, _ := GenesisVM(t, false, genesisJSONApricotPhase0, configJSON, "")
+   assert.Equal(t, vm.config.RPCTxFeeCap, txFeeCap, "Tx Fee Cap should be set")
+   assert.Equal(t, vm.config.NetAPIEnabled, netAPIEnabled, "Net API Enabled should be set")
+   assert.Equal(t, vm.config.EthAPIs(), enabledEthAPIs, "EnabledEthAPIs should be set")
+   assert.NoError(t, vm.Shutdown())
+}

func TestVMConfigDefaults(t *testing.T) {
+   txFeeCap := float64(11)
-   netAPIEnabled := true
-   configJSON := fmt.Sprintf("{\"rpc-tx-fee-cap\": %g, \"net-api-enabled\": %t}", txFeeCap, netAPIEnabled)
+   enabledEthAPIs := []string{"internal-private-debug"}
+   configJSON := fmt.Sprintf("{\"rpc-tx-fee-cap\": %g, \"eth-apis\": %s}", txFeeCap, fmt.Sprintf("[%q]", enabledEthAPIs[0]))
+   vm, _, _ := GenesisVM(t, false, genesisJSONApricotPhase0, configJSON, "")
+   var vmConfig Config
+   vmConfig.SetDefaults()
+   vmConfig.RPCTxFeeCap = txFeeCap
+   vmConfig.NetAPIEnabled = netAPIEnabled
+   vmConfig.EnabledEthAPIs = enabledEthAPIs
+   assert.Equal(t, vmConfig, vm.config, "VM Config should match default with overrides")
+   assert.NoError(t, vm.Shutdown())
+}

@@ -336,6 +362,11 @@ func TestVMUpgrades(t *testing.T) {
+   genesis: genesisJSONApricotPhase4,
+   expectedGasPrice: big.NewInt(0),
+   },
+   {
+       name: "Apricot Phase 5",
+       genesis: genesisJSONApricotPhase5,
+       expectedGasPrice: big.NewInt(0),
+   },
+   for _, test := range genesisTests {
+       t.Run(test.name, func(t *testing.T) {
@@ -394,2519 +425,63 @@ func TestVMUpgrades(t *testing.T) {
+}

-// Simple test to ensure we can issue an import transaction followed by an export transaction
-// and they will be indexed correctly when accepted.
- func TestIssueAtomicTx(t *testing.T) {
+   importAmount := uint64(50000000)
+   issuer, vm, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase2, "", "", map[ids.ShortID]uint64{
+       testShortIDAddr[0]: importAmount,
+   })
+   defer func() {
+       if err := vm.Shutdown(); err != nil {
+           t.Fatal(err)
+       }
+   }()
+   importTx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
+   if err != nil {
+       t.Fatal(err)
+   }
+   if err := vm.issueTx(importTx, true /*=local*/); err != nil {
+       t.Fatal(err)
+   }
+   <-issuer
+   blk, err := vm.BuildBlock()

```

```

- if err != nil {
-     t.Fatal(err)
- }
-
- if err := blk.Verify(); err != nil {
-     t.Fatal(err)
- }
-
- if status := blk.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
- }
-
- if err := vm.SetPreference(blk.ID()); err != nil {
-     t.Fatal(err)
- }
-
- if err := blk.Accept(); err != nil {
-     t.Fatal(err)
- }
-
- if status := blk.Status(); status != choices.Accepted {
-     t.Fatalf("Expected status of accepted block to be %s, but found %s", choices.Accepted, status)
- }
-
- if lastAcceptedID, err := vm.LastAccepted(); err != nil {
-     t.Fatal(err)
- } else if lastAcceptedID != blk.ID() {
-     t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk.ID(), lastAcceptedID)
- }
-
- exportTx, err := vm.newExportTx(vm.ctx.AVAXAssetID, importAmount-(2*params.AvalancheAtomicTxFee), vm.ctx.XChainID, testShortIDAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
- if err != nil {
-     t.Fatal(err)
- }
-
- if err := vm.issueTx(exportTx, true /*=local*/); err != nil {
-     t.Fatal(err)
- }
-
- <-issuer
-
- blk2, err := vm.BuildBlock()
- if err != nil {
-     t.Fatal(err)
- }
-
- if err := blk2.Verify(); err != nil {
-     t.Fatal(err)
- }
-
- if status := blk2.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
- }
-
- if err := blk2.Accept(); err != nil {
-     t.Fatal(err)
- }
-
- if status := blk2.Status(); status != choices.Accepted {
-     t.Fatalf("Expected status of accepted block to be %s, but found %s", choices.Accepted, status)
- }
-
- if lastAcceptedID, err := vm.LastAccepted(); err != nil {
-     t.Fatal(err)
- } else if lastAcceptedID != blk2.ID() {
-     t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk2.ID(), lastAcceptedID)
- }
-
- // Check that both atomic transactions were indexed as expected.
- indexedImportTx, status, height, err := vm.getAtomicTx(importTx.ID())
- assert.NoError(t, err)
- assert.Equal(t, Accepted, status)
- assert.Equal(t, uint64(1), height, "expected height of indexed import tx to be 1")
- assert.Equal(t, indexedImportTx.ID(), importTx.ID(), "expected ID of indexed import tx to match original txID")
-
- indexedExportTx, status, height, err := vm.getAtomicTx(exportTx.ID())
- assert.NoError(t, err)
- assert.Equal(t, Accepted, status)
- assert.Equal(t, uint64(2), height, "expected height of indexed export tx to be 2")
- assert.Equal(t, indexedExportTx.ID(), exportTx.ID(), "expected ID of indexed import tx to match original txID")
-}
-
-func TestBuildEthTxBlock(t *testing.T) {
-    importAmount := uint64(20000000)
-    issuer, vm, dbManager, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase2, "{\"pruning-enabled\":true}", "", map[ids.ShortID]uint64{
-        testShortIDAddrs[0]: importAmount,
-    })
-
-    defer func() {
-        if err := vm.Shutdown(); err != nil {
-            t.Fatal(err)
-        }
-    }()
-
-    newTxPoolHeadChan := make(chan core.NewTxPoolReorgEvent, 1)
-    vm.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan)
-
-    key, err := accountKeystore.NewKey(rand.Reader)
-    if err != nil {
-        t.Fatal(err)
-    }
-
-    importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
-    if err != nil {
-        t.Fatal(err)
-    }
-
-    if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-        t.Fatal(err)
-    }
-
-    <-issuer
-
-    blk1, err := vm.BuildBlock()
-    if err != nil {
-        t.Fatal(err)
-    }
-
-    if err := blk1.Verify(); err != nil {
-        t.Fatal(err)
-    }
-
-    if status := blk1.Status(); status != choices.Processing {
-        t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-    }
-
-    if err := vm.SetPreference(blk1.ID()); err != nil {
-        t.Fatal(err)
-    }
-
-    if err := blk1.Accept(); err != nil {
-        t.Fatal(err)
-    }
-
-    if status := blk1.Status(); status != choices.Accepted {
-        t.Fatalf("Expected status of accepted block to be %s, but found %s", choices.Accepted, status)
-    }
-
-    if lastAcceptedID, err := vm.LastAccepted(); err != nil {
-        t.Fatal(err)
-    } else if lastAcceptedID != blk1.ID() {
-        t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk1.ID(), lastAcceptedID)
-    }
-
-    // Check that both atomic transactions were indexed as expected.
-    indexedImportTx, status, height, err := vm.getAtomicTx(importTx.ID())
-    assert.NoError(t, err)
-    assert.Equal(t, Accepted, status)
-    assert.Equal(t, uint64(1), height, "expected height of indexed import tx to be 1")
-    assert.Equal(t, indexedImportTx.ID(), importTx.ID(), "expected ID of indexed import tx to match original txID")
-
-    indexedExportTx, status, height, err := vm.getAtomicTx(exportTx.ID())
-    assert.NoError(t, err)
-    assert.Equal(t, Accepted, status)
-    assert.Equal(t, uint64(2), height, "expected height of indexed export tx to be 2")
-    assert.Equal(t, indexedExportTx.ID(), exportTx.ID(), "expected ID of indexed import tx to match original txID")
-}

```

```

- newHead := <-newTxPoolHeadChan
- if newHead.Head.Hash() != common.Hash(blk1.ID()) {
-     t.Fatalf("Expected new block to match")
- }
-
- txs := make([]*types.Transaction, 10)
- for i := 0; i < 10; i++ {
-     tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-     signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
-     if err != nil {
-         t.Fatal(err)
-     }
-     txs[i] = signedTx
- }
- errs := vm.chain.AddRemoteTxsSync(txs)
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add tx at index %d: %s", i, err)
-     }
- }
-
- <-issuer
-
- blk2, err := vm.BuildBlock()
- if err != nil {
-     t.Fatal(err)
- }
-
- if err := blk2.Verify(); err != nil {
-     t.Fatal(err)
- }
-
- if status := blk2.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
- }
-
- if err := blk2.Accept(); err != nil {
-     t.Fatal(err)
- }
-
- newHead = <-newTxPoolHeadChan
- if newHead.Head.Hash() != common.Hash(blk2.ID()) {
-     t.Fatalf("Expected new block to match")
- }
-
- if status := blk2.Status(); status != choices.Accepted {
-     t.Fatalf("Expected status of accepted block to be %s, but found %s", choices.Accepted, status)
- }
-
- lastAcceptedID, err := vm.LastAccepted()
- if err != nil {
-     t.Fatal(err)
- }
-
- if lastAcceptedID != blk2.ID() {
-     t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk2.ID(), lastAcceptedID)
- }
-
- ethBlk1 := blk1.(*chain.BlockWrapper).Block.(*Block).ethBlock
- if ethBlk1.Root := ethBlk1.Root(); !vm.chain.BlockChain().HasState(ethBlk1.Root) {
-     t.Fatalf("Expected blk1 state root to not yet be pruned after blk2 was accepted because of tip buffer")
- }
-
- // Clear the cache and ensure that GetBlock returns internal blocks with the correct status
- vm.State.Flush()
- blk2Refreshed, err := vm.GetBlockInternal(blk2.ID())
- if err != nil {
-     t.Fatal(err)
- }
-
- if status := blk2Refreshed.Status(); status != choices.Accepted {
-     t.Fatalf("Expected refreshed blk2 to be Accepted, but found status: %s", status)
- }
-
- blk1RefreshedID := blk2Refreshed.Parent()
- blk1Refreshed, err := vm.GetBlockInternal(blk1RefreshedID)
- if err != nil {
-     t.Fatal(err)
- }
-
- if status := blk1Refreshed.Status(); status != choices.Accepted {
-     t.Fatalf("Expected refreshed blk1 to be Accepted, but found status: %s", status)
- }
-
- if blk1Refreshed.ID() != blk1.ID() {
-     t.Fatalf("Found unexpected blkID for parent of blk2")
- }
-
- restartedVM := &VM{}
- if err := restartedVM.Initialize(
-     NewContext(),
-     dbManager,
-     []byte(genesisJSONApricotPhase2),
-     []byte(""),
-     []byte("{\"pruning-enabled\":true"}),
-     issuer,
-     []*engCommon.Fx{},
-     nil,
- ); err != nil {
-     t.Fatal(err)
- }
-
- // State root should not have been committed and discarded on restart
- if ethBlk1.Root := ethBlk1.Root(); restartedVM.chain.BlockChain().HasState(ethBlk1.Root) {
-     t.Fatalf("Expected blk1 state root to be pruned after blk2 was accepted on top of it in pruning mode")
- }
-
- // State root should be committed when accepted tip on shutdown
- ethBlk2 := blk2.(*chain.BlockWrapper).Block.(*Block).ethBlock
- if ethBlk2.Root := ethBlk2.Root(); !restartedVM.chain.BlockChain().HasState(ethBlk2.Root) {
-     t.Fatalf("Expected blk2 state root to not be pruned after shutdown (last accepted tip should be committed)")
- }
-
-}
-
-func TestConflictingImportTxs(t *testing.T) {
-     importAmount := uint64(10000000)
-     issuer, vm, _, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-         testShortIDAddrs[1]: importAmount,
-         testShortIDAddrs[2]: importAmount,
-     })
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     conflictKey, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTxs := make([]*Tx, 0, 3)
-     conflictTxs := make([]*Tx, 0, 3)
-     for i, key := range testKeys {
-         importTx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[i], initialBaseFee, []crypto.PrivateKeySECP256K1R{key})

```



```

-         if err != nil {
-             t.Fatal(err)
-         }
-         importTx = append(importTx, importTx)
-
-         conflictTx, err := vm.newImportTx(vm.ctx.XChainID, conflictKey.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{key})
-         if err != nil {
-             t.Fatal(err)
-         }
-         conflictTx = append(conflictTx, conflictTx)
-     }
-
-     expectedParentBlkID, err := vm.LastAccepted()
-     if err != nil {
-         t.Fatal(err)
-     }
-     for i, tx := range importTx {
-         if err := vm.issueTx(tx, true /*=local*/); err != nil {
-             t.Fatal(err)
-         }
-
-         <-issuer
-
-         blk, err := vm.BuildBlock()
-         if err != nil {
-             t.Fatal(err)
-         }
-
-         if err := blk.Verify(); err != nil {
-             t.Fatal(err)
-         }
-
-         if status := blk.Status(); status != choices.Processing {
-             t.Fatalf("Expected status of built block %d to be %s, but found %s", i, choices.Processing, status)
-         }
-
-         if parentID := blk.Parent(); parentID != expectedParentBlkID {
-             t.Fatalf("Expected parent to have blockID %s, but found %s", expectedParentBlkID, parentID)
-         }
-
-         expectedParentBlkID = blk.ID()
-         if err := vm.SetPreference(blk.ID()); err != nil {
-             t.Fatal(err)
-         }
-     }
-
-     for i, tx := range conflictTx {
-         if err := vm.issueTx(tx, true /*=local*/); err == nil {
-             t.Fatalf("Expected issueTx to fail due to conflicting transaction")
-         }
-         // Force issue transaction directly to the mempool
-         if err := vm.mempool.ForceAddTx(tx); err != nil {
-             t.Fatal(err)
-         }
-
-         <-issuer
-
-         _, err = vm.BuildBlock()
-         // The new block is verified in BuildBlock, so
-         // BuildBlock should fail due to an attempt to
-         // double spend an atomic UTX0.
-         if err == nil {
-             t.Fatalf("Block verification should have failed in BuildBlock %d due to double spending atomic UTX0", i)
-         }
-     }
+func TestConfigureLogLevel(t *testing.T) {
+    configTests := []struct {
+        name          string
+        logConfig      string
+        genesisJSON    string
+        upgradeJSON    string
+        expectedErr    string
+    }{
+        {
+            name:          "Log level info",
+            logConfig:     "{\"log-level\": \"info\"}",
+            genesisJSON:   genesisJSONApricotPhase2,
+            upgradeJSON:  "",
+            expectedErr:  "",
+        },
+        {
+            name:          "Invalid log level",
+            logConfig:     "{\"log-level\": \"cchain\"}",
+            genesisJSON:   genesisJSONApricotPhase3,
+            upgradeJSON:  "",
+            expectedErr:   "failed to initialize logger due to",
+        },
+    }
-}
-
-// Regression test to ensure that after accepting block A
-// then calling SetPreference on block B (when it becomes preferred)
-// and the head of a longer chain (block D) does not corrupt the
-// canonical chain.
-// A
-// / \
-// B  C
-//   |
-//   D
-
-func TestSetPreferenceRace(t *testing.T) {
-    // Create two VMs which will agree on block A and then
-    // build the two distinct preferred chains above
-    importAmount := uint64(1000000000)
-    issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "{\"pruning-enabled\":true}", "", map[ids.ShortID]uint64{
-        testShortIDAddrs[0]: importAmount,
-    })
-    issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "{\"pruning-enabled\":true}", "", map[ids.ShortID]uint64{
-        testShortIDAddrs[0]: importAmount,
-    })
-    for _, test := range configTests {
-        t.Run(test.name, func(t *testing.T) {
-            vm := &VM{
-                ctx, dbManager, genesisBytes, issuer, _ := setupGenesis(t, test.genesisJSON)
-                appSender := &engCommon.SenderTest{}
-                appSender.CantSendAppGossip = true
-                appSender.SendAppGossipF = func([]byte) error { return nil }
-                err := vm.Initialize(
-                    ctx,
-                    dbManager,
-                    genesisBytes,
-                    []byte(""),
-                    []byte(test.logConfig),
-                    issuer,
-                    []*engCommon.Fx{},
-                    appSender,
-                )
-                if len(test.expectedErr) == 0 && err != nil {
-                    t.Fatal(err)
-                } else if len(test.expectedErr) > 0 {
-                    if err == nil {
-                        t.Fatalf("initialize should have failed due to %s", test.expectedErr)
-                    } else if !strings.Contains(err.Error(), test.expectedErr) {
-                        t.Fatalf("Expected initialize to fail due to %s, but failed with %s", test.expectedErr, err.Error())
-                    }
-                }
-            }
-        })
-    }
-}
-
-defer func() {

```

```

-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     })
-
-     newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
-     newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer1
-
-     vm1BlkA, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := vm1BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if status := vm1BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-
-     if err := vm2BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-
-     if status := vm2BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkA.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-
-     if err := vm2BlkA.Accept(); err != nil {
-         t.Fatalf("VM2 failed to accept block: %s", err)
-     }
-
-     newHead := <-newTxPoolHeadChan1
-     if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Create list of 10 successive transactions to build block A on vm1
-     // and to be split into two separate blocks on VM2
-     txs := make([]*types.Transaction, 10)
-     for i := 0; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-
-     var errs []error
-
-     // Add the remote transactions, build the block, and set VM1's preference for block A
-     errs = vm1.chain.AddRemoteTxsSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer1
-
-     vm1BlkB, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkB.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := vm1BlkB.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     // Split the transactions over two blocks, and set VM2's preference to them in sequence
-     // after building each block
-     // Block C
-     errs = vm2.chain.AddRemoteTxsSync(txs[0:5])
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer2
-
-     vm2BlkC, err := vm2.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build BlkC on VM2: %s", err)
-     }

```

```

-     if err := vm2BlkC.Verify(); err != nil {
-         t.Fatalf("BlkC failed verification on VM2: %s", err)
-     }
-
-     if status := vm2BlkC.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block C to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm2.SetPreference(vm2BlkC.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkC.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Block D
-     errs = vm2.chain.AddRemoteTxsSync(txs[5:10])
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer2
-     vm2BlkD, err := vm2.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build BlkD on VM2: %s", err)
-     }
-
-     if err := vm2BlkD.Verify(); err != nil {
-         t.Fatalf("BlkD failed verification on VM2: %s", err)
-     }
-
-     if status := vm2BlkD.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block D to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm2.SetPreference(vm2BlkD.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     // VM1 receives blkC and blkD from VM1
-     // and happens to call SetPreference on blkD without ever calling SetPreference
-     // on blkC
-     // Here we parse them in reverse order to simulate receiving a chain from the tip
-     // back to the last accepted block as would typically be the case in the consensus
-     // engine
-     vm1BlkD, err := vm1.ParseBlock(vm2BlkD.Bytes())
-     if err != nil {
-         t.Fatalf("VM1 errored parsing blkD: %s", err)
-     }
-     vm1BlkC, err := vm1.ParseBlock(vm2BlkC.Bytes())
-     if err != nil {
-         t.Fatalf("VM1 errored parsing blkC: %s", err)
-     }
-
-     // The blocks must be verified in order. This invariant is maintained
-     // in the consensus engine.
-     if err := vm1BlkC.Verify(); err != nil {
-         t.Fatalf("VM1 BlkC failed verification: %s", err)
-     }
-     if err := vm1BlkD.Verify(); err != nil {
-         t.Fatalf("VM1 BlkD failed verification: %s", err)
-     }
-
-     // Set VM1's preference to blockD, skipping blockC
-     if err := vm1.SetPreference(vm1BlkD.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     // Accept the longer chain on both VMs and ensure there are no errors
-     // VM1 Accepts the blocks in order
-     if err := vm1BlkC.Accept(); err != nil {
-         t.Fatalf("VM1 BlkC failed on accept: %s", err)
-     }
-     if err := vm1BlkD.Accept(); err != nil {
-         t.Fatalf("VM1 BlkC failed on accept: %s", err)
-     }
-
-     // VM2 Accepts the blocks in order
-     if err := vm2BlkC.Accept(); err != nil {
-         t.Fatalf("VM2 BlkC failed on accept: %s", err)
-     }
-     if err := vm2BlkD.Accept(); err != nil {
-         t.Fatalf("VM2 BlkC failed on accept: %s", err)
-     }
-
-     log.Info("Validating canonical chain")
-     // Verify the Canonical Chain for Both VMs
-     if err := vm2.chain.ValidateCanonicalChain(); err != nil {
-         t.Fatalf("VM2 failed canonical chain verification due to: %s", err)
-     }
-
-     if err := vm1.chain.ValidateCanonicalChain(); err != nil {
-         t.Fatalf("VM1 failed canonical chain verification due to: %s", err)
-     }
- }
-
- func TestConflictingTransitiveAncestryWithGap(t *testing.T) {
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     key0 := testKeys[0]
-     addr0 := key0.PublicKey().Address()
-
-     key1 := testKeys[1]
-     addr1 := key1.PublicKey().Address()
-
-     importAmount := uint64(1000000000)
-
-     issuer, vm, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "",
-         map[ids.ShortID]uint64{
-             addr0: importAmount,
-             addr1: importAmount,
-         })
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan := make(chan core.NewTxPoolReorgEvent, 1)
-     vm.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan)
-
-     importTx0A, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{key0})
-     if err != nil {
-         t.Fatal(err)
-     }
-     // Create a conflicting transaction

```

```

importTx0B, err := vm.newImportTx(vm.ctx.XChainID, testEthAddr[2], initialBaseFee, []*crypto.PrivateKeySECP256K1R{key0})
if err != nil {
    t.Fatal(err)
}

if err := vm.issueTx(importTx0A, true /*=local*/); err != nil {
    t.Fatalf("Failed to issue importTx0A: %s", err)
}

<-issuer

blk0, err := vm.BuildBlock()
if err != nil {
    t.Fatalf("Failed to build block with import transaction: %s", err)
}

if err := blk0.Verify(); err != nil {
    t.Fatalf("Block failed verification: %s", err)
}

if err := vm.SetPreference(blk0.ID()); err != nil {
    t.Fatal(err)
}

newHead := <-newTxPoolHeadChan
if newHead.Head.Hash() != common.Hash(blk0.ID()) {
    t.Fatalf("Expected new block to match")
}

tx := types.NewTransaction(0, key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
if err != nil {
    t.Fatal(err)
}

// Add the remote transactions, build the block, and set VML's preference for block A
errs := vm.chain.AddRemoteTxSync([]*types.Transaction{signedTx})
for i, err := range errs {
    if err != nil {
        t.Fatalf("Failed to add transaction to VML at index %d: %s", i, err)
    }
}

<-issuer

blk1, err := vm.BuildBlock()
if err != nil {
    t.Fatalf("Failed to build blk1: %s", err)
}

if err := blk1.Verify(); err != nil {
    t.Fatalf("blk1 failed verification due to %s", err)
}

if err := vm.SetPreference(blk1.ID()); err != nil {
    t.Fatal(err)
}

importTx1, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{key1})
if err != nil {
    t.Fatalf("Failed to issue importTx1 due to: %s", err)
}

if err := vm.issueTx(importTx1, true /*=local*/); err != nil {
    t.Fatal(err)
}

<-issuer

blk2, err := vm.BuildBlock()
if err != nil {
    t.Fatalf("Failed to build block with import transaction: %s", err)
}

if err := blk2.Verify(); err != nil {
    t.Fatalf("Block failed verification: %s", err)
}

if err := vm.SetPreference(blk2.ID()); err != nil {
    t.Fatal(err)
}

if err := vm.issueTx(importTx0B, true /*=local*/); err == nil {
    t.Fatalf("Should not have been able to issue import tx with conflict")
}
// Force issue transaction directly into the mempool
if err := vm.mempool.ForceAddTx(importTx0B); err != nil {
    t.Fatal(err)
}

<-issuer

_, err = vm.BuildBlock()
if err == nil {
    t.Fatal("Shouldn't have been able to build an invalid block")
}
-}

-func TestBonusBlocksTx(t *testing.T) {
    issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJ50NApricotPhase0, "", "")

    defer func() {
        if err := vm.Shutdown(); err != nil {
            t.Fatal(err)
        }
    }()

    importAmount := uint64(100000000)
    utxoID := avax.UTXOID{TxID: ids.GenerateTestID()}

    utxo := &avax.UTXO{
        UTXOID: utxoID,
        Asset:  avax.Asset{ID: vm.ctx.AVAXAssetID},
        Out:    &secp256k1fx.TransferOutput{
            Amt: importAmount,
            OutputOwners: secp256k1fx.OutputOwners{
                Threshold: 1,
                Addrs:      []ids.ShortID{testKeys[0].PublicKey().Address()},
            },
        },
    }

    utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
    if err != nil {
        t.Fatal(err)
    }

    xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
    inputID := utxo.InputID()
    if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
        Key:  inputID[:],
        Value: utxoBytes,
        Traits: [][]byte{
            testKeys[0].PublicKey().Address().Bytes(),
        },
    }}}}); err != nil {

```

```

-         t.Fatal(err)
-     }
-
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     bonusBlocks.Add(blk.ID())
-
-     // Remove the UTXOs from shared memory, so that non-bonus blocks will fail verification
-     if err := vm.ctx.SharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.XChainID: {RemoveRequests: [][]byte{inputID[:]}}}); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Accepted {
-         t.Fatalf("Expected status of accepted block to be %, but found %s", choices.Accepted, status)
-     }
-
-     lastAcceptedID, err := vm.LastAccepted()
-     if err != nil {
-         t.Fatal(err)
-     }
-     if lastAcceptedID != blk.ID() {
-         t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk.ID(), lastAcceptedID)
-     }
- }
-
- // Regression test to ensure that a VM that accepts block A and B
- // will not attempt to orphan either when verifying blocks C and D
- // from another VM (which have a common ancestor under the finalized
- // frontier).
- // A
- // / \
- // B   C
- //
- // verifies block B and C, then Accepts block B. Then we test to ensure
- // that the VM defends against any attempt to set the preference or to
- // accept block C, which should be an orphaned block at this point and
- // get rejected.
- func TestReorgProtection(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "{\"pruning-enabled\":false}", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-     issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "{\"pruning-enabled\":false}", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
-     newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer1
-
-     vm1BlkA, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := vm1BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if status := vm1BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-     if err := vm2BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-     if status := vm2BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %, but found %s", choices.Processing, status)
-     }
- }

```

```

-     if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkA.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-     if err := vm2BlkA.Accept(); err != nil {
-         t.Fatalf("VM2 failed to accept block: %s", err)
-     }
-
-     newHead := <-newTxPoolHeadChan1
-     if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Create list of 10 successive transactions to build block A on vm1
-     // and to be split into two separate blocks on VM2
-     txs := make([]*types.Transaction, 10)
-     for i := 0; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-
-     var errs []error
-
-     // Add the remote transactions, build the block, and set VM1's preference for block A
-     errs = vm1.chain.AddRemoteTxSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer1
-
-     vm1BlkB, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkB.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := vm1BlkB.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     // Split the transactions over two blocks, and set VM2's preference to them in sequence
-     // after building each block
-     // Block C
-     errs = vm2.chain.AddRemoteTxSync(txs[0:5])
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer2
-     vm2BlkC, err := vm2.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build BlkC on VM2: %s", err)
-     }
-
-     if err := vm2BlkC.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-     if status := vm2BlkC.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
-     }
-
-     vm1BlkC, err := vm1.ParseBlock(vm2BlkC.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-
-     if err := vm1BlkC.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     // Accept B, such that block C should get Rejected.
-     if err := vm1BlkB.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-
-     // The below (setting preference blocks that have a common ancestor
-     // with the preferred chain lower than the last finalized block)
-     // should NEVER happen. However, the VM defends against this
-     // just in case.
-     if err := vm1.SetPreference(vm1BlkC.ID()); !strings.Contains(err.Error(), "cannot orphan finalized block") {
-         t.Fatalf("Unexpected error when setting preference that would trigger reorg: %s", err)
-     }
-
-     if err := vm1BlkC.Accept(); !strings.Contains(err.Error(), "expected accepted block to have parent") {
-         t.Fatalf("Unexpected error when setting block at finalized height: %s", err)
-     }
- }
-
- // Regression test to ensure that a VM that accepts block C while preferring
- // block B will trigger a reorg.
- // A
- // / \
- // B   C
- func TestNonCanonicalAccept(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-     issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()

```

```

- newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
- vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
- newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
- vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
- key, err := accountKeystore.NewKey(rand.Reader)
- if err != nil {
-     t.Fatalf(err)
- }
-
- importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
- if err != nil {
-     t.Fatalf(err)
- }
-
- if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-     t.Fatalf(err)
- }
-
- <-issuer1
-
- vm1BlkA, err := vm1.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build block with import transaction: %s", err)
- }
-
- if err := vm1BlkA.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
-
- if status := vm1BlkA.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
- }
-
- if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-     t.Fatalf(err)
- }
-
- vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }
- if err := vm2BlkA.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM2: %s", err)
- }
- if status := vm2BlkA.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
- }
- if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-     t.Fatalf(err)
- }
-
- if err := vm1BlkA.Accept(); err != nil {
-     t.Fatalf("VM1 failed to accept block: %s", err)
- }
- if err := vm2BlkA.Accept(); err != nil {
-     t.Fatalf("VM2 failed to accept block: %s", err)
- }
-
- newHead := <-newTxPoolHeadChan1
- if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-     t.Fatalf("Expected new block to match")
- }
- newHead = <-newTxPoolHeadChan2
- if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-     t.Fatalf("Expected new block to match")
- }
-
- // Create list of 10 successive transactions to build block A on vm1
- // and to be split into two separate blocks on VM2
- txs := make([]*types.Transaction, 10)
- for i := 0; i < 10; i++ {
-     tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-     signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-     if err != nil {
-         t.Fatalf(err)
-     }
-     txs[i] = signedTx
- }
-
- var errs []error
-
- // Add the remote transactions, build the block, and set VM1's preference for block A
- errs = vm1.chain.AddRemoteTxSync(txs)
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-     }
- }
-
- <-issuer1
-
- vm1BlkB, err := vm1.BuildBlock()
- if err != nil {
-     t.Fatalf(err)
- }
-
- if err := vm1BlkB.Verify(); err != nil {
-     t.Fatalf(err)
- }
-
- if status := vm1BlkB.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
- }
-
- if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
-     t.Fatalf(err)
- }
-
- vm1.chain.BlockChain().GetVMConfig().AllowUnfinalizedQueries = true
-
- blkBHeight := vm1BlkB.Height()
- blkBHash := vm1BlkB.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
- if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkBHash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkBHash.Hex(), b.Hash().Hex())
- }
-
- errs = vm2.chain.AddRemoteTxSync(txs[0:5])
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-     }
- }
-
- <-issuer2
- vm2BlkC, err := vm2.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build BlkC on VM2: %s", err)
- }
-
- vm1BlkC, err := vm1.ParseBlock(vm2BlkC.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }

```

```

-     }
-
-     if err := vm1BlkC.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if err := vm1BlkC.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-
-     blkCHash := vm1BlkC.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
-     if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkCHash {
-         t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkCHash.Hex(), b.Hash().Hex())
-     }
- }
-
- // Regression test to ensure that a VM that verifies block B, C, then
- // D (preferring block B) does not trigger a reorg through the re-verification
- // of block C or D.
- //   A
- //  / \
- // B   C
- //   |
- //   D
- func TestStickyPreference(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddr0: importAmount,
-     })
-     issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddr0: importAmount,
-     })
-
-     defer func() {
-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
-     newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer1
-
-     vm1BlkA, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := vm1BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if status := vm1BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-     if err := vm2BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-     if status := vm2BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
-     }
-     if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkA.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-
-     if err := vm2BlkA.Accept(); err != nil {
-         t.Fatalf("VM2 failed to accept block: %s", err)
-     }
-
-     newHead := <-newTxPoolHeadChan1
-     if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Create list of 10 successive transactions to build block A on vm1
-     // and to be split into two separate blocks on VM2
-     txs := make([]*types.Transaction, 10)
-     for i := 0; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-
-     var errs []error
-
-     // Add the remote transactions, build the block, and set VM1's preference for block A
-     errs = vm1.chain.AddRemoteTxsSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-         }
-     }
- }
-
- <-issuer1

```



```

- vm1BlkB, err := vm1.BuildBlock()
- if err != nil {
-     t.Fatal(err)
- }
-
- if err := vm1BlkB.Verify(); err != nil {
-     t.Fatal(err)
- }
-
- if status := vm1BlkB.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
- }
-
- if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
-     t.Fatal(err)
- }
-
- vm1.chain.BlockChain().GetVMConfig().AllowUnfinalizedQueries = true
-
- blkBHeight := vm1BlkB.Height()
- blkBHash := vm1BlkB.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
- if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkBHash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkBHash.Hex(), b.Hash().Hex())
- }
-
- errs = vm2.chain.AddRemoteTxSync(txs[0:5])
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-     }
- }
-
- <-issuer2
- vm2BlkC, err := vm2.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build BlkC on VM2: %s", err)
- }
-
- if err := vm2BlkC.Verify(); err != nil {
-     t.Fatalf("BlkC failed verification on VM2: %s", err)
- }
-
- if status := vm2BlkC.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block C to be %, but found %s", choices.Processing, status)
- }
-
- if err := vm2.SetPreference(vm2BlkC.ID()); err != nil {
-     t.Fatal(err)
- }
-
- newHead = <-newTxPoolHeadChan2
- if newHead.Head.Hash() != common.Hash(vm2BlkC.ID()) {
-     t.Fatalf("Expected new block to match")
- }
-
- errs = vm2.chain.AddRemoteTxSync(txs[5:])
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-     }
- }
-
- <-issuer2
- vm2BlkD, err := vm2.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build BlkD on VM2: %s", err)
- }
-
- // Parse blocks produced in vm2
- vm1BlkC, err := vm1.ParseBlock(vm2BlkC.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }
- blkChash := vm1BlkC.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
-
- vm1BlkD, err := vm1.ParseBlock(vm2BlkD.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }
- blkDHeight := vm1BlkD.Height()
- blkDHash := vm1BlkD.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
-
- // Should be no-ops
- if err := vm1BlkC.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
- if err := vm1BlkD.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
- if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkBHash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkBHash.Hex(), b.Hash().Hex())
- }
- if b := vm1.chain.GetBlockByNumber(blkDHeight); b != nil {
-     t.Fatalf("expected block at %d to be nil but got %s", blkDHeight, b.Hash().Hex())
- }
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkBHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkBHash.Hex(), b.Hash().Hex())
- }
-
- // Should still be no-ops on re-verify
- if err := vm1BlkC.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
- if err := vm1BlkD.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
- if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkBHash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkBHash.Hex(), b.Hash().Hex())
- }
- if b := vm1.chain.GetBlockByNumber(blkDHeight); b != nil {
-     t.Fatalf("expected block at %d to be nil but got %s", blkDHeight, b.Hash().Hex())
- }
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkBHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkBHash.Hex(), b.Hash().Hex())
- }
-
- // Should be queryable after setting preference to side chain
- if err := vm1.SetPreference(vm1BlkD.ID()); err != nil {
-     t.Fatal(err)
- }
-
- if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkChash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkChash.Hex(), b.Hash().Hex())
- }
- if b := vm1.chain.GetBlockByNumber(blkDHeight); b.Hash() != blkDHash {
-     t.Fatalf("expected block at %d to have hash %s but got %s", blkDHeight, blkDHash.Hex(), b.Hash().Hex())
- }
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkDHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkDHash.Hex(), b.Hash().Hex())
- }
-
- // Attempt to accept out of order
- if err := vm1BlkD.Accept(); !strings.Contains(err.Error(), "expected accepted block to have parent") {

```

```

-         t.Fatalf("unexpected error when accepting out of order block: %s", err)
-     }
-
-     // Accept in order
-     if err := vm1BlkC.Accept(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-     if err := vm1BlkD.Accept(); err != nil {
-         t.Fatalf("Block failed acceptance on VM1: %s", err)
-     }
-
-     // Ensure queryable after accepting
-     if b := vm1.chain.GetBlockByNumber(blkBHeight); b.Hash() != blkCHash {
-         t.Fatalf("expected block at %d to have hash %s but got %s", blkBHeight, blkCHash.Hex(), b.Hash().Hex())
-     }
-     if b := vm1.chain.GetBlockByNumber(blkDHeight); b.Hash() != blkDHash {
-         t.Fatalf("expected block at %d to have hash %s but got %s", blkDHeight, blkDHash.Hex(), b.Hash().Hex())
-     }
-     if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkDHash {
-         t.Fatalf("expected current block to have hash %s but got %s", blkDHash.Hex(), b.Hash().Hex())
-     }
- }
-
- // Regression test to ensure that a VM that prefers block B is able to parse
- // block C but unable to parse block D because it names B as an uncle, which
- // are not supported.
- // A
- // / \
- // B   C
- // |
- // D
- func TestUncleBlock(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-     issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
-     newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer1
-
-     vm1BlkA, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := vm1BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if status := vm1BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-     if err := vm2BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-     if status := vm2BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
-     }
-     if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkA.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-
-     if err := vm2BlkA.Accept(); err != nil {
-         t.Fatalf("VM2 failed to accept block: %s", err)
-     }
-
-     newHead := <-newTxPoolHeadChan1
-     if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     txs := make([]*types.Transaction, 10)
-     for i := 0; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-
-     var errs []error
-
-     errs = vm1.chain.AddRemoteTxsSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-         }
-     }

```

```

    }
    <-issuer1

    vm1BlkB, err := vm1.BuildBlock()
    if err != nil {
        t.Fatal(err)
    }

    if err := vm1BlkB.Verify(); err != nil {
        t.Fatal(err)
    }

    if status := vm1BlkB.Status(); status != choices.Processing {
        t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
    }

    if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
        t.Fatal(err)
    }

    errs = vm2.chain.AddRemoteTxSync(txs[0:5])
    for i, err := range errs {
        if err != nil {
            t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
        }
    }

    <-issuer2
    vm2BlkC, err := vm2.BuildBlock()
    if err != nil {
        t.Fatalf("Failed to build BlkC on VM2: %s", err)
    }

    if err := vm2BlkC.Verify(); err != nil {
        t.Fatalf("BlkC failed verification on VM2: %s", err)
    }

    if status := vm2BlkC.Status(); status != choices.Processing {
        t.Fatalf("Expected status of built block C to be %, but found %s", choices.Processing, status)
    }

    if err := vm2.SetPreference(vm2BlkC.ID()); err != nil {
        t.Fatal(err)
    }

    newHead = <-newTxPoolHeadChan2
    if newHead.Head.Hash() != common.Hash(vm2BlkC.ID()) {
        t.Fatalf("Expected new block to match")
    }

    errs = vm2.chain.AddRemoteTxSync(txs[5:10])
    for i, err := range errs {
        if err != nil {
            t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
        }
    }

    <-issuer2
    vm2BlkD, err := vm2.BuildBlock()
    if err != nil {
        t.Fatalf("Failed to build BlkD on VM2: %s", err)
    }

    // Create uncle block from blkD
    blkDEthBlock := vm2BlkD.(*chain.BlockWrapper).Block.(*Block).ethBlock
    uncles := []types.Header{vm1BlkB.(*chain.BlockWrapper).Block.(*Block).ethBlock.Header()}
    uncleBlockHeader := types.CopyHeader(blkDEthBlock.Header())
    uncleBlockHeader.UncleHash = types.CalcUncleHash(uncles)

    uncleEthBlock := types.NewBlock(
        uncleBlockHeader,
        blkDEthBlock.Transactions(),
        uncles,
        nil,
        new(trie.Trie),
        blkDEthBlock.ExtData(),
        false,
    )
    uncleBlock := &Block{
        vm:      vm2,
        ethBlock: uncleEthBlock,
        id:      ids.ID(uncleEthBlock.Hash()),
    }
    if err := uncleBlock.Verify(); !errors.Is(err, errUnclesUnsupported) {
        t.Fatalf("VM2 should have failed with %q but got %q", errUnclesUnsupported, err.Error())
    }
    if _, err := vm1.ParseBlock(vm2BlkC.Bytes()); err != nil {
        t.Fatalf("VM1 errored parsing blkC: %s", err)
    }
    if _, err := vm1.ParseBlock(uncleBlock.Bytes()); !errors.Is(err, errUnclesUnsupported) {
        t.Fatalf("VM1 should have failed with %q but got %q", errUnclesUnsupported, err.Error())
    }
}

-}

-// Regression test to ensure that a VM that is not able to parse a block that
-// contains no transactions.
-func TestEmptyBlock(t *testing.T) {
    importAmount := uint64(1000000000)
    issuer, vm, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
        testShortIDAddr[0]: importAmount,
    })

    defer func() {
        if err := vm.Shutdown(); err != nil {
            t.Fatal(err)
        }
    }()

    key, err := accountKeystore.NewKey(rand.Reader)
    if err != nil {
        t.Fatal(err)
    }

    importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []crypto.PrivateKeySECP256K1R{testKeys[0]})
    if err != nil {
        t.Fatal(err)
    }

    if err := vm.issueTx(importTx, true /*=local*//); err != nil {
        t.Fatal(err)
    }

    <-issuer

    blk, err := vm.BuildBlock()
    if err != nil {
        t.Fatalf("Failed to build block with import transaction: %s", err)
    }

    // Create empty block from blkA
    ethBlock := blk.(*chain.BlockWrapper).Block.(*Block).ethBlock

    emptyEthBlock := types.NewBlock(

```

```

-         types.CopyHeader(ethBlock.Header()),
-         nil,
-         nil,
-         nil,
-         new(trie.Trie),
-         nil,
-         false,
-     )
-
-     if len(emptyEthBlock.ExtData()) != 0 || emptyEthBlock.Header().ExtDataHash != (common.Hash{}) {
-         t.Fatalf("emptyEthBlock should not have any extra data")
-     }
-
-     emptyBlock := &Block{
-         vm:      vm,
-         ethBlock: emptyEthBlock,
-         id:       ids.ID(emptyEthBlock.Hash()),
-     }
-
-     if _, err := vm.ParseBlock(emptyBlock.Bytes()); !errors.Is(err, errEmptyBlock) {
-         t.Fatalf("VM should have failed with errEmptyBlock but got %s", err.Error())
-     }
-     if err := emptyBlock.Verify(); !errors.Is(err, errEmptyBlock) {
-         t.Fatalf("block should have failed verification with errEmptyBlock but got %s", err.Error())
-     }
- }
-
- // Regression test to ensure that a VM that verifies block B, C, then
- // D (preferring block B) reorgs when C and then D are accepted.
- //   A
- //   / \
- //  B   C
- //   |
- //   D
- func TestAcceptReorg(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer1, vm1, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-     issuer2, vm2, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm1.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-
-         if err := vm2.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan1 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm1.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan1)
-     newTxPoolHeadChan2 := make(chan core.NewTxPoolReorgEvent, 1)
-     vm2.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan2)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-     importTx, err := vm1.newImportTx(vm1.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R(testKeys[0]))
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer1
-
-     vm1BlkA, err := vm1.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := vm1BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM1: %s", err)
-     }
-
-     if status := vm1BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm1.SetPreference(vm1BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     vm2BlkA, err := vm2.ParseBlock(vm1BlkA.Bytes())
-     if err != nil {
-         t.Fatalf("Unexpected error parsing block from vm2: %s", err)
-     }
-     if err := vm2BlkA.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM2: %s", err)
-     }
-
-     if status := vm2BlkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of block on VM2 to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm2.SetPreference(vm2BlkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm1BlkA.Accept(); err != nil {
-         t.Fatalf("VM1 failed to accept block: %s", err)
-     }
-     if err := vm2BlkA.Accept(); err != nil {
-         t.Fatalf("VM2 failed to accept block: %s", err)
-     }
-
-     newHead := <-newTxPoolHeadChan1
-     if newHead.Head.Hash() != common.Hash(vm1BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-     newHead = <-newTxPoolHeadChan2
-     if newHead.Head.Hash() != common.Hash(vm2BlkA.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Create list of 10 successive transactions to build block A on vm1
-     // and to be split into two separate blocks on VM2
-     txs := make([]types.Transaction, 10)
-     for i := 0; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm1.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-
-     // Add the remote transactions, build the block, and set VM1's preference
-     // for block B

```

```

-     errs := vm1.chain.AddRemoteTxSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add transaction to VM1 at index %d: %s", i, err)
-         }
-     }
- }
-
- <-issuer1
-
- vm1BlkB, err := vm1.BuildBlock()
- if err != nil {
-     t.Fatal(err)
- }
-
- if err := vm1BlkB.Verify(); err != nil {
-     t.Fatal(err)
- }
-
- if status := vm1BlkB.Status(); status != choices.Processing {
-     t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
- }
-
- if err := vm1.SetPreference(vm1BlkB.ID()); err != nil {
-     t.Fatal(err)
- }
-
- errs = vm2.chain.AddRemoteTxSync(txs[0:5])
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-     }
- }
-
- <-issuer2
-
- vm2BlkC, err := vm2.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build BlkC on VM2: %s", err)
- }
-
- if err := vm2BlkC.Verify(); err != nil {
-     t.Fatalf("BlkC failed verification on VM2: %s", err)
- }
-
- if err := vm2.SetPreference(vm2BlkC.ID()); err != nil {
-     t.Fatal(err)
- }
-
- newHead = <-newTxPoolHeadChan2
- if newHead.Head.Hash() != common.Hash(vm2BlkC.ID()) {
-     t.Fatalf("Expected new block to match")
- }
-
- errs = vm2.chain.AddRemoteTxSync(txs[5:])
- for i, err := range errs {
-     if err != nil {
-         t.Fatalf("Failed to add transaction to VM2 at index %d: %s", i, err)
-     }
- }
-
- <-issuer2
-
- vm2BlkD, err := vm2.BuildBlock()
- if err != nil {
-     t.Fatalf("Failed to build BlkD on VM2: %s", err)
- }
-
- // Parse blocks produced in vm2
- vm1BlkC, err := vm1.ParseBlock(vm2BlkC.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }
-
- vm1BlkD, err := vm1.ParseBlock(vm2BlkD.Bytes())
- if err != nil {
-     t.Fatalf("Unexpected error parsing block from vm2: %s", err)
- }
-
- if err := vm1BlkC.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
-
- if err := vm1BlkD.Verify(); err != nil {
-     t.Fatalf("Block failed verification on VM1: %s", err)
- }
-
- blkBHash := vm1BlkB.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkBHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkBHash.Hex(), b.Hash().Hex())
- }
-
- if err := vm1BlkC.Accept(); err != nil {
-     t.Fatal(err)
- }
-
- blkCHash := vm1BlkC.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkCHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkCHash.Hex(), b.Hash().Hex())
- }
-
- if err := vm1BlkB.Reject(); err != nil {
-     t.Fatal(err)
- }
-
- if err := vm1BlkD.Accept(); err != nil {
-     t.Fatal(err)
- }
-
- blkDHash := vm1BlkD.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
- if b := vm1.chain.BlockChain().CurrentBlock(); b.Hash() != blkDHash {
-     t.Fatalf("expected current block to have hash %s but got %s", blkDHash.Hex(), b.Hash().Hex())
- }
- }
-}
-
- func TestFutureBlock(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer, vm, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
- }

```

```

-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blkA, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     // Create empty block from blkA
-     blkAEthBlock := blkA.(*chain.BlockWrapper).Block.(*Block).ethBlock
-
-     modifiedHeader := types.CopyHeader(blkAEthBlock.Header())
-     // Set the VM's clock to the time of the produced block
-     vm.clock.Set(time.Unix(int64(modifiedHeader.Time), 0))
-     // Set the modified time to exceed the allowed future time
-     modifiedTime := modifiedHeader.Time + uint64(maxFutureBlockTime.Seconds()+1)
-     modifiedHeader.Time = modifiedTime
-     modifiedBlock := types.NewBlock(
-         modifiedHeader,
-         nil,
-         nil,
-         nil,
-         new(trie.Trie),
-         blkAEthBlock.ExtData(),
-         false,
-     )
-
-     futureBlock := &Block{
-         vm:      vm,
-         ethBlock: modifiedBlock,
-         id:       ids.ID(modifiedBlock.Hash()),
-     }
-
-     if err := futureBlock.Verify(); err == nil {
-         t.Fatal("Future block should have failed verification due to block timestamp too far in the future")
-     } else if !strings.Contains(err.Error(), "block timestamp is too far in the future") {
-         t.Fatalf("Expected error to be block timestamp too far in the future but found %s", err)
-     }
- }
-
- })
-
- // Regression test to ensure we can build blocks if we are starting with the
- // Apricot Phase 1 ruleset in genesis.
- func TestBuildApricotPhase1Block(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer, vm, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase1, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan := make(chan core.NewTxPoolReorgEvent, 1)
-     vm.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan)
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     newHead := <-newTxPoolHeadChan
-     if newHead.Head.Hash() != common.Hash(blk.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     txs := make([]*types.Transaction, 10)
-     for i := 0; i < 5; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-     for i := 5; i < 10; i++ {
-         tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.ApricotPhase1MinGasPrice), nil)
-         signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
-         if err != nil {
-             t.Fatal(err)
-         }
-         txs[i] = signedTx
-     }
-     errs := vm.chain.AddRemoteTxsSync(txs)
-     for i, err := range errs {
-         if err != nil {
-             t.Fatalf("Failed to add tx at index %d: %s", i, err)
-         }
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {

```

```

-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Accepted {
-         t.Fatalf("Expected status of accepted block to be %, but found %s", choices.Accepted, status)
-     }
-
-     lastAcceptedID, err := vm.LastAccepted()
-     if err != nil {
-         t.Fatal(err)
-     }
-     if lastAcceptedID != blk.ID() {
-         t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk.ID(), lastAcceptedID)
-     }
-
-     // Confirm all txs are present
-     ethBlkTxs := vm.chain.GetBlockByNumber(2).Transactions()
-     for i, tx := range txs {
-         if len(ethBlkTxs) <= i {
-             t.Fatalf("missing transactions expected: %d but found: %d", len(txs), len(ethBlkTxs))
-         }
-         if ethBlkTxs[i].Hash() != tx.Hash() {
-             t.Fatalf("expected tx at index %d to have hash: %x but has: %x", i, txs[i].Hash(), tx.Hash())
-         }
-     }
- }
-
-}
-
-func TestLastAcceptedBlockNumberAllow(t *testing.T) {
-     importAmount := uint64(1000000000)
-     issuer, vm, _, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase0, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: importAmount,
-     })
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     key, err := accountKeystore.NewKey(rand.Reader)
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatalf("Failed to build block with import transaction: %s", err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatalf("Block failed verification on VM: %s", err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     blkHeight := blk.Height()
-     blkHash := blk.(*chain.BlockWrapper).Block.(*Block).ethBlock.Hash()
-
-     vm.chain.BlockChain().GetVMConfig().AllowUnfinalizedQueries = true
-
-     ctx := context.Background()
-     b, err := vm.chain.APIBackend().BlockByNumber(ctx, rpc.BlockNumber(blkHeight))
-     if err != nil {
-         t.Fatal(err)
-     }
-     if b.Hash() != blkHash {
-         t.Fatalf("expected block at %d to have hash %s but got %s", blkHeight, blkHash.Hex(), b.Hash().Hex())
-     }
-
-     vm.chain.BlockChain().GetVMConfig().AllowUnfinalizedQueries = false
-
-     _, err = vm.chain.APIBackend().BlockByNumber(ctx, rpc.BlockNumber(blkHeight))
-     if !errors.Is(err, eth.ErrUnfinalizedData) {
-         t.Fatalf("expected ErrUnfinalizedData but got %s", err.Error())
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatalf("VM failed to accept block: %s", err)
-     }
-
-     if b := vm.chain.GetBlockByNumber(blkHeight); b.Hash() != blkHash {
-         t.Fatalf("expected block at %d to have hash %s but got %s", blkHeight, blkHash.Hex(), b.Hash().Hex())
-     }
- }
-
-}
-
-// Builds [blkA] with a virtuous import transaction and [blkB] with a separate import transaction
-// that does not conflict. Accepts [blkB] and rejects [blkA], then asserts that the virtuous atomic
-// transaction in [blkA] is correctly re-issued into the atomic transaction mempool.
-func TestReissueAtomicTx(t *testing.T) {
-     issuer, vm, _, _, _ := GenesisVMWithUTXOs(t, true, genesisJSONApricotPhase1, "", "", map[ids.ShortID]uint64{
-         testShortIDAddrs[0]: 10000000,
-         testShortIDAddrs[1]: 10000000,
-     })
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     genesisBlkID, err := vm.LastAccepted()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     importTx, err := vm.newImportTx(vm.ctx.XChainID, testEthAddrs[0], initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
-     if err != nil {
-         t.Fatal(err)
-     }

```

```

-     if err := vm.issueTx(importTx, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blkA, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blkA.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blkA.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     // SetPreference to parent before rejecting (will rollback state to genesis
-     // so that atomic transaction can be reissued, otherwise current block will
-     // conflict with UTX0 to be reissued)
-     if err := vm.SetPreference(genesisBlkID); err != nil {
-         t.Fatal(err)
-     }
-
-     // Rejecting [blkA] should cause [importTx] to be re-issued into the mempool.
-     if err := blkA.Reject(); err != nil {
-         t.Fatal(err)
-     }
-
-     // Sleep for a minimum of two seconds to ensure that [blkB] will have a different timestamp
-     // than [blkA] so that the block will be unique. This is necessary since we have marked [blkA]
-     // as Rejected.
-     time.Sleep(2 * time.Second)
-     <-issuer
-     blkB, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if blkB.Height() != blkA.Height() {
-         t.Fatalf("Expected blkB (%d) to have the same height as blkA (%d)", blkB.Height(), blkA.Height())
-     }
-
-     if status := blkA.Status(); status != choices.Rejected {
-         t.Fatalf("Expected status of blkA to be %, but found %s", choices.Rejected, status)
-     }
-
-     if status := blkB.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of blkB to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := blkB.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blkB.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of blkC to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blkB.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blkB.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blkB.Status(); status != choices.Accepted {
-         t.Fatalf("Expected status of accepted block to be %, but found %s", choices.Accepted, status)
-     }
-
-     if lastAcceptedID, err := vm.LastAccepted(); err != nil {
-         t.Fatal(err)
-     } else if lastAcceptedID != blkB.ID() {
-         t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blkB.ID(), lastAcceptedID)
-     }
-
-     // Check that [importTx] has been indexed correctly after [blkB] is accepted.
-     _, height, err := vm.getAcceptedAtomicTx(importTx.ID())
-     if err != nil {
-         t.Fatal(err)
-     } else if height != blkB.Height() {
-         t.Fatalf("Expected indexed height of import tx to be %d, but found %d", blkB.Height(), height)
-     }
- }
-
- func TestAtomicTxFailsEVMStateTransferBuildBlock(t *testing.T) {
-     issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJSONApricotPhase1, "", "")
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     exportTxS := createExportTxOptions(t, vm, issuer, sharedMemory)
-     exportTx1, exportTx2 := exportTxS[0], exportTxS[1]
-
-     if err := vm.issueTx(exportTx1, true /*=local*/); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-     exportBlk1, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := exportBlk1.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.SetPreference(exportBlk1.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := vm.issueTx(exportTx2, true /*=local*/); err == nil {
-         t.Fatalf("Should have failed to issue due to an invalid export tx")
-     }
-
-     if err := vm.mempool.AddTx(exportTx2); err == nil {
-         t.Fatalf("Should have failed to add because conflicting")
-     }
-
-     // Manually add transaction to mempool to bypass validation
-     if err := vm.mempool.ForceAddTx(exportTx2); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     _, err = vm.BuildBlock()
-     if err == nil {
-         t.Fatalf("BuildBlock should have returned an error due to invalid export transaction")
-     }
- }
-
- }

```



```

- func TestBuildInvalidBlockHead(t *testing.T) {
-     issuer, vm, _, _, _ := GenesisVM(t, true, genesisJSONApricotPhase0, "", "")
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     key0 := testKeys[0]
-     addr0 := key0.PublicKey().Address()
-
-     // Create the transaction
-     utx := &UnsignedImportTx{
-         NetworkID:    vm.ctx.NetworkID,
-         BlockchainID:   vm.ctx.ChainID,
-         Outs: []EVMOutput{
-             {
-                 Address: common.Address(addr0),
-                 Amount:  1 * units.Avox,
-                 AssetID: vm.ctx.AVAXAssetID,
-             },
-         },
-         ImportedInputs: []*avax.TransferableInput{
-             {
-                 Asset: avax.Asset{ID: vm.ctx.AVAXAssetID},
-                 In: &secp256k1fx.TransferInput{
-                     Amt: 1 * units.Avox,
-                     Input: secp256k1fx.Input{
-                         SigIndices: []uint32{0},
-                     },
-                 },
-             },
-         },
-         SourceChain: vm.ctx.XChainID,
-     }
-     tx := &Tx{UnsignedAtomicTx: utx}
-     if err := tx.Sign(vm.codec, [][*crypto.PrivateKeySECP256K1R]{key0}); err != nil {
-         t.Fatal(err)
-     }
-
-     currentBlock := vm.chain.BlockChain().CurrentBlock()
-
-     // Verify that the transaction fails verification when attempting to issue
-     // it into the atomic mempool.
-     if err := vm.issueTx(tx, true /*=local*/); err == nil {
-         t.Fatal("Should have failed to issue invalid transaction")
-     }
-
-     // Force issue the transaction directly to the mempool
-     if err := vm.mempool.AddTx(tx); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     if _, err := vm.BuildBlock(); err == nil {
-         t.Fatalf("Unexpectedly created a block")
-     }
-
-     newCurrentBlock := vm.chain.BlockChain().CurrentBlock()
-
-     if currentBlock.Hash() != newCurrentBlock.Hash() {
-         t.Fatal("current block changed")
-     }
- }
-
- func TestConfigureLogLevel(t *testing.T) {
-     configTests := []struct {
-         name          string
-         logConfig      string
-         genesisJSON    string
-         upgradeJSON    string
-         expectedErr    string
-     }{
-         {
-             name:          "Log level info",
-             logConfig:      "{\"log-level\": \"info\"}",
-             genesisJSON:    genesisJSONApricotPhase2,
-             upgradeJSON:    "",
-             expectedErr:    "",
-         },
-         {
-             name:          "Invalid log level",
-             logConfig:      "{\"log-level\": \"cchain\"}",
-             genesisJSON:    genesisJSONApricotPhase3,
-             upgradeJSON:    "",
-             expectedErr:    "failed to initialize logger due to",
-         },
-     }
-     for _, test := range configTests {
-         t.Run(test.name, func(t *testing.T) {
-             vm := &VM{
-                 ctx, dbManager, genesisBytes, issuer, _ := setupGenesis(t, test.genesisJSON)
-                 appSender := &engCommon.SenderTest{
-                     appSender.CantSendAppGossip = true
-                     appSender.SendAppGossipF = func([]byte) error { return nil }
-                 }
-                 err := vm.Initialize(
-                     ctx,
-                     dbManager,
-                     genesisBytes,
-                     []byte(""),
-                     []byte(test.logConfig),
-                     issuer,
-                     []*engCommon.Fx{
-                         appSender,
-                     },
-                 )
-                 if len(test.expectedErr) == 0 && err != nil {
-                     t.Fatal(err)
-                 } else if len(test.expectedErr) > 0 {
-                     if err == nil {
-                         t.Fatalf("initialize should have failed due to %s", test.expectedErr)
-                     } else if !strings.Contains(err.Error(), test.expectedErr) {
-                         t.Fatalf("Expected initialize to fail due to %s, but failed with %s", test.expectedErr, err.Error())
-                     }
-                 }
-             }
-
-             // If the VM was not initialized, do not attempt to shut it down
-             if err == nil {
-                 shutdownChan := make(chan error, 1)
-                 shutdownFunc := func() {
-                     err := vm.Shutdown()
-                     shutdownChan <- err
-                 }
-                 go shutdownFunc()
-
-                 // If the VM was not initialized, do not attempt to shut it down
-                 if err == nil {
-                     shutdownChan := make(chan error, 1)
-                     shutdownFunc := func() {
-                         err := vm.Shutdown()
-                         shutdownChan <- err
-                     }
-                     go shutdownFunc()
-
-                     shutdownTimeout := 50 * time.Millisecond
-                     ticker := time.NewTicker(shutdownTimeout)
-
-                     @@ -2922,248 +497,3 @@ func TestConfigureLogLevel(t *testing.T) {

```

```

    })
}
}

-// Regression test to ensure we can build blocks if we are starting with the
-// Apricot Phase 4 ruleset in genesis.
-func TestBuildApricotPhase4Block(t *testing.T) {
    issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJSONApricotPhase4, "", "")

    defer func() {
        if err := vm.Shutdown(); err != nil {
            t.Fatal(err)
        }
    }()

    newTxPoolHeadChan := make(chan core.NewTxPoolReorgEvent, 1)
    vm.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan)

    key, err := accountKeystore.NewKey(rand.Reader)
    if err != nil {
        t.Fatal(err)
    }

    importAmount := uint64(1000000000)
    utxoID := avax.UTXOID{TxID: ids.GenerateTestID()}

    utxo := &avax.UTXO{
        UTXOID: utxoID,
        Asset:  avax.Asset{ID: vm.ctx.AVAXAssetID},
        Out: &secp256k1fx.TransferOutput{
            Amt: importAmount,
            OutputOwners: secp256k1fx.OutputOwners{
                Threshold: 1,
                Addrs:      []ids.ShortID{testKeys[0].PublicKey().Address()},
            },
        },
    },
}
    utxoBytes, err := vm.codec.Marshal(codecVersion, utxo)
    if err != nil {
        t.Fatal(err)
    }

    xChainSharedMemory := sharedMemory.NewSharedMemory(vm.ctx.XChainID)
    inputID := utxo.InputID()
    if err := xChainSharedMemory.Apply(map[ids.ID]*atomic.Requests{vm.ctx.ChainID: {PutRequests: []*atomic.Element{{
        Key: inputID[:],
        Value: utxoBytes,
        Traits: [][]byte{
            testKeys[0].PublicKey().Address().Bytes(),
        }},
    }}}}); err != nil {
        t.Fatal(err)
    }

    importTx, err := vm.newImportTx(vm.ctx.XChainID, key.Address, initialBaseFee, []*crypto.PrivateKeySECP256K1R{testKeys[0]})
    if err != nil {
        t.Fatal(err)
    }

    if err := vm.issueTx(importTx, true /*=local*); err != nil {
        t.Fatal(err)
    }

    <-issuer

    blk, err := vm.BuildBlock()
    if err != nil {
        t.Fatal(err)
    }

    if err := blk.Verify(); err != nil {
        t.Fatal(err)
    }

    if status := blk.Status(); status != choices.Processing {
        t.Fatalf("Expected status of built block to be %s, but found %s", choices.Processing, status)
    }

    if err := vm.SetPreference(blk.ID()); err != nil {
        t.Fatal(err)
    }

    if err := blk.Accept(); err != nil {
        t.Fatal(err)
    }

    ethBlk := blk.(*chain.BlockWrapper).Block.(*Block).ethBlock
    if eBlockGasCost := ethBlk.BlockGasCost(); eBlockGasCost == nil || eBlockGasCost.Cmp(common.Big0) != 0 {
        t.Fatalf("expected blockGasCost to be 0 but got %d", eBlockGasCost)
    }
    if eExtDataGasUsed := ethBlk.ExtDataGasUsed(); eExtDataGasUsed == nil || eExtDataGasUsed.Cmp(big.NewInt(1230)) != 0 {
        t.Fatalf("expected extDataGasUsed to be 1000 but got %d", eExtDataGasUsed)
    }
    minRequiredTip, err := dummy.MinRequiredTip(vm.chainConfig, ethBlk.Header())
    if err != nil {
        t.Fatal(err)
    }
    if minRequiredTip == nil || minRequiredTip.Cmp(common.Big0) != 0 {
        t.Fatalf("expected minRequiredTip to be 0 but got %d", minRequiredTip)
    }

    newHead := <-newTxPoolHeadChan
    if newHead.Head.Hash() != common.Hash(blk.ID()) {
        t.Fatalf("Expected new block to match")
    }

    txs := make([]*types.Transaction, 10)
    for i := 0; i < 5; i++ {
        tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.LaunchMinGasPrice), nil)
        signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
        if err != nil {
            t.Fatal(err)
        }
        txs[i] = signedTx
    }
    for i := 5; i < 10; i++ {
        tx := types.NewTransaction(uint64(i), key.Address, big.NewInt(10), 21000, big.NewInt(params.ApricotPhase1MinGasPrice), nil)
        signedTx, err := types.SignTx(tx, types.NewEIP155Signer(vm.chainID), key.PrivateKey)
        if err != nil {
            t.Fatal(err)
        }
        txs[i] = signedTx
    }
    errs := vm.chain.AddRemoteTxs(txs)
    for i, err := range errs {
        if err != nil {
            t.Fatalf("Failed to add tx at index %d: %s", i, err)
        }
    }

    <-issuer

    blk, err := vm.BuildBlock()
    if err != nil {

```

```

-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     ethBlk = blk.(*chain.BlockWrapper).Block.(*Block).ethBlock
-     if ethBlk.BlockGasCost() == nil || ethBlk.BlockGasCost().Cmp(big.NewInt(100)) < 0 {
-         t.Fatalf("expected blockGasCost to be at least 100 but got %d", ethBlk.BlockGasCost())
-     }
-     if ethBlk.ExtDataGasUsed() == nil || ethBlk.ExtDataGasUsed().Cmp(common.Big0) != 0 {
-         t.Fatalf("expected extDataGasUsed to be 0 but got %d", ethBlk.ExtDataGasUsed())
-     }
-     minRequiredTip, err = dummy.MinRequiredTip(vm.chainConfig, ethBlk.Header())
-     if err != nil {
-         t.Fatal(err)
-     }
-     if minRequiredTip == nil || minRequiredTip.Cmp(big.NewInt(0.05*params.GWei)) < 0 {
-         t.Fatalf("expected minRequiredTip to be at least 0.05 gwei but got %d", minRequiredTip)
-     }
-
-     if status := blk.Status(); status != choices.Accepted {
-         t.Fatalf("Expected status of accepted block to be %, but found %s", choices.Accepted, status)
-     }
-
-     lastAcceptedID, err := vm.LastAccepted()
-     if err != nil {
-         t.Fatal(err)
-     }
-     if lastAcceptedID != blk.ID() {
-         t.Fatalf("Expected last accepted blockID to be the accepted block: %s, but found %s", blk.ID(), lastAcceptedID)
-     }
-
-     // Confirm all txs are present
-     ethBlkTxs := vm.chain.GetBlockByNumber(2).Transactions()
-     for i, tx := range txs {
-         if len(ethBlkTxs) <= i {
-             t.Fatalf("missing transactions expected: %d but found: %d", len(txs), len(ethBlkTxs))
-         }
-         if ethBlkTxs[i].Hash() != tx.Hash() {
-             t.Fatalf("expected tx at index %d to have hash: %x but has: %x", i, txs[i].Hash(), tx.Hash())
-         }
-     }
- }
- }
-
- // This is a regression test to ensure that if two consecutive atomic transactions fail verification
- // in onFinalizeAndAssemble it will not cause a panic due to calling RevertToSnapshot(revid) on the
- // same revision ID twice.
- func TestConsecutiveAtomicTransactionsRevertSnapshot(t *testing.T) {
-     issuer, vm, _, sharedMemory, _ := GenesisVM(t, true, genesisJSONApricotPhase1, "", "")
-
-     defer func() {
-         if err := vm.Shutdown(); err != nil {
-             t.Fatal(err)
-         }
-     }()
-
-     newTxPoolHeadChan := make(chan core.NewTxPoolReorgEvent, 1)
-     vm.chain.GetTxPool().SubscribeNewReorgEvent(newTxPoolHeadChan)
-
-     // Create three conflicting import transactions
-     importTxs := createImportTxOptions(t, vm, sharedMemory)
-
-     // Issue the first import transaction, build, and accept the block.
-     if err := vm.issueTx(importTxs[0], true); err != nil {
-         t.Fatal(err)
-     }
-
-     <-issuer
-
-     blk, err := vm.BuildBlock()
-     if err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Verify(); err != nil {
-         t.Fatal(err)
-     }
-
-     if status := blk.Status(); status != choices.Processing {
-         t.Fatalf("Expected status of built block to be %, but found %s", choices.Processing, status)
-     }
-
-     if err := vm.SetPreference(blk.ID()); err != nil {
-         t.Fatal(err)
-     }
-
-     if err := blk.Accept(); err != nil {
-         t.Fatal(err)
-     }
-
-     newHead := <-newTxPoolHeadChan
-     if newHead.Head.Hash() != common.Hash(blk.ID()) {
-         t.Fatalf("Expected new block to match")
-     }
-
-     // Add the two conflicting transactions directly to the mempool, so that two consecutive transactions
-     // will fail verification when build block is called.
-     vm.mempool.AddTx(importTxs[1])
-     vm.mempool.AddTx(importTxs[2])
-
-     if _, err := vm.BuildBlock(); err == nil {
-         t.Fatal("Expected build block to fail due to empty block")
-     }
- }
- }
-
diff --git a/plugin/main.go b/plugin/main.go
index 96cd428f..44588435 100644
--- a/plugin/main.go
+++ b/plugin/main.go
@@ -9,9 +9,10 @@ import (
     "github.com/hashicorp/go-plugin"
 
     "github.com/ava-labs/avalanchego/vms/rpcchainvm"
+    "github.com/flare-foundation/flare/utlils/ulimit"
+    "github.com/flare-foundation/flare/vms/rpcchainvm"
 
     "github.com/ava-labs/coreth/plugin/evm"
+    "github.com/flare-foundation/coreth/plugin/evm"
 )
 
 func main() {
@@ -24,6 +25,10 @@ func main() {
     @ -24,6 +25,10 @@ func main() {
         fmt.Println(evm.Version)
         os.Exit(0)
     }
}

```

```

    }
    if err := ulimit.Set(ulimit.DefaultFDLimit); err != nil {
        fmt.Printf("failed to set fd limit correctly due to: %s", err)
        os.Exit(1)
    }
    plugin.Serve(&plugin.ServeConfig{
        HandshakeConfig: rpcchainvm.Handshake,
        Plugins: map[string]plugin.Plugin{
diff --git a/rpc/errors.go b/rpc/errors.go
index a43fa9bd..d8b597c1 100644
--- a/rpc/errors.go
+++ b/rpc/errors.go
@@ -64,6 +64,7 @@ var (
     _ Error = new(invalidRequestError)
     _ Error = new(invalidMessageError)
     _ Error = new(invalidParamsError)
+    _ Error = new(CustomError)
 )

const defaultErrorCode = -32000
@@ -111,3 +112,12 @@ type invalidParamsError struct{ message string }
func (e *invalidParamsError) ErrorCode() int { return -32002 }

func (e *invalidParamsError) Error() string { return e.message }
+
+type CustomError struct {
+    Code      int
+    ValidationError string
+}
+
+func (e *CustomError) ErrorCode() int { return e.Code }
+
+func (e *CustomError) Error() string { return e.ValidationError }
diff --git a/rpc/handler.go b/rpc/handler.go
index abafa663..82d9d664 100644
--- a/rpc/handler.go
+++ b/rpc/handler.go
@@ -36,6 +36,7 @@ import (
    "time"

    "github.com/ethereum/go-ethereum/log"
    "github.com/ethereum/go-ethereum/metrics"
    "golang.org/x/time/rate"
)

@@ -83,6 +84,8 @@ type handler struct {
    type callProc struct {
        ctx      context.Context
        notifiers []*Notifier
        callStart time.Time
        procStart time.Time
    }

    func newHandler(connCtx context.Context, conn jsonWriter, idgen func() ID, reg *serviceRegistry) *handler {
@@ -262,15 +265,15 @@ func (h *handler) awaitLimit(ctx context.Context) {
    timer.Stop()
}

-// consumeLimit removes the time since [startTime] from the rate limiter. It is
+// consumeLimit removes the time since [procStart] from the rate limiter. It is
// assumed that the rate limiter is full.
-func (h *handler) consumeLimit(startTime time.Time) {
+func (h *handler) consumeLimit(procStart time.Time) {
    if h.limiter == nil {
        return
    }

    stopTime := time.Now()
    processingTime := stopTime.Sub(startTime)
+    processingTime := stopTime.Sub(procStart)
    if processingTime > h.deadlineContext {
        processingTime = h.deadlineContext
    }
@@ -293,11 +296,17 @@ func (h *handler) startCallProc(fn func(*callProc)) {
    }
    defer h.callWG.Done()

    // Capture the time before we await for processing
    callStart := time.Now()
    h.awaitLimit(ctx)
    startTime := time.Now()

    // If we are not limiting CPU, [procStart] will be identical to
    // [callStart]
    procStart := time.Now()
    defer cancel()
    fn(&callProc{ctx: ctx})
    h.consumeLimit(startTime)

    fn(&callProc{ctx: ctx, callStart: callStart, procStart: procStart})
    h.consumeLimit(procStart)
}
if h.limiter == nil {
    go callFn()
@@ -309,7 +318,7 @@ func (h *handler) startCallProc(fn func(*callProc)) {
// handleImmediate executes non-call messages. It returns false if the message is a
// call or requires a reply.
func (h *handler) handleImmediate(msg *jsonrpcMessage) bool {
-    start := time.Now()
+    execStart := time.Now()
    switch {
    case msg.isNotification():
        if strings.HasSuffix(msg.Method, notificationMethodSuffix) {
@@ -319,7 +328,7 @@ func (h *handler) handleImmediate(msg *jsonrpcMessage) bool {
@@ -319,7 +328,7 @@ func (h *handler) handleImmediate(msg *jsonrpcMessage) bool {
        return false
    case msg.isResponse():
        h.handleResponse(msg)
-        h.Log.Trace("Handled RPC response", "reqid", idForLog(msg.ID), "t", time.Since(start))
+        h.Log.Trace("Handled RPC response", "reqid", idForLog(msg.ID), "duration", time.Since(execStart))
        return true
    default:
        return false
}
@@ -367,16 +376,16 @@ func (h *handler) handleResponse(msg *jsonrpcMessage) {

// handleCallMsg executes a call message and returns the answer.
func (h *handler) handleCallMsg(ctx *callProc, msg *jsonrpcMessage) *jsonrpcMessage {
-    start := time.Now()
    // [callStart] is the time the message was enqueued for handler processing
    callStart := ctx.callStart
    // [procStart] is the time the message cleared the [limiter] and began to be
    // processed by the handler
    procStart := ctx.procStart
    // [execStart] is the time the message began to be executed by the handler
    //
    // Note: This can be different than the executionStart in [startCallProc] as
    // the goroutine that handles execution may not be executed right away.
    execStart := time.Now()

    switch {
    case msg.isNotification():
        h.handleCall(ctx, msg)
-        h.Log.Debug("Served "+msg.Method, "t", time.Since(start))
+        h.Log.Debug("Served "+msg.Method, "execTime", time.Since(execStart), "procTime", time.Since(procStart), "totalTime", time.Since(callStart))
    }
    return nil
}

```

```

case msg.isCall():
    resp := h.handleCall(ctx, msg)
    var ctx []interface{}
    -   ctx = append(ctx, "reqid", idForLog[msg.ID], "t", time.Since(start))
+   ctx = append(ctx, "reqid", idForLog[msg.ID], "execTime", time.Since(execStart), "procTime", time.Since(procStart), "totalTime", time.Since(callStart))
    if resp.Error != nil {
        ctx = append(ctx, "err", resp.Error.Message)
        if resp.Error.Data != nil {
@@ -425,7 +444,9 @@ func (h *handler) handleCall(cp *callProc, msg *jsonrpcMessage) *jsonrpcMessage
            successfulRequestGauge.Inc(1)
        }
        rpcServingTimer.UpdateSince(start)
    -   newRPCServingTimer(msg.Method, answer.Error == nil).UpdateSince(start)
+   if metrics.EnabledExpensive {
+       newRPCServingTimer(msg.Method, answer.Error == nil).UpdateSince(start)
+   }
    }
    return answer
}

diff --git a/rpc/types.go b/rpc/types.go
index 96015551..4fd6ded5 100644
--- a/rpc/types.go
+++ b/rpc/types.go
@@ -44,6 +44,7 @@ type API struct {
    Version    string    // api version for DApp's
    Service    interface{} // receiver instance which holds the methods
    Public     bool      // indication if the methods must be considered safe for public use
+   Name       string    // Name of the API
}

// ServerCodec implements reading, parsing and writing RPC messages for the server side of
diff --git a/scripts/build.sh b/scripts/build.sh
index 98c19678..a4250631 100755
--- a/scripts/build.sh
+++ b/scripts/build.sh
@@ -4,7 +4,11 @@ set -o errexit
set -o nounset
set -o pipefail

+<<<<<<< HEAD
+# Coreth root directory
+=====
+# Avalanche root directory
+>>>>>>> upstream-v0.8.5-rc.2
+CORETH_PATH=$( cd "${dirname "${BASH_SOURCE[0]}"}" ); cd .. && pwd )

# Load the versions
@@ -27,4 +31,4 @@ coreth_commit=${CORETH_COMMIT:-$( git rev-list -1 HEAD )}

# Build Coreth, which is run as a subprocess
echo "Building Coreth Version: $coreth_version; GitCommit: $coreth_commit"
-go build -ldflags "-X github.com/ava-labs/coreth/plugin/evm.GitCommit=$coreth_commit -X github.com/ava-labs/coreth/plugin/evm.Version=$coreth_version" -o "$binary_path" "plugin/*.go"
+go build -ldflags "-X github.com/flare-foundation/coreth/plugin/evm.GitCommit=$coreth_commit -X github.com/flare-foundation/coreth/plugin/evm.Version=$coreth_version" -o "$binary_path" "plugin/*.go"
diff --git a/scripts/build_image.sh b/scripts/build_image.sh
deleted file mode 100755
index 6de76735..00000000
--- a/scripts/build_image.sh
+++ /dev/null
@@ -1,20 +0,0 @@
-#!/usr/bin/env bash
-
--set -o errexit
--set -o nounset
--set -o pipefail
-
-# Avalanche root directory
-CORETH_PATH=$( cd "${dirname "${BASH_SOURCE[0]}"}" ); cd .. && pwd )
-
-# Load the versions
--source "$CORETH_PATH"/scripts/versions.sh
-
-# Load the constants
--source "$CORETH_PATH"/scripts/constants.sh
-
--echo "Building Docker Image: $dockerhub_repo:$build_image_id based of $avalanche_version"
--docker build -t "$dockerhub_repo:$build_image_id" "$CORETH_PATH" -f "$CORETH_PATH/Dockerfile" \
-  --build-arg AVALANCHE_VERSION="$avalanche_version" \
-  --build-arg CORETH_COMMIT="$coreth_commit" \
-  --build-arg CURRENT_BRANCH="$current_branch"
diff --git a/scripts/constants.sh b/scripts/constants.sh
index 262f1a48..69a04370 100644
--- a/scripts/constants.sh
+++ b/scripts/constants.sh
@@ -4,20 +4,12 @@
@@ -4,20 +4,12 @@ GOPATH=$(go env GOPATH)"

# Set binary location
-binary_path=${CORETH_BINARY_PATH:-"$GOPATH/src/github.com/ava-labs/avalanchego/build/plugins/evm"}
-
-# Avalabs docker hub
-dockerhub_repo="avaplatform/avalanchego"
+binary_path=${CORETH_BINARY_PATH:-"$GOPATH/src/github.com/flare-foundation/flare/build/plugins/evm"}

# Current branch
current_branch=${CURRENT_BRANCH:-$(git describe --tags --exact-match 2> /dev/null || git symbolic-ref -q --short HEAD || git rev-parse --short HEAD)}
echo "Using branch: ${current_branch}"

-# Image build id
-# Use an abbreviated version of the full commit to tag the image.
-
# WARNING: this will use the most recent commit even if there are un-committed changes present
coreth_commit=${git --git-dir="$CORETH_PATH/.git" rev-parse HEAD}
coreth_commit_id=${coreth_commit::8}

-build_image_id=${BUILD_IMAGE_ID:-"$avalanche_version-$coreth_commit_id"}
diff --git a/scripts/lint.sh b/scripts/lint.sh
new file mode 100755
index 00000000..f41d6b56
--- /dev/null
+++ b/scripts/lint.sh
@@ -0,0 +1,7 @@
+#!/usr/bin/env bash
+
++set -o errexit
++set -o nounset
++set -o pipefail
+
+golangci-lint run --path-prefix=. --timeout 3m
diff --git a/scripts/versions.sh b/scripts/versions.sh
index 6f2cc893..5a3f2696 100644
--- a/scripts/versions.sh
+++ b/scripts/versions.sh
@@ -1,6 +1,6 @@
@@ -1,6 +1,6 @@ #!/usr/bin/env bash

# Set up the versions to be used
-coreth_version=${CORETH_VERSION:-'v0.7.4'}
+coreth_version=${CORETH_VERSION:-'v0.3.1'}
# Don't export them as they're used in the context of other calls
-avalanche_version=${AVALANCHE_VERSION:-'v1.6.4'}
+flare_version=${FLARE_VERSION:-'v0.5.1'}
diff --git a/signer/core/apitypes/types.go b/signer/core/apitypes/types.go
index cab7f9cf..60afd52f 100644
--- a/signer/core/apitypes/types.go

```

```

+++ b/signer/core/apitypes/types.go
@@ -32,9 +32,9 @@ import (
    "math/big"
    "strings"

-    "github.com/ava-labs/coreth/core/types"
    "github.com/ethereum/go-ethereum/common"
    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/flare-foundation/coreth/core/types"
)

type ValidationInfo struct {
diff --git a/tests/init.go b/tests/init.go
index d08bca7e..25e33865 100644
--- a/tests/init.go
+++ b/tests/init.go
@@ -31,7 +31,7 @@ import (
    "math/big"
    "sort"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"
)

// Forks table defines supported forks and their chain config.
@@ -208,6 +208,22 @@ var Forks = map[string]*params.ChainConfig{
    ApricotPhase3BlockTimestamp: big.NewInt(0),
    ApricotPhase4BlockTimestamp: big.NewInt(0),
    },
+    "ApricotPhase5": {
+        ChainID:                big.NewInt(1),
+        HomesteadBlock:          big.NewInt(0),
+        EIP150Block:              big.NewInt(0),
+        EIP155Block:              big.NewInt(0),
+        EIP158Block:              big.NewInt(0),
+        ByzantiumBlock:           big.NewInt(0),
+        ConstantinopleBlock:      big.NewInt(0),
+        PetersburgBlock:          big.NewInt(0),
+        IstanbulBlock:           big.NewInt(0),
+        ApricotPhase1BlockTimestamp: big.NewInt(0),
+        ApricotPhase2BlockTimestamp: big.NewInt(0),
+        ApricotPhase3BlockTimestamp: big.NewInt(0),
+        ApricotPhase4BlockTimestamp: big.NewInt(0),
+        ApricotPhase5BlockTimestamp: big.NewInt(0),
+    },
}

// Returns the set of defined fork names
diff --git a/tests/init_test.go b/tests/init_test.go
index 1ddbfc6..5888be03 100644
--- a/tests/init_test.go
+++ b/tests/init_test.go
@@ -40,7 +40,7 @@ import (
    "strings"
    "testing"

-    "github.com/ava-labs/coreth/params"
+    "github.com/flare-foundation/coreth/params"
)

func readJSON(reader io.Reader, value interface{}) error {
diff --git a/tests/state_test_util.go b/tests/state_test_util.go
index 6e33efa5..73ae6c1f 100644
--- a/tests/state_test_util.go
+++ b/tests/state_test_util.go
@@ -34,17 +34,17 @@ import (
    "strconv"
    "strings"

-    "github.com/ava-labs/coreth/core"
-    "github.com/ava-labs/coreth/core/state"
-    "github.com/ava-labs/coreth/core/state/snapshot"
-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ava-labs/coreth/core/vm"
-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/params"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/hexutil"
+    "github.com/ethereum/go-ethereum/common/math"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/core"
+    "github.com/flare-foundation/coreth/core/state"
+    "github.com/flare-foundation/coreth/core/state/snapshot"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/flare-foundation/coreth/core/vm"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/params"
)

// StateTest checks transaction processing without block context.
diff --git a/trie/database.go b/trie/database.go
index 4f59c68e..c5e59f49 100644
--- a/trie/database.go
+++ b/trie/database.go
@@ -36,12 +36,12 @@ import (
    "time"

    "github.com/VictoriaMetrics/fastcache"
-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/metrics"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
)

var (
diff --git a/trie/database_test.go b/trie/database_test.go
index bc3d3e9d..626ddfe0 100644
--- a/trie/database_test.go
+++ b/trie/database_test.go
@@ -29,8 +29,8 @@ package trie
import (
    "testing"

-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
)

// Tests that the trie database returns a missing trie node error if attempting
diff --git a/trie/iterator.go b/trie/iterator.go
index 4ab0863c..dbba17ce 100644
--- a/trie/iterator.go
+++ b/trie/iterator.go
@@ -31,9 +31,9 @@ import (
    "container/heap"
    "errors"

-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/rlp"

```

```

+         "github.com/flare-foundation/coreth/ethdb"
+     )

    // Iterator is a key-value trie iterator that traverses a Trie.
@@ -285,7 +285,7 @@ func (it *nodeIterator) seek(prefix []byte) error {
+     }
+ }

-// init initializes the the iterator.
+// init initializes the iterator.
func (it *nodeIterator) init() (*nodeIteratorState, error) {
    root := it.trie.Hash()
    state := &nodeIteratorState{node: it.trie.root, index: -1}
diff --git a/trie/iterator_test.go b/trie/iterator_test.go
index 7ef711c3..fe13af8b 100644
--- a/trie/iterator_test.go
+++ b/trie/iterator_test.go
@@ -33,10 +33,10 @@ @@ import (
    "math/rand"
    "testing"

-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
+ )

func TestIterator(t *testing.T) {
diff --git a/trie/proof.go b/trie/proof.go
index 37ce9ed3..4ded339a 100644
--- a/trie/proof.go
+++ b/trie/proof.go
@@ -31,11 +31,11 @@ @@ import (
    "errors"
    "fmt"

-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
+ )

// Prove constructs a merkle proof for key. The result contains all encoded nodes
@@ -482,12 +482,17 @@ @@ func VerifyRangeProof(rootHash common.Hash, firstKey []byte, lastKey []byte, key
    if len(keys) != len(values) {
        return false, fmt.Errorf("inconsistent proof data, keys: %d, values: %d", len(keys), len(values))
    }

-    // Ensure the received batch is monotonic increasing.
+    // Ensure the received batch is monotonic increasing and contains no deletions
    for i := 0; i < len(keys)-1; i++ {
        if bytes.Compare(keys[i], keys[i+1]) >= 0 {
            return false, errors.New("range is not monotonically increasing")
        }
    }

+    for _, value := range values {
+        if len(value) == 0 {
+            return false, errors.New("range contains deletion")
+        }
+    }

    // Special case, there is no edge proof at all. The given range is expected
    // to be the whole leaf-set in the trie.
    if proof == nil {
diff --git a/trie/proof_test.go b/trie/proof_test.go
index f37c2060..3471d671 100644
--- a/trie/proof_test.go
+++ b/trie/proof_test.go
@@ -35,9 +35,9 @@ @@ import (
    "testing"
    "time"

-    "github.com/ava-labs/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
+ )

func init() {
@@ -823,6 +823,85 @@ @@ func TestBloomedProof(t *testing.T) {
    }
}

+// TestEmptyValueRangeProof tests normal range proof with both edge proofs
+// as the existent proof, but with an extra empty value included, which is a
+// noop technically, but practically should be rejected.
+func TestEmptyValueRangeProof(t *testing.T) {
+    trie, values := randomTrie(512)
+    var entries entrySlice
+    for _, kv := range values {
+        entries = append(entries, kv)
+    }
+    sort.Sort(entries)
+
+    // Create a new entry with a slightly modified key
+    mid := len(entries) / 2
+    key := common.CopyBytes(entries[mid-1].k)
+    for n := len(key) - 1; n >= 0; n-- {
+        if key[n] < 0xff {
+            key[n]++
+            break
+        }
+    }
+
+    noop := &kv{key, []byte{}, false}
+    entries = append(append(append([]*kv{}, entries[:mid]...), noop), entries[mid:]...)
+
+    start, end := 1, len(entries)-1
+
+    proof := memorydb.New()
+    if err := trie.Prove(entries[start].k, 0, proof); err != nil {
+        t.Fatalf("Failed to prove the first node %v", err)
+    }
+    if err := trie.Prove(entries[end-1].k, 0, proof); err != nil {
+        t.Fatalf("Failed to prove the last node %v", err)
+    }
+
+    var keys [][]byte
+    var vals [][]byte
+    for i := start; i < end; i++ {
+        keys = append(keys, entries[i].k)
+        vals = append(vals, entries[i].v)
+    }
+
+    _, err := VerifyRangeProof(trie.Hash(), keys[0], keys[len(keys)-1], keys, vals, proof)
+    if err == nil {
+        t.Fatalf("Expected failure on noop entry")
+    }
+}
+
+// TestAllElementsEmptyValueRangeProof tests the range proof with all elements,
+// but with an extra empty value included, which is a noop technically, but
+// practically should be rejected.
+func TestAllElementsEmptyValueRangeProof(t *testing.T) {

```

```

+     trie, values := randomTrie(512)
+     var entries entrySlice
+     for _, kv := range values {
+         entries = append(entries, kv)
+     }
+     sort.Sort(entries)
+
+     // Create a new entry with a slightly modified key
+     mid := len(entries) / 2
+     key := common.CopyBytes(entries[mid-1].k)
+     for n := len(key) - 1; n >= 0; n-- {
+         if key[n] < 0xff {
+             key[n]++
+             break
+         }
+     }
+     noop := &kv{key, []byte{}, false}
+     entries = append(append(append([]*kv{}, entries[:mid]...), noop), entries[mid:]...)
+
+     var keys [][]byte
+     var vals [][]byte
+     for i := 0; i < len(entries); i++ {
+         keys = append(keys, entries[i].k)
+         vals = append(vals, entries[i].v)
+     }
+     _, err := VerifyRangeProof(trie.Hash(), nil, nil, keys, vals, nil)
+     if err == nil {
+         t.Fatalf("Expected failure on noop entry")
+     }
+ }
+ }
+
+ // mutateByte changes one byte in b.
+ func mutateByte(b []byte) {
+     for r := mrand.Intn(len(b)); ; {
diff --git a/trie/secure_trie.go b/trie/secure_trie.go
index 4f6a556e..0d324e15 100644
--- a/trie/secure_trie.go
+++ b/trie/secure_trie.go
@@ -29,10 +29,10 @@ package trie
import (
    "fmt"

-    "github.com/ava-labs/coreth/core/types"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/log"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/core/types"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
)

// SecureTrie wraps a trie with key hashing. In a secure trie, all
diff --git a/trie/secure_trie_test.go b/trie/secure_trie_test.go
index c9c29d70..642e7d2a 100644
--- a/trie/secure_trie_test.go
+++ b/trie/secure_trie_test.go
@@ -32,9 +32,9 @@ import (
    "sync"
    "testing"

-    "github.com/ava-labs/coreth/ethdb/memorydb"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/crypto"
)

func newEmptySecure() *SecureTrie {
diff --git a/trie/stacktrie.go b/trie/stacktrie.go
index 248b4e4e..33445f8f 100644
--- a/trie/stacktrie.go
+++ b/trie/stacktrie.go
@@ -35,10 +35,10 @@ import (
    "io"
    "sync"

-    "github.com/ava-labs/coreth/ethdb"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/log"
-    "github.com/ethereum/go-ethereum/rlp"
+    "github.com/flare-foundation/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/rlp"
)

var ErrCommitDisabled = errors.New("no database for committing")
@@ -64,12 +64,11 @@ func returnToPool(st *StackTrie) {
// in order. Once it determines that a subtree will no longer be inserted
// into, it will hash it and free up the memory it uses.
type StackTrie struct {
-    nodeType uint8 // node type (as in branch, ext, leaf)
-    val []byte // value contained by this node if it's a leaf
-    key []byte // key chunk covered by this (full)ext node
-    keyOffset int // offset of the key chunk inside a full key
-    children [16]*StackTrie // list of children (for fullnodes and exts)
-    db ethdb.KeyValueWriter // Pointer to the commit db, can be nil
+    nodeType uint8 // node type (as in branch, ext, leaf)
+    val []byte // value contained by this node if it's a leaf
+    key []byte // key chunk covered by this (leaf)ext node
+    children [16]*StackTrie // list of children (for branch and exts)
+    db ethdb.KeyValueWriter // Pointer to the commit db, can be nil
}

// NewStackTrie allocates and initializes an empty trie.
@@ -100,15 +99,13 @@ func (st *StackTrie) MarshalBinary() (data []byte, err error) {
    w := bufio.NewWriter(&b)

    if err := gob.NewEncoder(w).Encode(struct {
-        NodeType uint8
-        Val []byte
-        Key []byte
-        KeyOffset uint8
+        NodeType uint8
+        Val []byte
+        Key []byte
    }) {
        st.nodeType,
        st.val,
        st.key,
-        uint8(st.keyOffset),
    }); err != nil {
        return nil, err
    }
}

@@ -136,16 +133,14 @@ func (st *StackTrie) UnmarshalBinary(data []byte) error {
    func (st *StackTrie) unmarshalBinary(r io.Reader) error {
        var dec struct {
-            NodeType uint8
-            Val []byte
-            Key []byte
-            KeyOffset uint8
+            NodeType uint8
+            Val []byte
+            Key []byte
        }
        gob.NewDecoder(r).Decode(&dec)
        st.nodeType = dec.NodeType
        st.val = dec.Val
        st.key = dec.Key
    }
}

```



```

-     st.keyOffset = int(dec.KeyOffset)

    var hasChild = make([]byte, 1)
    for i := range st.children {
@@ -170,20 +165,18 @@ func (st *StackTrie) setDb(db ethdb.KeyValueWriter) {
    }
}

-func newLeaf(ko int, key, val []byte, db ethdb.KeyValueWriter) *StackTrie {
+func newLeaf(key, val []byte, db ethdb.KeyValueWriter) *StackTrie {
    st := stackTrieFromPool(db)
    st.nodeType = leafNode
-    st.keyOffset = ko
-    st.key = append(st.key, key[ko:]...)
+    st.key = append(st.key, key...)
+    st.val = val
    return st
}

-func newExt(ko int, key []byte, child *StackTrie, db ethdb.KeyValueWriter) *StackTrie {
+func newExt(key []byte, child *StackTrie, db ethdb.KeyValueWriter) *StackTrie {
    st := stackTrieFromPool(db)
    st.nodeType = extNode
-    st.keyOffset = ko
-    st.key = append(st.key, key[ko:]...)
+    st.key = append(st.key, key...)
+    st.children[0] = child
    return st
}

@@ -221,17 +214,18 @@ func (st *StackTrie) Reset() {
    st.children[i] = nil
}
st.nodeType = emptyNode
- st.keyOffset = 0
}

// Helper function that, given a full key, determines the index
// at which the chunk pointed by st.keyOffset is different from
// the same chunk in the full key.
func (st *StackTrie) getDiffIndex(key []byte) int {
-    diffindex := 0
-    for ; diffindex < len(st.key) && st.key[diffindex] == key[st.keyOffset+diffindex]; diffindex++ {
+    for idx, nibble := range st.key {
+        if nibble != key[idx] {
+            return idx
+        }
    }
-    return diffindex
+    return len(st.key)
}

// Helper function to that inserts a (key, value) pair into
@@ -239,7 +233,7 @@ func (st *StackTrie) getDiffIndex(key []byte) int {
func (st *StackTrie) insert(key, value []byte) {
    switch st.nodeType {
    case branchNode: /* Branch */
-        idx := int(key[st.keyOffset])
+        idx := int(key[0])
        // Unresolve elder siblings
        for i := idx - 1; i >= 0; i-- {
            if st.children[i] != nil {
@@ -251,10 +245,10 @@ func (st *StackTrie) insert(key, value []byte) {
        }
        // Add new child
        if st.children[idx] == nil {
-            st.children[idx] = stackTrieFromPool(st.db)
-            st.children[idx].keyOffset = st.keyOffset + 1
+            st.children[idx] = newLeaf(key[1:], value, st.db)
        } else {
            st.children[idx].insert(key[1:], value)
        }
-        st.children[idx].insert(key, value)
    case extNode: /* Ext */
        // Compare both key chunks and see where they differ
        diffidx := st.getDiffIndex(key)
@@ -267,7 +261,7 @@ func (st *StackTrie) insert(key, value []byte) {
        if diffidx == len(st.key) {
            // Ext key and key segment are identical, recurse into
            // the child node.
-            st.children[0].insert(key, value)
+            st.children[0].insert(key[diffidx:], value)
            return
        }
        // Save the original part. Depending if the break is
@@ -276,7 +270,7 @@ func (st *StackTrie) insert(key, value []byte) {
        // node directly.
        var n *StackTrie
        if diffidx < len(st.key)-1 {
-            n = newExt(diffidx+1, st.key, st.children[0], st.db)
+            n = newExt(st.key[diffidx+1:], st.children[0], st.db)
        } else {
            // Break on the last byte, no need to insert
            // an extension node: reuse the current node
@@ -298,15 +292,14 @@ func (st *StackTrie) insert(key, value []byte) {
            // node.
            st.children[0] = stackTrieFromPool(st.db)
            st.children[0].nodeType = branchNode
-            st.children[0].keyOffset = st.keyOffset + diffidx
            p := st.children[0]

            // Create a leaf for the inserted part
-            o := newLeaf(st.keyOffset+diffidx+1, key, value, st.db)
+            o := newLeaf(key[diffidx+1:], value, st.db)

            // Insert both child leaves where they belong:
            origIdx := st.key[diffidx]
            newIdx := key[diffidx+st.keyOffset]
+            newIdx := key[diffidx]
            p.children[origIdx] = n
            p.children[newIdx] = o
            st.key = st.key[:diffidx]
@@ -340,7 +333,6 @@ func (st *StackTrie) insert(key, value []byte) {
        st.nodeType = extNode
        st.children[0] = NewStackTrie(st.db)
        st.children[0].nodeType = branchNode
-        st.children[0].keyOffset = st.keyOffset + diffidx
        p := st.children[0]
    }
}

@@ -349,11 +341,11 @@ func (st *StackTrie) insert(key, value []byte) {
    // The child leave will be hashed directly in order to
    // free up some memory.
    origIdx := st.key[diffidx]
-    p.children[origIdx] = newLeaf(diffidx+1, st.key, st.val, st.db)
+    p.children[origIdx] = newLeaf(st.key[diffidx+1:], st.val, st.db)
    p.children[origIdx].hash()

-    newIdx := key[diffidx+st.keyOffset]
-    p.children[newIdx] = newLeaf(p.keyOffset+1, key, value, st.db)
+    newIdx := key[diffidx]
+    p.children[newIdx] = newLeaf(key[diffidx+1:], value, st.db)

    // Finally, cut off the key part that has been passed
    // over to the children.

```

```
@@ -361,7 +353,7 @@ func (st *StackTrie) insert(key, value []byte) {
    st.val = nil
    case emptyNode: /* Empty */
        st.nodeType = leafNode
-       st.key = key[st.keyOffset:]
+       st.key = key
        st.val = value
    case hashedNode:
        panic("trying to insert into hash")
diff --git a/trie/stacktrie_test.go b/trie/stacktrie_test.go
index 189e9d41..1a5b2005 100644
--- a/trie/stacktrie_test.go
+++ b/trie/stacktrie_test.go
@@ -31,11 +31,171 @@ import (
    "math/big"
    "testing"

-   "github.com/ava-labs/coreth/ethdb/memorydb"
-   "github.com/ethereum/go-ethereum/common"
-   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/flare-foundation/coreth/ethdb/memorydb"
)

+func TestStackTrieInsertAndHash(t *testing.T) {
+    type KeyValueHash struct {
+        K string // Hex string for key.
+        V string // Value, directly converted to bytes.
+        H string // Expected root hash after insert of (K, V) to an existing trie.
+    }
+    tests := [][]KeyValueHash{
+        { // {0:0, 7:0, f:0}
+            {"00", "v", 0, 0, "5cb26357b95bb9af08475be00243ceb68ade0b66b5cd816b0c18a18c612d2d21"},
+            {"70", "v", 0, 1, "8ff64309574f7a437a7ad1628e690eb7663cfde10676f8a904a8c8291dbc1603"},
+            {"f0", "v", 0, 2, "9e3a01bd8d43efb8e9d4b5506648150b8e3ed1cae596f84ee28e01a72635470"},
+        },
+        { // {1:0cc, e:1:fc, e:fc}}
+            {"10cc", "v", 1, 0, "233e9b257843f3dfdb1cce6676cdf9e595ac96ee1b55031434d852bc7ac9185"},
+            {"elfc", "v", 1, 1, "39c5e908ae83d0c78520c7cbda0b3782daf594700e44546e93def8f049cca95"},
+            {"eefc", "v", 1, 2, "d789567559fd76fe5b7d9c42f3750f942502ac1c7f2a466e2f690ec4b6c2a7c"},
+        },
+        { // {b:{a:ac, b:ac}, d:acc}
+            {"baac", "v", 2, 0, "8be1c86ba7ec4c61e14c1a9b75055e0464c2633ae66a055a24e75450156a5d42"},
+            {"bbac", "v", 2, 1, "8495159b9895a7d88d973171d737c8aaace6fe6ac02a4769fff1bc43bccccc4cc"},
+            {"dacc", "v", 2, 2, "9bfcfc56220a27328deb9dc6ee23d46c9ebc9c69e78acdafa2c7040602c63ca"},
+        },
+        { // {0:0cccc, 2:456{0:0, 2:2}
+            {"00cccc", "v", 3, 0, "e57dc2785b99ce9205000cb41b32ebaea7ac3e158952b44c87d186e6d190a6530"},
+            {"245600", "v", 3, 1, "0335354adbd360a45c1871a842452287721b64b4234dfe08760b243523c998db"},
+            {"245622", "v", 3, 2, "9e6832db0dca2b5cf81c0e0727bfde6a6fc39d5de33e5720bccacc183c162104e"},
+        },
+        { // {1:4567{1:1c, 3:3c}, 3:0cccccc}
+            {"1456711c", "v", 4, 0, "f2389e78d98fed99f3e63d6d1623c1d4d9e8c91cb1d585de81fbc7c0e60d3529"},
+            {"1456733c", "v", 4, 1, "101189b3fab852be97a0120c03d95eefc984d3ed639f2328527de6def55a9c0"},
+            {"30cccccc", "v", 4, 2, "3780ce111f98d15751dfde1eb21080efc7d3914b429e5c84c64db637c55405b3"},
+        },
+        { // 8800{1:f, 2:e, 3:d}
+            {"88001f", "v", 5, 0, "e817db50d84f341d443c6f6593cafda093fc85e773a762421d47daa6ac993bd5"},
+            {"88002e", "v", 5, 1, "d6e3e6047bdc110edd296a4d63c030aec451bee9d0875bc5a198eee8cda34f68"},
+            {"88003d", "v", 5, 2, "b6bdf8298c703342188e5f7f84921a402042d0e5fb059969dd53a6b6b1fb989e"},
+        },
+        { // 0{1:fc, 2:ec, 4:dc}
+            {"01fc", "v", 6, 0, "693268f2ca80d32b015f61cd2c4dba5a47a6b52a1c34f8e6945fad684e7a0d5"},
+            {"02ec", "v", 6, 1, "e24ddd44469310c2b785a204618874bf486d2f7822603a9b8dc58d6524d5de"},
+            {"04dc", "v", 6, 2, "33fc259629187bbe54b92f82f8cd80883b91a12e41a9456b84fc155321e334db7"},
+        },
+        { // f{0:fdccc, f:ff{0:f, f:f}}
+            {"f0fdccc", "v", 7, 0, "b0966b5aa469a3e292bc5f7cfa6c396ae7a657255eef552ea7e12f996de795b90"},
+            {"fffff0f", "v", 7, 1, "3b1ca154ec2a3d96d8d77b0ddef0abf4e40a53a64eb03cecf78da9ec43799fa3d0"},
+            {"fffffff", "v", 7, 2, "e75463041f1be8252781be0ace579a44ea4387bfb5b2739f4607af676f7719678"},
+        },
+        { // ff{0:f{0:f, f:f}, f:fcc}
+            {"ff0f0f", "v", 8, 0, "0928af9b14718ec8262ab89df430f1e5fbf66fac0fed037aff2b6767ae8c8684"},
+            {"ff0fff", "v", 8, 1, "d870f4d3ce26b0bf86912810a1960693630c20a48ba56be0ad04bc3e9ddb01e6"},
+            {"fffffcc", "v", 8, 2, "4239f10dd9d9915ecf2e047d6a576bdc1733ed77a30830f1bf29dea7f7d8e966f"},
+        },
+        {
+            {"123d", "x", 0, "fc453d88b6f128a77c448669710497380fa4588abbea9f78f4c20c80daa797d0"},
+            {"123e", "x", 1, "5af48f2d8a9a015c1ff7fa8b8c7f6b676233bd320e8fb57fd7933622badd2cec"},
+            {"123f", "x", 2, "1164d7299964e74ac40d761f9189b2a3987fae95980d0f7e29d3aaf3eae9e15"},
+        },
+        {
+            {"123d", "x", 0, "fc453d88b6f128a77c448669710497380fa4588abbea9f78f4c20c80daa797d0"},
+            {"123e", "x", 1, "5af48f2d8a9a015c1ff7fa8b8c7f6b676233bd320e8fb57fd7933622badd2cec"},
+            {"124a", "x", 2, "661a96a669869d76b7231380da0649d013301425fbae9d5c5fae6405aa31cfce"},
+        },
+        {
+            {"123d", "x", 0, "fc453d88b6f128a77c448669710497380fa4588abbea9f78f4c20c80daa797d0"},
+            {"123e", "x", 1, "5af48f2d8a9a015c1ff7fa8b8c7f6b676233bd320e8fb57fd7933622badd2cec"},
+            {"13aa", "x", 2, "6590120e1fd3ffda1a90e8de5bb10750b61079b0776cca414dd79a24e4d4356"},
+        },
+        {
+            {"123d", "x", 0, "fc453d88b6f128a77c448669710497380fa4588abbea9f78f4c20c80daa797d0"},
+            {"123e", "x", 1, "5af48f2d8a9a015c1ff7fa8b8c7f6b676233bd320e8fb57fd7933622badd2cec"},
+            {"2aaa", "x", 2, "f869b40e0c55eace1918332ef91563616fbf0755e2b946119679f7ef8e44b514"},
+        },
+        {
+            {"1234da", "x", 0, "1c4b4462e9f56a80ca0f5d77c0d632c41b0102290930343cf1791e971a045a79"},
+            {"1234ea", "x", 1, "2f502917f3ba7d328c21c8b45ee0f160652e68450332c166d4ad02d1afe31862"},
+            {"1234fa", "x", 2, "4f4e368ab367090d5bc3db5f7729f8bd60df84de309b4633a6b69ab66142c0"},
+        },
+        {
+            {"1234da", "x", 0, "1c4b4462e9f56a80ca0f5d77c0d632c41b0102290930343cf1791e971a045a79"},
+            {"1234ea", "x", 1, "2f502917f3ba7d328c21c8b45ee0f160652e68450332c166d4ad02d1afe31862"},
+            {"1235aa", "x", 2, "21840121d11a91ac8bbad9a5d06a4f902a5c8d56a47b85600ba813814b7bfc9b"},
+        },
+        {
+            {"1234da", "x", 0, "1c4b4462e9f56a80ca0f5d77c0d632c41b0102290930343cf1791e971a045a79"},
+            {"1234ea", "x", 1, "2f502917f3ba7d328c21c8b45ee0f160652e68450332c166d4ad02d1afe31862"},
+            {"124aaa", "x", 2, "ea040dddf6ae3fbd1524bdcc19c0ab15810159962006632027fa5cf21e4441e"},
+        },
+        {
+            {"1234da", "x", 0, "1c4b4462e9f56a80ca0f5d77c0d632c41b0102290930343cf1791e971a045a79"},
+            {"1234ea", "x", 1, "2f502917f3ba7d328c21c8b45ee0f160652e68450332c166d4ad02d1afe31862"},
+            {"13aaaa", "x", 2, "e4be66c67e44f2dd8ba36036e45a44ff68f8d5294272b1911a45f886a34507"},
+        },
+        {
+            {"1234da", "x", 0, "1c4b4462e9f56a80ca0f5d77c0d632c41b0102290930343cf1791e971a045a79"},
+            {"1234ea", "x", 1, "2f502917f3ba7d328c21c8b45ee0f160652e68450332c166d4ad02d1afe31862"},
+            {"2aaaaa", "x", 2, "5f5989b820ff576b7d49e77bb64f26602294f6c42a1a3bec669cd9e8dc8ec9"},
+        },
+        {
+            {"000000", "x", 0, "3b32b7af0bddc7940e7364ee18b5a59702c1825e469452c8483b9c4e0218b55a"},
+            {"1234da", "x", 1, "3ab152a1285dca31945566f872c1cc2f17a707040eda32aeae46a5e91033dde2"},
+            {"1234ea", "x", 2, "0cccc87f96ddef55563c1b3be3c64ff66a644333cd9cd99852cb53b6412b9b8"},
+            {"1234fa", "x", 3, "65bb3aafea8121111d693fffe34881c14d27b128fd113fa120961f251fe28428d"},
+        },
+        {
+            {"000000", "x", 0, "3b32b7af0bddc7940e7364ee18b5a59702c1825e469452c8483b9c4e0218b55a"},
+            {"1234da", "x", 1, "3ab152a1285dca31945566f872c1cc2f17a707040eda32aeae46a5e91033dde2"},
+            {"1234ea", "x", 2, "0cccc87f96ddef55563c1b3be3c64ff66a644333cd9cd99852cb53b6412b9b8"},
+            {"1235aa", "x", 3, "f670e4d2547c533c5f21e0045442e2ecb733f347ad6d29ef36e0f5ba31bb11a8"},
+        },
+        {
+            {"000000", "x", 0, "3b32b7af0bddc7940e7364ee18b5a59702c1825e469452c8483b9c4e0218b55a"},
+        },
+    }
}
```

```

+         {"1234da", "x", 1, "3ab152a1285dca31945566f872c1cc2f17a770440eda32aeee46a5e91033dde2"},
+         {"1234ea", "x", 2, "0cccc87f96ddef55563c1b3be3c64ff6a644333c3d9cd99852cb53b6412b9b8"},
+         {"124aaa", "x", 3, "c17464123050a9a6f29b5574bb2f92f6d305c1794976b475b7fb0316b6335598"},
+     },
+     {
+         {"000000", "x", 0, "3b32b7af0bddc7940e7364ee18b5a59702c1825e469452c8483b9c4e0218b55a"},
+         {"1234da", "x", 1, "3ab152a1285dca31945566f872c1cc2f17a770440eda32aeee46a5e91033dde2"},
+         {"1234ea", "x", 2, "0cccc87f96ddef55563c1b3be3c64ff6a644333c3d9cd99852cb53b6412b9b8"},
+         {"13aaaa", "x", 3, "aa8301be8cb52ea5cd249f5feb79fb4315ee8de2140c604033f4b3ffff78f0105"},
+     },
+     {
+         {"0000", "x", 0, "cb8c09ad07ae882136f602b3f21f8733a9f5a78f1d2525a8d24d1c13258000b2"},
+         {"123d", "x", 1, "8f09663deb02f08958136410dc48565e077f76bb6c9d8c84d35fc8913a657d31"},
+         {"123e", "x", 2, "0d230561e398c579e09a9f7b69ceaf7d3970f5a436fdb28b68b7a37c5bdd6b80"},
+         {"123f", "x", 3, "80f7bad1893ca57e3443bb3305a517723a74d3ba831bcaca22a170645eb7aafb"},
+     },
+     {
+         {"0000", "x", 0, "cb8c09ad07ae882136f602b3f21f8733a9f5a78f1d2525a8d24d1c13258000b2"},
+         {"123d", "x", 1, "8f09663deb02f08958136410dc48565e077f76bb6c9d8c84d35fc8913a657d31"},
+         {"123e", "x", 2, "0d230561e398c579e09a9f7b69ceaf7d3970f5a436fdb28b68b7a37c5bdd6b80"},
+         {"124a", "x", 3, "383bc1bb4f019e6bc4da3751509ea709b58dd1ac46081670834bae072f3e9557"},
+     },
+     {
+         {"0000", "x", 0, "cb8c09ad07ae882136f602b3f21f8733a9f5a78f1d2525a8d24d1c13258000b2"},
+         {"123d", "x", 1, "8f09663deb02f08958136410dc48565e077f76bb6c9d8c84d35fc8913a657d31"},
+         {"123e", "x", 2, "0d230561e398c579e09a9f7b69ceaf7d3970f5a436fdb28b68b7a37c5bdd6b80"},
+         {"13aa", "x", 3, "ff0dc70ce2e5db90ee42a4c2ad12139596b890e90eb4e16526ab38fa465b35cf"},
+     }
+ }
+ st := NewStackTrie(nil)
+ for i, test := range tests {
+     // The StackTrie does not allow Insert(), Hash(), Insert(), ...
+     // so we will create new trie for every sequence length of inserts.
+     for l := 1; l <= len(test); l++ {
+         st.Reset()
+         for j := 0; j < l; j++ {
+             kv := &test[j]
+             if err := st.TryUpdate(common.FromHex(kv.K), []byte(kv.V)); err != nil {
+                 t.Fatal(err)
+             }
+         }
+         expected := common.HexToHash(test[l-1].H)
+         if h := st.Hash(); h != expected {
+             t.Errorf("%d(%d): root hash mismatch: %x, expected %x", i, l, h, expected)
+         }
+     }
+ }
+}
+
+func TestSizeBug(t *testing.T) {
+    st := NewStackTrie(nil)
+    nt, _ := New(common.Hash{}, NewDatabase(memorydb.New()))
+diff --git a/trie/sync.go b/trie/sync.go
index 85b1f5a5..af5e2821 100644
--- a/trie/sync.go
+++ b/trie/sync.go
@@ -30,10 +30,11 @@ import (
     "errors"
     "fmt"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/common/prque"

+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
 )

// ErrNotRequested is returned by the trie sync when it's requested to process a
diff --git a/trie/sync_bloom.go b/trie/sync_bloom.go
index 51c318c4..fafc9839 100644
--- a/trie/sync_bloom.go
+++ b/trie/sync_bloom.go
@@ -33,12 +33,14 @@ import (
     "sync/atomic"
     "time"

-    "github.com/ava-labs/coreth/core/rawdb"
-    "github.com/ava-labs/coreth/ethdb"
+    bloomfilter "github.com/holiman/bloomfilter/v2"

+    "github.com/ethereum/go-ethereum/common"
+    "github.com/ethereum/go-ethereum/log"
+    "github.com/ethereum/go-ethereum/metrics"
-    bloomfilter "github.com/holiman/bloomfilter/v2"

+    "github.com/flare-foundation/coreth/core/rawdb"
+    "github.com/flare-foundation/coreth/ethdb"
 )

var (
diff --git a/trie/sync_test.go b/trie/sync_test.go
index d3bbdbd9..1f3839d2 100644
--- a/trie/sync_test.go
+++ b/trie/sync_test.go
@@ -27,12 +27,8 @@ package trie

import (
-    "bytes"
-    "testing"

-    "github.com/ava-labs/coreth/ethdb/memorydb"
-    "github.com/ethereum/go-ethereum/common"
-    "github.com/ethereum/go-ethereum/crypto"
+    "github.com/flare-foundation/coreth/ethdb/memorydb"
)

// makeTestTrie create a sample test trie to test node-wise reconstruction.
@@ -65,424 +61,3 @@ func makeTestTrie() (*Database, *SecureTrie, map[string][]byte) {
    // Return the generated trie
    return triedb, trie, content
}

-// checkTrieContents cross references a reconstructed trie with an expected data
-// content map.
-func checkTrieContents(t *testing.T, db *Database, root []byte, content map[string][]byte) {
-    // Check root availability and trie contents
-    trie, err := NewSecure(common.BytesToHash(root), db)
-    if err != nil {
-        t.Fatalf("failed to create trie at %x: %v", root, err)
-    }
-    if err := checkTrieConsistency(db, common.BytesToHash(root)); err != nil {
-        t.Fatalf("inconsistent trie at %x: %v", root, err)
-    }
-    for key, val := range content {
-        if have := trie.Get([]byte(key)); !bytes.Equal(have, val) {
-            t.Errorf("entry %x: content mismatch: have %x, want %x", key, have, val)
-        }
-    }
-}
-
-// checkTrieConsistency checks that all nodes in a trie are indeed present.

```

```

-func checkTrieConsistency(db *Database, root common.Hash) error {
-    // Create and iterate a trie rooted in a subnode
-    trie, err := NewSecure(root, db)
-    if err != nil {
-        return nil // Consider a non existent state consistent
-    }
-    it := trie.NodeIterator(nil)
-    for it.Next(true) {
-    }
-    return it.Error()
-}
-
-// Tests that an empty trie is not scheduled for syncing.
-func TestEmptySync(t *testing.T) {
-    dbA := NewDatabase(memorydb.New())
-    dbB := NewDatabase(memorydb.New())
-    emptyA, _ := New(common.Hash{}, dbA)
-    emptyB, _ := New(emptyRoot, dbB)
-
-    for i, trie := range []*Trie{emptyA, emptyB} {
-        sync := NewSync(trie.Hash(), memorydb.New(), nil, NewSyncBloom(1, memorydb.New()))
-        if nodes, paths, codes := sync.Missing(1); len(nodes) != 0 || len(paths) != 0 || len(codes) != 0 {
-            t.Errorf("test %d: content requested for empty trie: %v, %v, %v", i, nodes, paths, codes)
-        }
-    }
-}
-
-// Tests that given a root hash, a trie can sync iteratively on a single thread,
-// requesting retrieval tasks and returning all of them in one go.
-func TestIterativeSyncIndividual(t *testing.T) { testIterativeSync(t, 1, false) }
-func TestIterativeSyncBatched(t *testing.T) { testIterativeSync(t, 100, false) }
-func TestIterativeSyncIndividualByPath(t *testing.T) { testIterativeSync(t, 1, true) }
-func TestIterativeSyncBatchedByPath(t *testing.T) { testIterativeSync(t, 100, true) }
-
-func testIterativeSync(t *testing.T, count int, bypath bool) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    nodes, paths, codes := sched.Missing(count)
-    var (
-        hashQueue []common.Hash
-        pathQueue []SyncPath
-    )
-    if !bypath {
-        hashQueue = append(append(hashQueue[:0], nodes...), codes...)
-    } else {
-        hashQueue = append(hashQueue[:0], codes...)
-        pathQueue = append(pathQueue[:0], paths...)
-    }
-    for len(hashQueue)+len(pathQueue) > 0 {
-        results := make([]SyncResult, len(hashQueue)+len(pathQueue))
-        for i, hash := range hashQueue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for hash %x: %v", hash, err)
-            }
-            results[i] = SyncResult{hash, data}
-        }
-        for i, path := range pathQueue {
-            data, _, err := srcTrie.TryGetNode(path[0])
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for path %x: %v", path, err)
-            }
-            results[len(hashQueue)+i] = SyncResult{crypto.Keccak256Hash(data), data}
-        }
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-
-        nodes, paths, codes = sched.Missing(count)
-        if !bypath {
-            hashQueue = append(append(hashQueue[:0], nodes...), codes...)
-        } else {
-            hashQueue = append(hashQueue[:0], codes...)
-            pathQueue = append(pathQueue[:0], paths...)
-        }
-    }
-    // Cross check that the two tries are in sync
-    checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-}
-
-// Tests that the trie scheduler can correctly reconstruct the state even if only
-// partial results are returned, and the others sent only later.
-func TestIterativeDelayedSync(t *testing.T) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    nodes, _, codes := sched.Missing(10000)
-    queue := append(append([]common.Hash{}, nodes...), codes...)
-
-    for len(queue) > 0 {
-        // Sync only half of the scheduled nodes
-        results := make([]SyncResult, len(queue)/2+1)
-        for i, hash := range queue[:len(results)] {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            results[i] = SyncResult{hash, data}
-        }
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-
-        nodes, _, codes = sched.Missing(10000)
-        queue = append(append(queue[len(results):], nodes...), codes...)
-    }
-    // Cross check that the two tries are in sync
-    checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-}

```

```

-
-// Tests that given a root hash, a trie can sync iteratively on a single thread,
-// requesting retrieval tasks and returning all of them in one go, however in a
-// random order.
-func TestIterativeRandomSyncIndividual(t *testing.T) { testIterativeRandomSync(t, 1) }
-func TestIterativeRandomSyncBatched(t *testing.T) { testIterativeRandomSync(t, 100) }
-
-func testIterativeRandomSync(t *testing.T, count int) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    queue := make(map[common.Hash]struct{})
-    nodes, _, codes := sched.Missing(count)
-    for _, hash := range append(nodes, codes...) {
-        queue[hash] = struct{}{}
-    }
-    for len(queue) > 0 {
-        // Fetch all the queued nodes in a random order
-        results := make([]SyncResult, 0, len(queue))
-        for hash := range queue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            results = append(results, SyncResult{hash, data})
-        }
-        // Feed the retrieved results back and queue new tasks
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-
-        queue = make(map[common.Hash]struct{})
-        nodes, _, codes = sched.Missing(count)
-        for _, hash := range append(nodes, codes...) {
-            queue[hash] = struct{}{}
-        }
-    }
-    // Cross check that the two tries are in sync
-    checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-}
-
-// Tests that the trie scheduler can correctly reconstruct the state even if only
-// partial results are returned (Even those randomly), others sent only later.
-func TestIterativeRandomDelayedSync(t *testing.T) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    queue := make(map[common.Hash]struct{})
-    nodes, _, codes := sched.Missing(10000)
-    for _, hash := range append(nodes, codes...) {
-        queue[hash] = struct{}{}
-    }
-    for len(queue) > 0 {
-        // Sync only half of the scheduled nodes, even those in random order
-        results := make([]SyncResult, 0, len(queue)/2+1)
-        for hash := range queue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            results = append(results, SyncResult{hash, data})
-
-            if len(results) >= cap(results) {
-                break
-            }
-        }
-        // Feed the retrieved results back and queue new tasks
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-        for _, result := range results {
-            delete(queue, result.Hash)
-        }
-        nodes, _, codes = sched.Missing(10000)
-        for _, hash := range append(nodes, codes...) {
-            queue[hash] = struct{}{}
-        }
-    }
-    // Cross check that the two tries are in sync
-    checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-}
-
-// Tests that a trie sync will not request nodes multiple times, even if they
-// have such references.
-func TestDuplicateAvoidanceSync(t *testing.T) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    nodes, _, codes := sched.Missing(0)
-    queue := append(append([]common.Hash{}, nodes...), codes...)
-    requested := make(map[common.Hash]struct{})
-
-    for len(queue) > 0 {
-        results := make([]SyncResult, len(queue))
-        for i, hash := range queue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            if _, ok := requested[hash]; ok {
-                t.Errorf("hash %x already requested once", hash)
-            }
-            requested[hash] = struct{}{}
-
-            results[i] = SyncResult{hash, data}
-        }

```

```

    }
    for _, result := range results {
        if err := sched.Process(result); err != nil {
            t.Fatalf("failed to process result %v", err)
        }
    }
    batch := diskdb.NewBatch()
    if err := sched.Commit(batch); err != nil {
        t.Fatalf("failed to commit data: %v", err)
    }
    batch.Write()

    nodes, _, codes = sched.Missing(0)
    queue = append(append(queue[:0], nodes...), codes...)
}
// Cross check that the two tries are in sync
checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-}

-// Tests that at any point in time during a sync, only complete sub-tries are in
-// the database.
-func TestIncompleteSync(t *testing.T) {
-    // Create a random trie to copy
-    srcDb, srcTrie, _ := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    var added []common.Hash
-
-    nodes, _, codes := sched.Missing(1)
-    queue := append(append([]common.Hash{}, nodes...), codes...)
-    for len(queue) > 0 {
-        // Fetch a batch of trie nodes
-        results := make([]SyncResult, len(queue))
-        for i, hash := range queue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            results[i] = SyncResult{hash, data}
-        }
-        // Process each of the trie nodes
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-        for _, result := range results {
-            added = append(added, result.Hash)
-            // Check that all known sub-tries in the synced trie are complete
-            if err := checkTrieConsistency(triedb, result.Hash); err != nil {
-                t.Fatalf("trie inconsistent: %v", err)
-            }
-        }
-        // Fetch the next batch to retrieve
-        nodes, _, codes = sched.Missing(1)
-        queue = append(append(queue[:0], nodes...), codes...)
-    }
-    // Sanity check that removing any node from the database is detected
-    for _, node := range added[1:] {
-        key := node.Bytes()
-        value, _ := diskdb.Get(key)
-
-        diskdb.Delete(key)
-        if err := checkTrieConsistency(triedb, added[0]); err == nil {
-            t.Fatalf("trie inconsistency not caught, missing: %x", key)
-        }
-        diskdb.Put(key, value)
-    }
-}
-
-// Tests that trie nodes get scheduled lexicographically when having the same
-// depth.
-func TestSyncOrdering(t *testing.T) {
-    // Create a random trie to copy
-    srcDb, srcTrie, srcData := makeTestTrie()
-
-    // Create a destination trie and sync with the scheduler, tracking the requests
-    diskdb := memorydb.New()
-    triedb := NewDatabase(diskdb)
-    sched := NewSync(srcTrie.Hash(), diskdb, nil, NewSyncBloom(1, diskdb))
-
-    nodes, paths, _ := sched.Missing(1)
-    queue := append([]common.Hash{}, nodes...)
-    reqs := append([]SyncPath{}, paths...)
-
-    for len(queue) > 0 {
-        results := make([]SyncResult, len(queue))
-        for i, hash := range queue {
-            data, err := srcDb.Node(hash)
-            if err != nil {
-                t.Fatalf("failed to retrieve node data for %x: %v", hash, err)
-            }
-            results[i] = SyncResult{hash, data}
-        }
-        for _, result := range results {
-            if err := sched.Process(result); err != nil {
-                t.Fatalf("failed to process result %v", err)
-            }
-        }
-        batch := diskdb.NewBatch()
-        if err := sched.Commit(batch); err != nil {
-            t.Fatalf("failed to commit data: %v", err)
-        }
-        batch.Write()
-
-        nodes, paths, _ = sched.Missing(1)
-        queue = append(queue[:0], nodes...)
-        reqs = append(reqs, paths...)
-    }
-    // Cross check that the two tries are in sync
-    checkTrieContents(t, triedb, srcTrie.Hash().Bytes(), srcData)
-
-    // Check that the trie nodes have been requested path-ordered
-    for i := 0; i < len(reqs)-1; i++ {
-        if len(reqs[i]) > 1 || len(reqs[i+1]) > 1 {
-            // In the case of the trie tests, there's no storage so the tuples
-            // must always be single items. 2-tuples should be tested in state.
-            t.Errorf("Invalid request tuples: len(%v) or len(%v) > 1", reqs[i], reqs[i+1])
-        }
-        if bytes.Compare(compactToHex(reqs[i][0]), compactToHex(reqs[i+1][0])) > 0 {
-            t.Errorf("Invalid request order: %v before %v", compactToHex(reqs[i][0]), compactToHex(reqs[i+1][0]))
-        }
-    }
-}
-}
diff --git a/trie/trie.go b/trie/trie.go

```

```
index ac682725..200d79e5 100644
--- a/trie/trie.go
+++ b/trie/trie.go
@@ -33,11 +33,11 @@ import (
    "fmt"
    "sync"

-   "github.com/ava-labs/coreth/core/types"
+   "github.com/ethereum/go-ethereum/common"
+   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/ethereum/go-ethereum/log"
+   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/core/types"
)

var (
diff --git a/trie/trie_test.go b/trie/trie_test.go
index 8f2447ff..40f64b97 100644
--- a/trie/trie_test.go
+++ b/trie/trie_test.go
@@ -38,13 +38,13 @@ import (
    "testing"
    "testing/quick"

-   "github.com/ava-labs/coreth/core/rawdb"
-   "github.com/ava-labs/coreth/ethdb"
-   "github.com/ava-labs/coreth/ethdb/memorydb"
+   "github.com/davecgh/go-spew/spew"
+   "github.com/ethereum/go-ethereum/common"
+   "github.com/ethereum/go-ethereum/crypto"
+   "github.com/ethereum/go-ethereum/rlp"
+   "github.com/flare-foundation/coreth/core/rawdb"
+   "github.com/flare-foundation/coreth/ethdb"
+   "github.com/flare-foundation/coreth/ethdb/memorydb"
+   "golang.org/x/crypto/sha3"
)
```